

# **Introduction**

## **COL331**

**Abhilash Jindal**

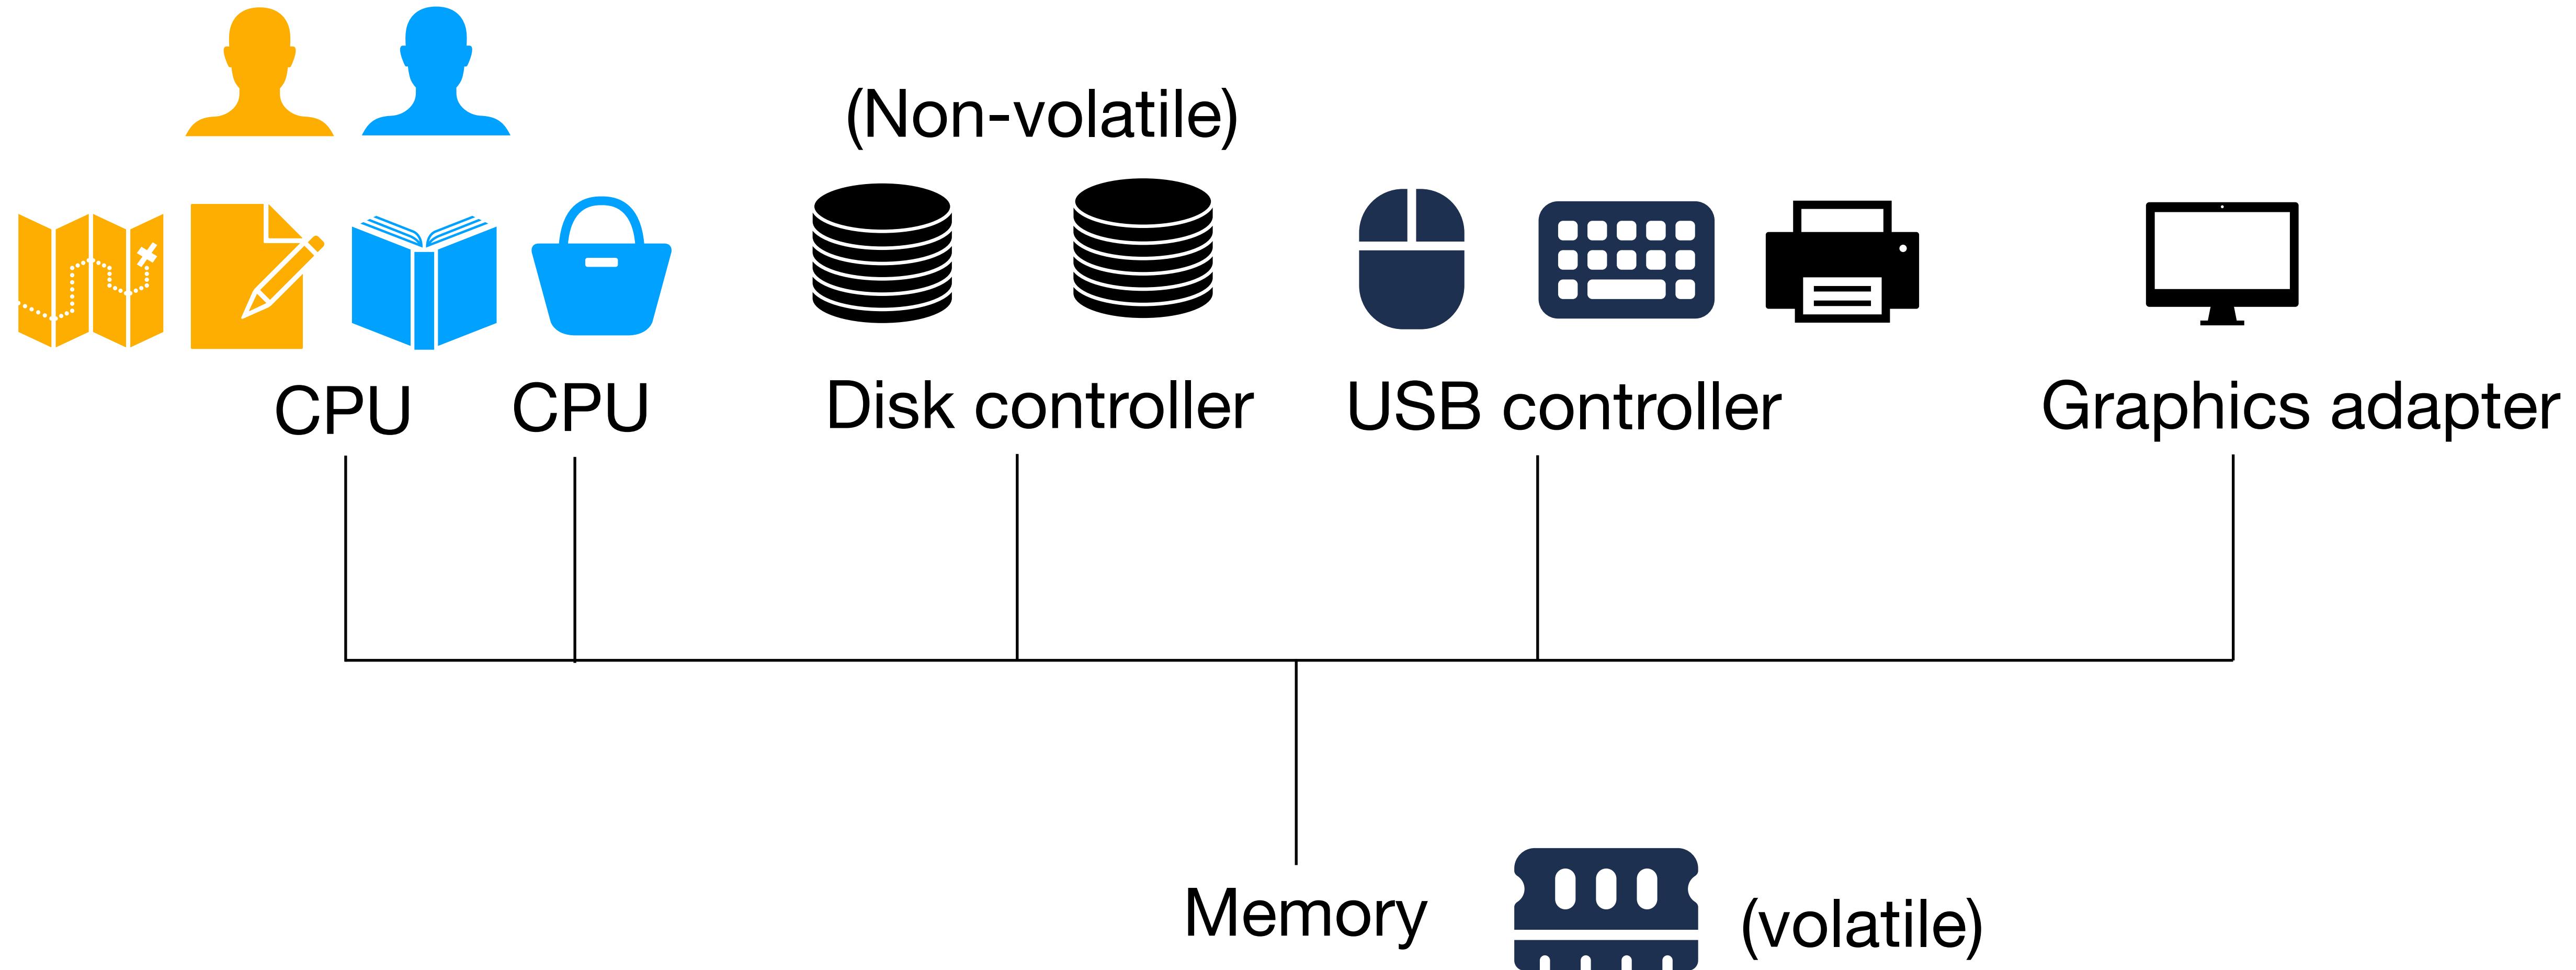
**Reference. OSTEP book: Chapter 2**

# Administrivia

- <http://abhilash-jindal.com/teach.html>
- Grading criteria, TAs, late policy, audit criteria, quizzes, labs, project, piazza link

# **Why does OS matter to a computer?**

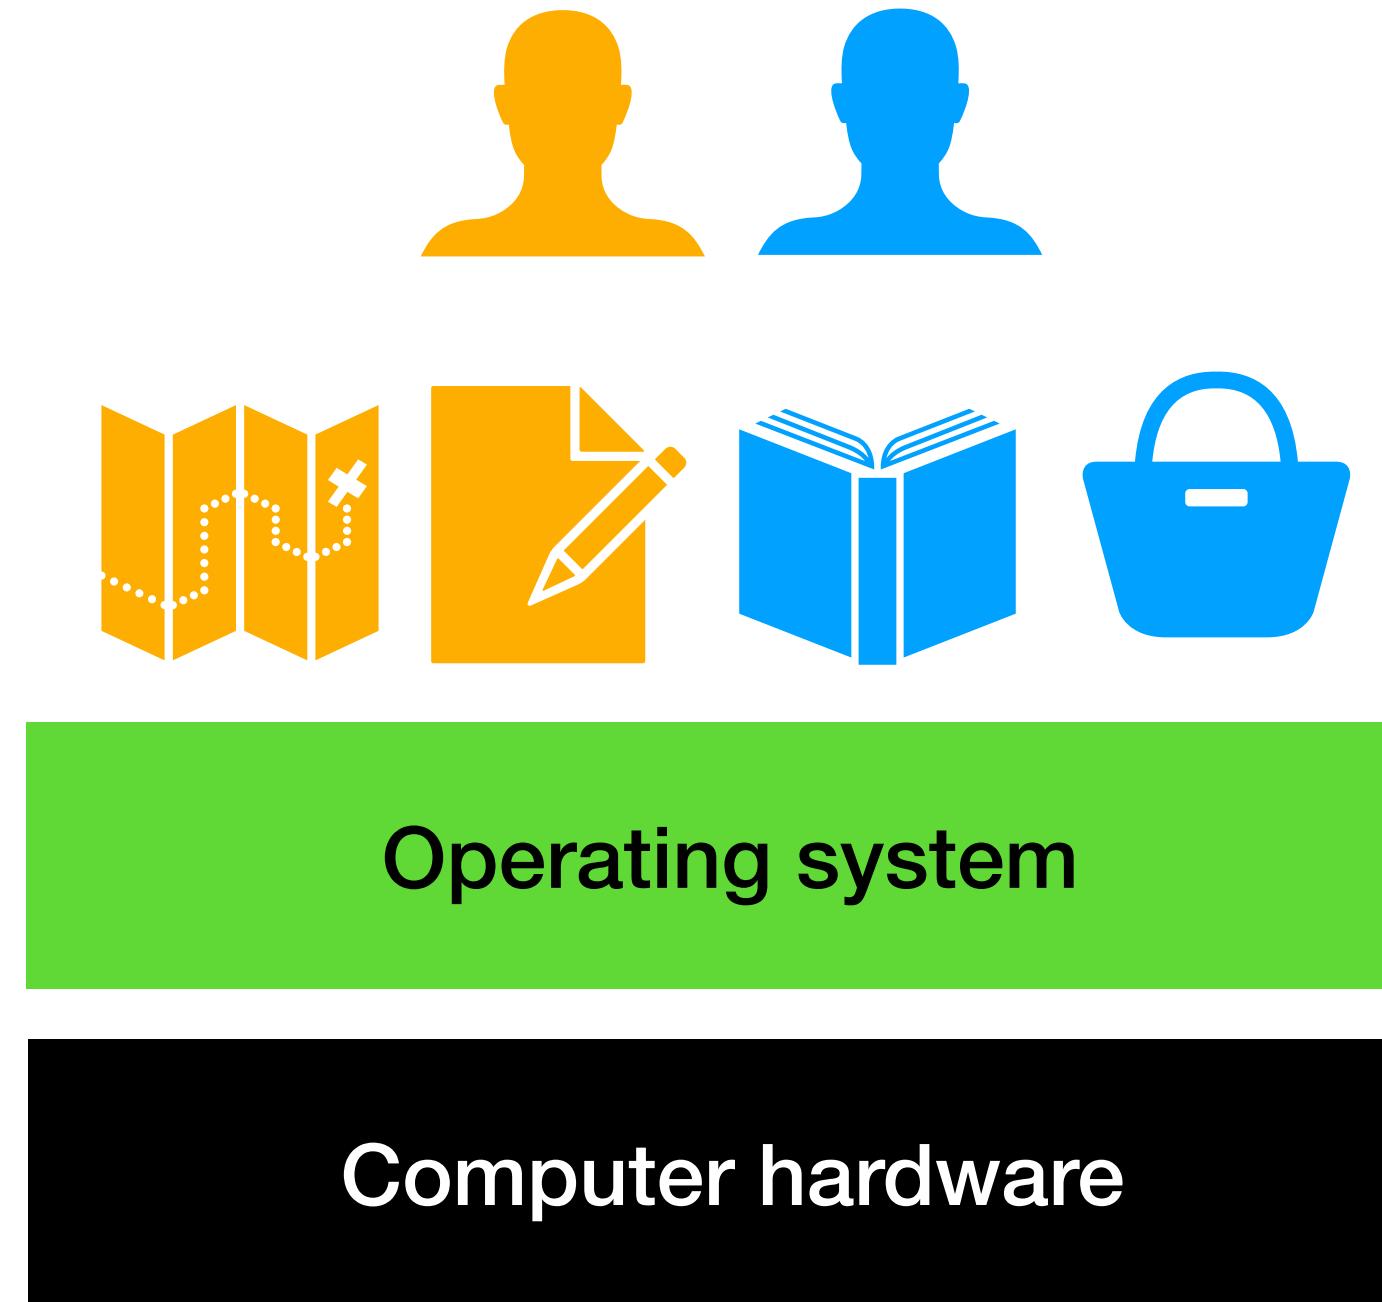
# Computer organization



# Purpose of an OS

- **Resource management**
- Provide higher-level services
- Protection and isolation

# Purpose of OS: Resource management



- Example: `cpu.c`
- Give the illusion of more CPUs than there are
- Multiplex the hardware

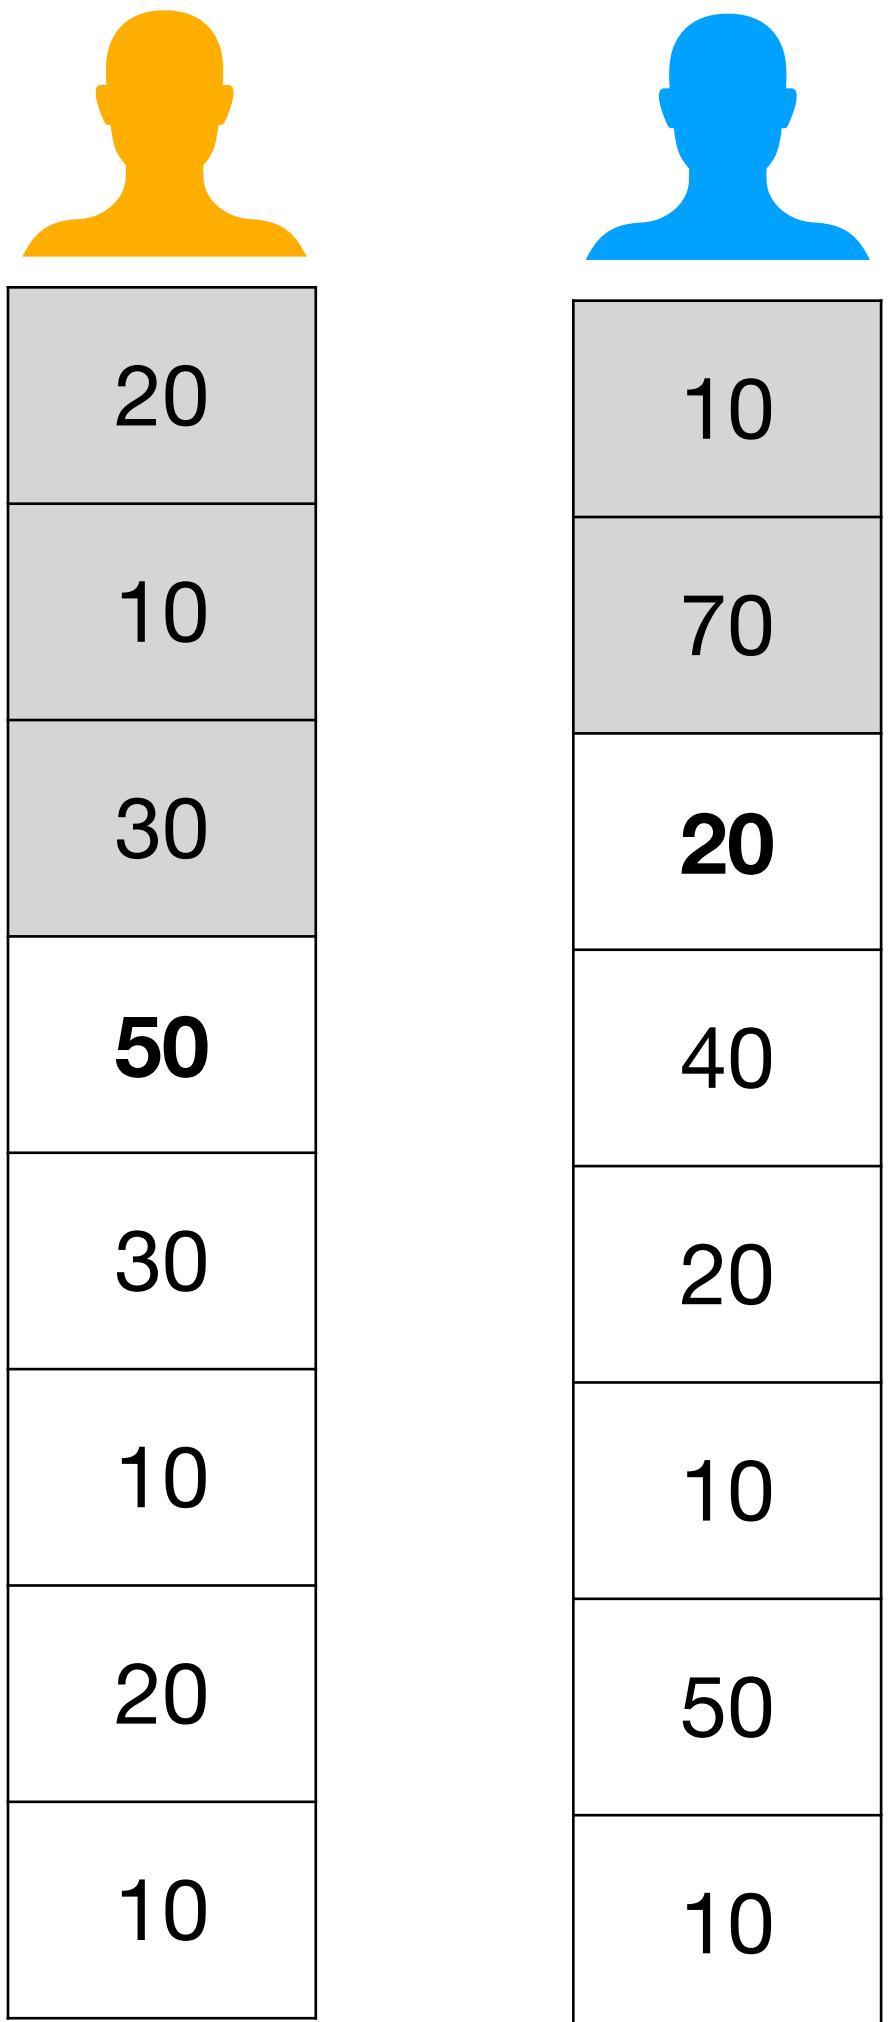
# Calculator analogy: Computing long sum



20
10
30
50
30
10
20
10

- $2 \ 0 =$  (move pointer to 10)
- $+ 1 \ 0 =$  (move pointer to 30)
- $+ 3 \ 0 =$  (move pointer to 50)
- $+ 5 \ 0 =$  (move pointer to 30)
- $+ 3 \ 0 =$  (move pointer to 10)
- $+ 1 \ 0 =$  (move pointer to 20)
- $+ 2 \ 0 =$  (move pointer to 10)

# Sharing the calculator



# Sharing the calculator

20	10
10	70
30	20
50	40
30	20
10	10
20	50
10	10

- Steps to share the calculator:
  - $20 + 10 = 30 + 30 = 60$
  - Write 60 in notebook, remember that we were done till 30, give calculator
  - $10 + 70 = 80$
  - Write 80 in notebook, remember that we were done till 70, give the calculator back

# Remember whatever is on the screen and give calculator?



20
10
30
50
30
10
20
10



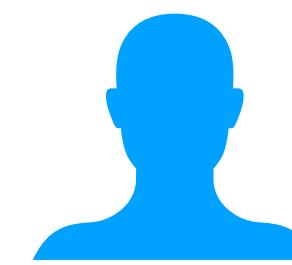
10
70
20
40
20
10
50
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

# Remember whatever is on the screen and give calculator?



20
10
30
50
30
10
20
10



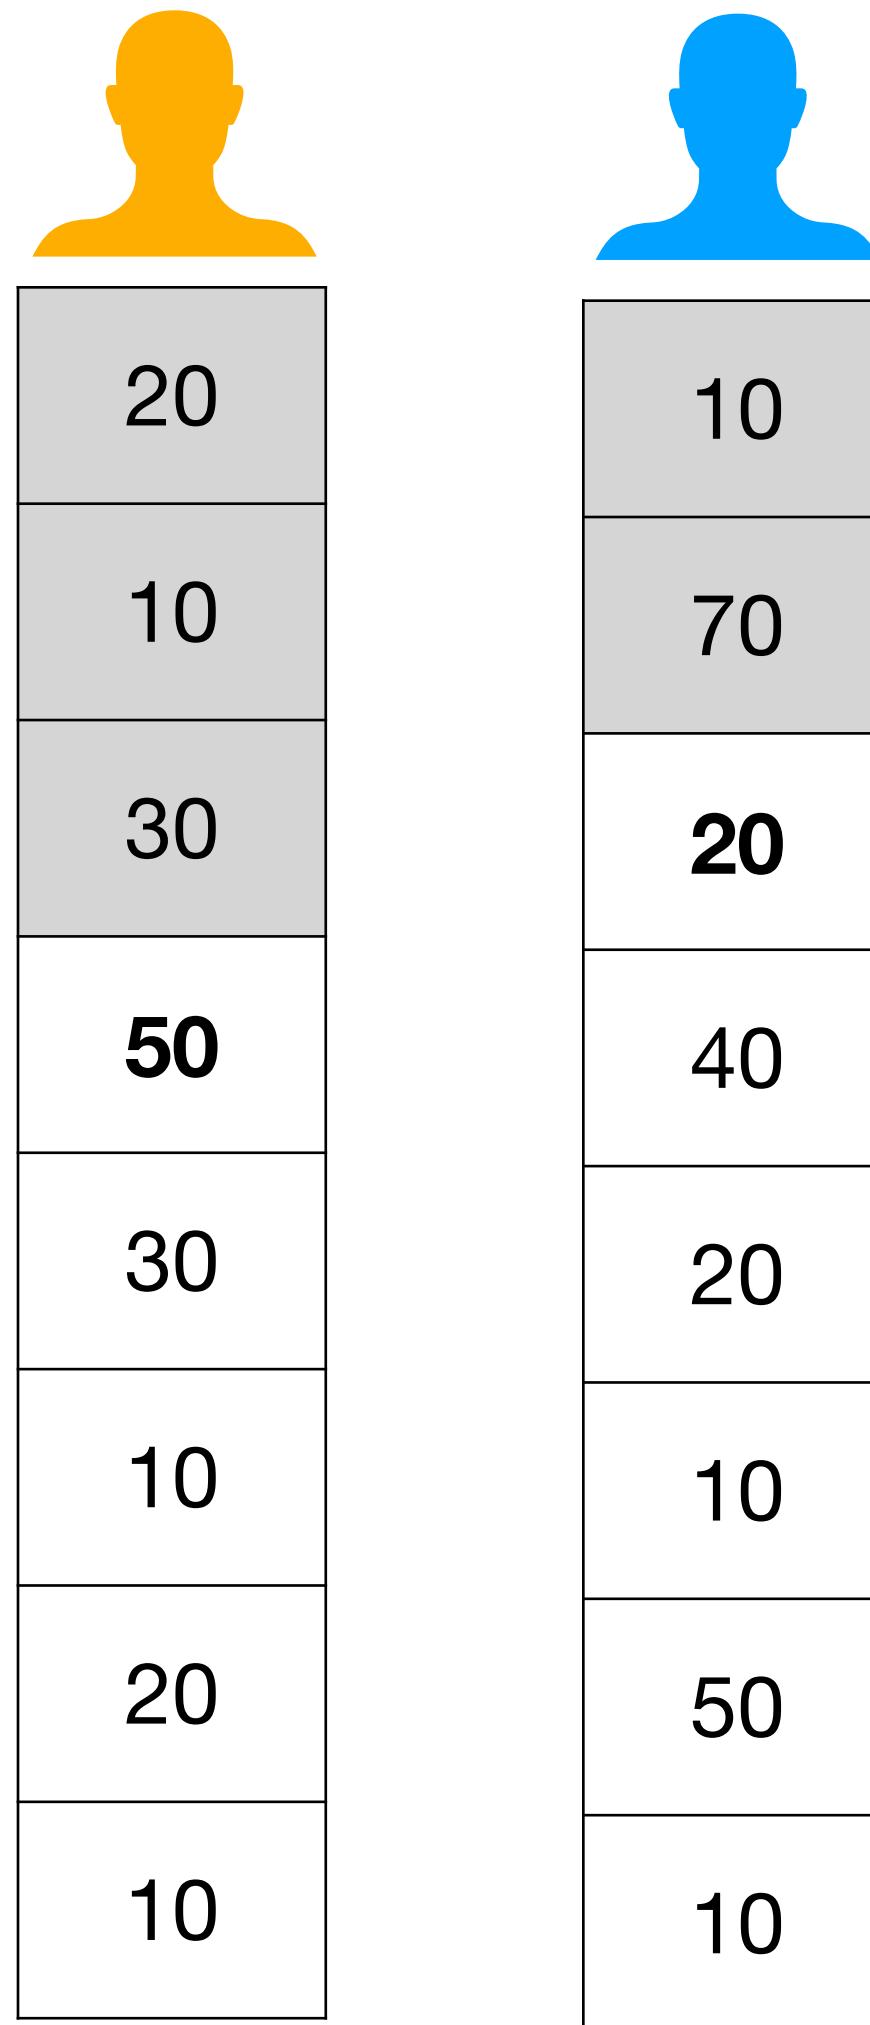
10
70
20
40
20
10
50
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

Can I give calculator here?

“Save 1 and remember pointer to be at 10”

# Remember whatever is on the screen and give calculator?



- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

Can I give calculator here?

“Save 1 and remember pointer to be at 10”

No! Sum would be wrong!

“+ xx = (move pointer)” has to be atomic

# **Resource manager: multiplexing CPU**

# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*

# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)

# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)
- What should happen when multiple processes want to run simultaneously? (policy)

# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)
- What should happen when multiple processes want to run simultaneously? (policy)
  - Fairness: One banker got more calculator time than others

# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)
- What should happen when multiple processes want to run simultaneously? (policy)
  - Fairness: One banker got more calculator time than others
    - Often need to break away from fairness. Game should get more CPU time than Dropbox to provide good user experience

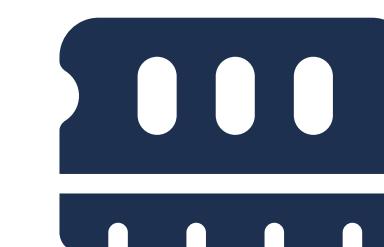
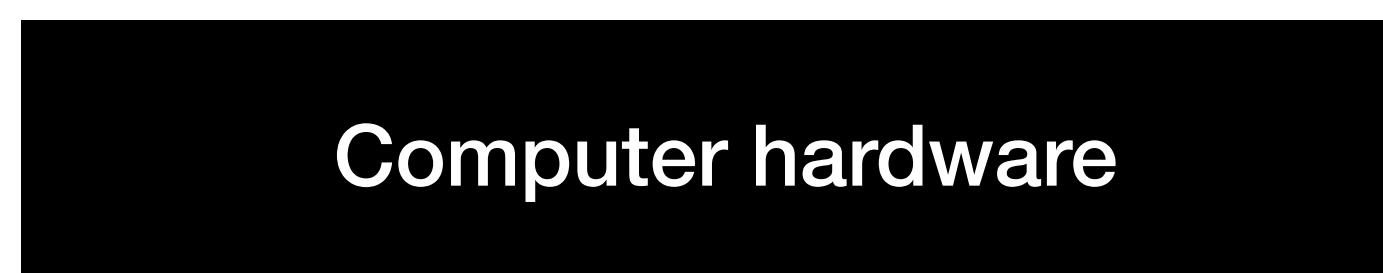
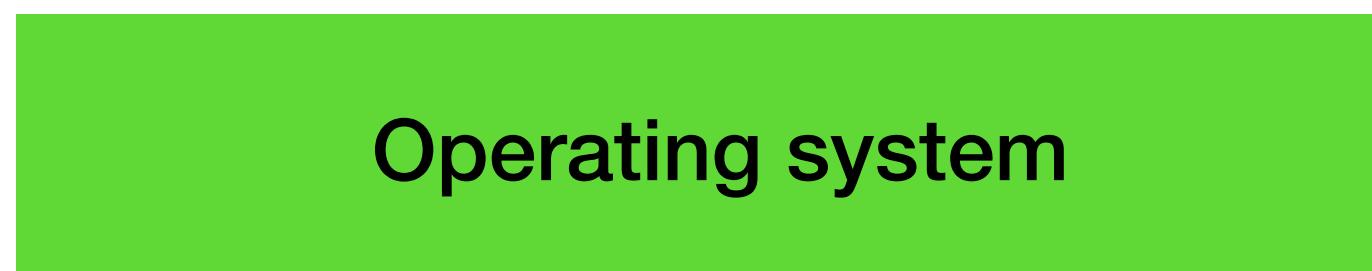
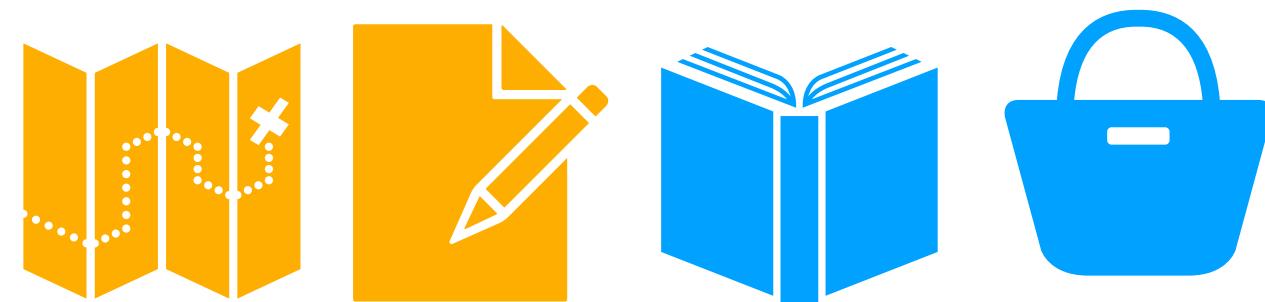
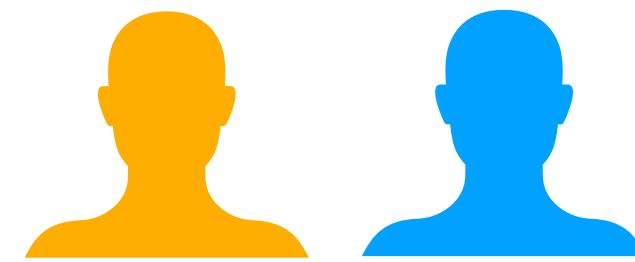
# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)
- What should happen when multiple processes want to run simultaneously? (policy)
  - Fairness: One banker got more calculator time than others
    - Often need to break away from fairness. Game should get more CPU time than Dropbox to provide good user experience
    - Starvation freedom: When there are multiple bankers, one banker never got the calculator

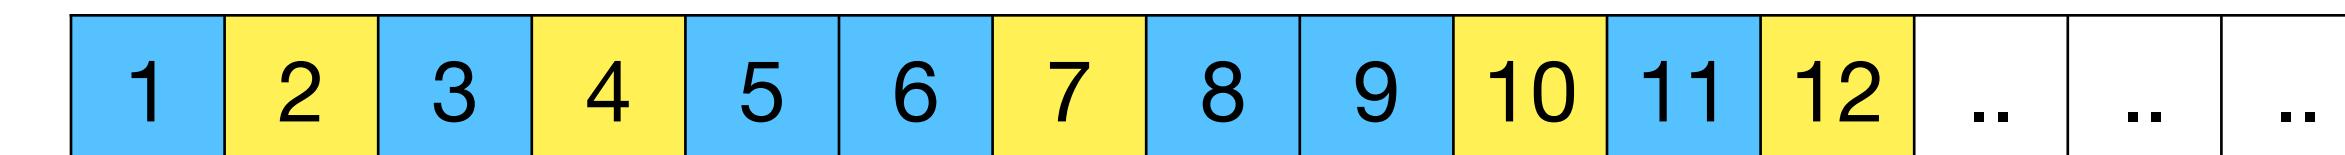
# Purpose of OS: Resource manager

*Different approaches to memory management:*

- Segment different memory portions to different processes
- Multiplex memory pages across processes

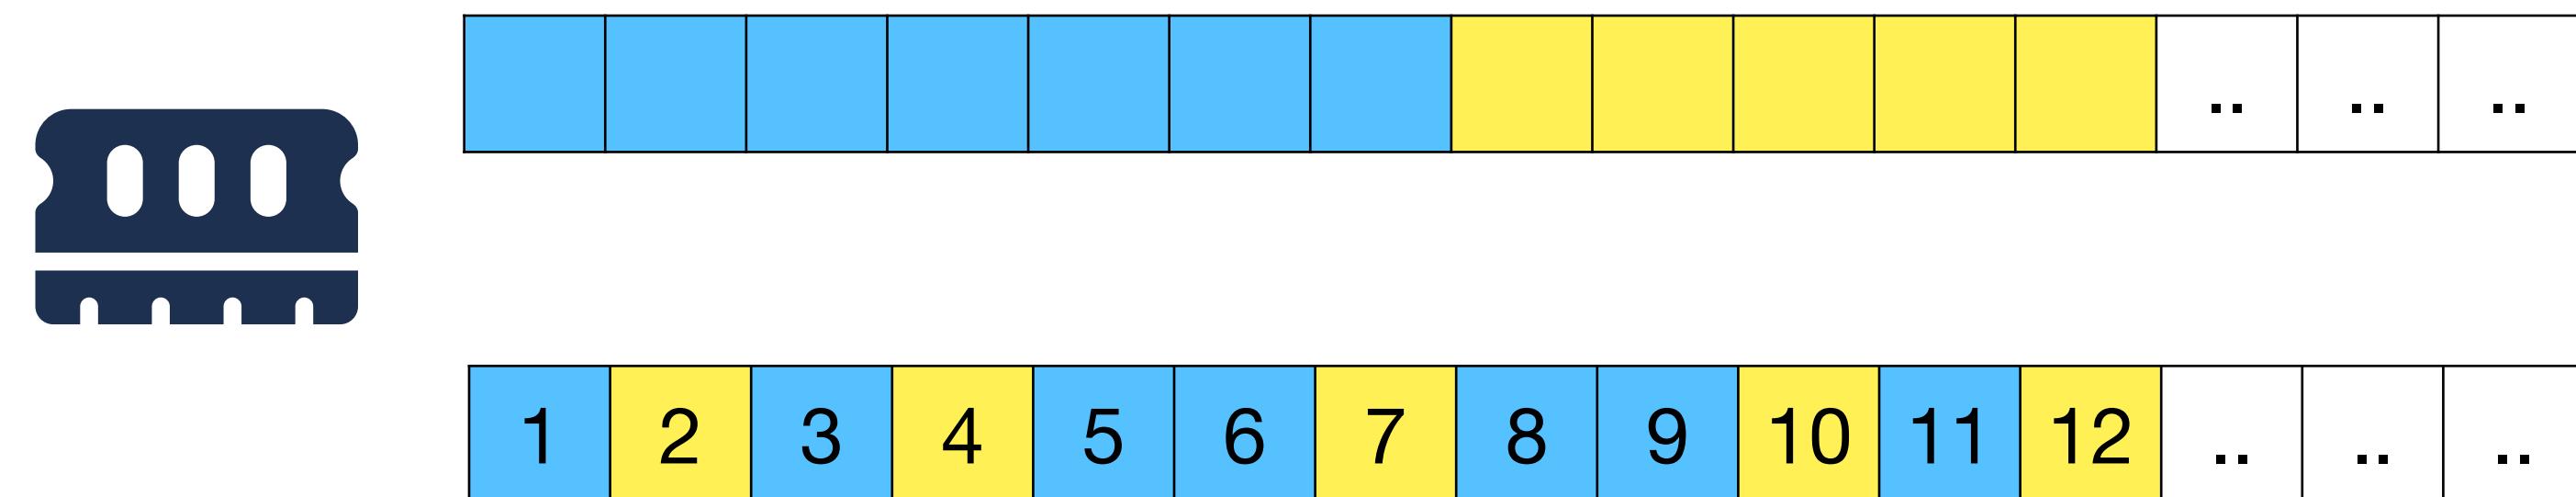


Or



# Memory management

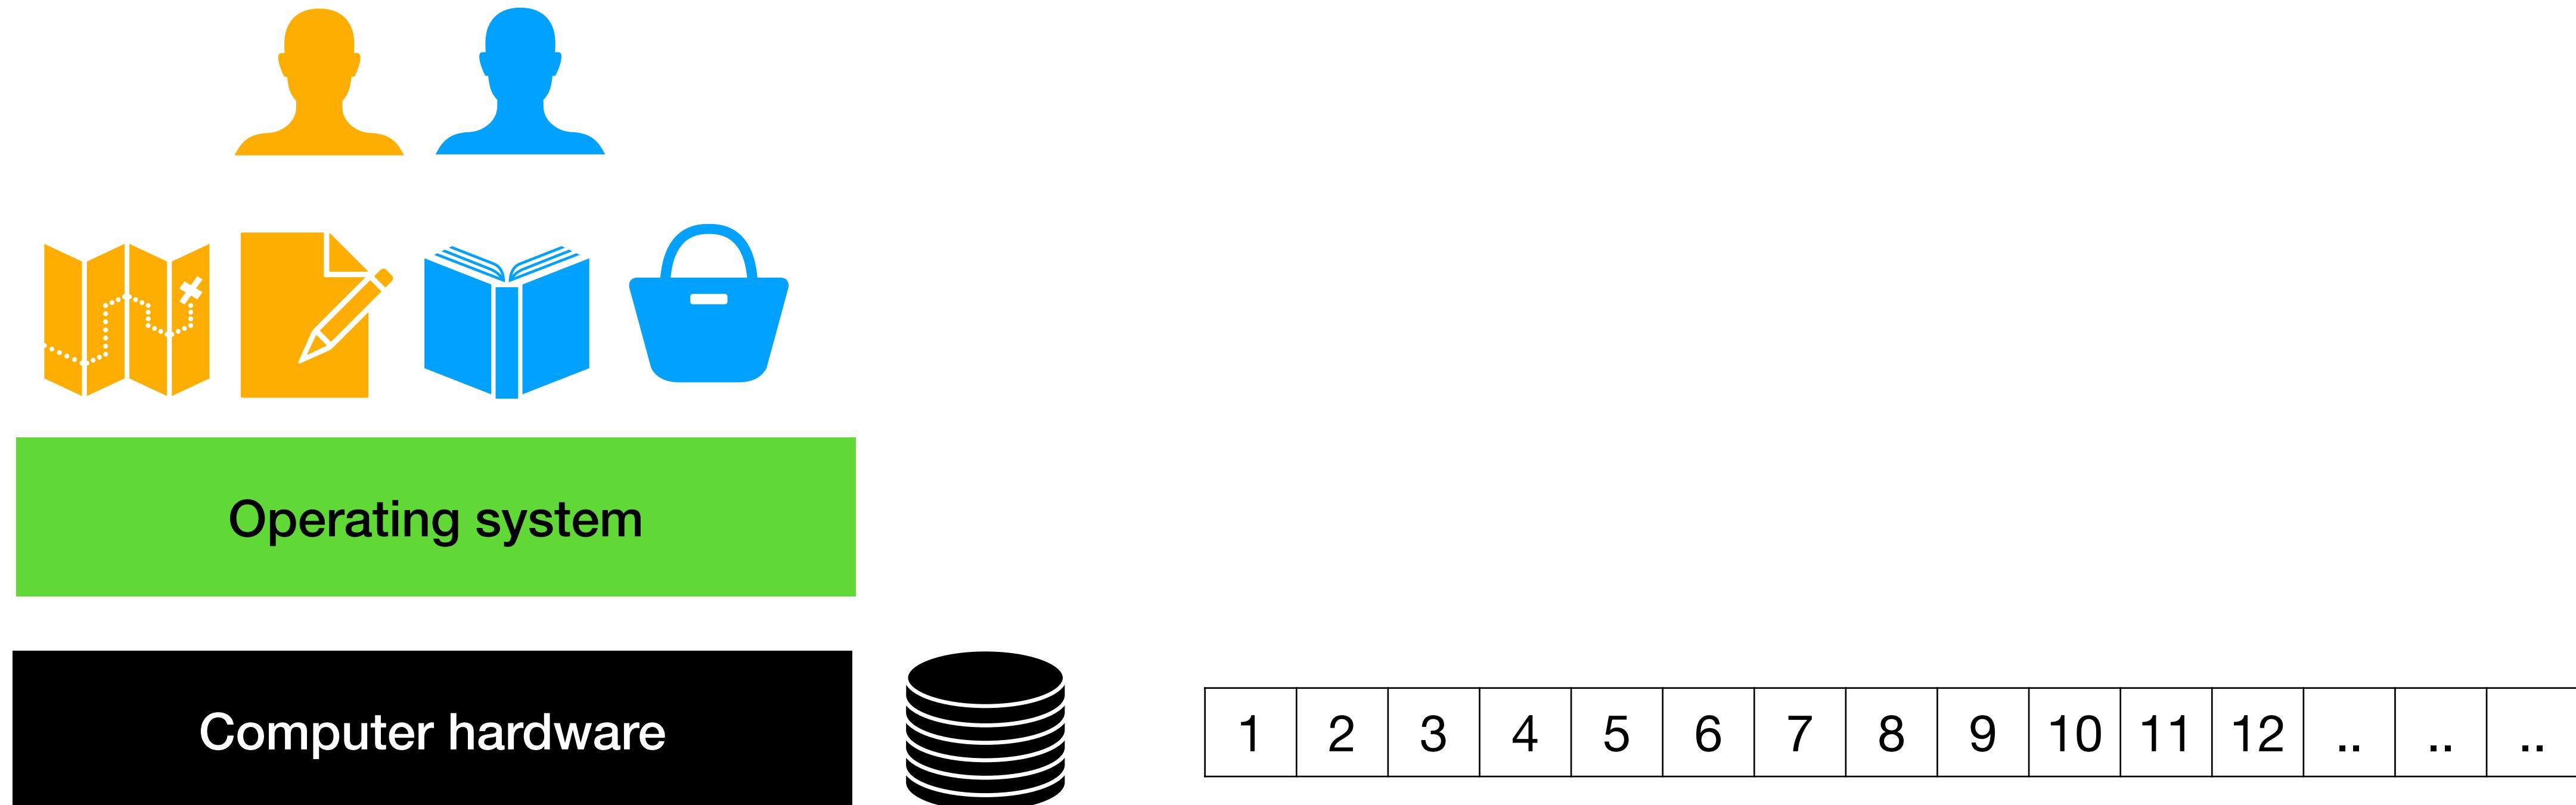
- Segmentation is cheap to implement
- But not flexible. What if a process needs more memory than what OS gave?
- Paging is complicated to implement
- Highly flexible



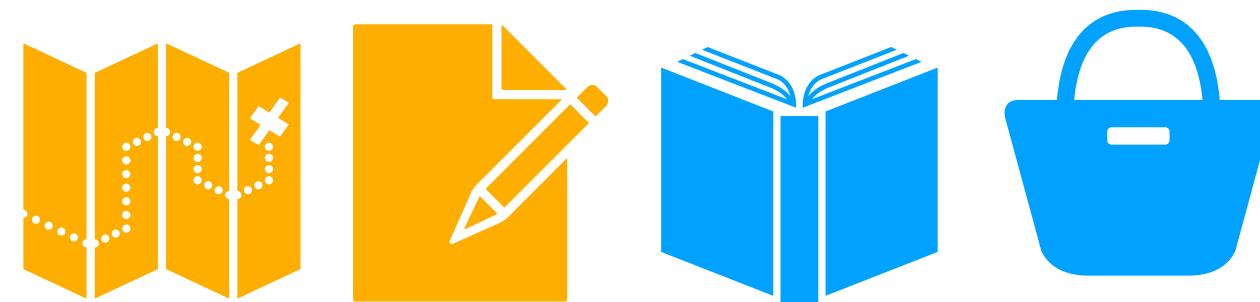
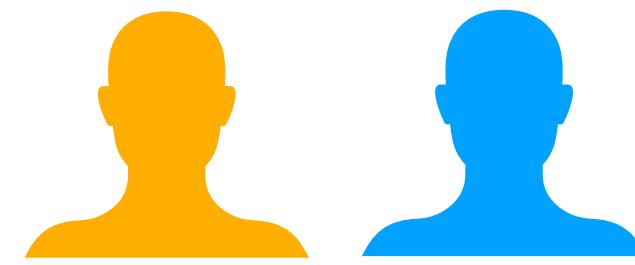
# Purpose of an OS

- Resource management
- **Provide higher-level services**
- Protection and isolation

# Purpose of OS: Provide higher-level services

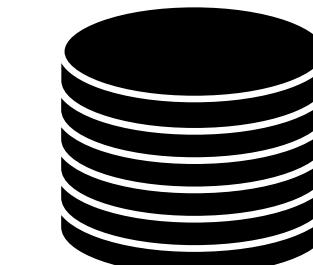


# Purpose of OS: Provide higher-level services



Operating system

Computer hardware

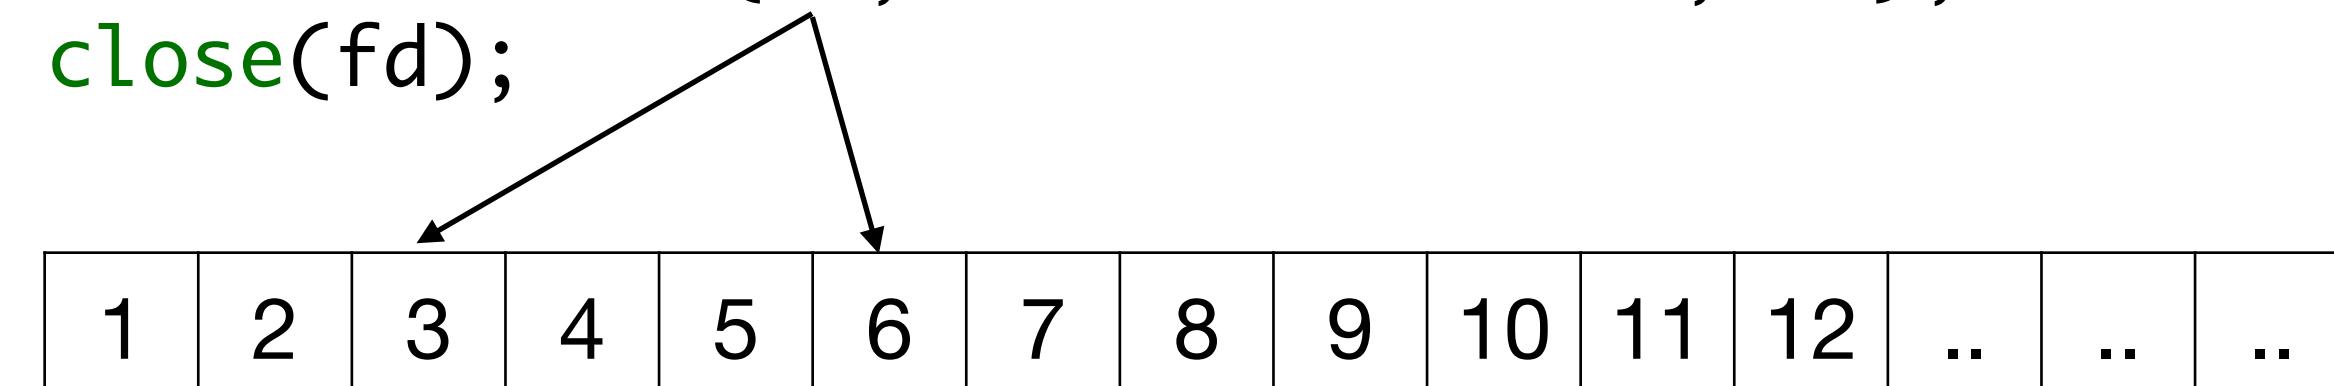


*Example: io.c*

Disk interface: List of blocks

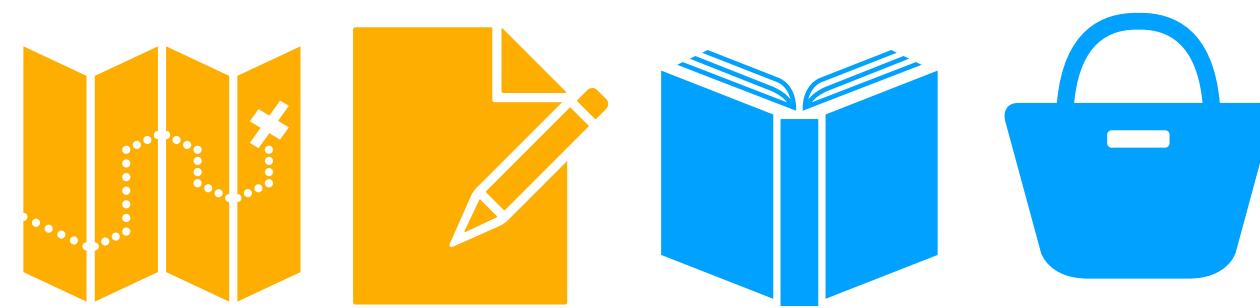
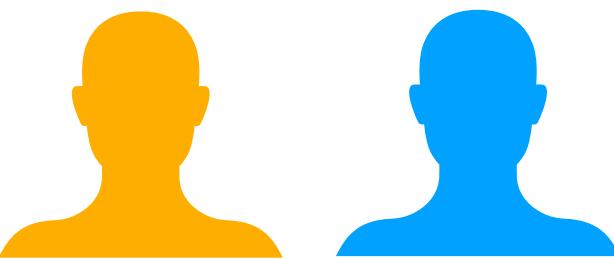
File system OS interface: Folders and files

```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```



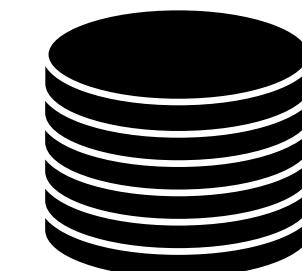
# Why file system?

Why not just multiplex disk blocks like memory?



Operating system

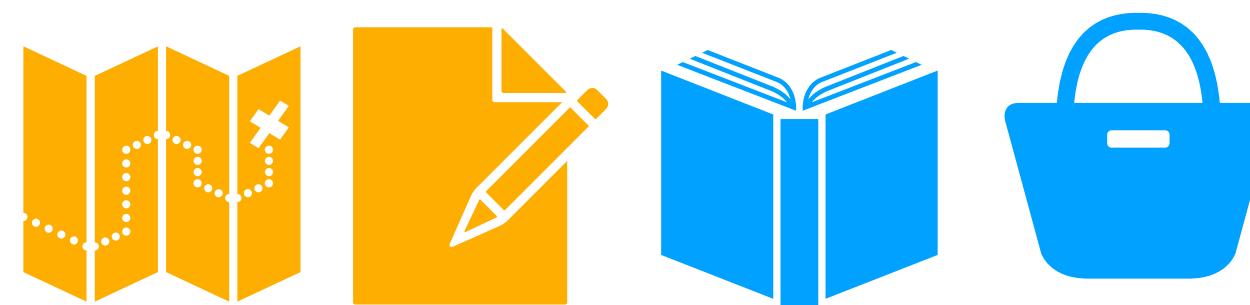
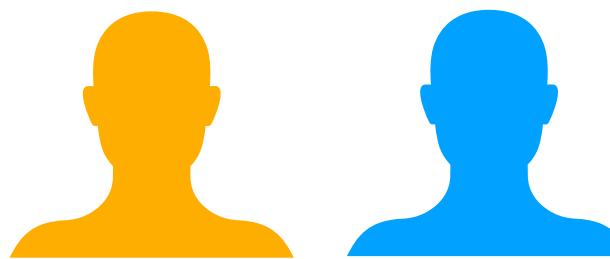
Computer hardware



# Why file system?

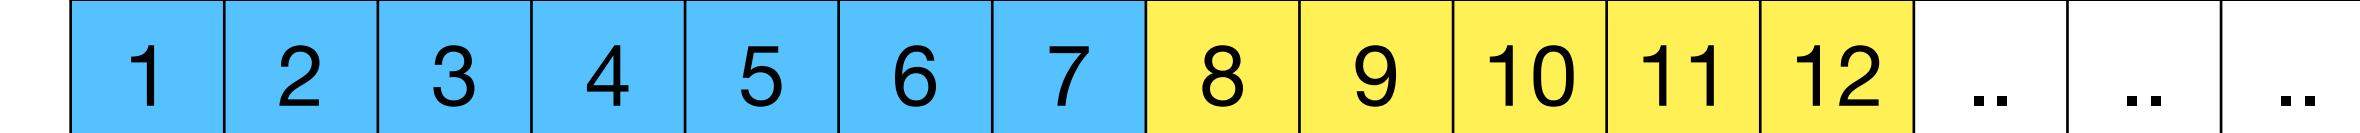
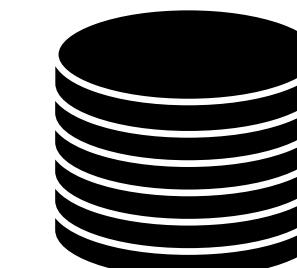
Why not just multiplex disk blocks like memory?

- Disk blocks live after programs exits, computer restarts



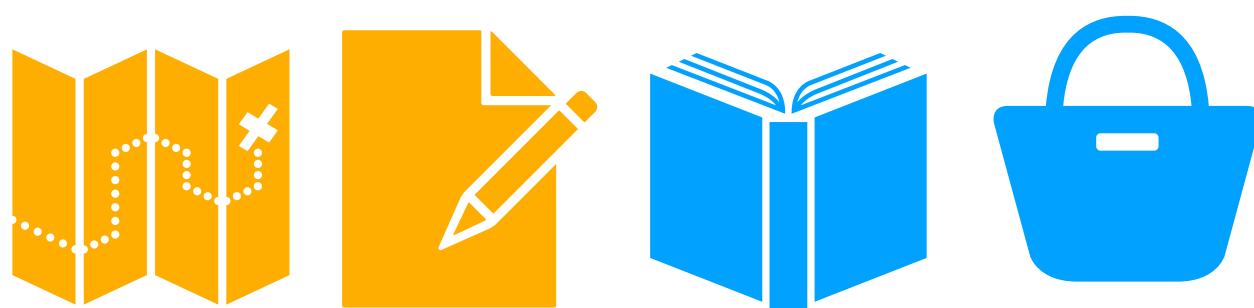
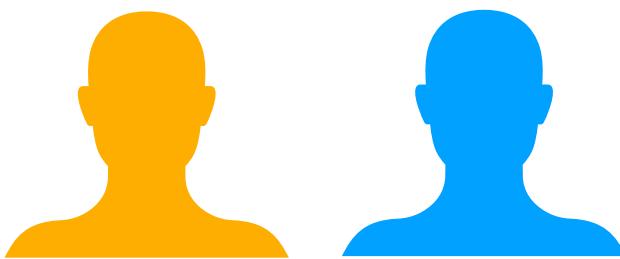
Operating system

Computer hardware



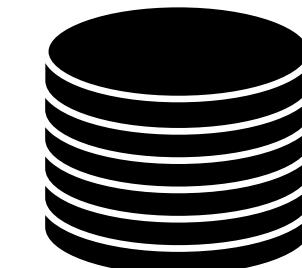
# Why file system?

Why not just multiplex disk blocks like memory?



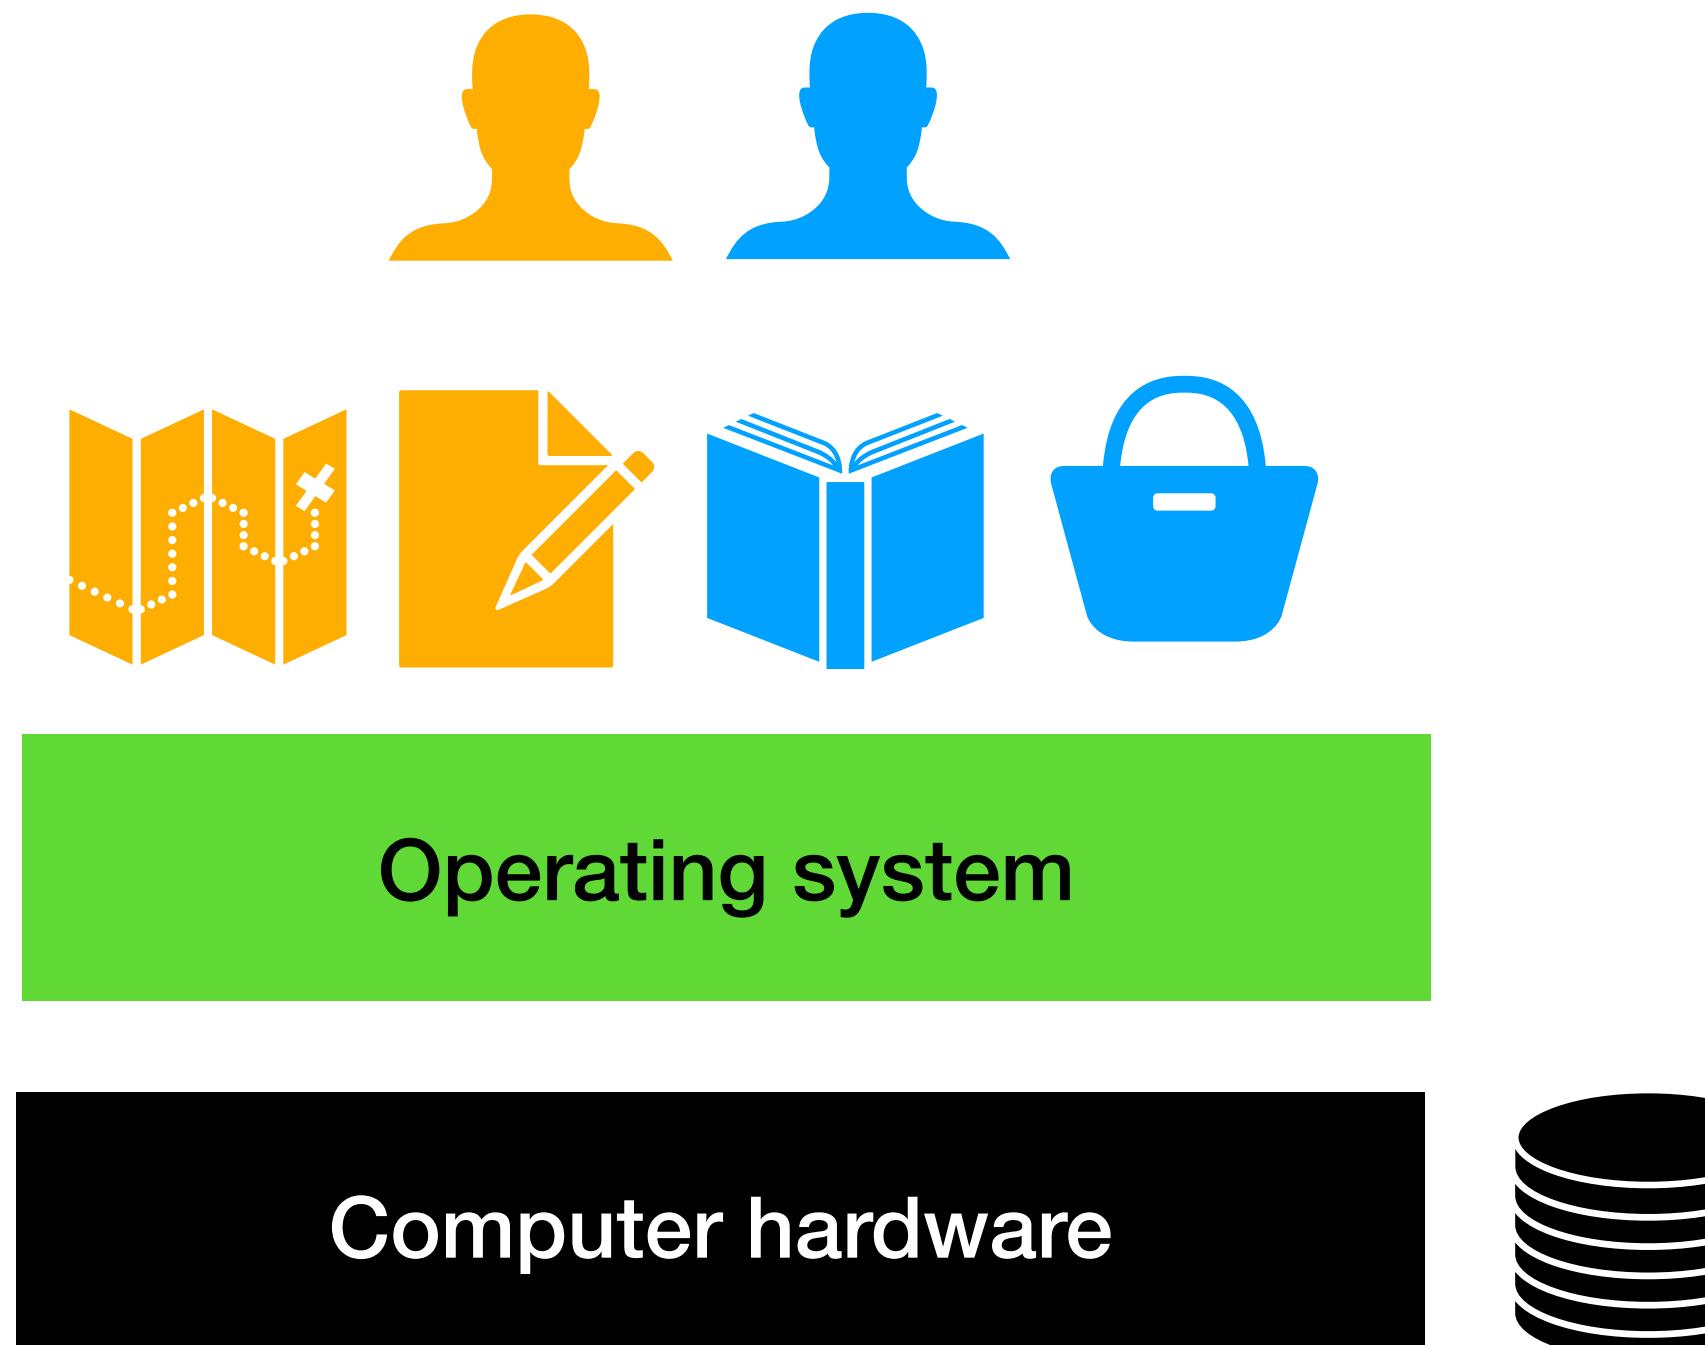
Operating system

Computer hardware

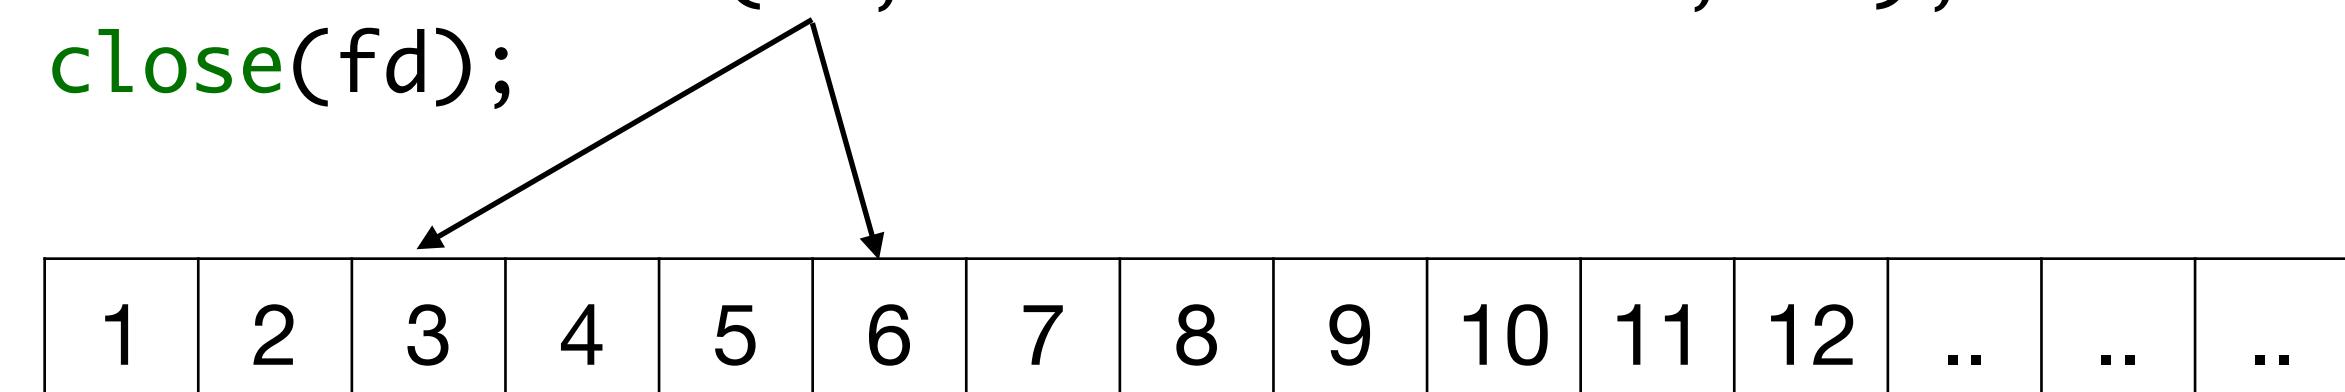


- Disk blocks live after programs exits, computer restarts
- Different programs read / write same file
  - vim writes io.c
  - gcc reads io.c, write io
  - We finally run io

# Higher level services provide portability

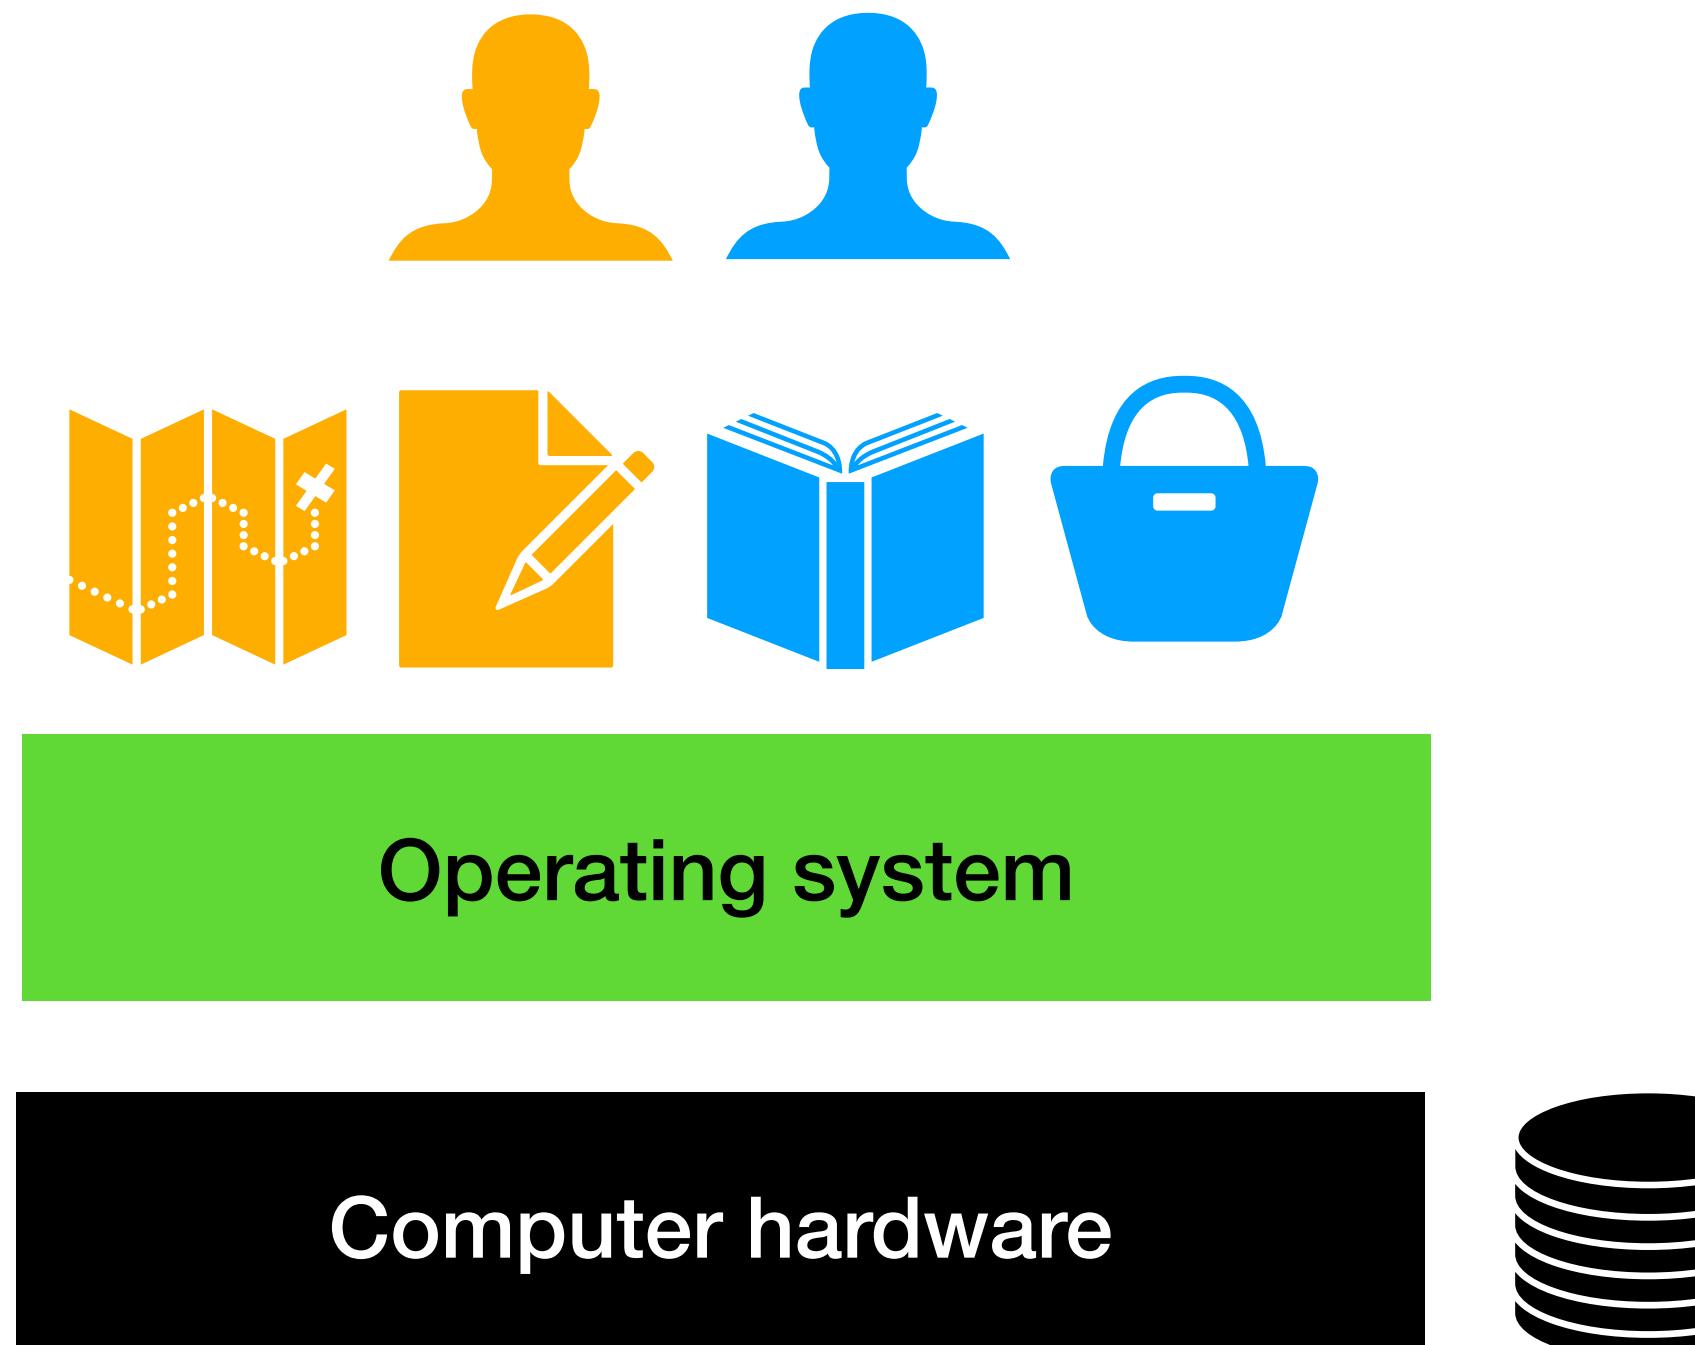


```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```



# Higher level services provide portability

- Abstract away hardware details

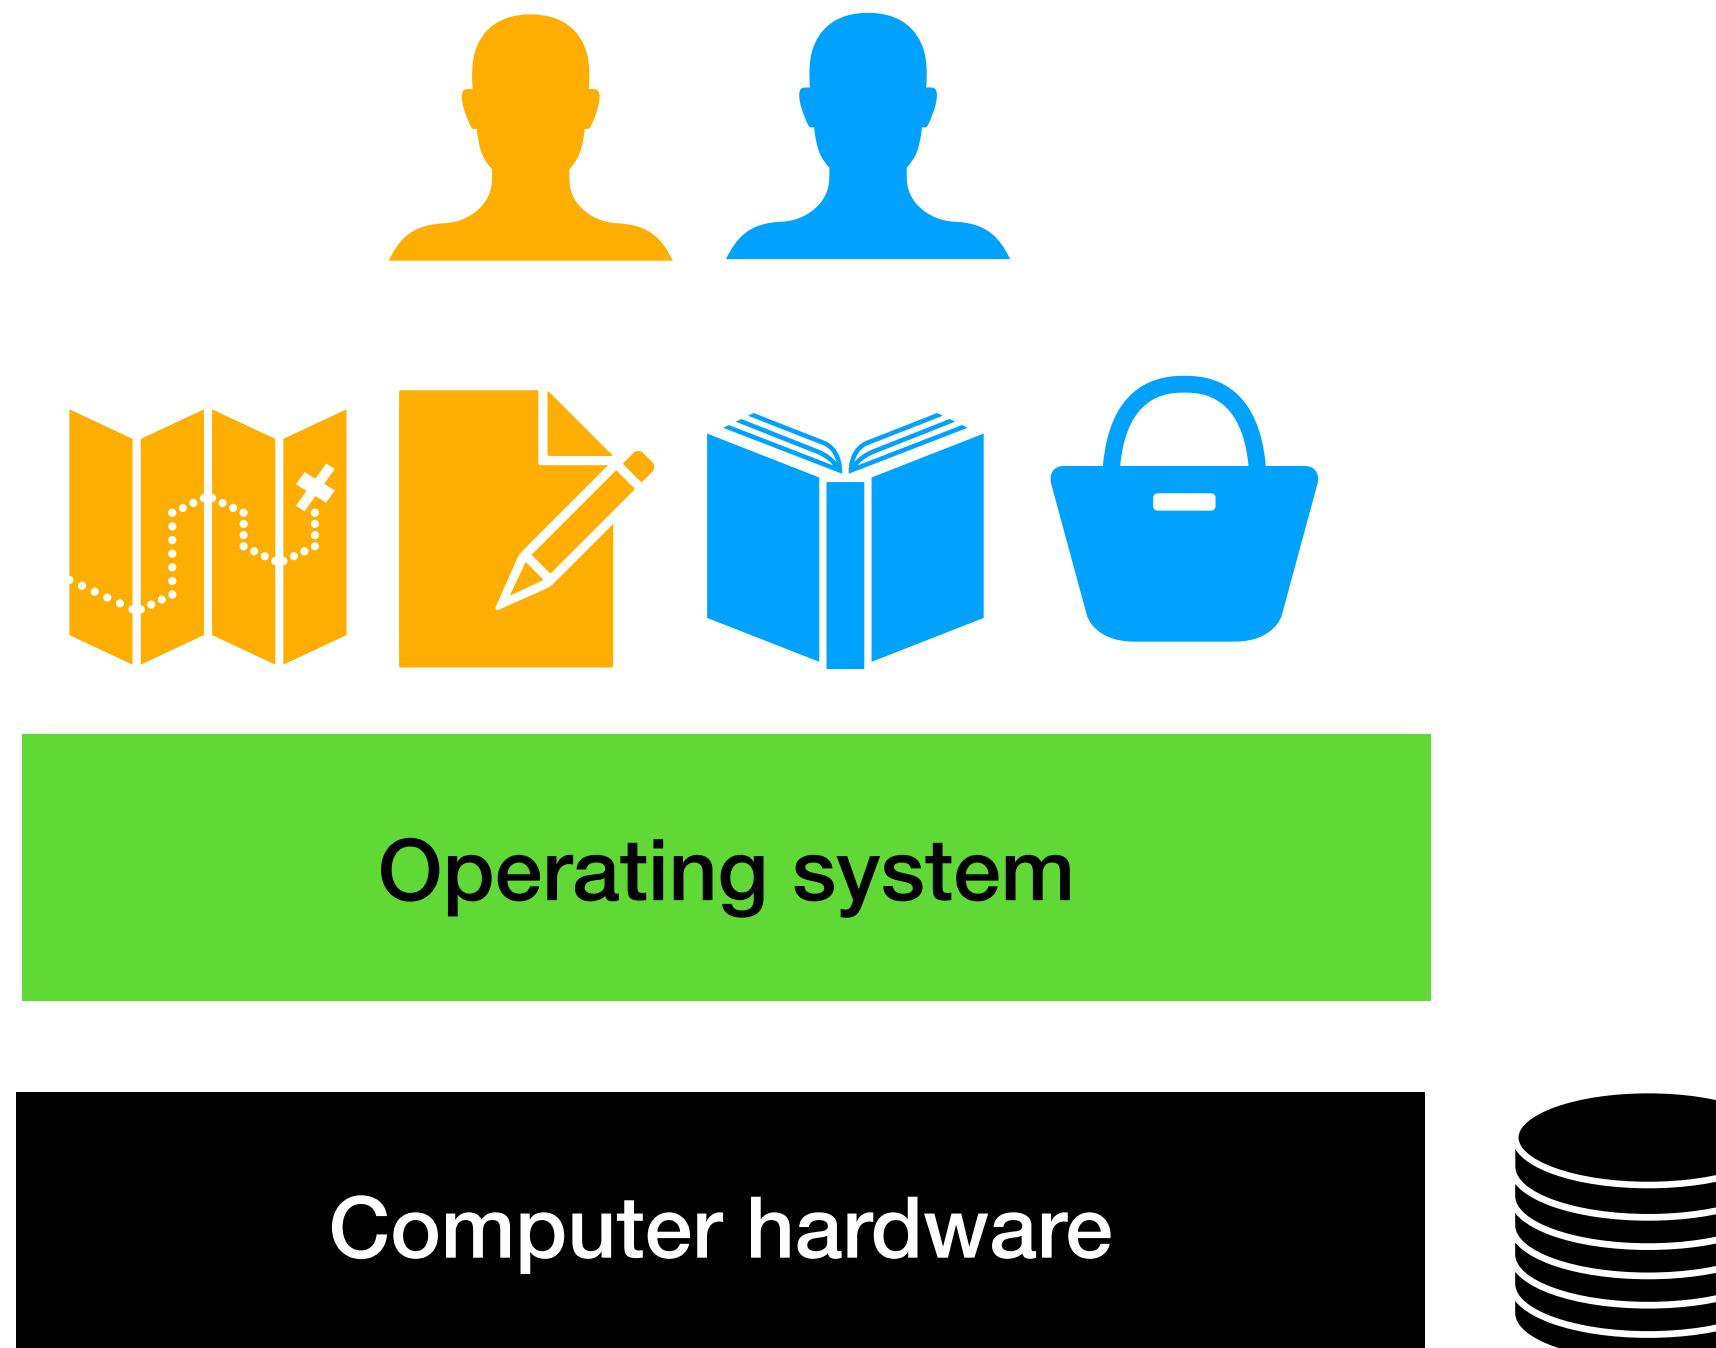


```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```

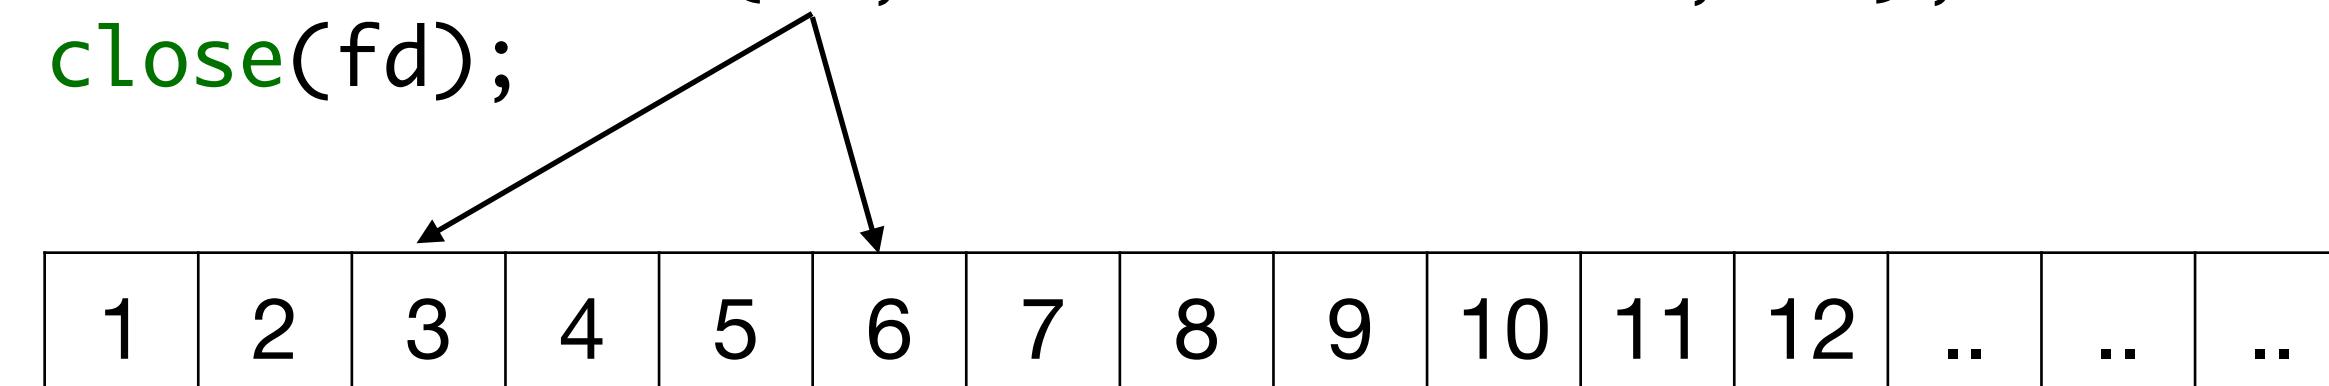


# Higher level services provide portability

- Abstract away hardware details
- Programs need not be rewritten when moving from hard-disk drive to solid state drive



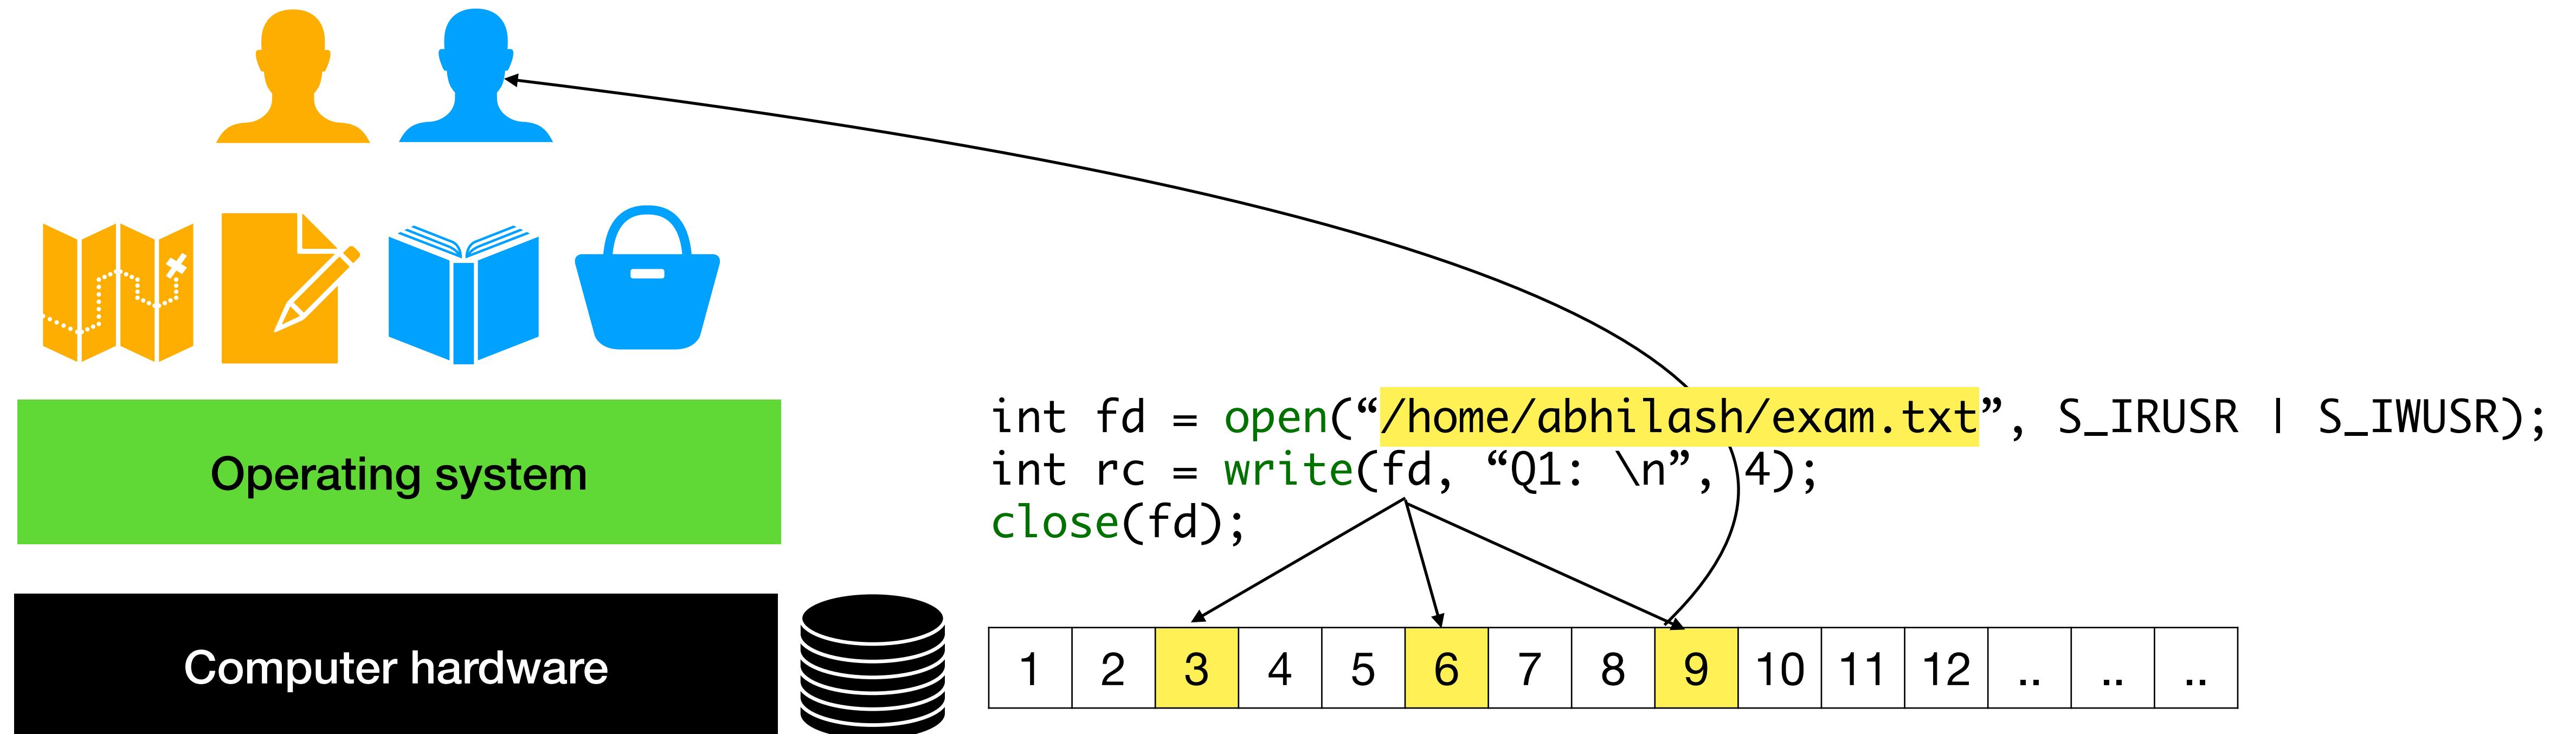
```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```



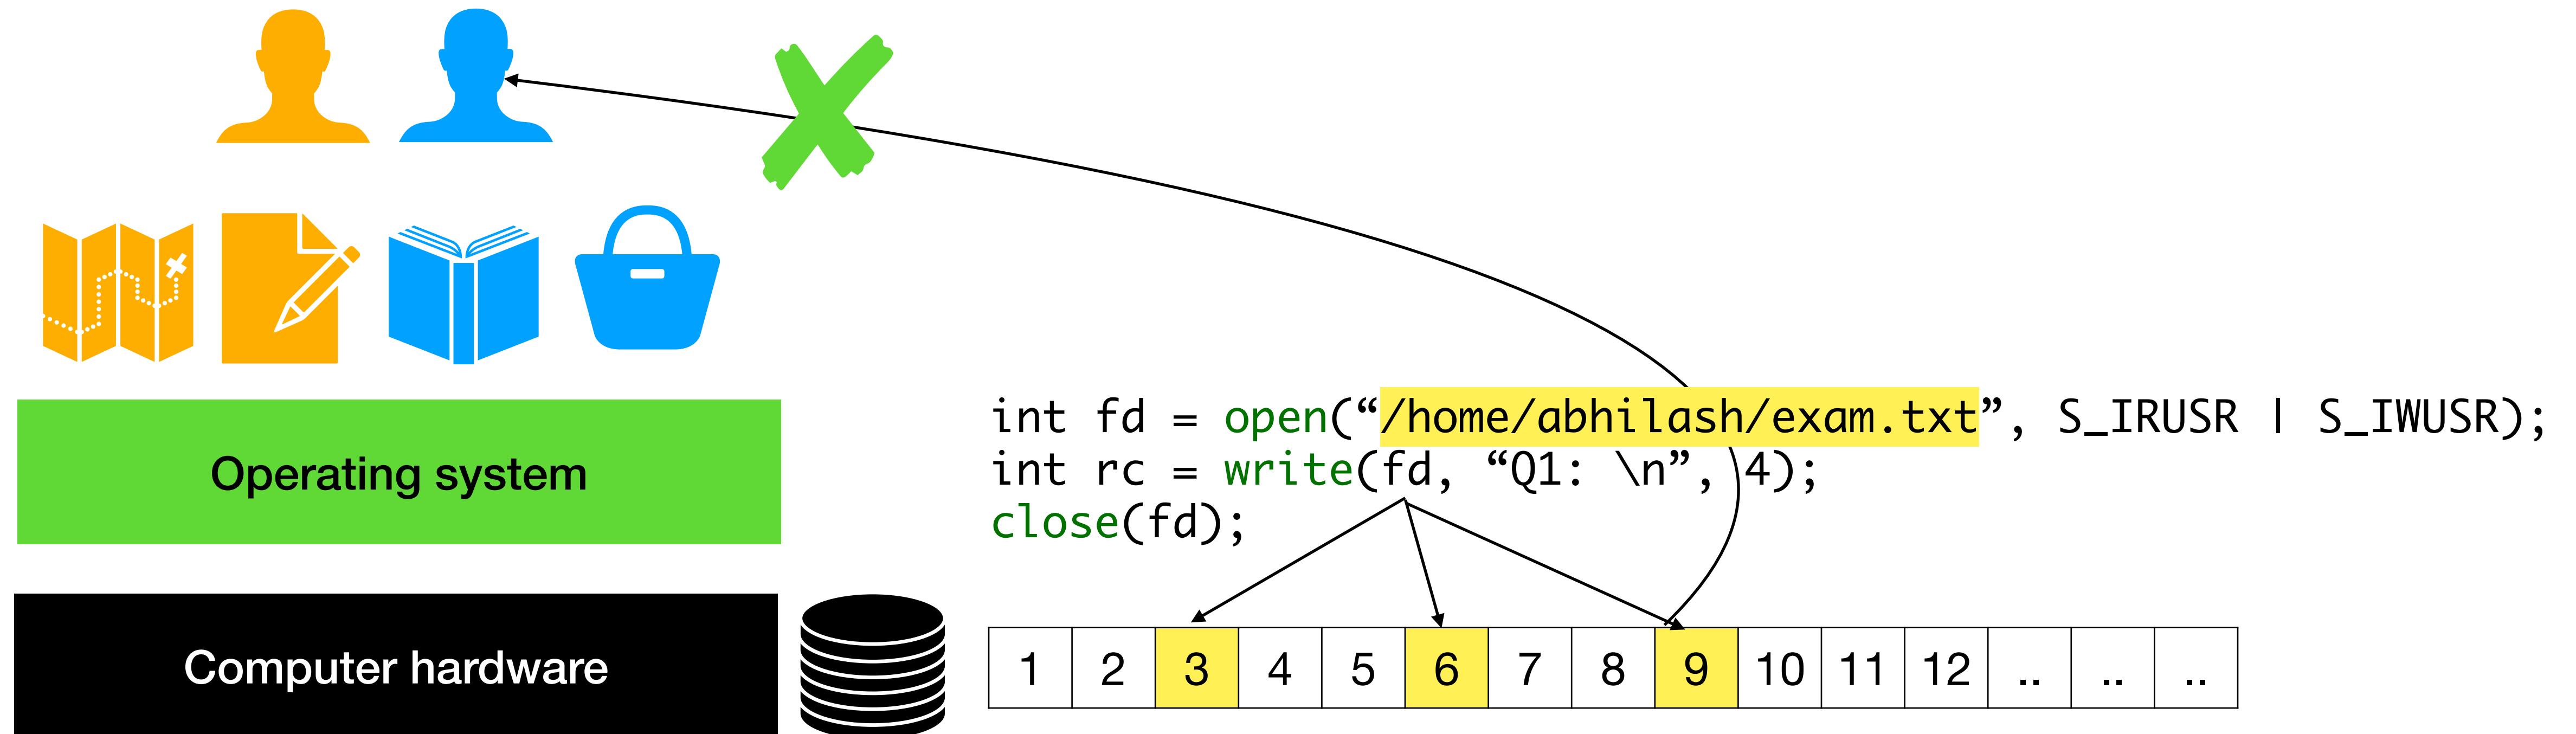
# Purpose of an OS

- Resource management
- Provide higher-level services
- **Protection and isolation**

# Purpose of OS: Protection

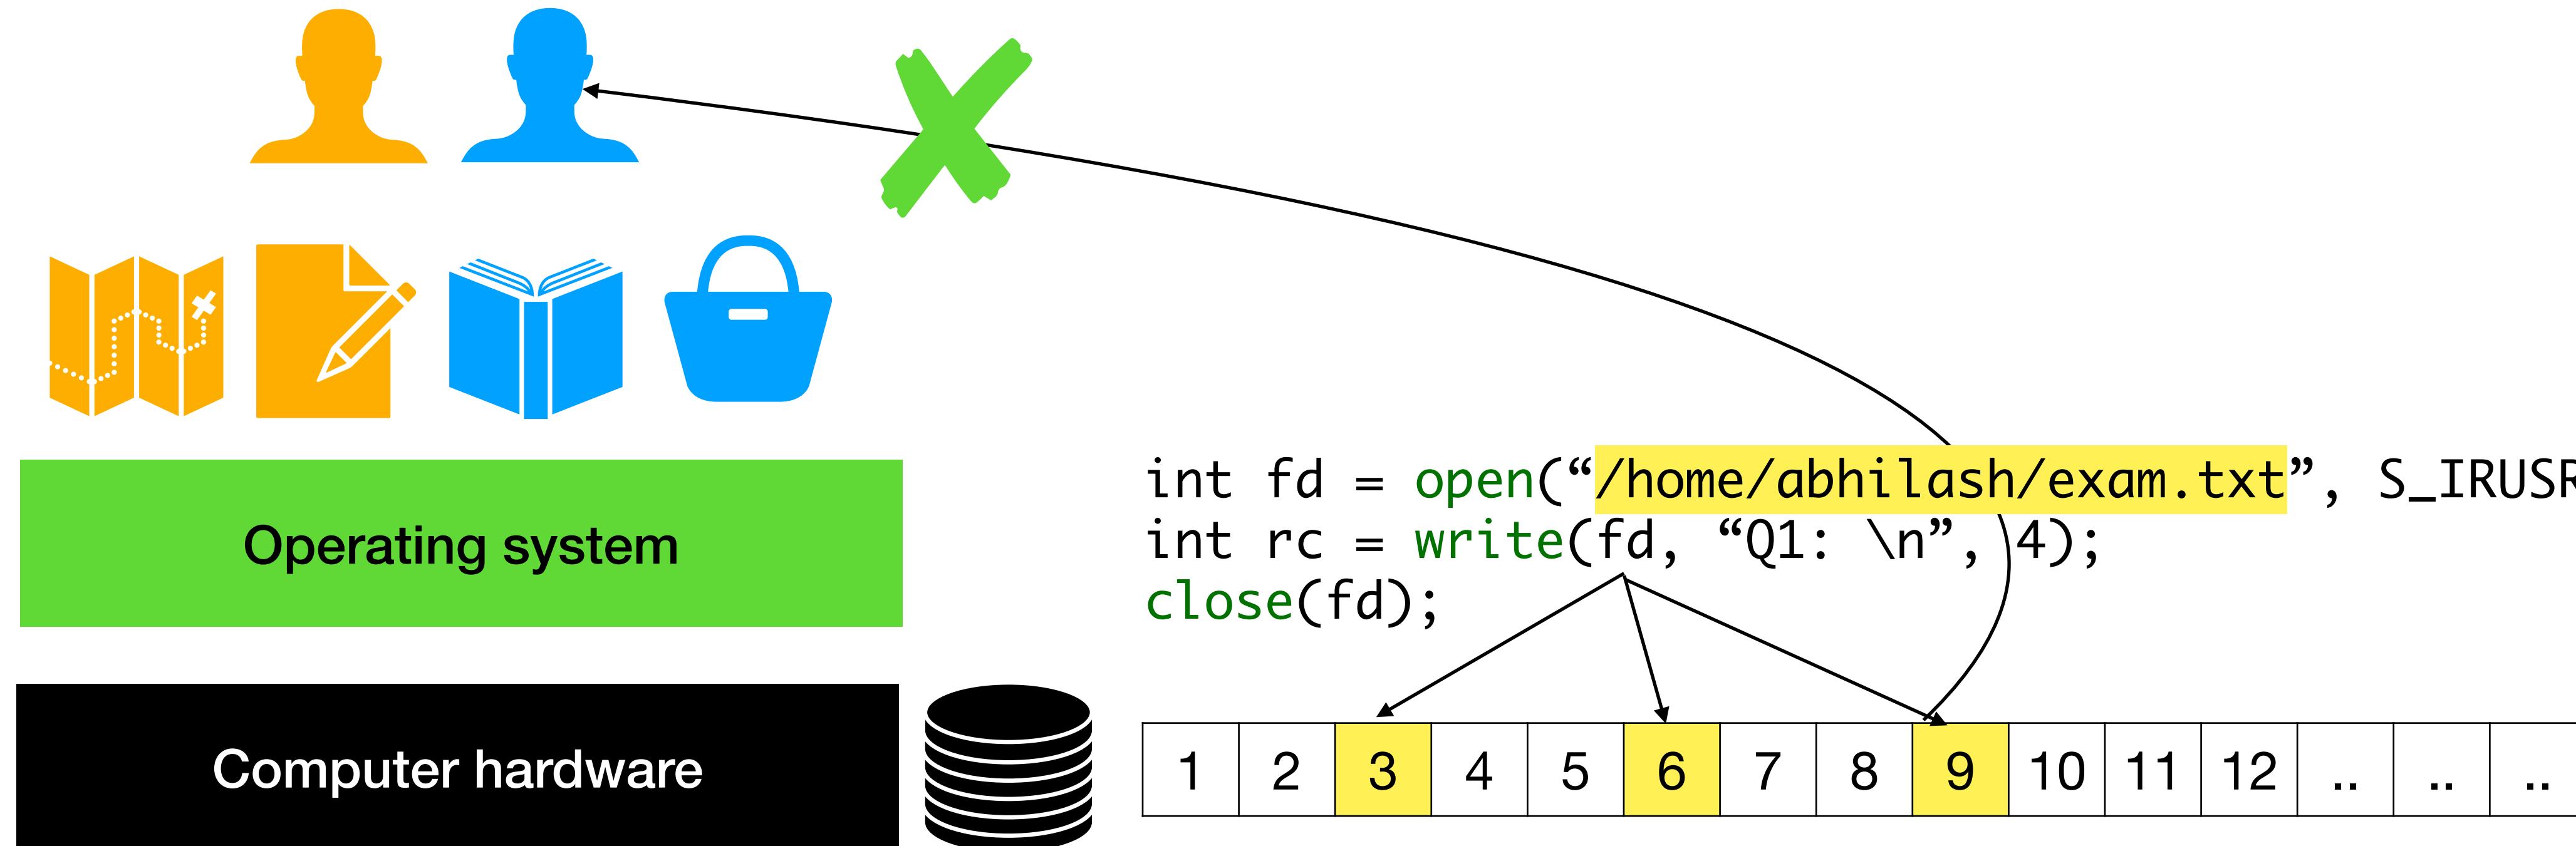


# Purpose of OS: Protection



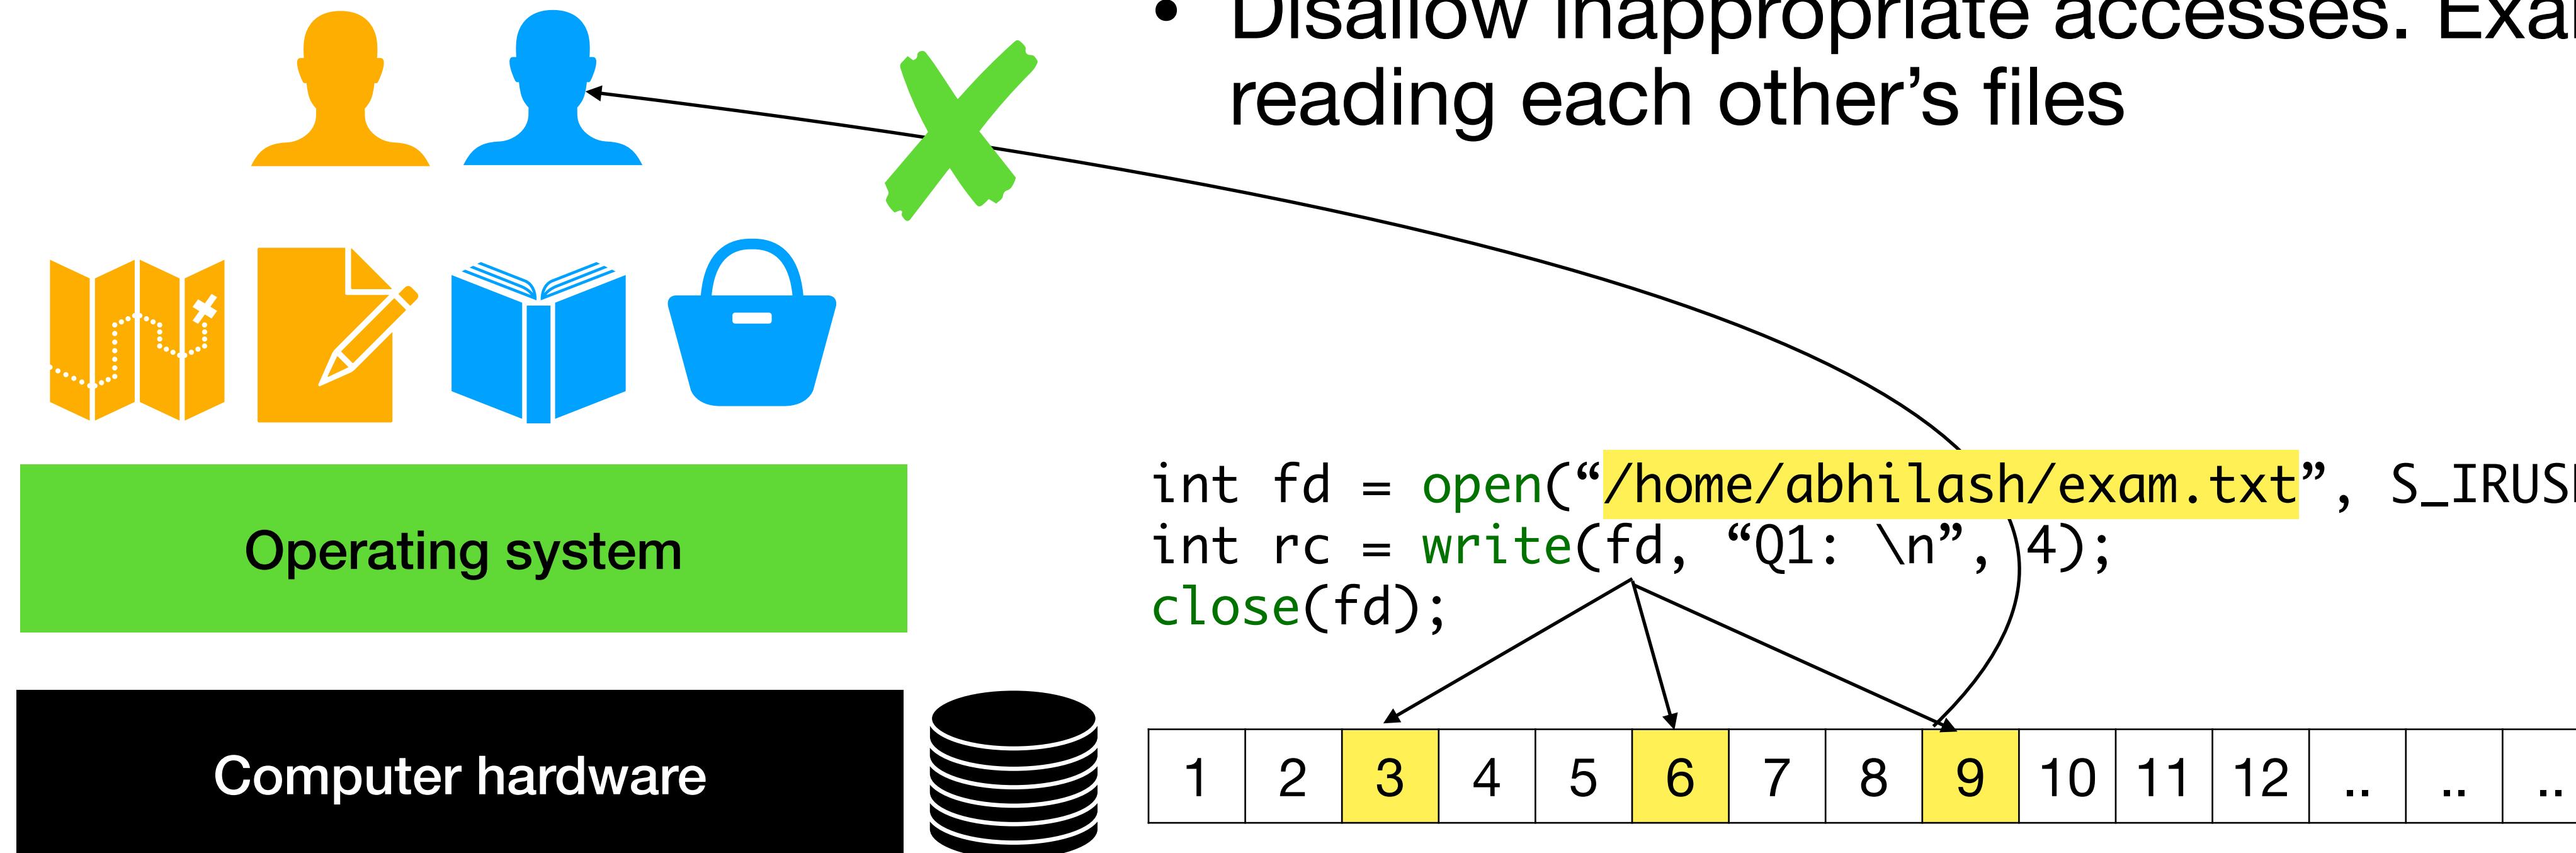
# Purpose of OS: Protection

- S\_IRUSR | S\_IWUSR: File can only be rw by user



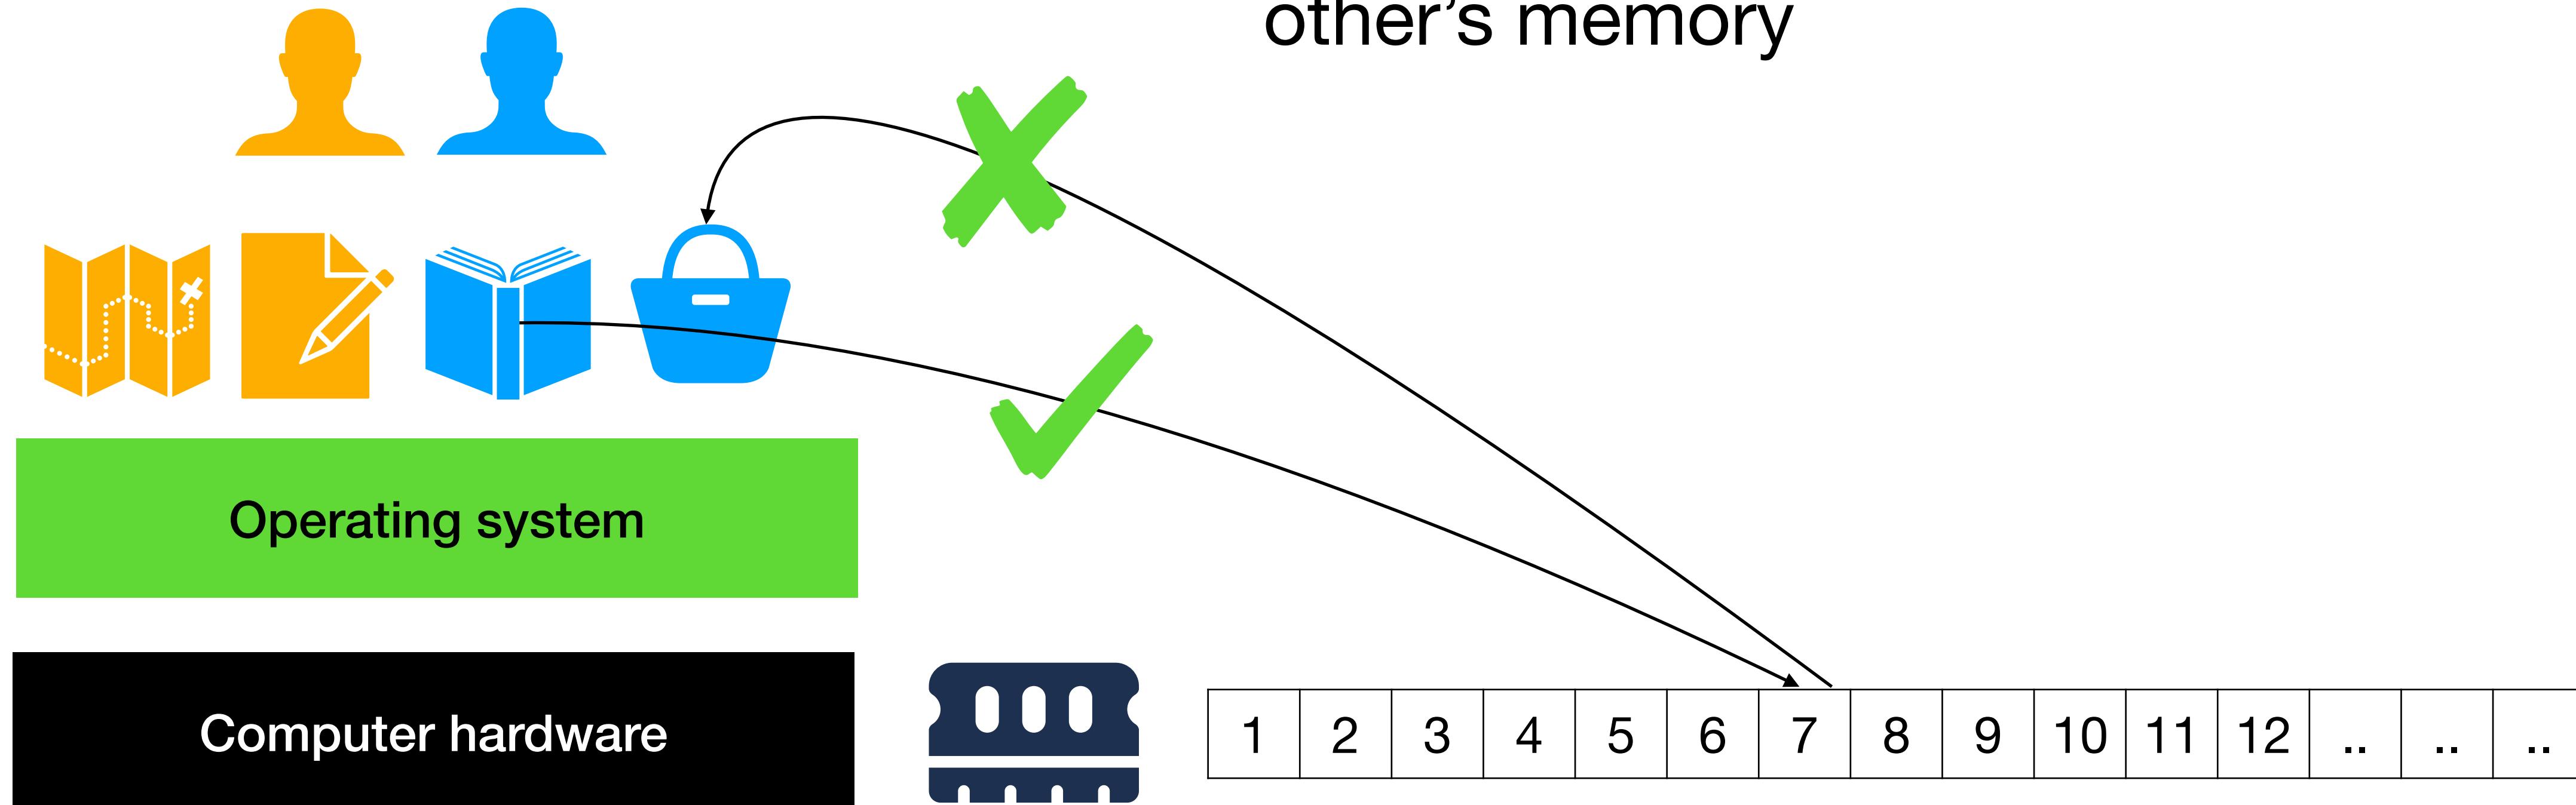
# Purpose of OS: Protection

- S\_IRUSR | S\_IWUSR: File can only be rw by user
- Disallow inappropriate accesses. Example: users reading each other's files

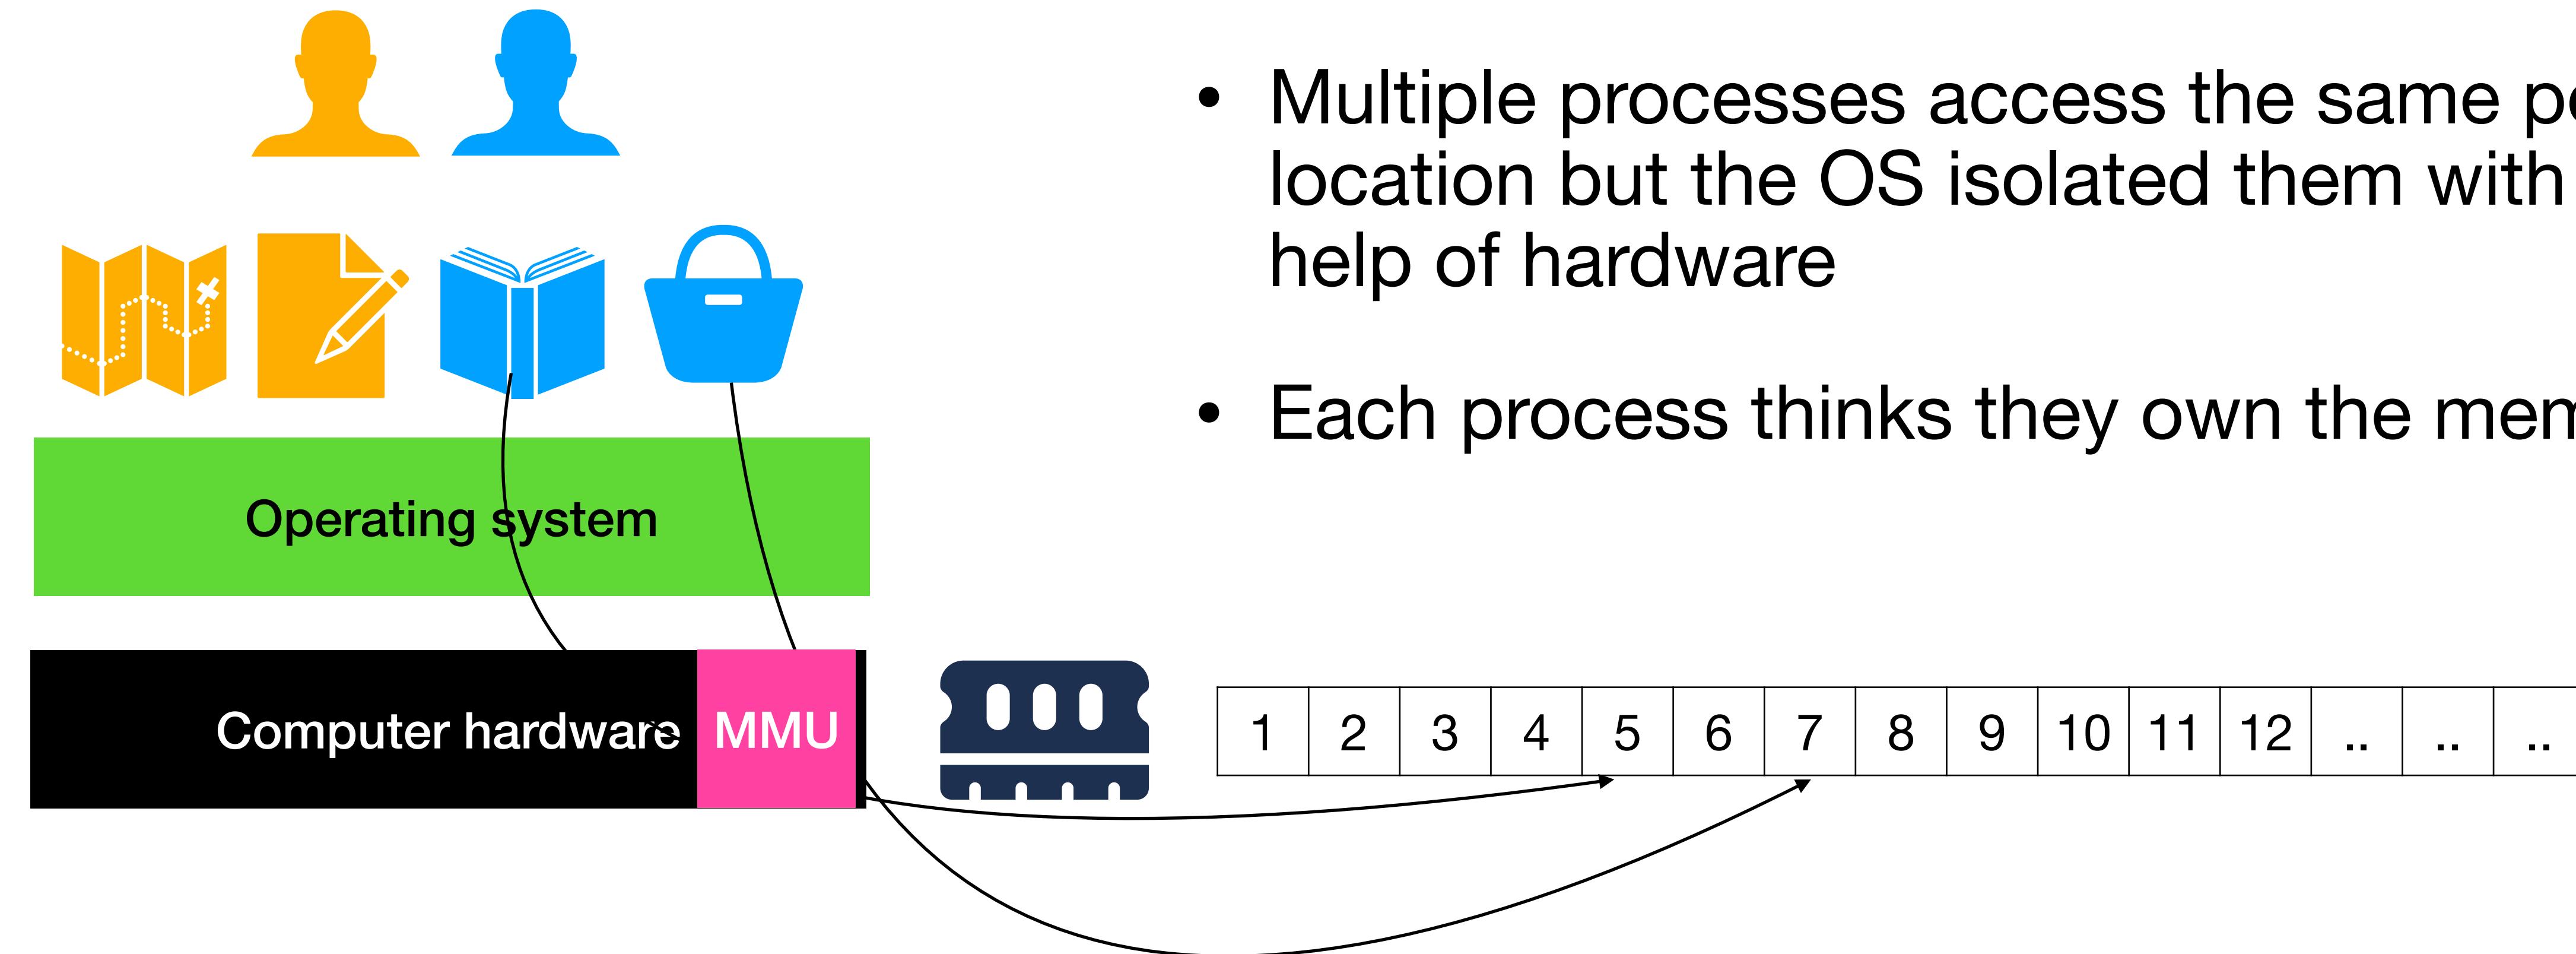


# Purpose of OS: Protection

- Disallow inappropriate accesses.  
Example: processes reading each other's memory



# Purpose of OS: Isolation



# **Course structure: OS in action**

# Course structure: OS in action

- We will build an OS (xv6) from scratch

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
  - Concurrency: data races, different types of locks

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
  - Concurrency: data races, different types of locks
  - Memory virtualization: memory hierarchy, address translation mechanism, demand paging, thrashing, fork system call

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
  - Concurrency: data races, different types of locks
  - Memory virtualization: memory hierarchy, address translation mechanism, demand paging, thrashing, fork system call
  - Shell: Pipes, IO redirection

# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
  - Concurrency: data races, different types of locks
  - Memory virtualization: memory hierarchy, address translation mechanism, demand paging, thrashing, fork system call
  - Shell: Pipes, IO redirection
  - Parallelism: Enable more CPUs, revisit locks

# Data races due to concurrency

<b>Thread 1</b>	<b>Thread 2</b>
Read counter = 0	
Write counter = 1	
	Read counter = 1
	Writer counter = 2
Read counter = 2	
	Read counter = 2
	Writer counter = 3
Writer counter = 3	

# Data races due to concurrency

./threads 100000

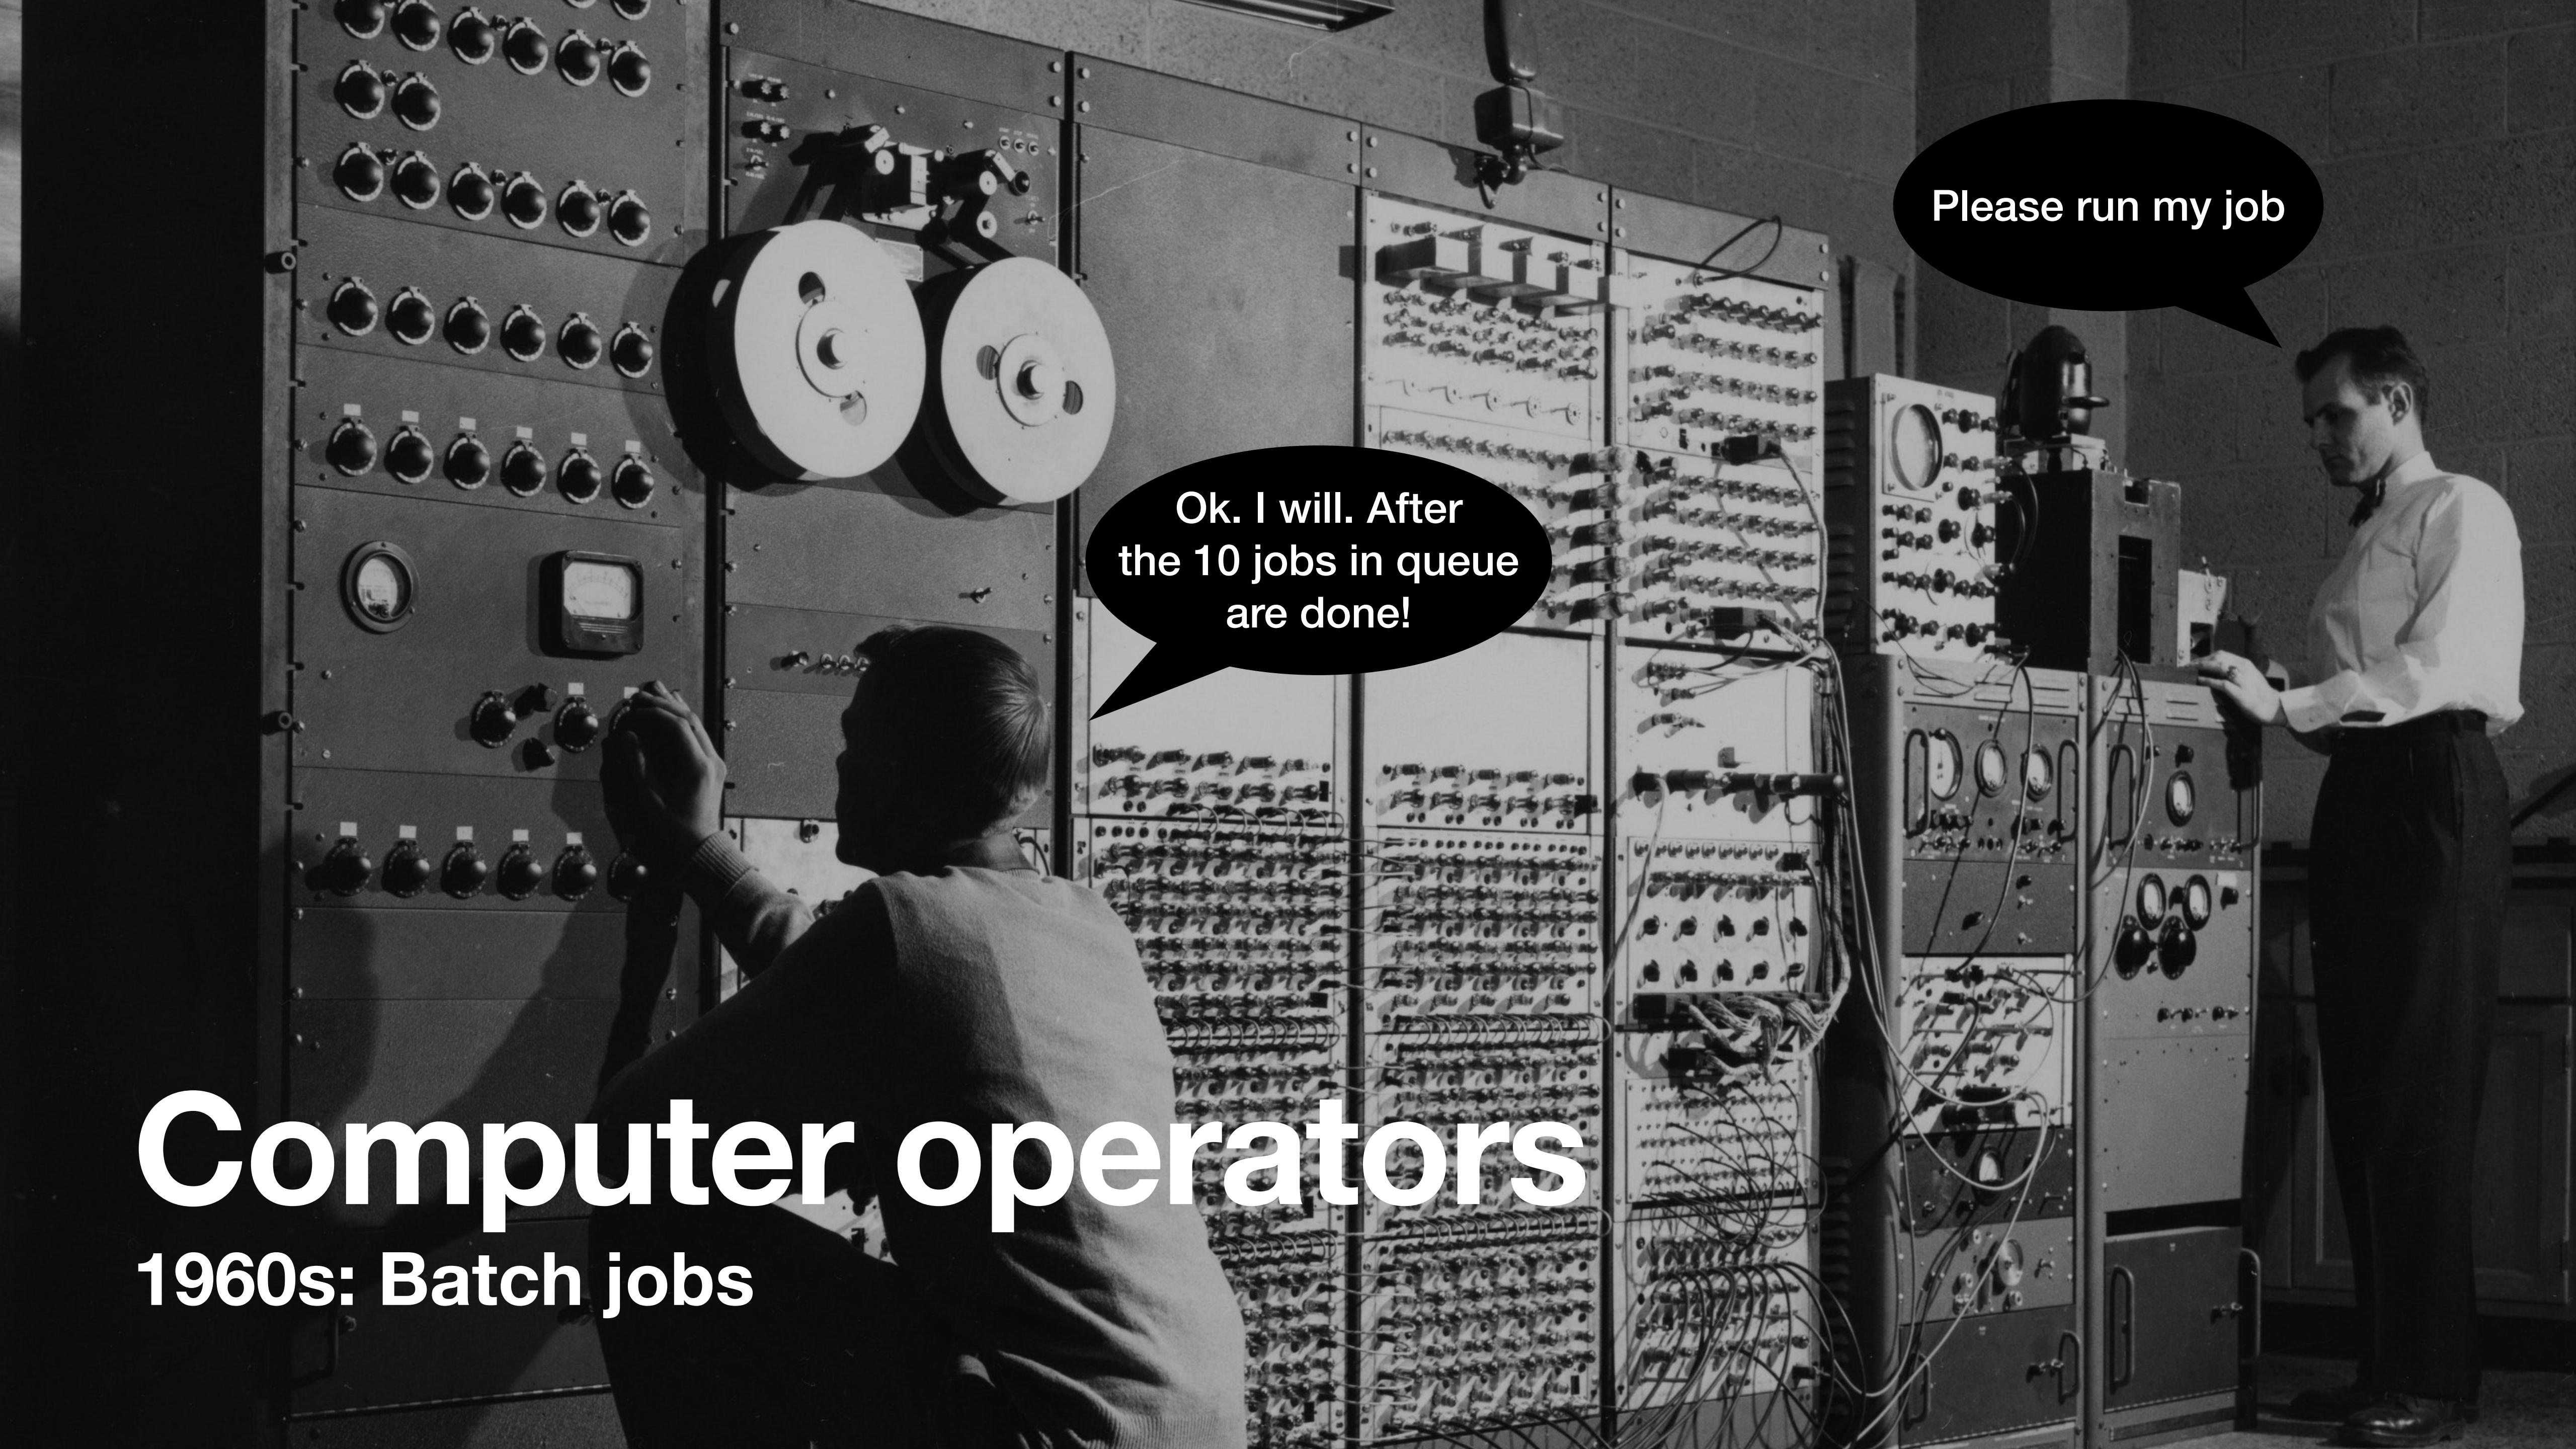
Thread 1	Thread 2
Read counter = 0	
Write counter = 1	
	Read counter = 1
	Writer counter = 2
Read counter = 2	
	Read counter = 2
	Writer counter = 3
Writer counter = 3	

**Why should I learn OS in 2026?  
Isn't it a solved problem?**

**We have indeed made good progress ..**

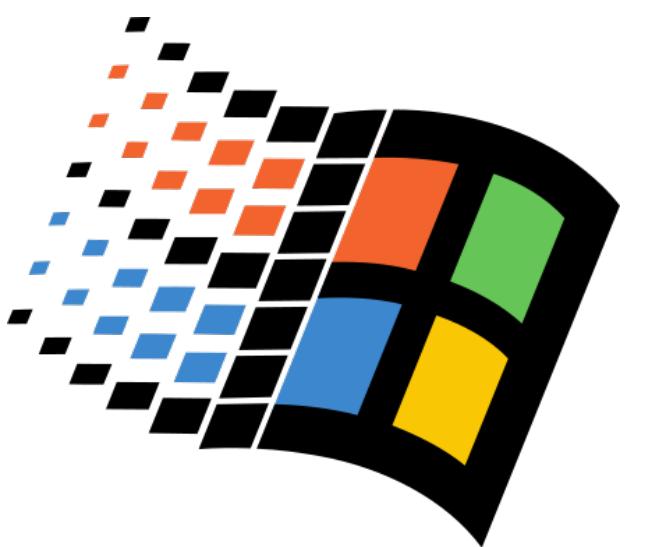
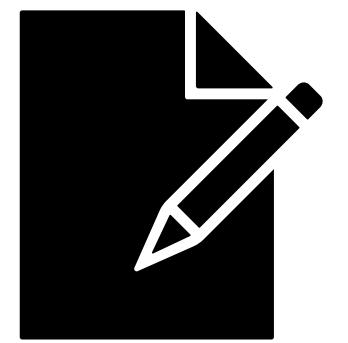
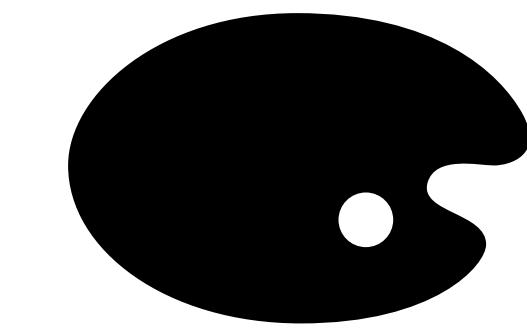
# Computer operators

## 1960s: Batch jobs



# Personal computers

1980s: Interactive jobs!

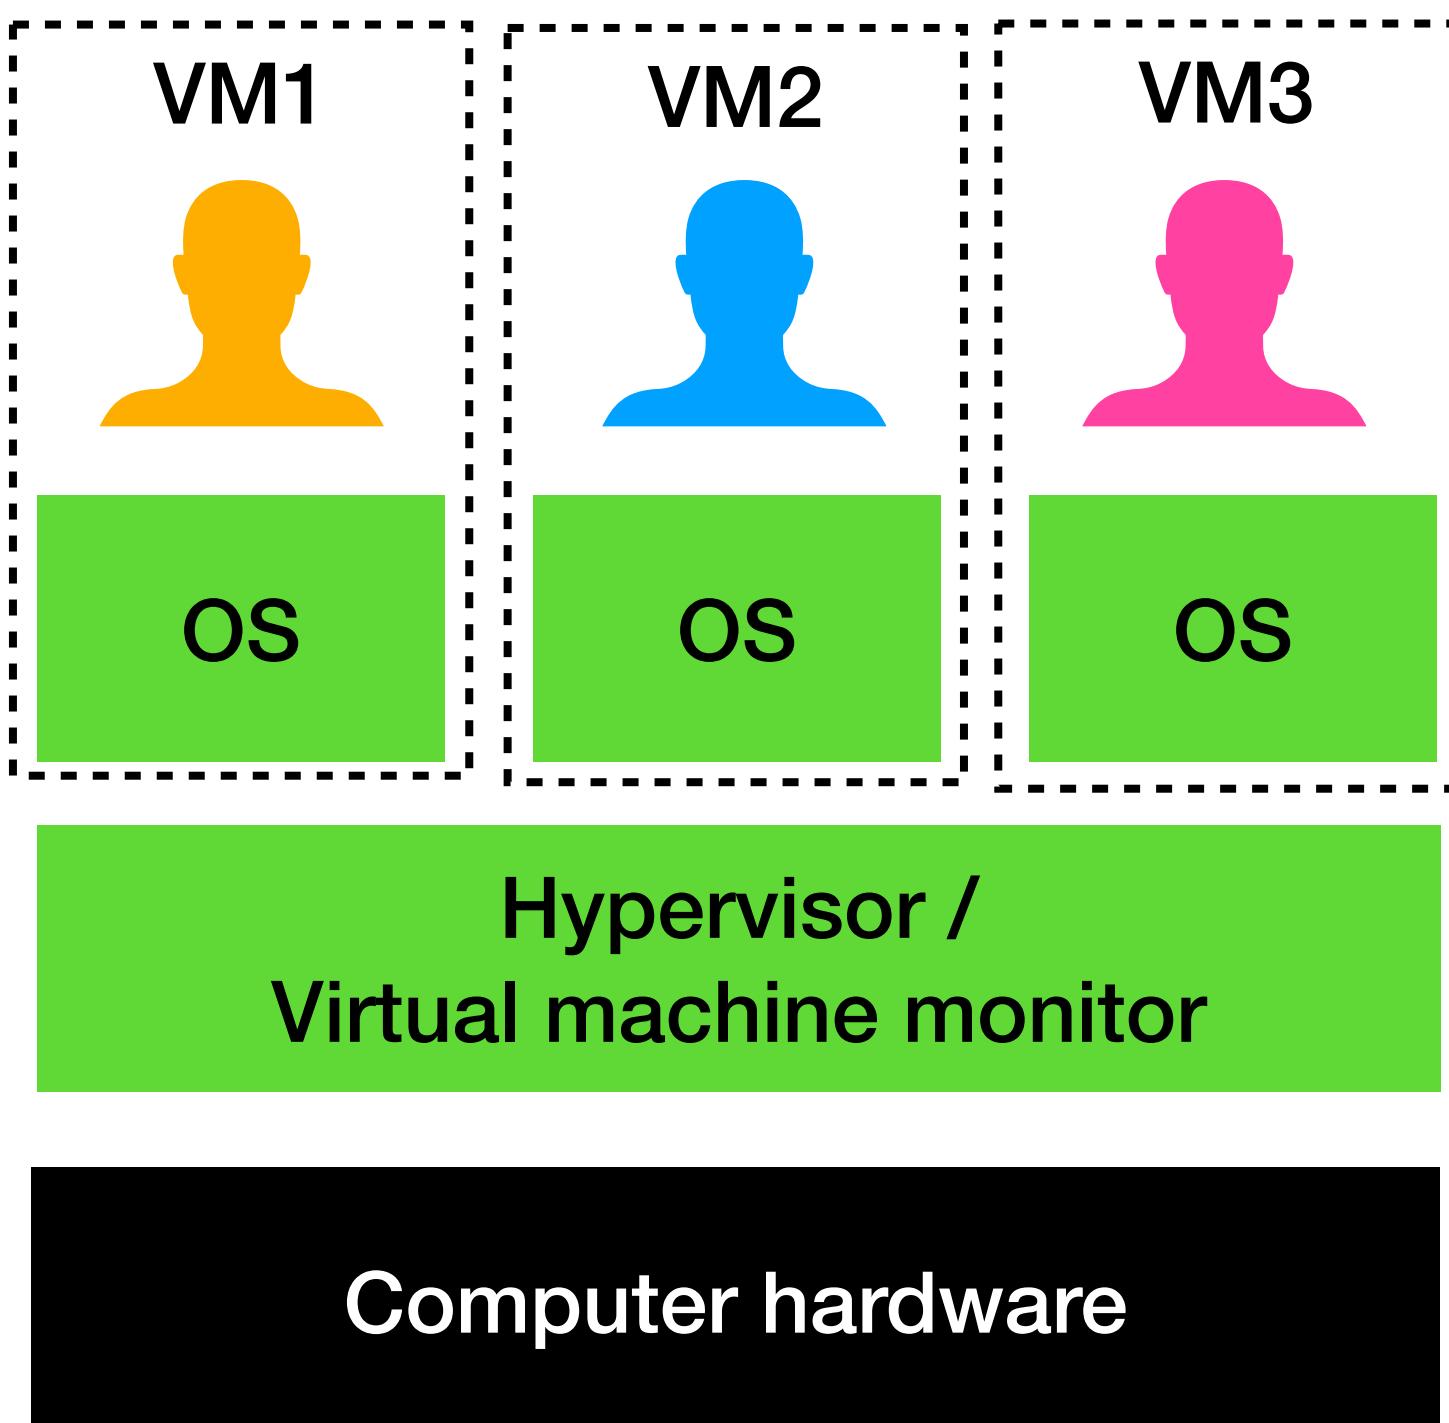


**UNIX.**<sup>®</sup>

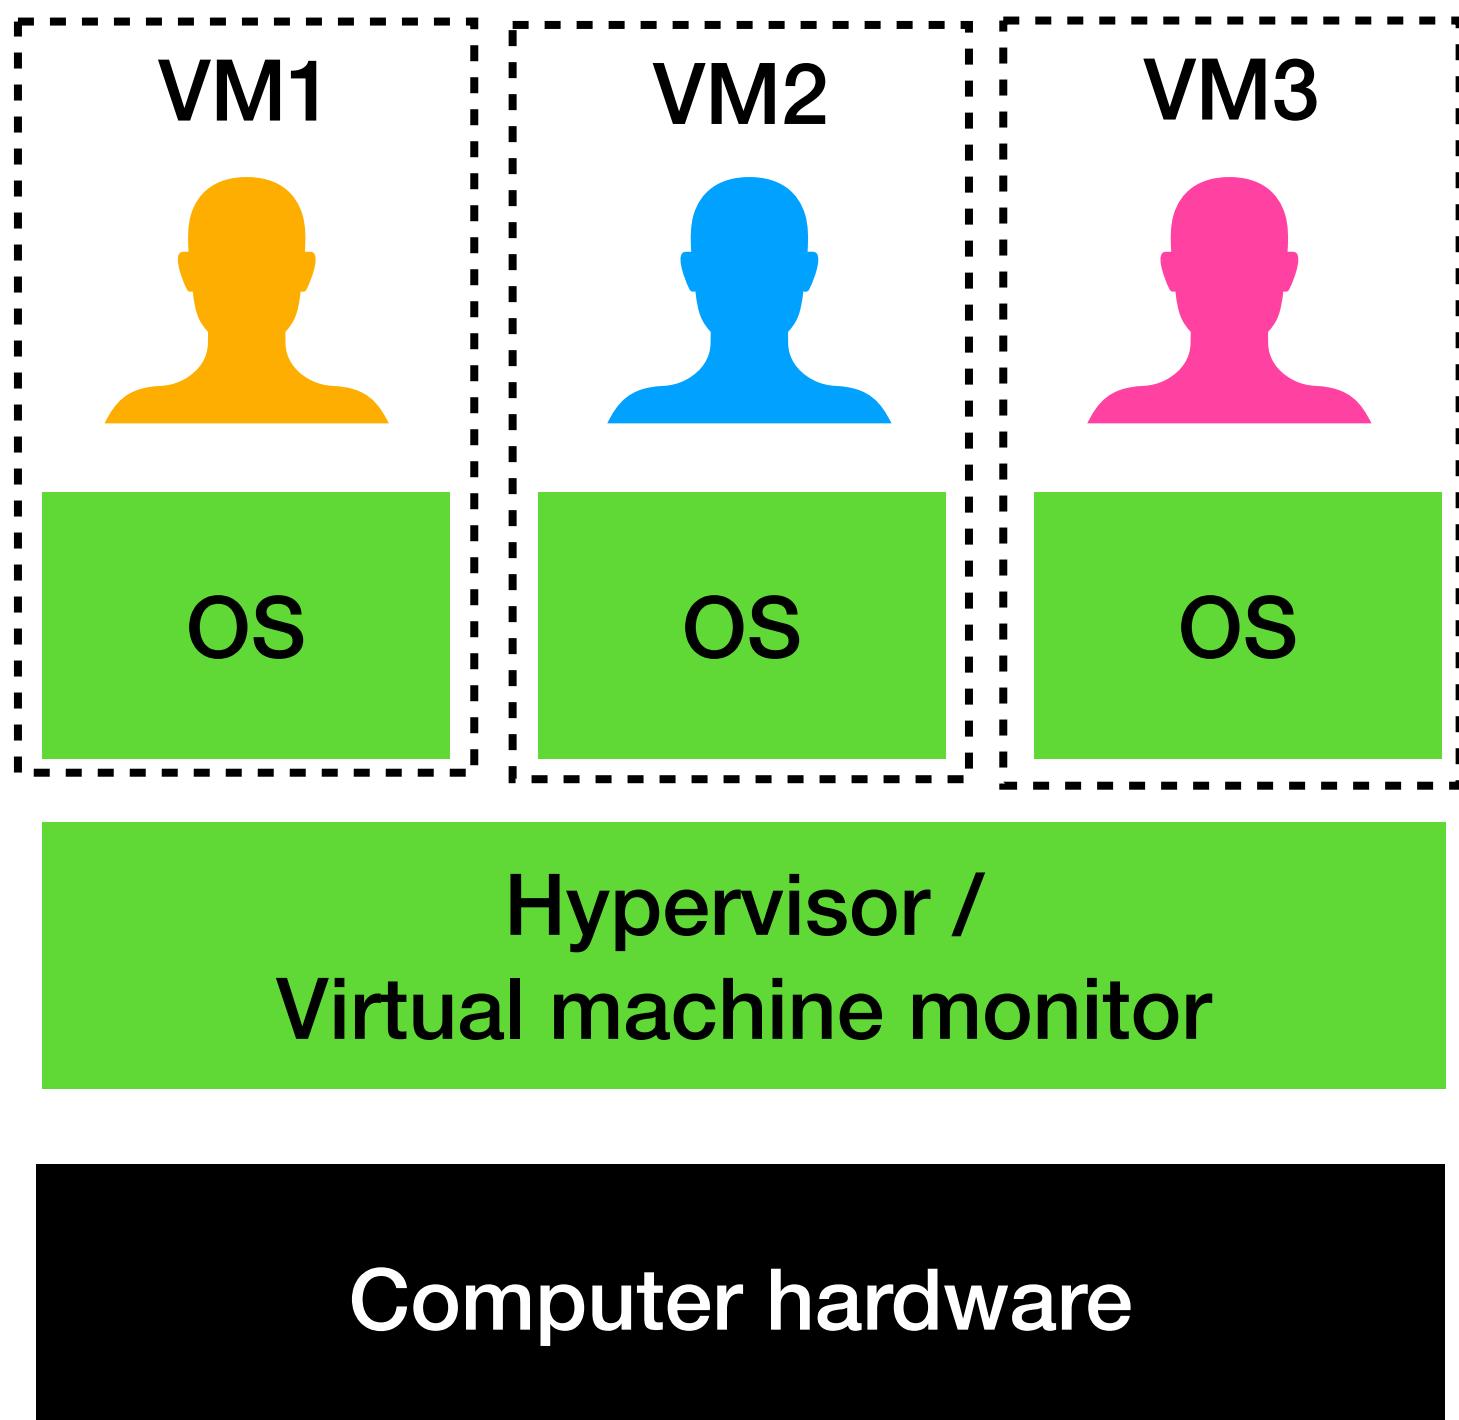
00011110 00011110 00011110 00011110 00011110 00011110



# 2000s: Cloud

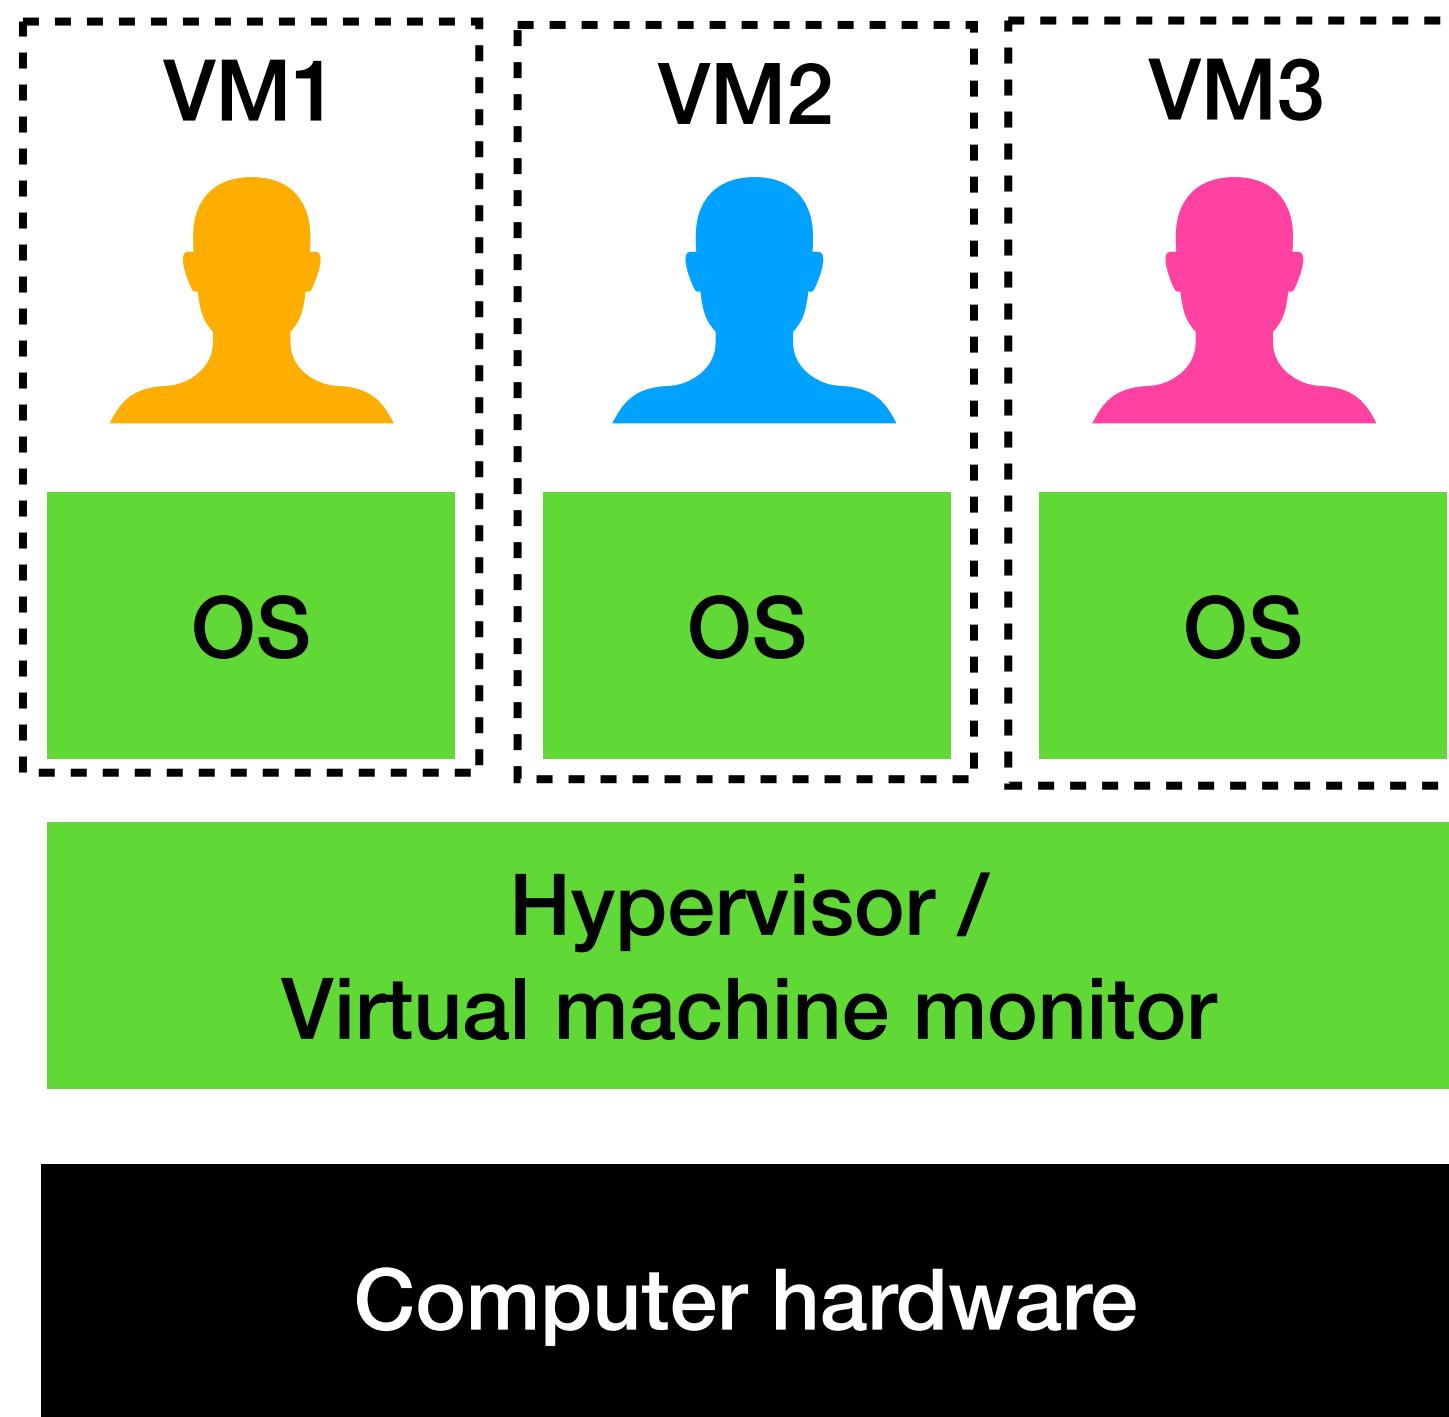


# 2000s: Cloud



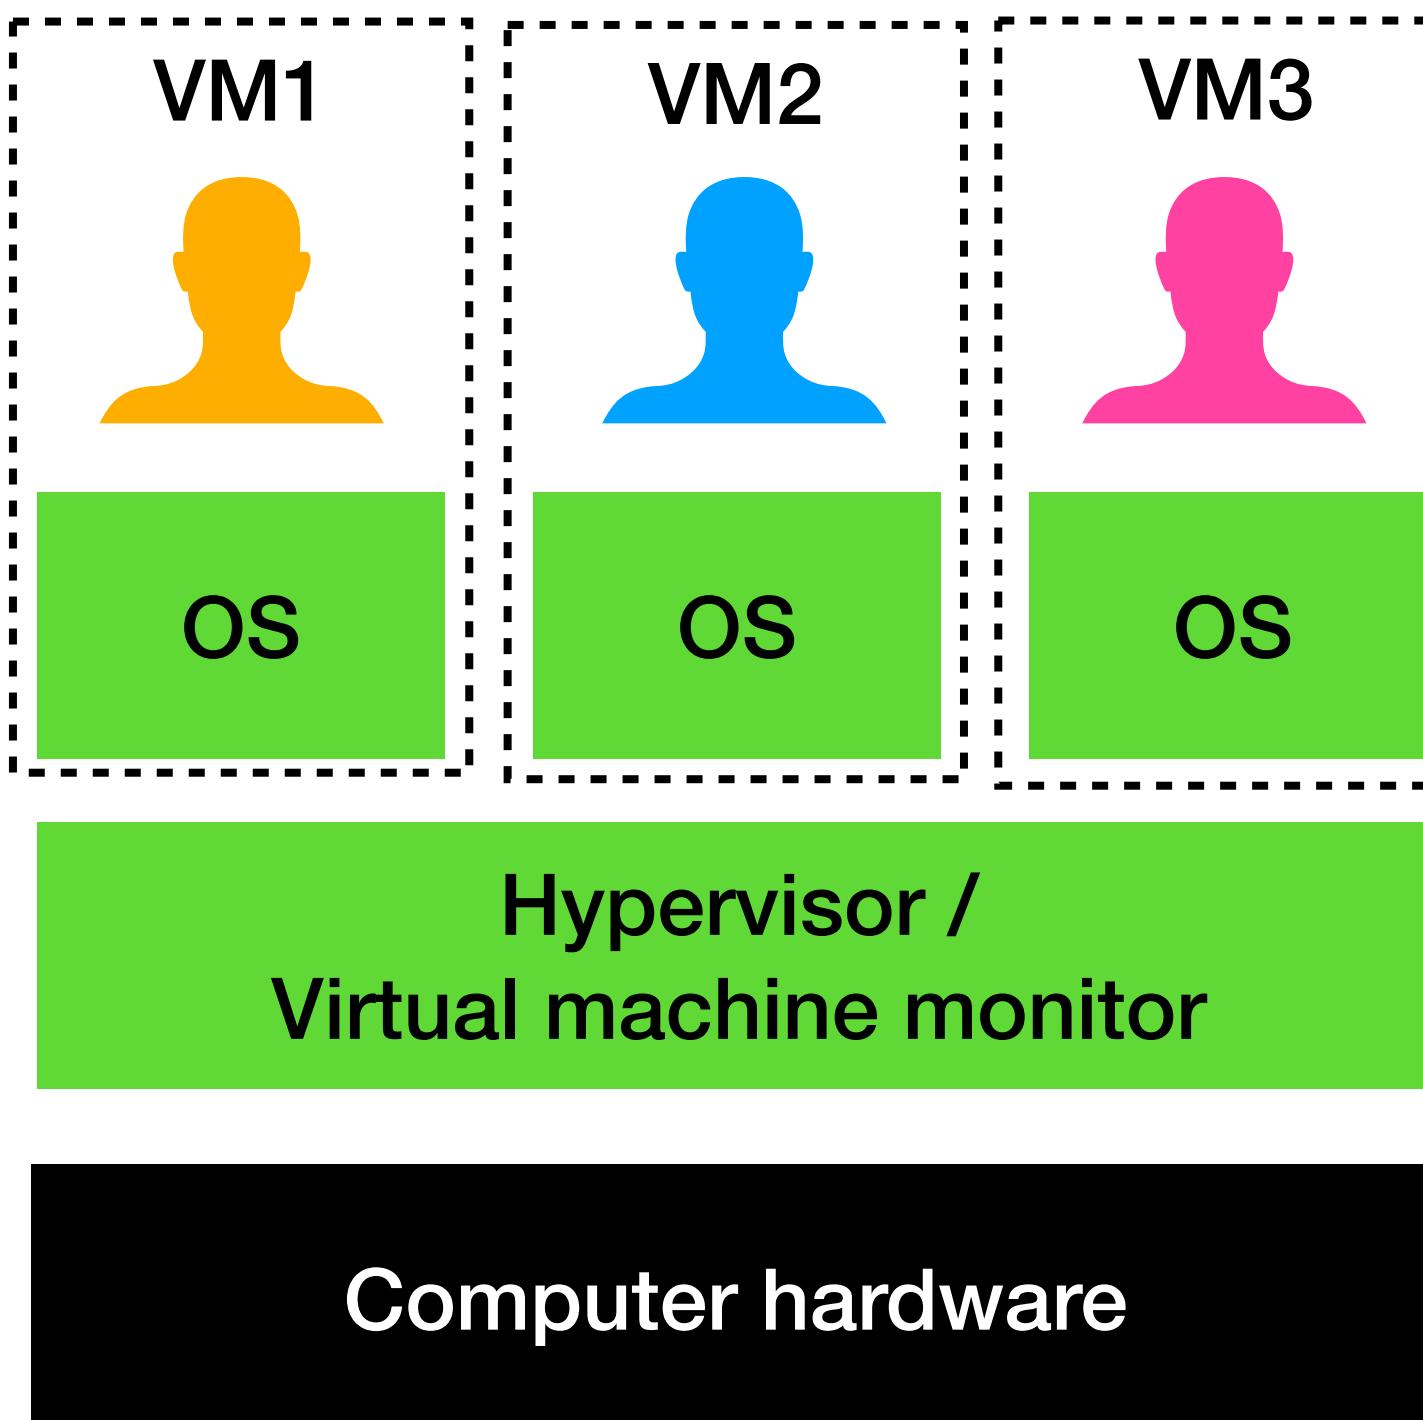
- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.

# 2000s: Cloud



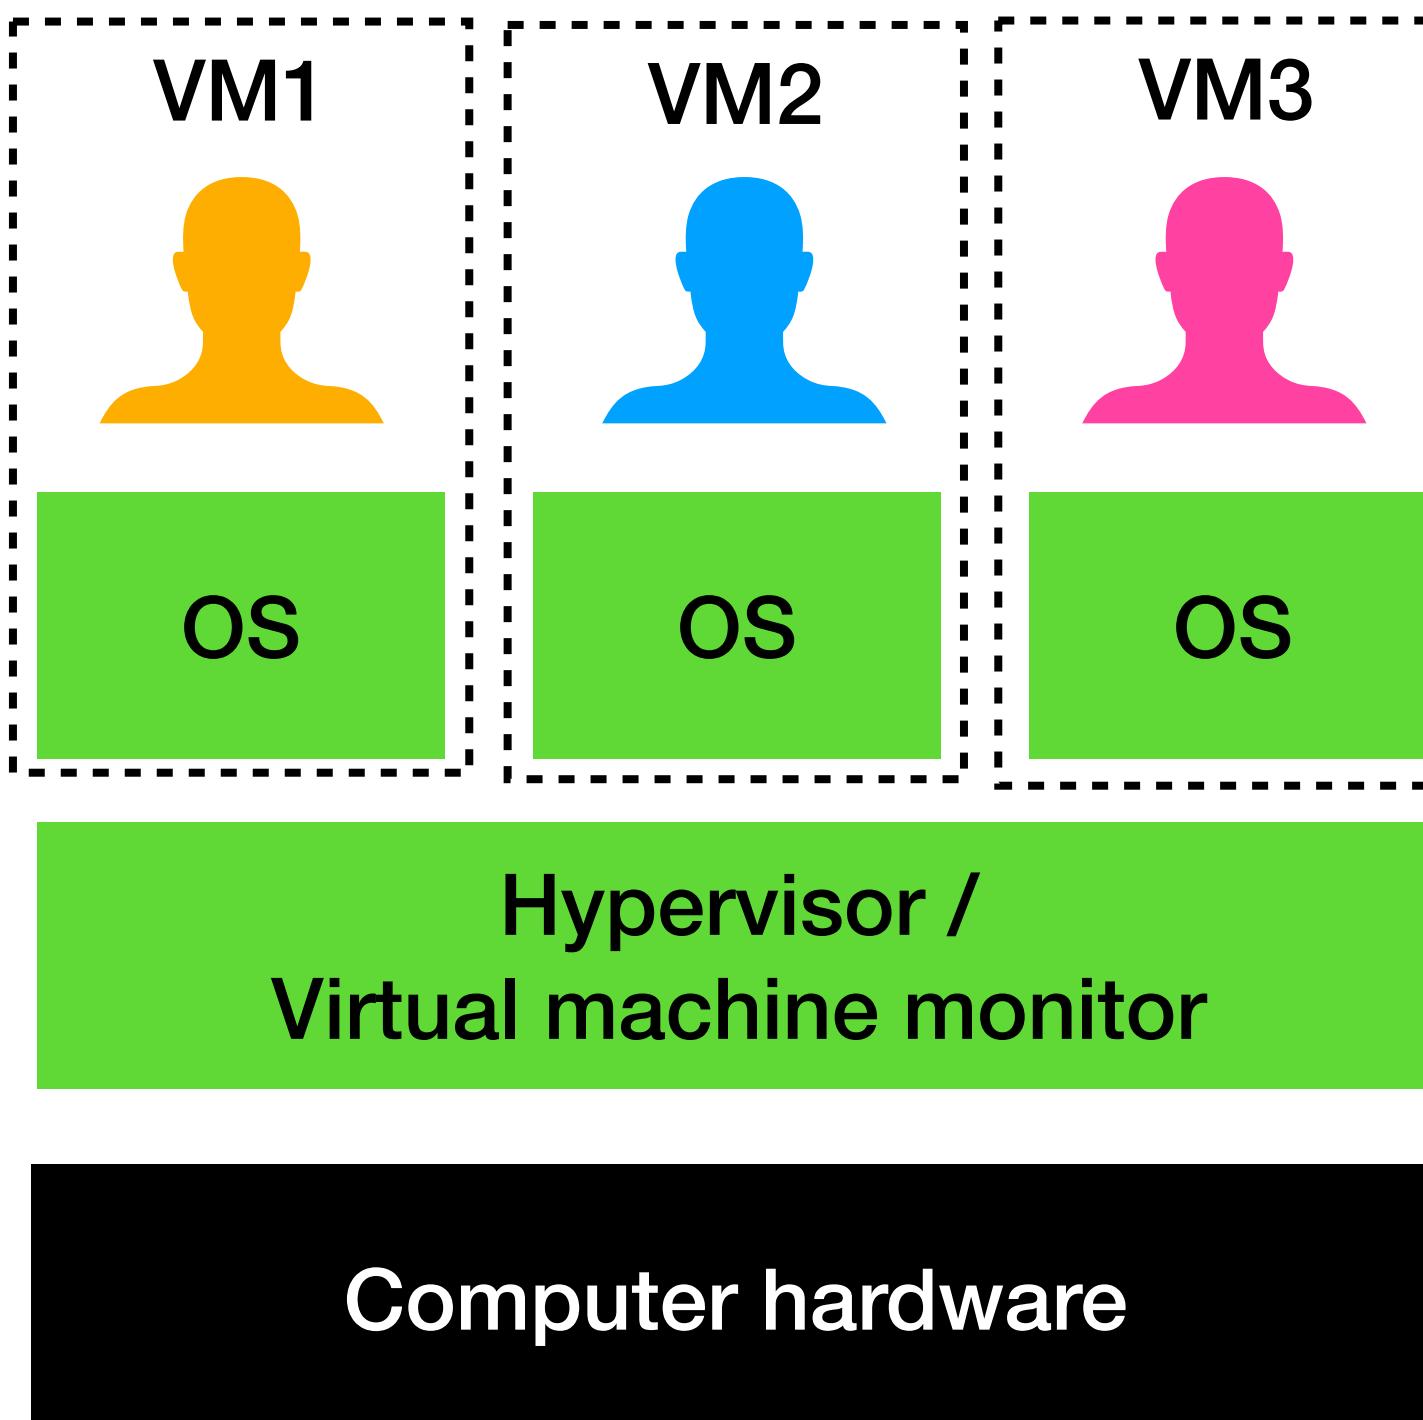
- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes

# 2000s: Cloud



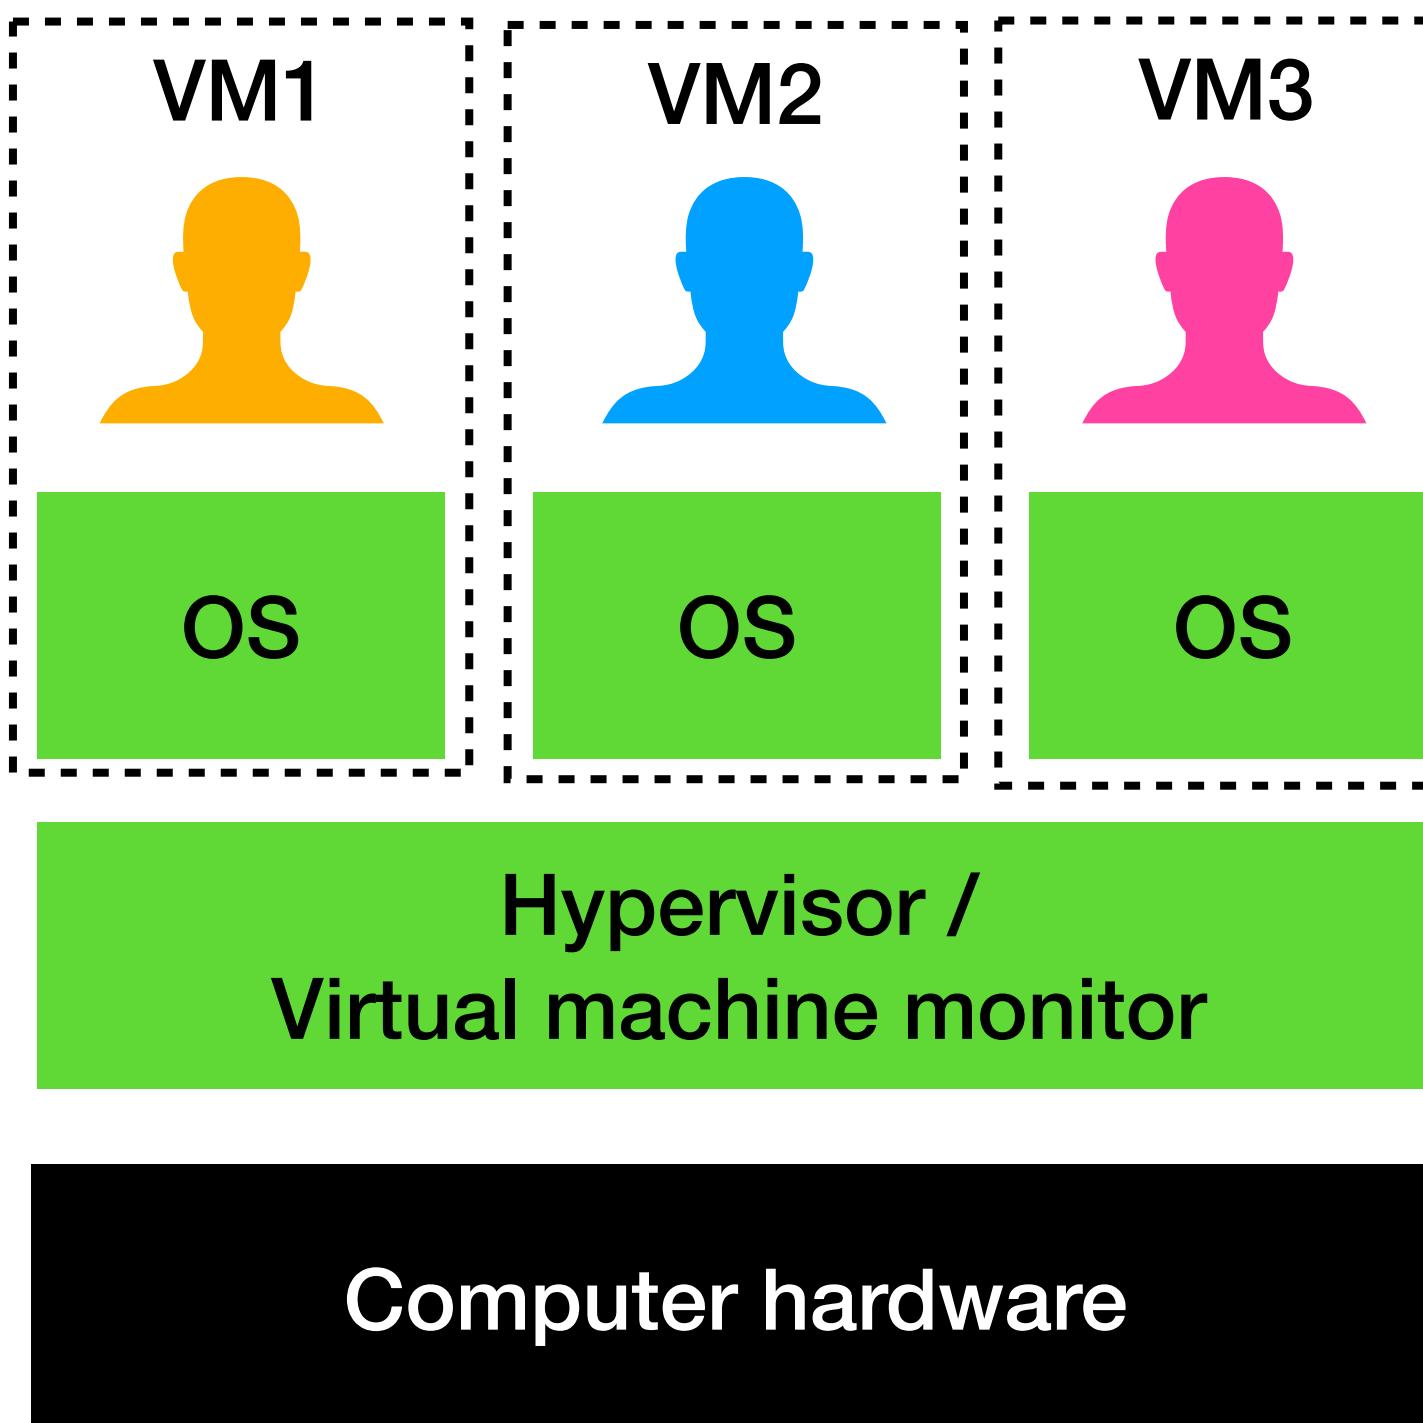
- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes
  - multiplexes hardware among OS

# 2000s: Cloud



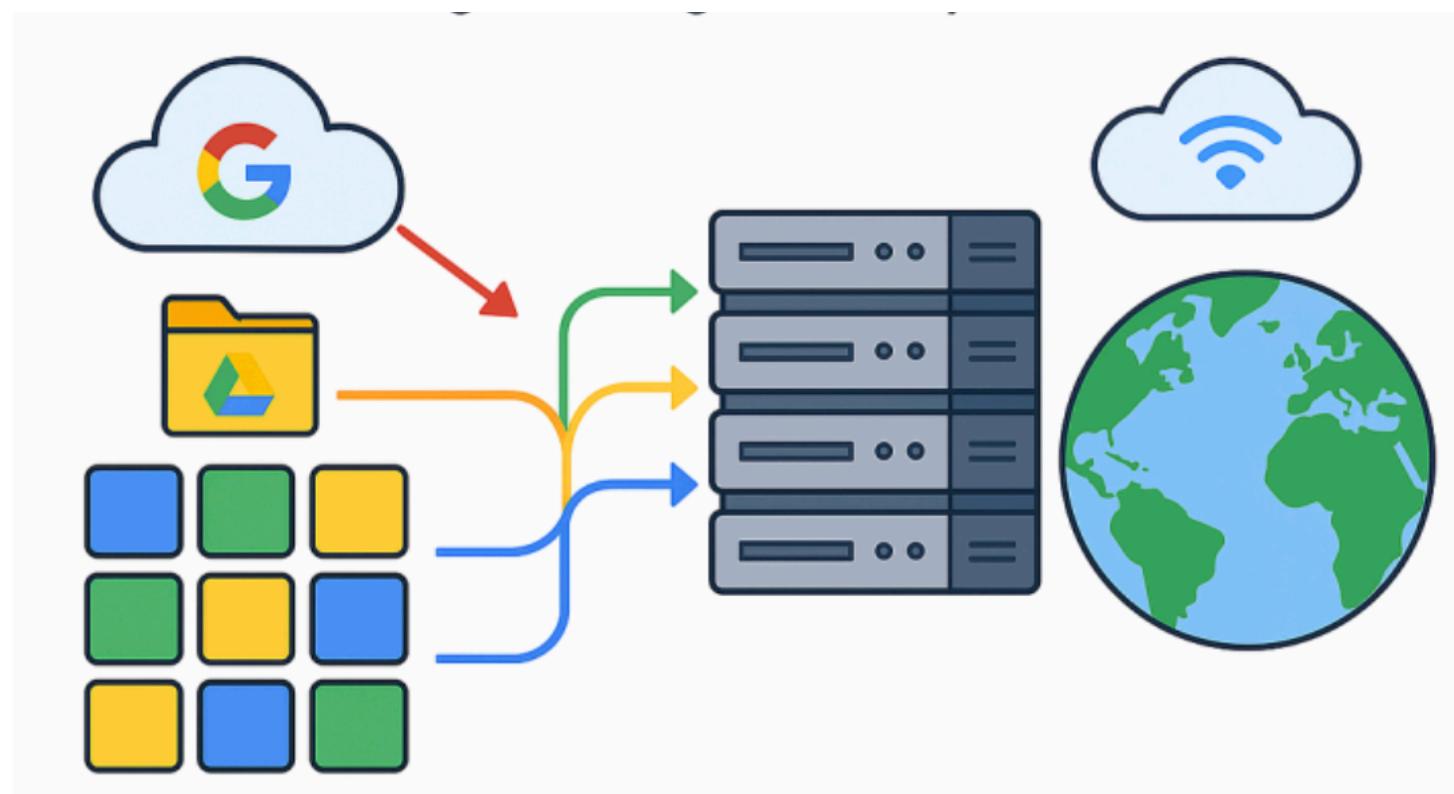
- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes
  - multiplexes hardware among OS
  - protects and isolates OS from each other

# 2000s: Cloud



- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes
  - multiplexes hardware among OS
  - protects and isolates OS from each other
- Hypervisors fundamentally enabled cloud computing

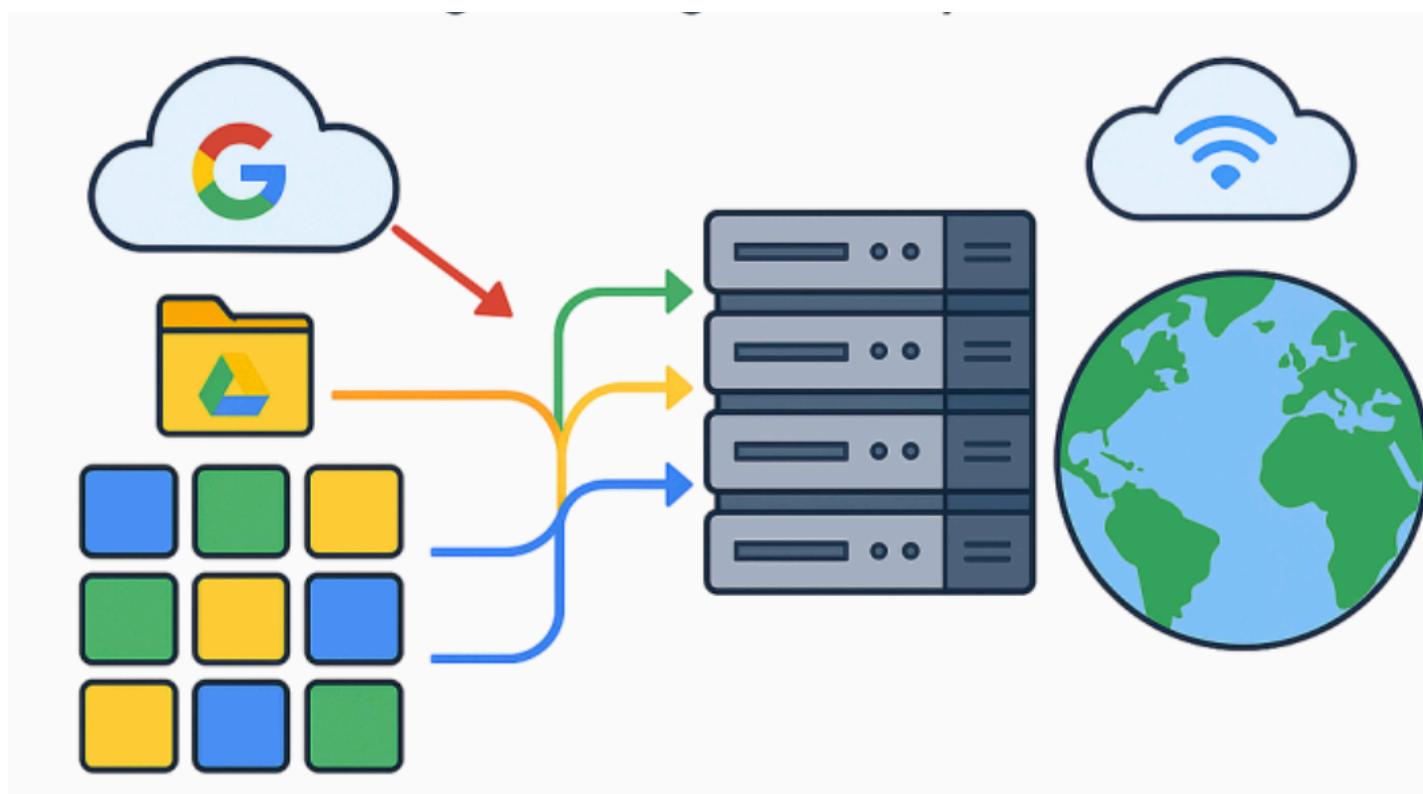
# 2000s: Distributed systems



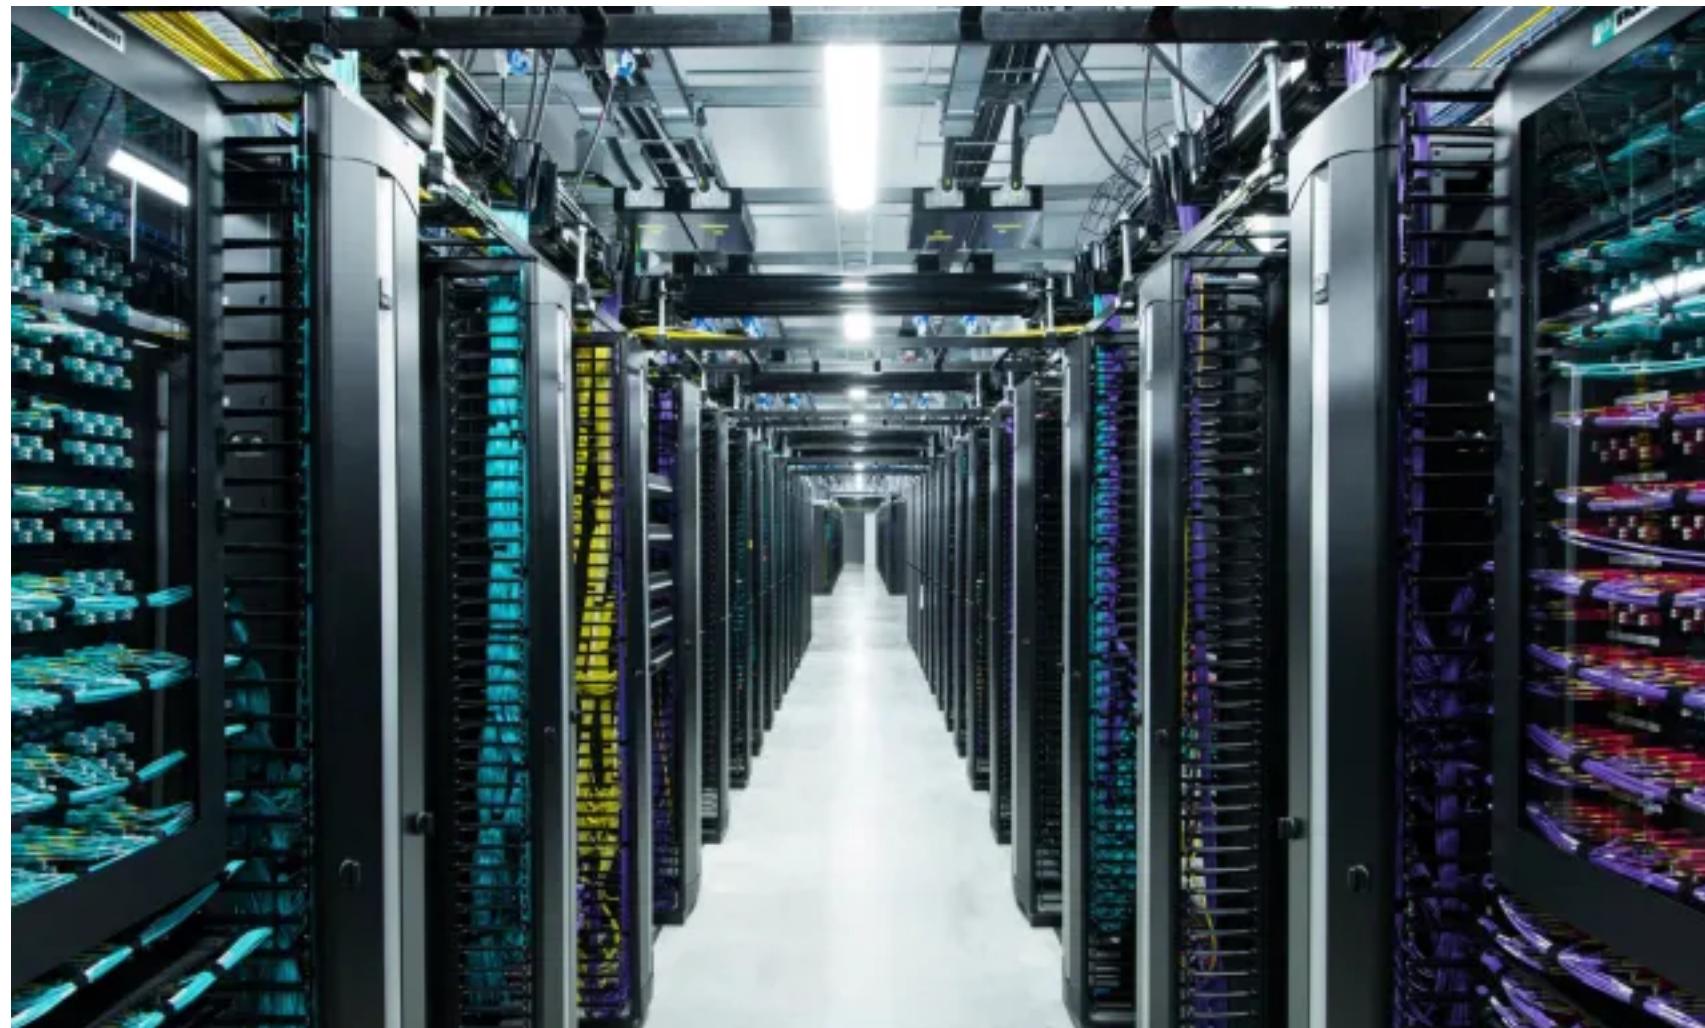
# 2000s: Distributed systems



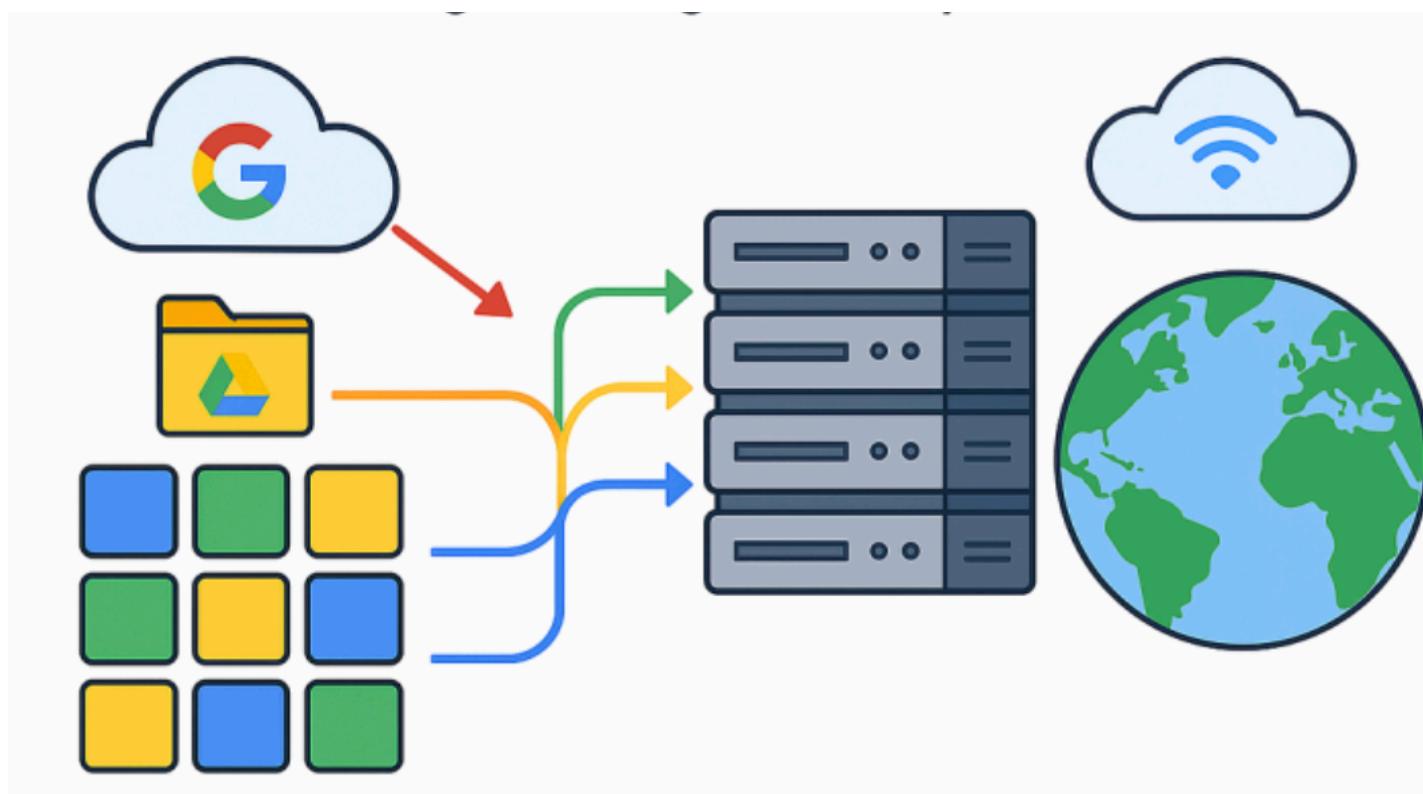
- My data does not fit in a single machine.  
Distributed file systems; e.g., crawled web data is stored in Google File System



# 2000s: Distributed systems



- My data does not fit in a single machine. Distributed file systems; e.g., crawled web data is stored in Google File System
- My “tasks” span 100s-1000s of machines; e.g., computing page rank, training large language models over the entire internet



# 2010s: Smartphones



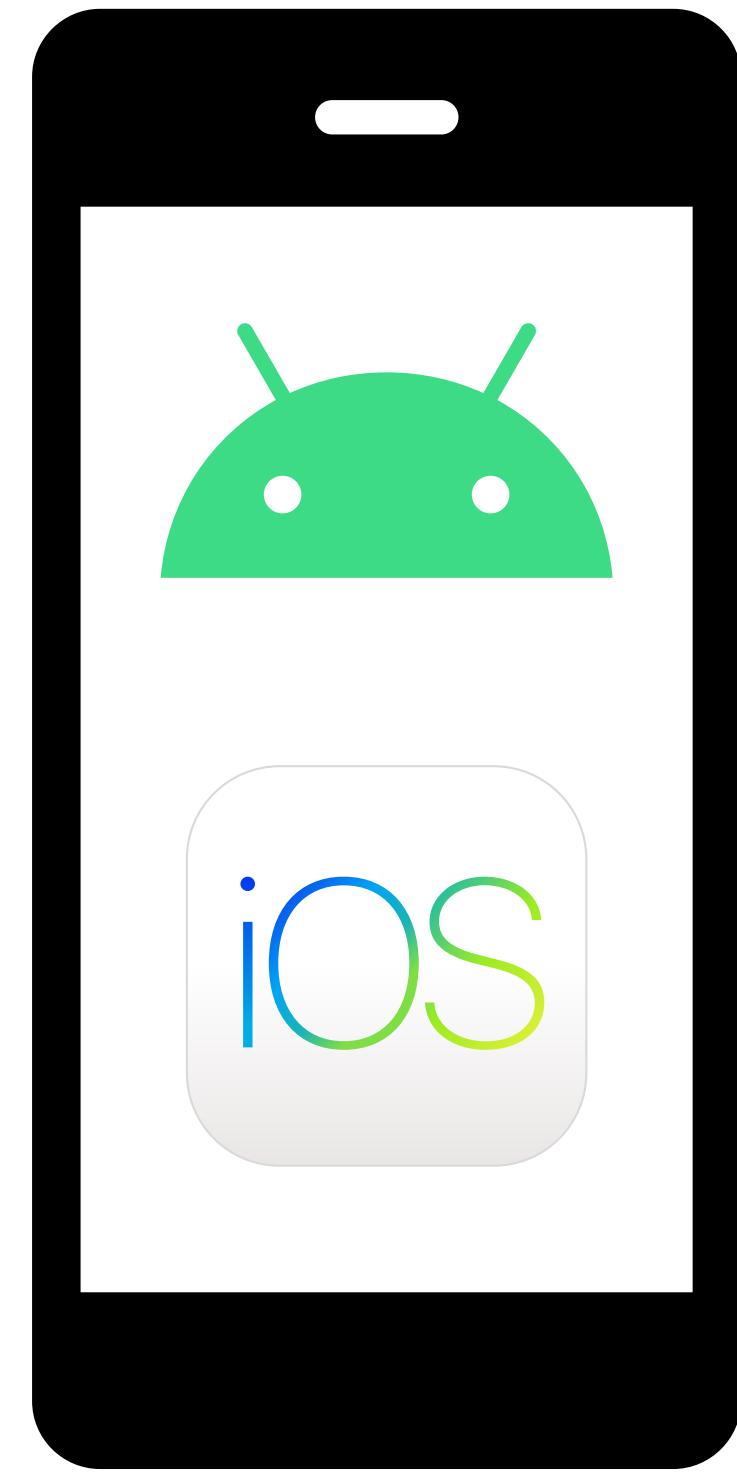
# 2010s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.



# 2010s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.
- Resource constraints: Power management, UI system, etc.



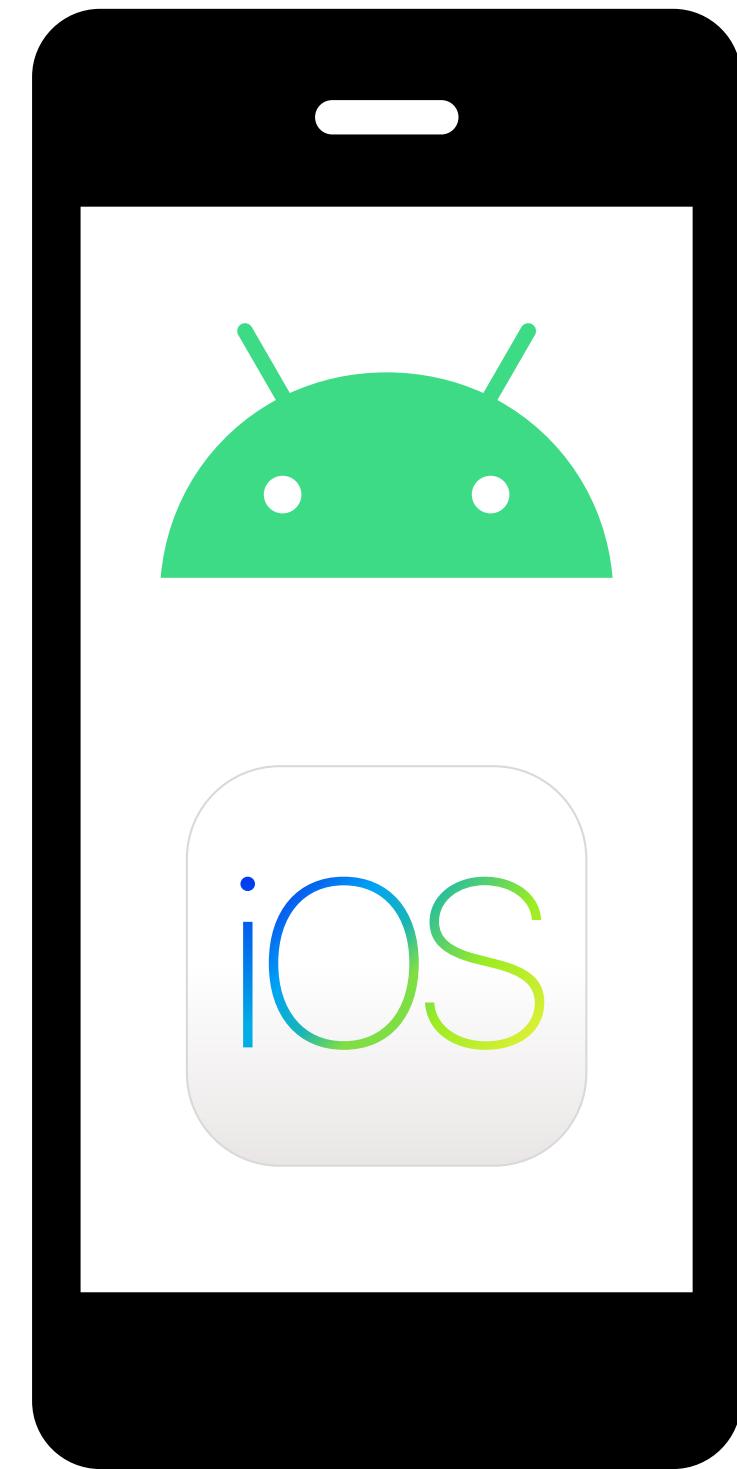
# 2010s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.
- Resource constraints: Power management, UI system, etc.
- Even higher-level services: voice recognition, augmented reality, etc.



# 2010s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.
- Resource constraints: Power management, UI system, etc.
- Even higher-level services: voice recognition, augmented reality, etc.
- Increased security concerns because of increase in sensitive data with rise of mobile banking, UPI etc. and because of moving devices



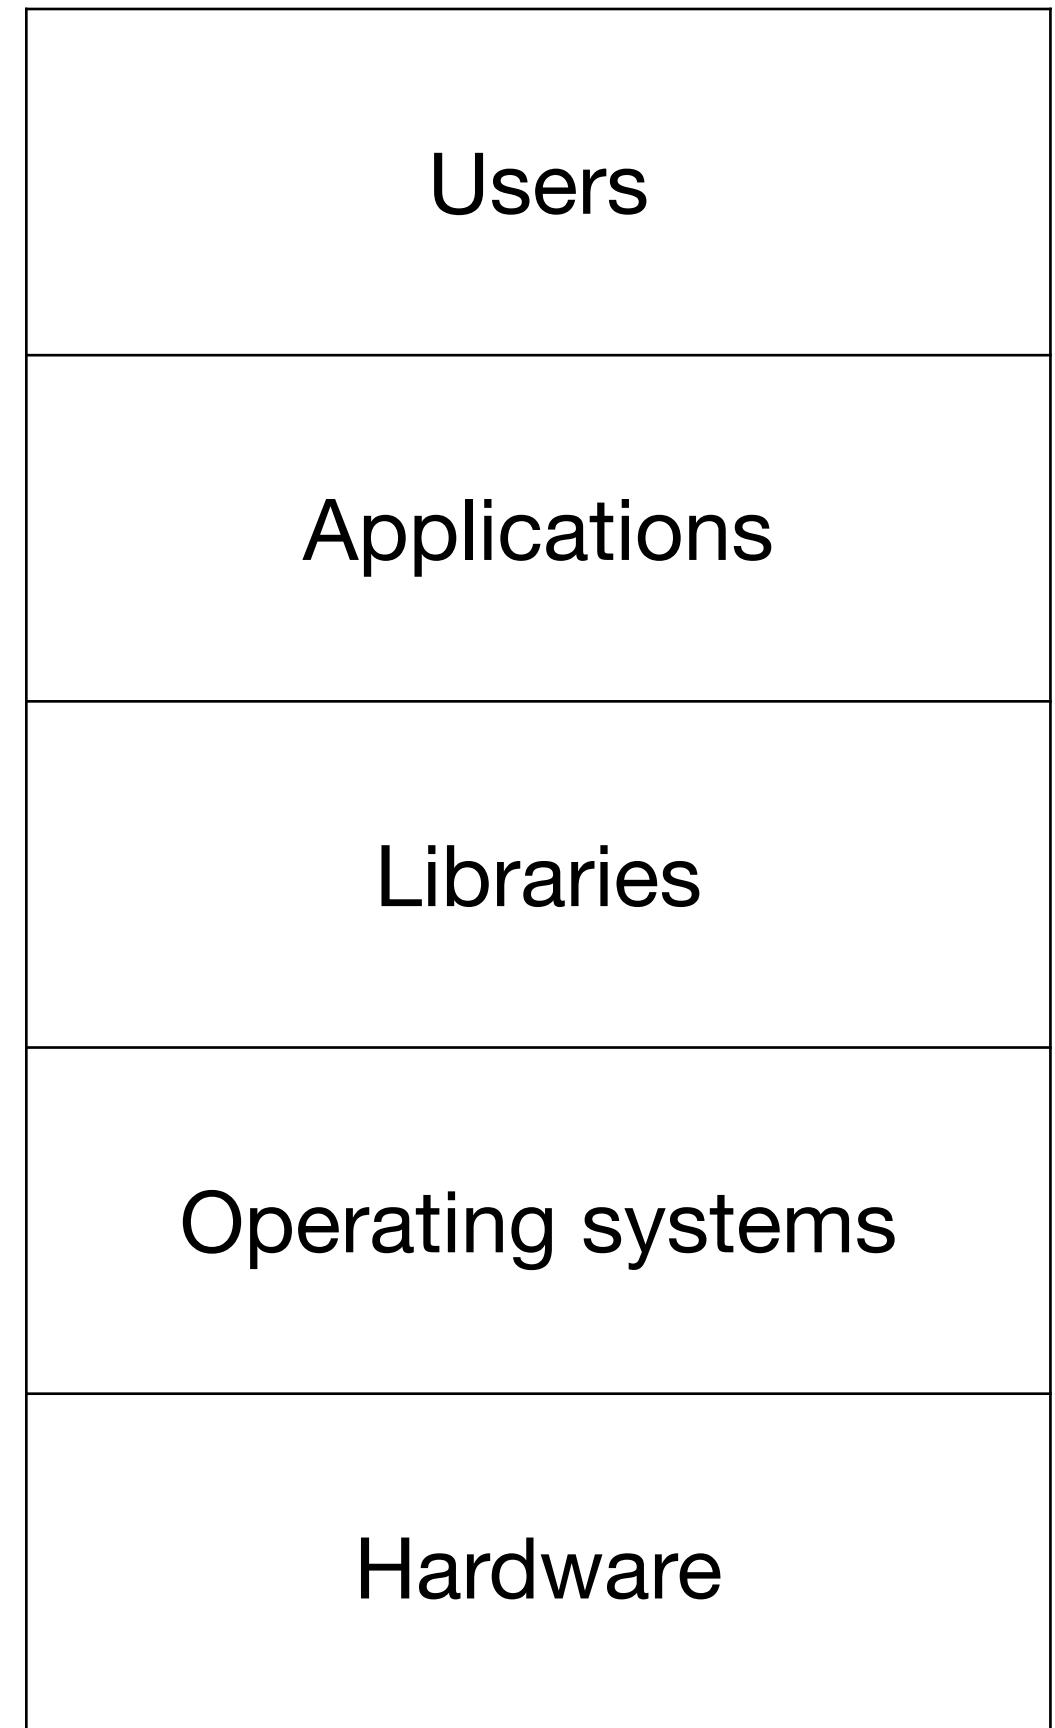


# 2010s: Cyber-physical systems

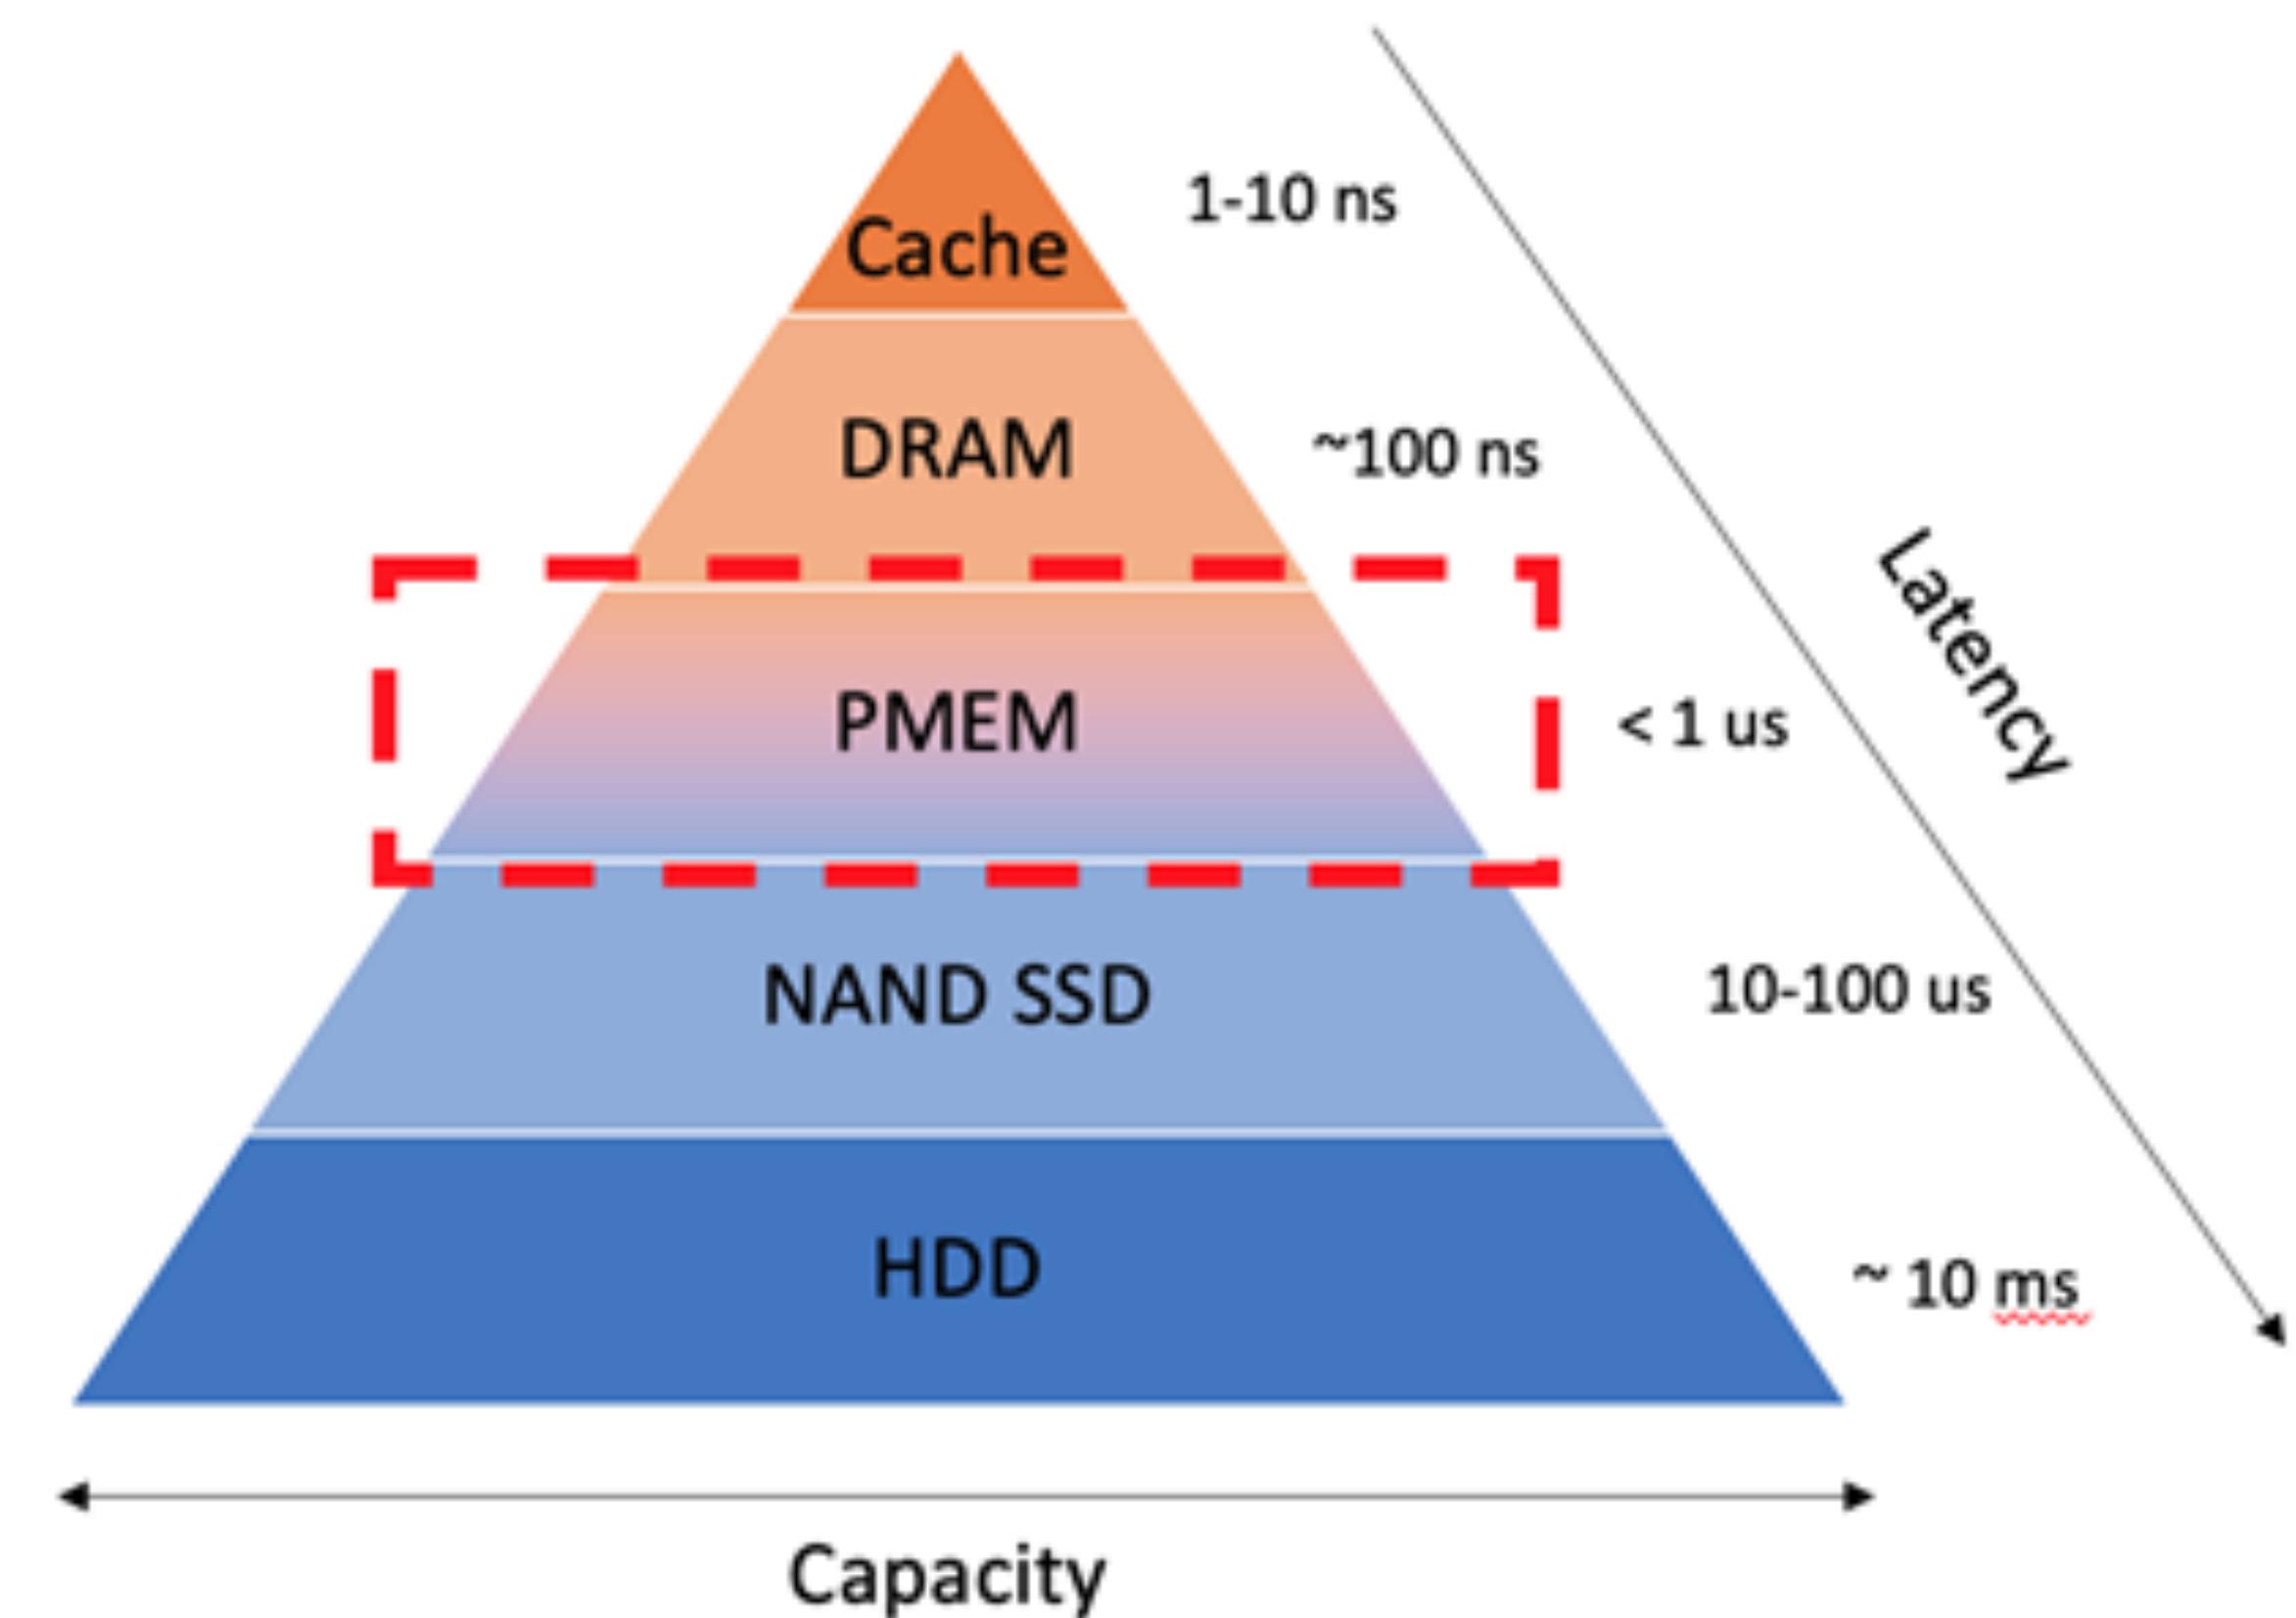
OS must not crash! Formally verified OS. Example: seL4

# Typical progression of systems research

- Systems optimise for the “common case”
  - If common case changes, we need to rethink OS design
- Macro examples
  - Personal computers: batch jobs to interactive jobs
  - Smartphones: resource constraints, new sensors
  - Cyber physical systems: risk of human life

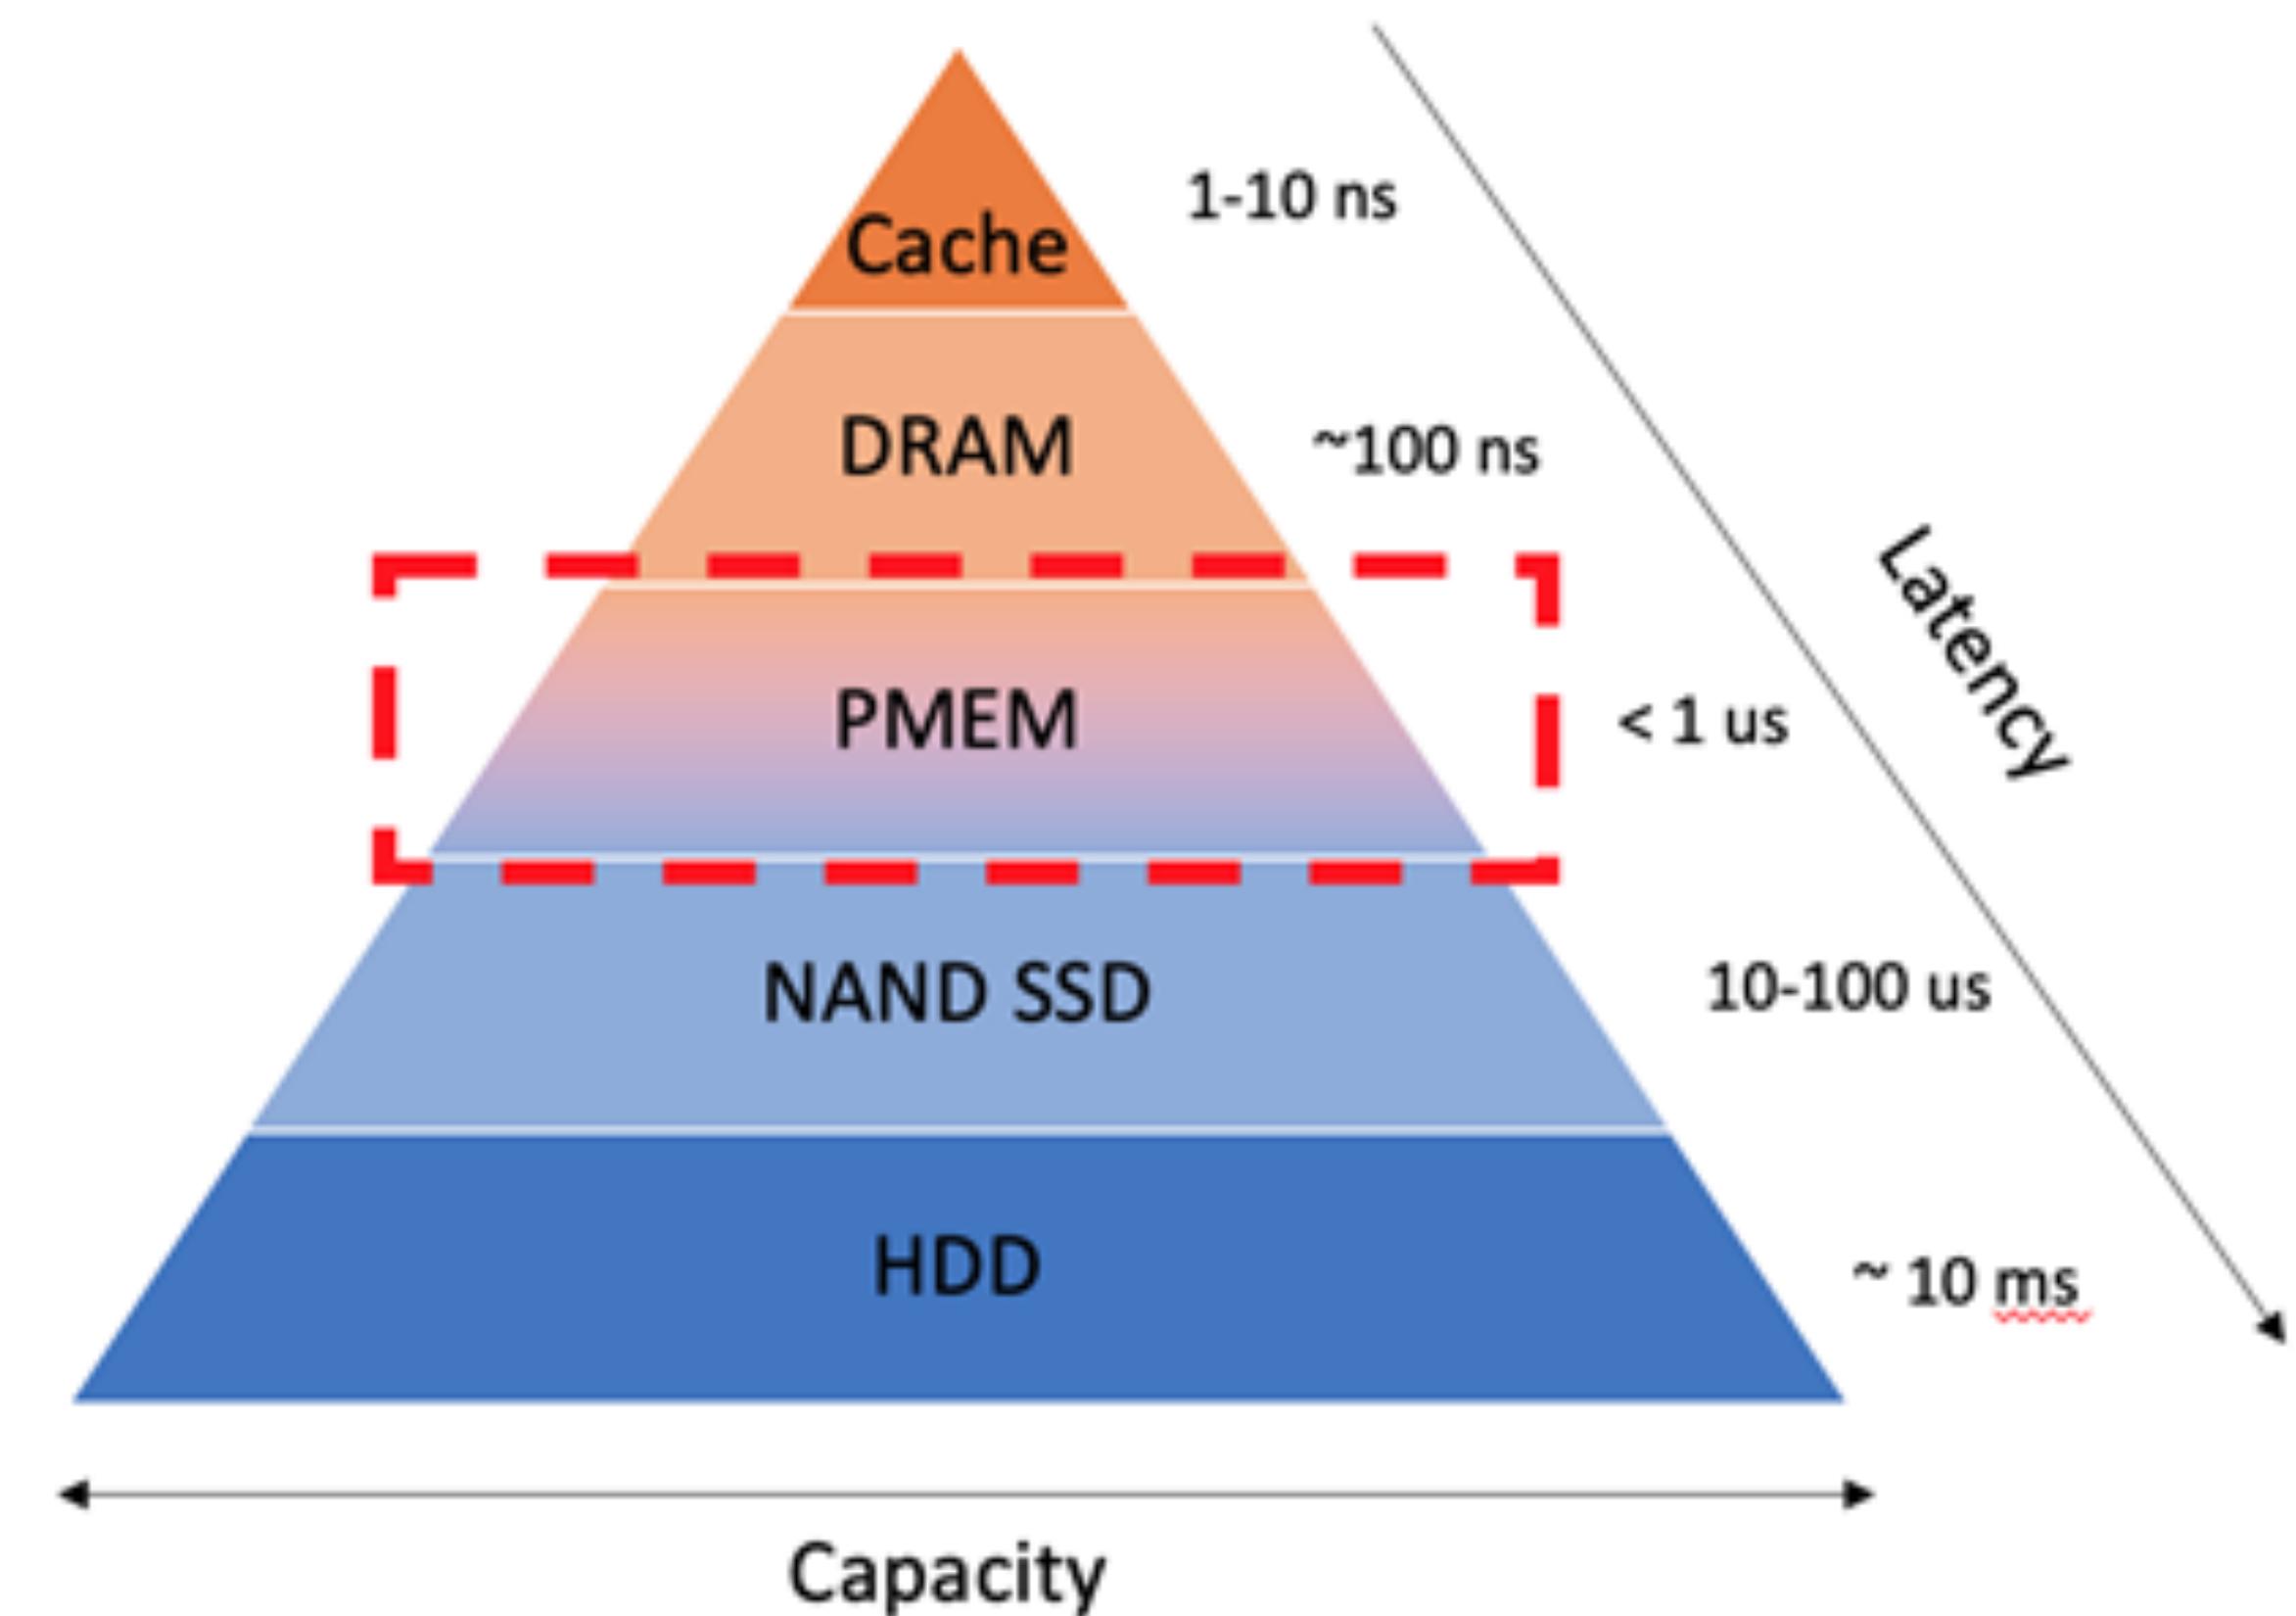


# More trends: persistent memory



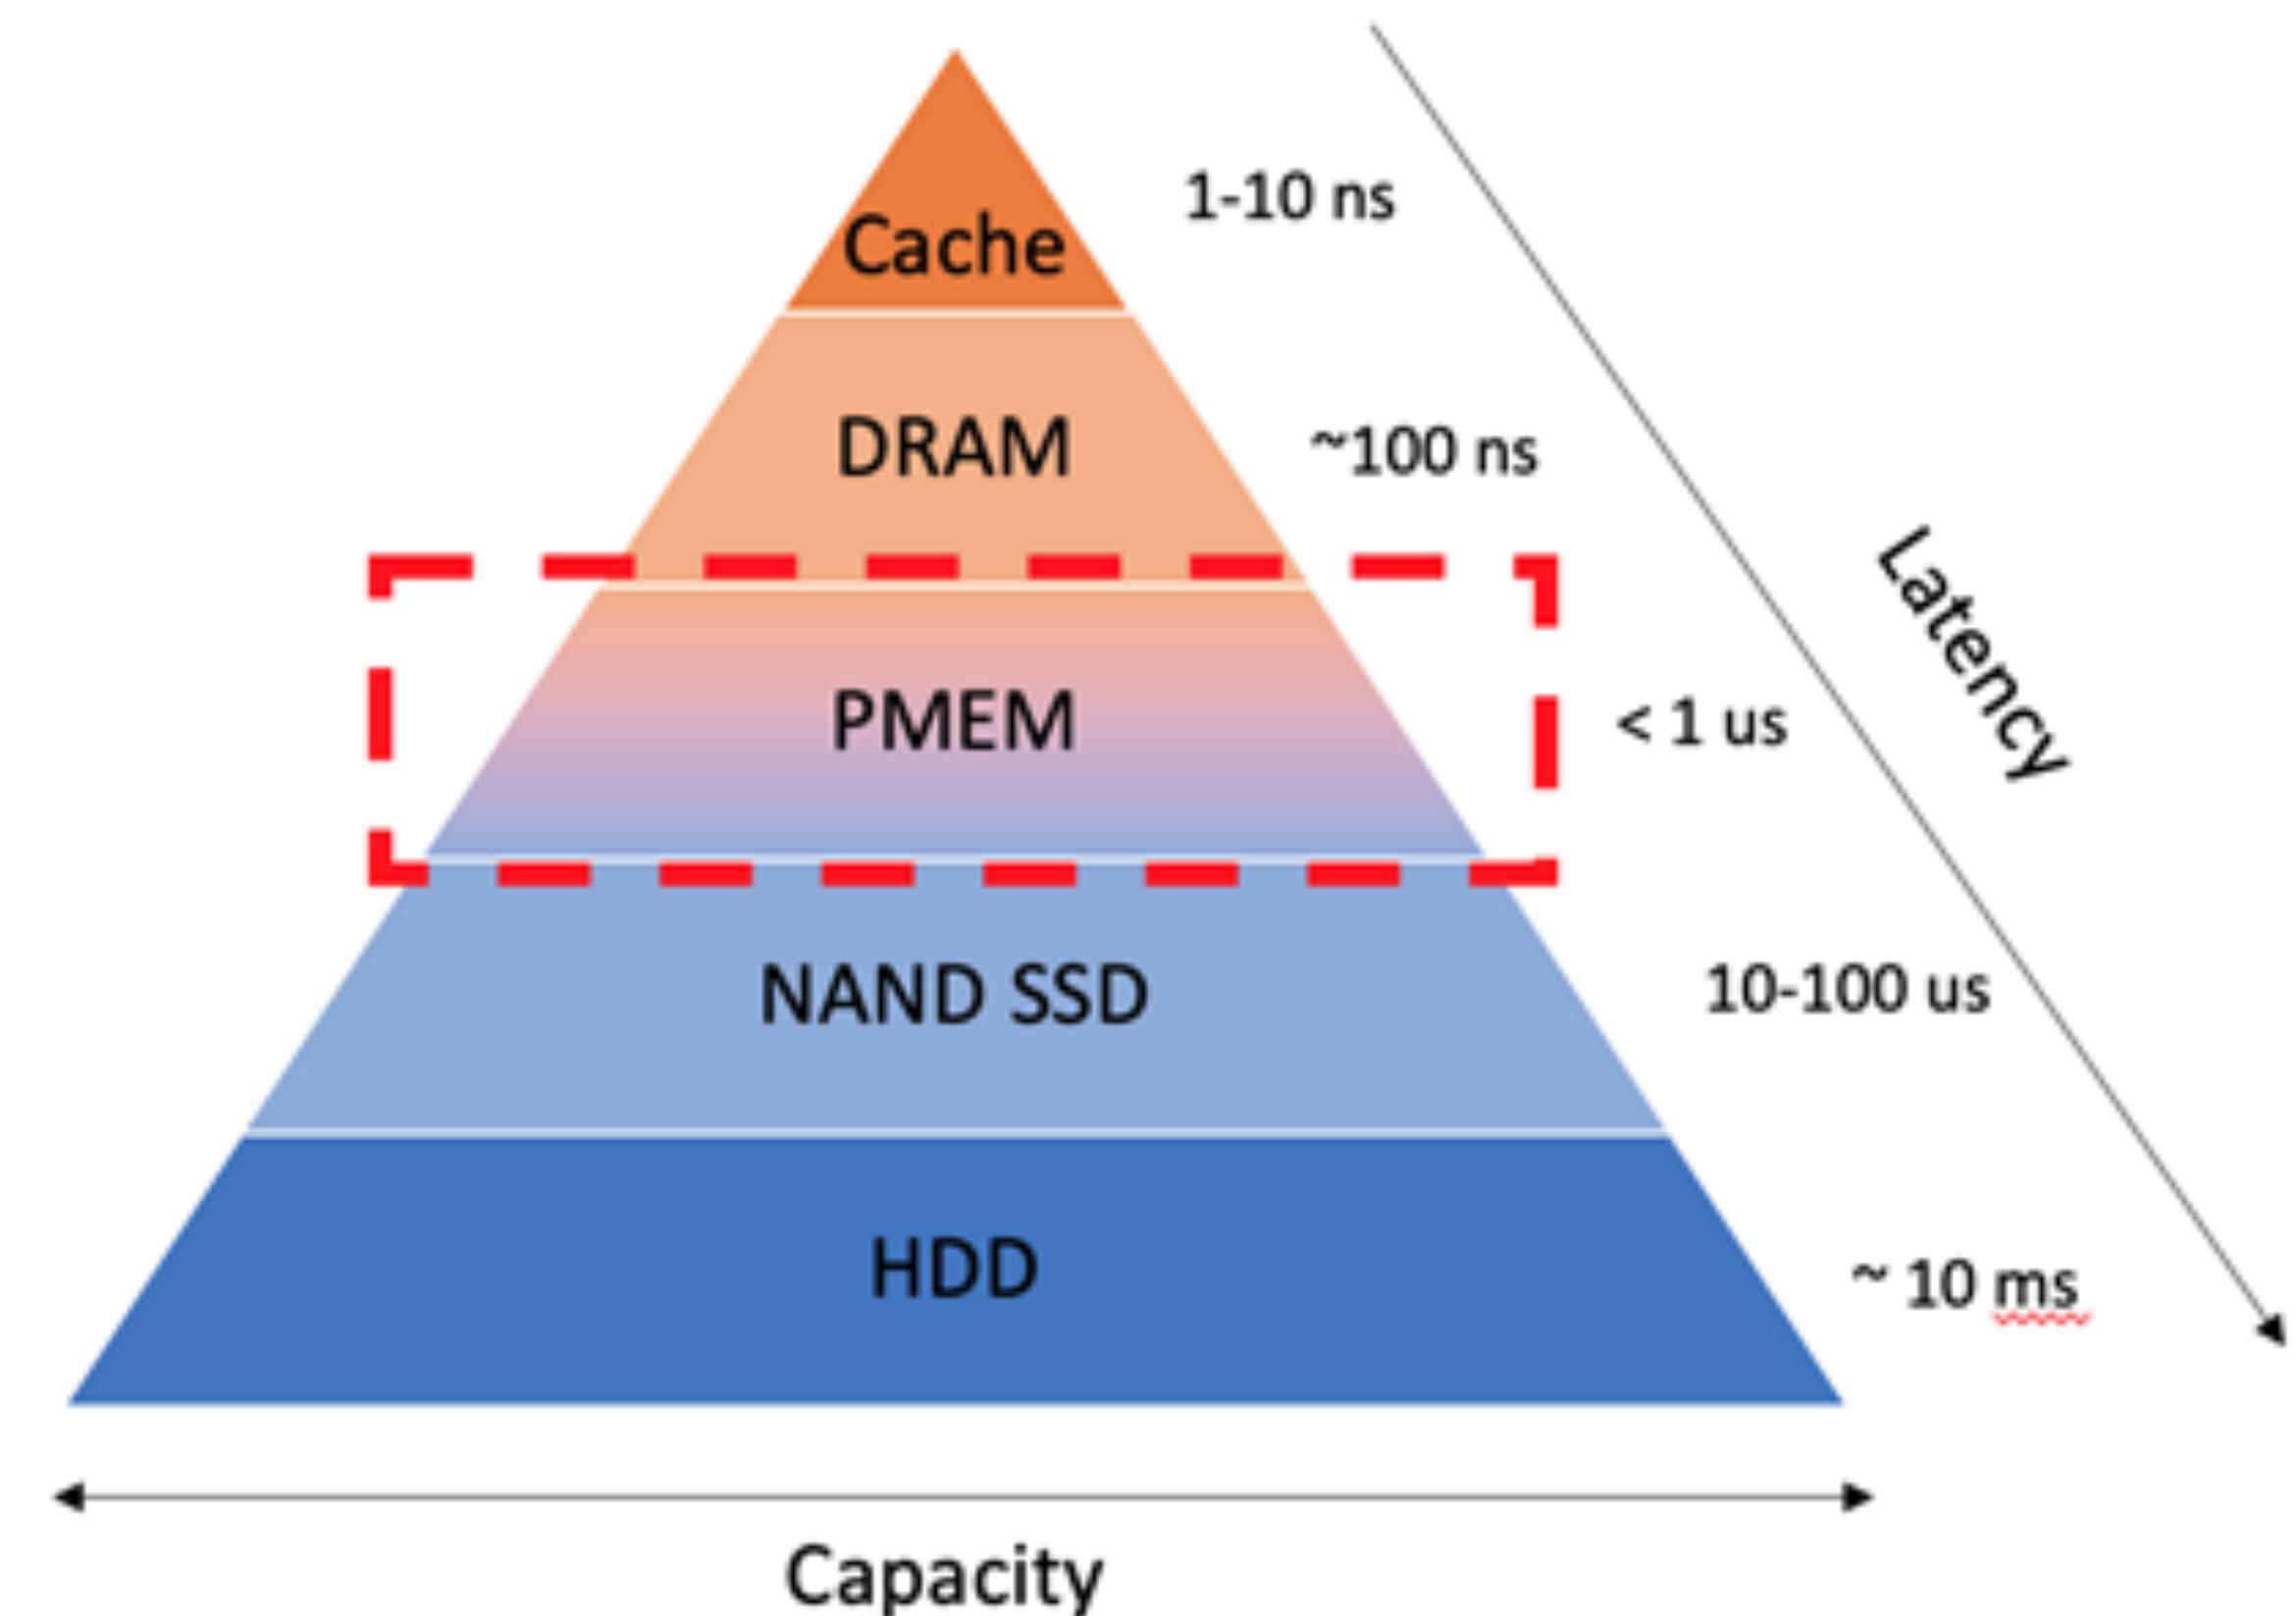
# More trends: persistent memory

- Memory: volatile but fast



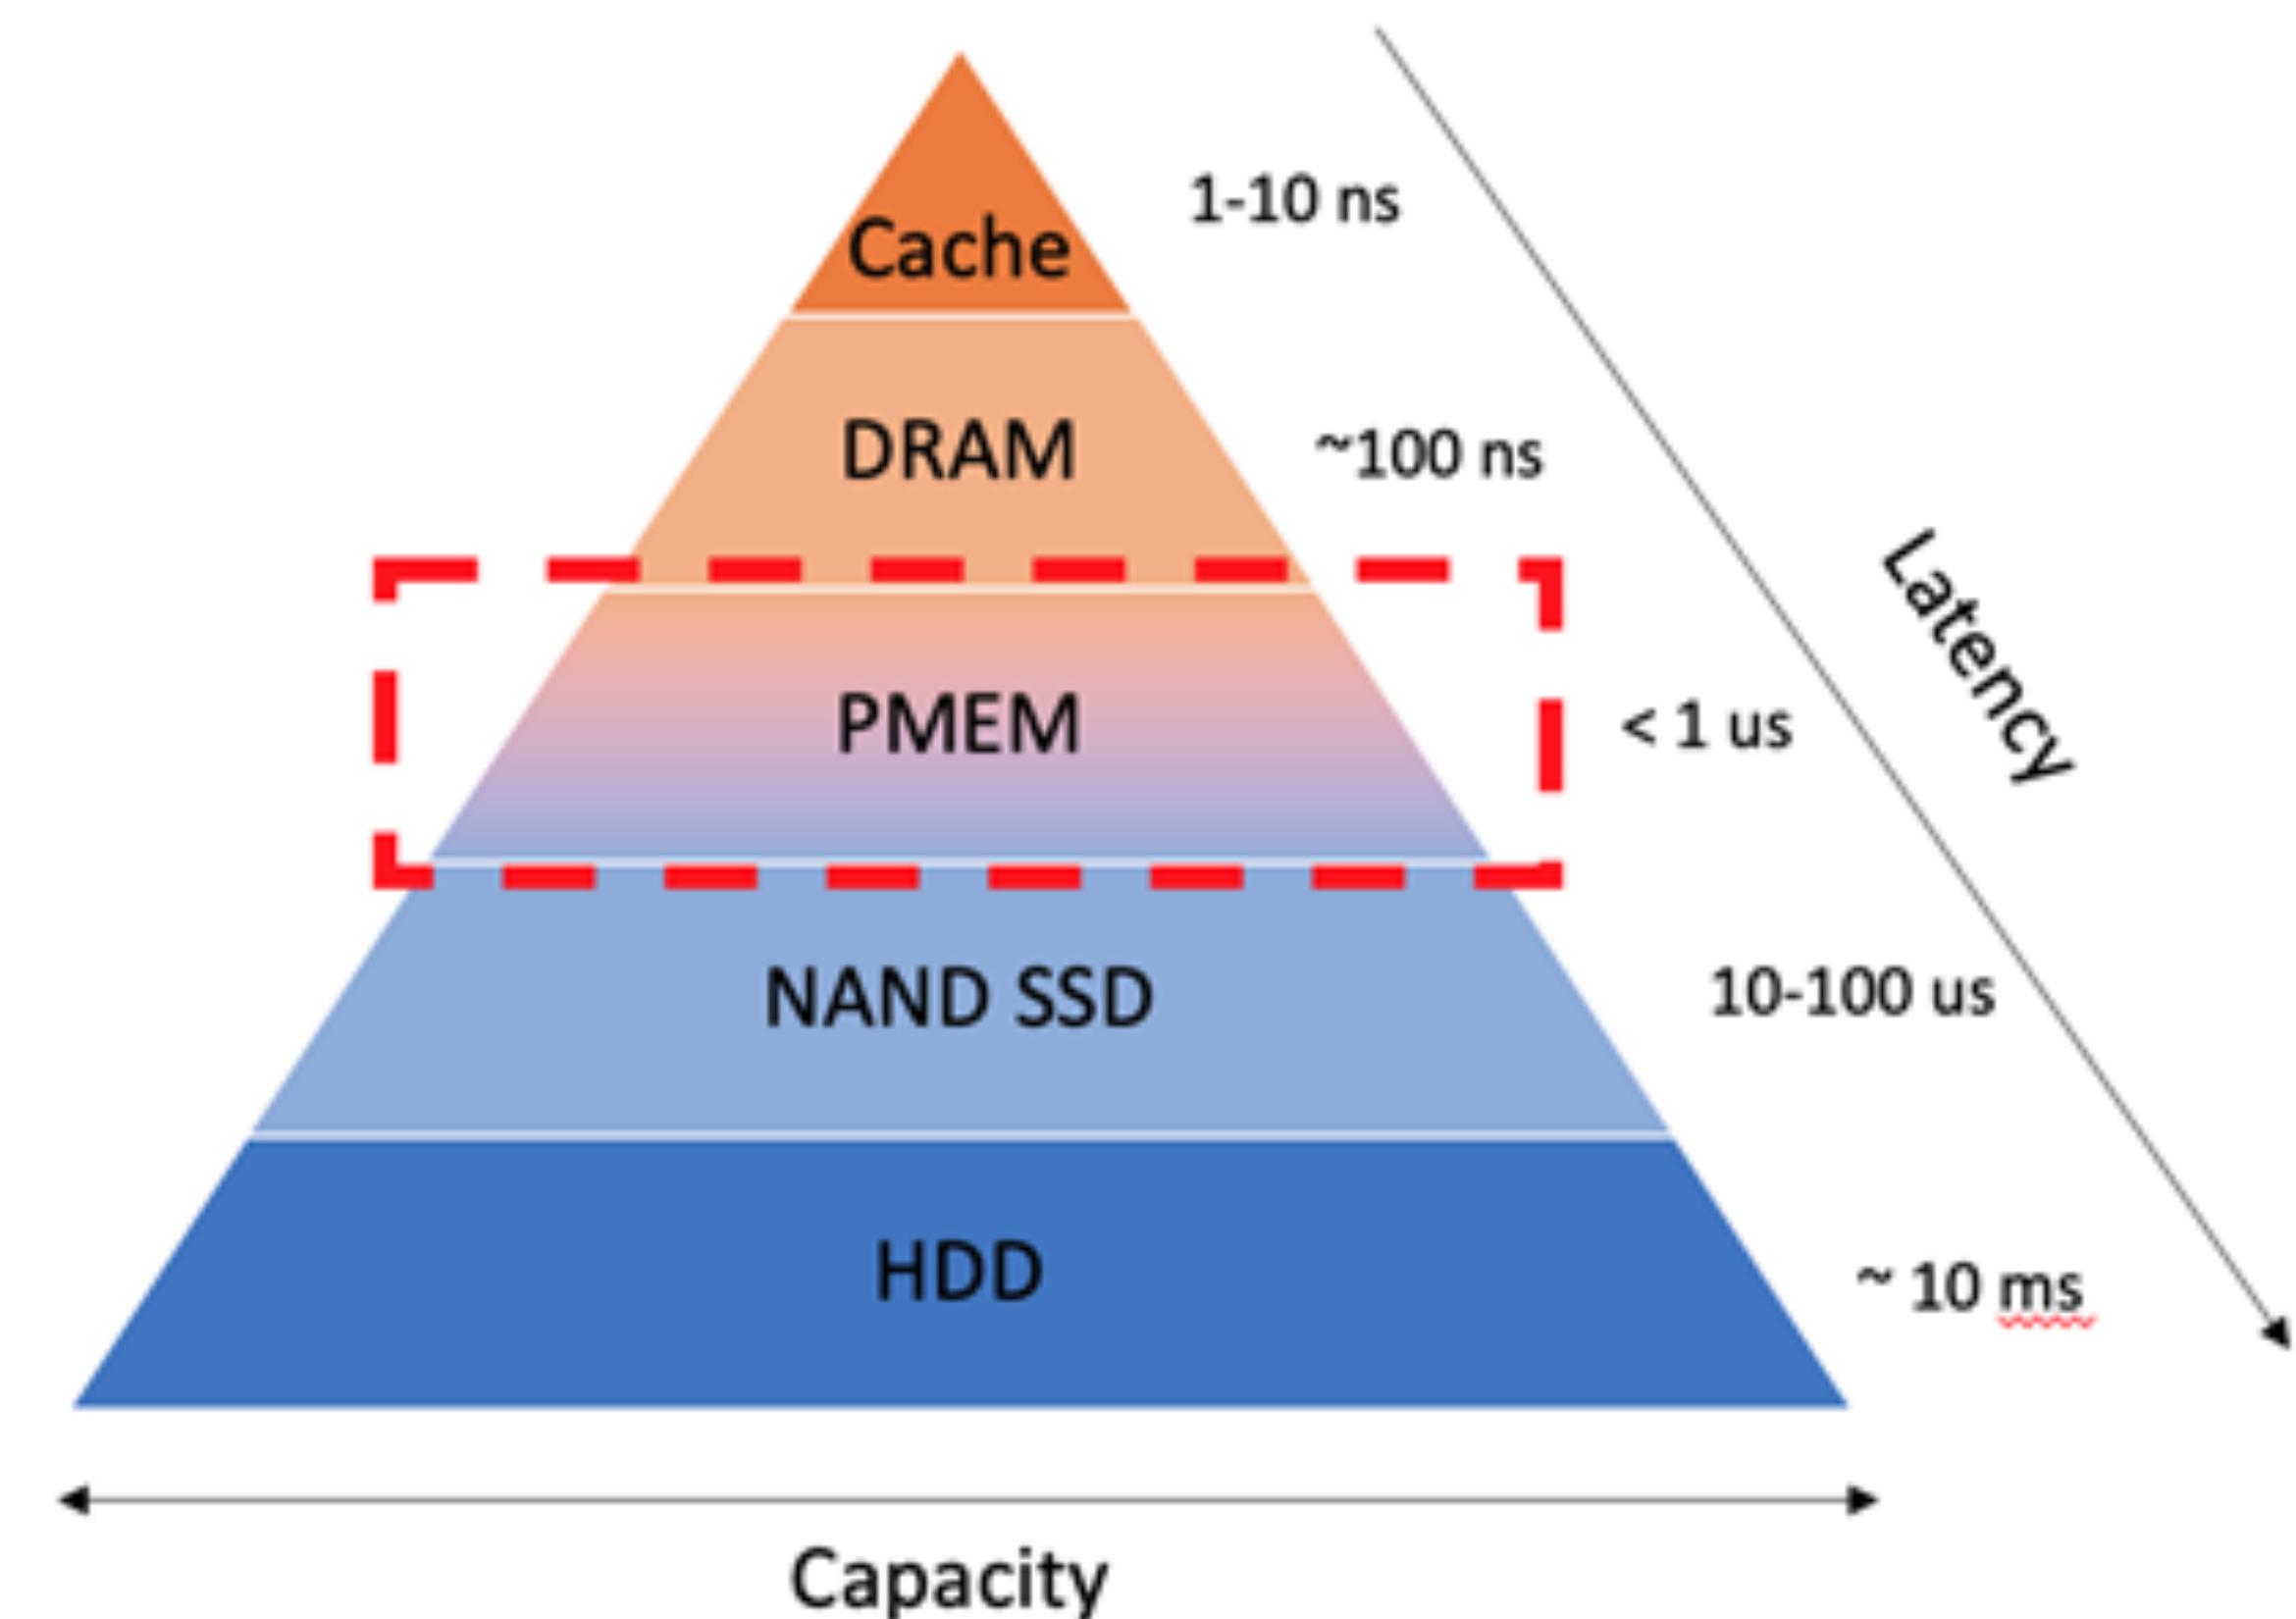
# More trends: persistent memory

- Memory: volatile but fast
- Disk: persistent but large



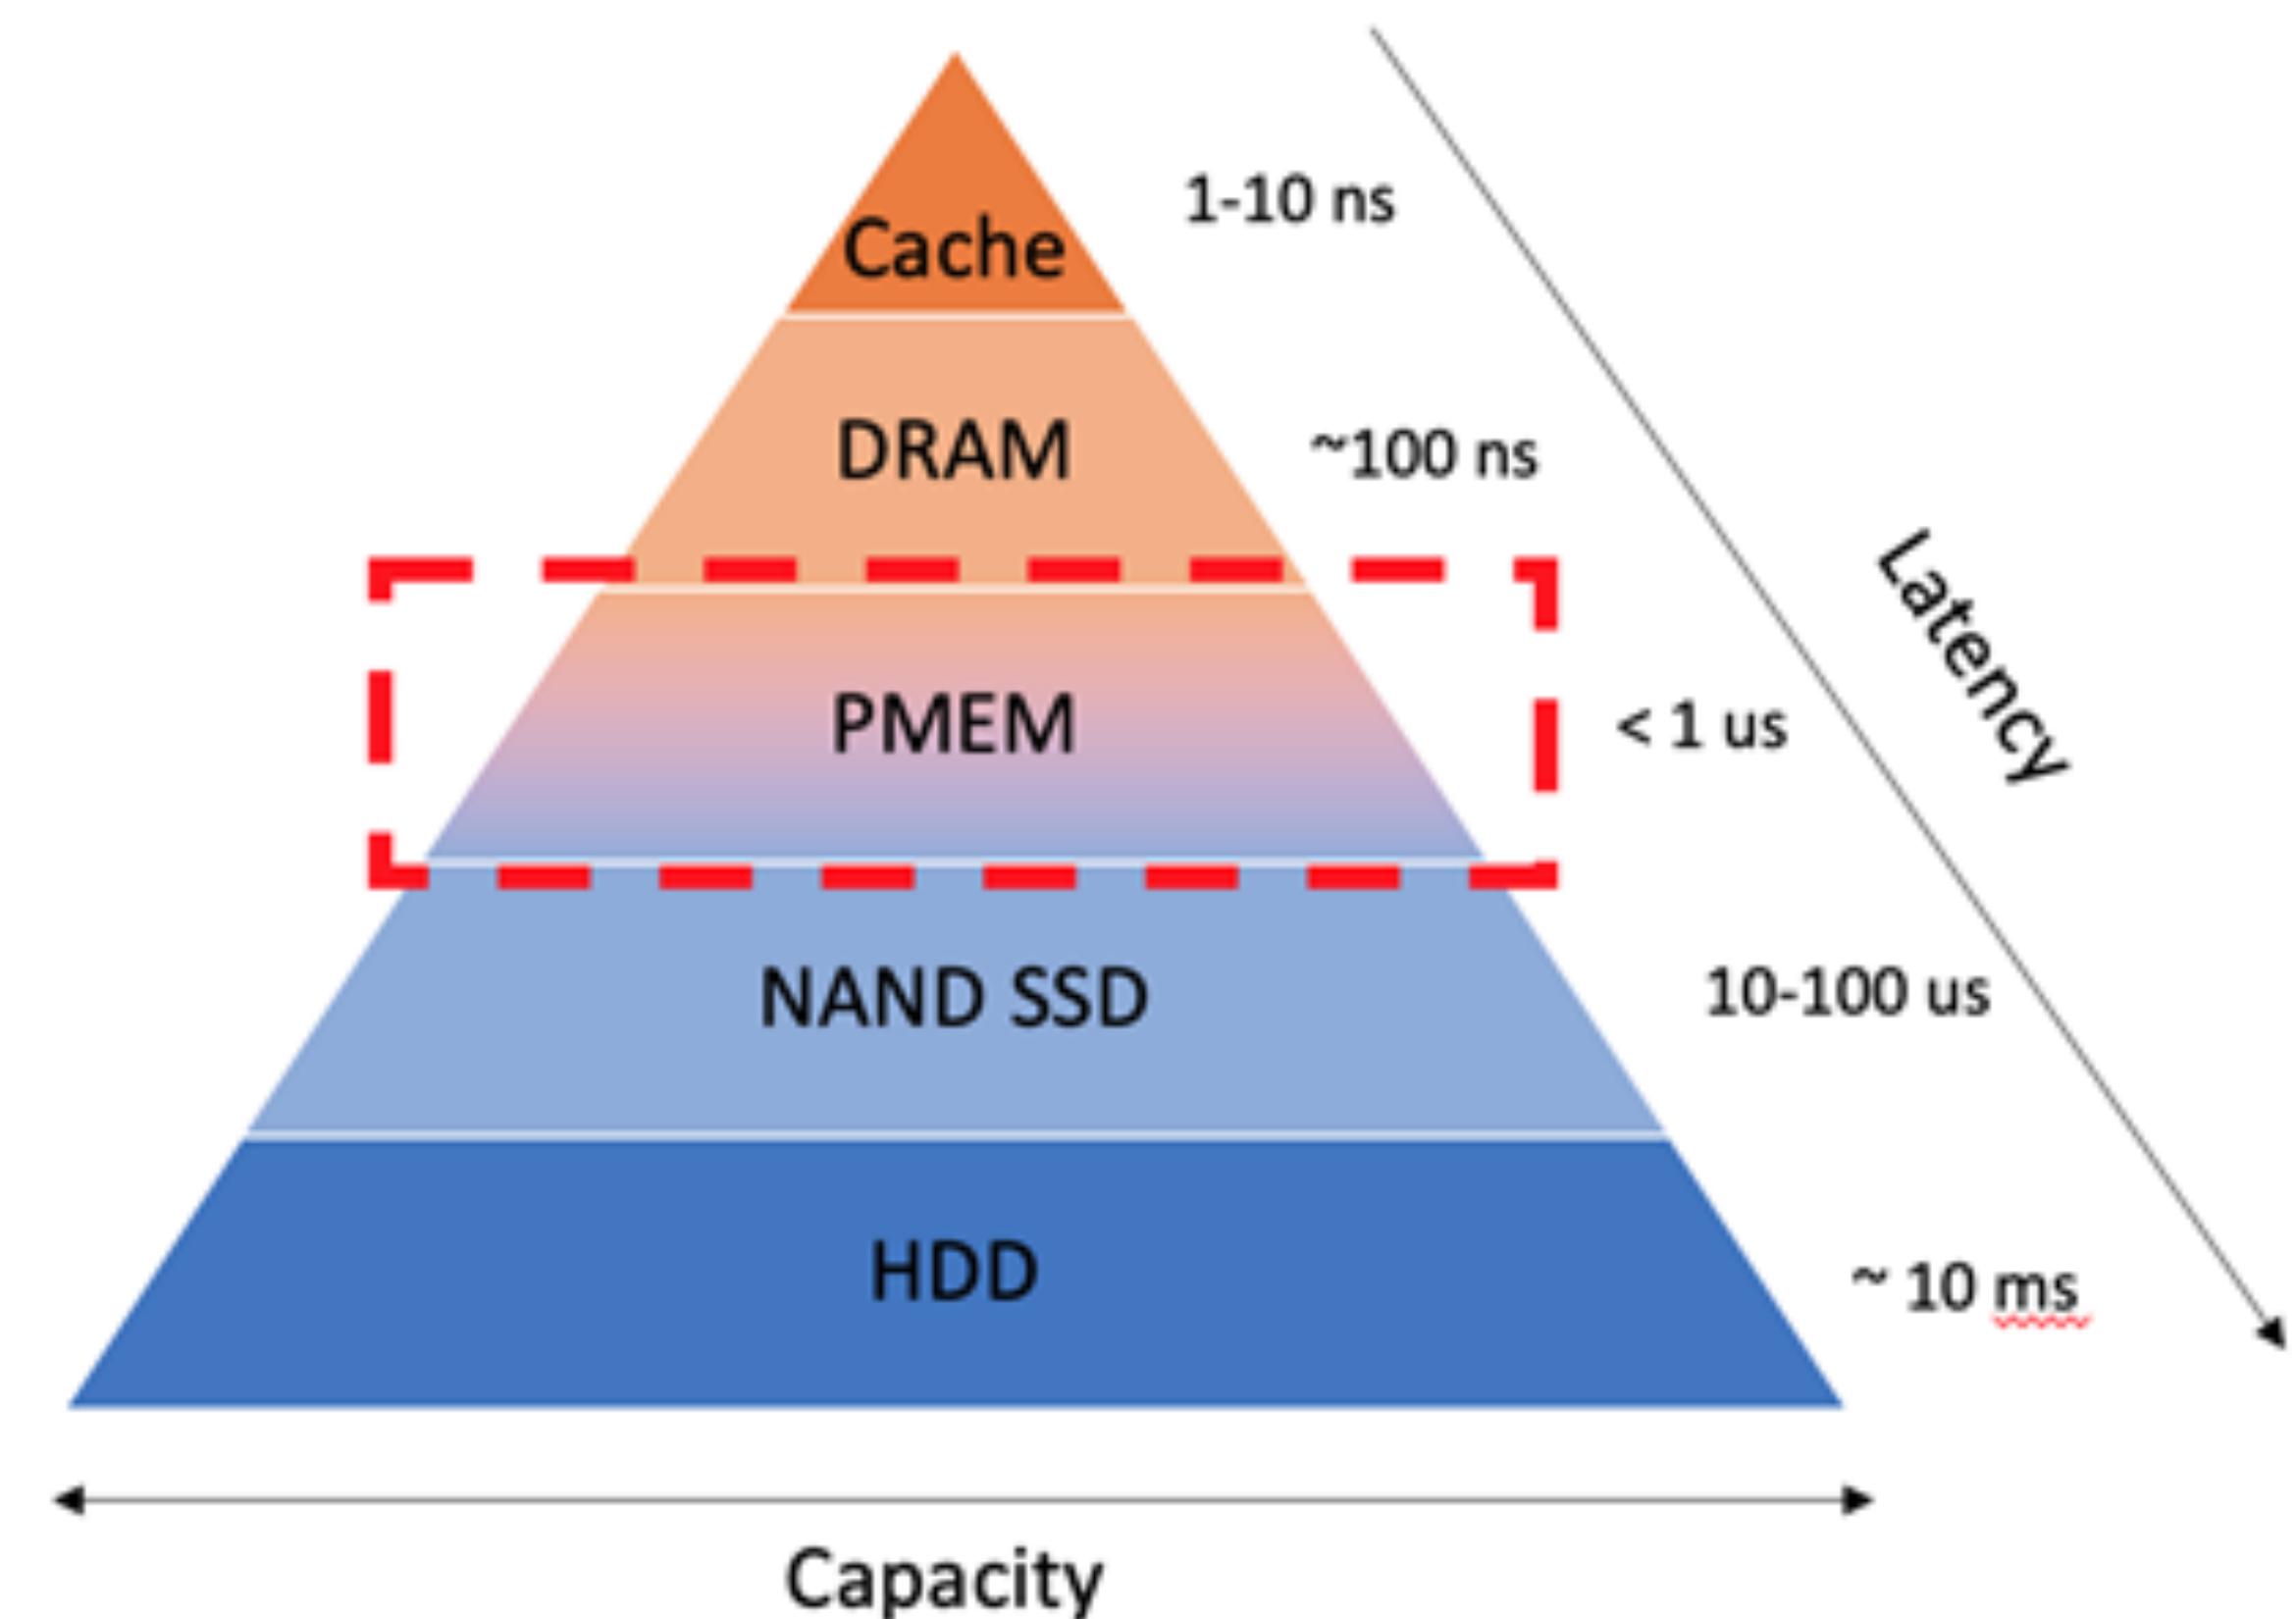
# More trends: persistent memory

- Memory: volatile but fast
- Disk: persistent but large
- PMEM: persistent and fast!



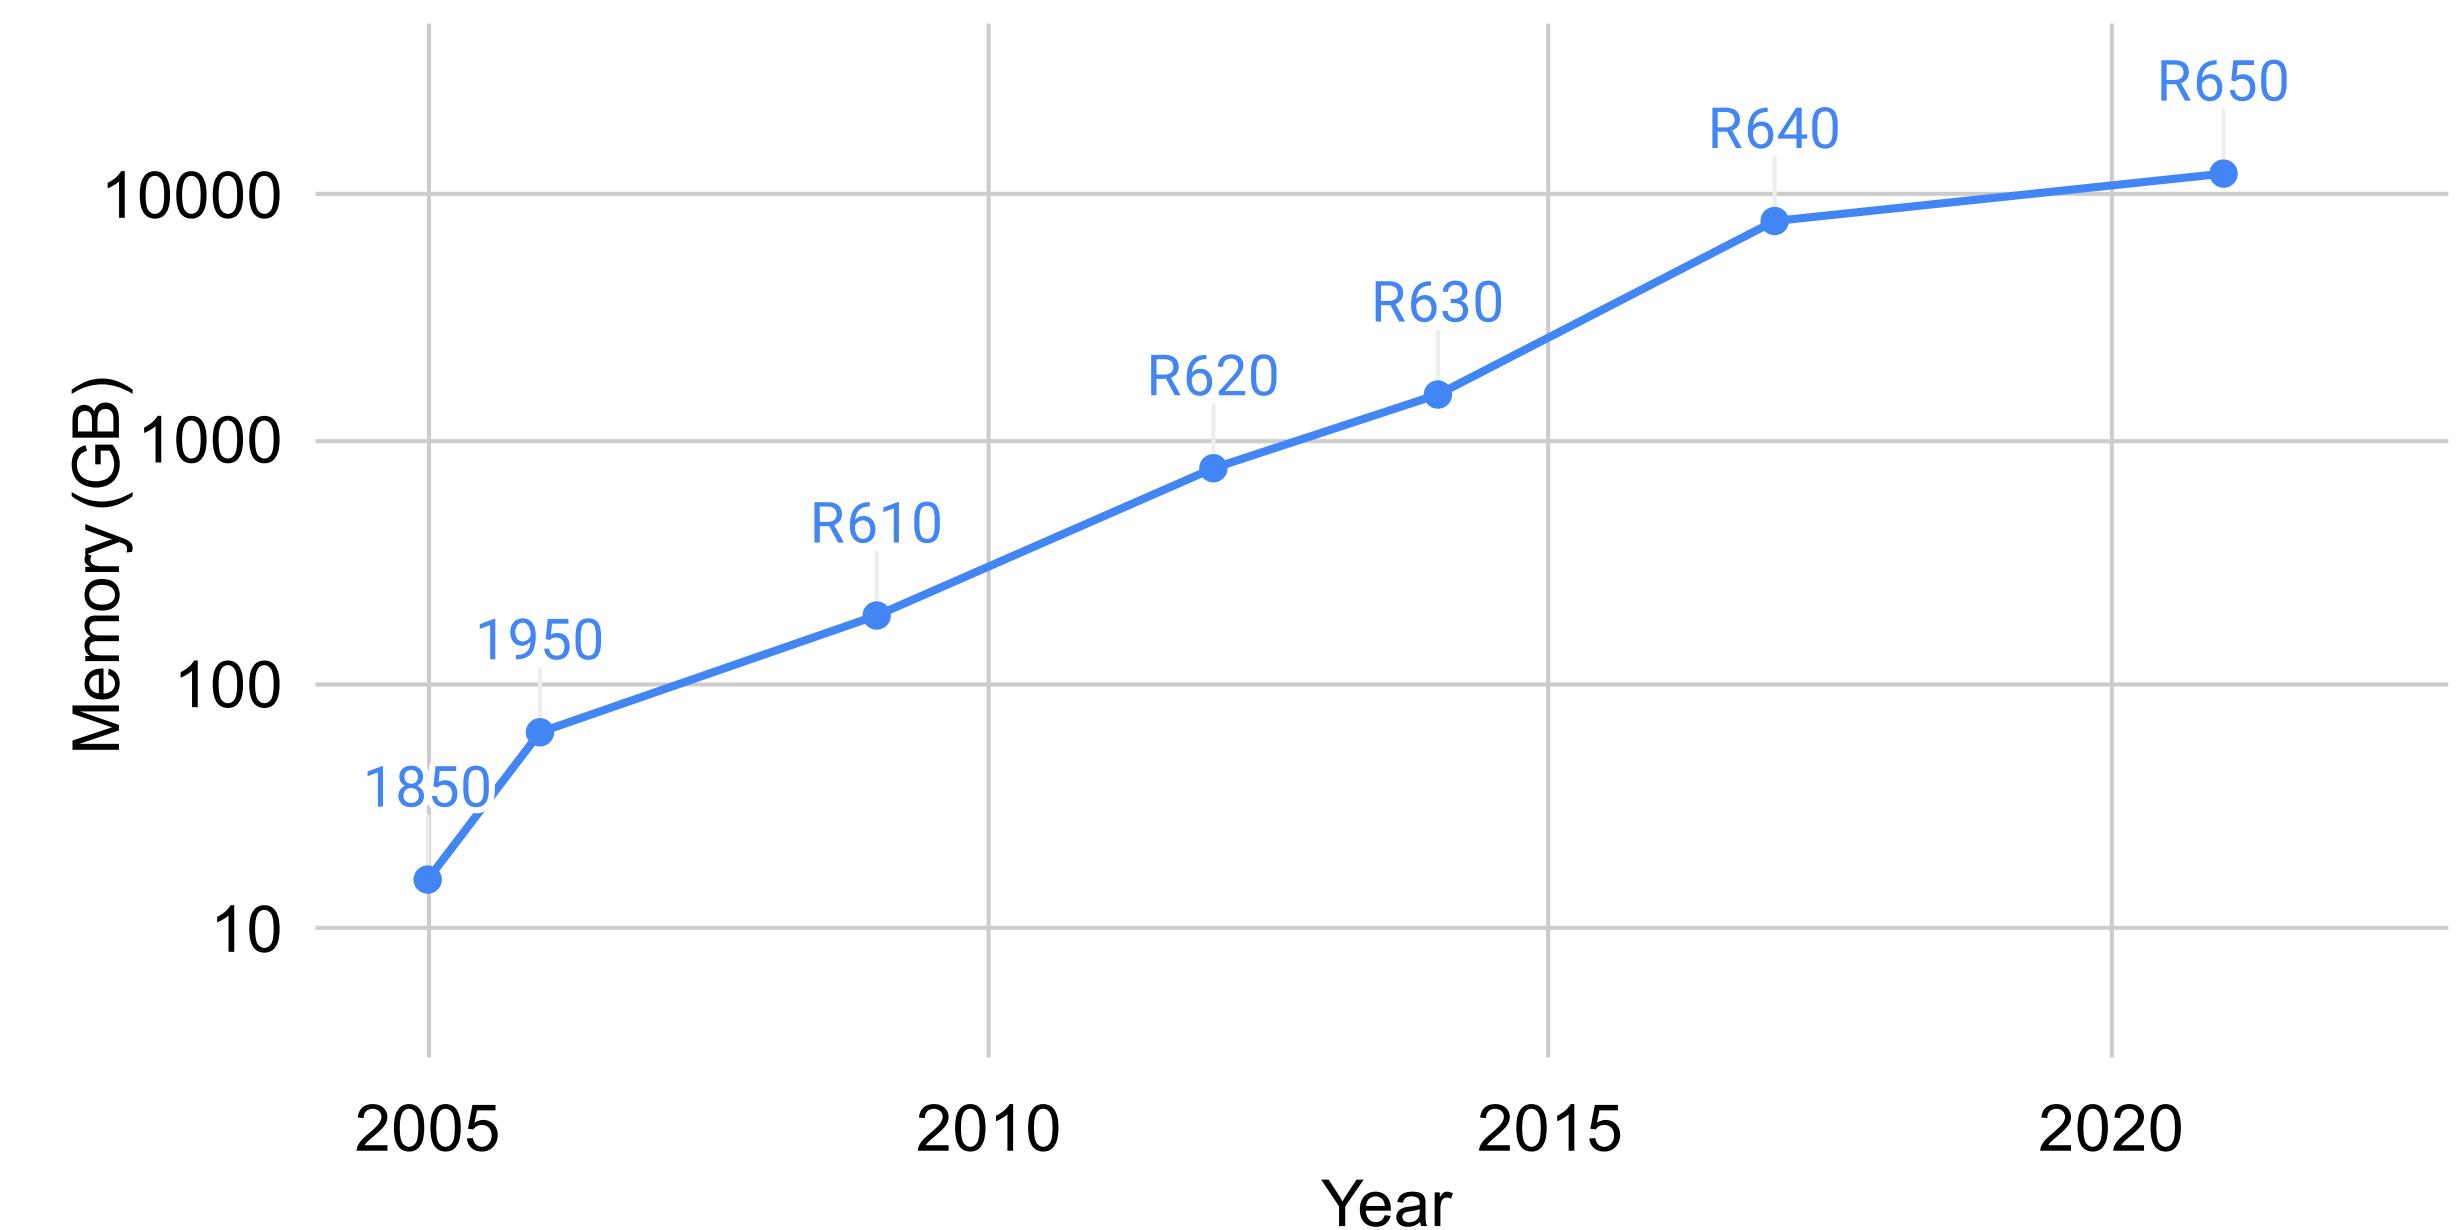
# More trends: persistent memory

- Memory: volatile but fast
- Disk: persistent but large
- PMEM: persistent and fast!
  - PMEM aware file systems



# More trends: big memory

- OS were designed when memory was scarce: few KBs
- You can now buy a server with 12TB of DRAM!
- Transparent huge pages
- Caching: Redis/memcached, etc.
- In memory compute: Spark, etc

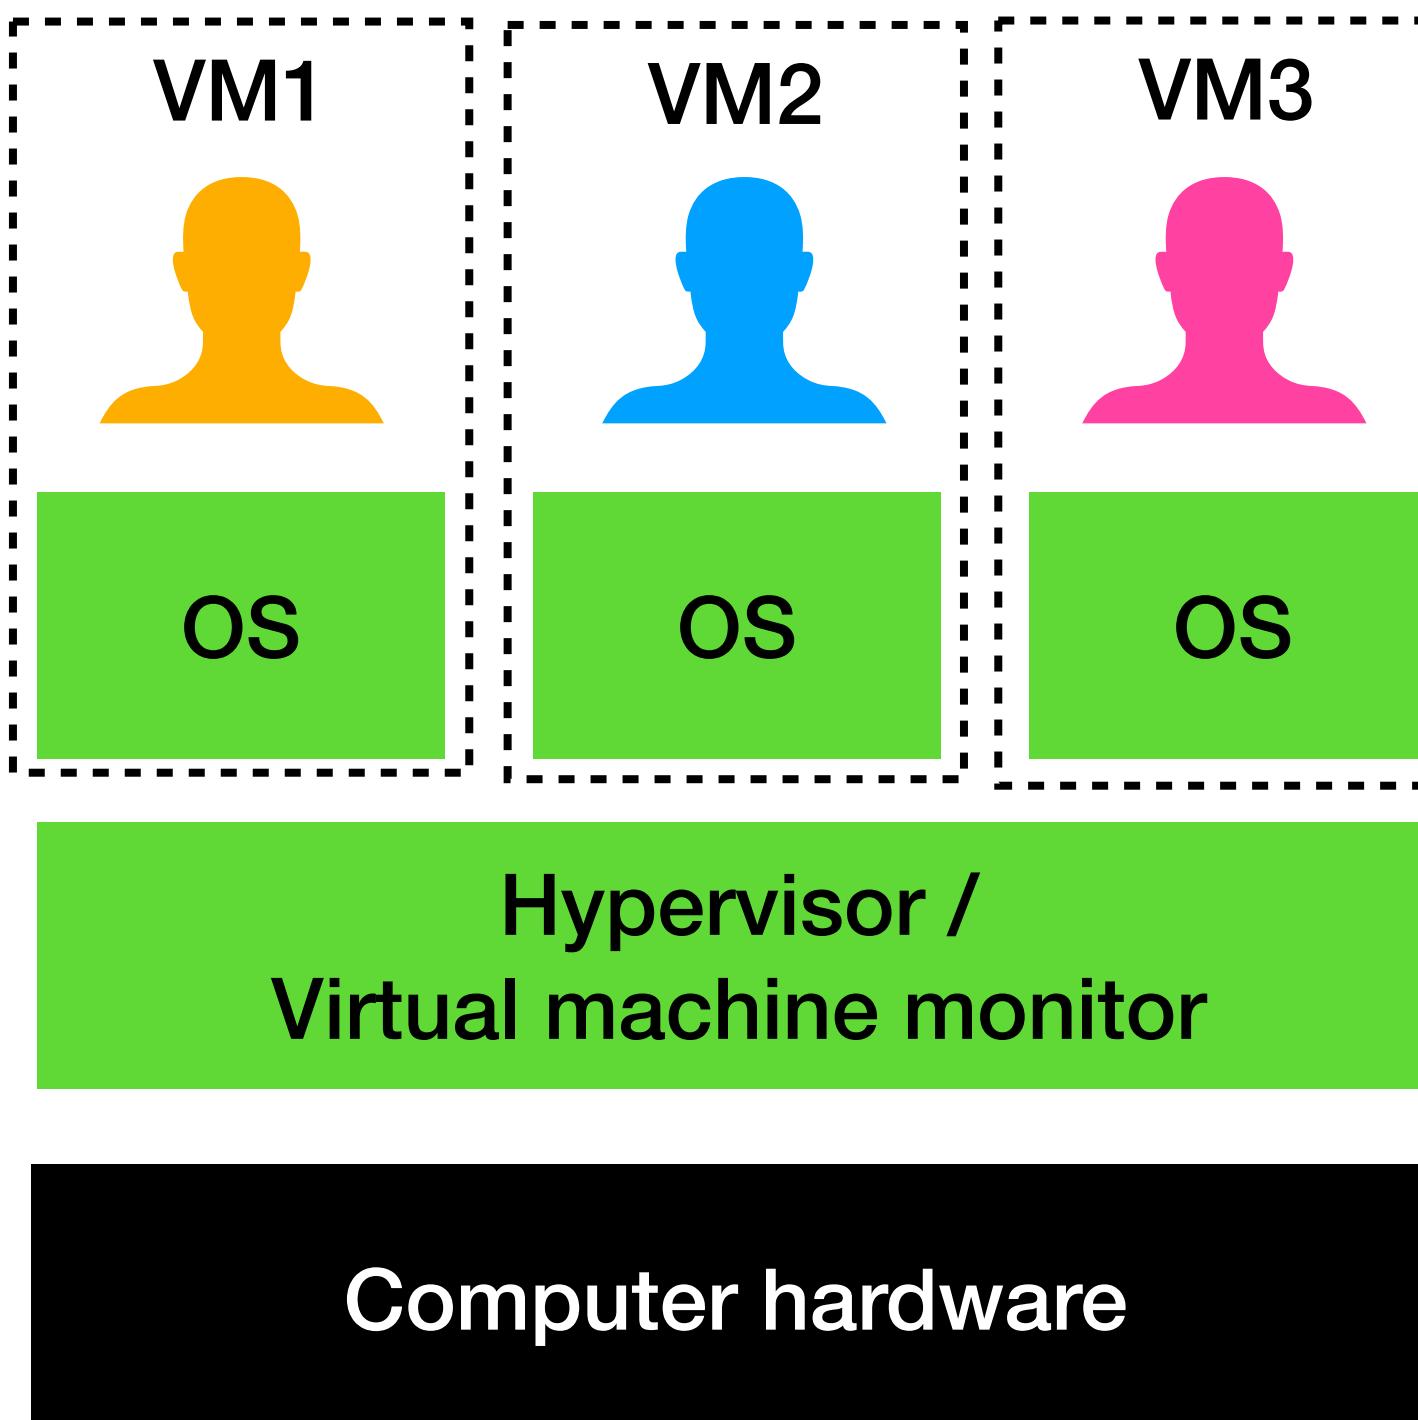


# More trends: fast networks

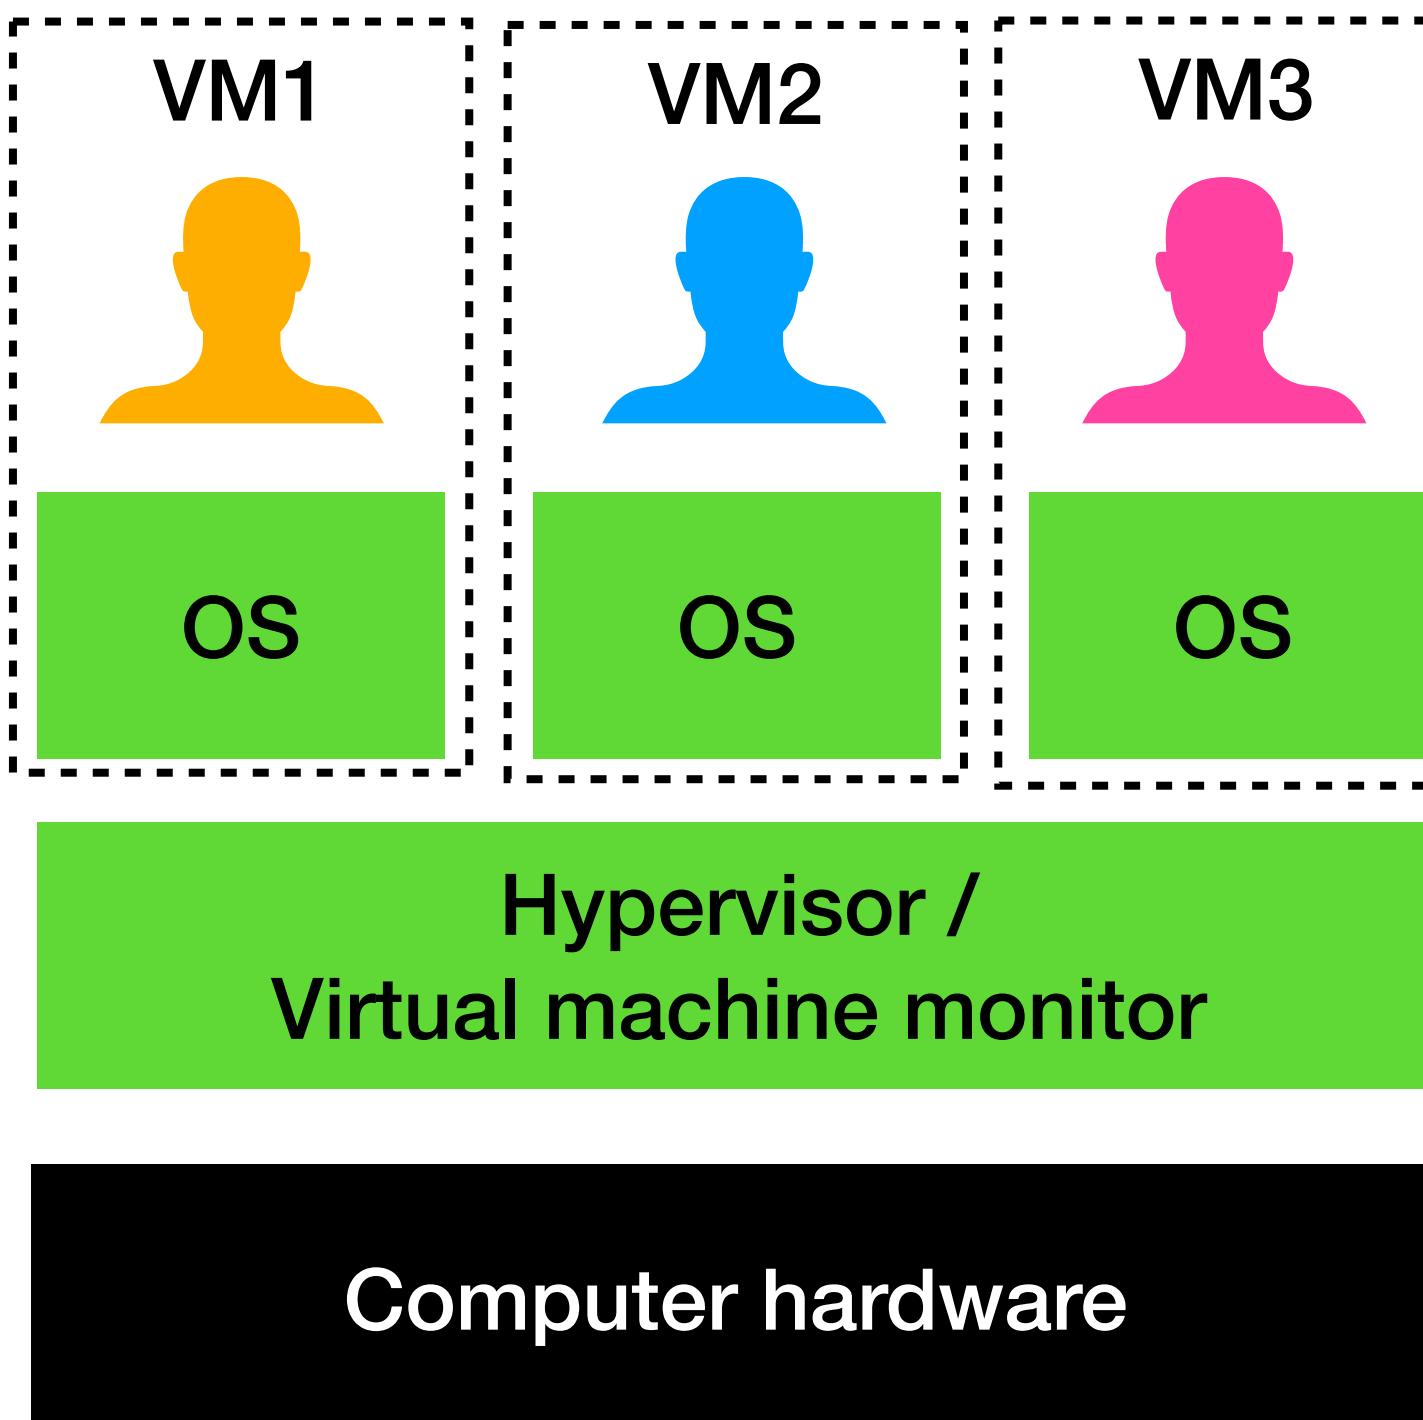
- OS typically assumed network is *much slower* than DRAM
  - Far memory

2020s	Latency	Bandwidth
DRAM	15ns	400 GBps
Ethernet	500ns	50 GBps

# More trends: Rise of Unikernels

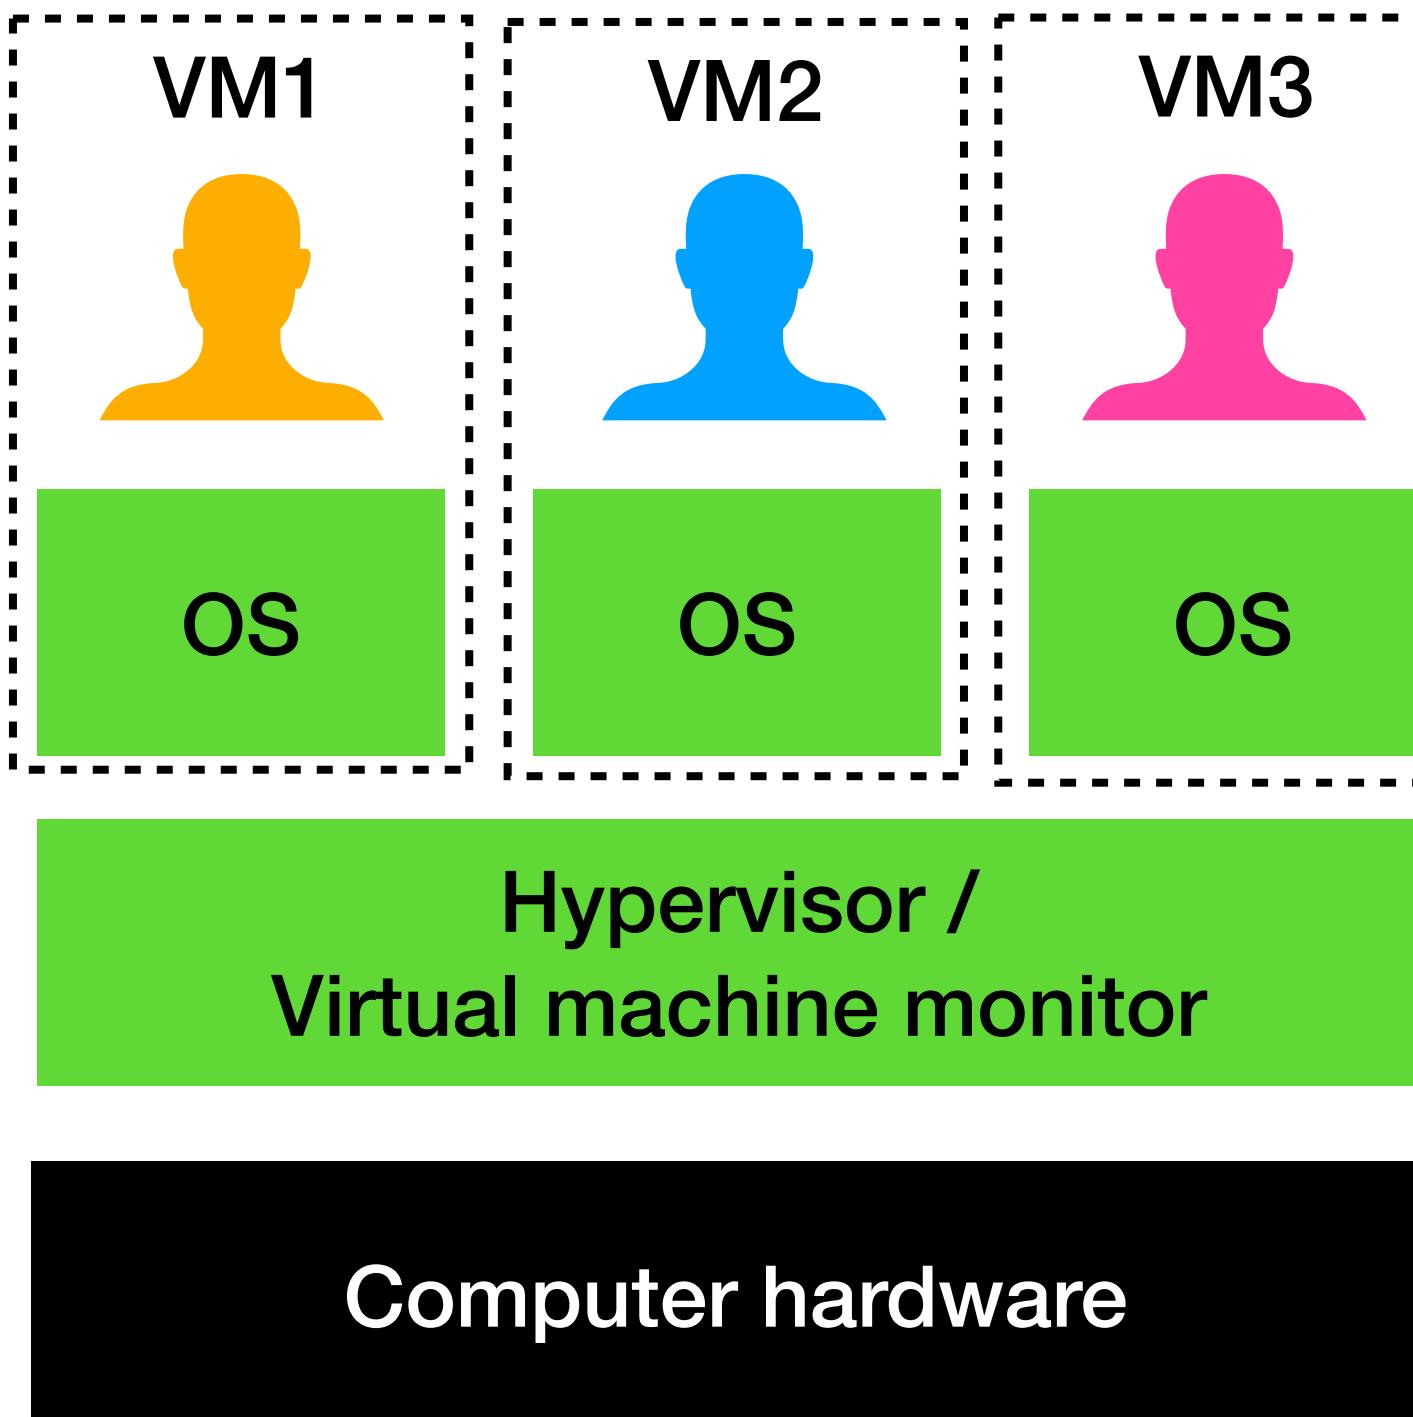


# More trends: Rise of Unikernels



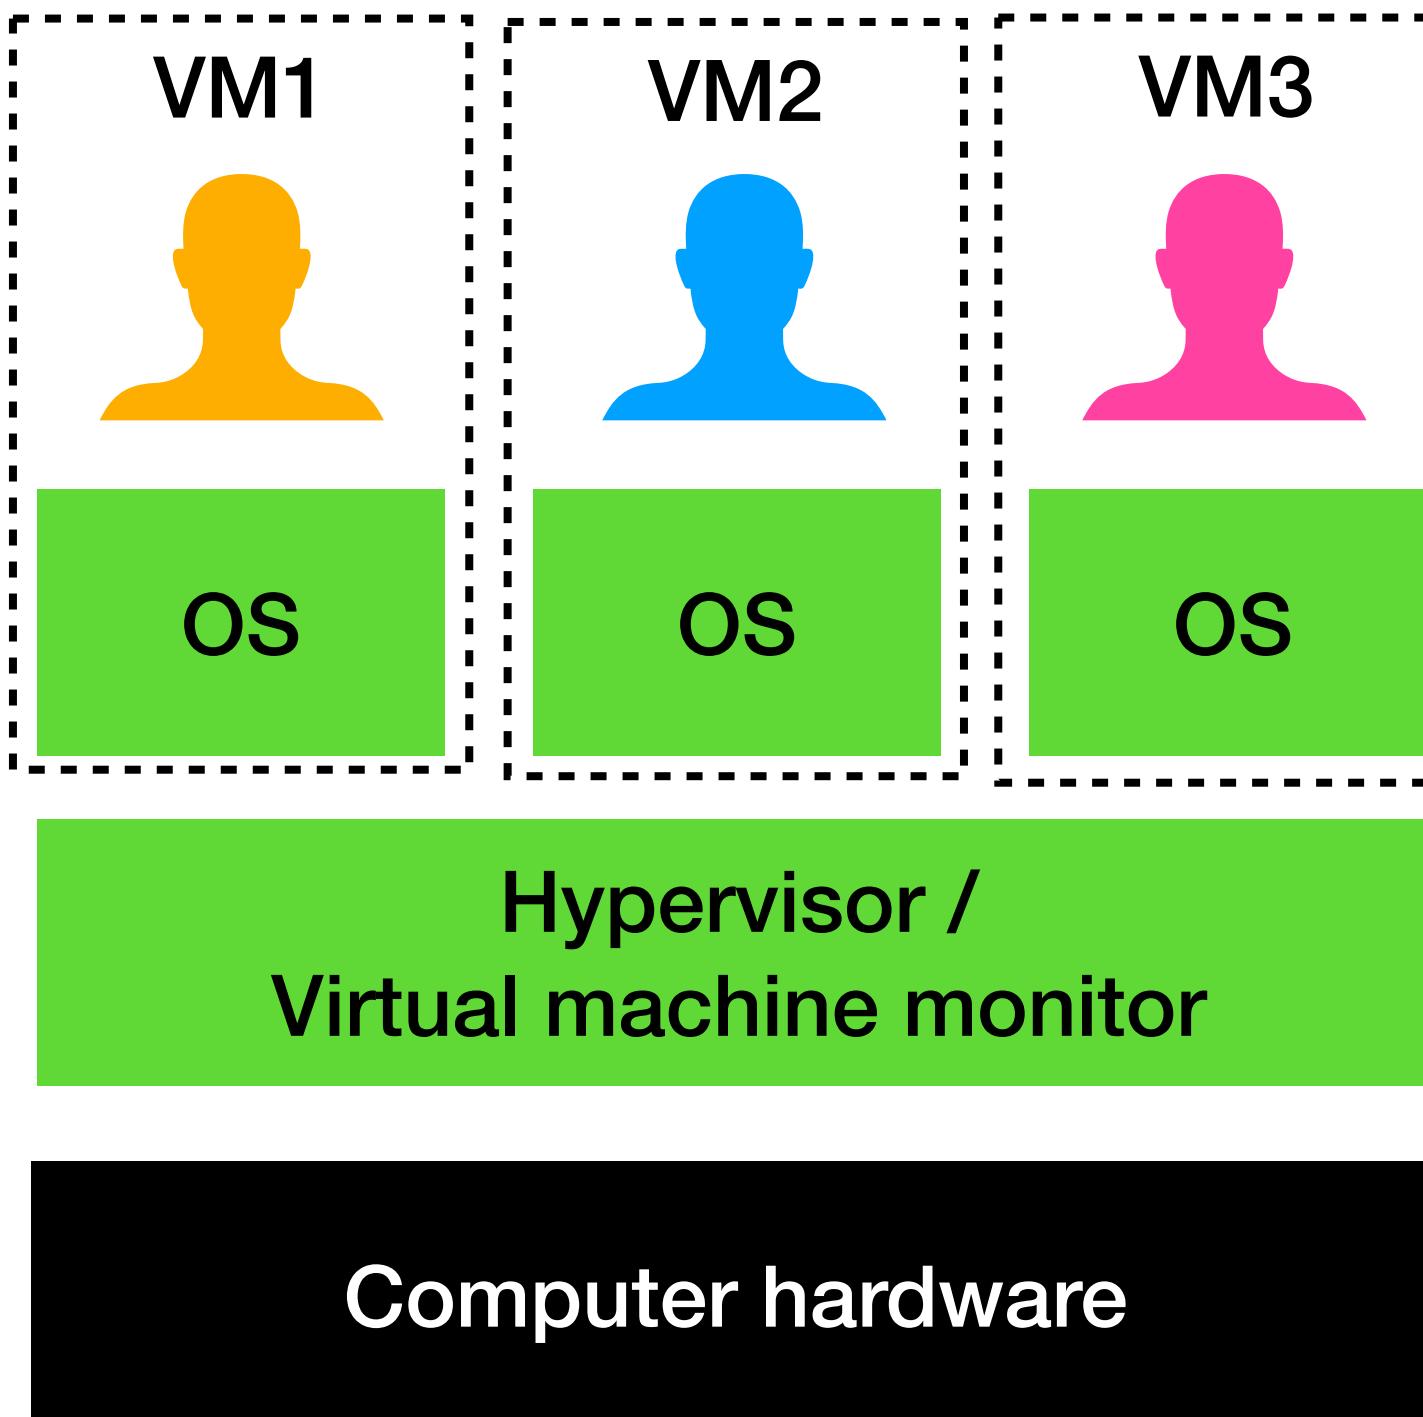
- OS optimises for common behaviours across all applications

# More trends: Rise of Unikernels



- OS optimises for common behaviours across all applications
- Each OS is now running only single application

# More trends: Rise of Unikernels

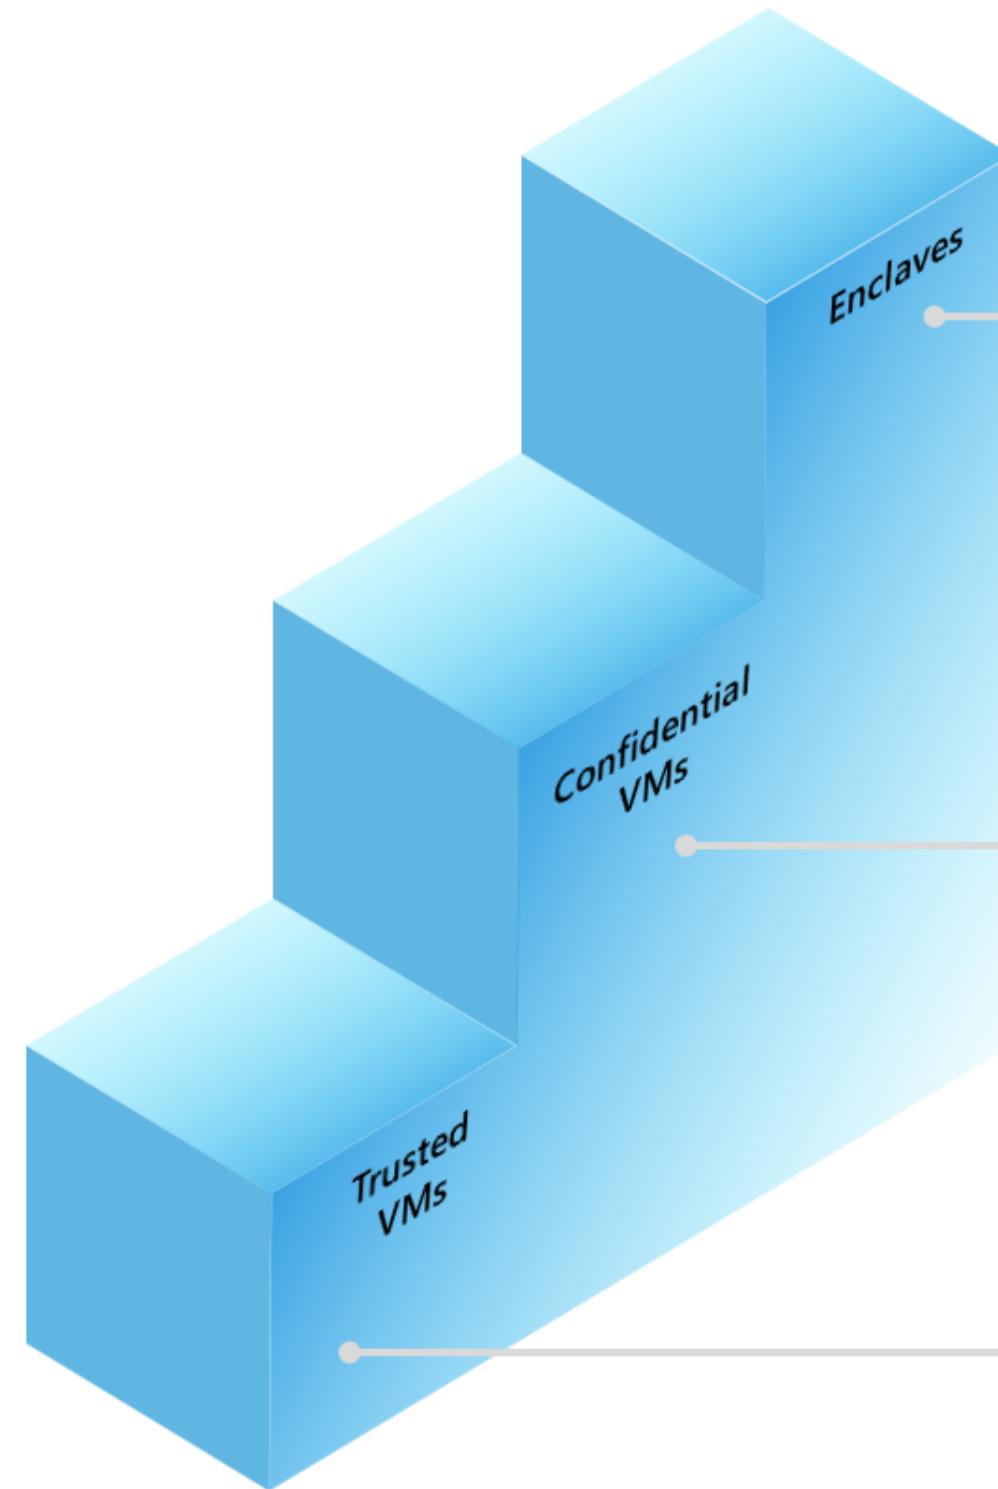


- OS optimises for common behaviours across all applications
- Each OS is now running only single application
- Unikernels optimise only for a single application

# Confidential computing

- Developers want to run their proprietary code/models on sensitive data on someone else's machine (cloud) where all software including (host) OS is untrusted

## Trusted Execution Environments (TEEs)



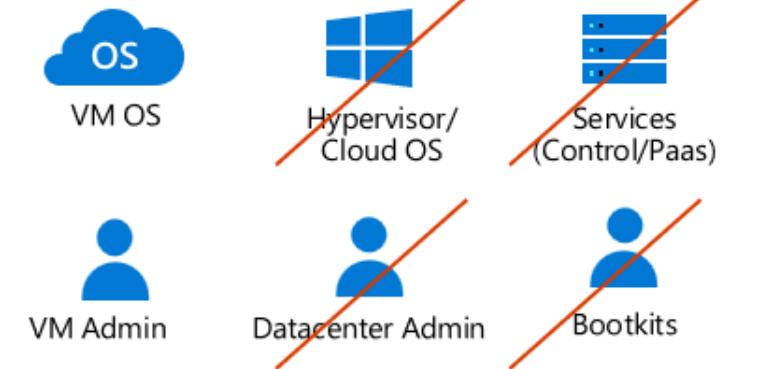
### Hardware enclaves with Intel SGX

 "I just trust my app code and the chip."



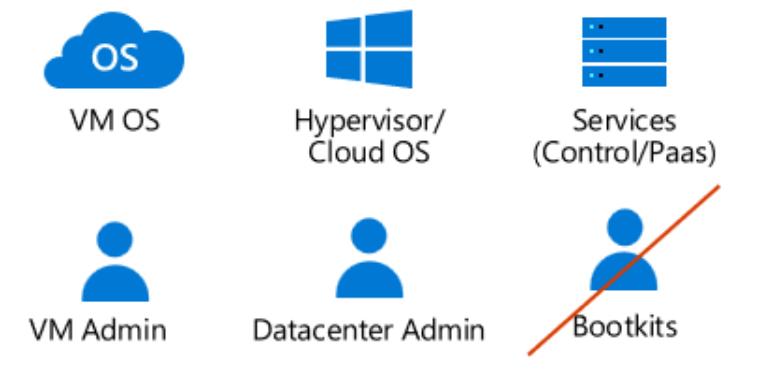
### Hardware Confidential VMs with AMD SEV-SNP & Intel TDX

 "Microsoft cannot touch my stuff in my VM."



### Trusted launch VMs

 "Only known, trusted code is running on my VM."



Trust

# More trends: ML/LLM Systems

# More trends: ML/LLM Systems

- ML training at scale
  - Pathways/TensorFlow  
(Google), Ray (Anyscale),  
...

# More trends: ML/LLM Systems

- ML training at scale
  - Pathways/TensorFlow (Google), Ray (Anyscale),  
...
- LLM inference at scale
  - VLLM, SGLang, etc.

# More trends: ML/LLM Systems

- ML training at scale
  - Pathways/TensorFlow (Google), Ray (Anyscale),  
...
- LLM inference at scale
  - VLLM, SGLang, etc.
- Agentic programs

# More trends: ML/LLM Systems

- ML training at scale
  - Pathways/TensorFlow (Google), Ray (Anyscale), ...
- LLM inference at scale
  - VLLM, SGLang, etc.
- Agentic programs

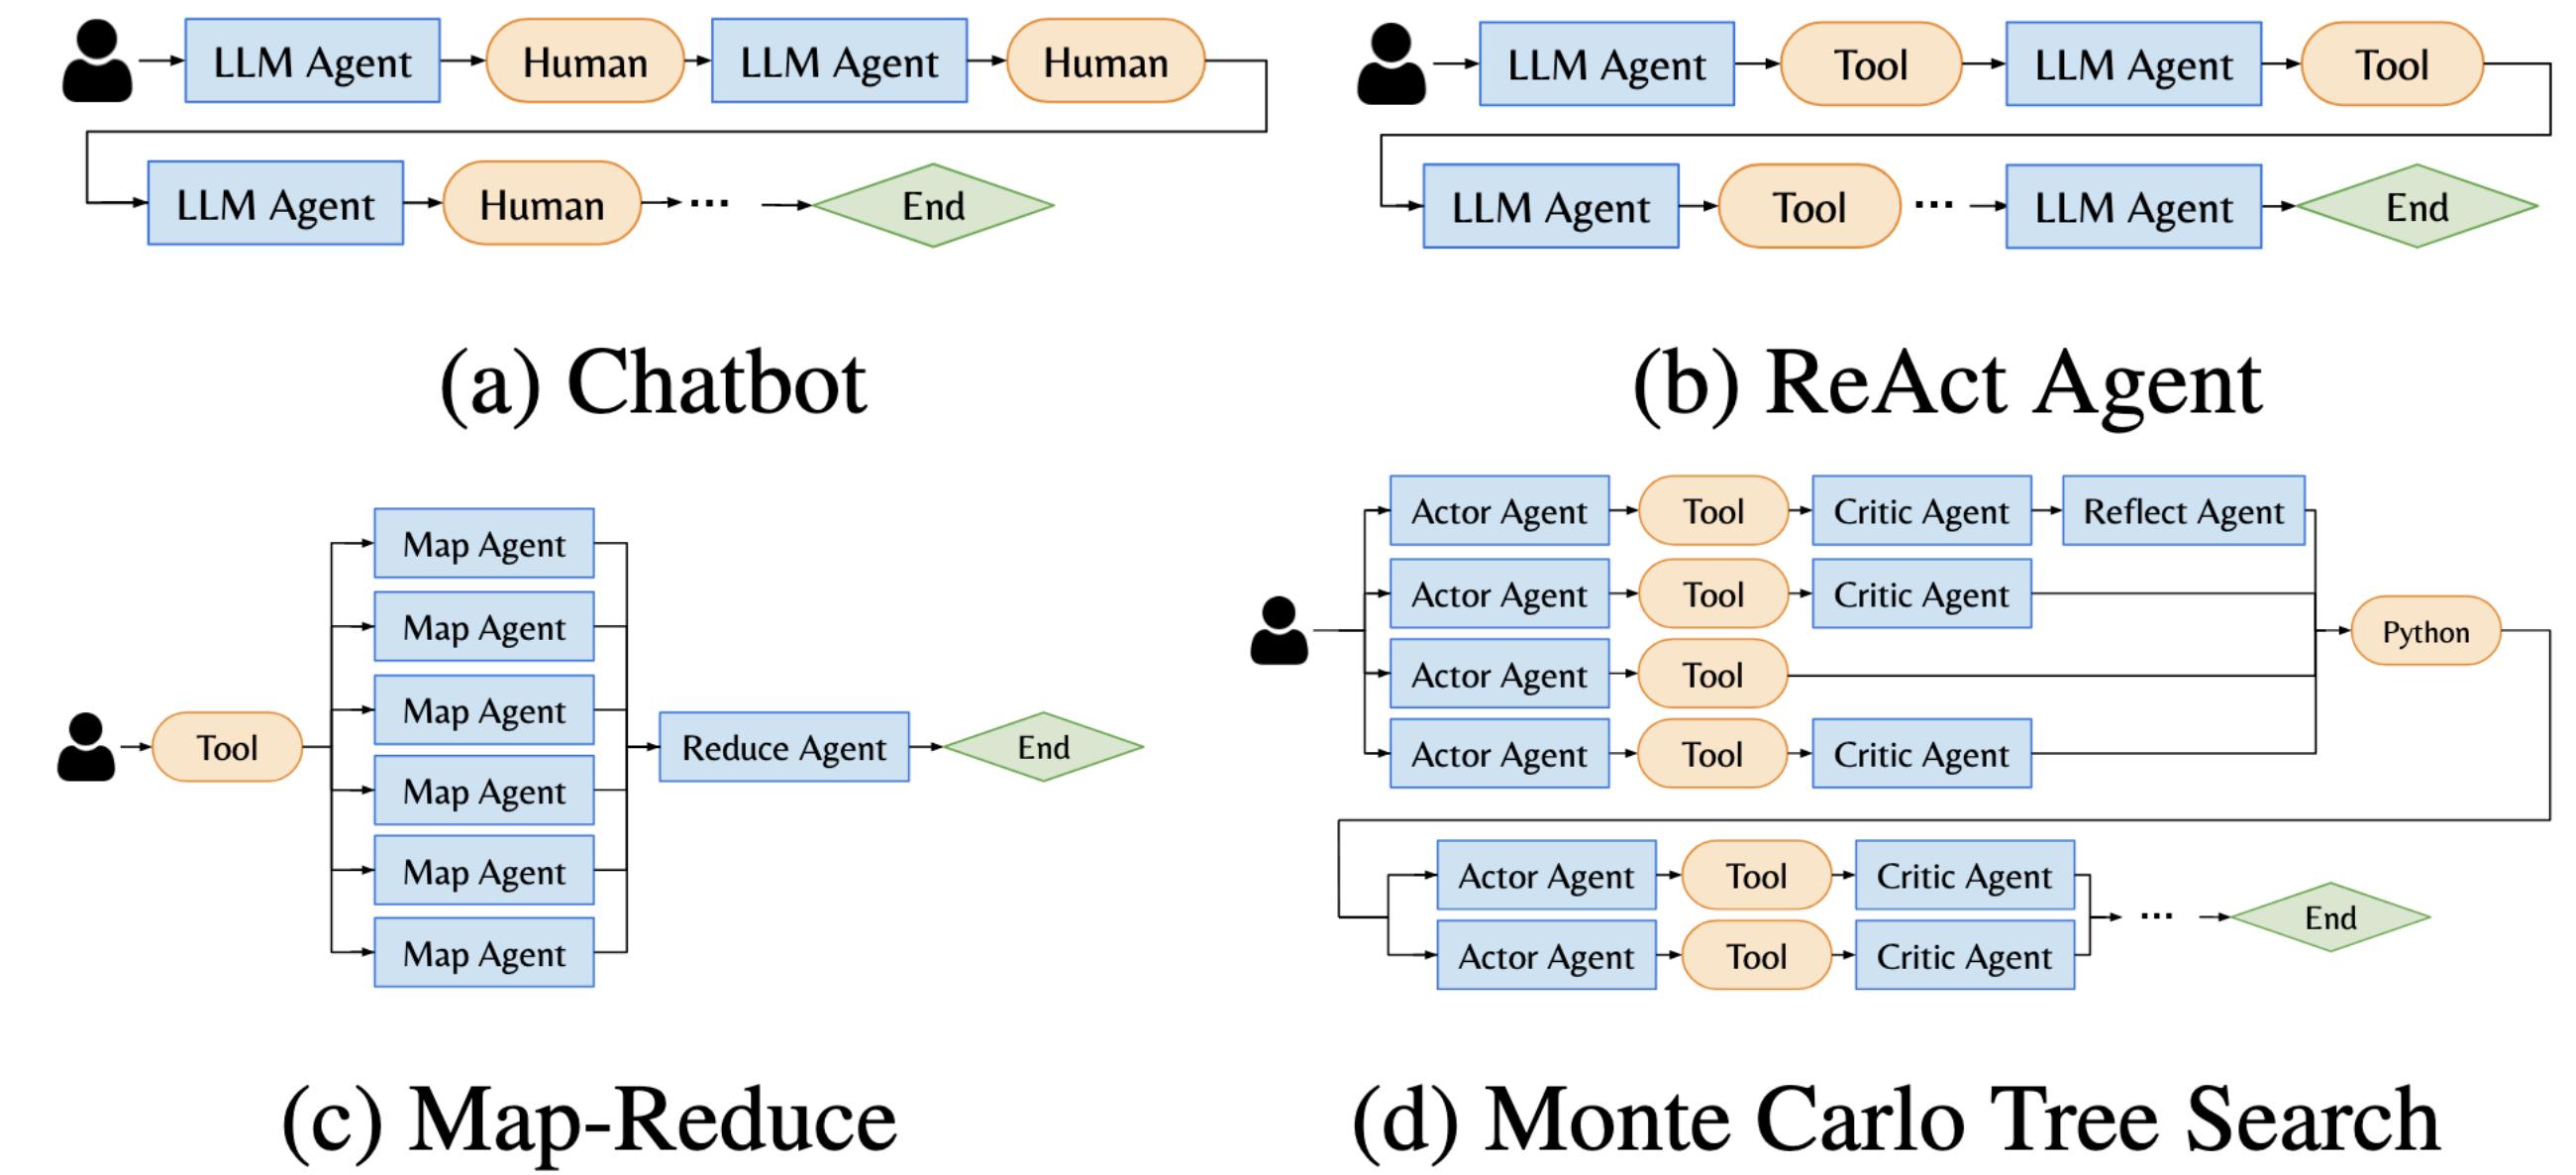


Figure 1: **Execution workflows for Agentic Programs.** Agentic programs are highly dynamic execution workflows that follow a directed acyclic graph (DAG). It consists of **LLM calls** from one or more LLM agents and **external interrupts** (i.e. tool calls, humans).

Image from Autellix paper

# Why should I care about learning OS?

If I don't want to do systems research

- OS is a study of abstraction. Absorb all the complexity away from the developer / the user.
- OS has to provide high performance. Manage resources, provide isolation and protection with minimal overheads.
- Principles are useful when designing any practical system