

# **Processes**

**Abhilash Jindal**

# Agenda

# Agenda

- Vocabulary (OSTEP Ch. 4)
  - What is a process? System calls? Scheduler? Address space?

# Agenda

- Vocabulary (OSTEP Ch. 4)
  - What is a process? System calls? Scheduler? Address space?
- Memory management (OSTEP Ch. 13-17)
  - How to manage and isolate memory? What are memory APIs? How are they implemented?

# Agenda

- Vocabulary (OSTEP Ch. 4)
  - What is a process? System calls? Scheduler? Address space?
- Memory management (OSTEP Ch. 13-17)
  - How to manage and isolate memory? What are memory APIs? How are they implemented?
- Processes in action (xv6 Ch. 3: system calls, x86 protection, trap handlers)
  - Process control block, user stack<>kernel stack, sys call handling

# Agenda

- Vocabulary (OSTEP Ch. 4)
  - What is a process? System calls? Scheduler? Address space?
- Memory management (OSTEP Ch. 13-17)
  - How to manage and isolate memory? What are memory APIs? How are they implemented?
- Processes in action (xv6 Ch. 3: system calls, x86 protection, trap handlers)
  - Process control block, user stack<>kernel stack, sys call handling
- Scheduling (xv6 Ch5: context switching, OSTEP Ch. 6-9)
  - Response time, throughput, fairness

# Process is a running program

- Load program from disk to memory
  - Exactly how we loaded OS
- Give control to the process. Jump cs, eip

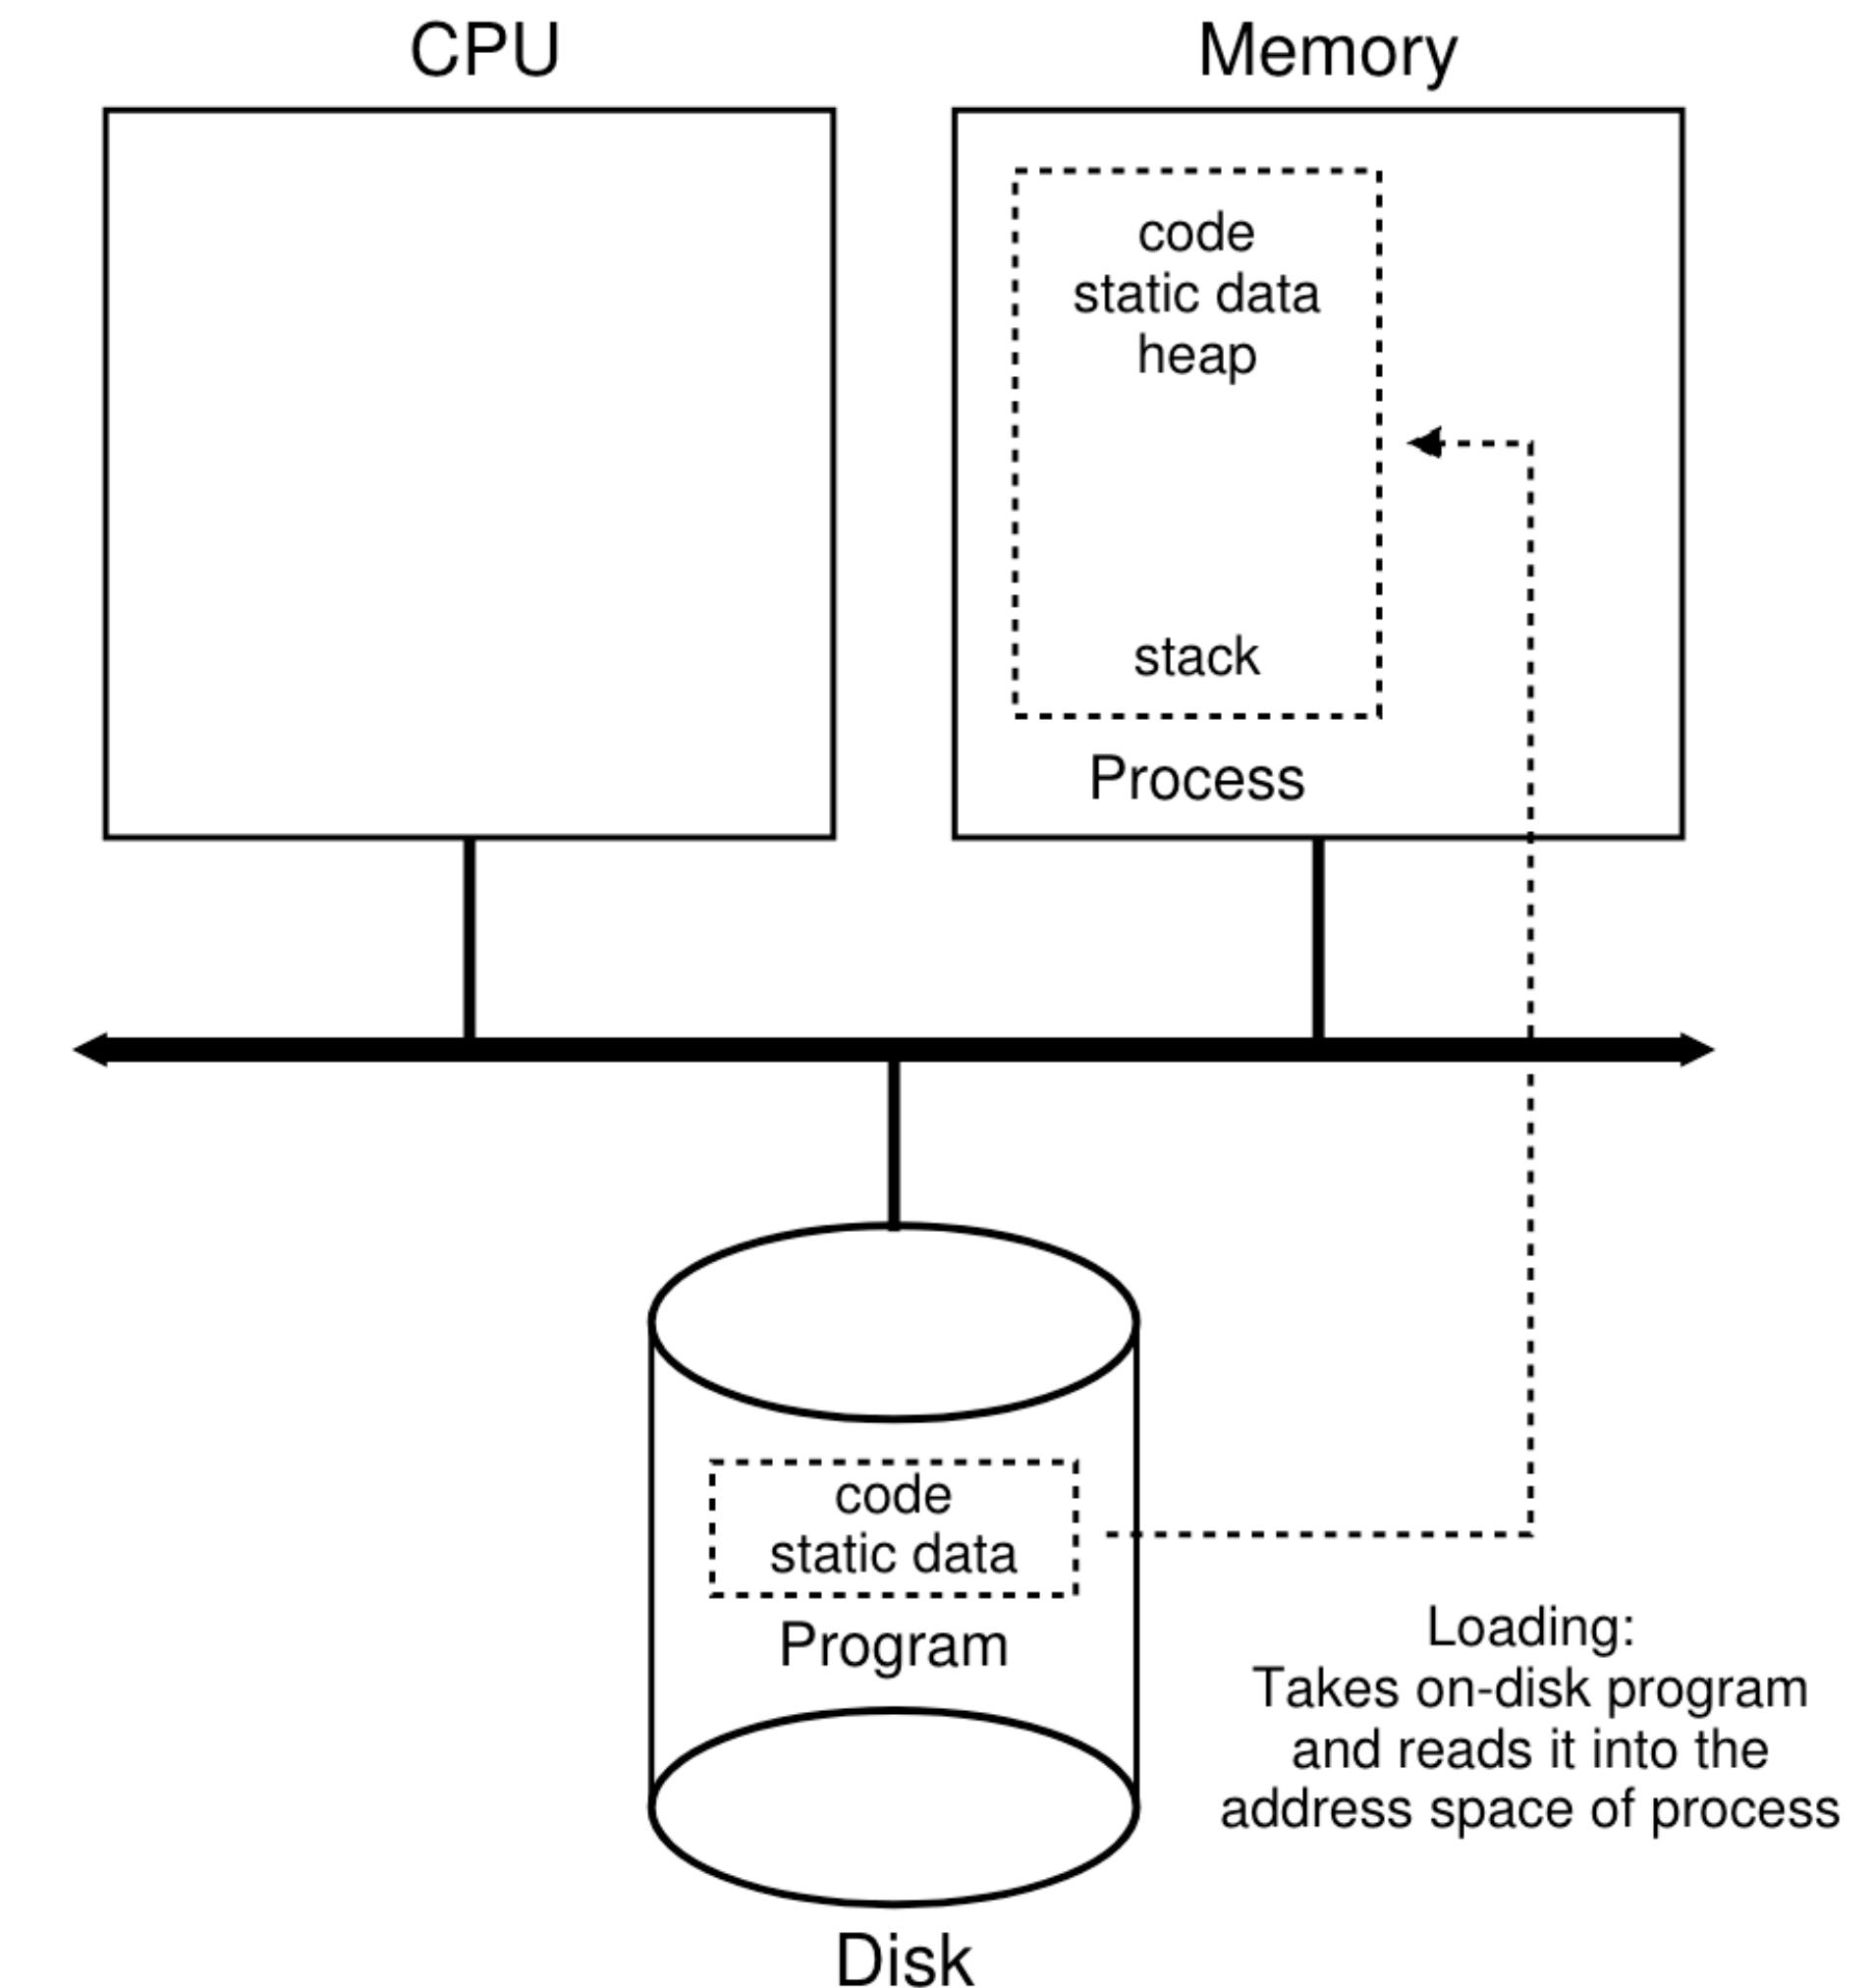


Figure 4.1: Loading: From Program To Process

# Processes can ask OS to do work for them

## System calls

```
$ strace cat /tmp/foo
...
openat(AT_FDCWD, "/tmp/foo", 0_RDONLY) = 3
read(3, "hi\n", 131072)                = 3
write(1, "hi\n", 3)                    = 3
...
```

# OS maintains process states

## Scheduler switches between processes

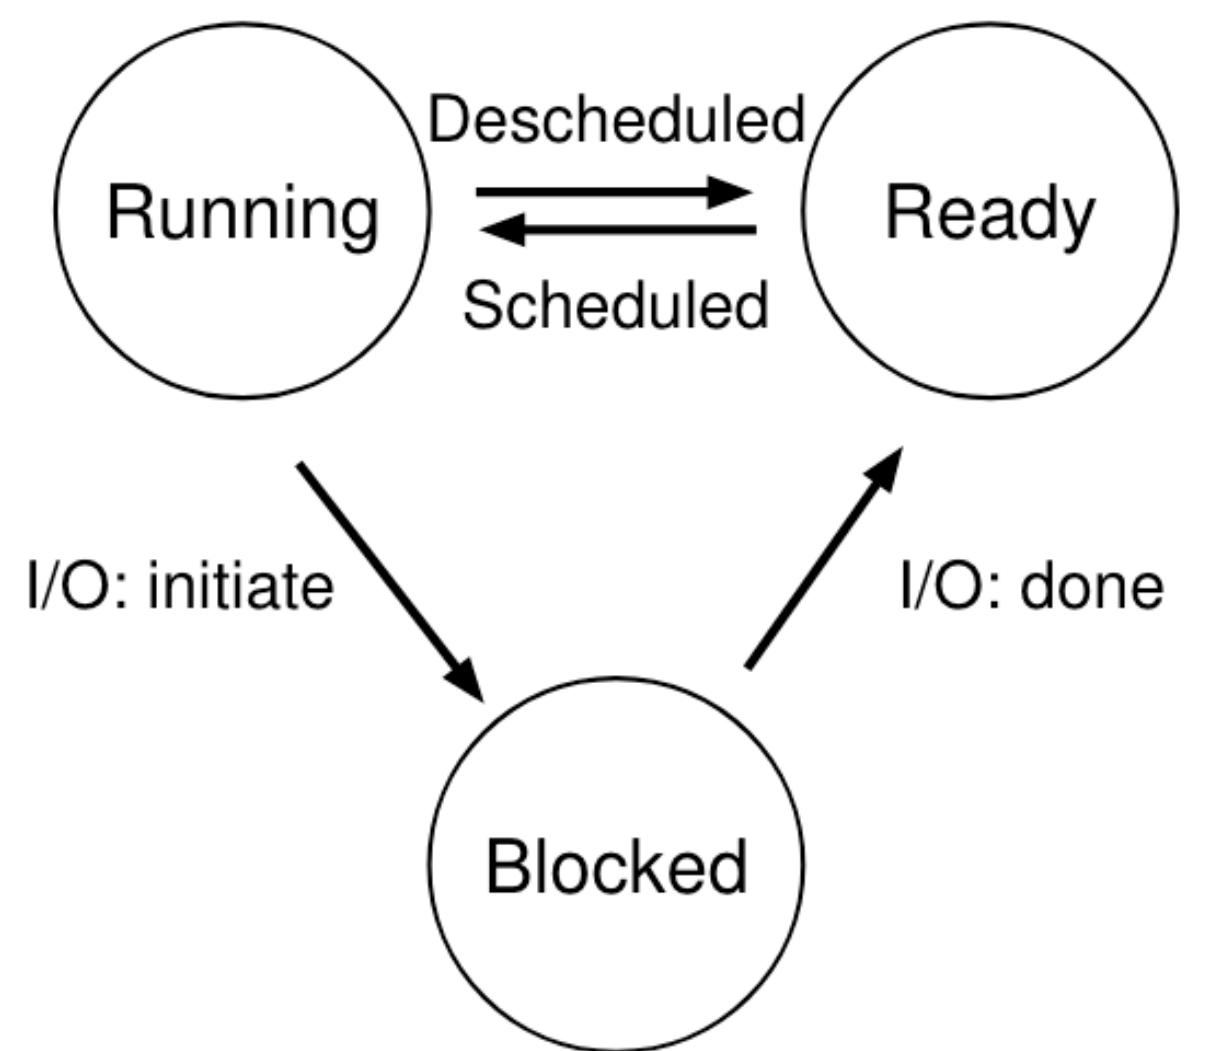


Figure 4.2: Process: State Transitions

# OS maintains process states

## Scheduler switches between processes

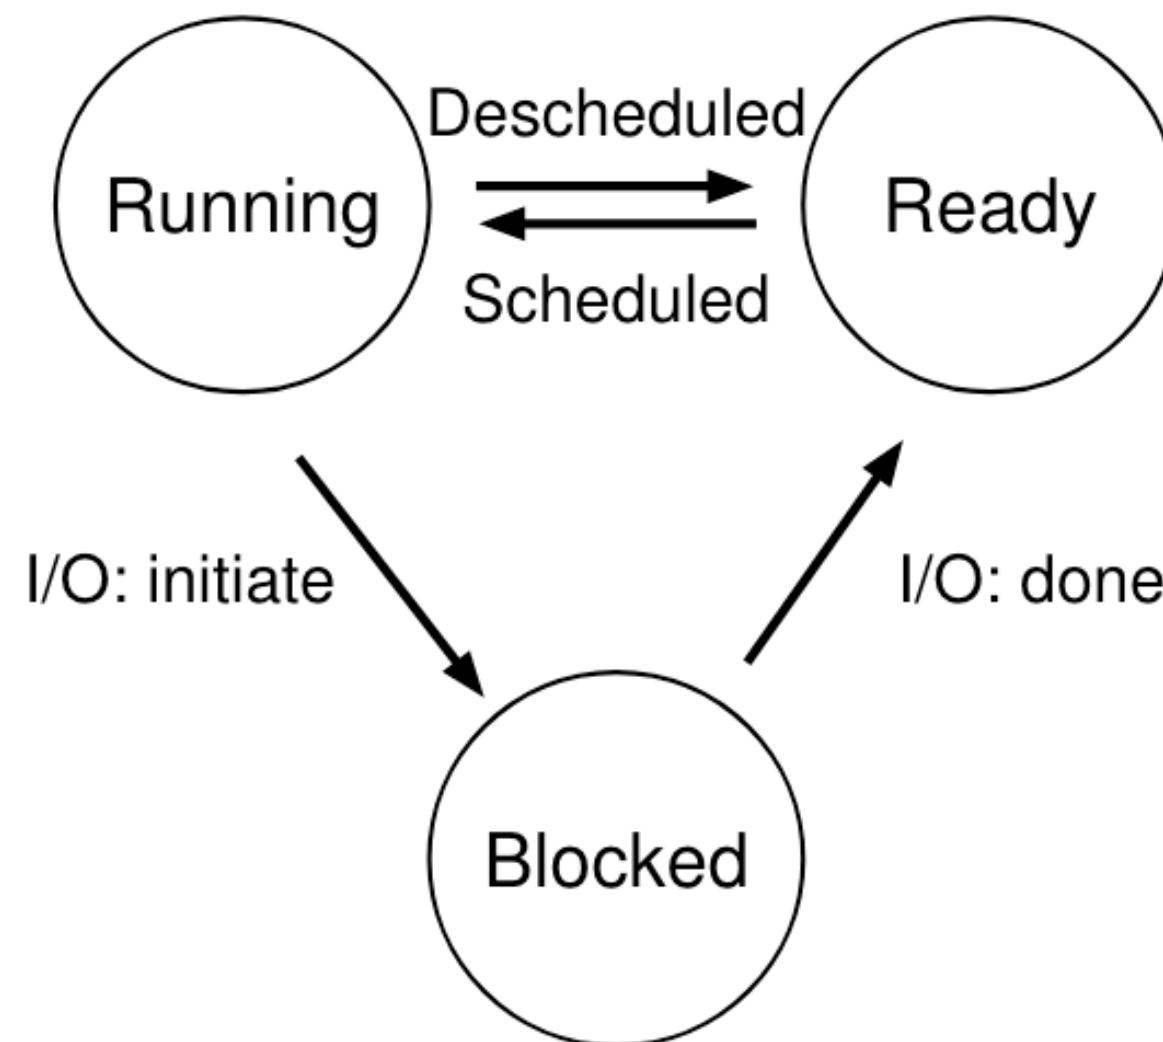


Figure 4.2: Process: State Transitions

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Blocked	Running	Process <sub>0</sub> initiates I/O
5	Blocked	Running	Process <sub>0</sub> is blocked, so Process <sub>1</sub> runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process <sub>1</sub> now done
9	Running	—	
10	Running	—	Process <sub>0</sub> now done

Figure 4.4: Tracing Process State: CPU and I/O

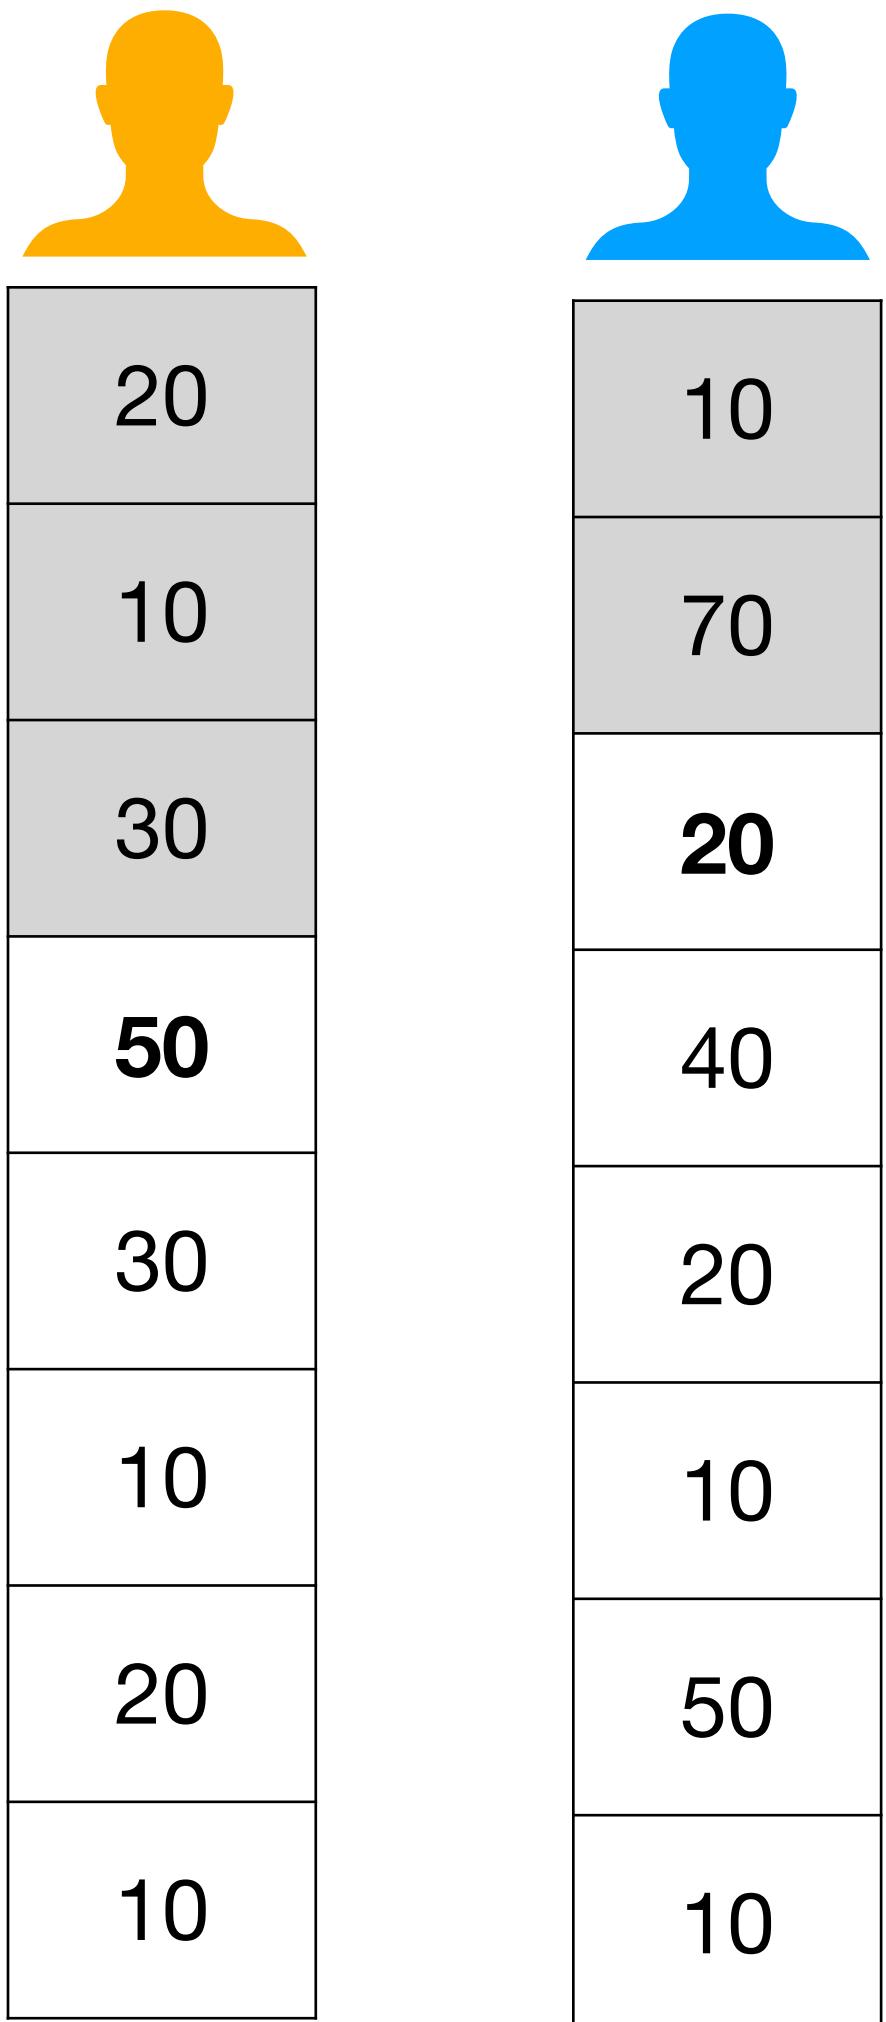
# Calculator analogy: Computing long sum



20
10
30
50
30
10
20
10

- $2 \ 0 =$  (move pointer to 10)
- $+ 1 \ 0 =$  (move pointer to 30)
- $+ 3 \ 0 =$  (move pointer to 50)
- $+ 5 \ 0 =$  (move pointer to 30)
- $+ 3 \ 0 =$  (move pointer to 10)
- $+ 1 \ 0 =$  (move pointer to 20)
- $+ 2 \ 0 =$  (move pointer to 10)

# Sharing the calculator



# Sharing the calculator

20	10
10	70
30	20
50	40
30	20
10	10
20	50
10	10

- Steps to share the calculator:
  - $20 + 10 = 30 + 30 = 60$
  - Write 60 in notebook, remember that we were done till 30, give calculator
  - $10 + 70 = 80$
  - Write 80 in notebook, remember that we were done till 70, give the calculator back

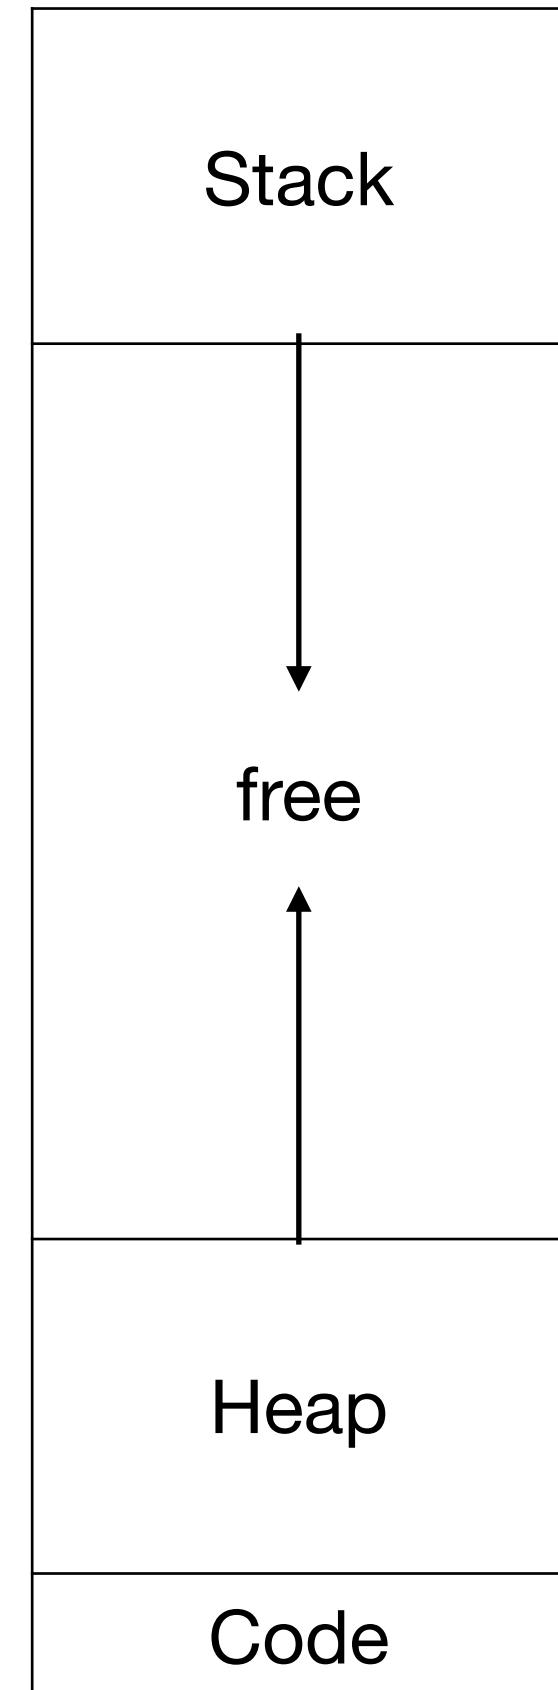
# **Memory isolation and management**

**OSTEP Ch. 13-17**

**Abhilash Jindal**

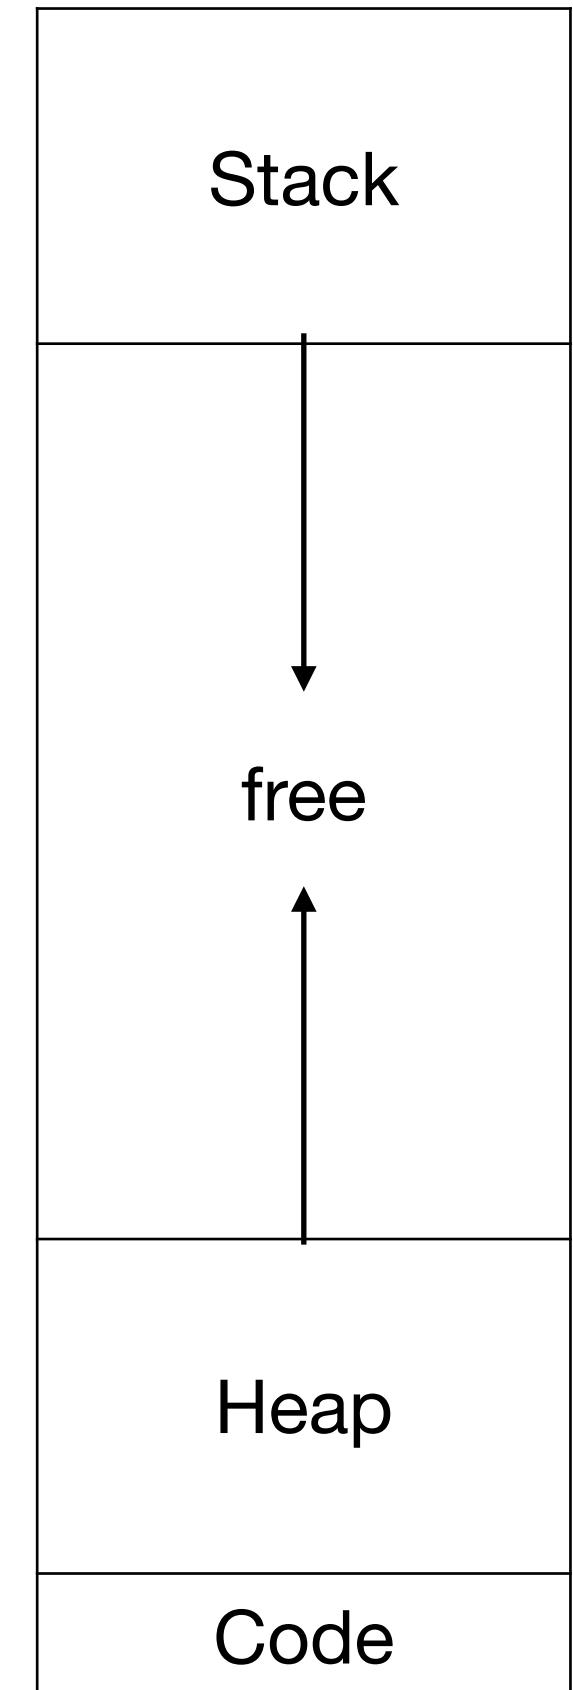
# Process Address Space

## Code, Heap, Stack



# Process Address Space

## Code, Heap, Stack



Process address space

# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

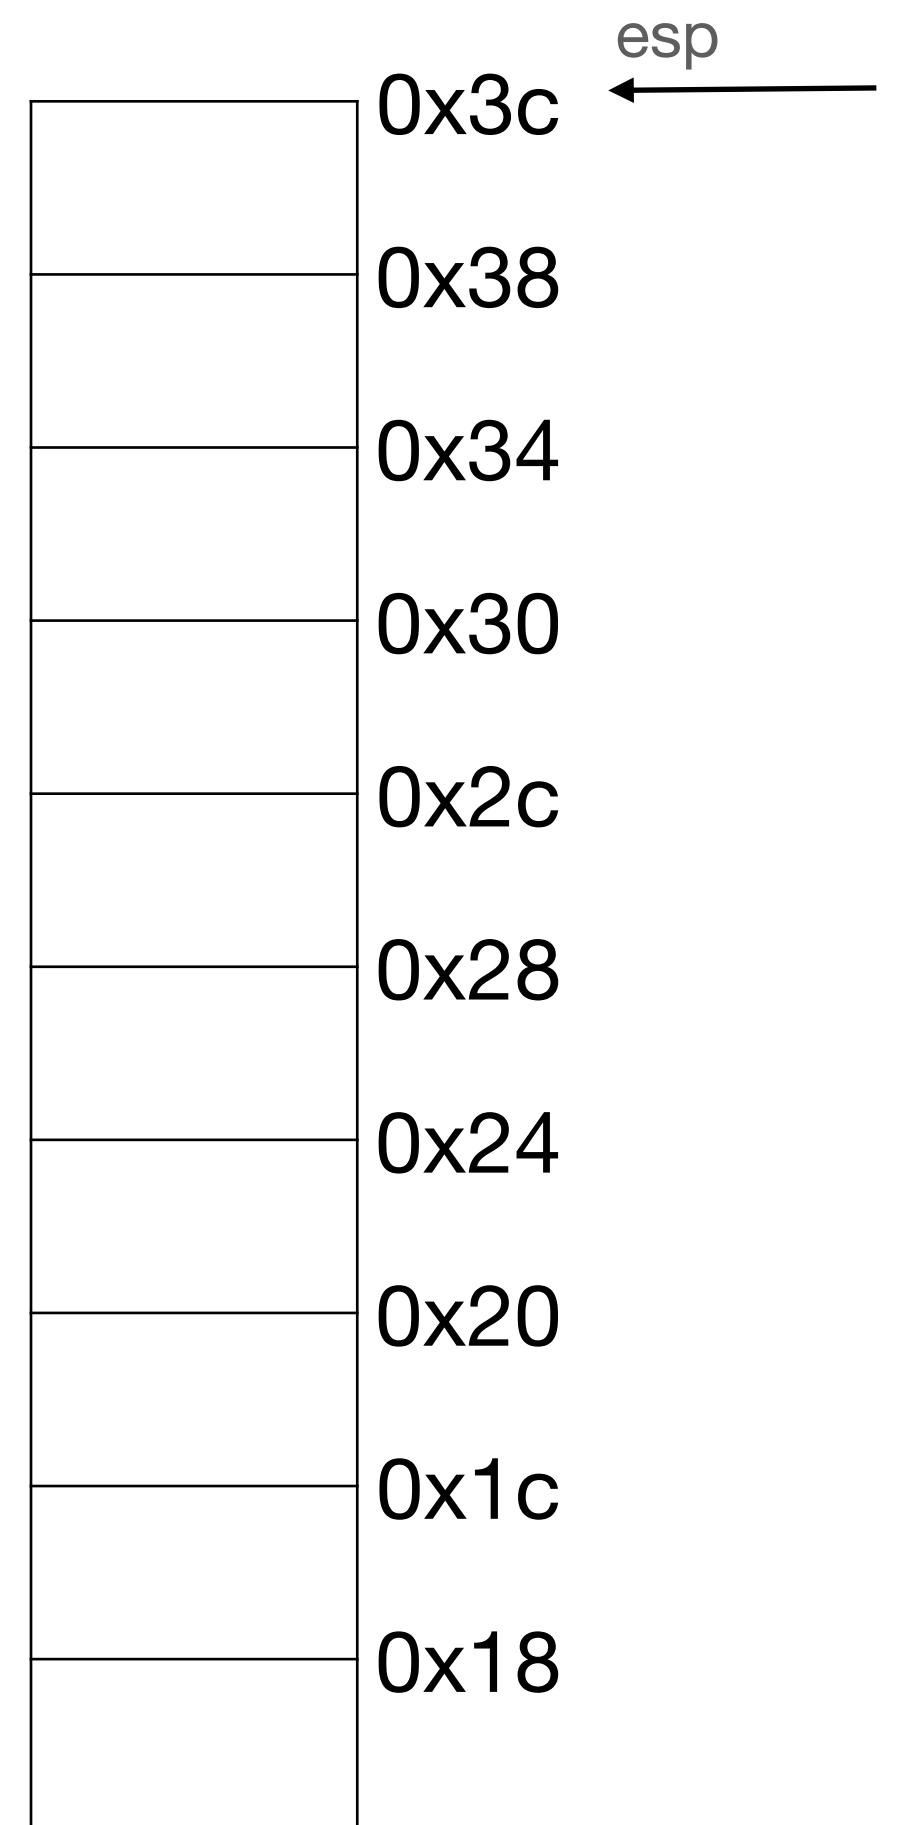
eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp

ebp →



# Function calling in action

## Stack

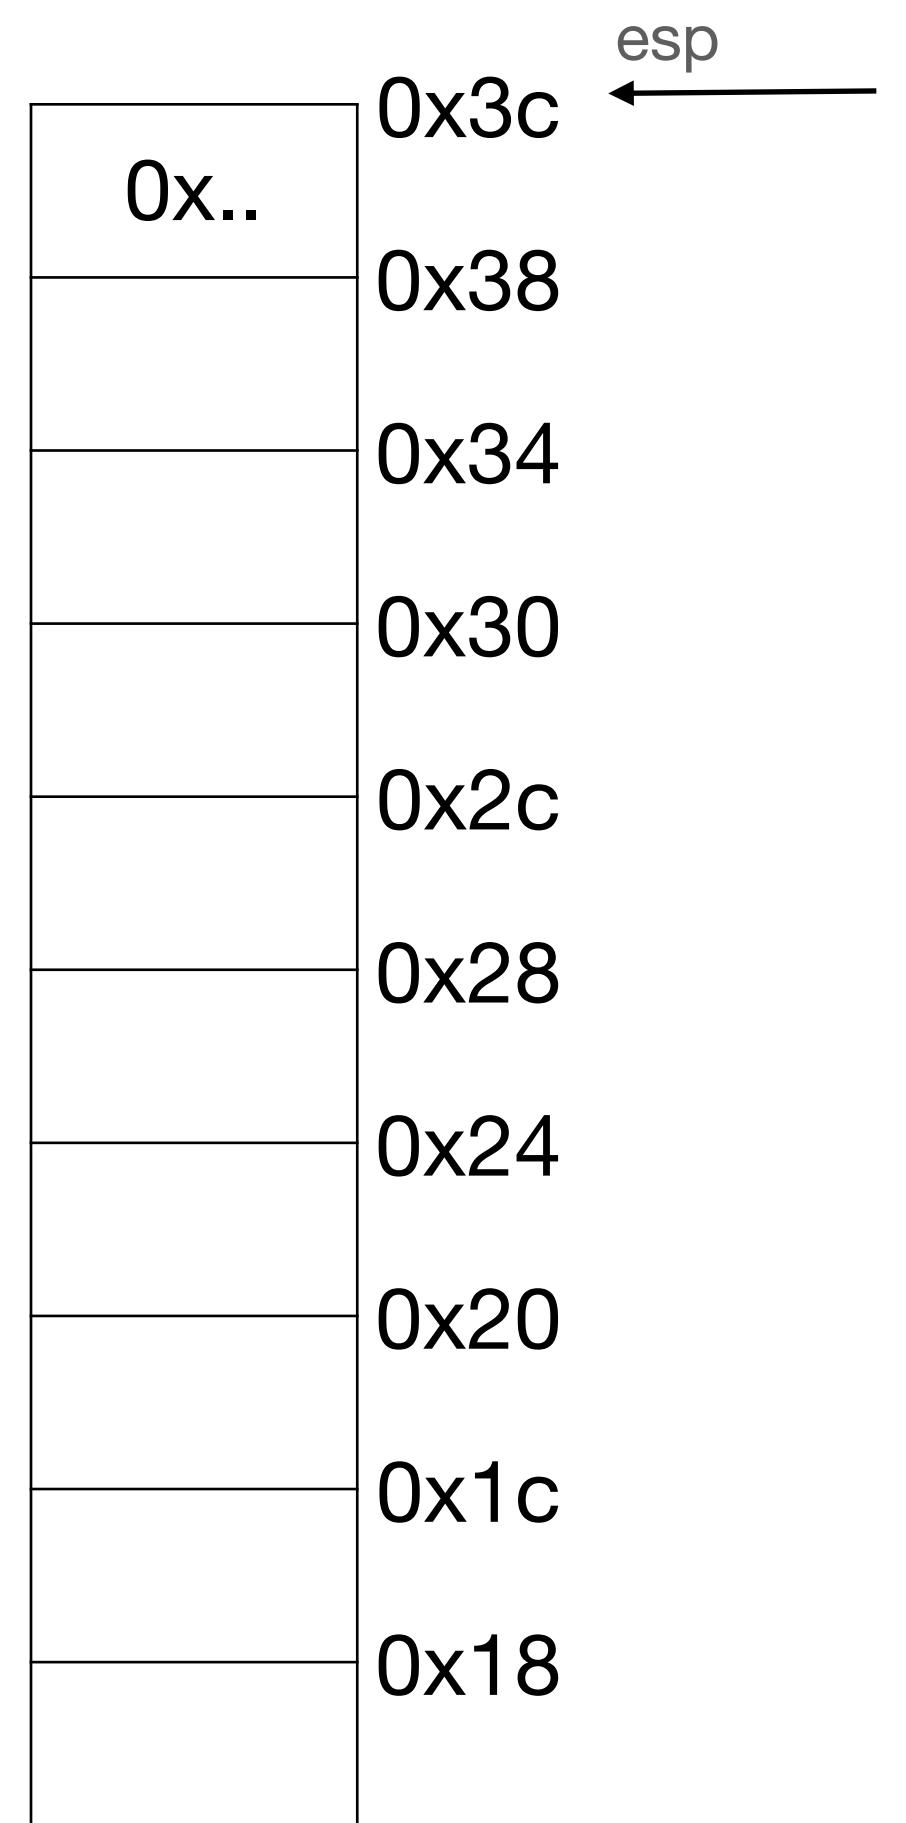
```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

ebp →



eip →

# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

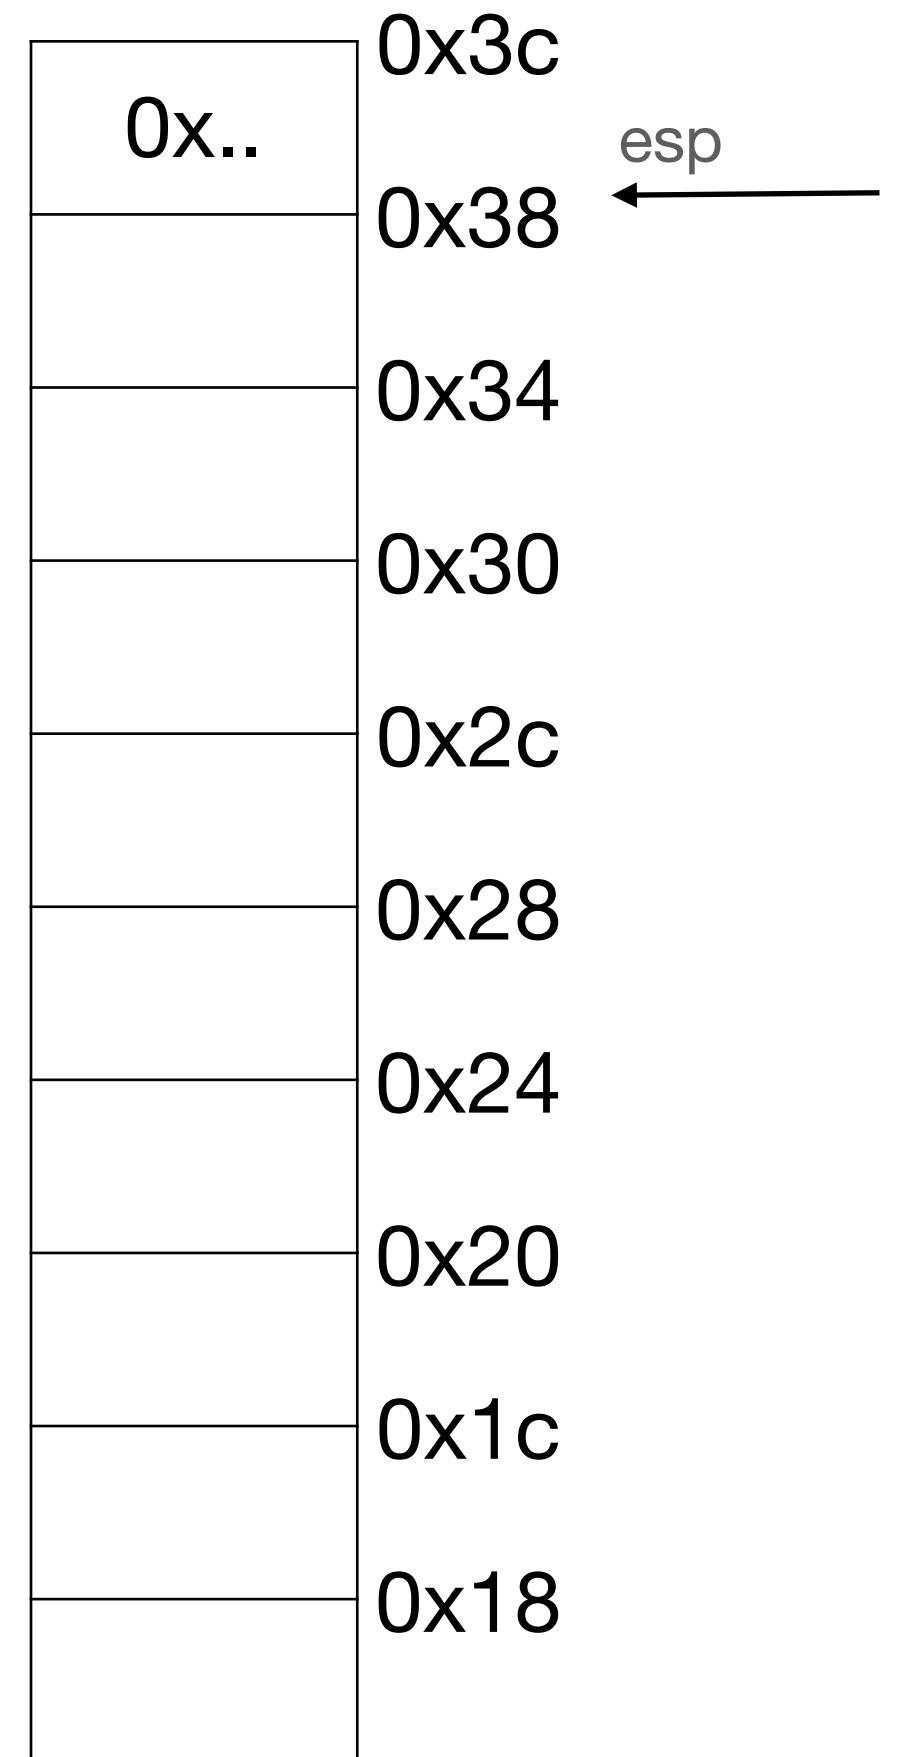
eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp

ebp →



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

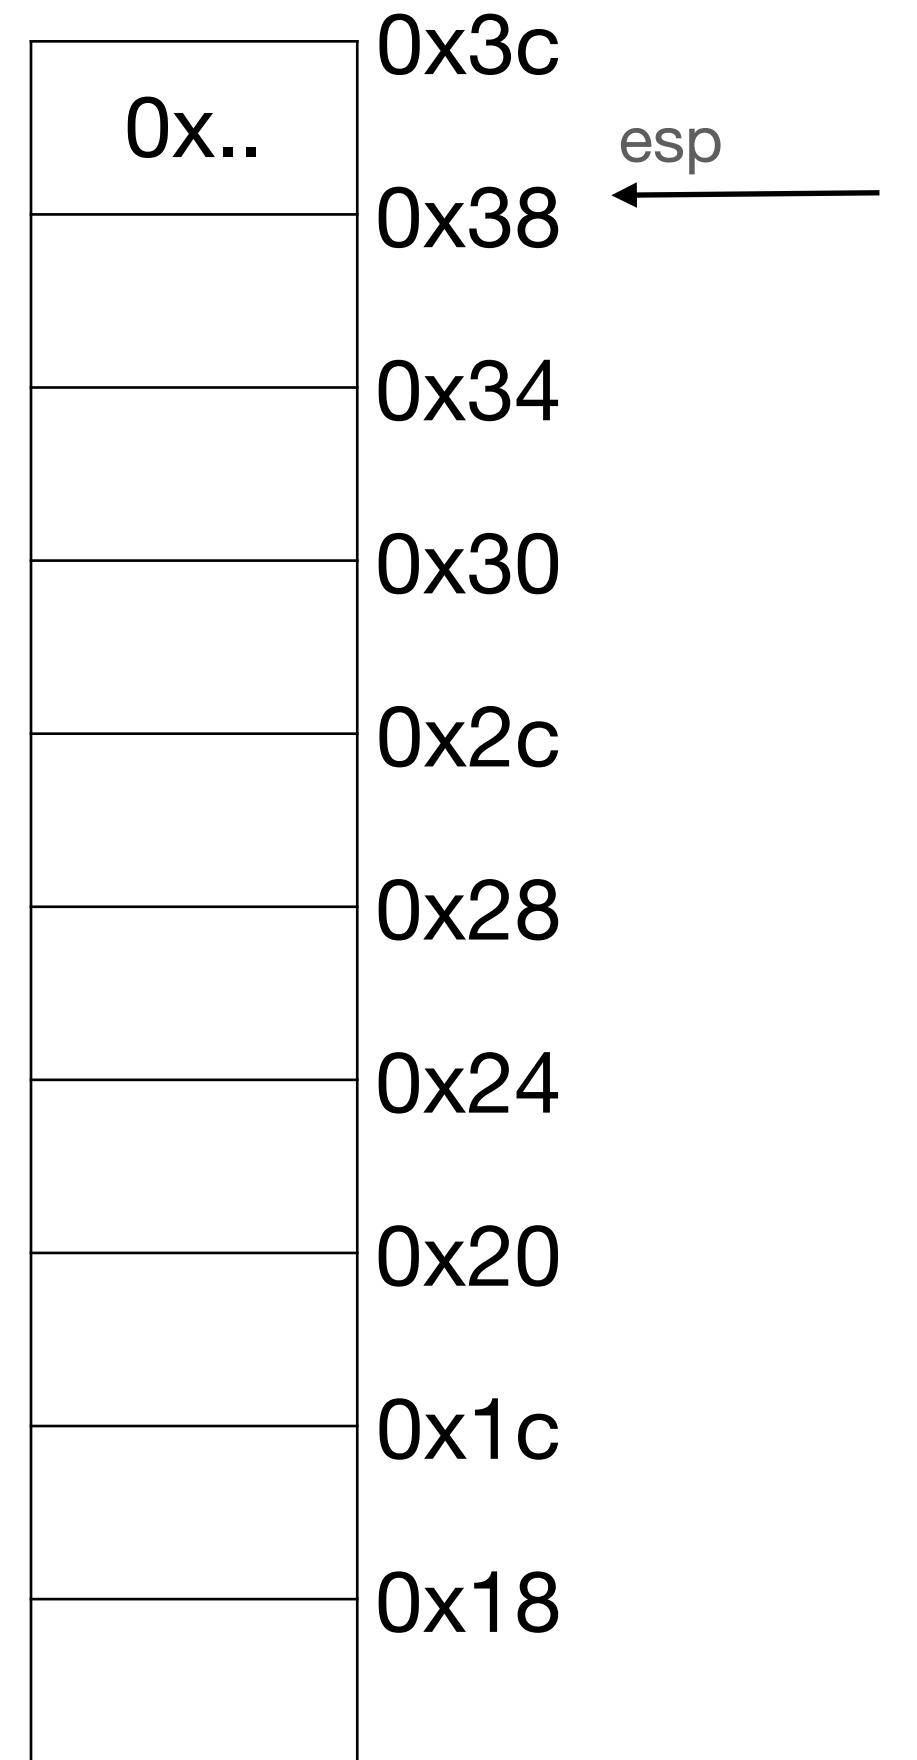
eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp

ebp →



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

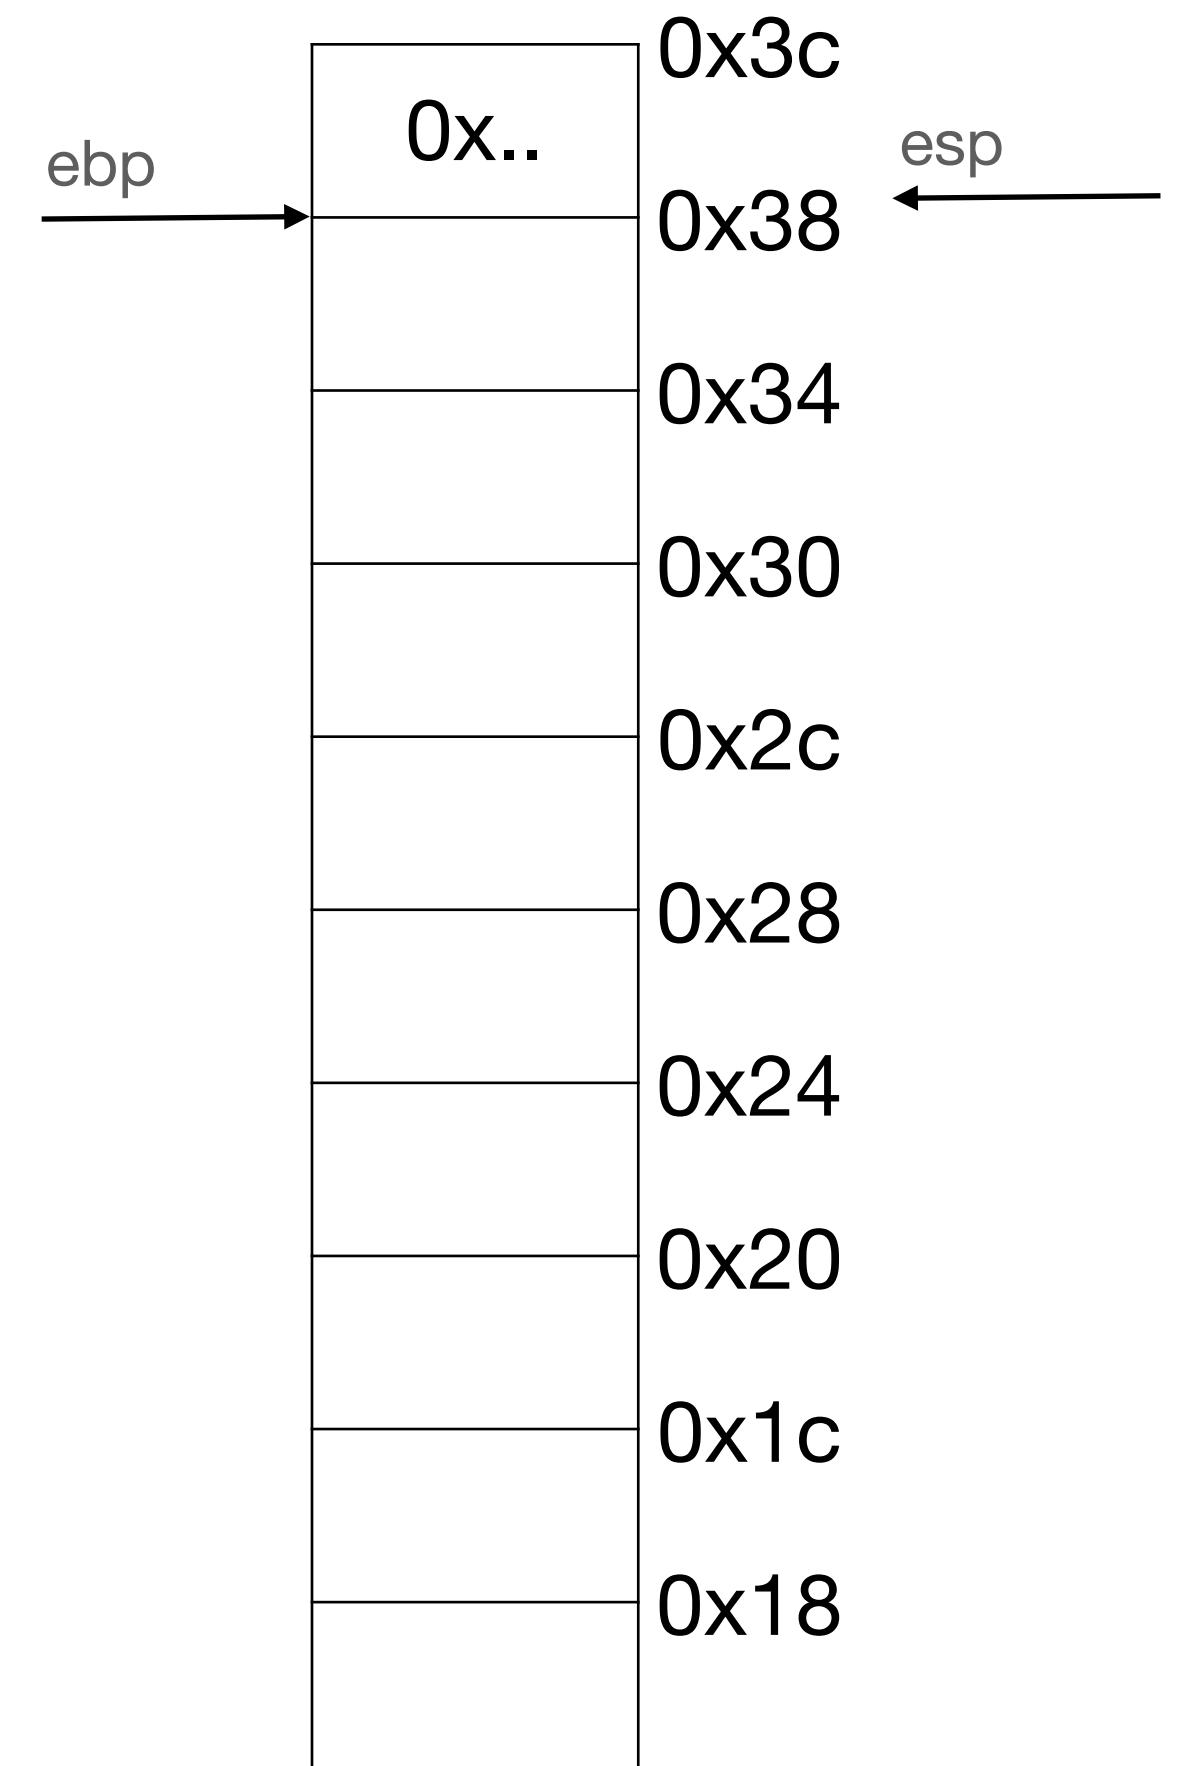
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

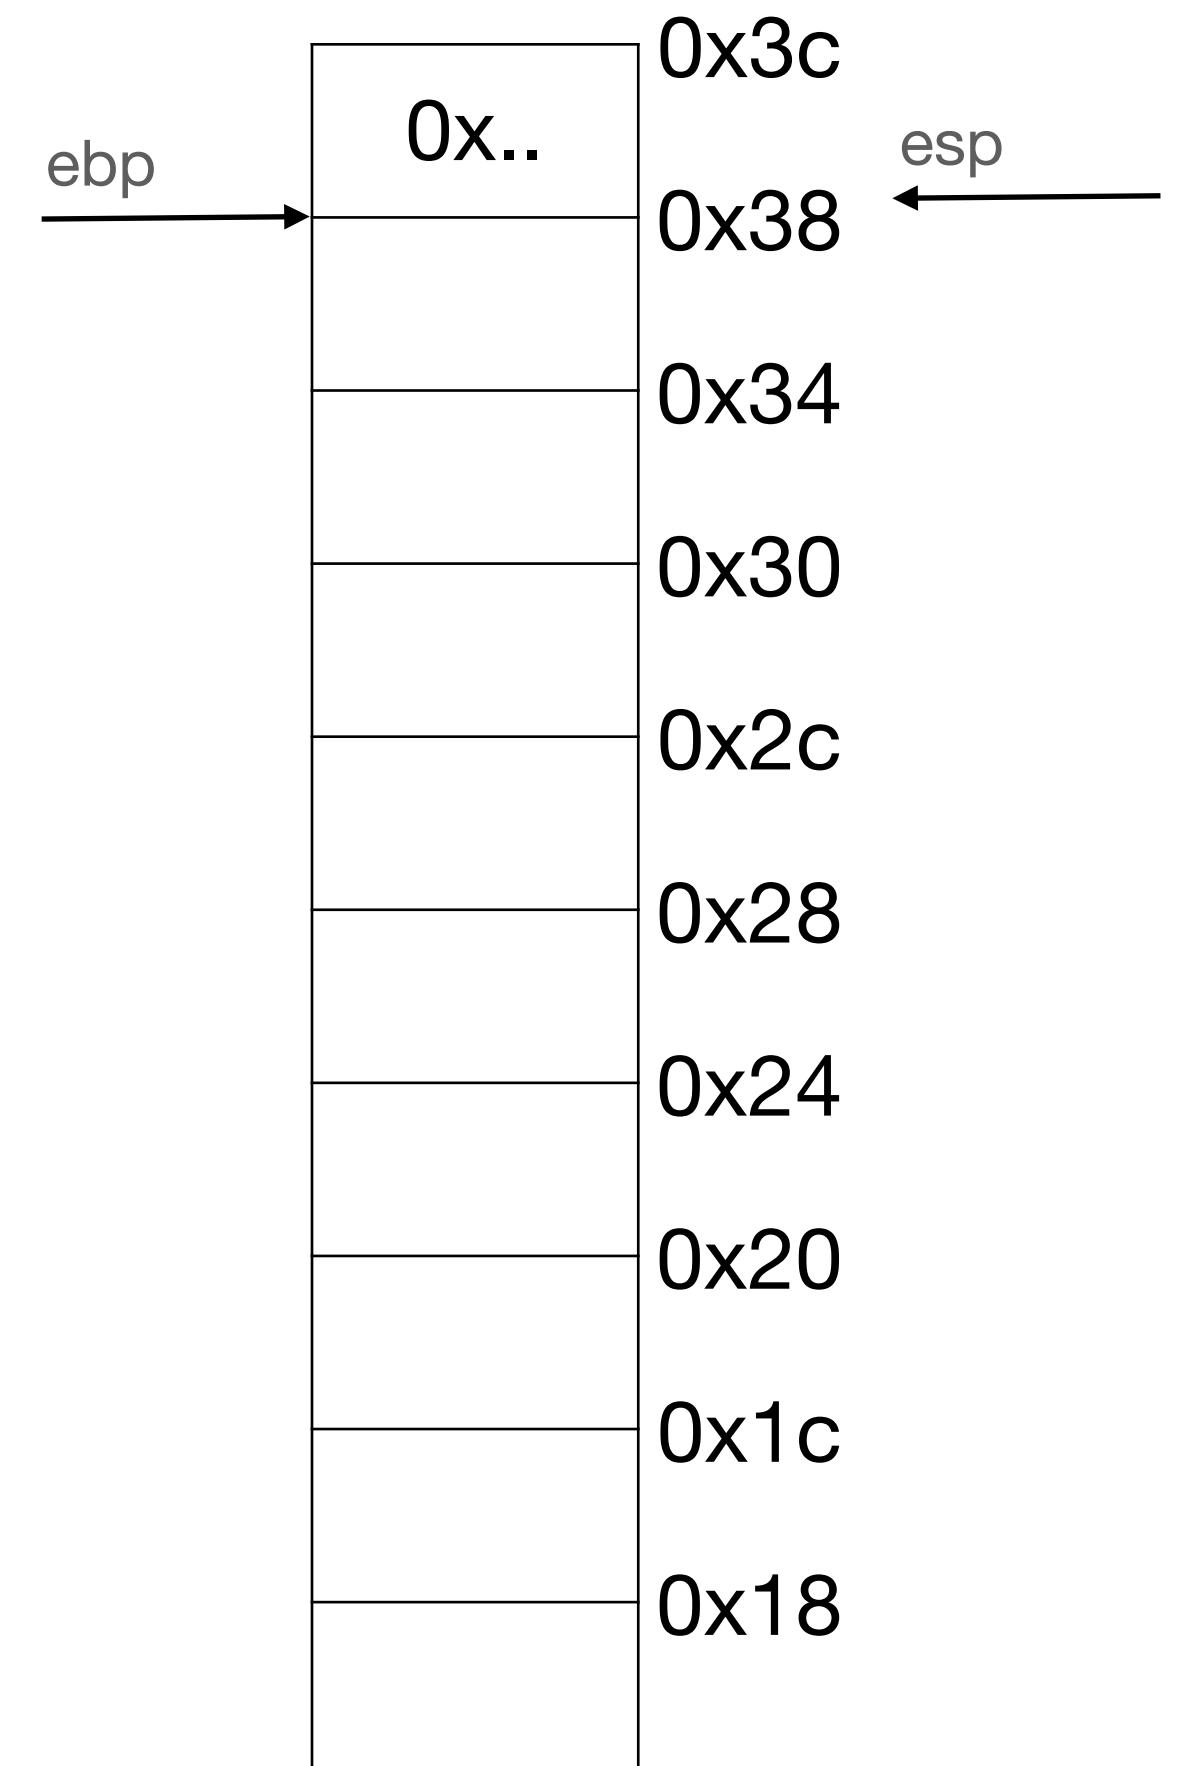
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

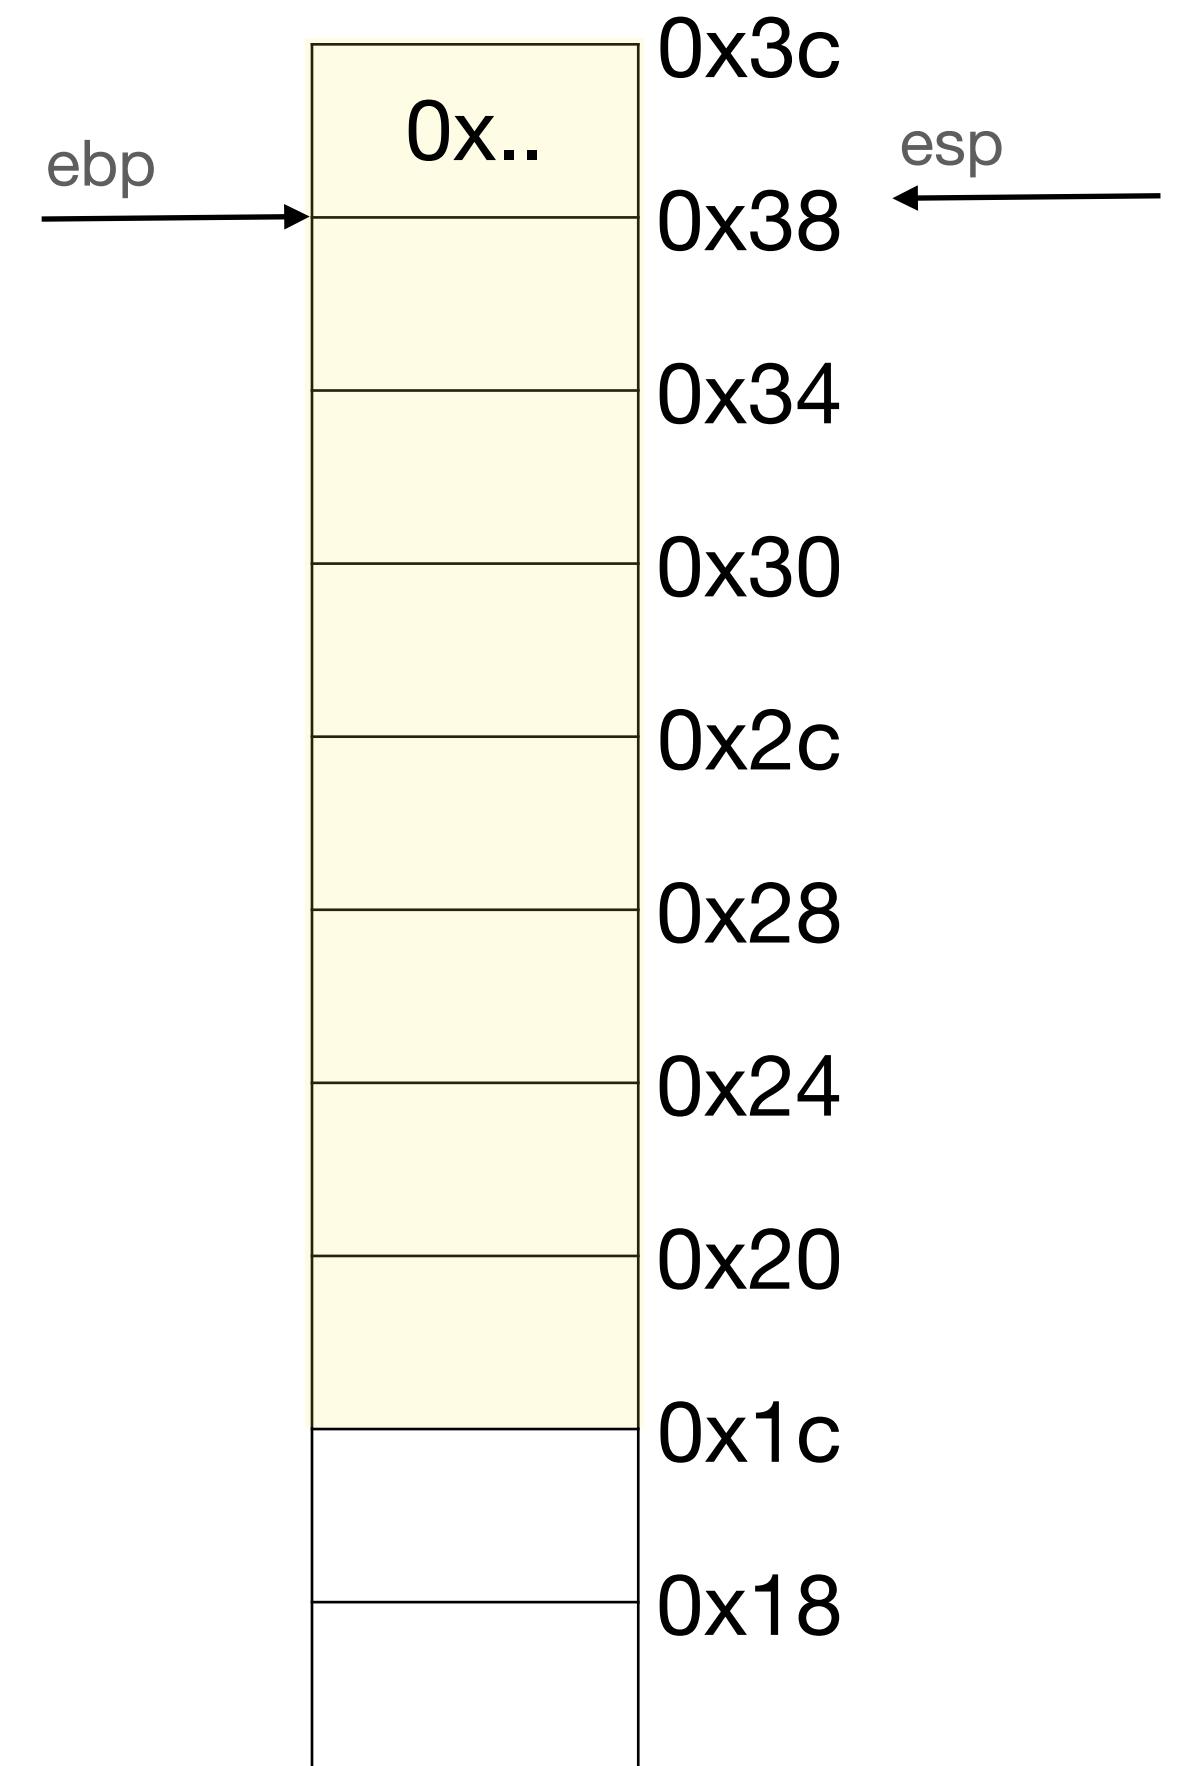
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

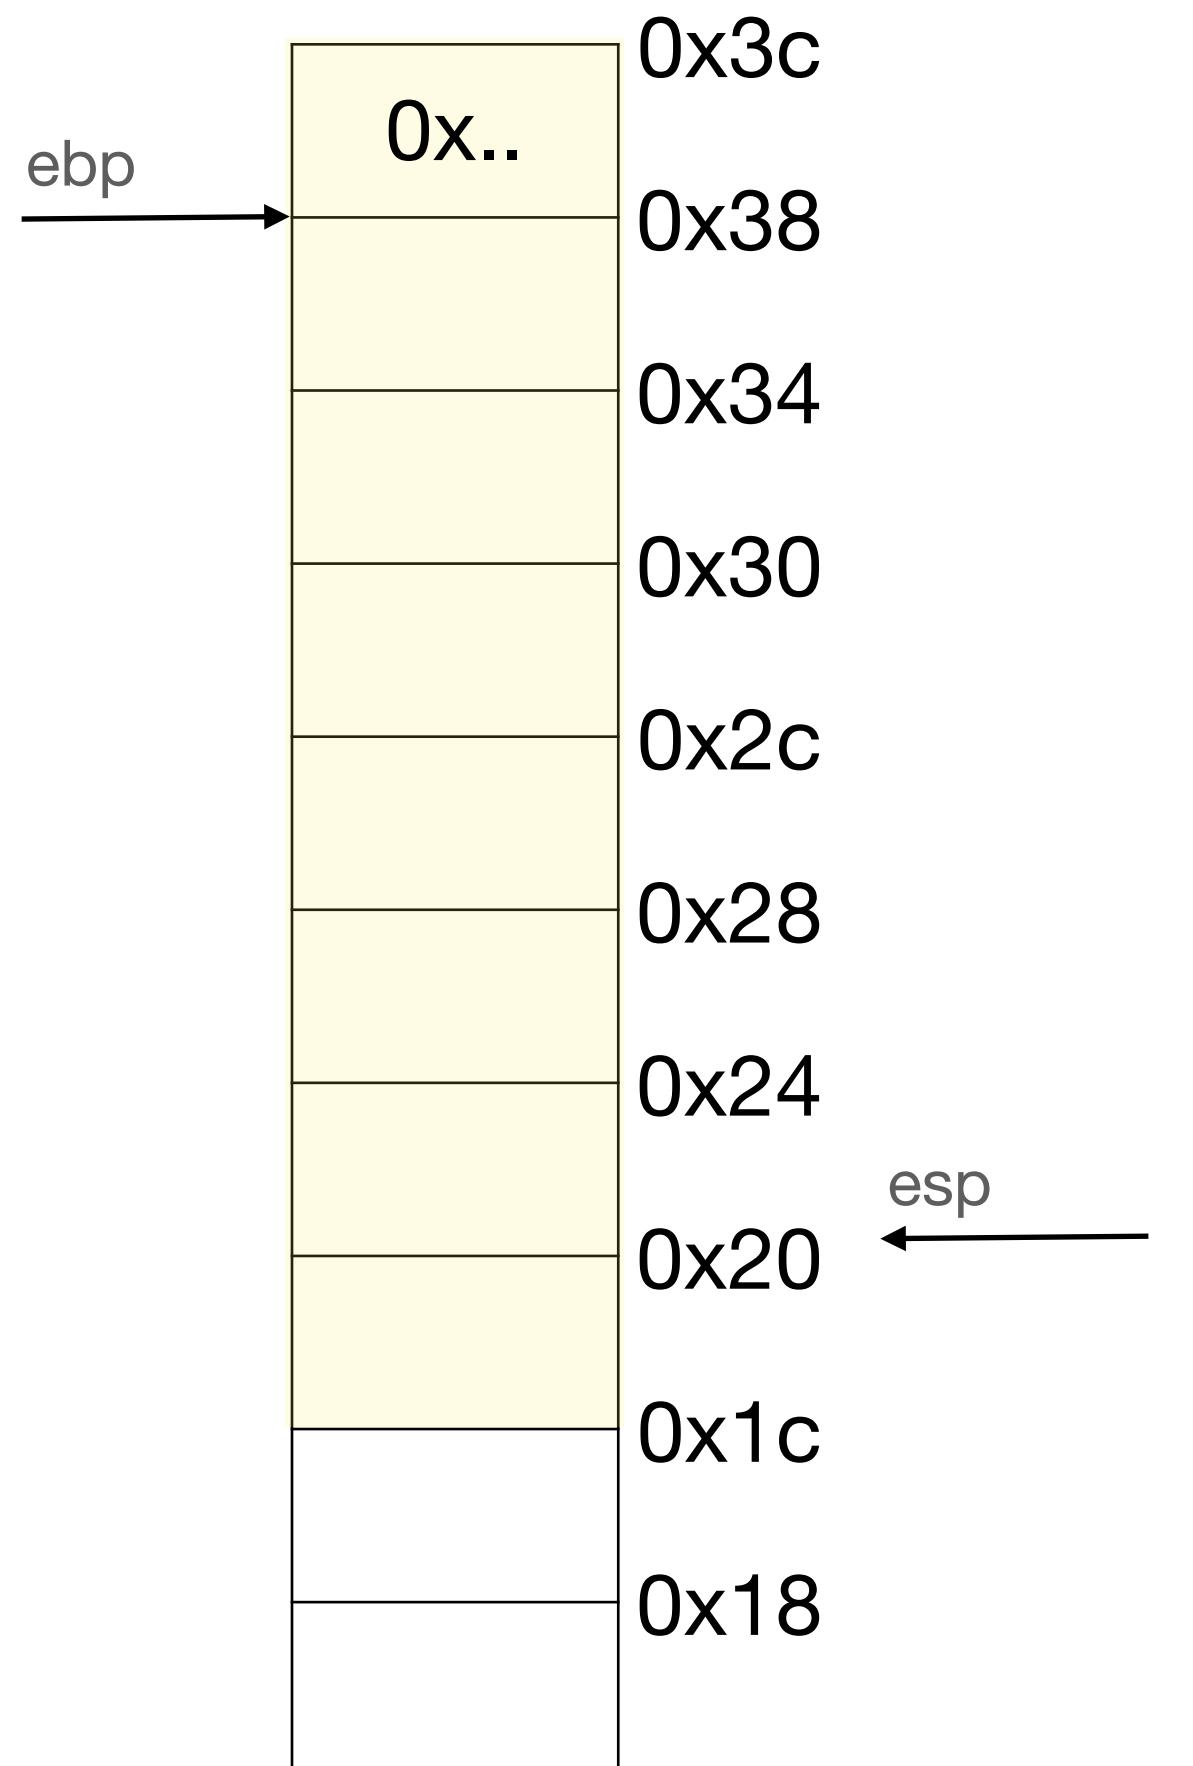
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

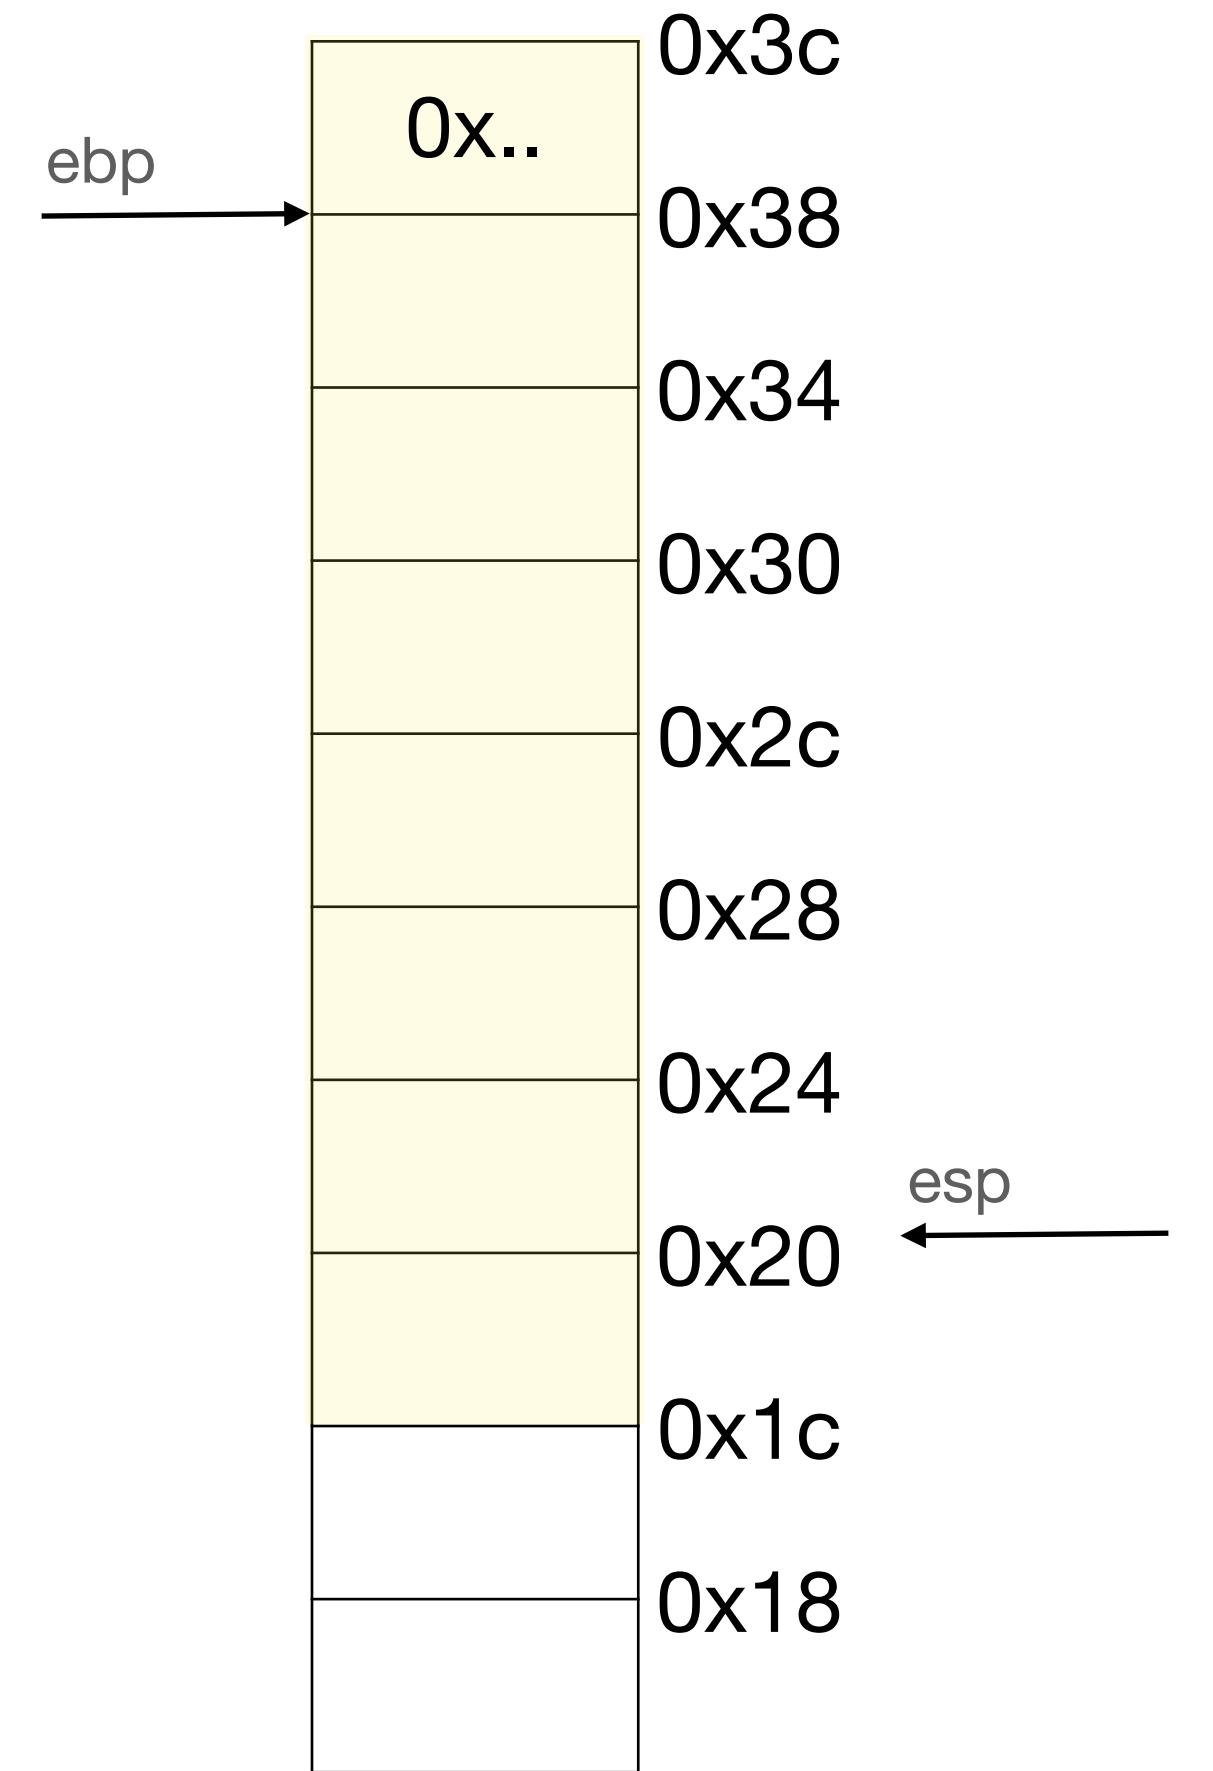
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

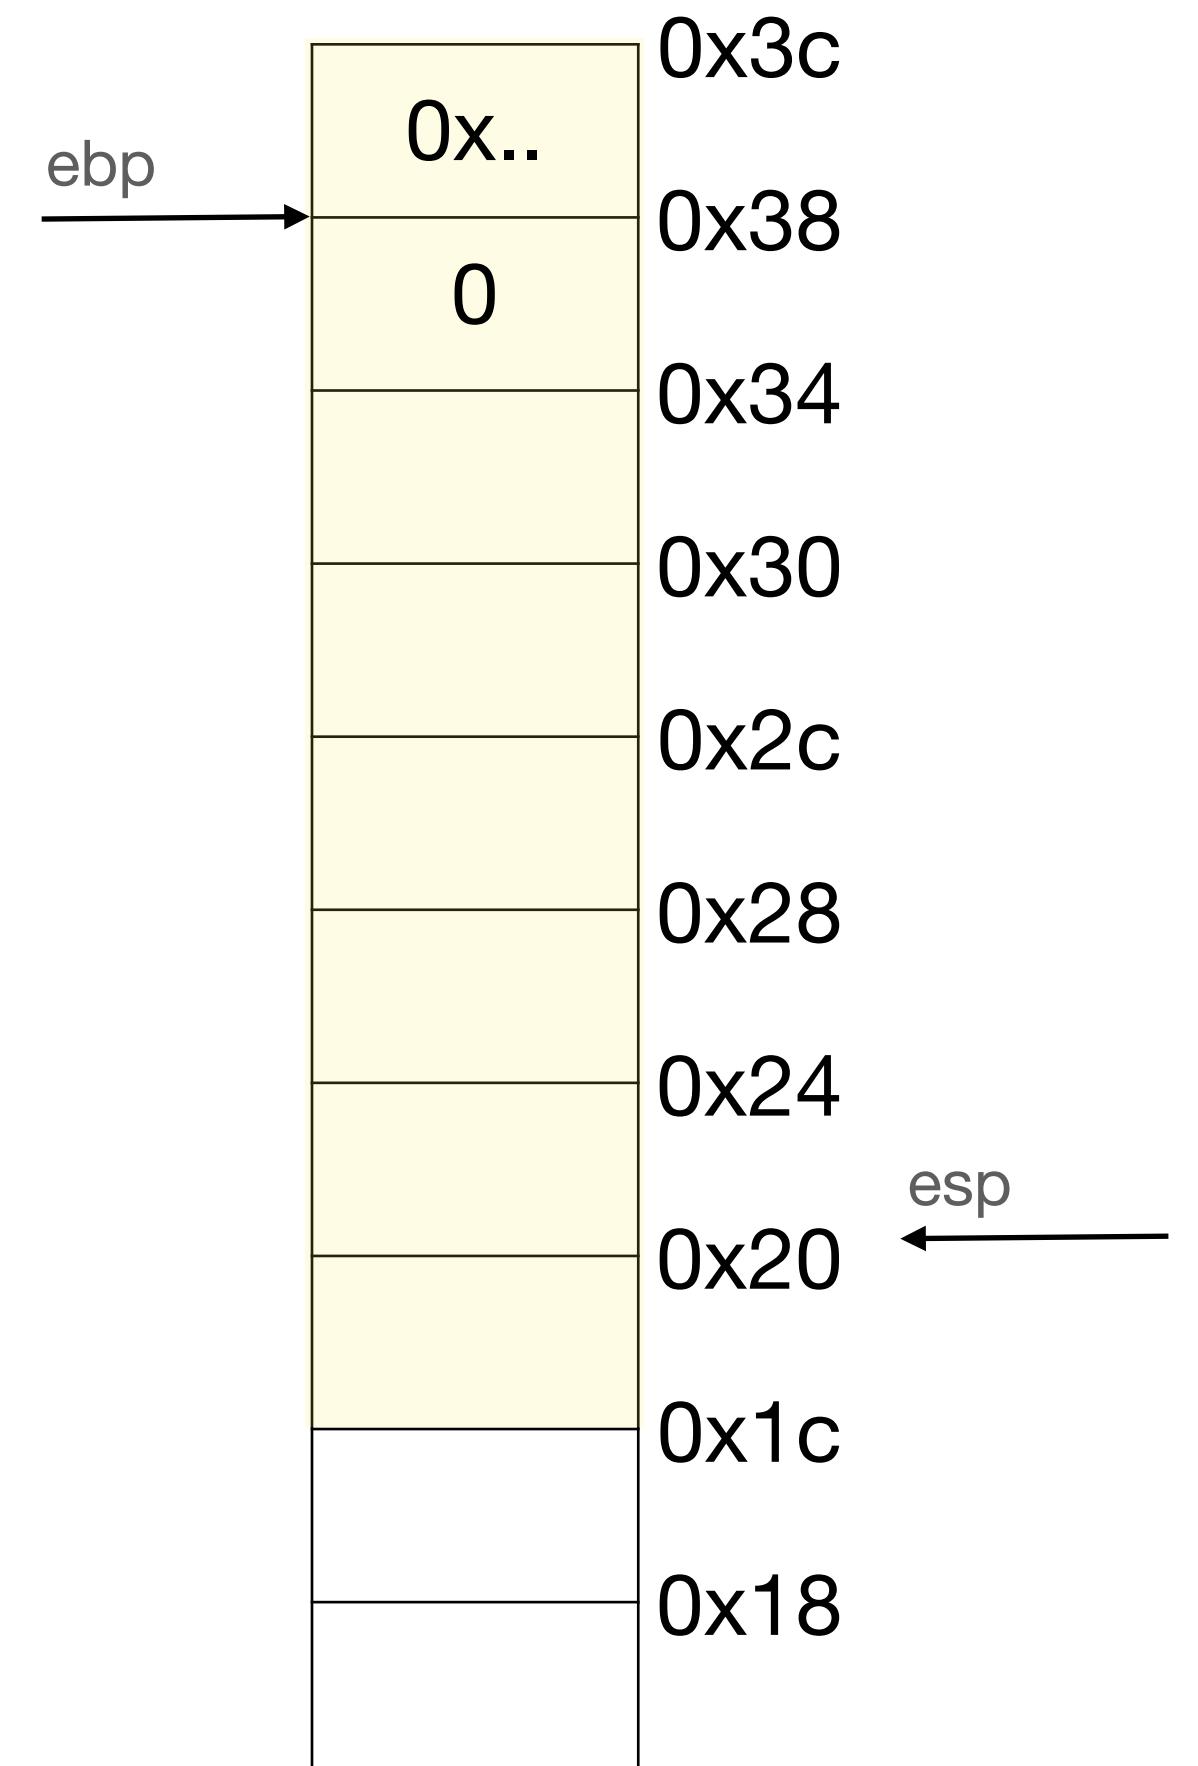
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

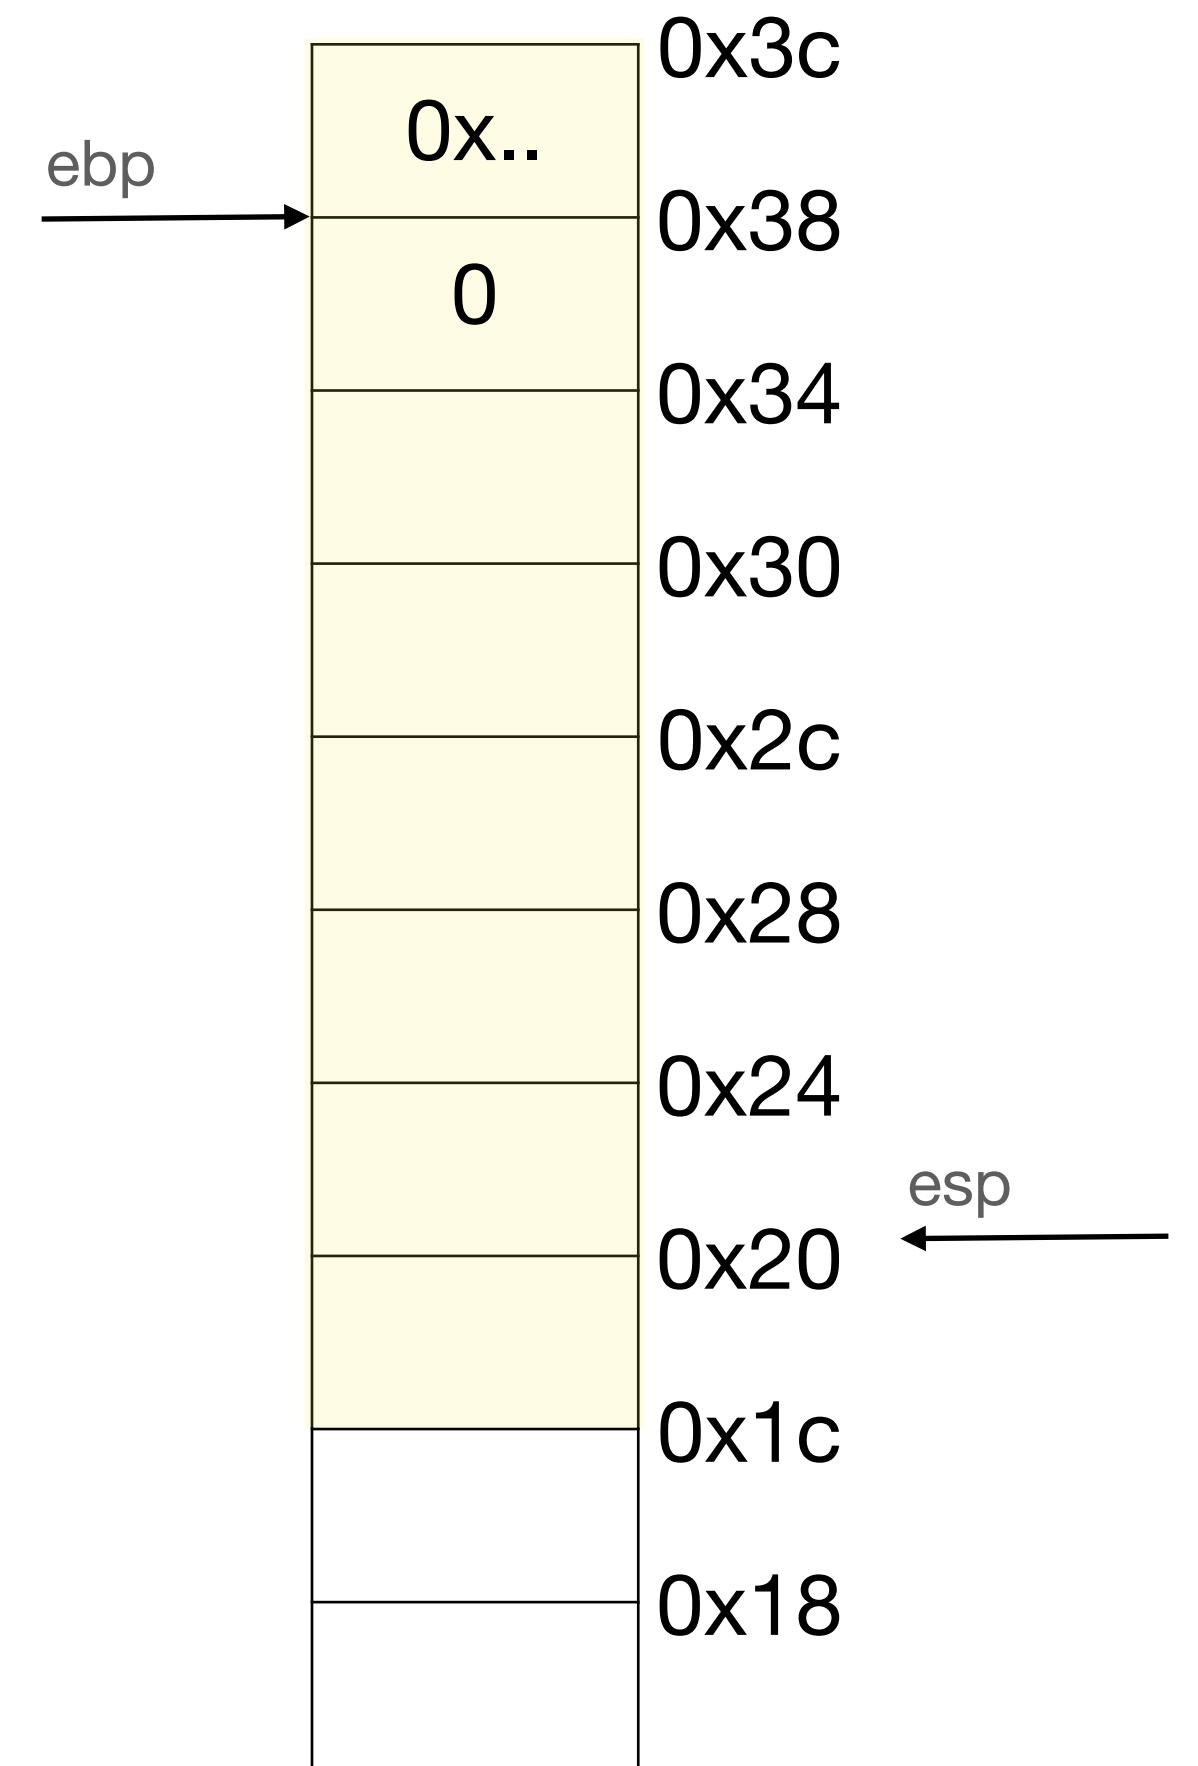
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

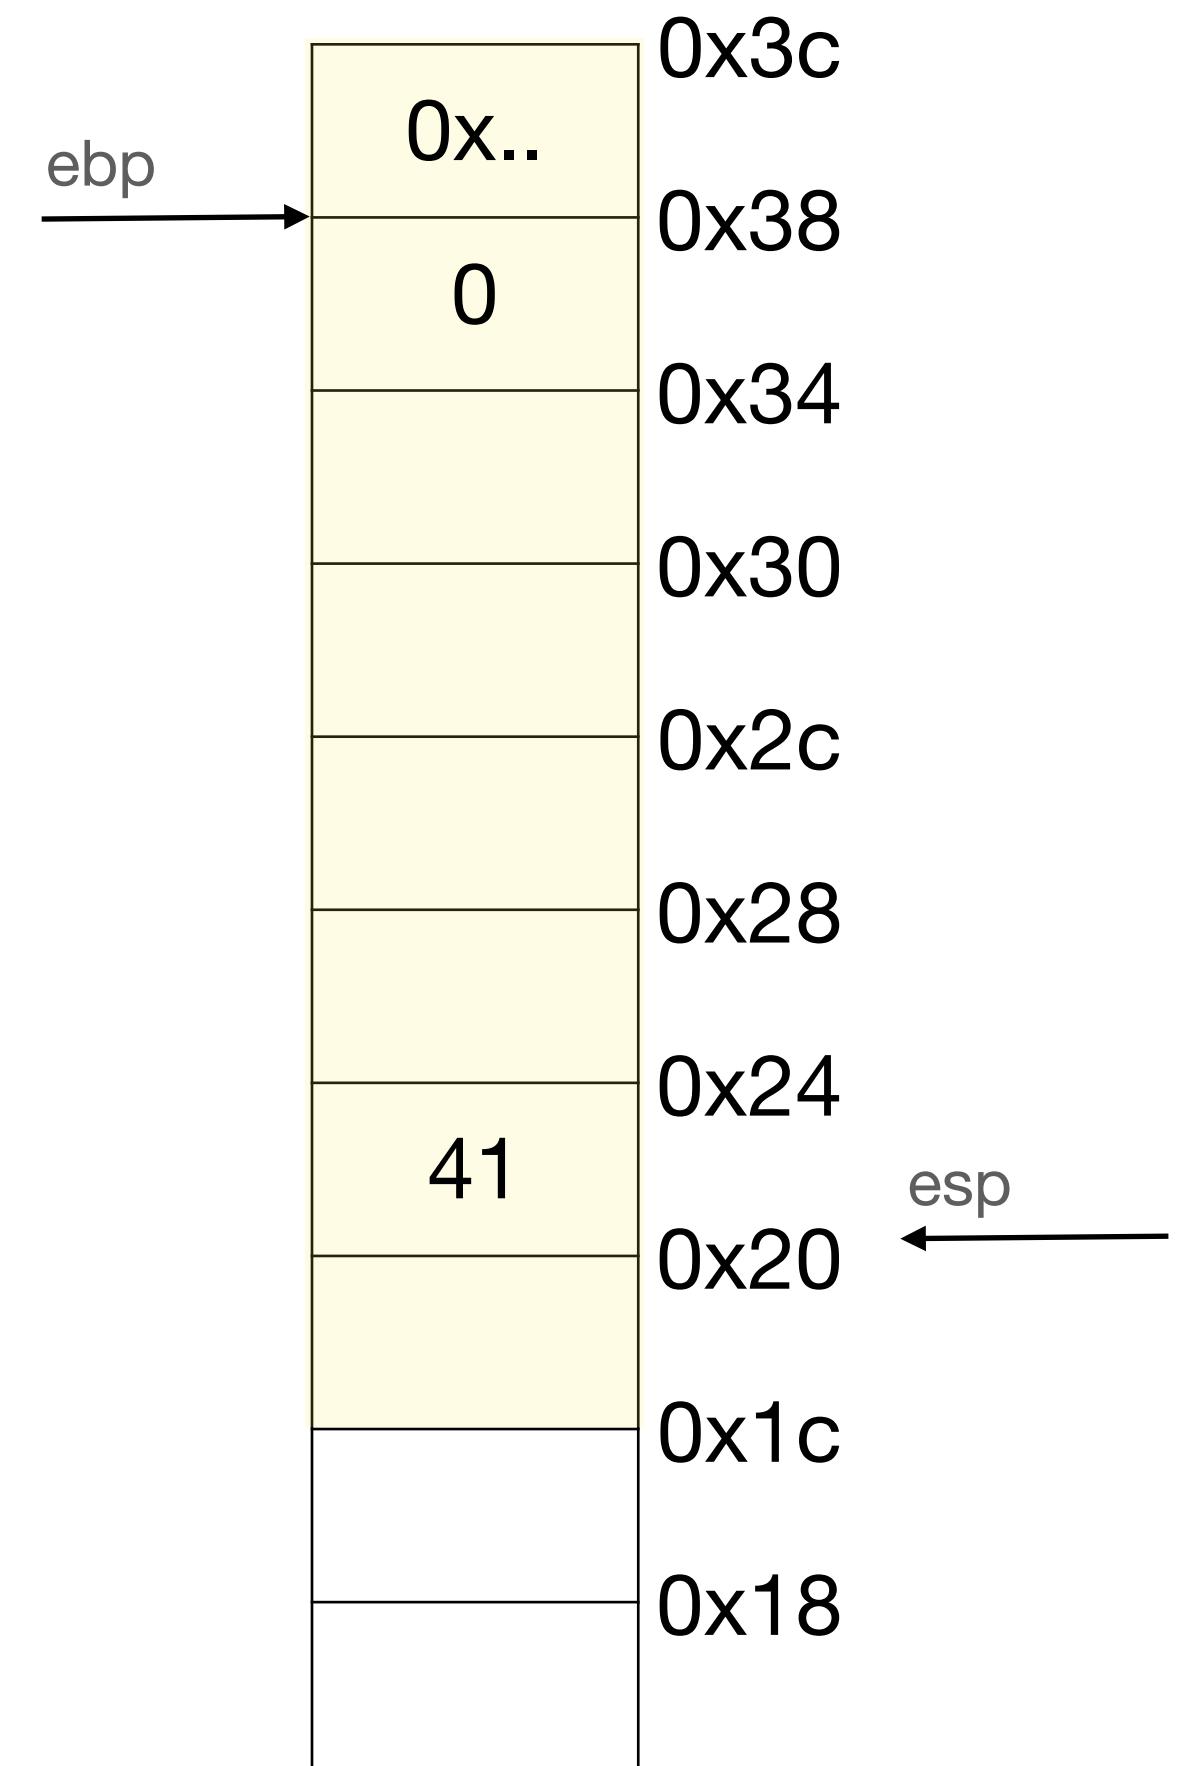
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

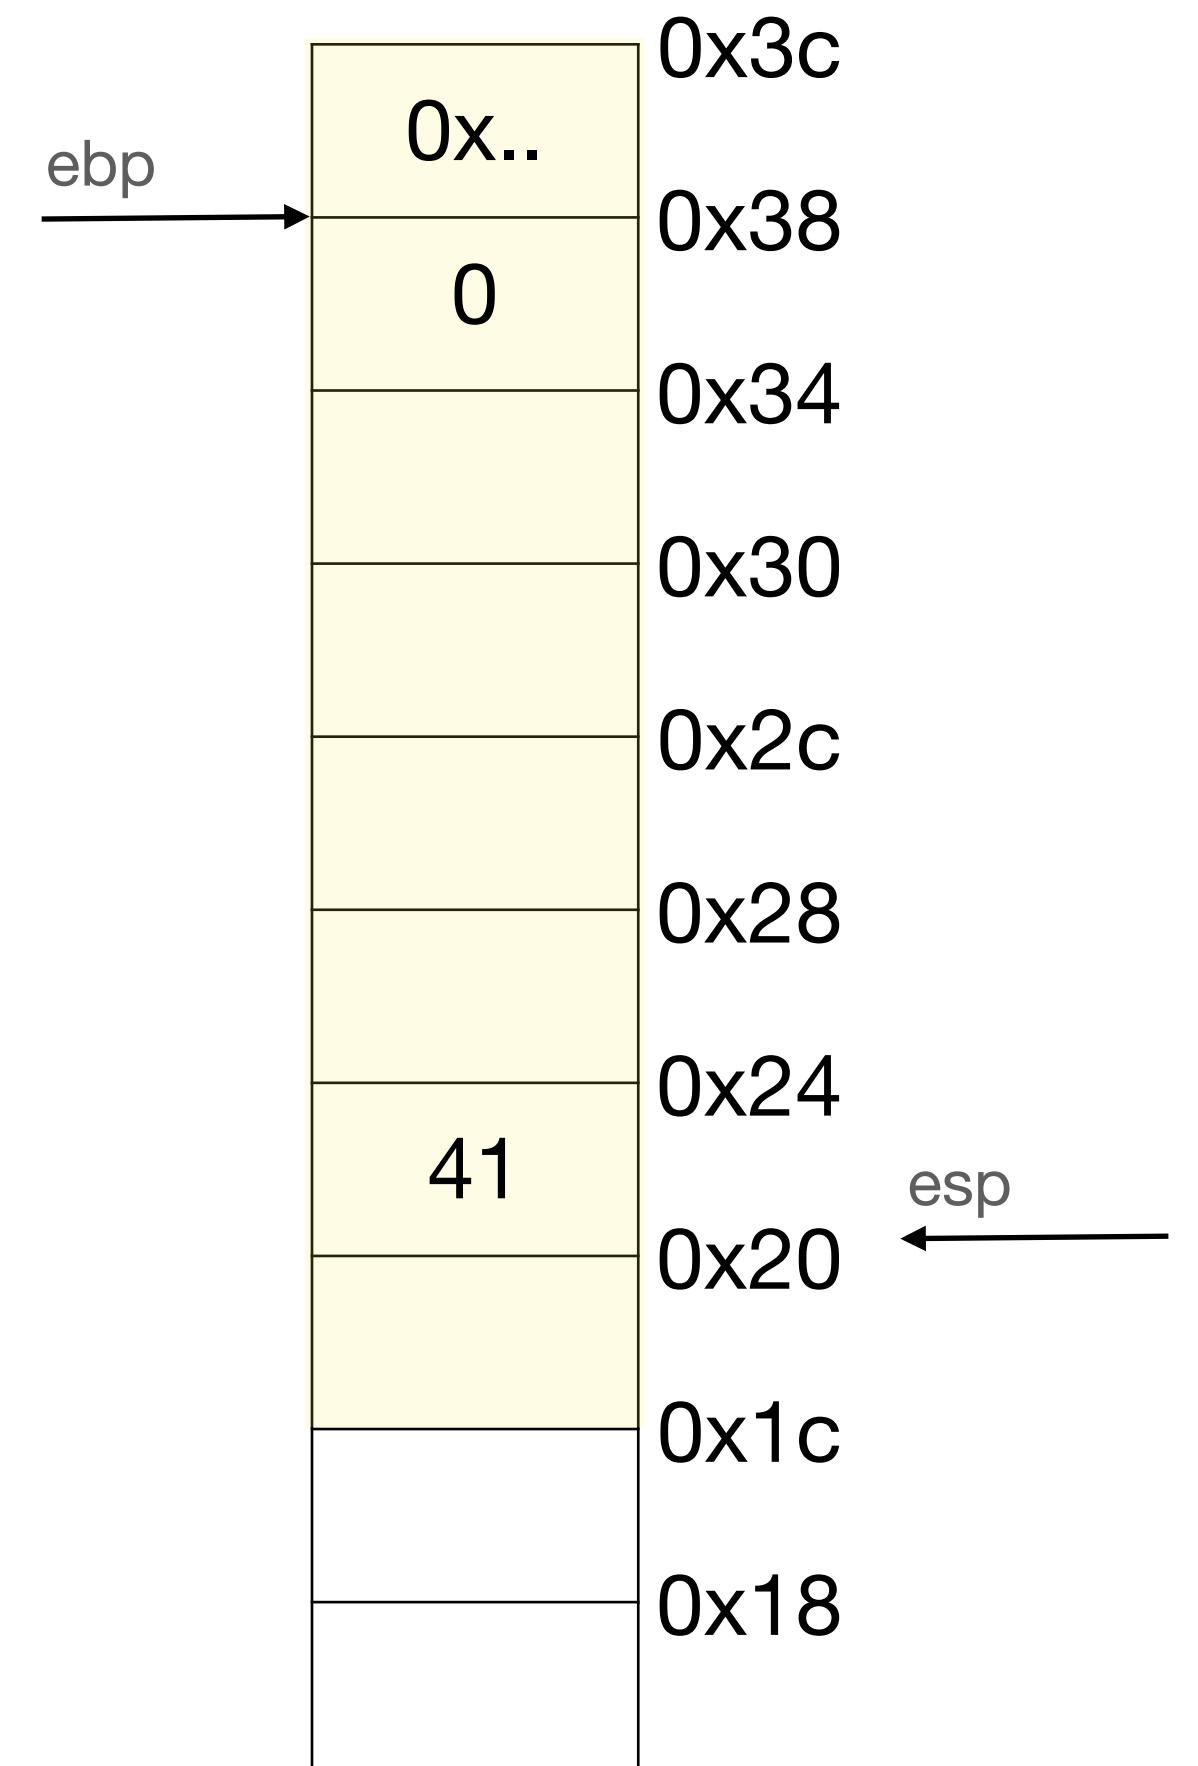
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

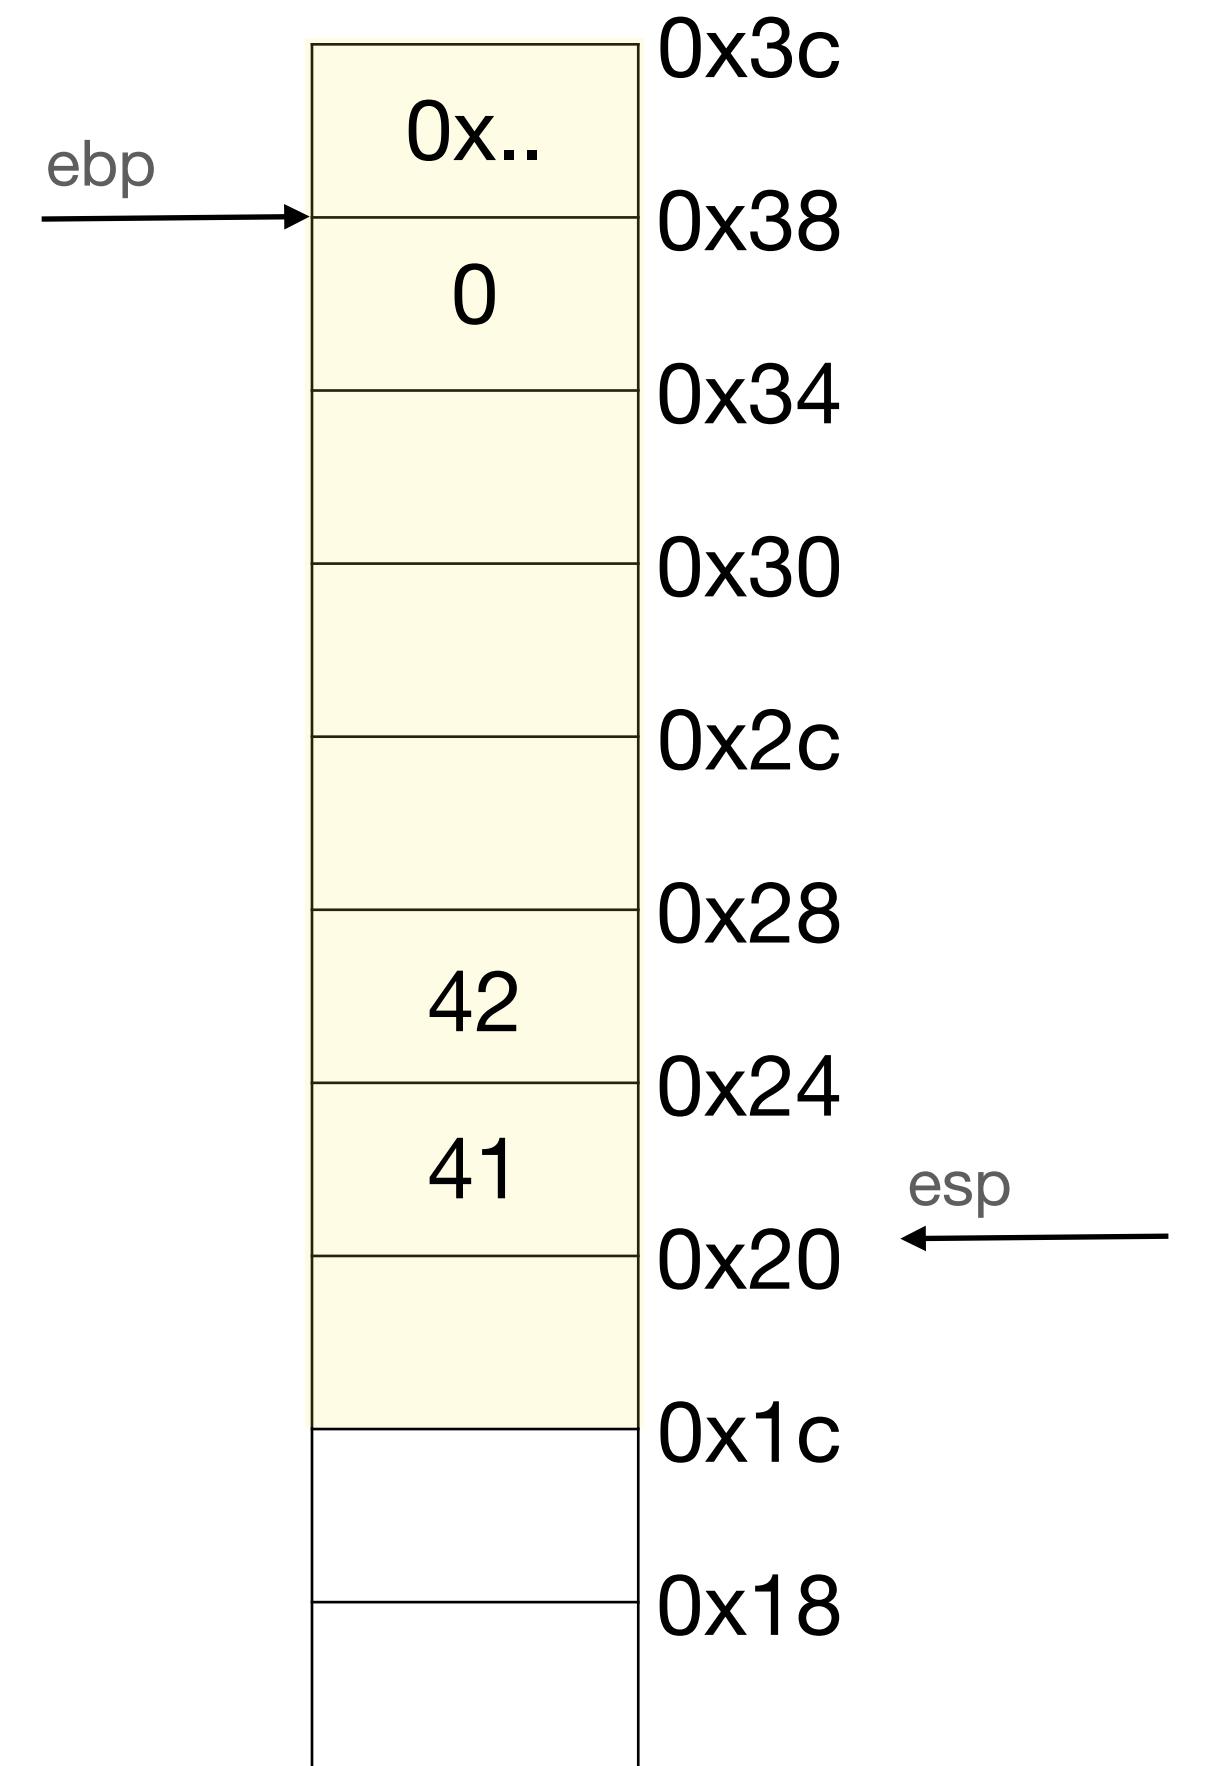
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

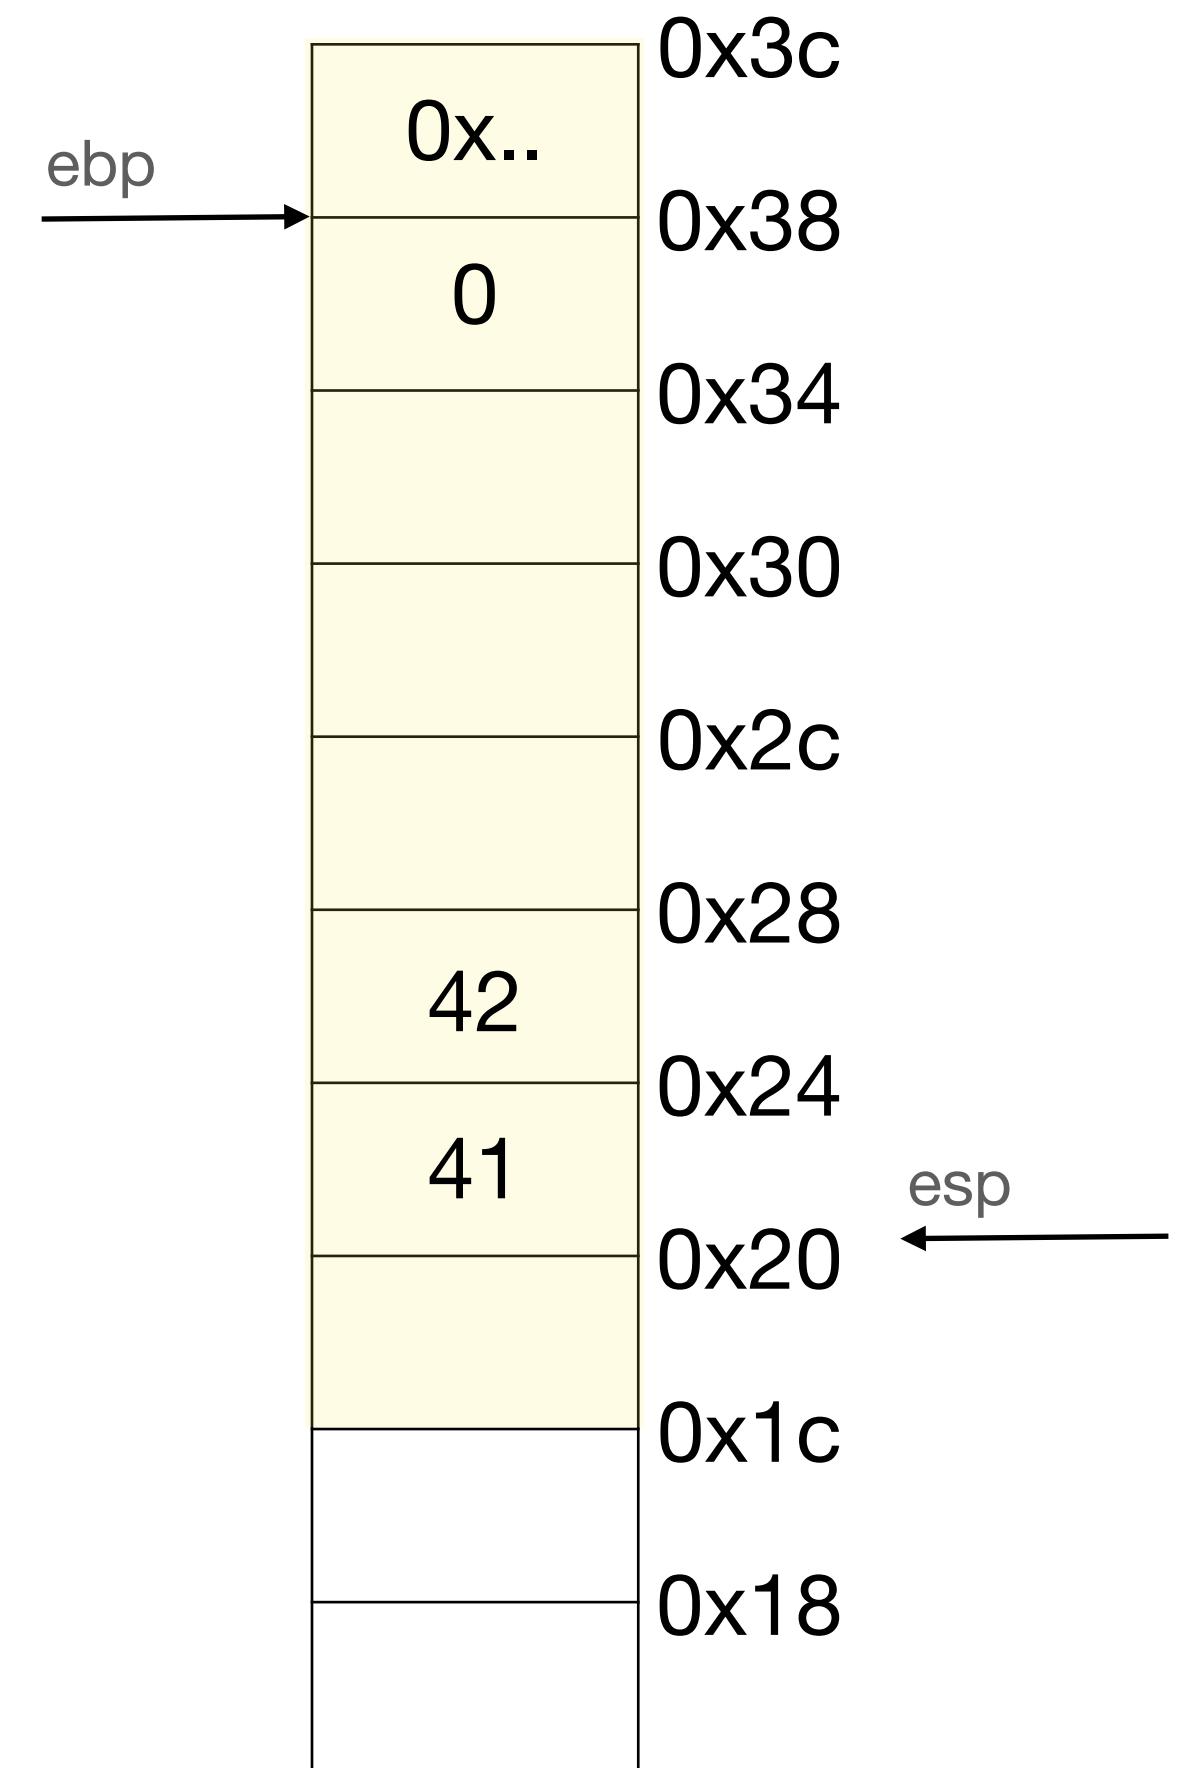
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

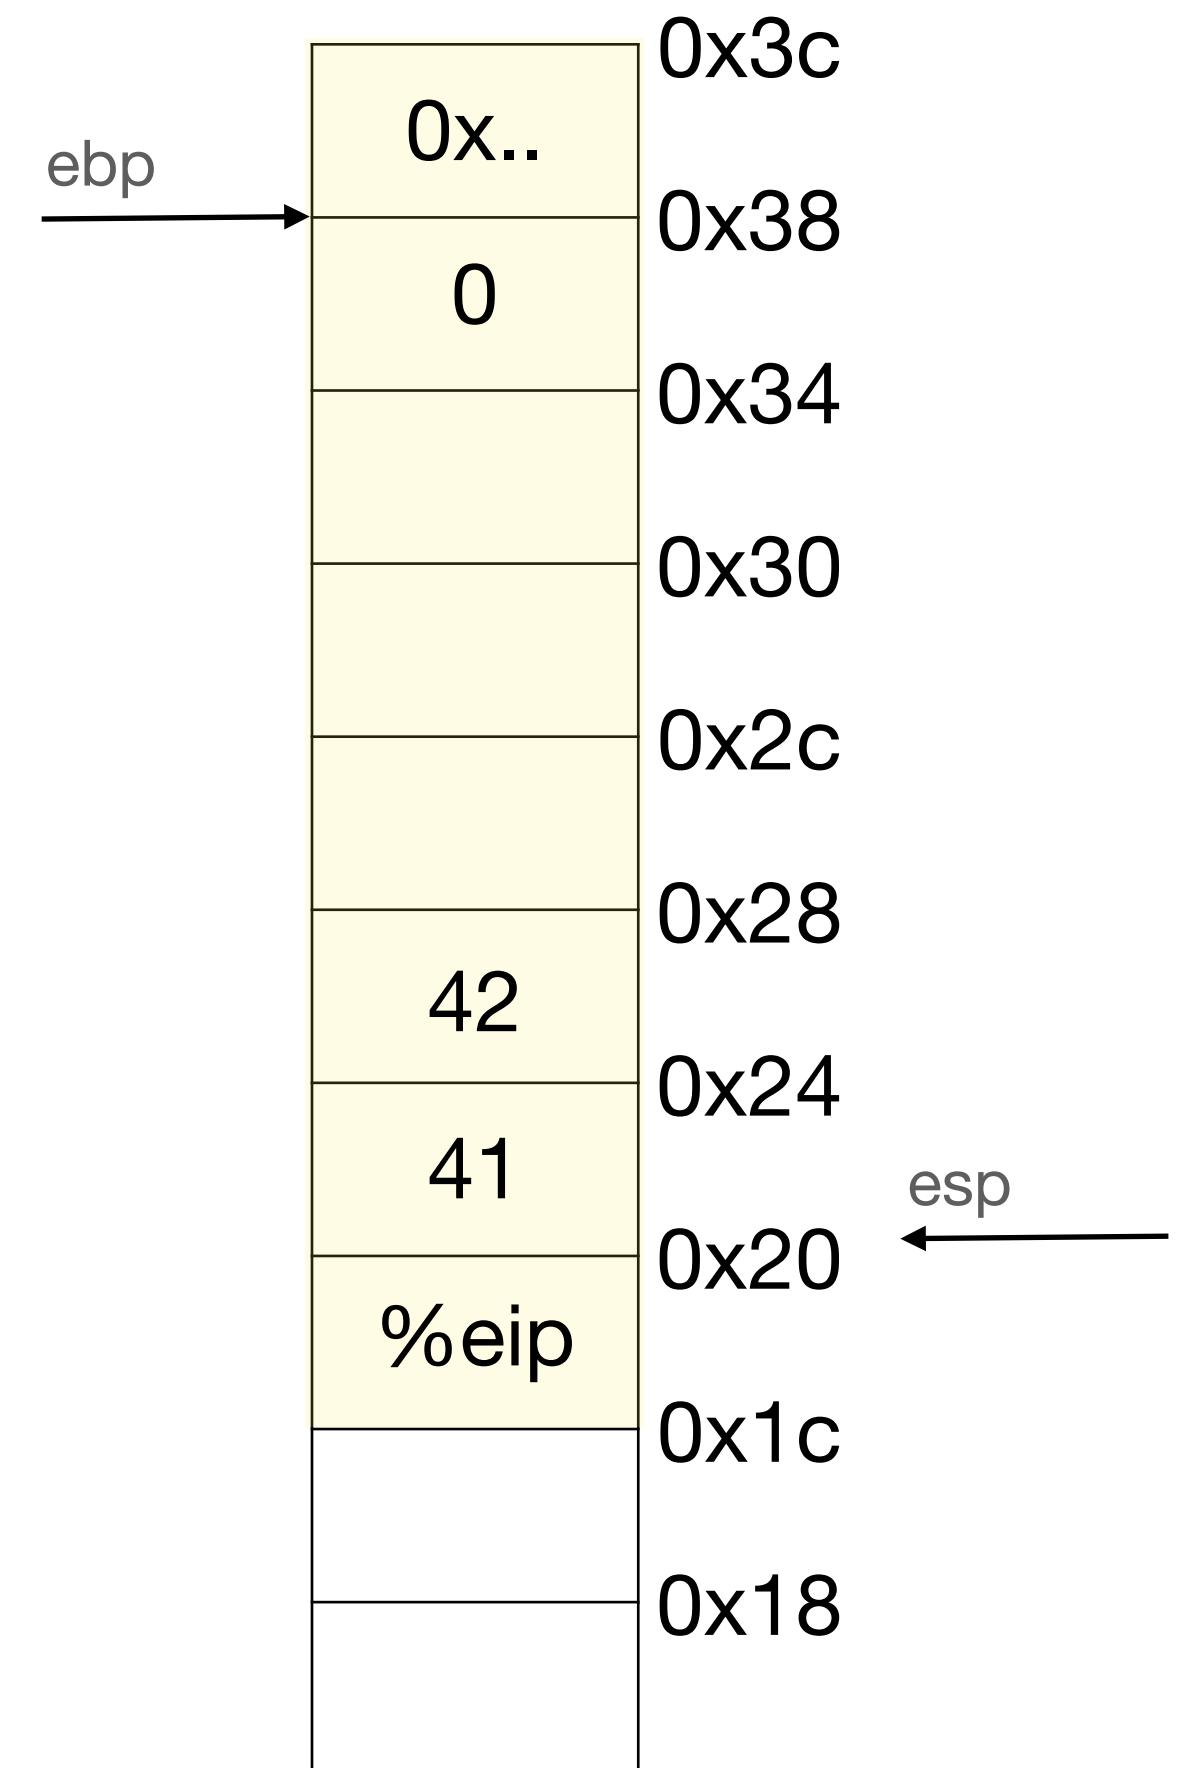
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

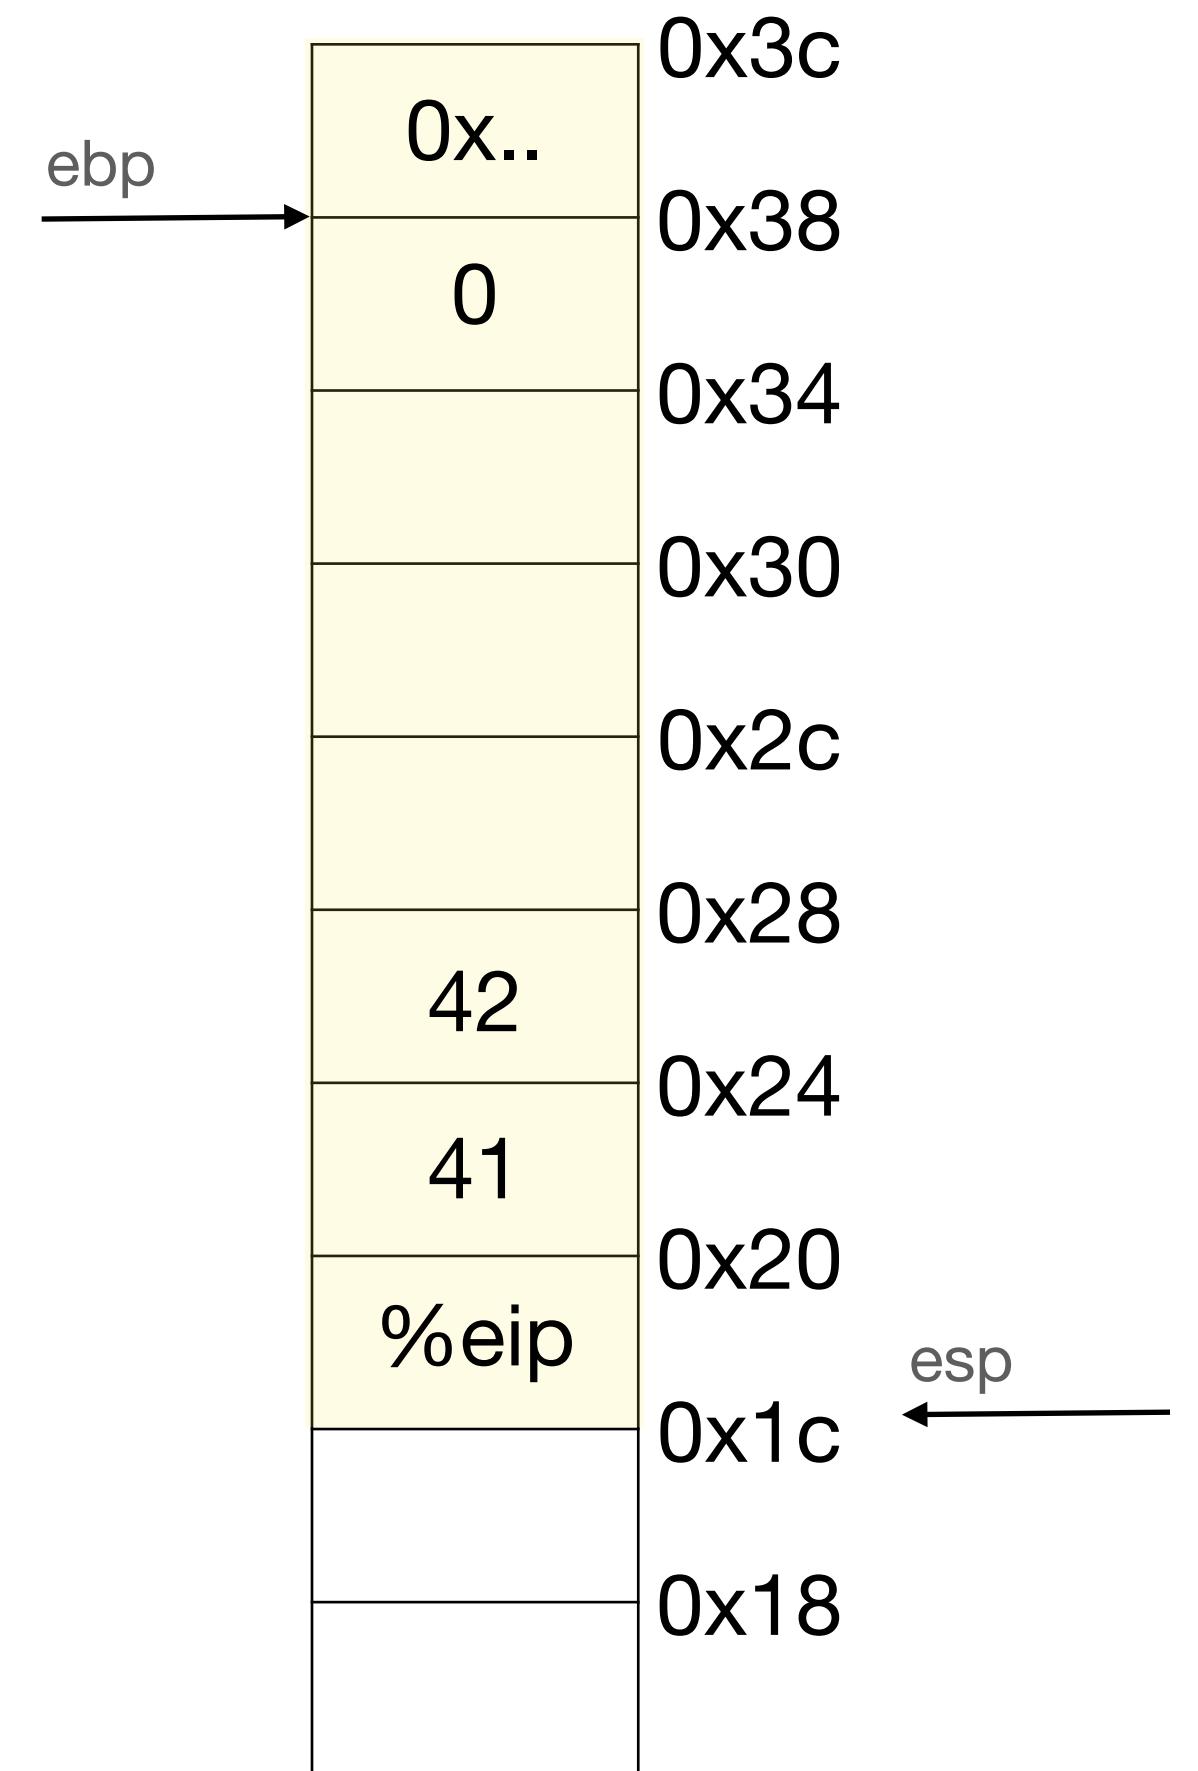
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

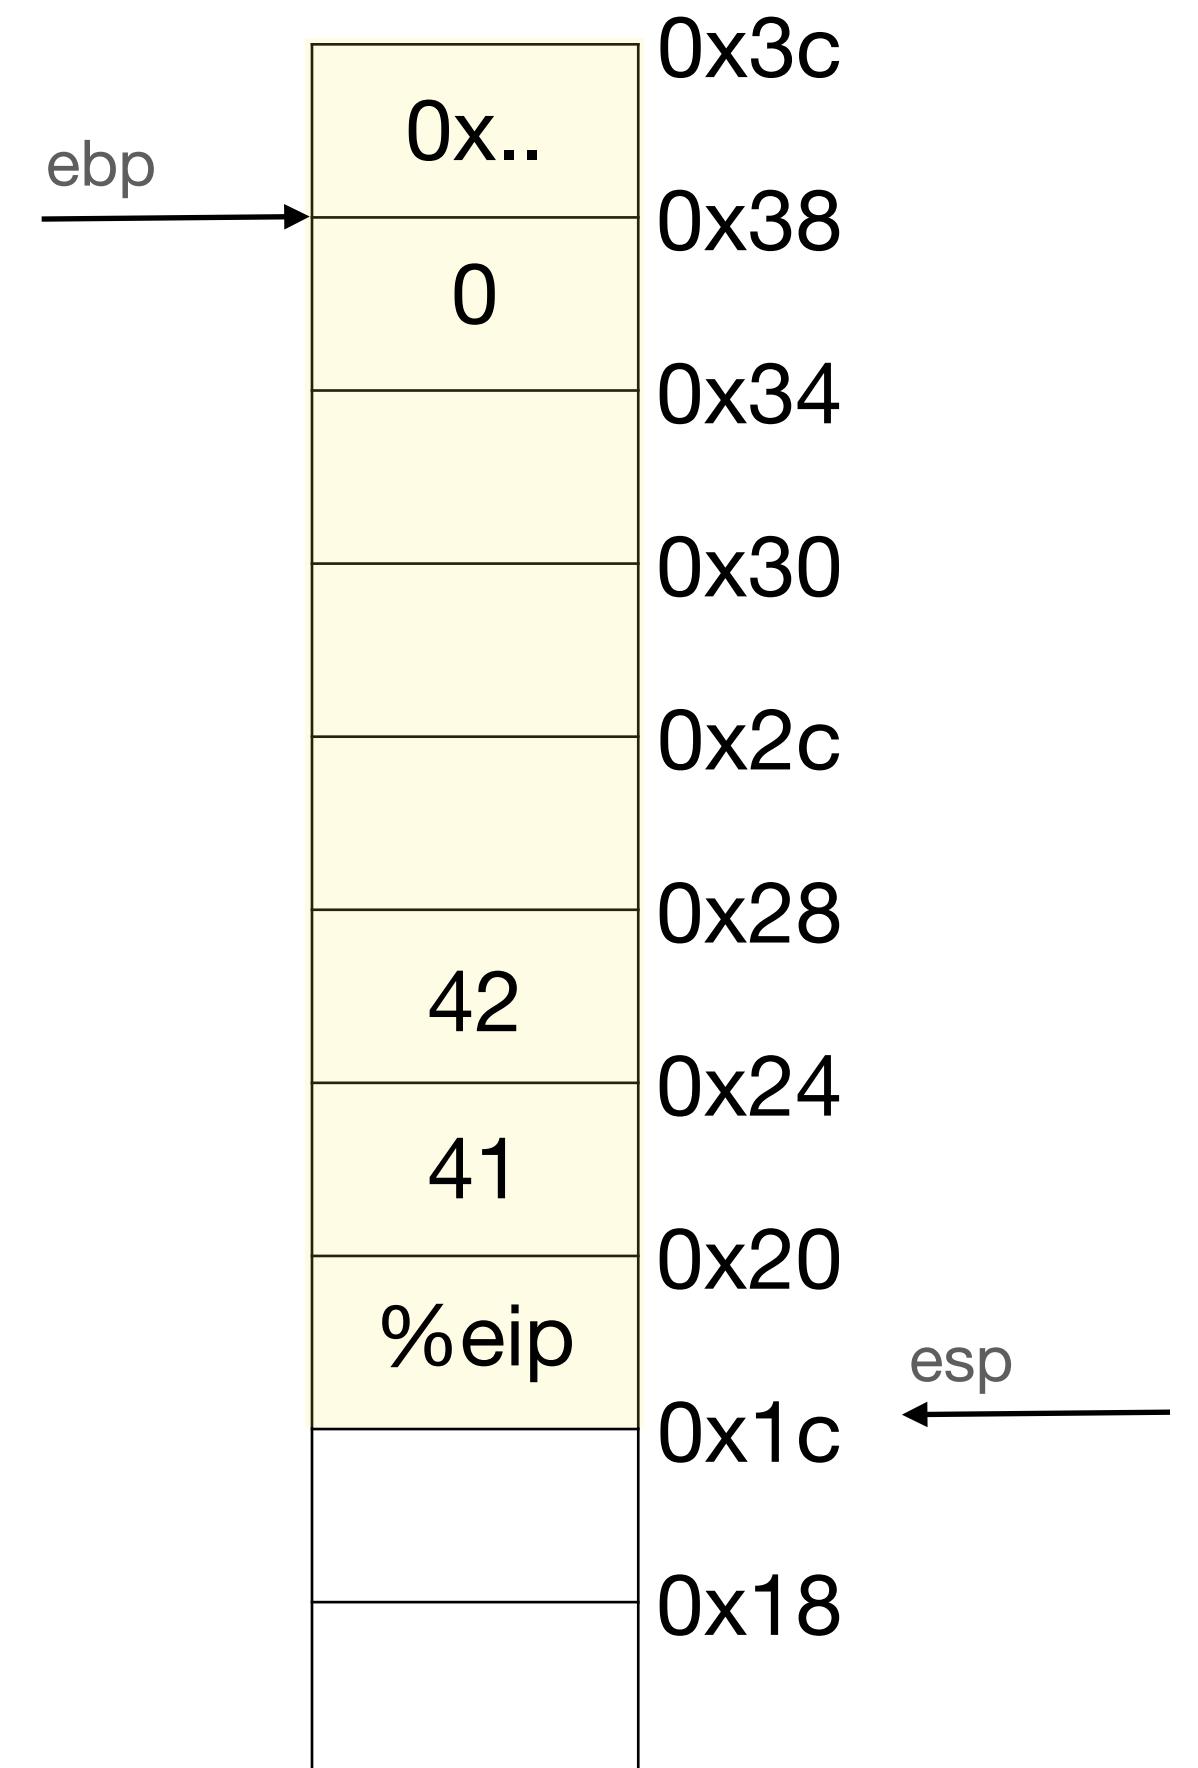
_eip →
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

.globl _main
.p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main

Save caller's base pointer
ebp = esp
eax = *(ebp + 8)
eax = eax + *(ebp + 12)
Restore caller's base pointer
change eip to return address

Save caller's base pointer
ebp = esp
esp = esp - 0x18
*(ebp-4)=0
*(esp) = 41
*(esp+4) = 42
Push current eip on to stack, jump to foo
esp = esp + 24 (Restore caller's esp)
Restore caller's ebp
```



# Function calling in action

## Stack

```
02.s

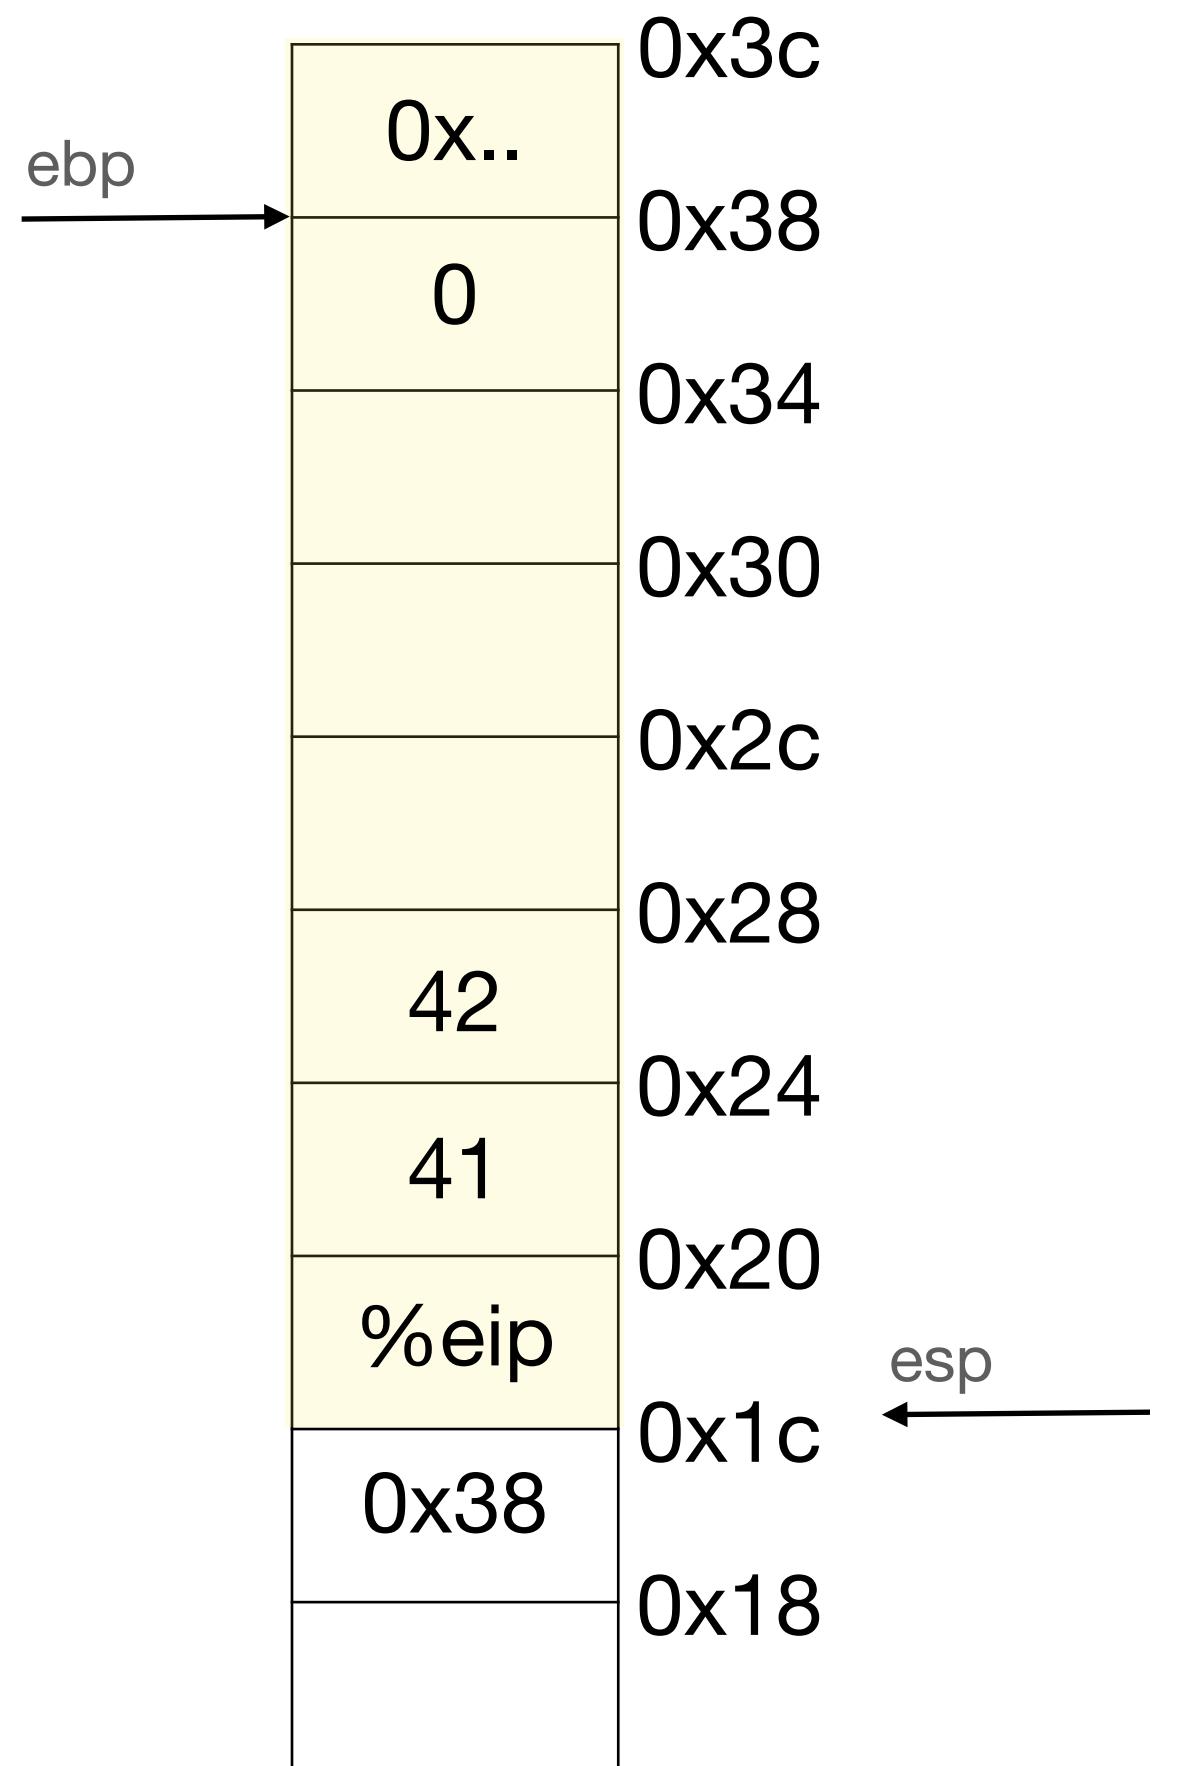
_eip →
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

.globl _main
.p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main

Save caller's base pointer
ebp = esp
eax = *(ebp + 8)
eax = eax + *(ebp + 12)
Restore caller's base pointer
change eip to return address

Save caller's base pointer
ebp = esp
esp = esp - 0x18
*(ebp-4)=0
*(esp) = 41
*(esp+4) = 42
Push current eip on to stack, jump to foo
esp = esp + 24 (Restore caller's esp)
Restore caller's ebp
```



# Function calling in action

## Stack

```
02.s

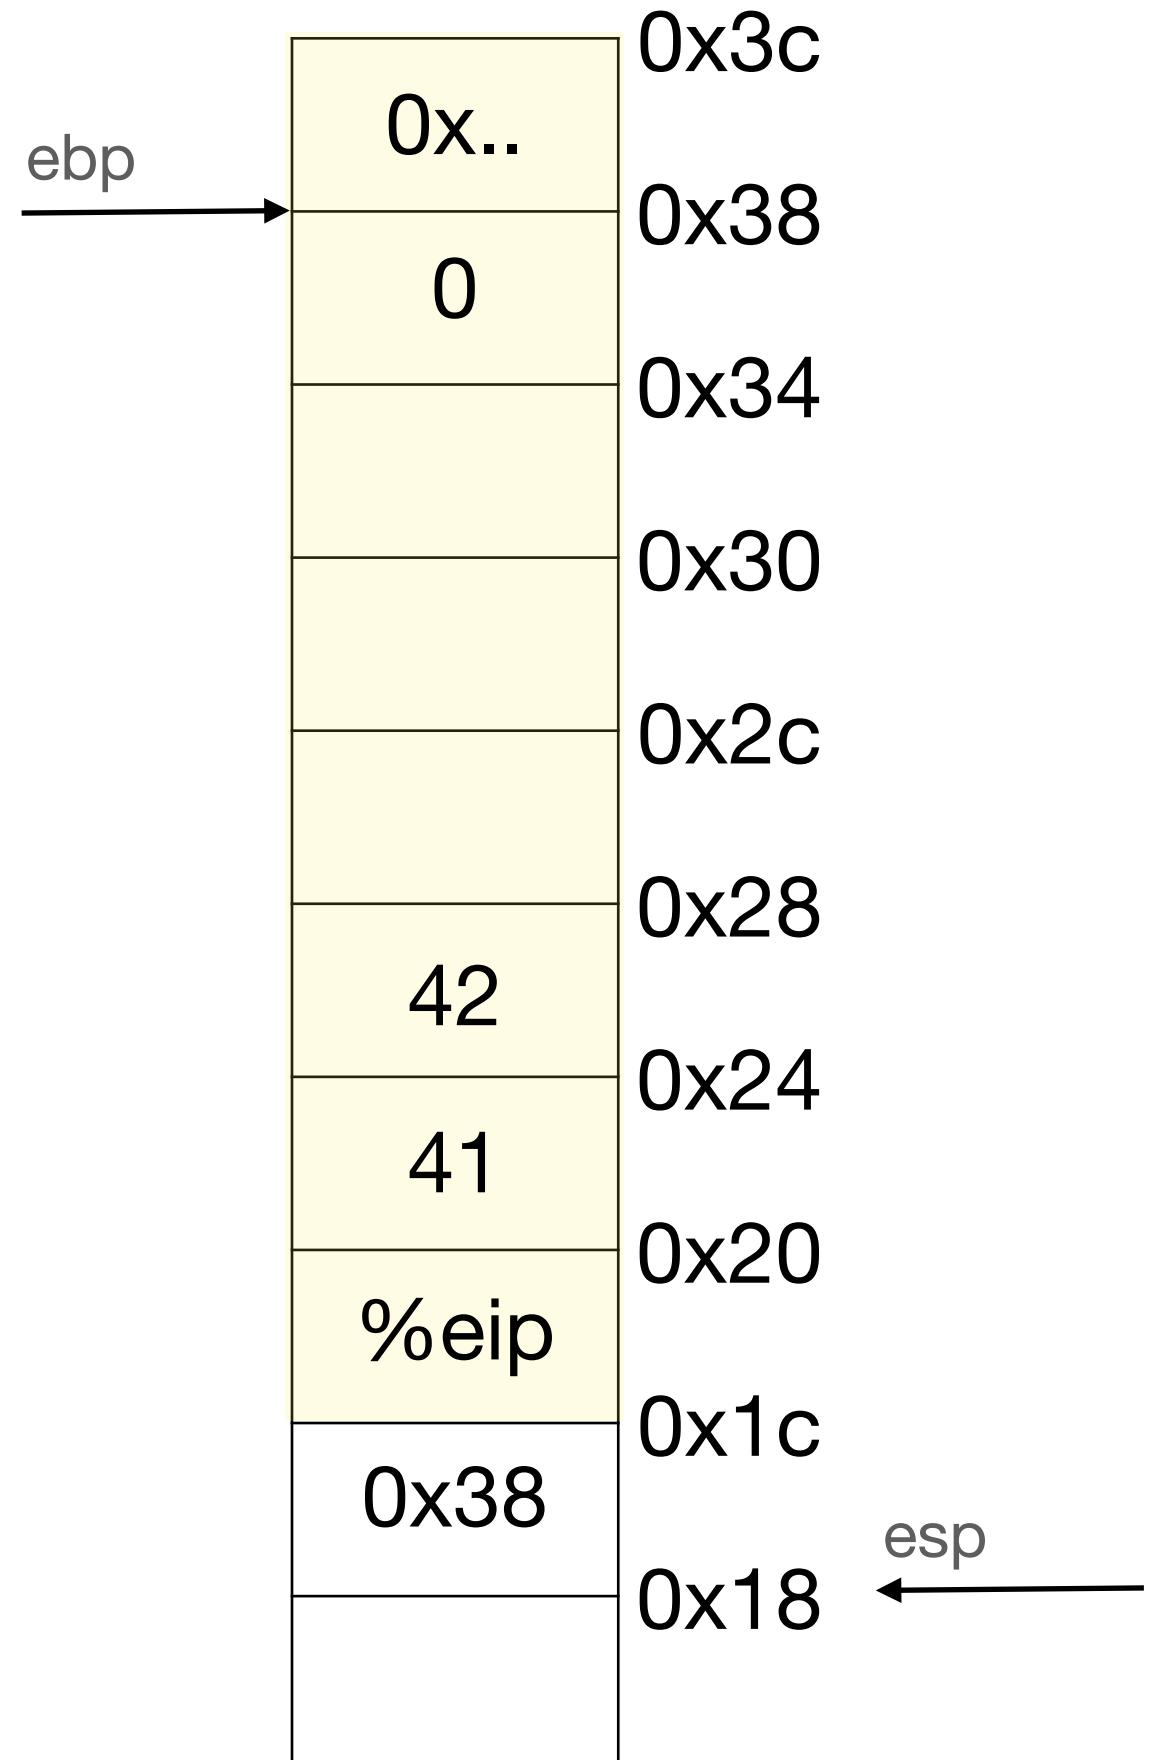
_eip →
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

.globl _main
.p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main

Save caller's base pointer
ebp = esp
eax = *(ebp + 8)
eax = eax + *(ebp + 12)
Restore caller's base pointer
change eip to return address

Save caller's base pointer
ebp = esp
esp = esp - 0x18
*(ebp-4)=0
*(esp) = 41
*(esp+4) = 42
Push current eip on to stack, jump to foo
esp = esp + 24 (Restore caller's esp)
Restore caller's ebp
```



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

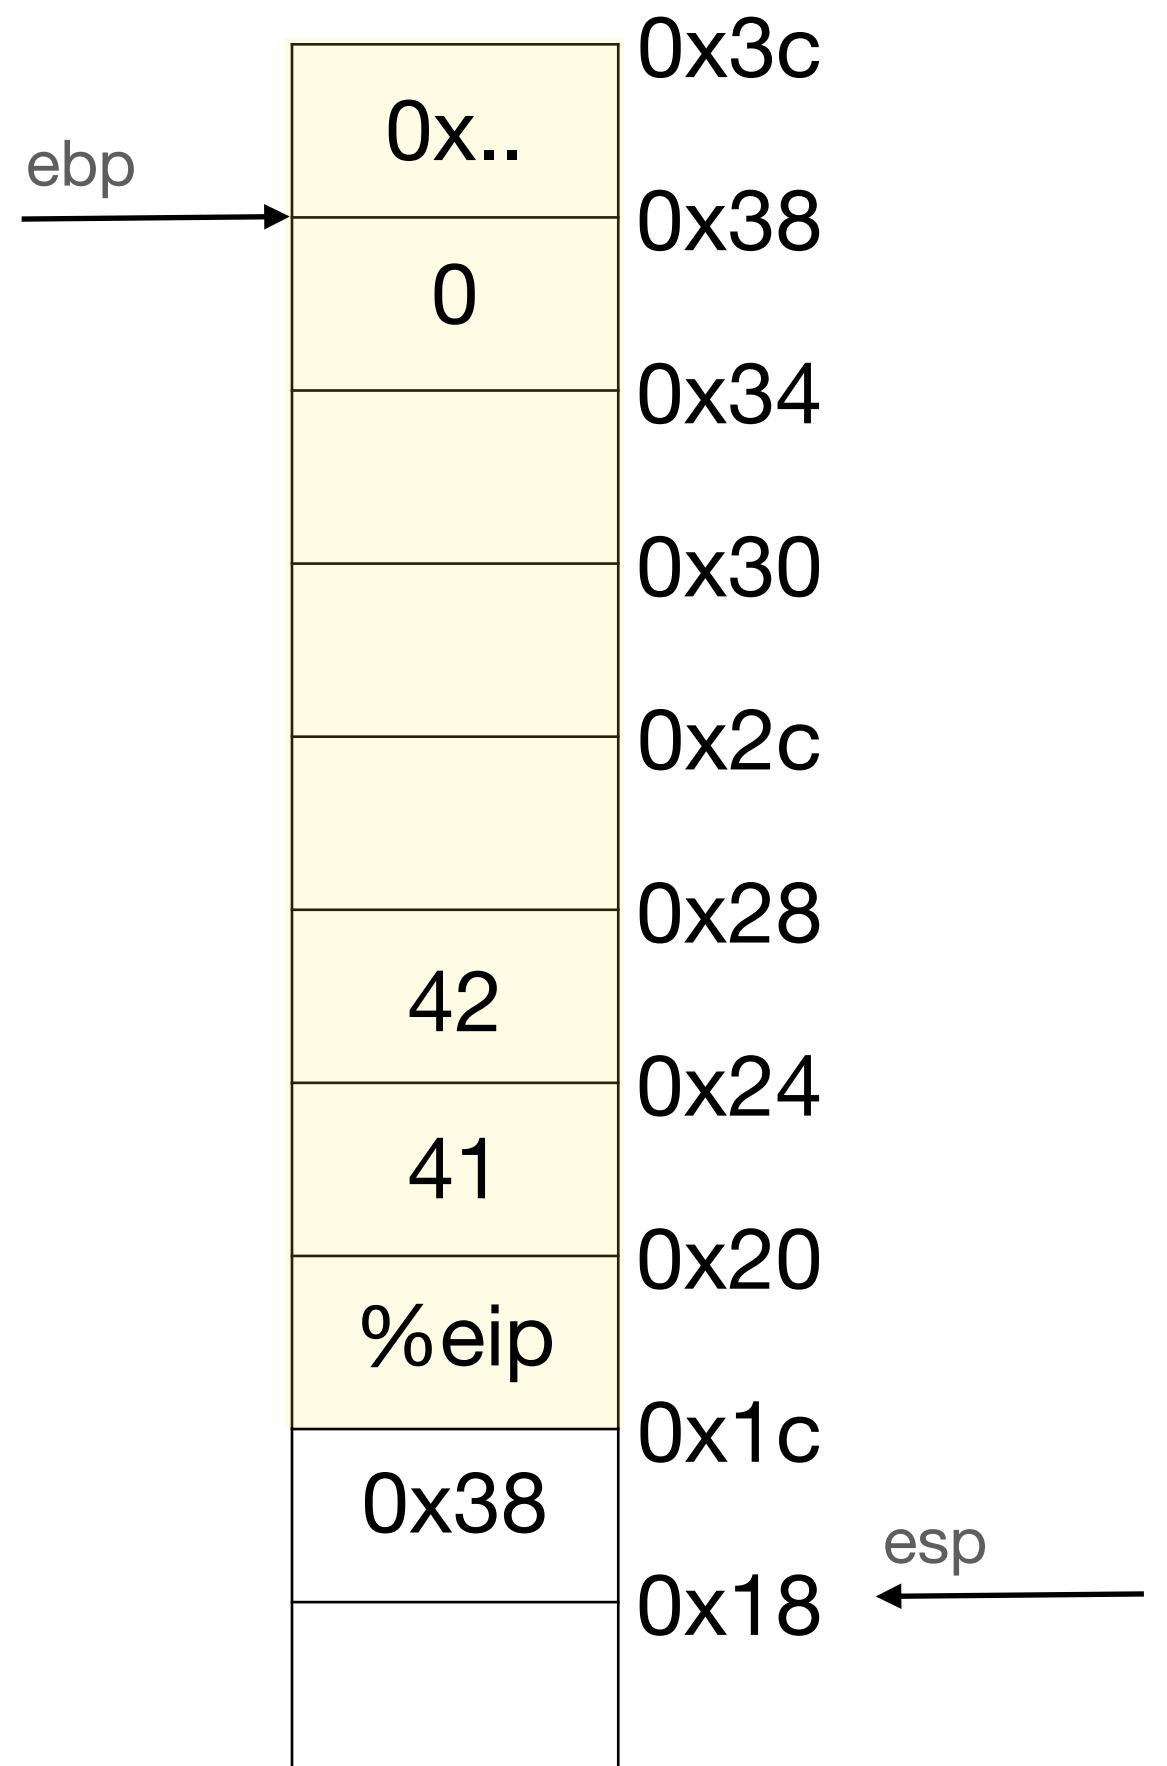
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

eip →

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

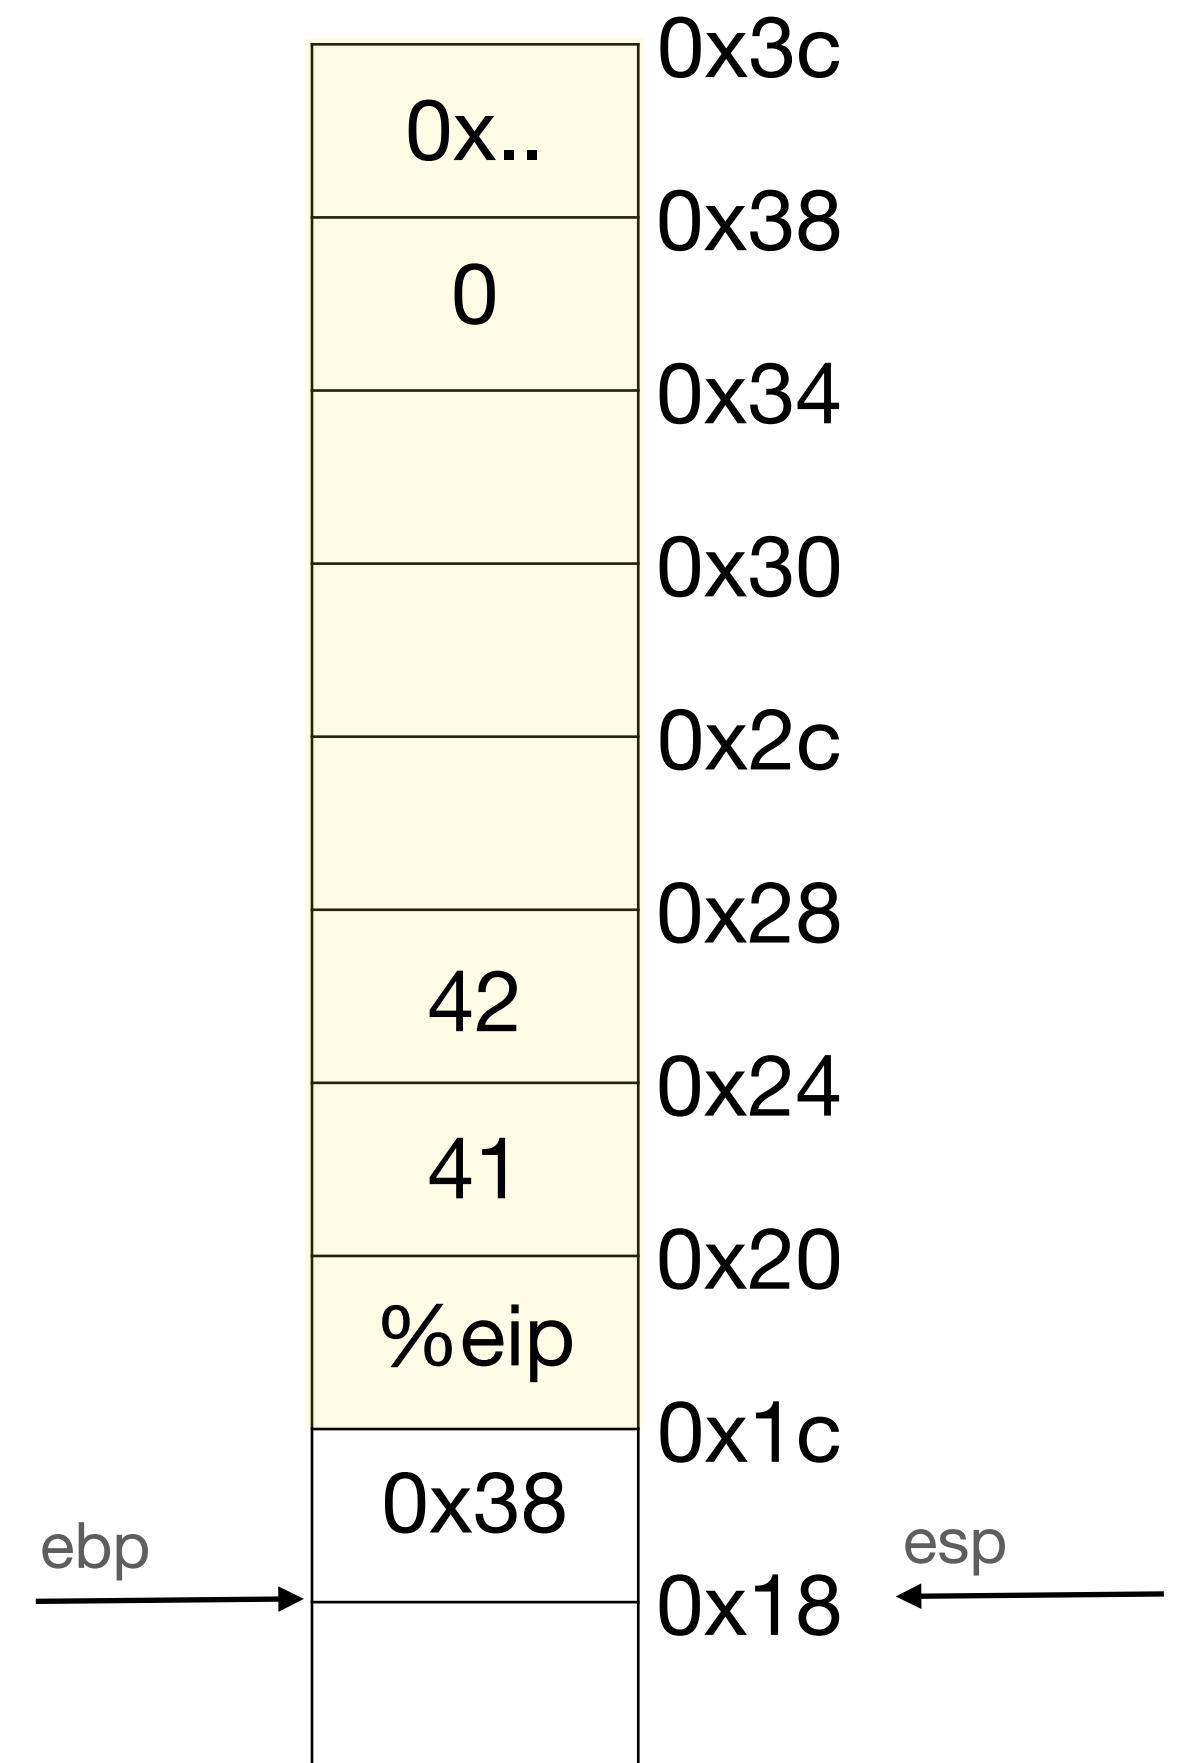
## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```



# Function calling in action

## Stack

```
02.s

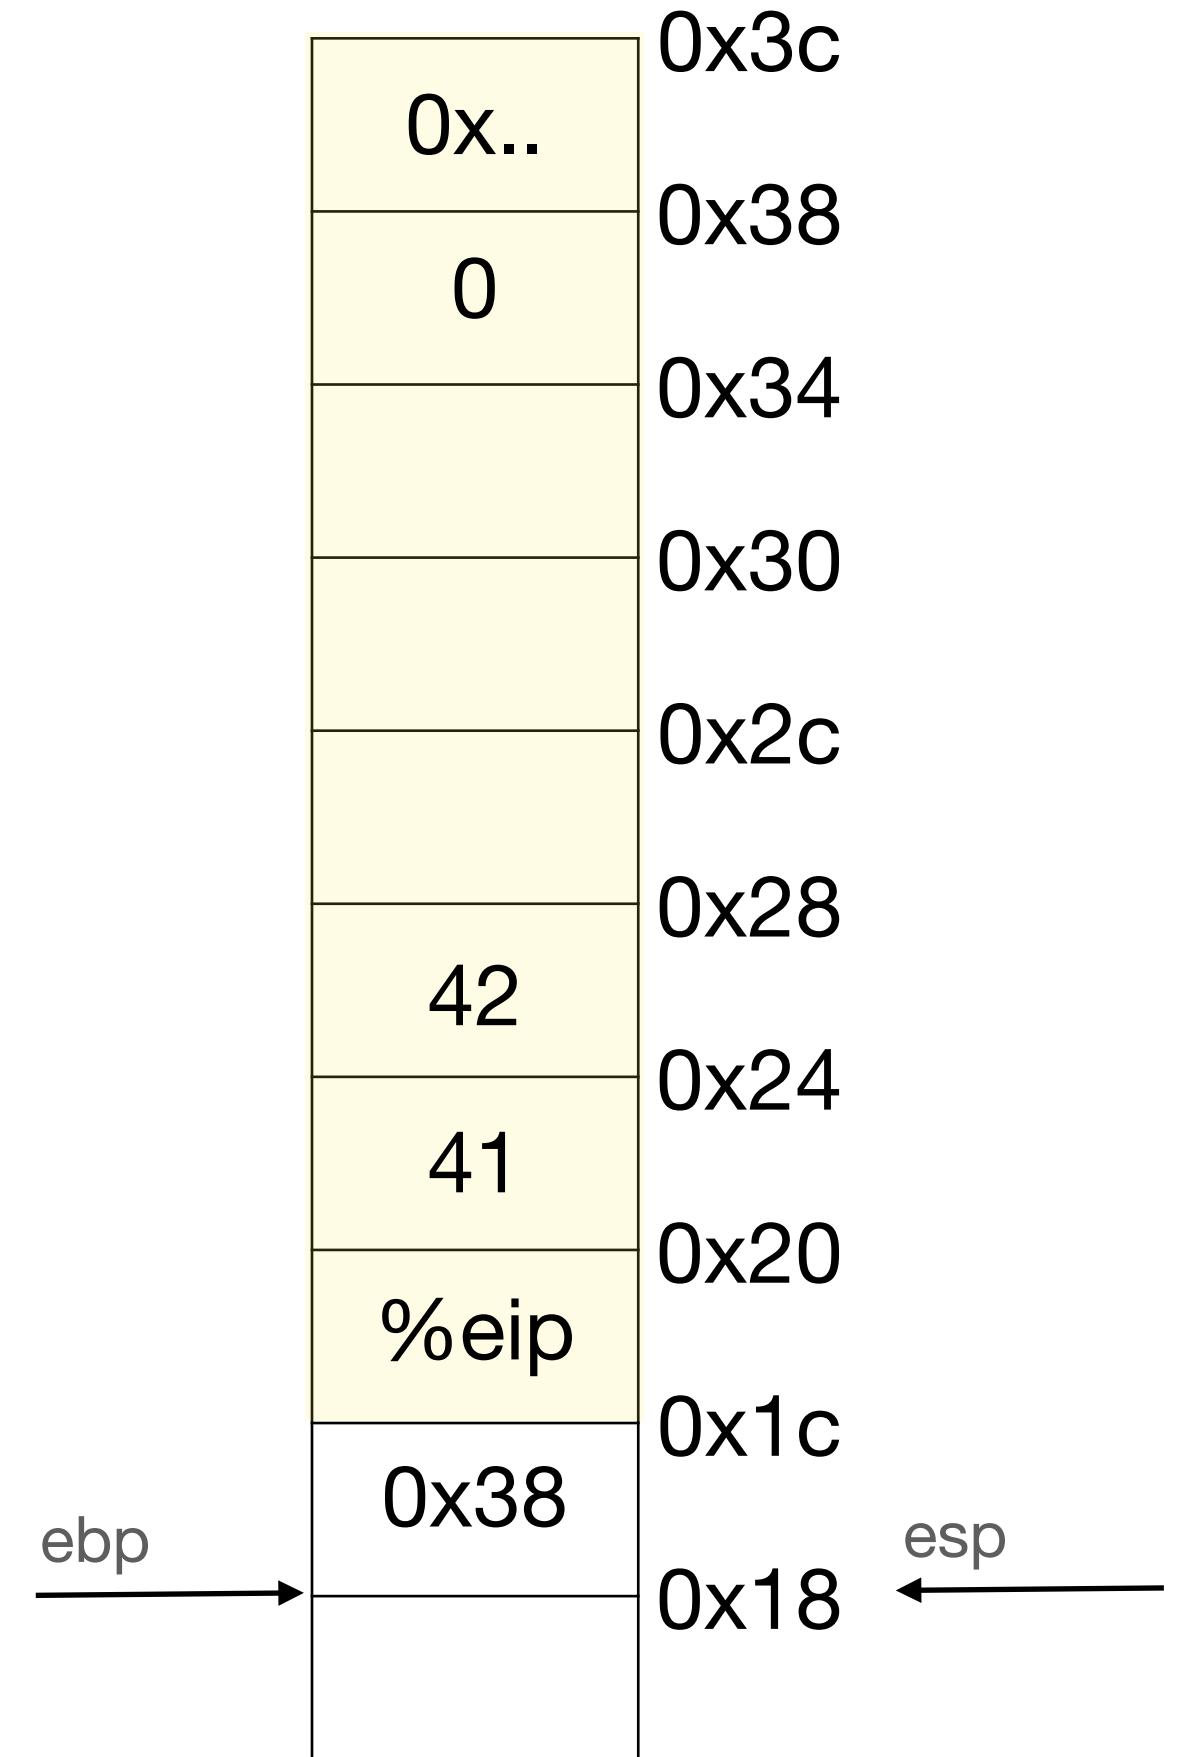
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

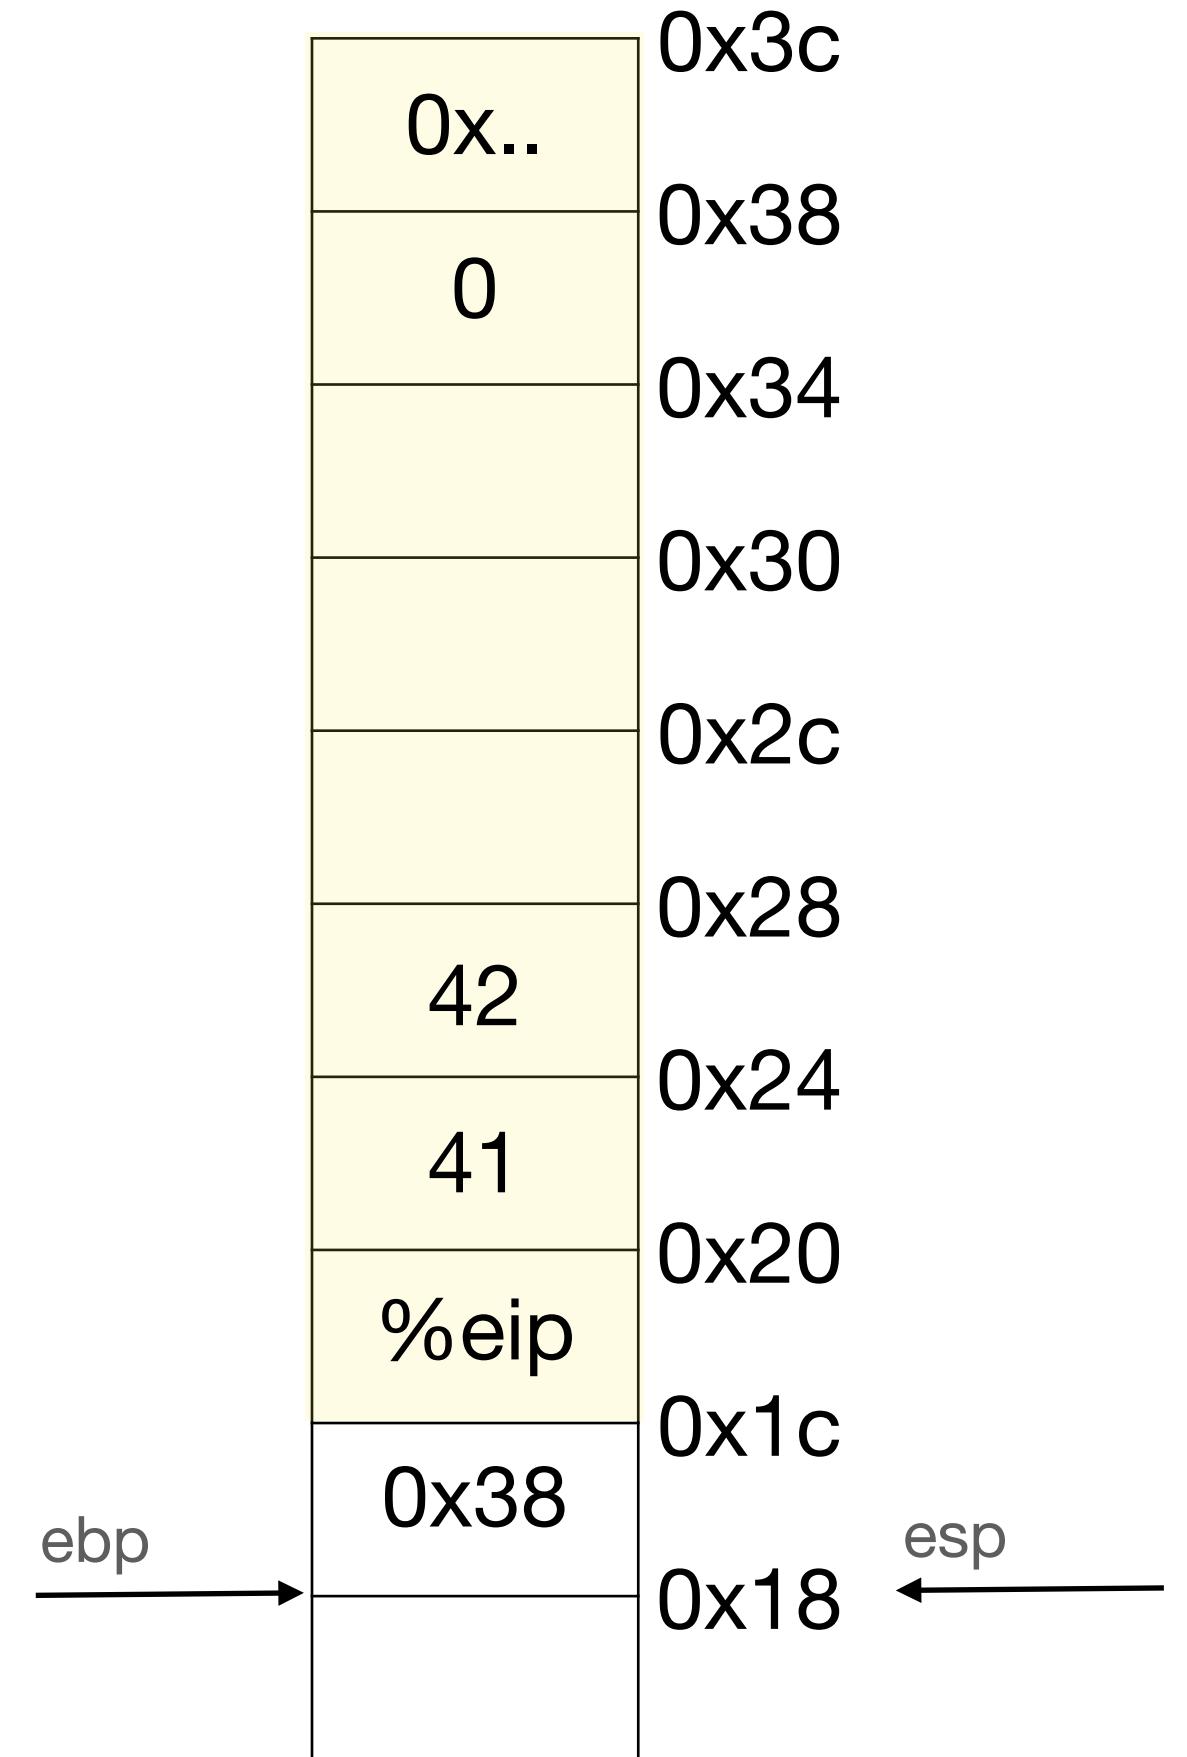
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

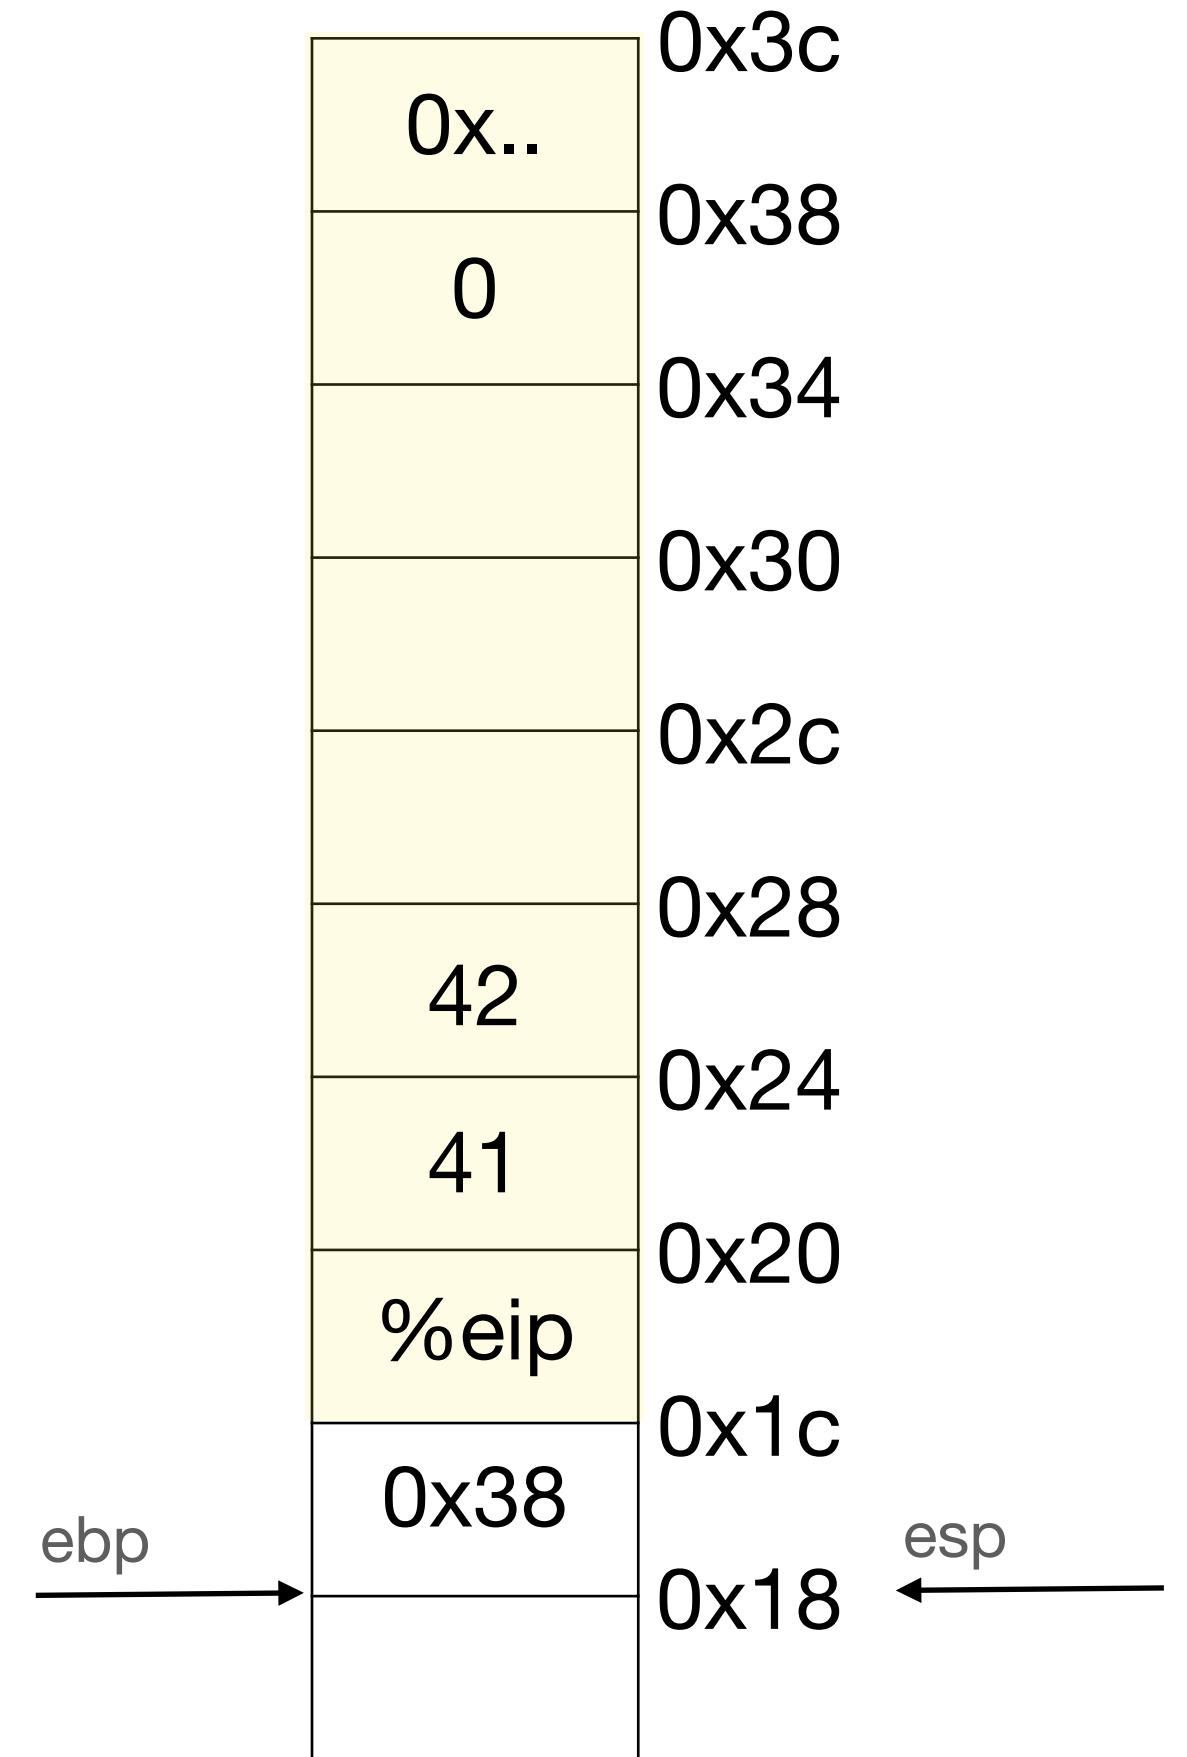
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

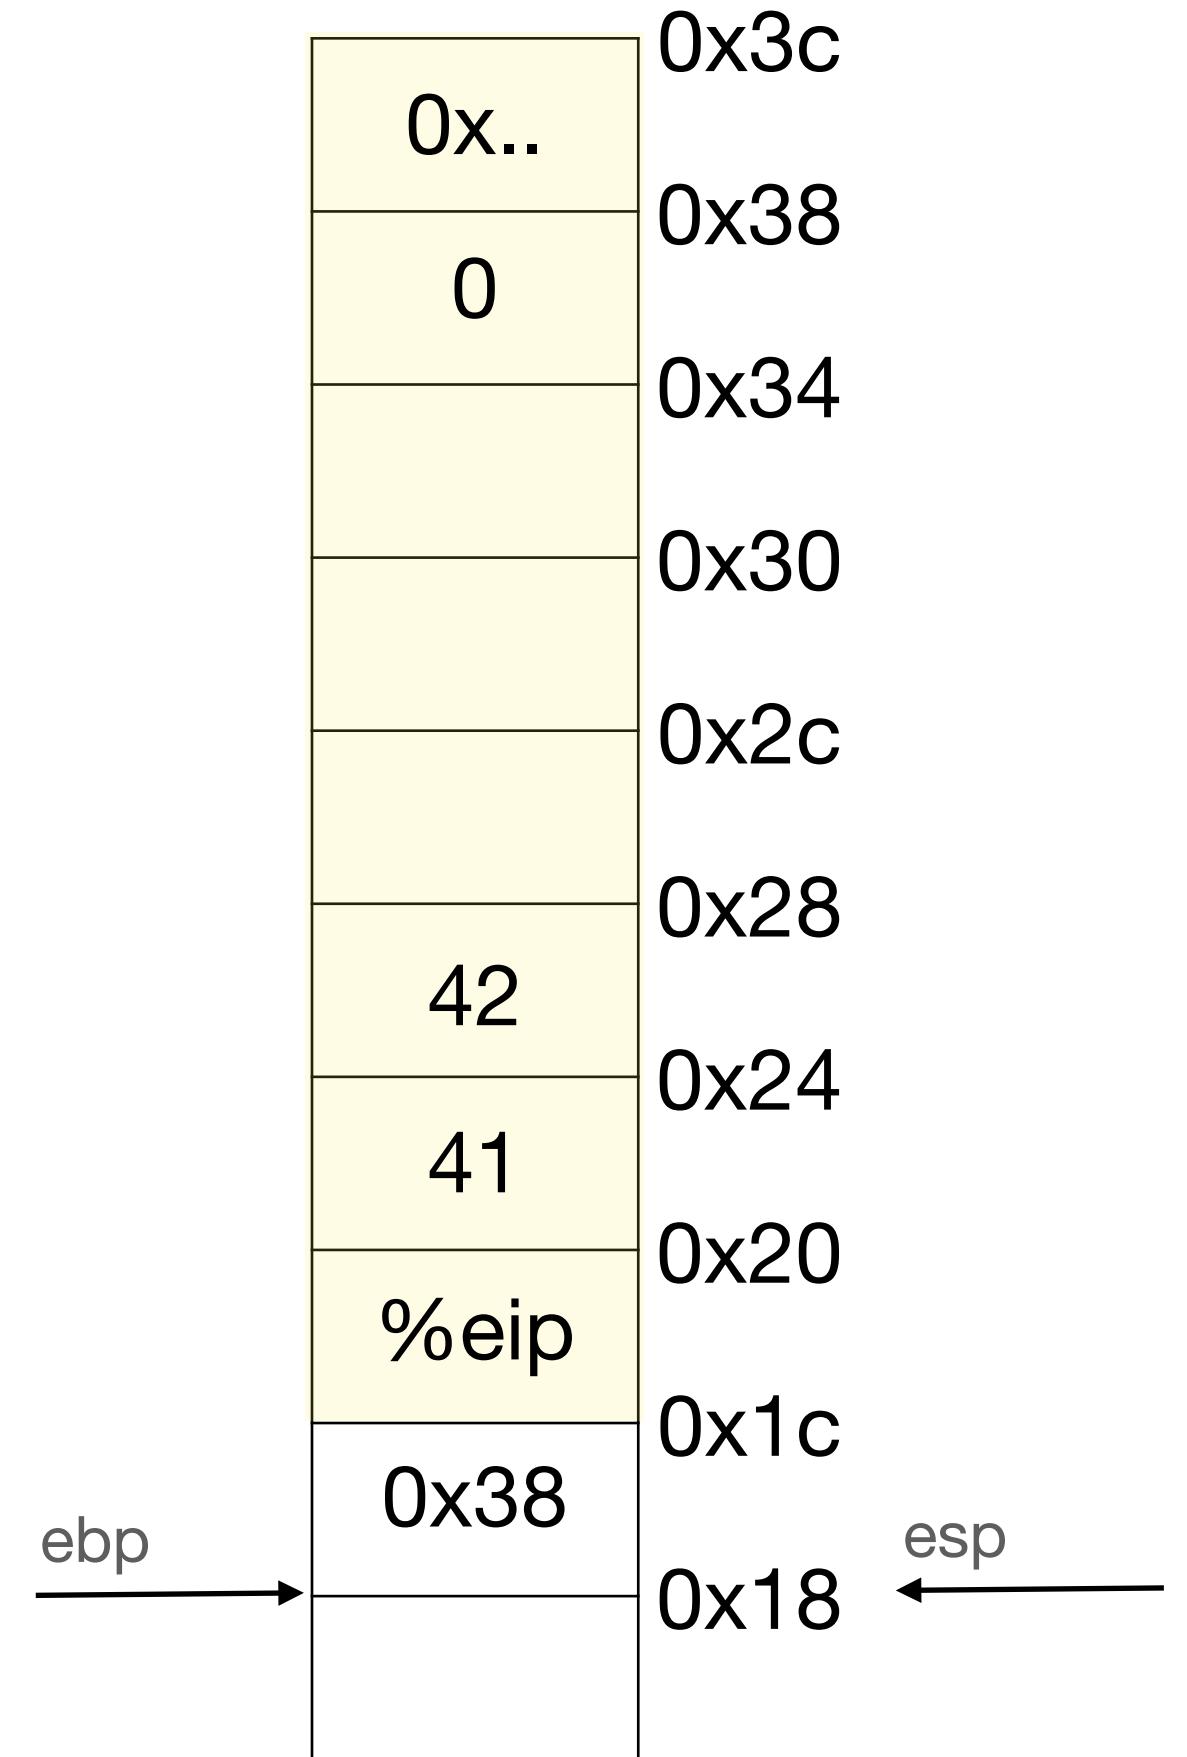
_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

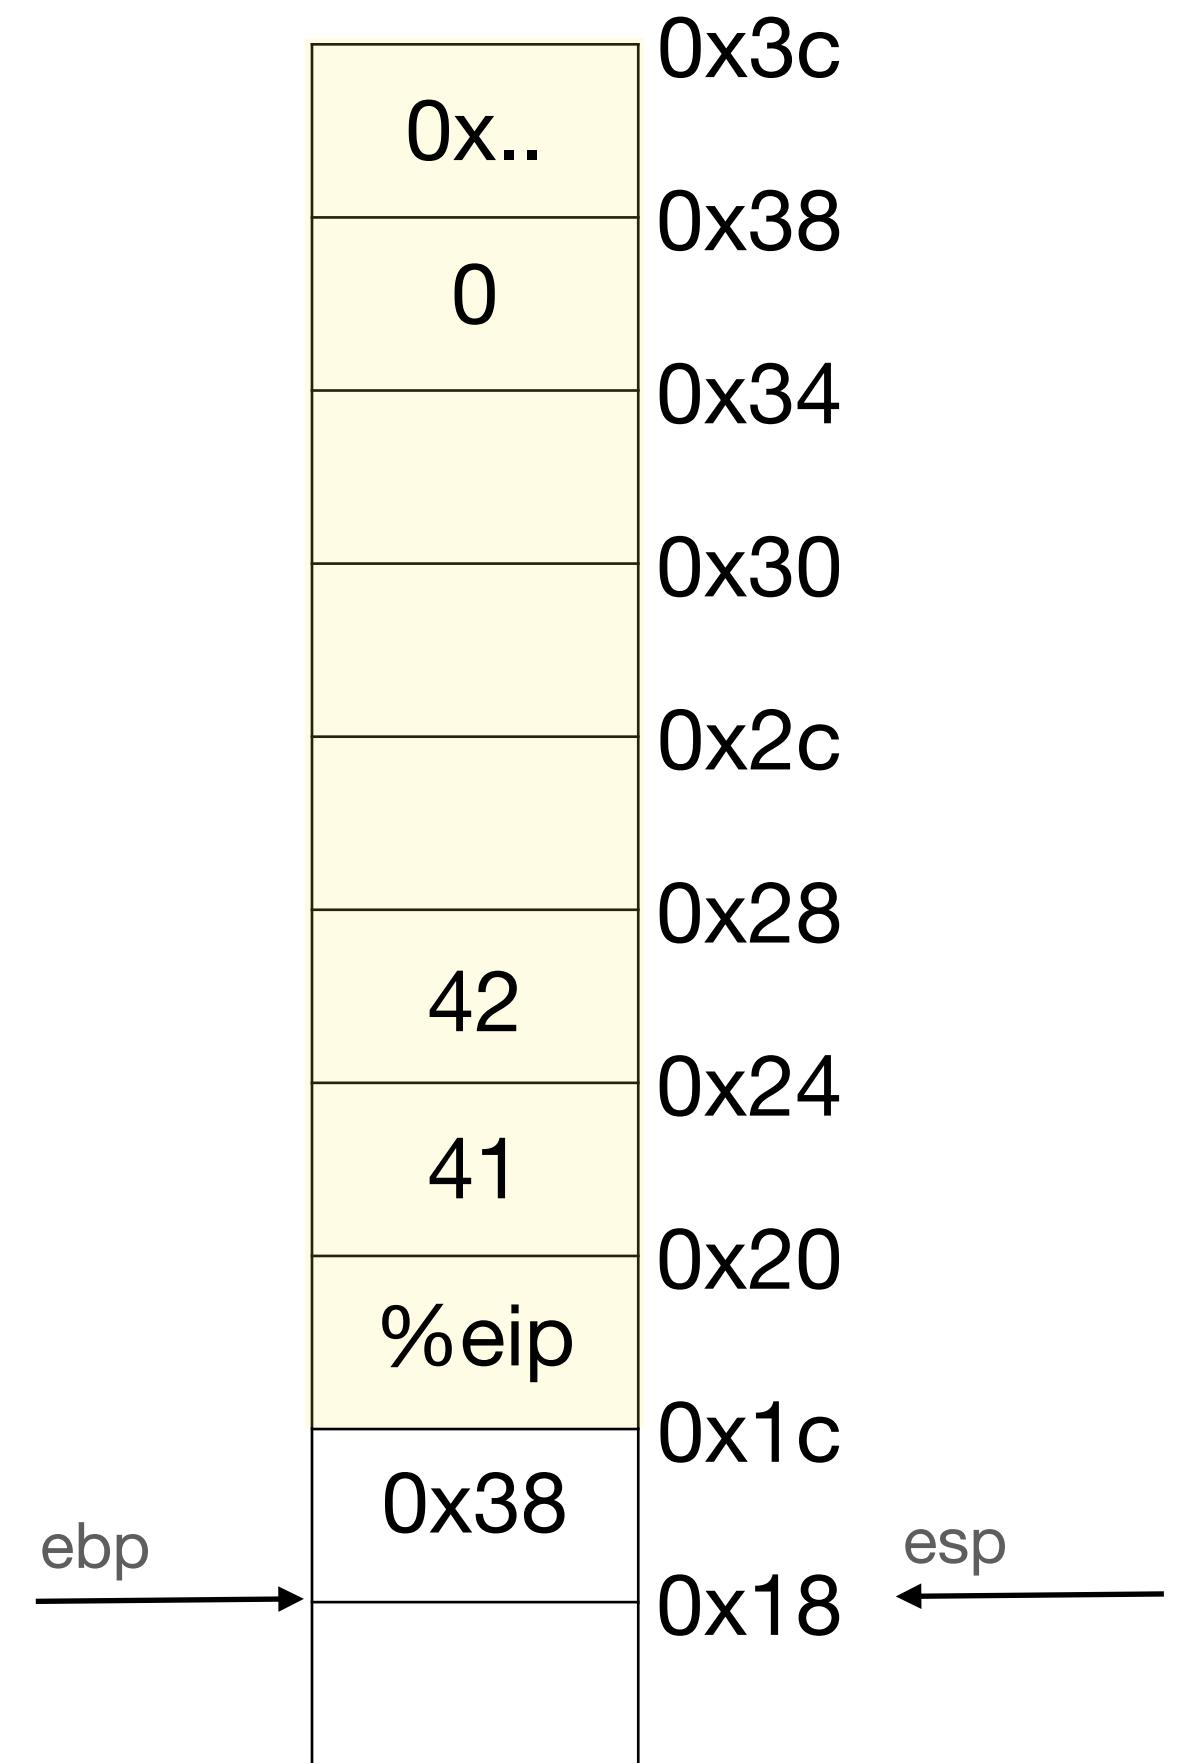
## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

## -- Begin function main
```



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

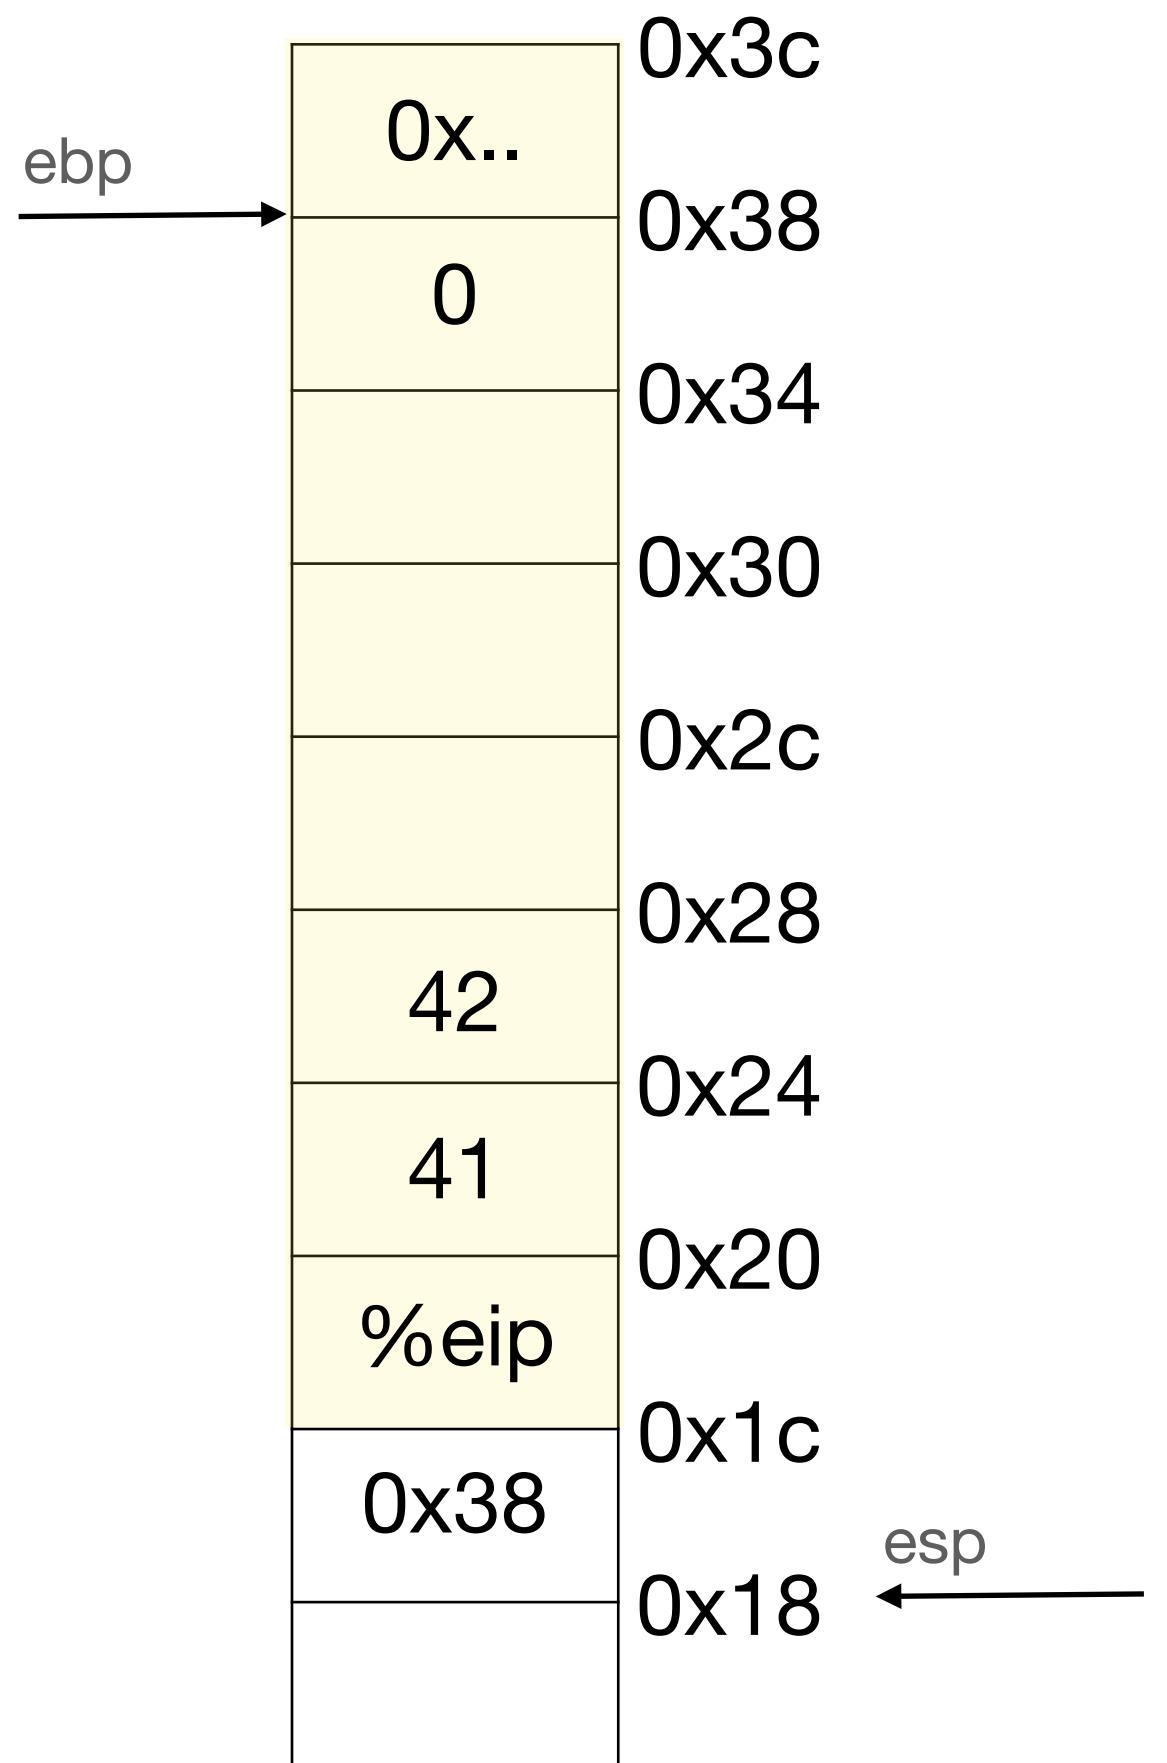
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

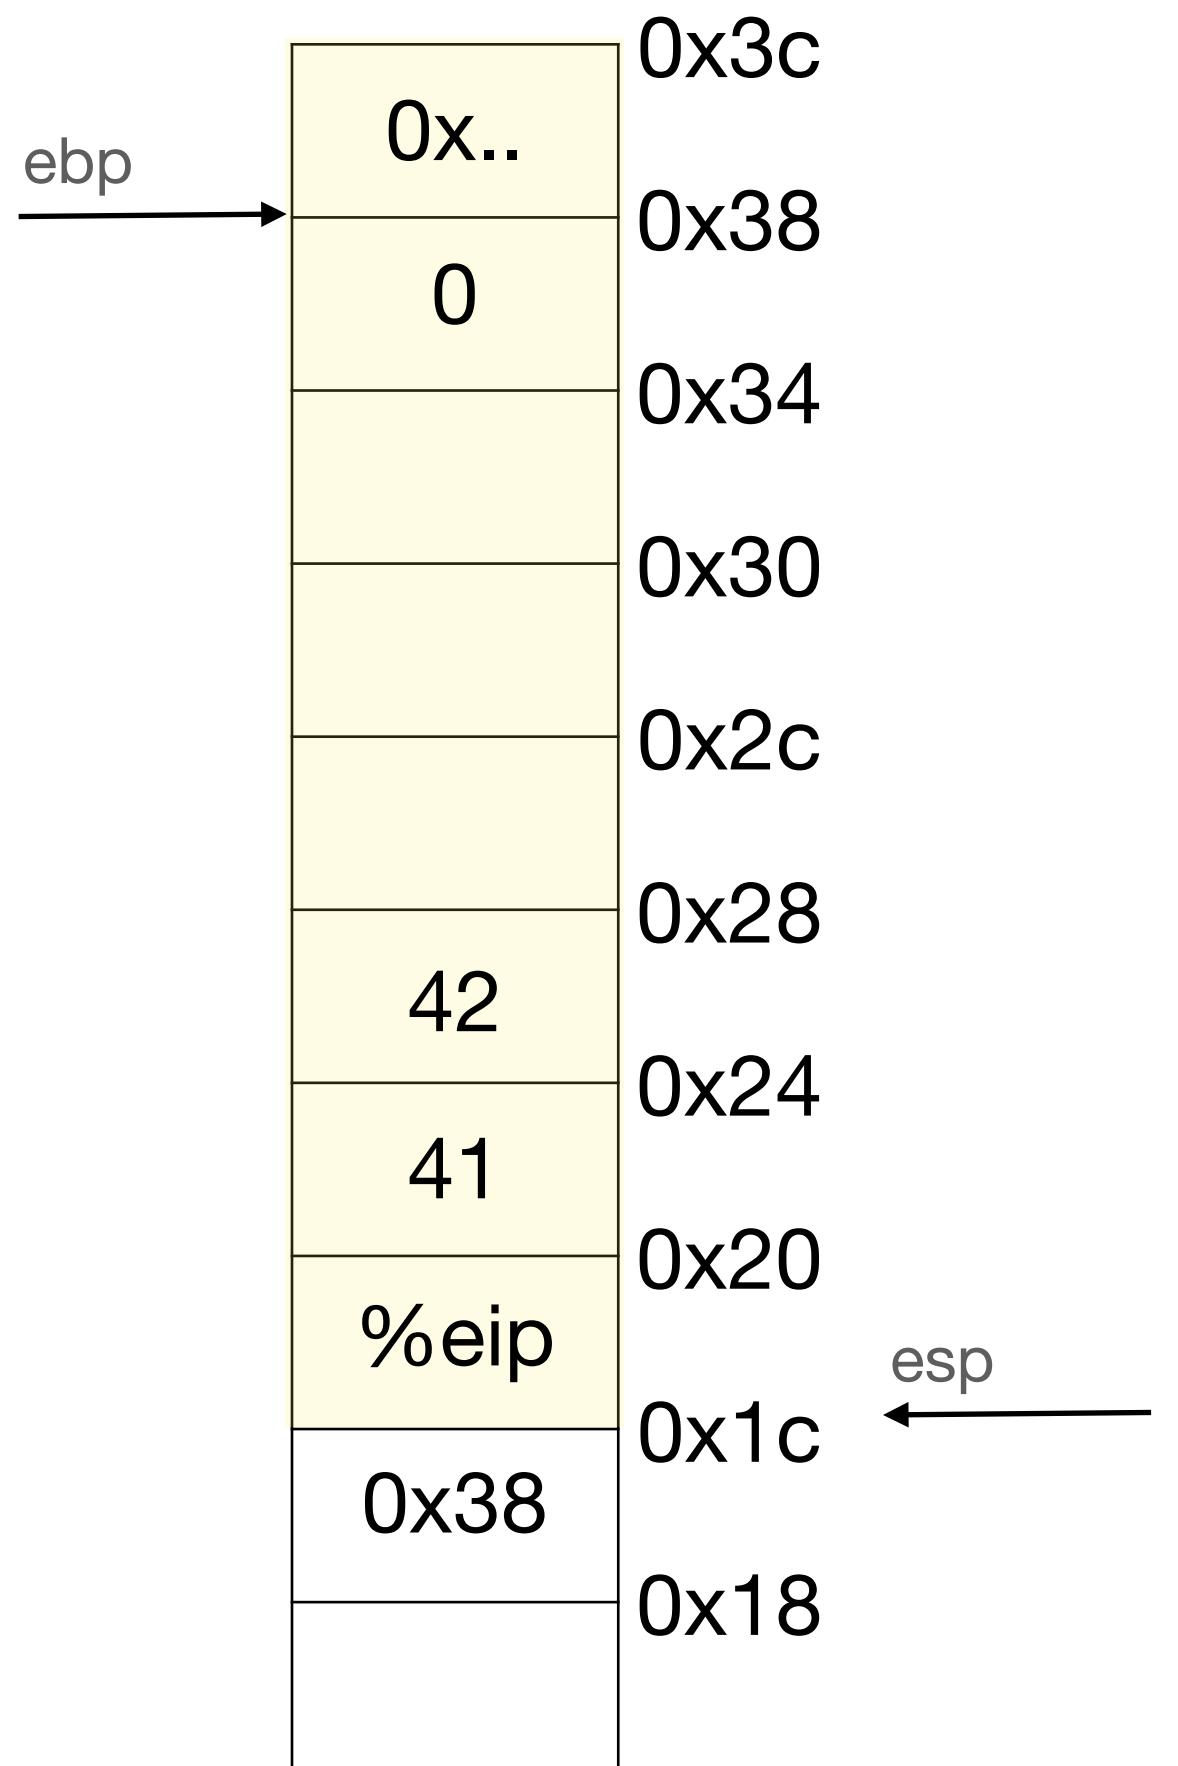
## -- Begin function main
```

Annotations for \_foo:

- pushl %ebp: Save caller's base pointer
- movl %esp, %ebp: ebp = esp
- movl 8(%ebp), %eax: eax = \*(ebp + 8)
- addl 12(%ebp), %eax: eax = eax + \*(ebp + 12)
- popl %ebp: Restore caller's base pointer
- retl: change eip to return address

Annotations for \_main:

- pushl %ebp: Save caller's base pointer
- movl %esp, %ebp: ebp = esp
- subl \$24, %esp: esp = esp - 0x18
- movl \$0, -4(%ebp): \*(ebp-4)=0
- movl \$41, (%esp): \*(esp) = 41
- movl \$42, 4(%esp): \*(esp+4) = 42
- calll \_foo: Push current eip on to stack, jump to foo
- addl \$24, %esp: esp = esp + 24 (Restore caller's esp)
- popl %ebp: Restore caller's ebp
- retl:



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

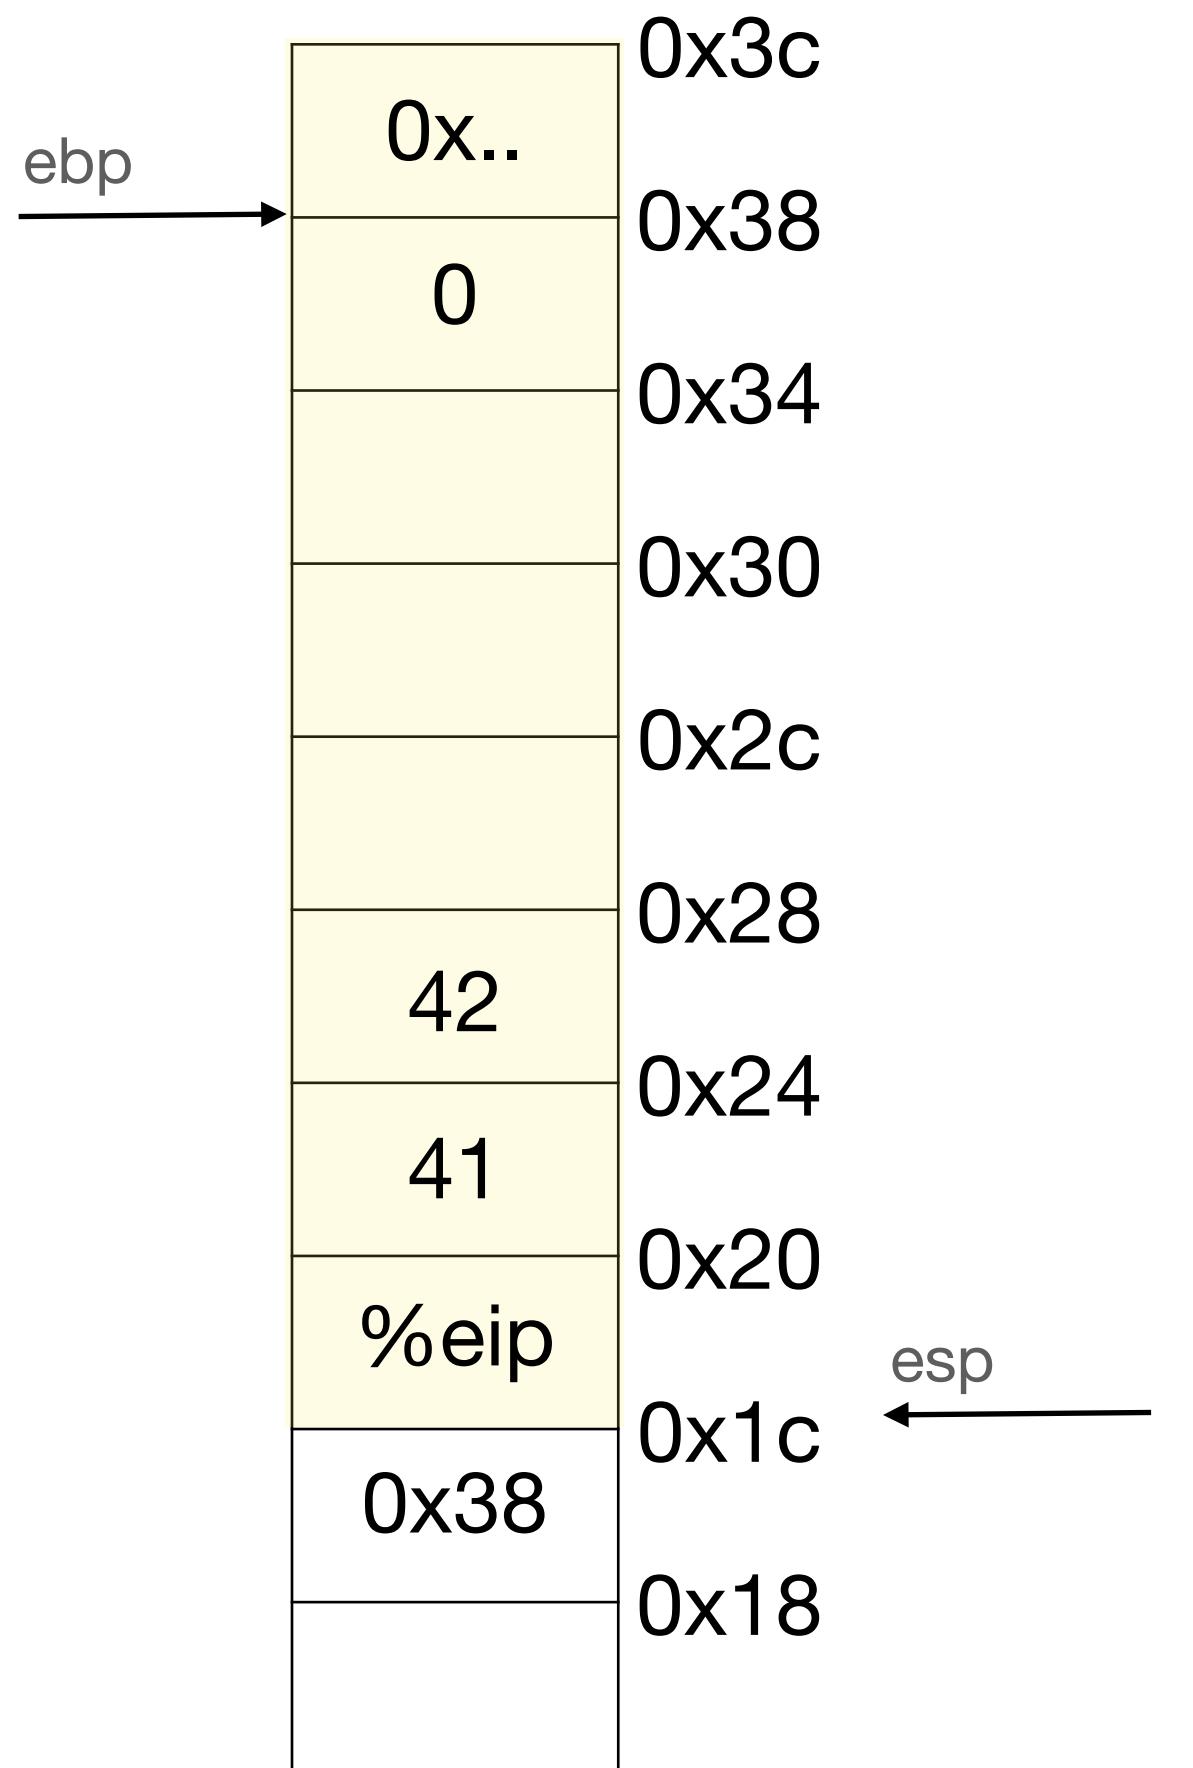
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip → retl
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

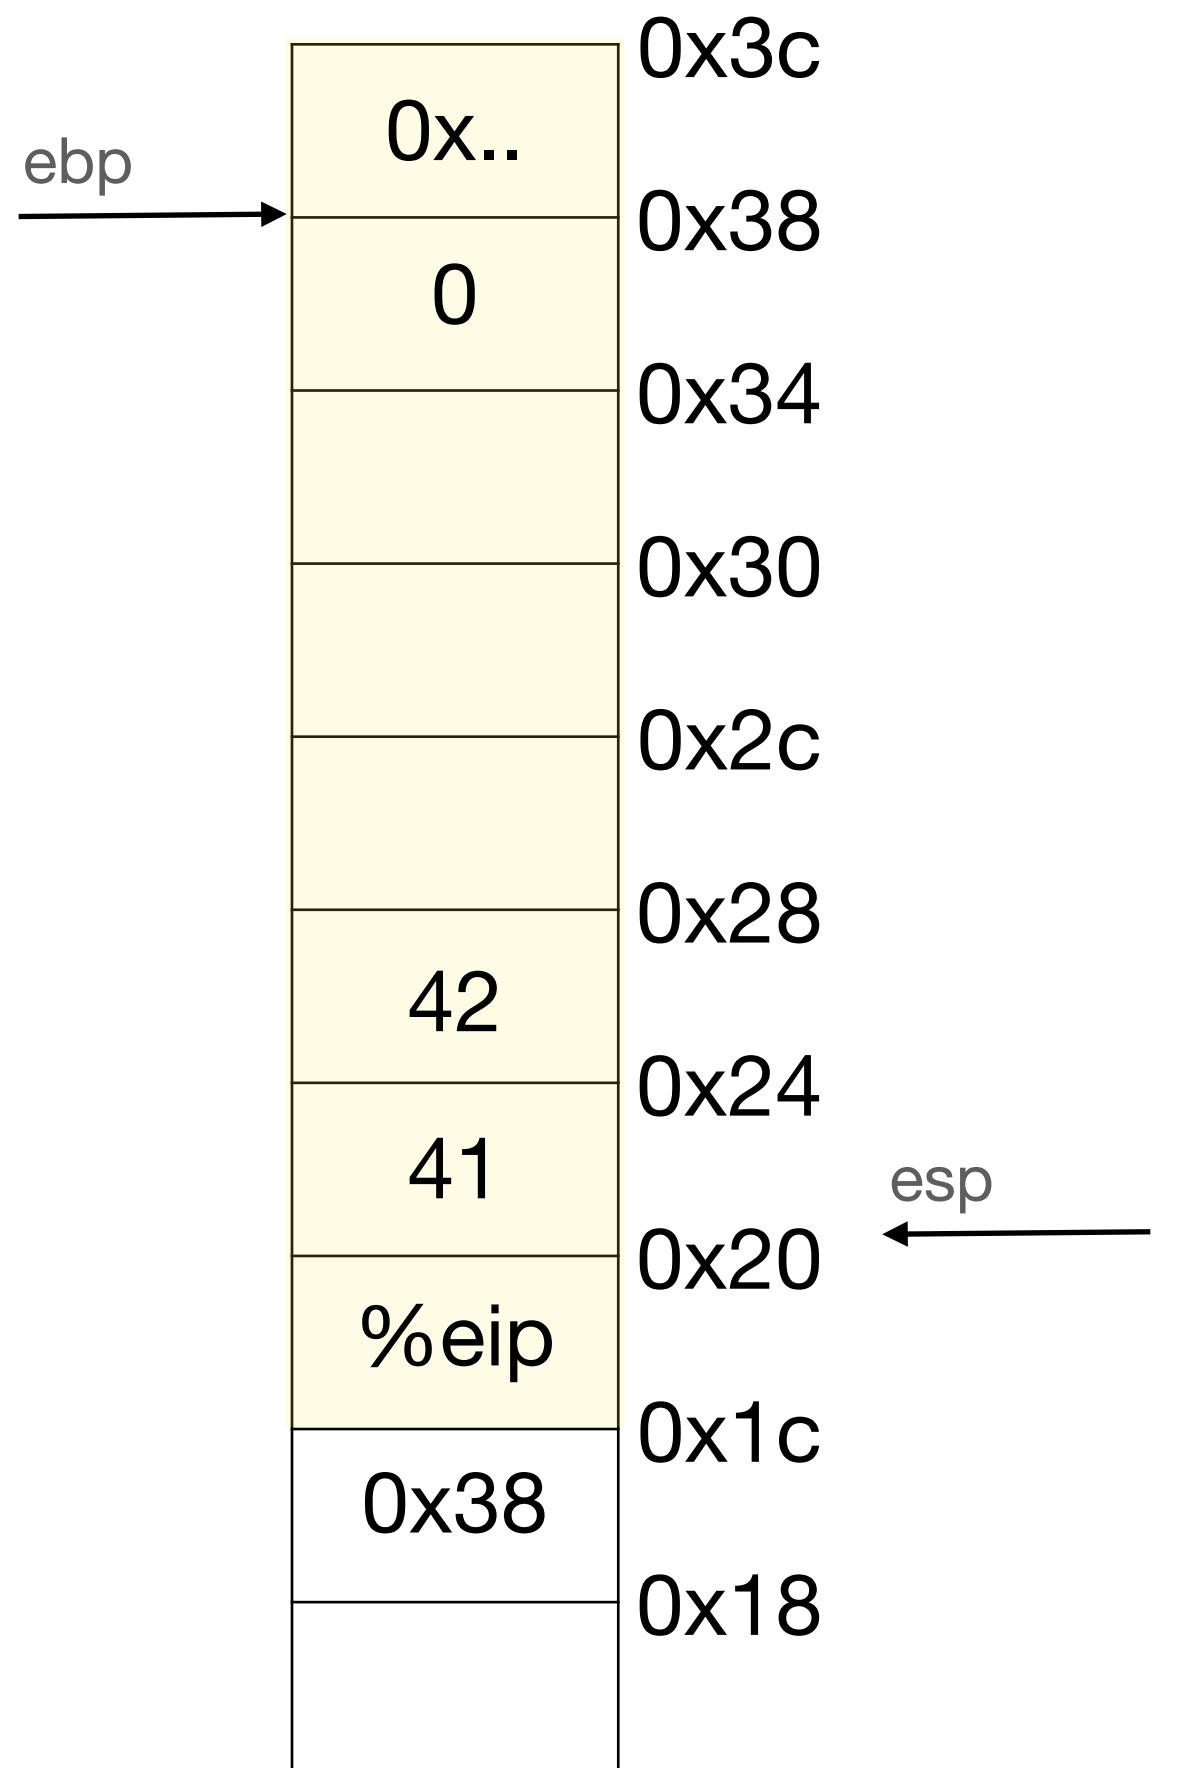
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip → retl
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

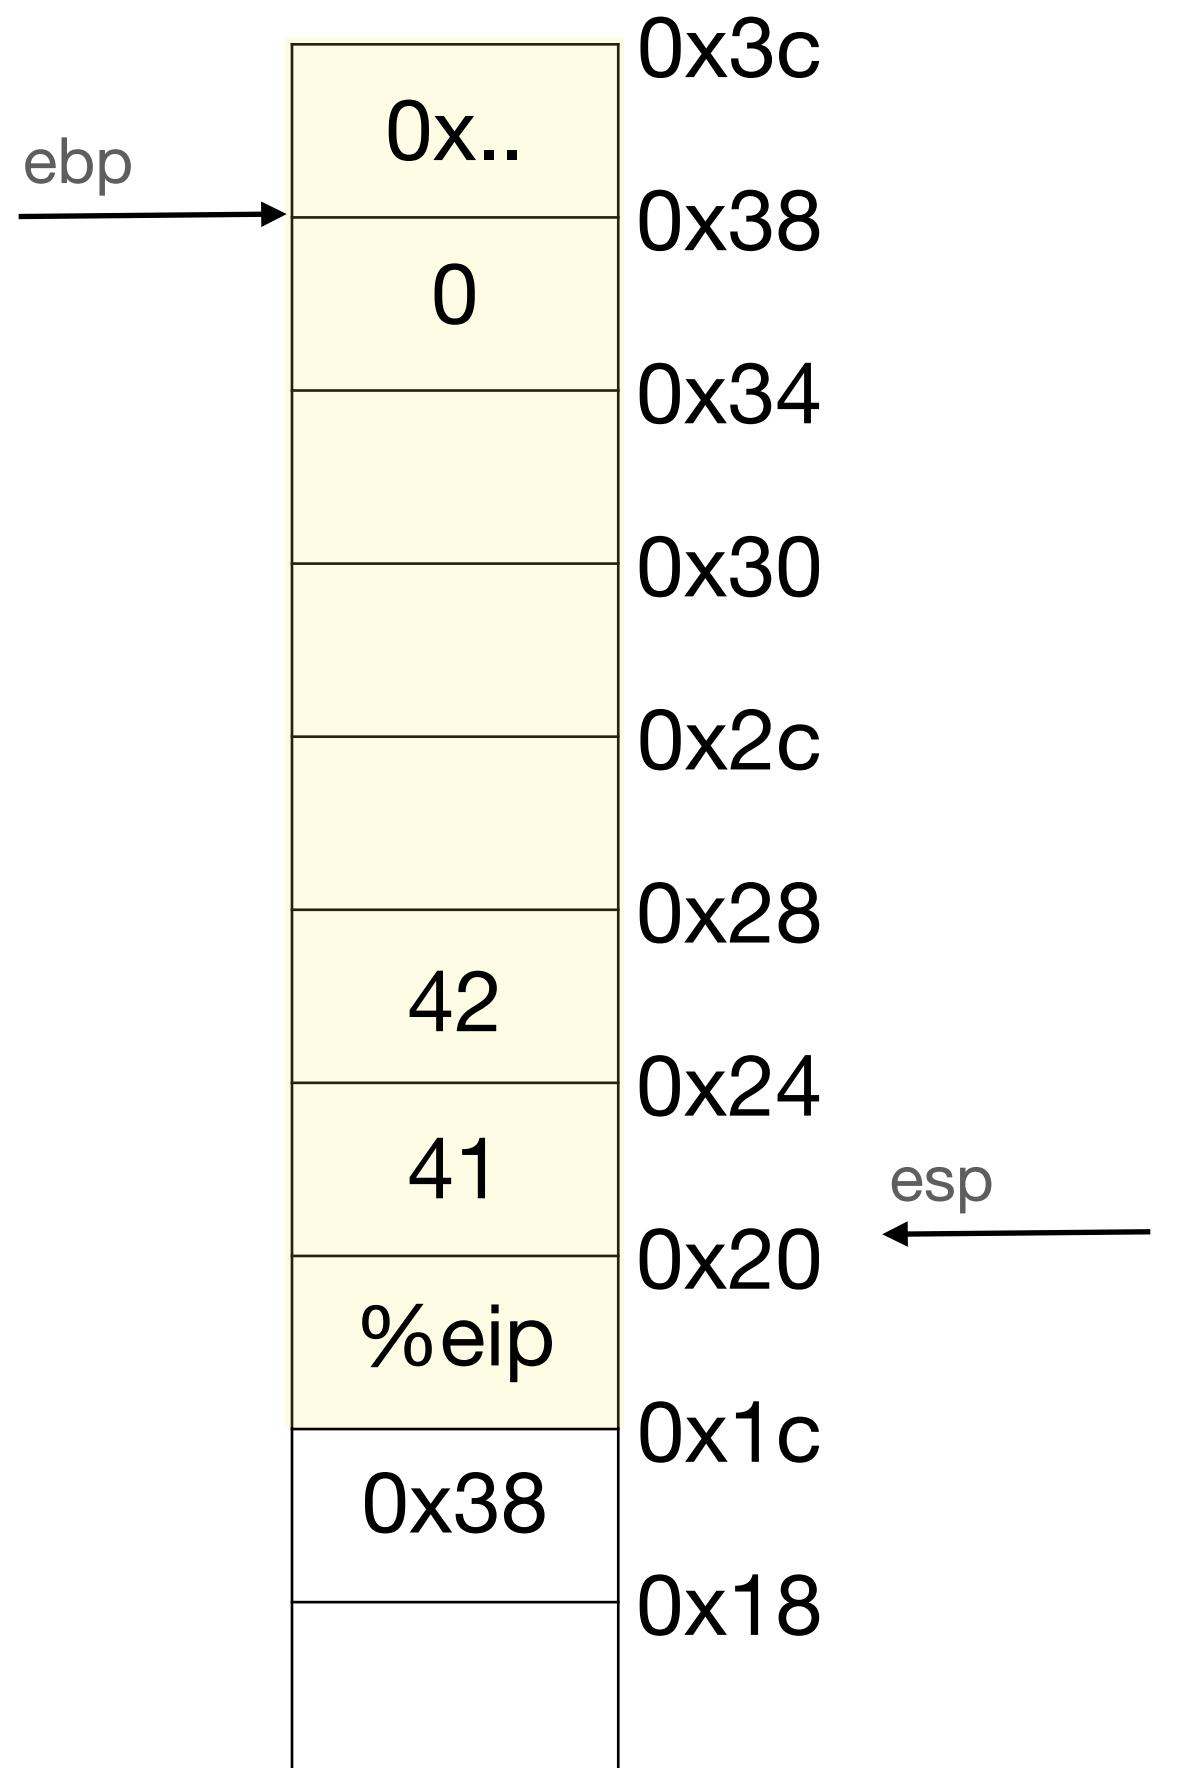
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

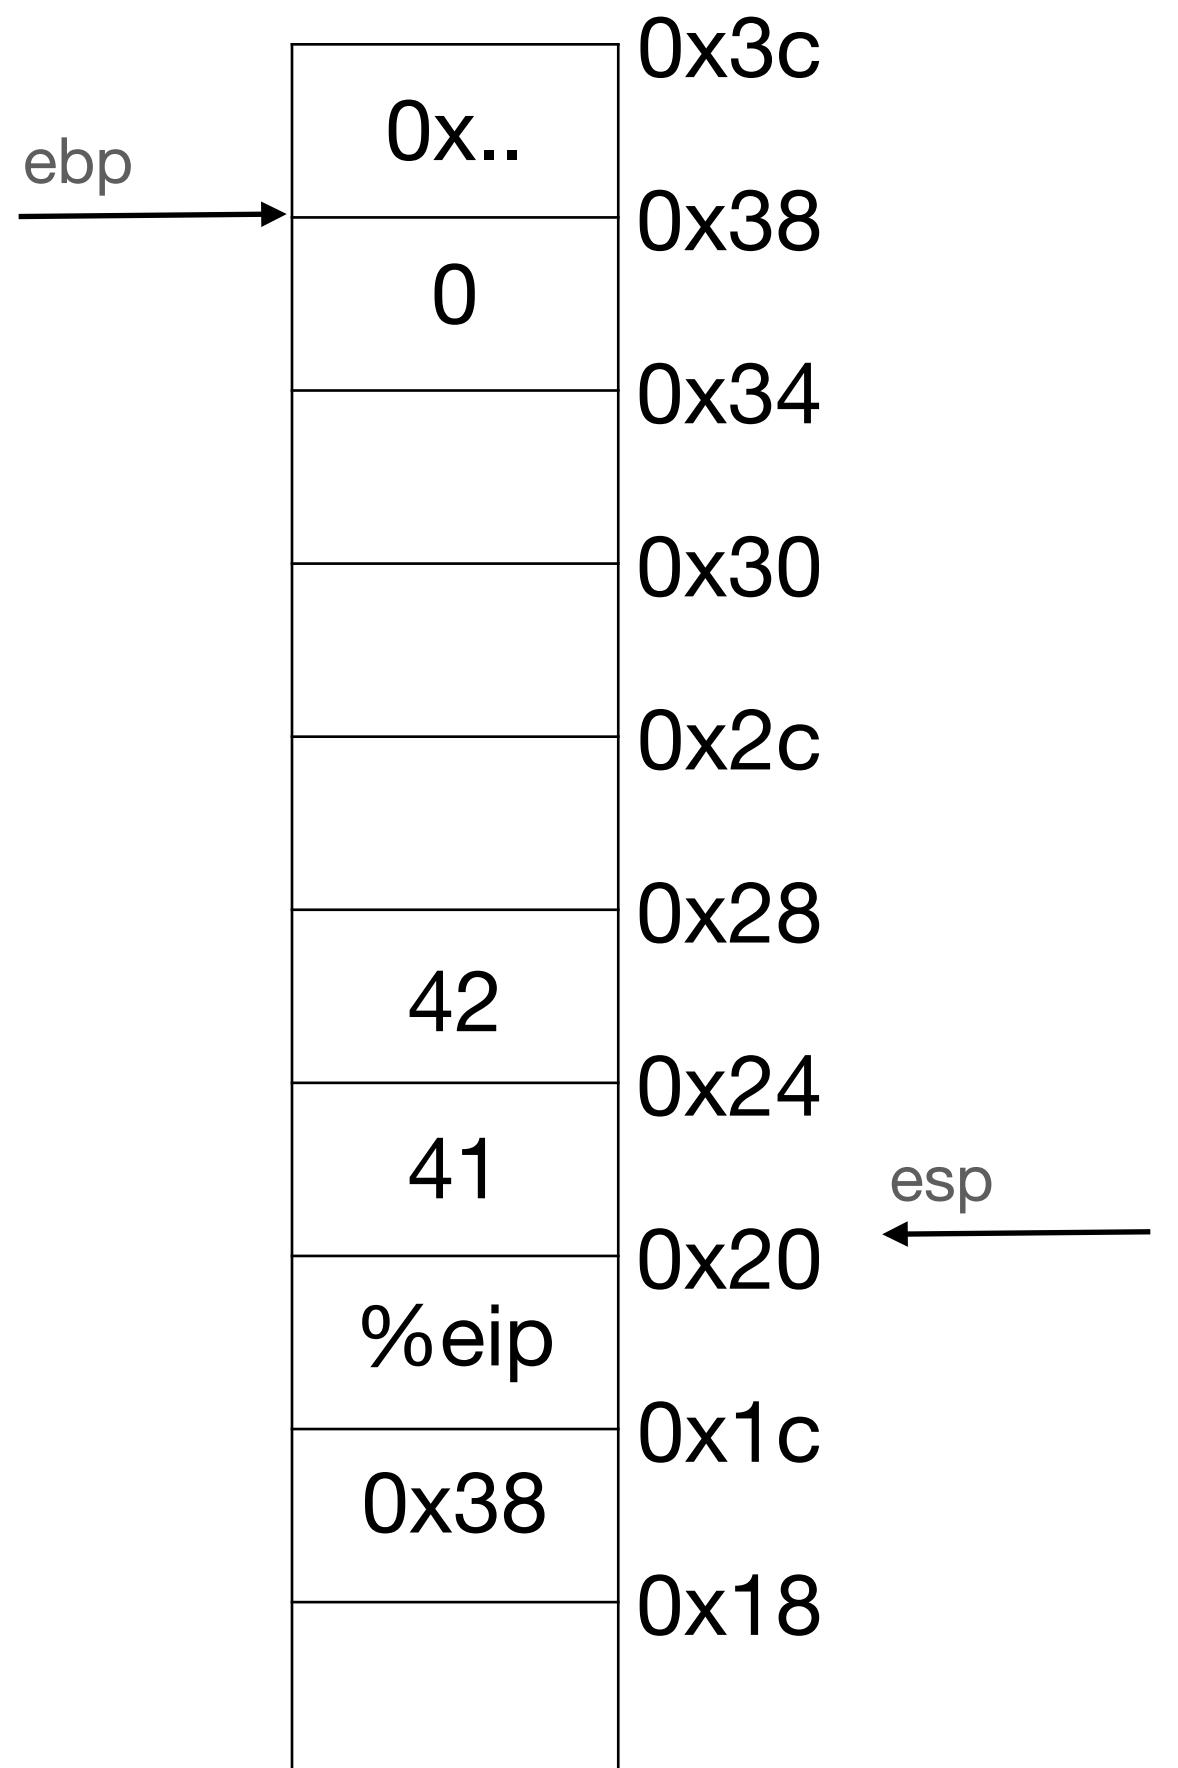
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

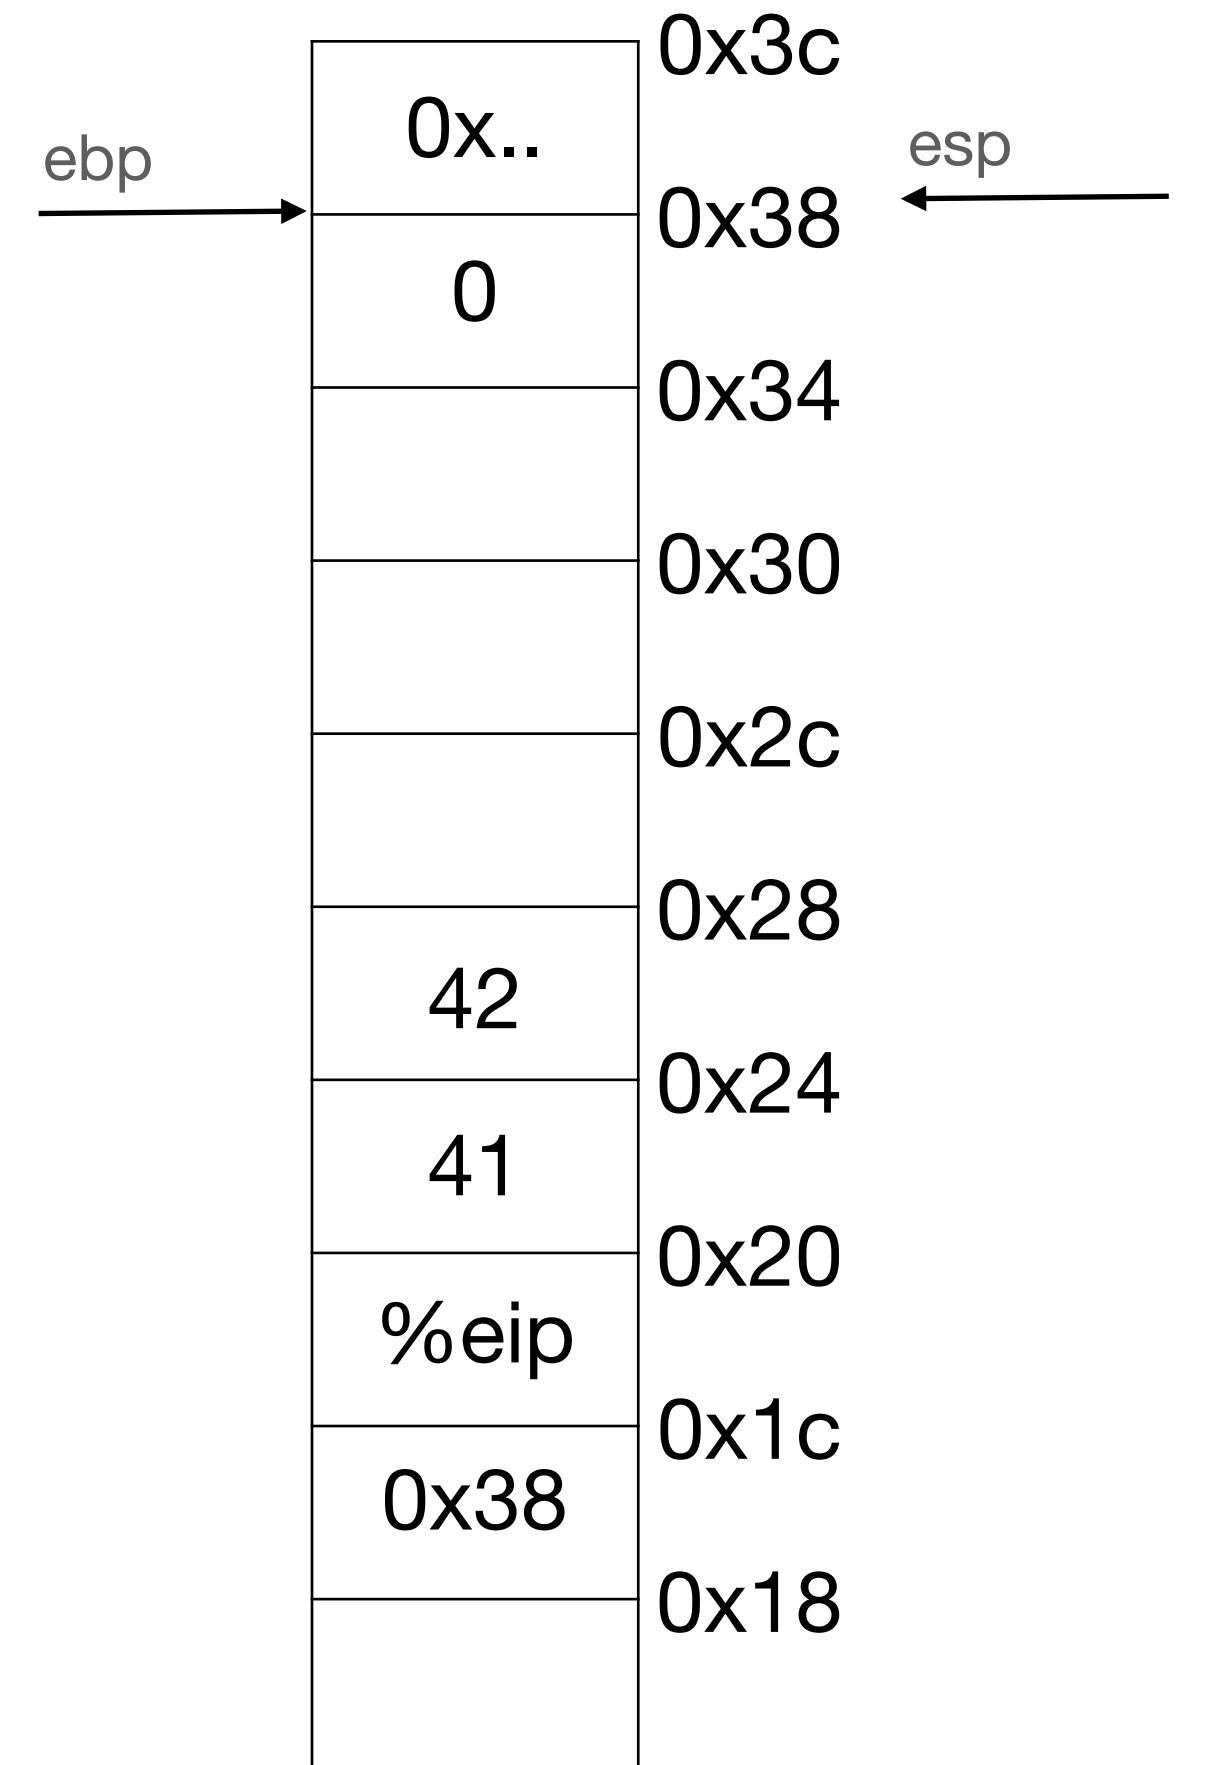
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

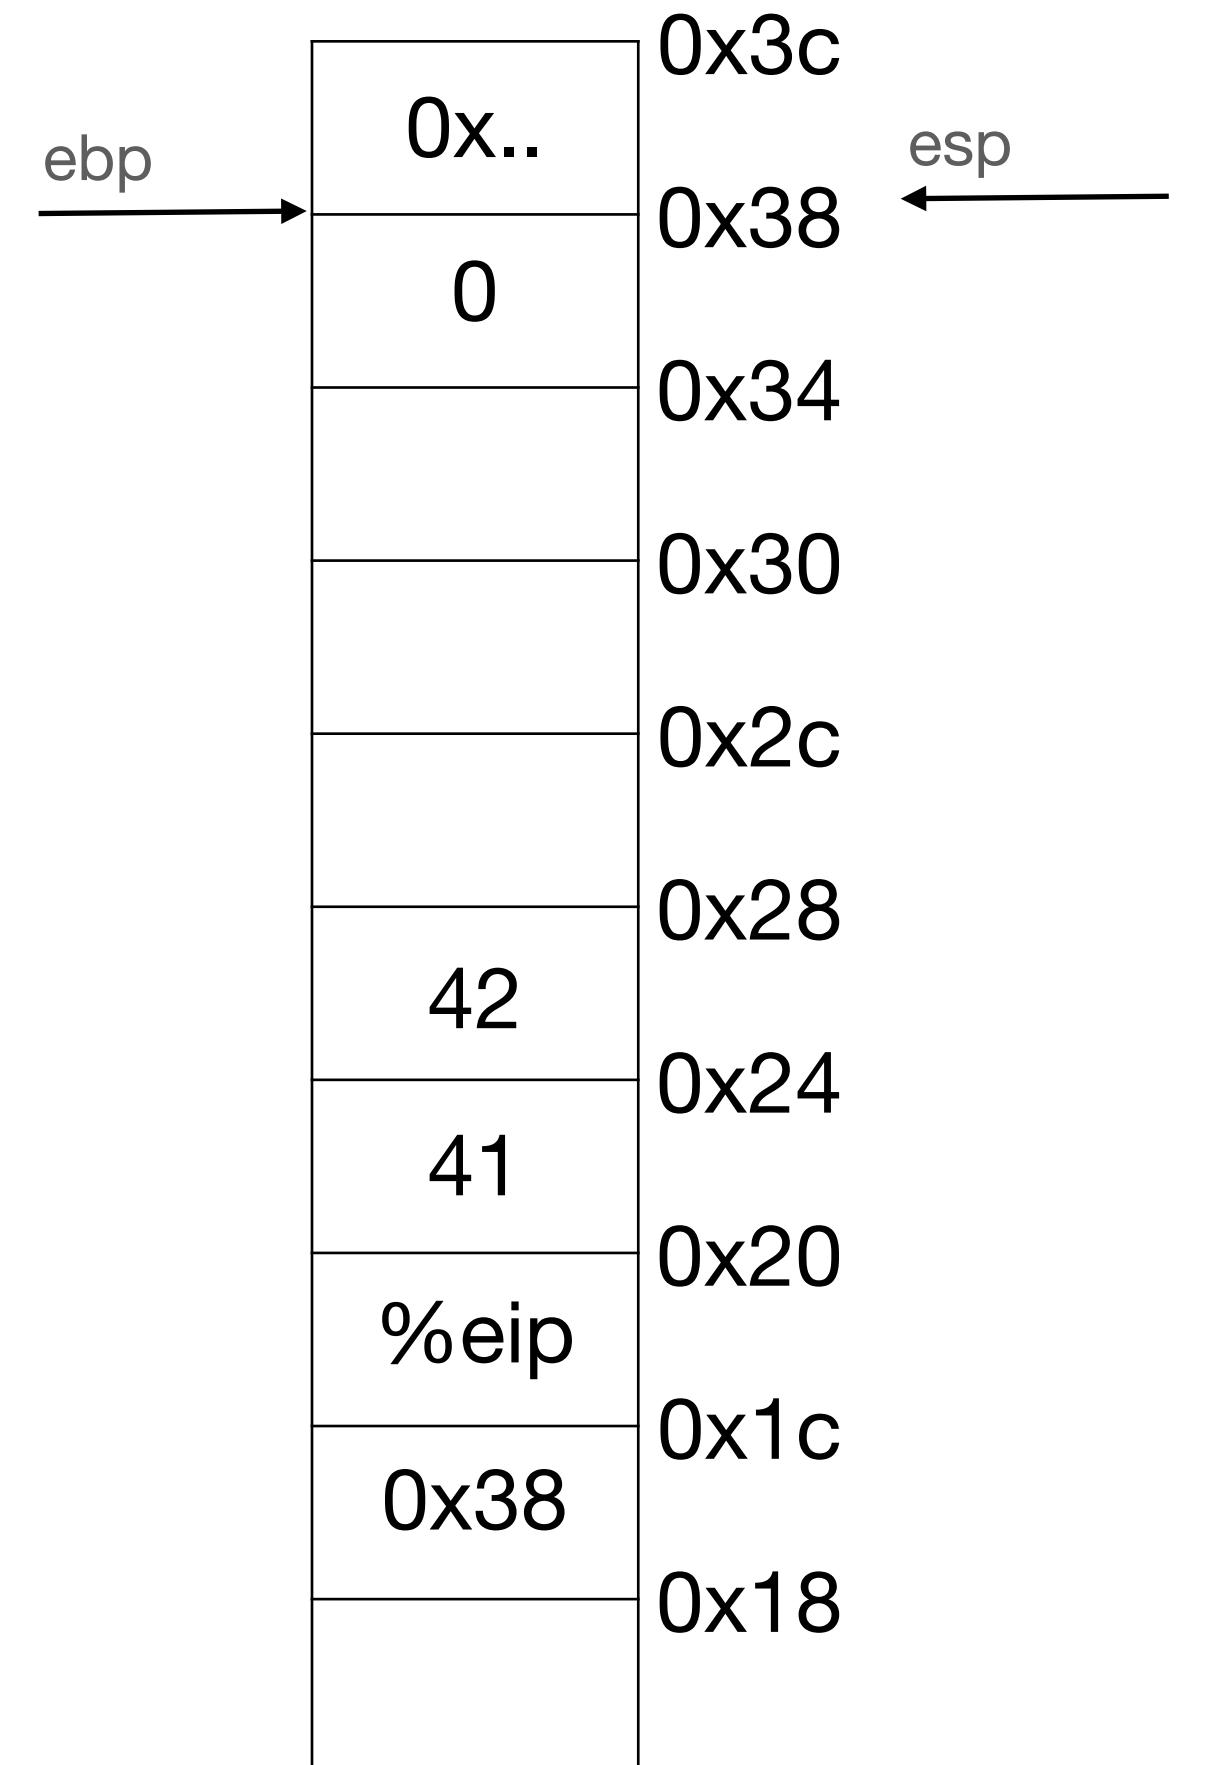
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

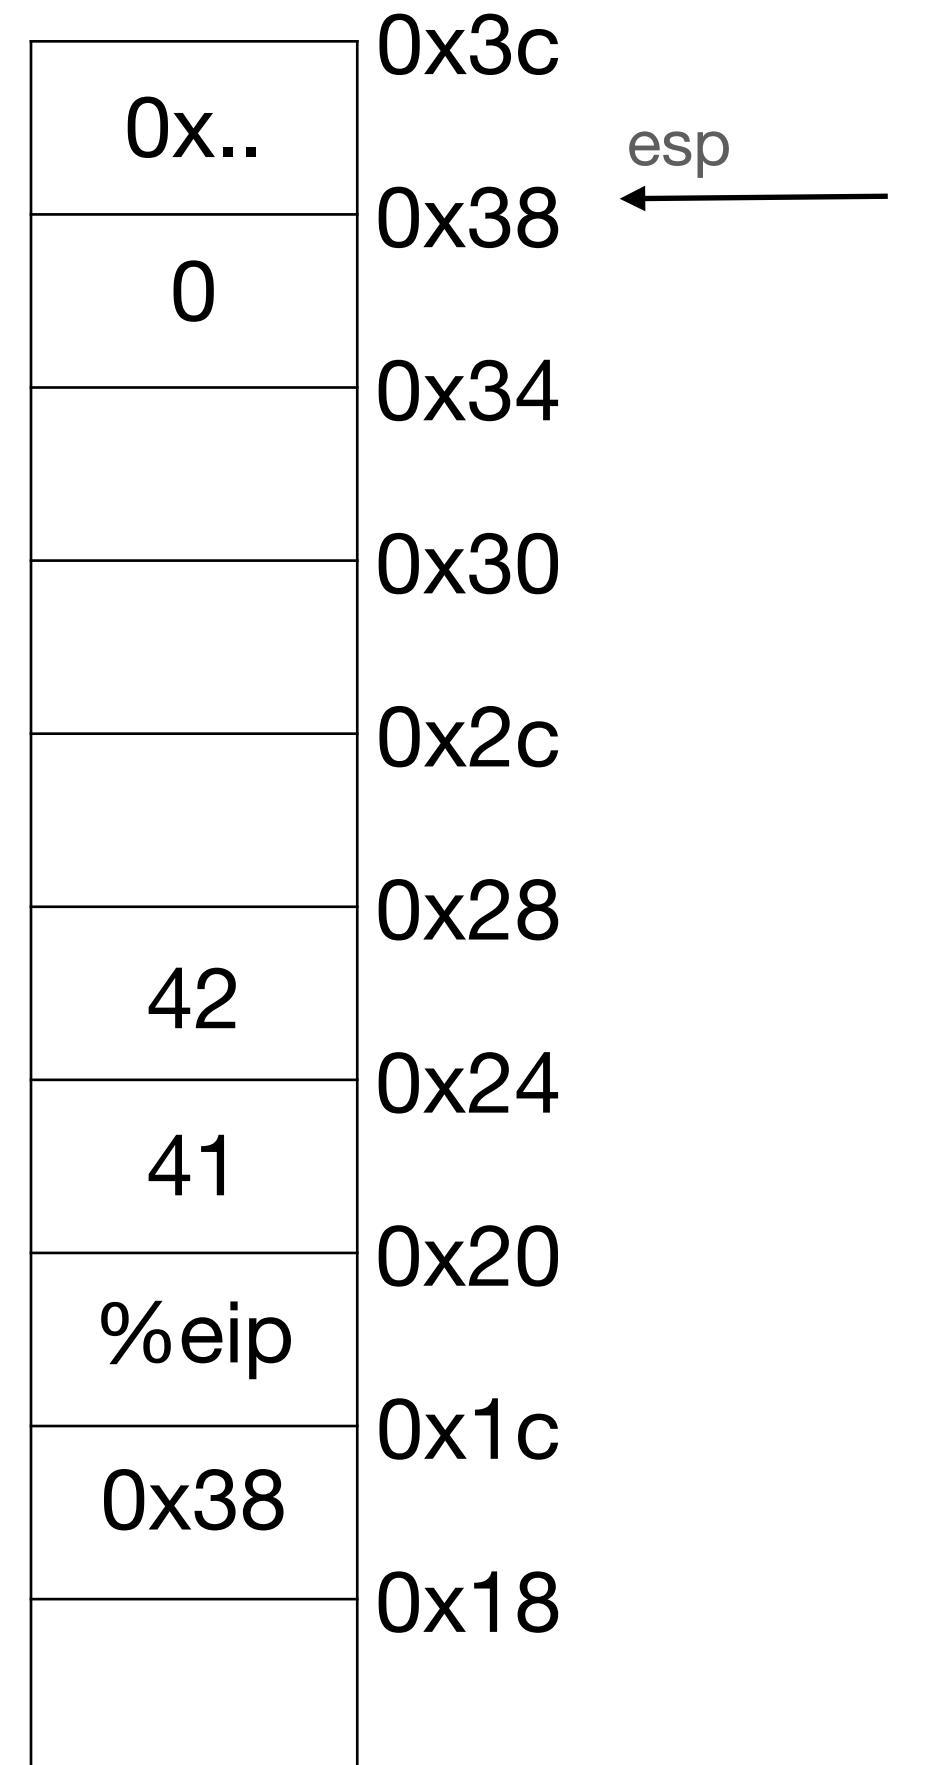
eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp

ebp →



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

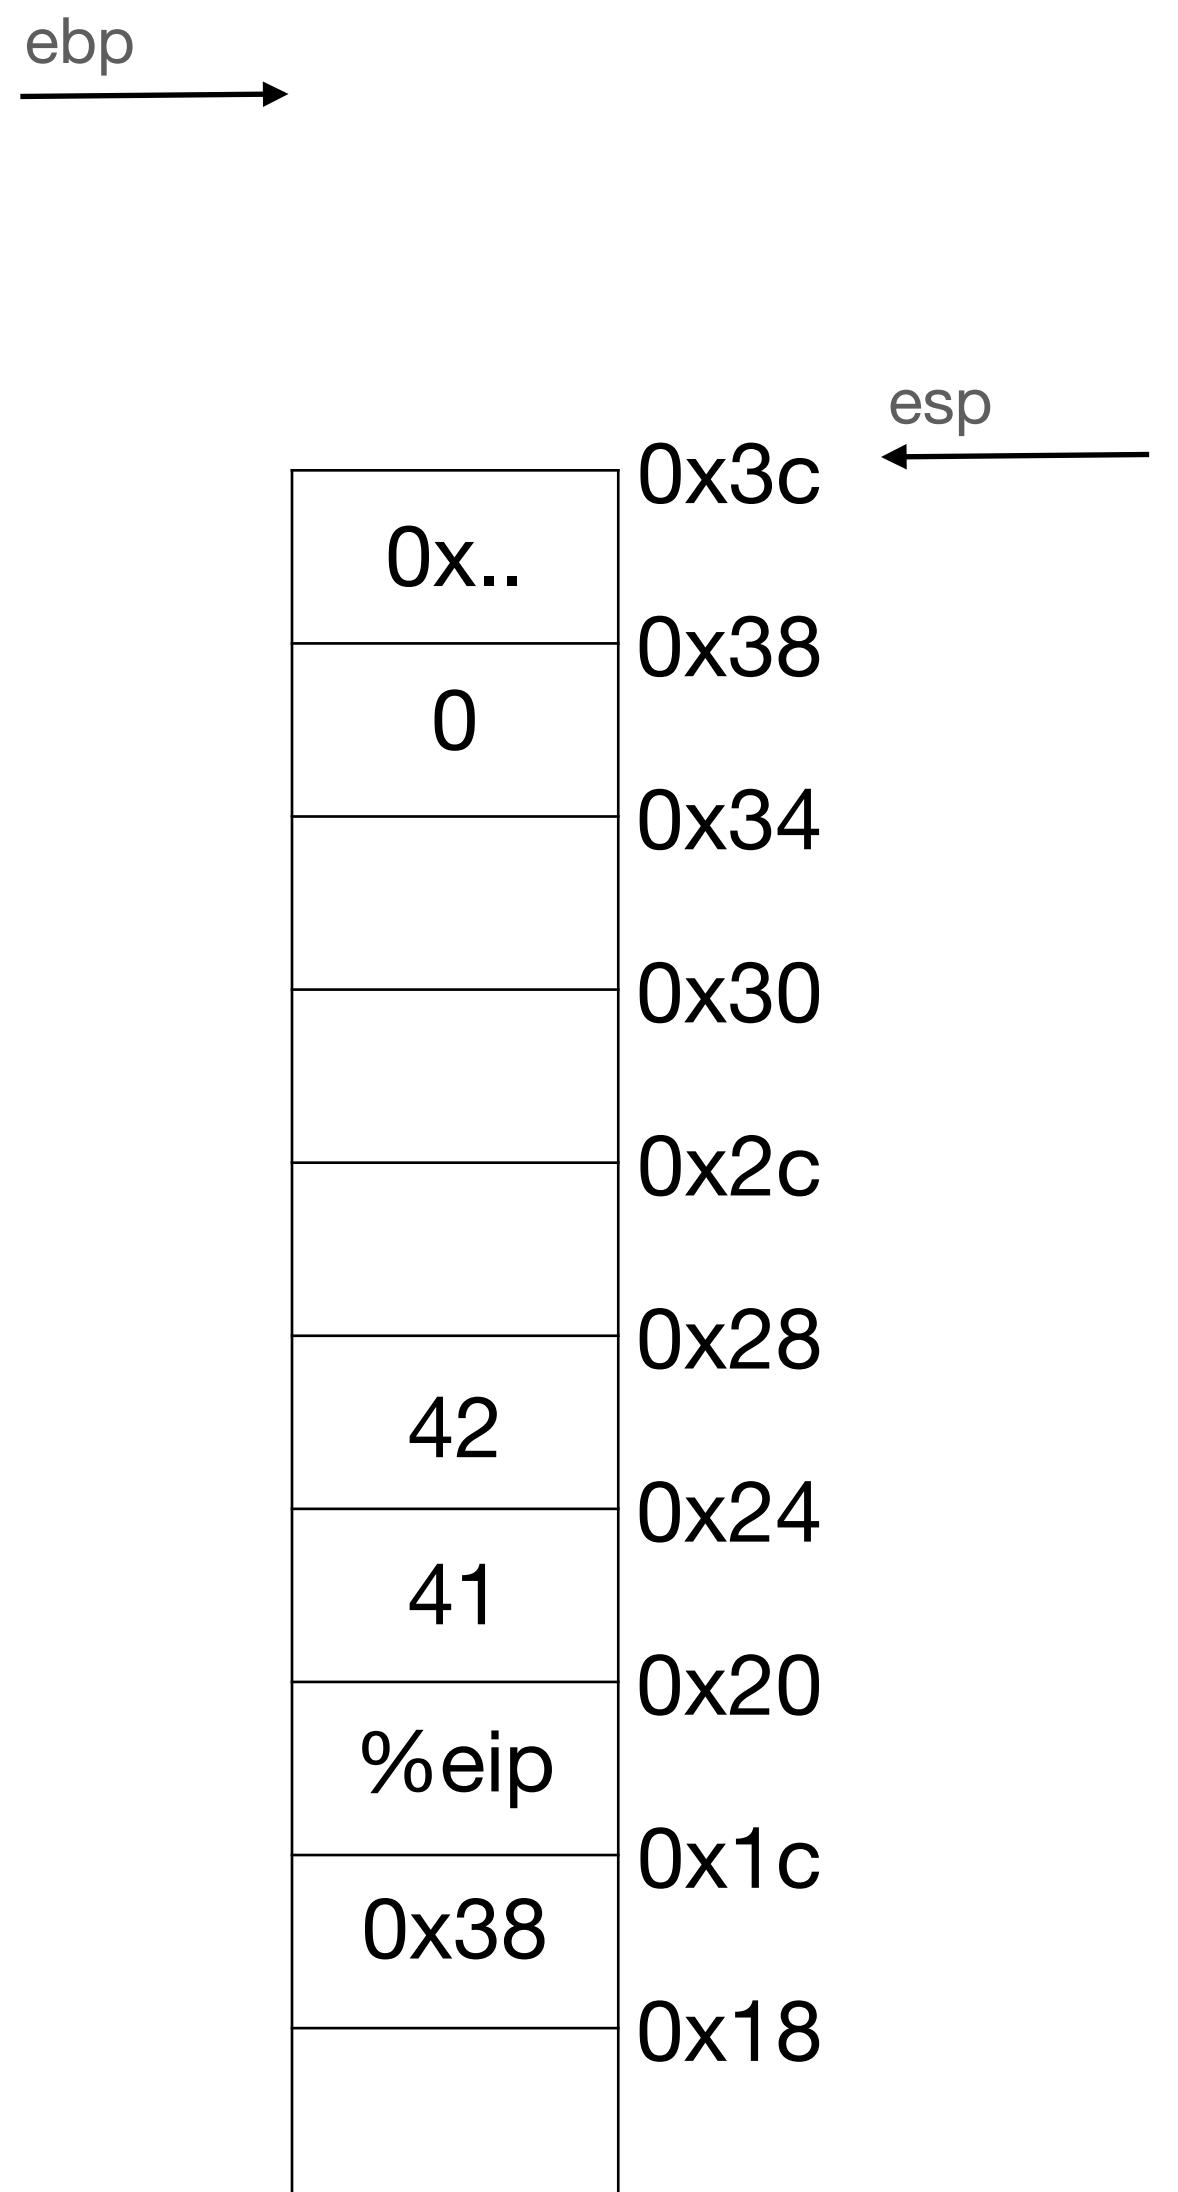
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip →
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp



# Function calling in action

## Stack

```
02.s

_foo:
    pushl %ebp
    movl %esp, %ebp
    movl 8(%ebp), %eax
    addl 12(%ebp), %eax
    popl %ebp
    retl

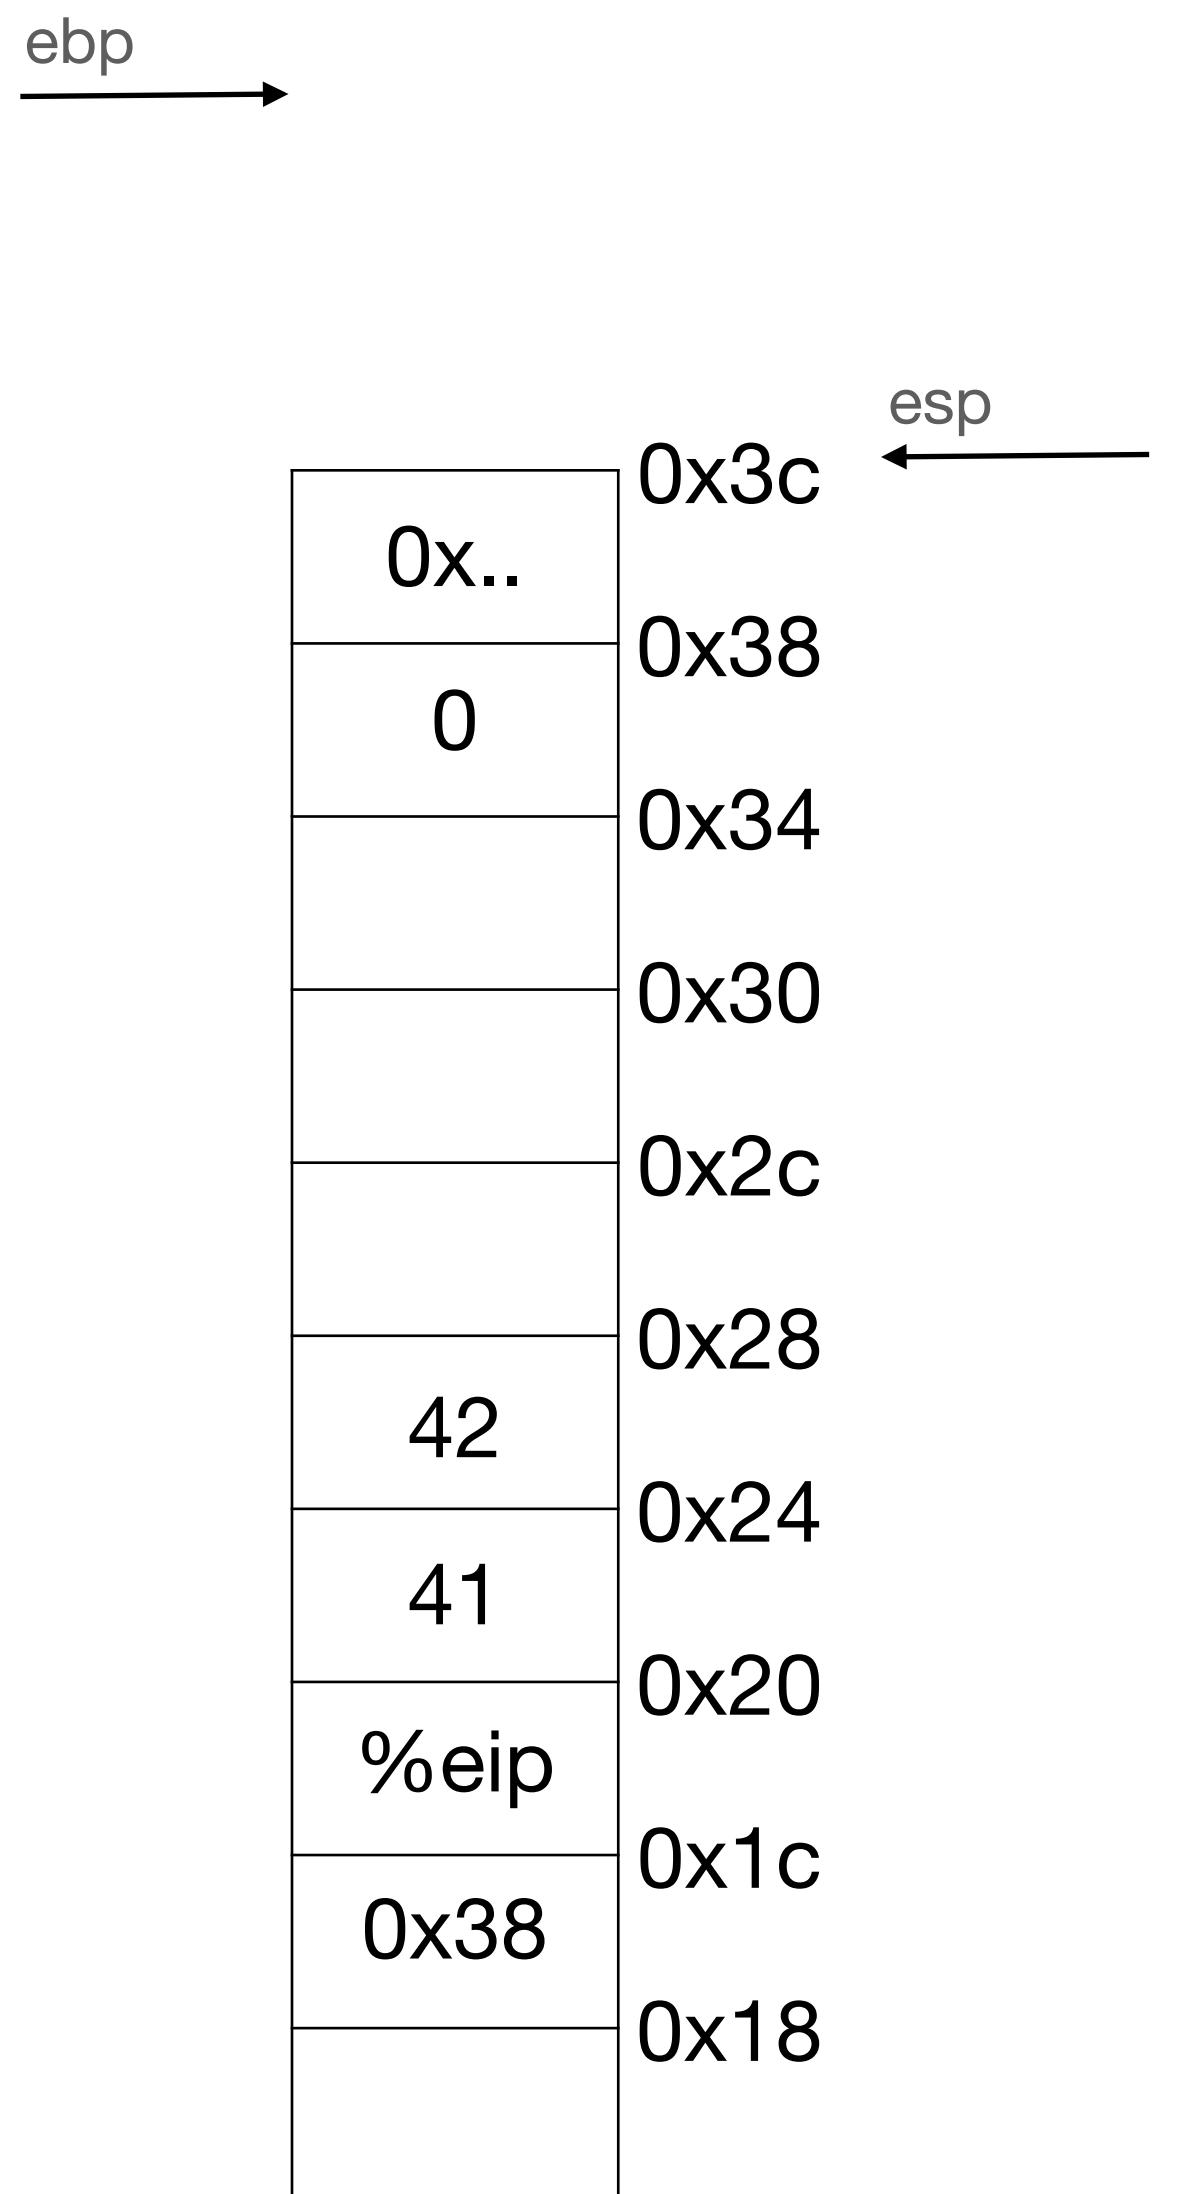
    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl %esp, %ebp
    subl $24, %esp
    movl $0, -4(%ebp)
    movl $41, (%esp)
    movl $42, 4(%esp)
    calll _foo
    addl $24, %esp
    popl %ebp
    retl

eip
```

Save caller's base pointer  
ebp = esp  
eax = \*(ebp + 8)  
eax = eax + \*(ebp + 12)  
Restore caller's base pointer  
change eip to return address

## -- Begin function main

Save caller's base pointer  
ebp = esp  
esp = esp - 0x18  
\*(ebp-4)=0  
\*(esp) = 41  
\*(esp+4) = 42  
Push current eip on to stack, jump to foo  
esp = esp + 24 (Restore caller's esp)  
Restore caller's ebp

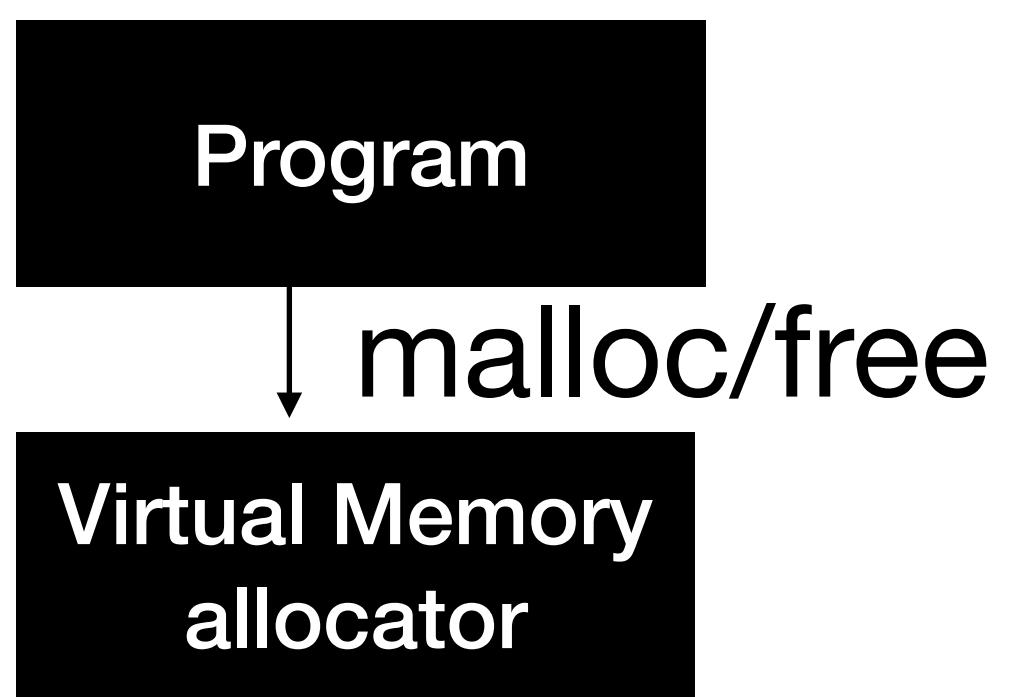


# Memory APIs and bugs

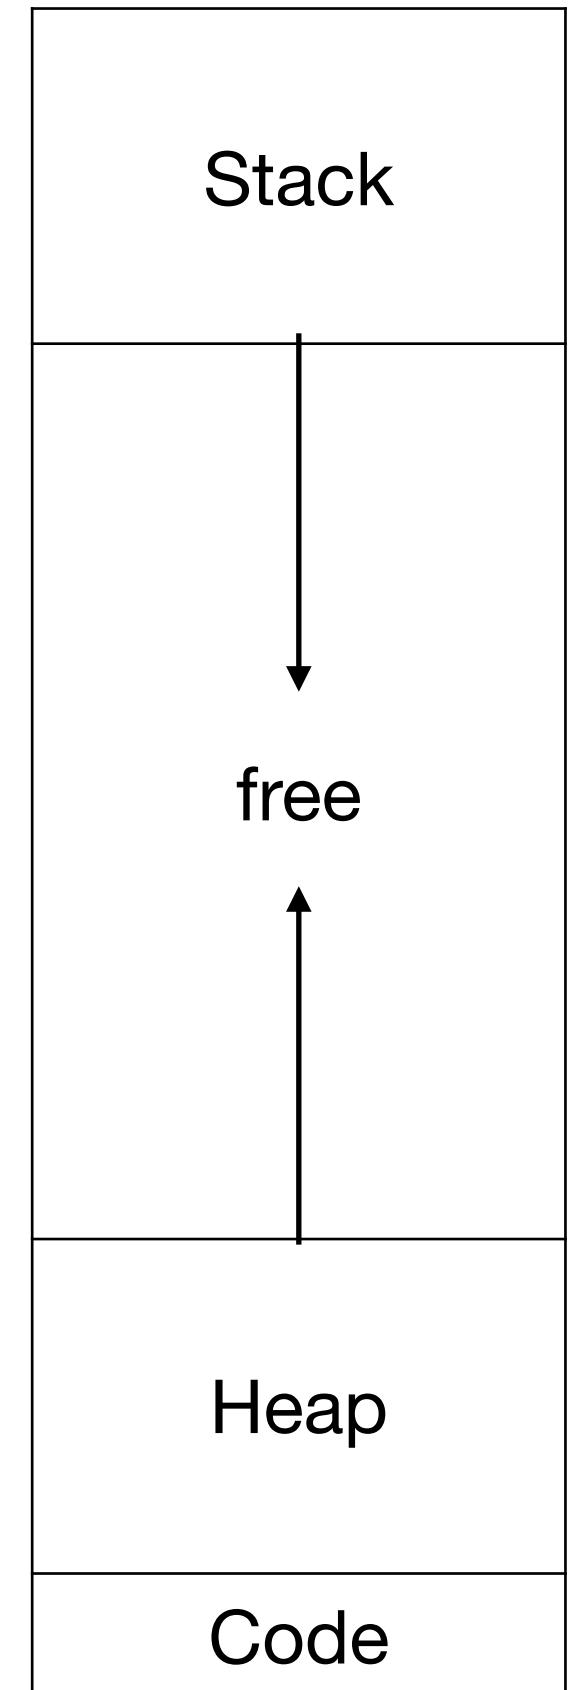
- malloc, free, va.c.
  - malloc is for dynamic allocation. Size is not known at compile time. Slower than stack allocations. Need to find free space.
- Null pointer dereference. null.c
- Memory leak. leak.c
- Buffer overflow. overflow.c
- Use after free. useafterfree.c
- Invalid free. invalidfree.c
- Double free. doublefree.c
- Uninitialised read. uninitread.c

# Memory allocator

Works with virtual memory

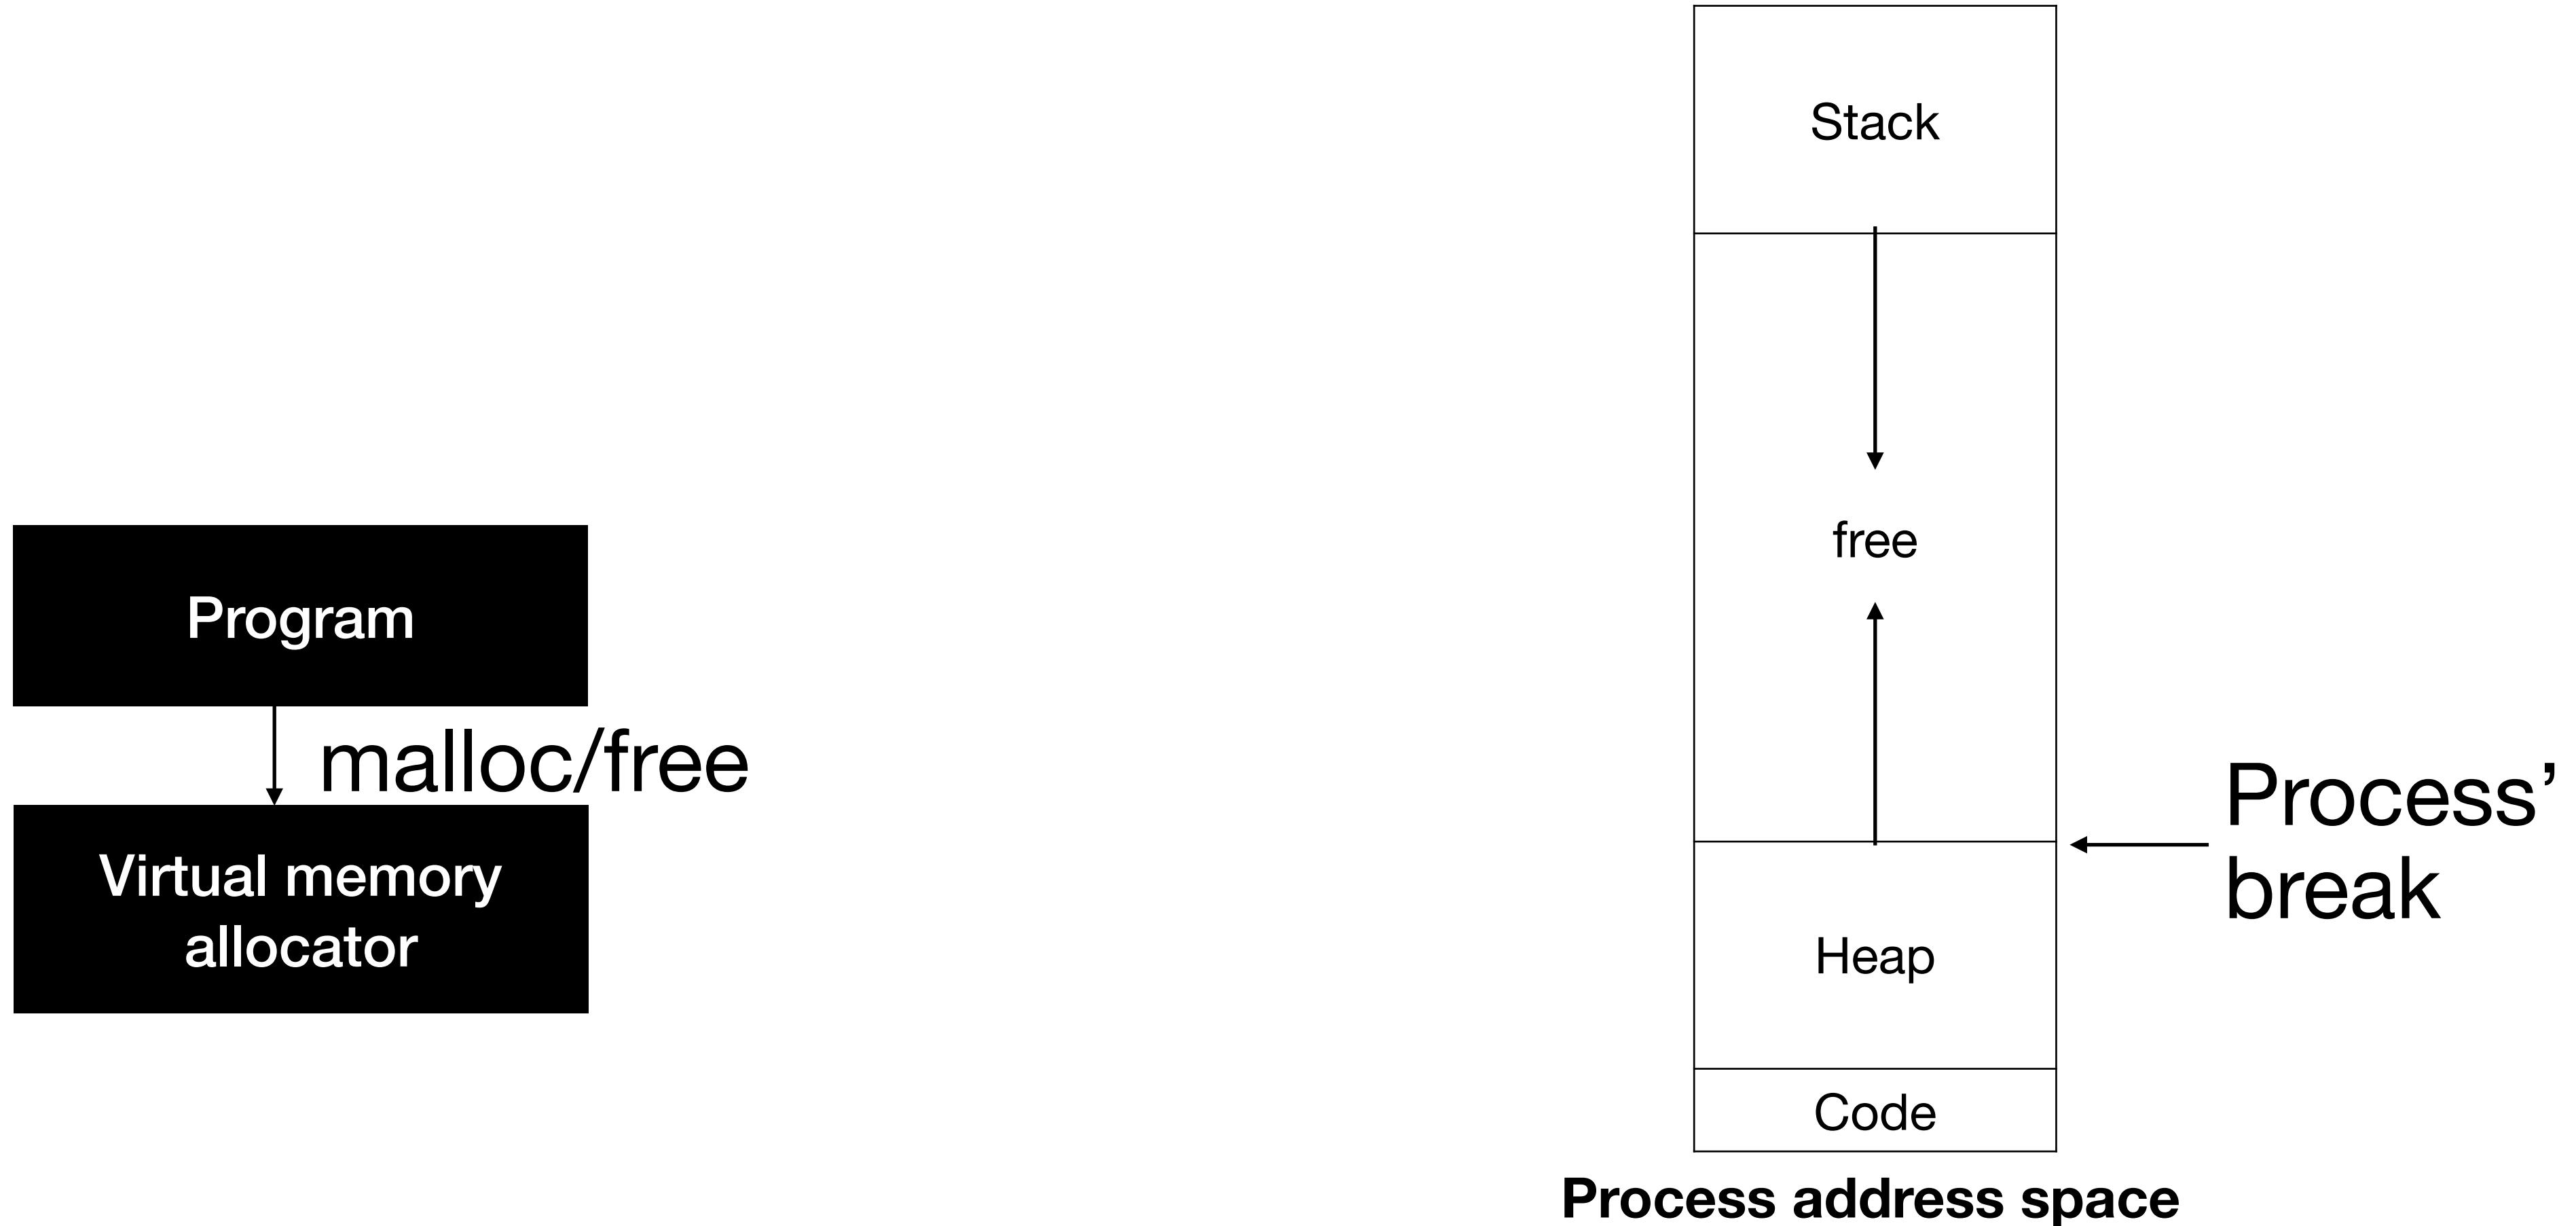


Manages heap memory



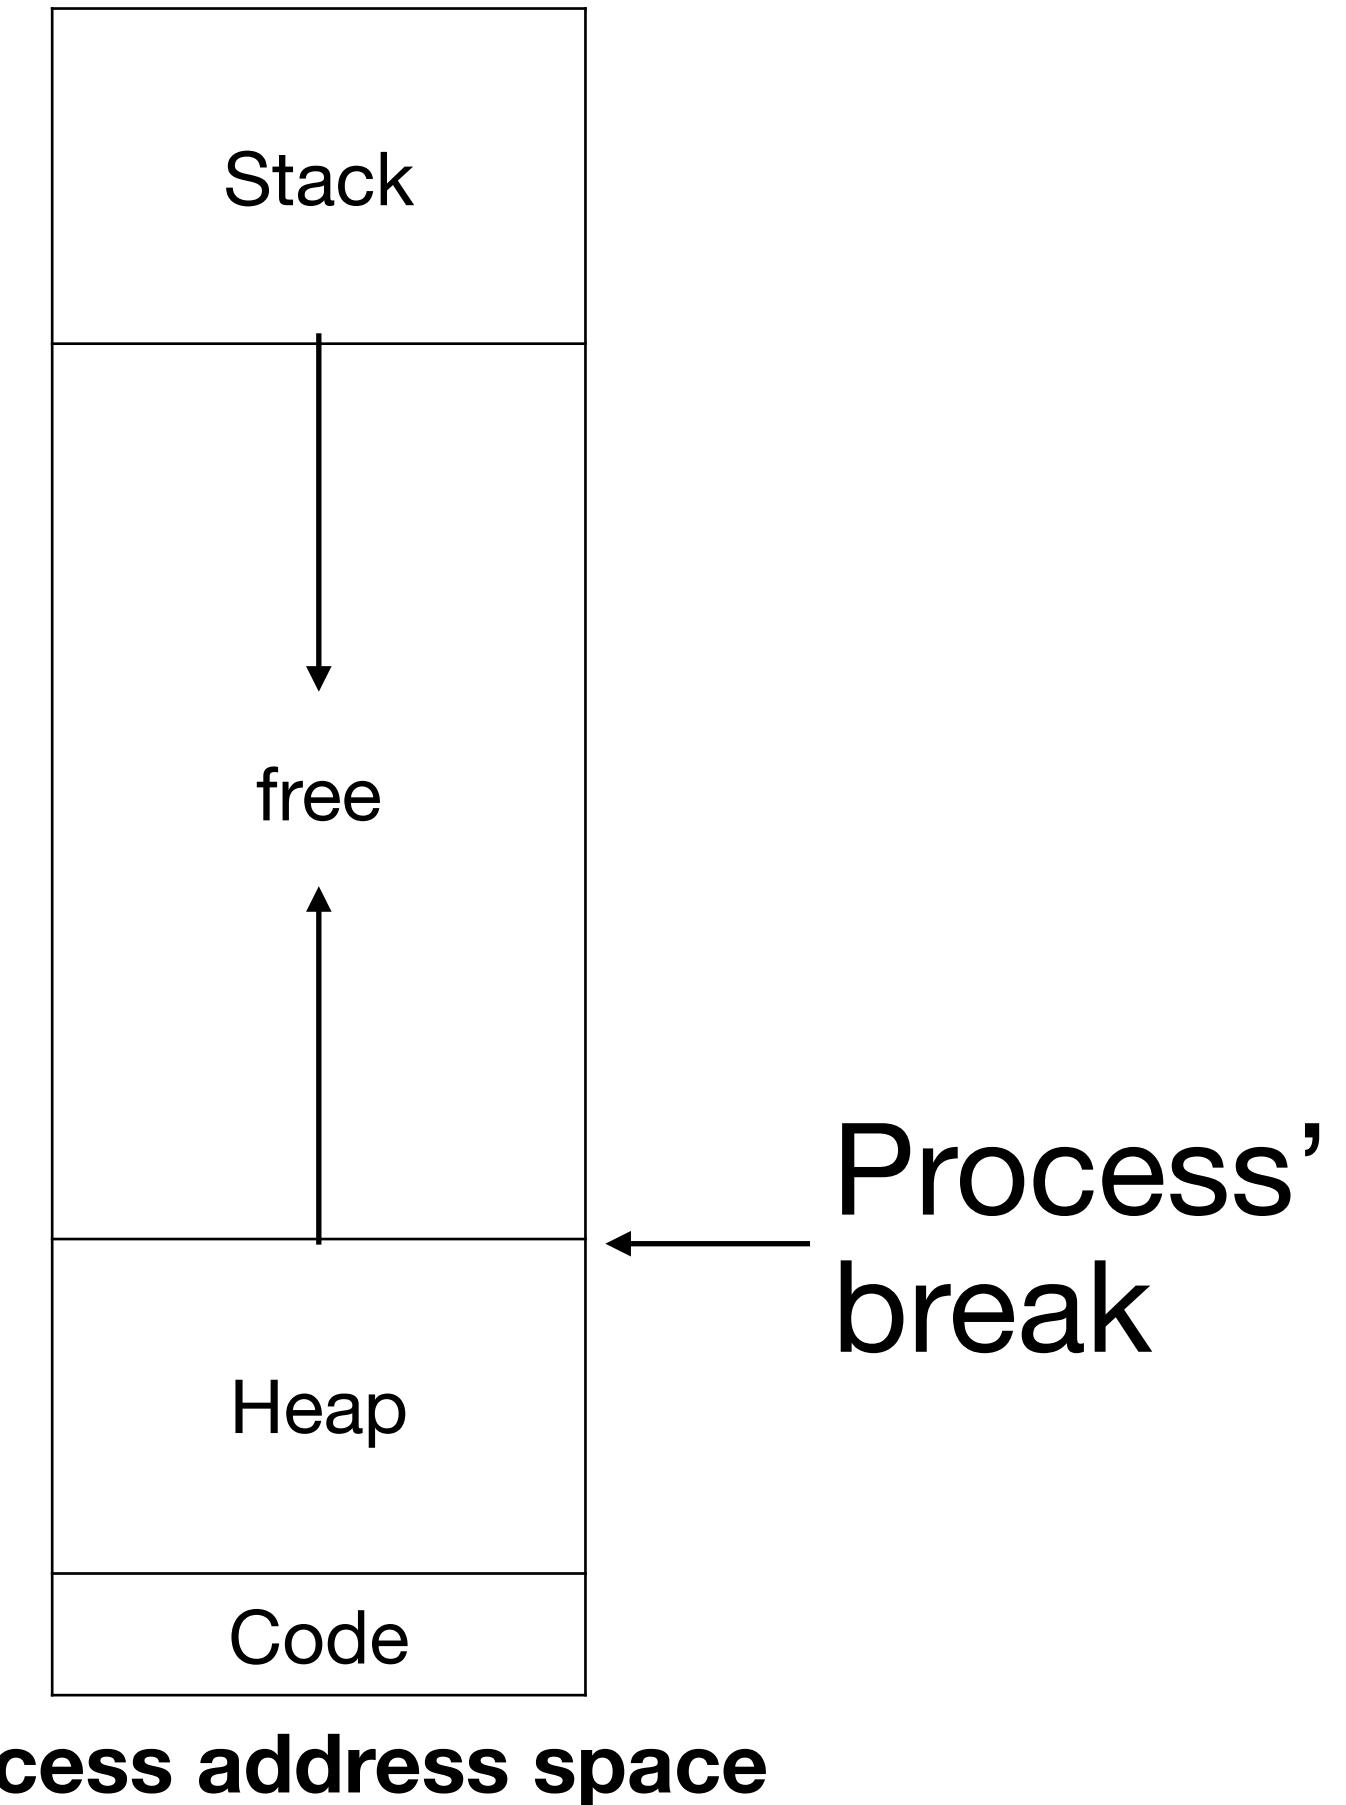
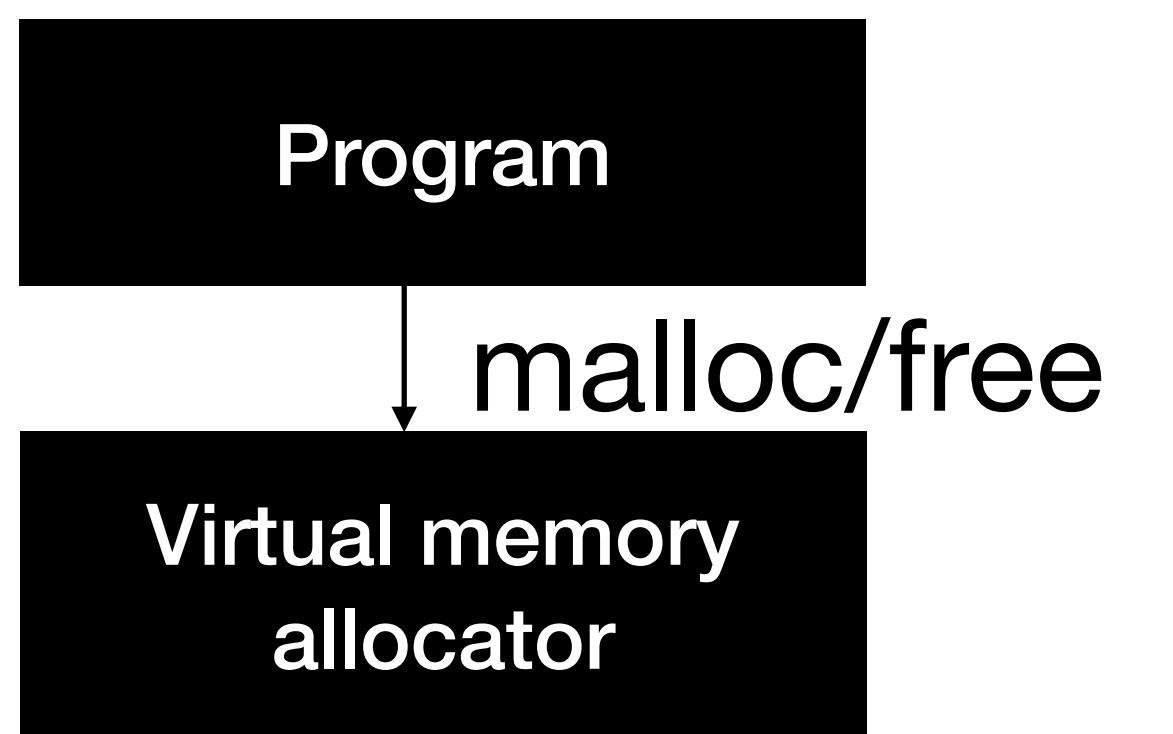
Process address space

# OS memory allocator



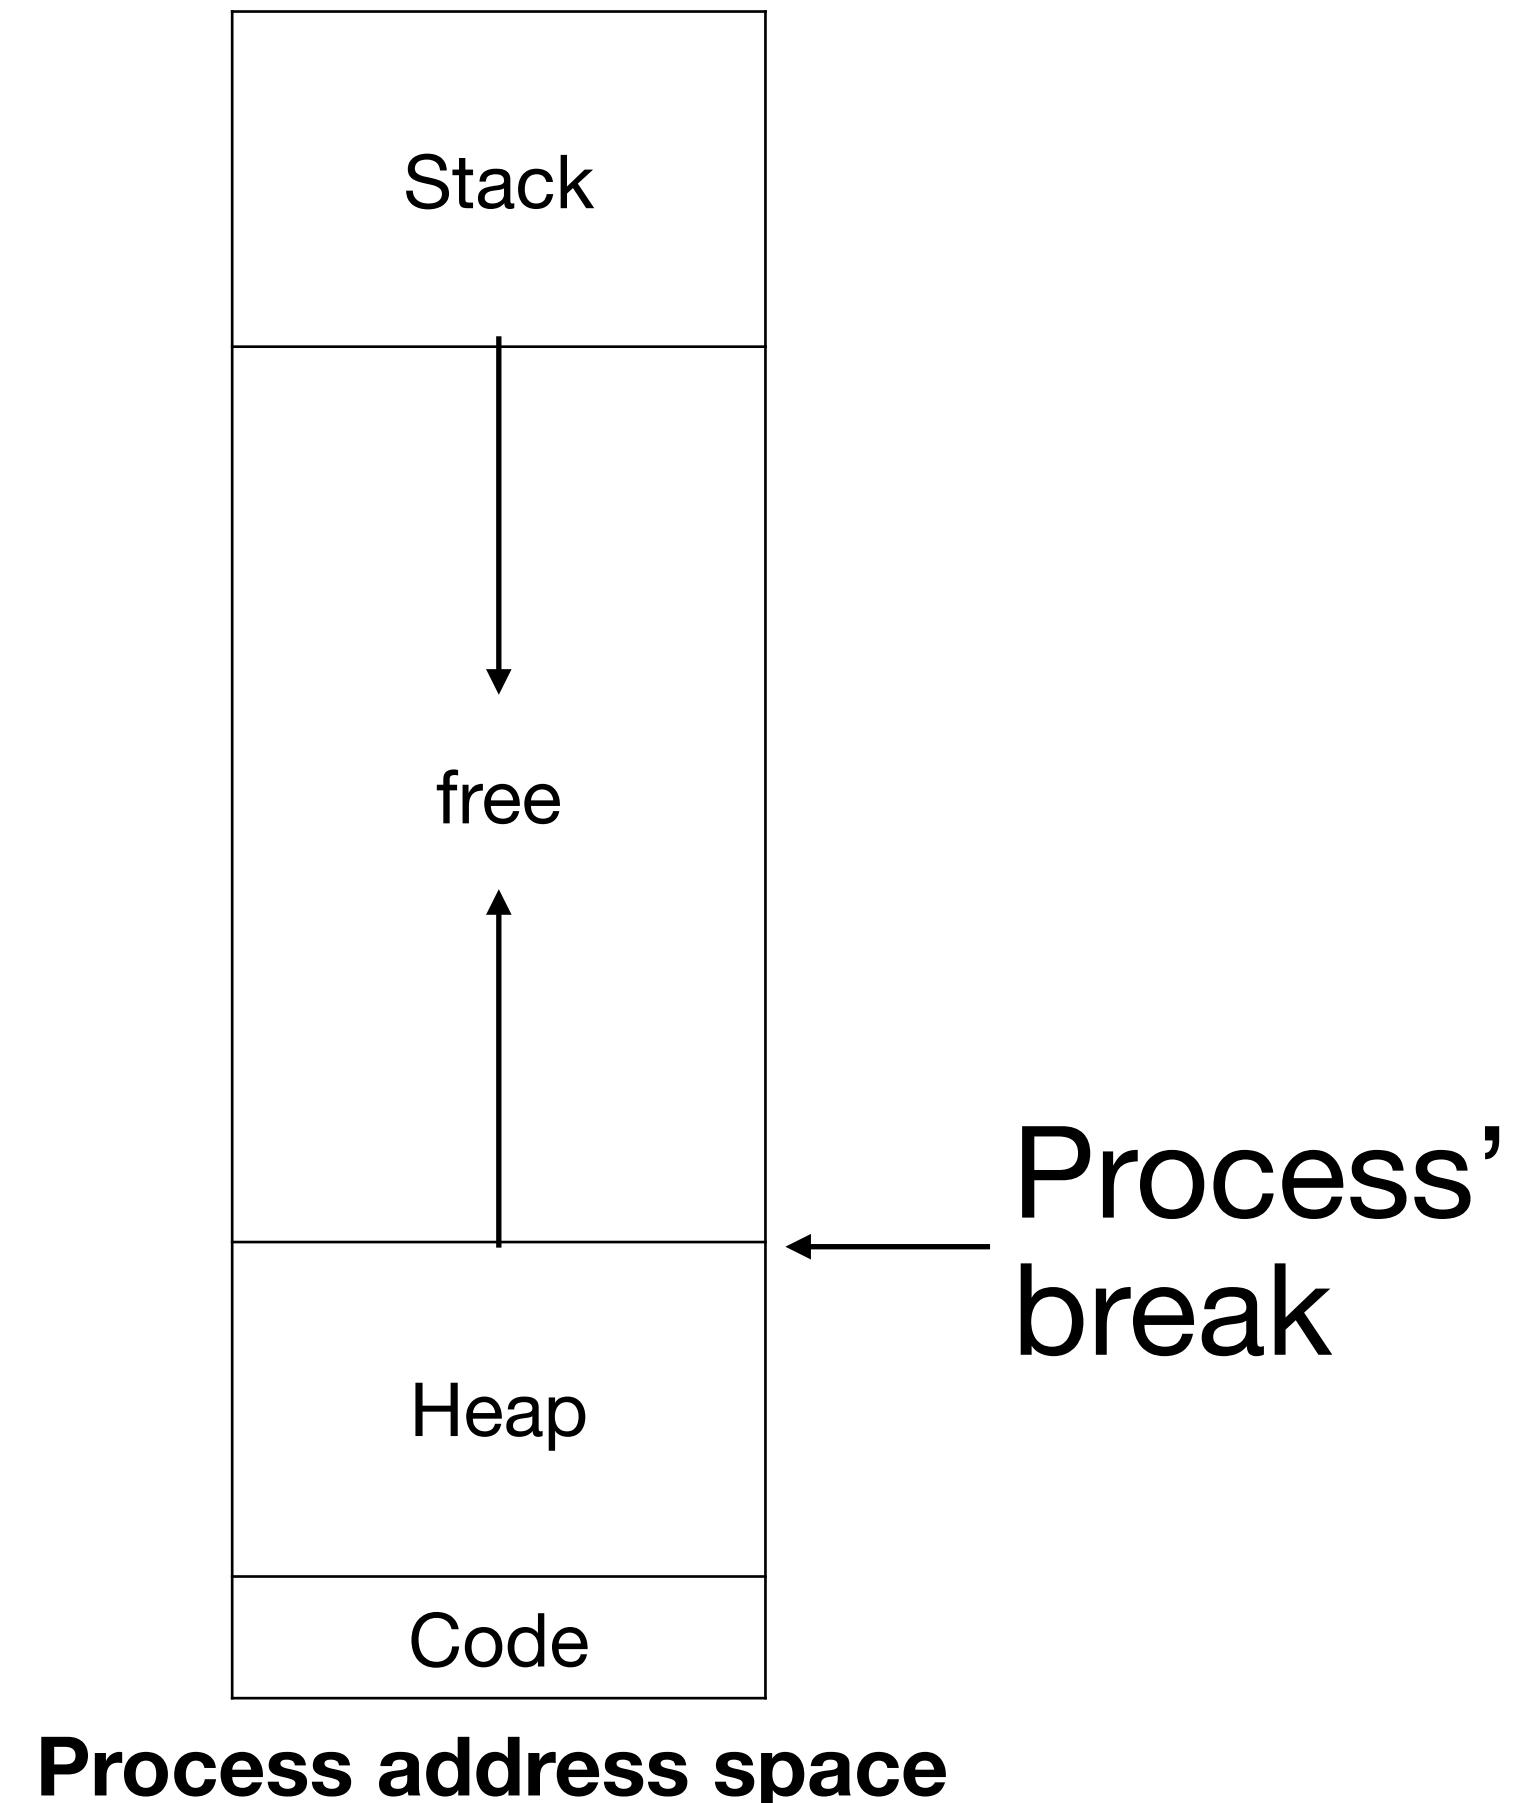
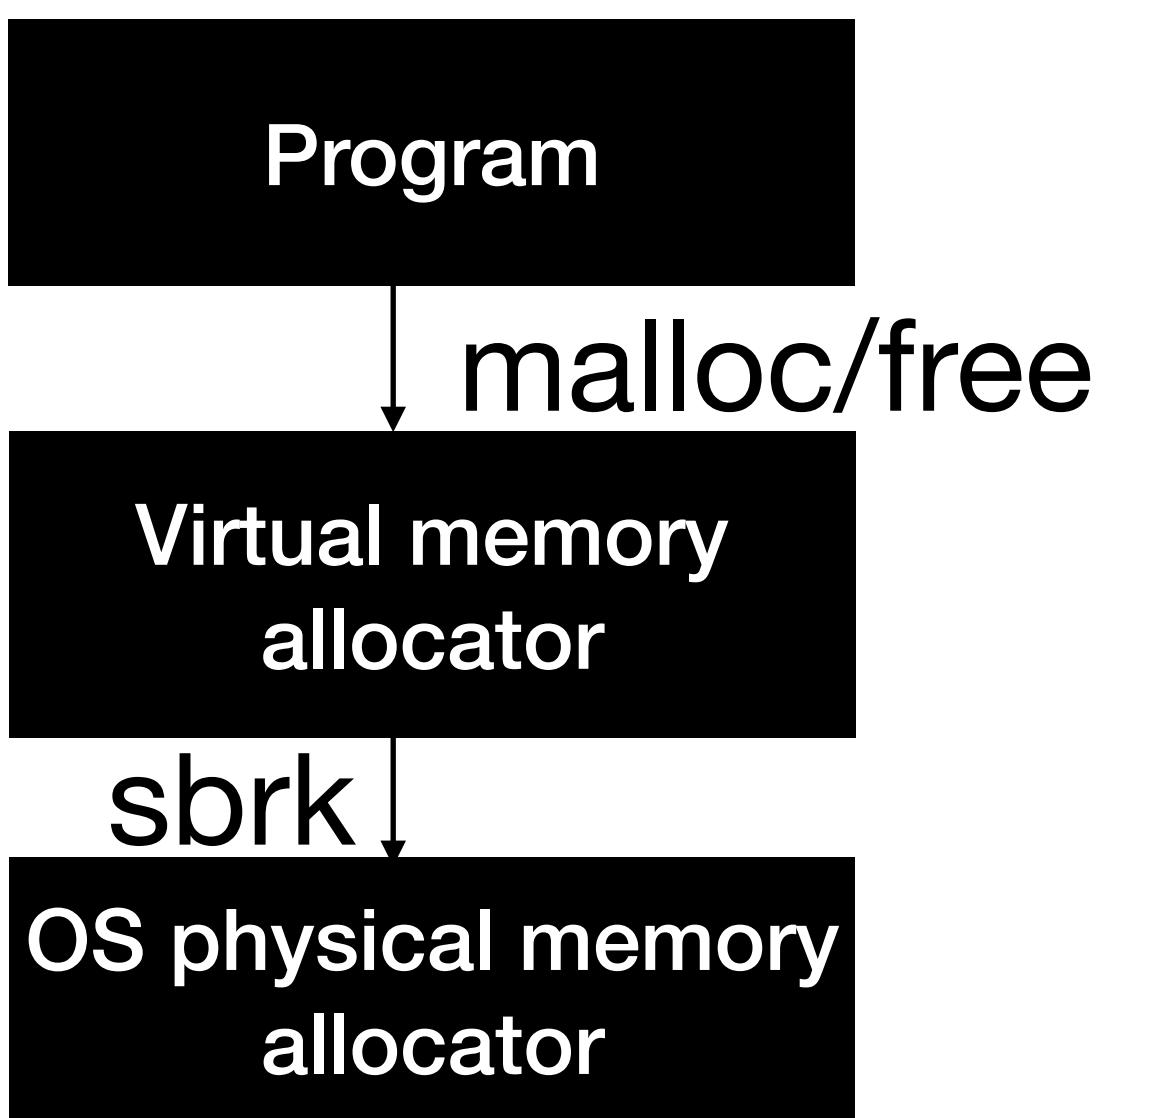
# OS memory allocator

- `sbrk(int increment)` increments process' break. *increment* can be negative.



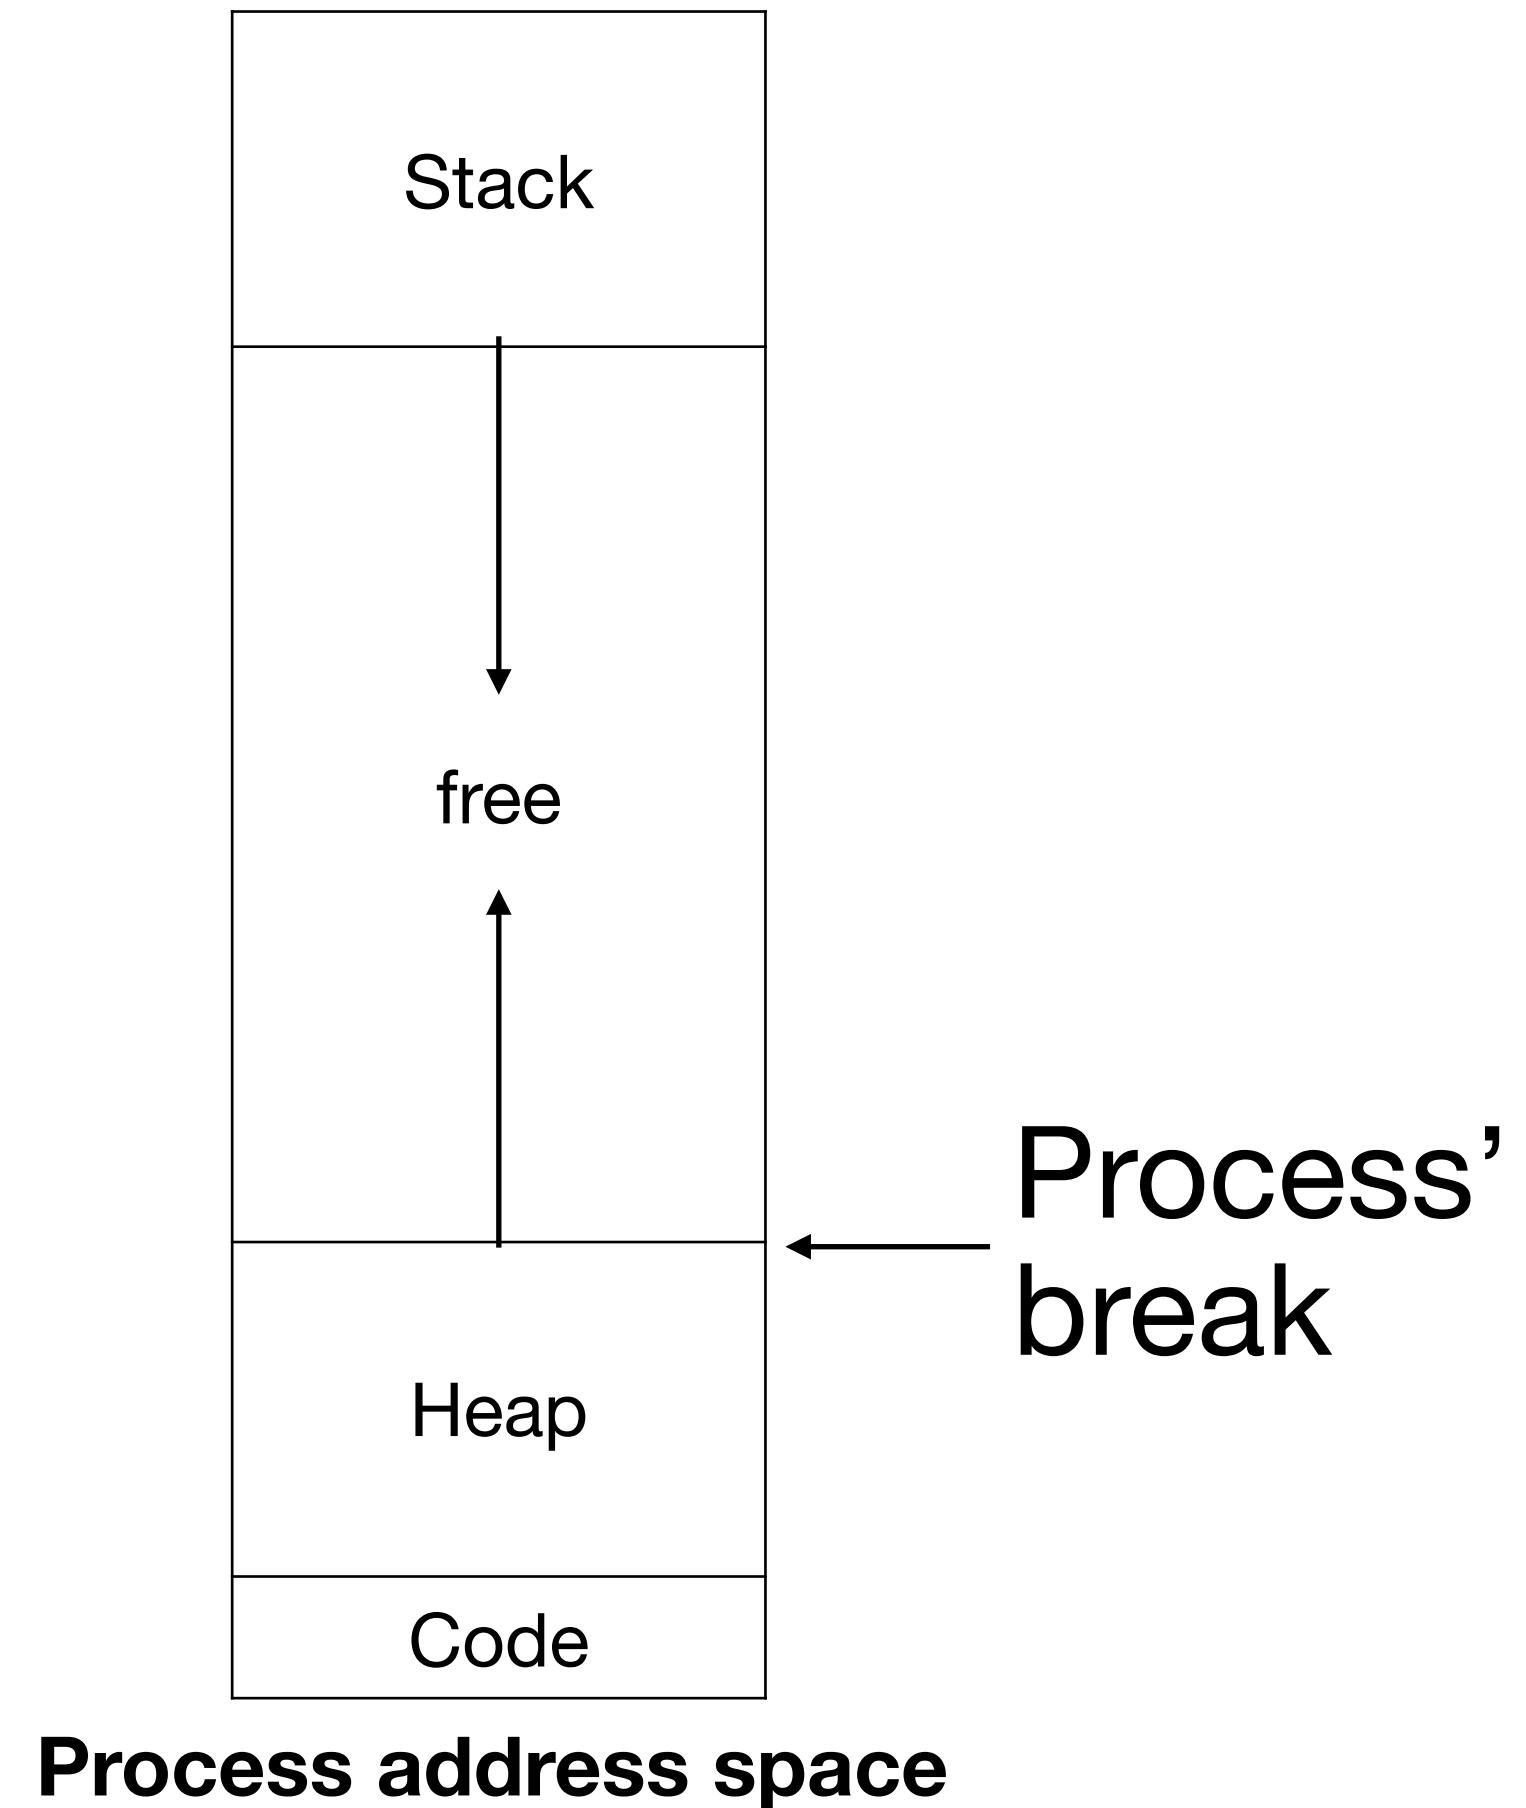
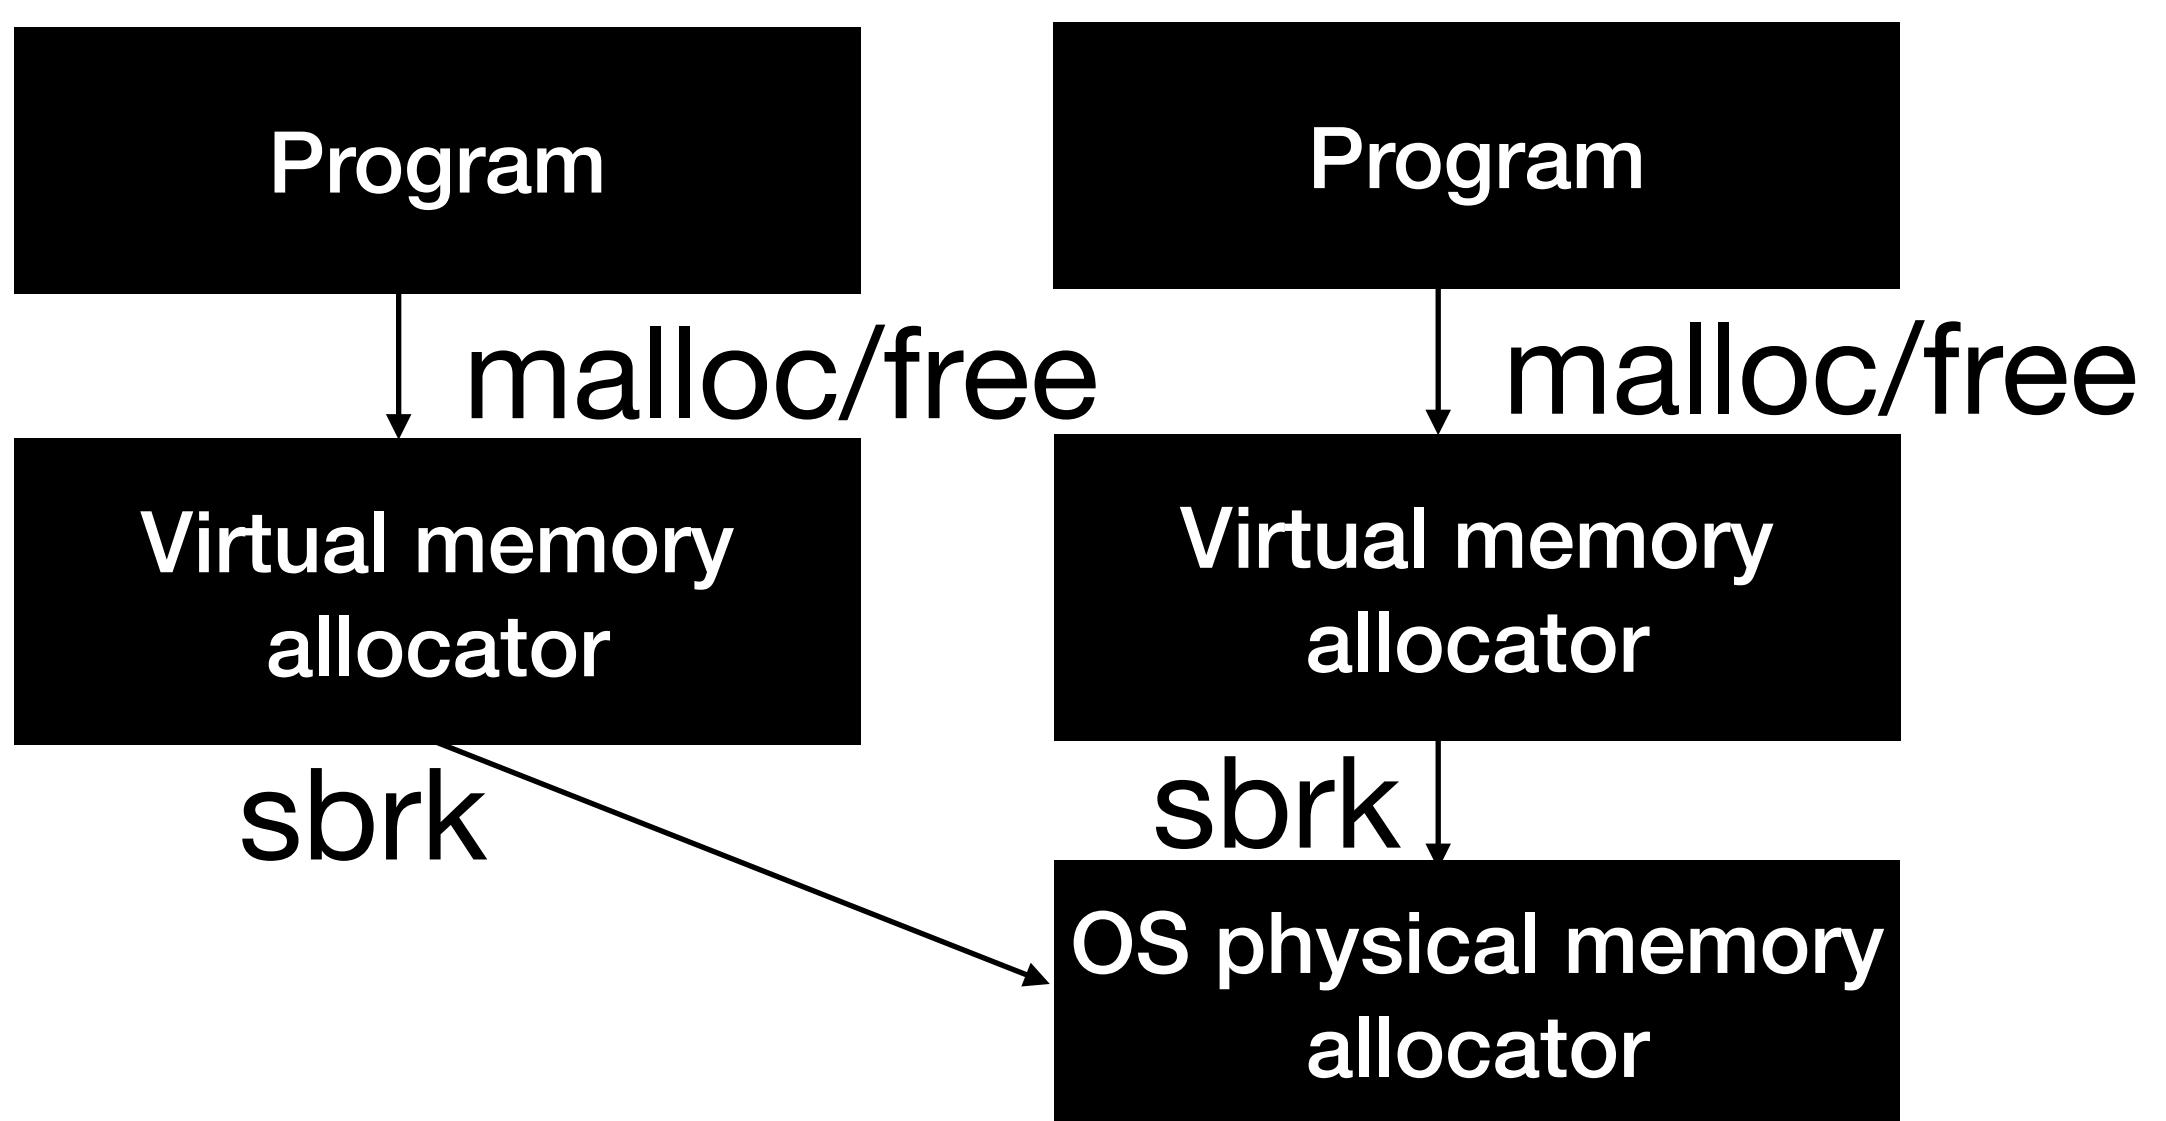
# OS memory allocator

- `sbrk(int increment)` increments process' break. *increment* can be negative.



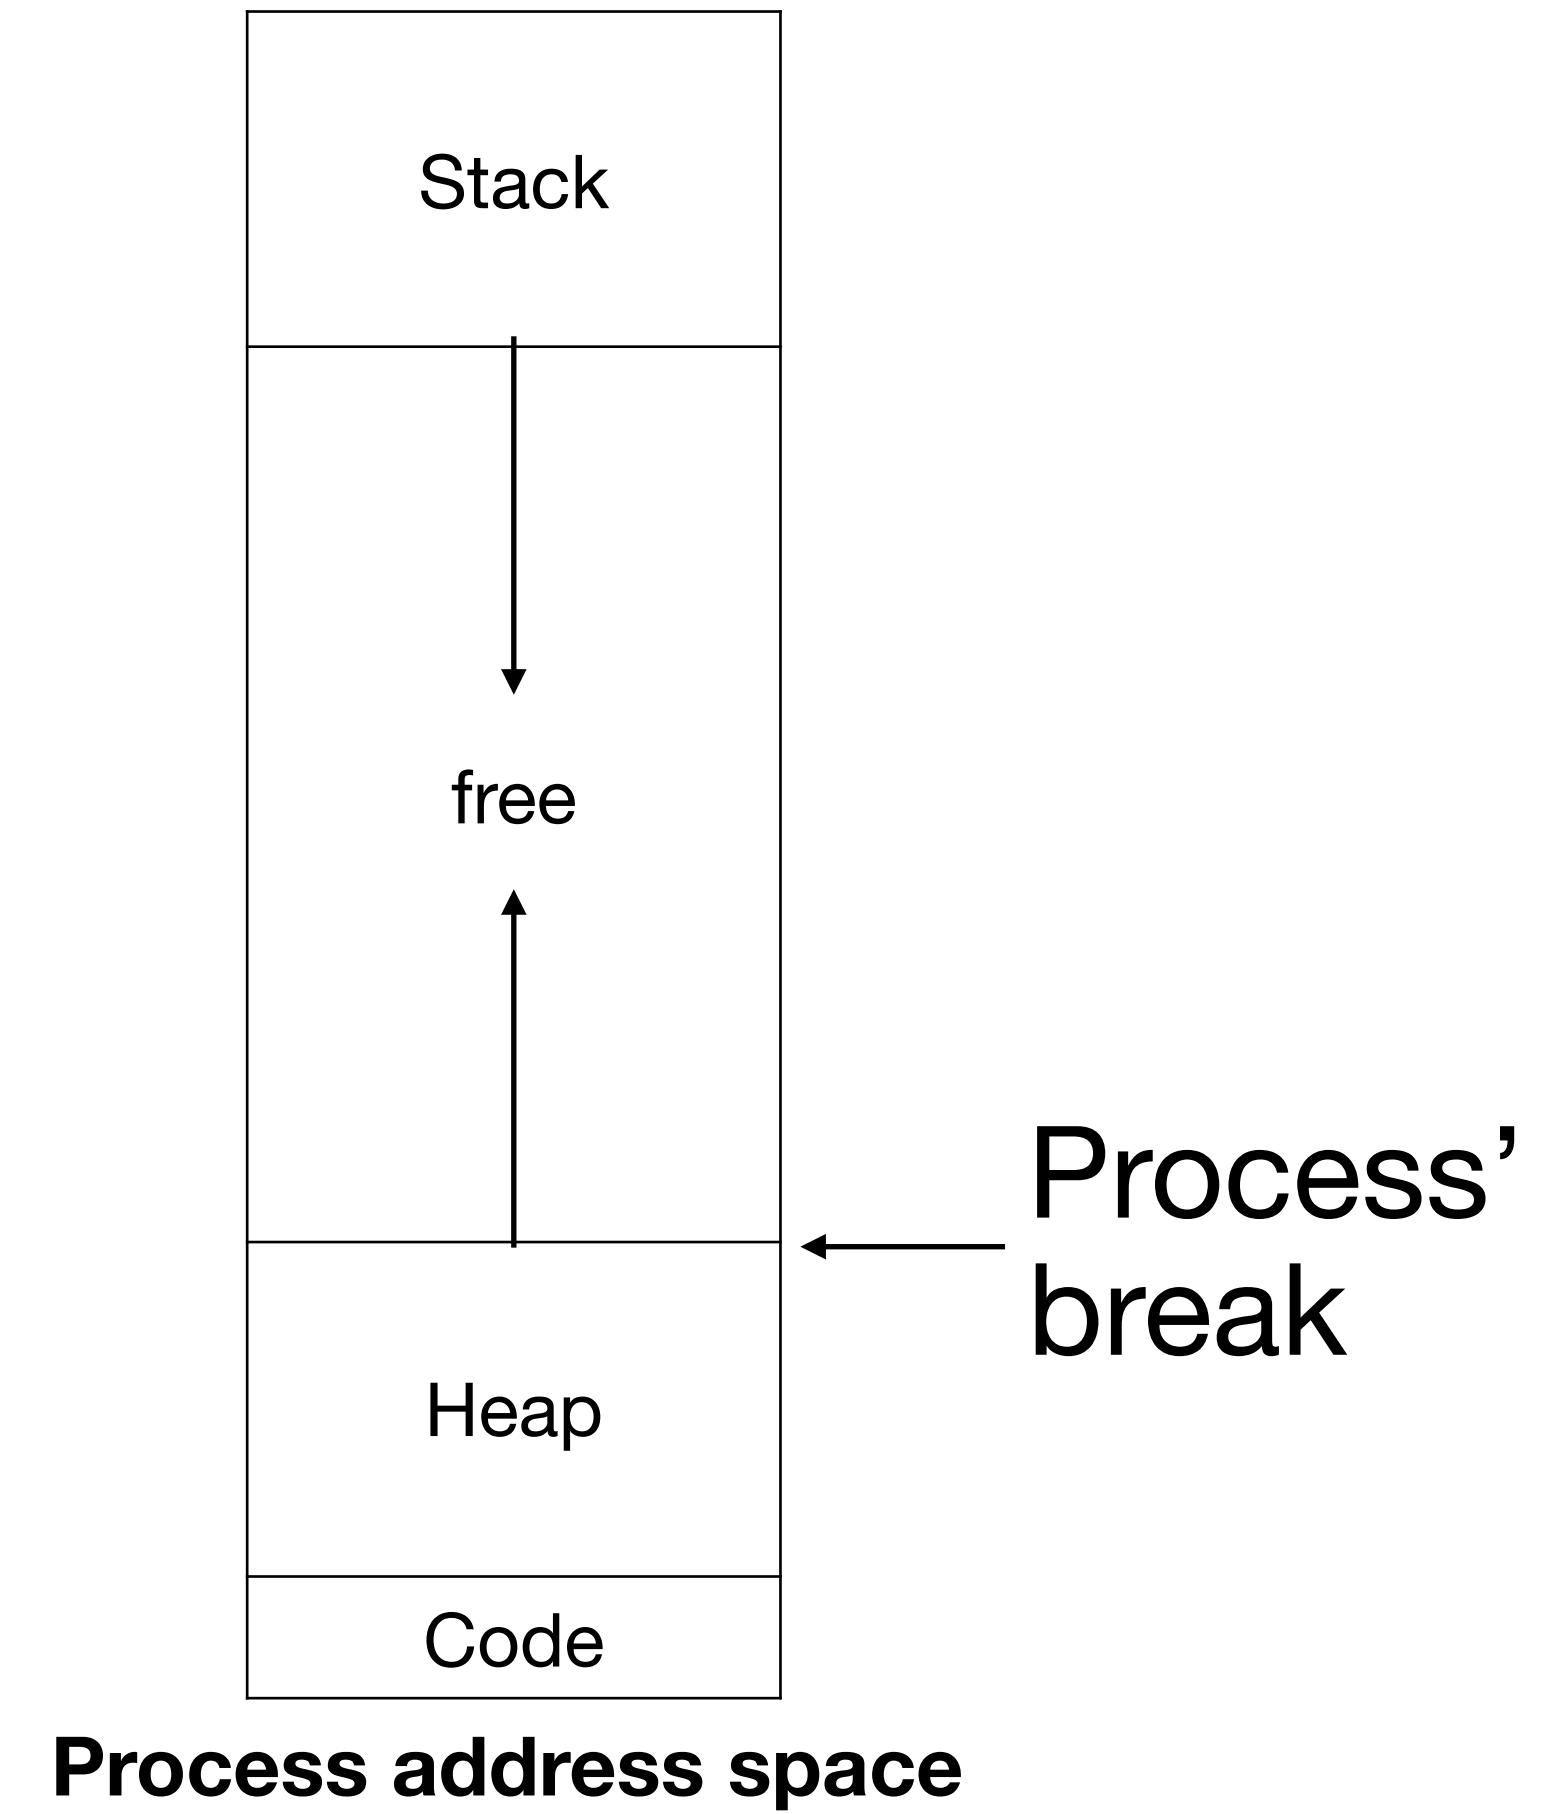
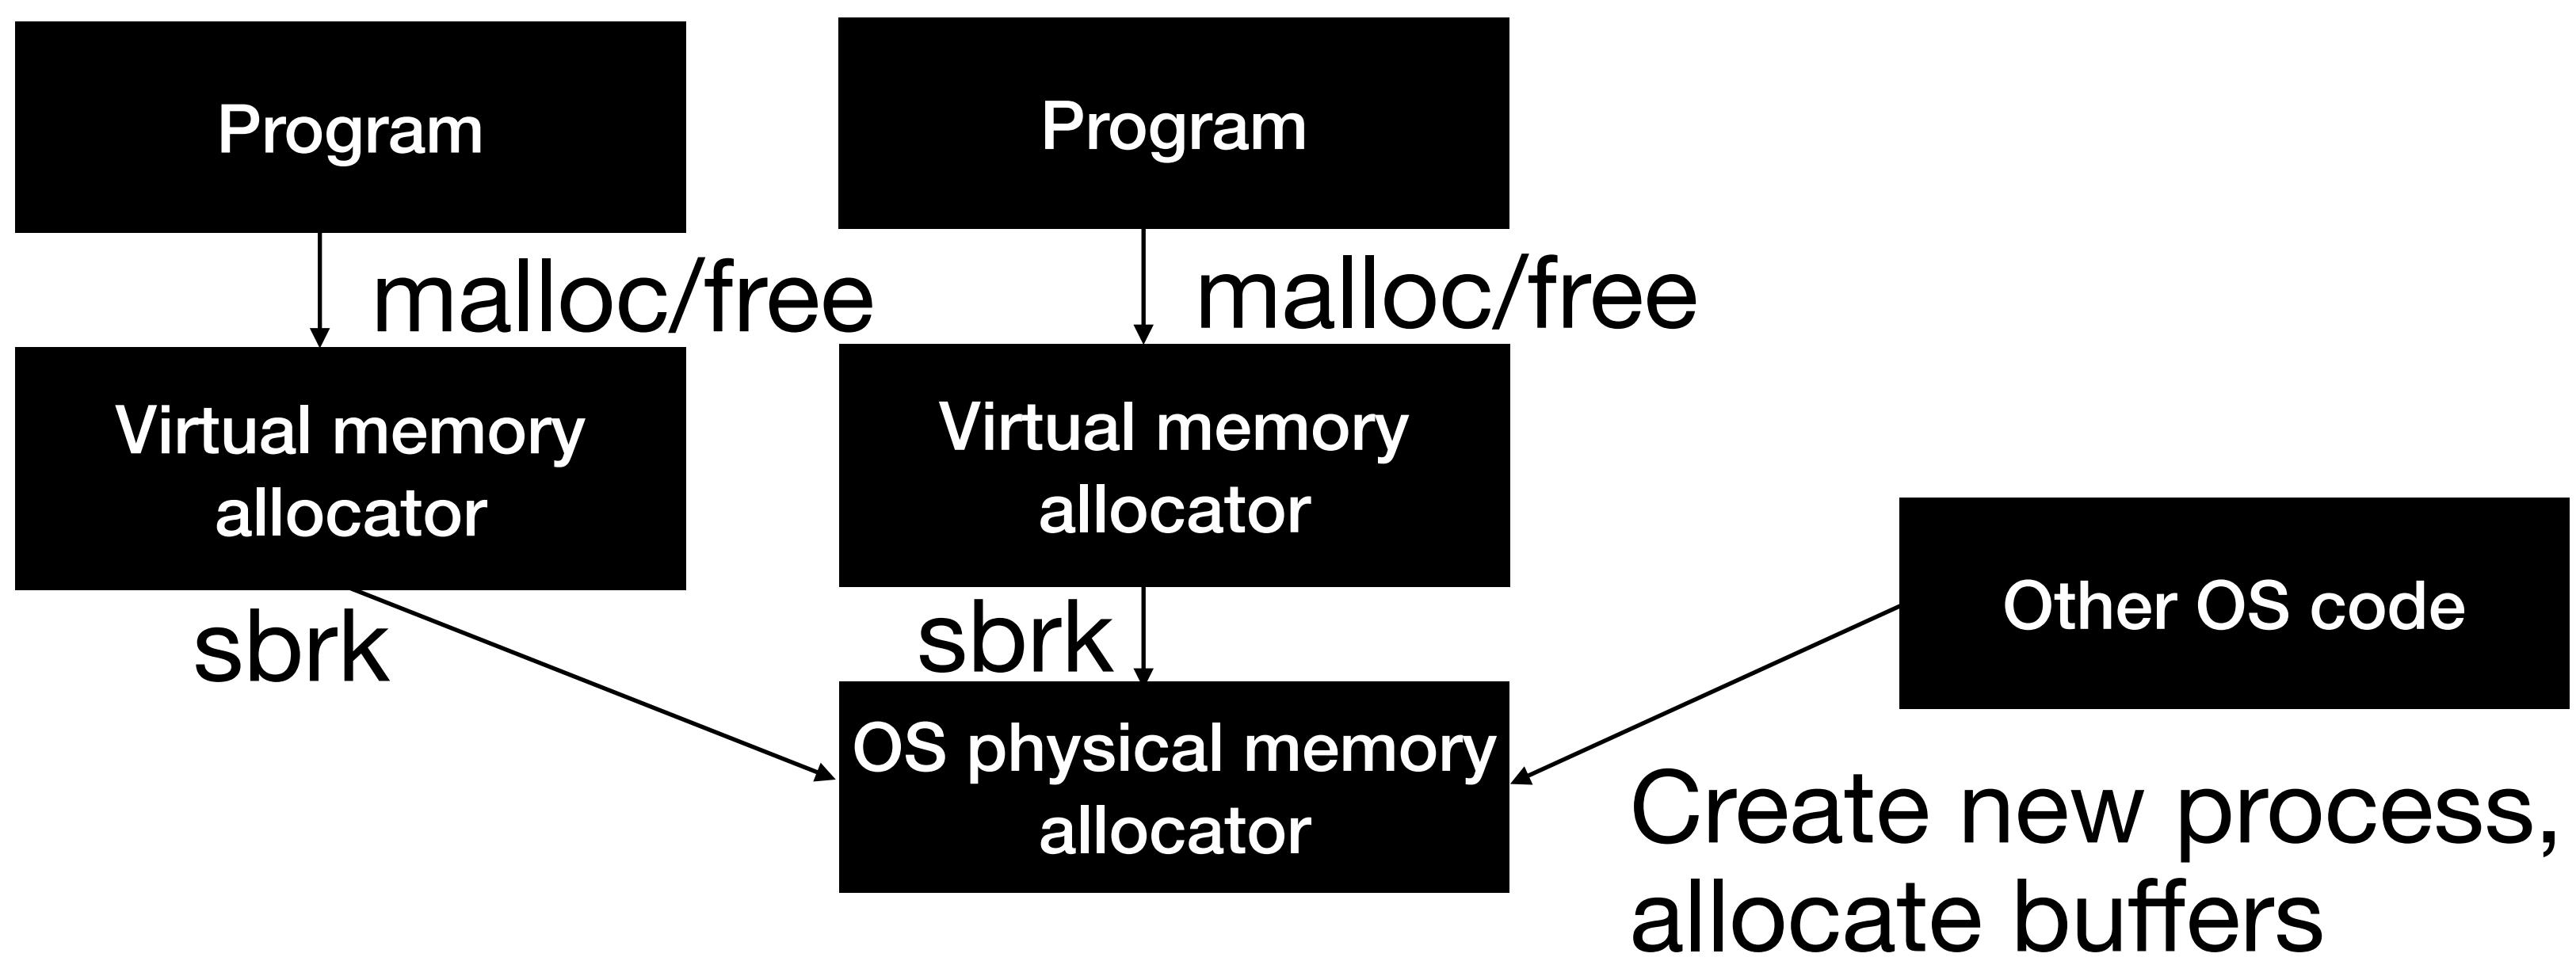
# OS memory allocator

- `sbrk(int increment)` increments process' break. *increment* can be negative.



# OS memory allocator

- `sbrk(int increment)` increments process' break. *increment* can be negative.



# Memory allocation

```
ptr=malloc(size_t size);
```

```
free(ptr);
```

# Memory allocation

```
ptr=malloc(size_t size);  
  
free(ptr);
```

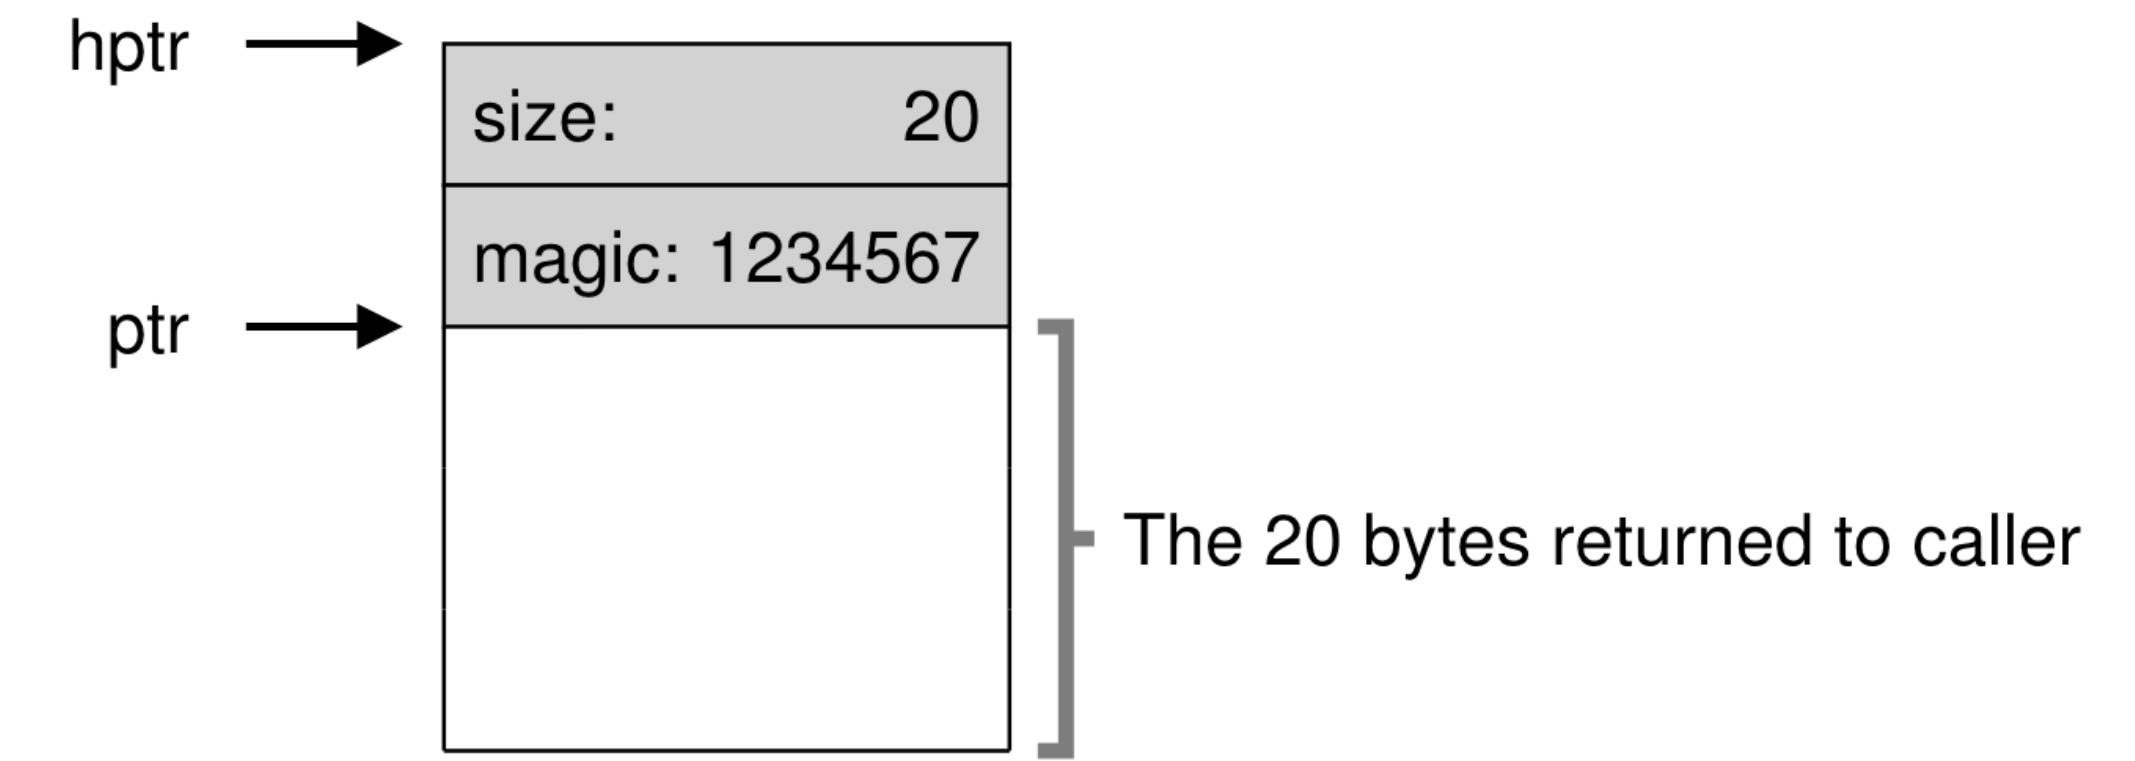


Figure 17.2: Specific Contents Of The Header

# Memory allocation

```
ptr=malloc(size_t size);  
free(ptr);
```

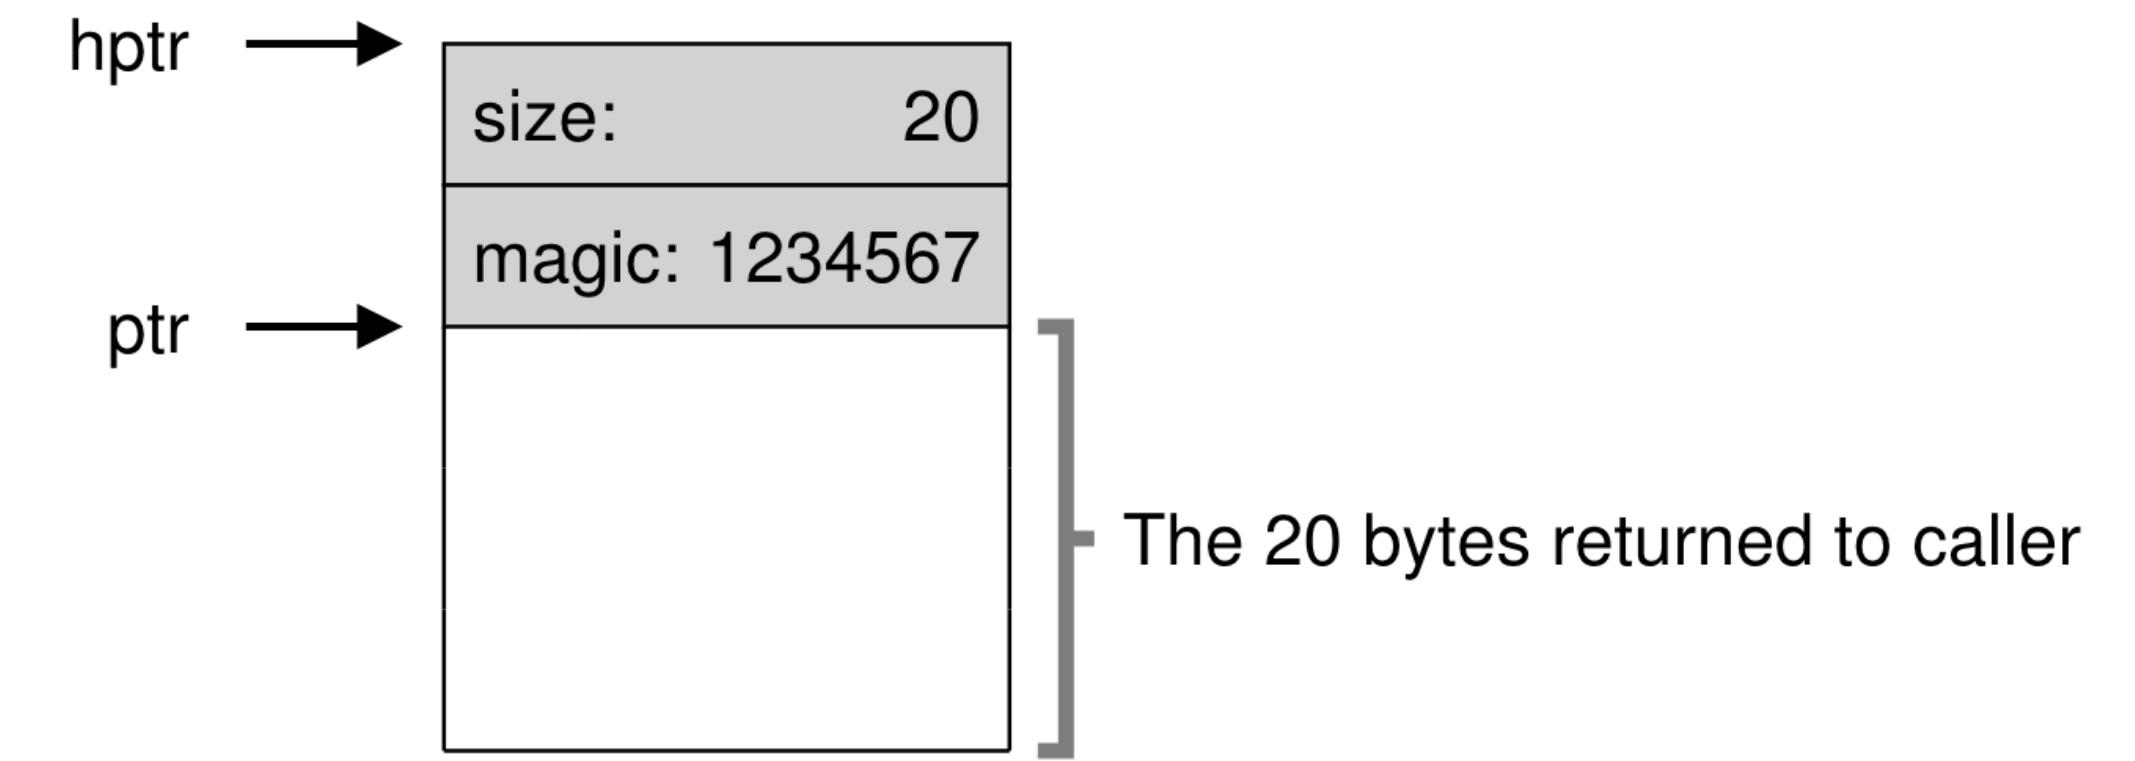
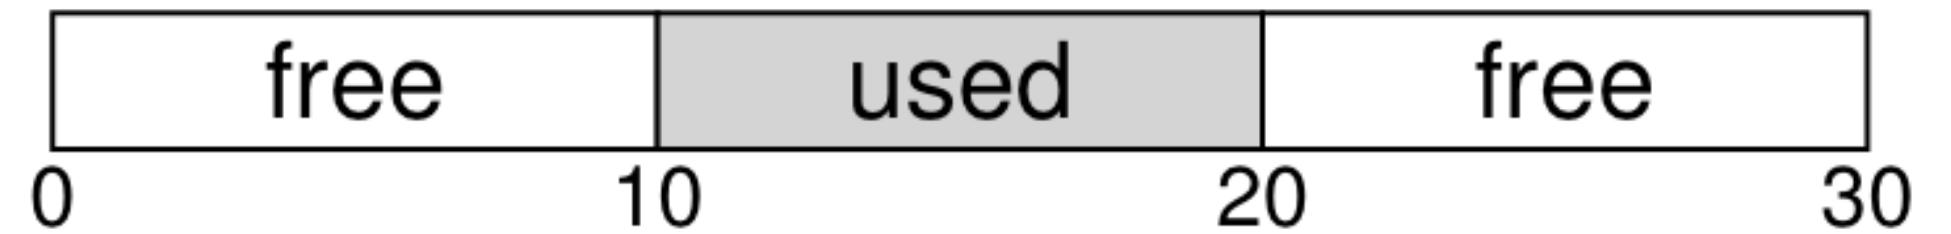


Figure 17.2: Specific Contents Of The Header

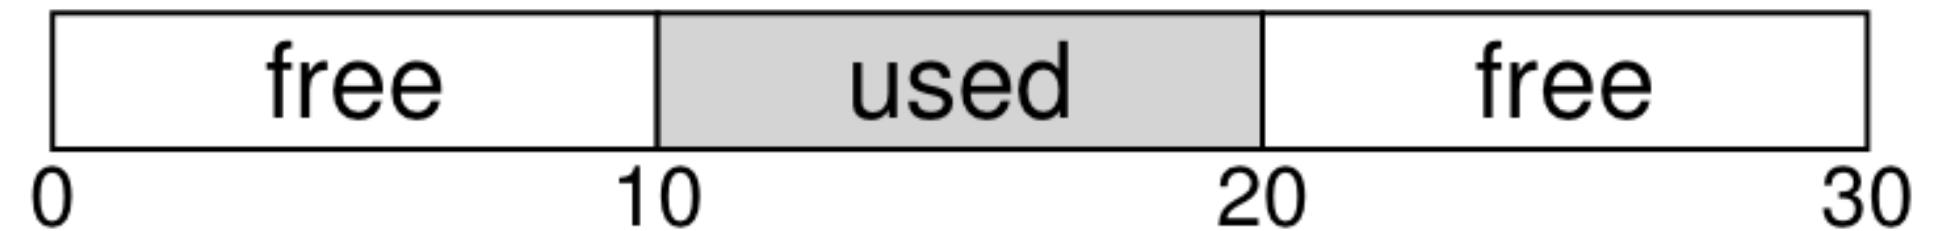
Stack allocations are faster than heap allocations. No need to find space.

# Memory allocator



Fragmented heap over time

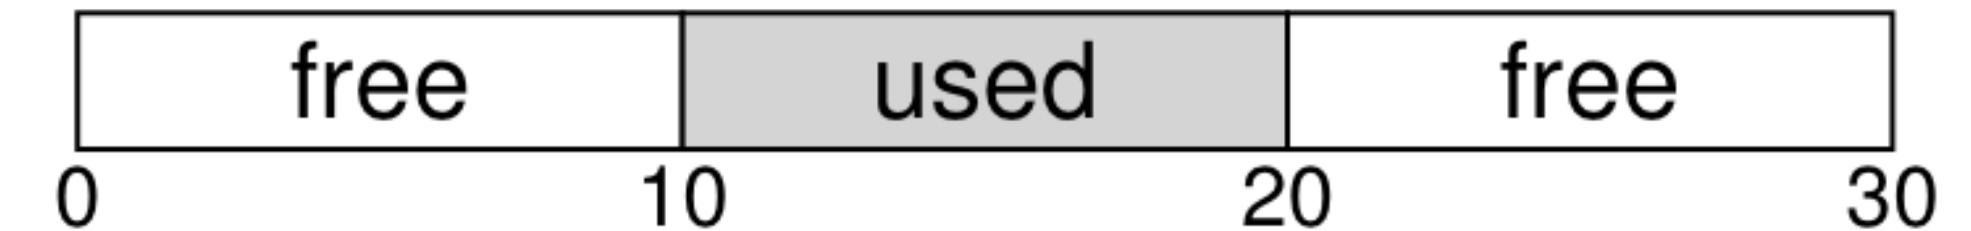
# Memory allocator



Fragmented heap over time

- Assumptions
  - Do not apriori know allocation size and order
  - Cannot move memory once it is allocated. Program might have the pointer to it.

# Memory allocator



Fragmented heap over time

- Assumptions
  - Do not apriori know allocation size and order
  - Cannot move memory once it is allocated. Program might have the pointer to it.
- Goals
  - Quickly satisfy variable-sized memory allocation requests. How to track free memory?
  - Minimize fragmentation

# Memory (de)allocation patterns

- Small mallocs can be frequent. Large mallocs are usually infrequent.
  - After malloc, program will initialise the memory area.
- “Clustered deaths”: Objects allocated together die together.

# Free list splitting and coalescing

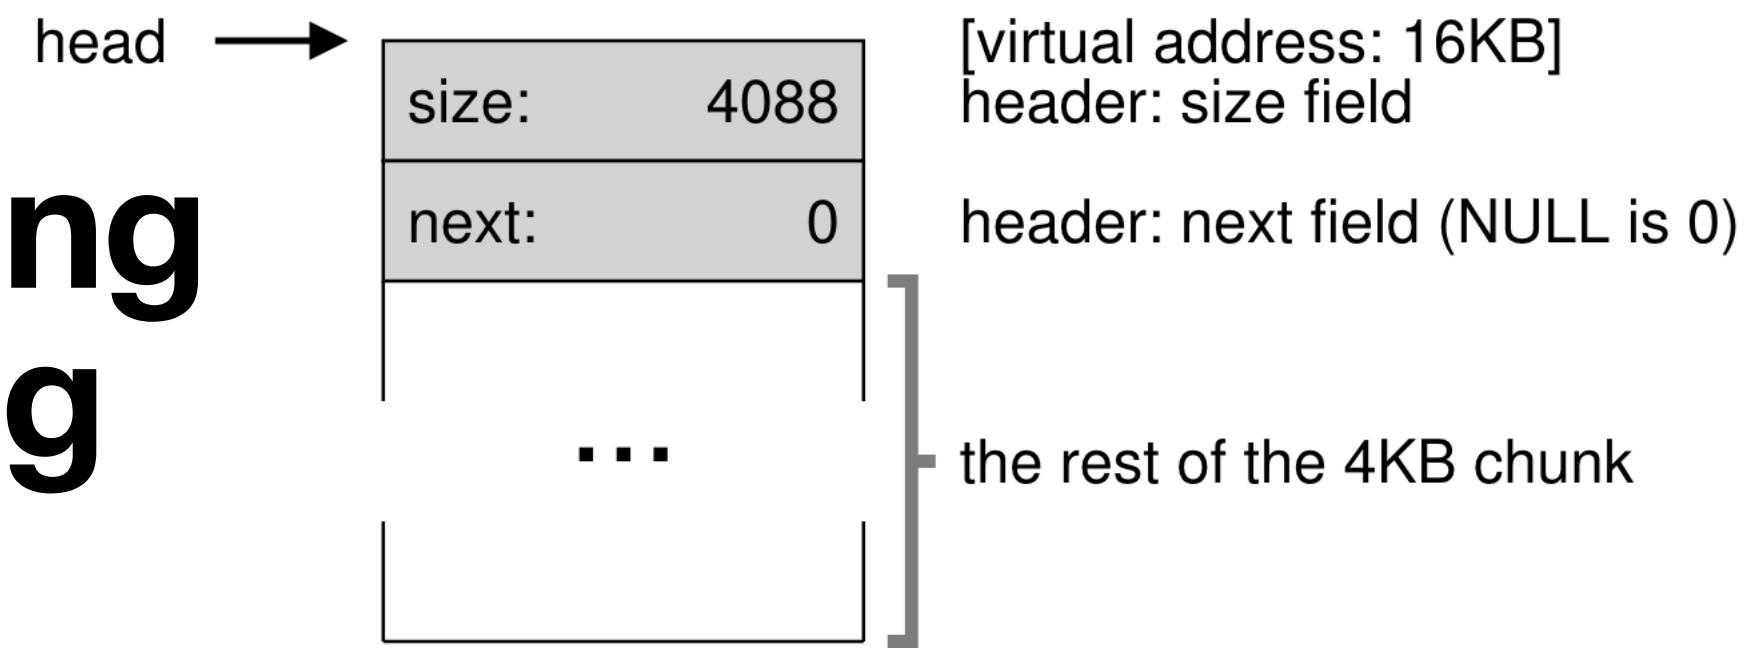


Figure 17.3: A Heap With One Free Chunk

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

# Free list splitting and coalescing

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

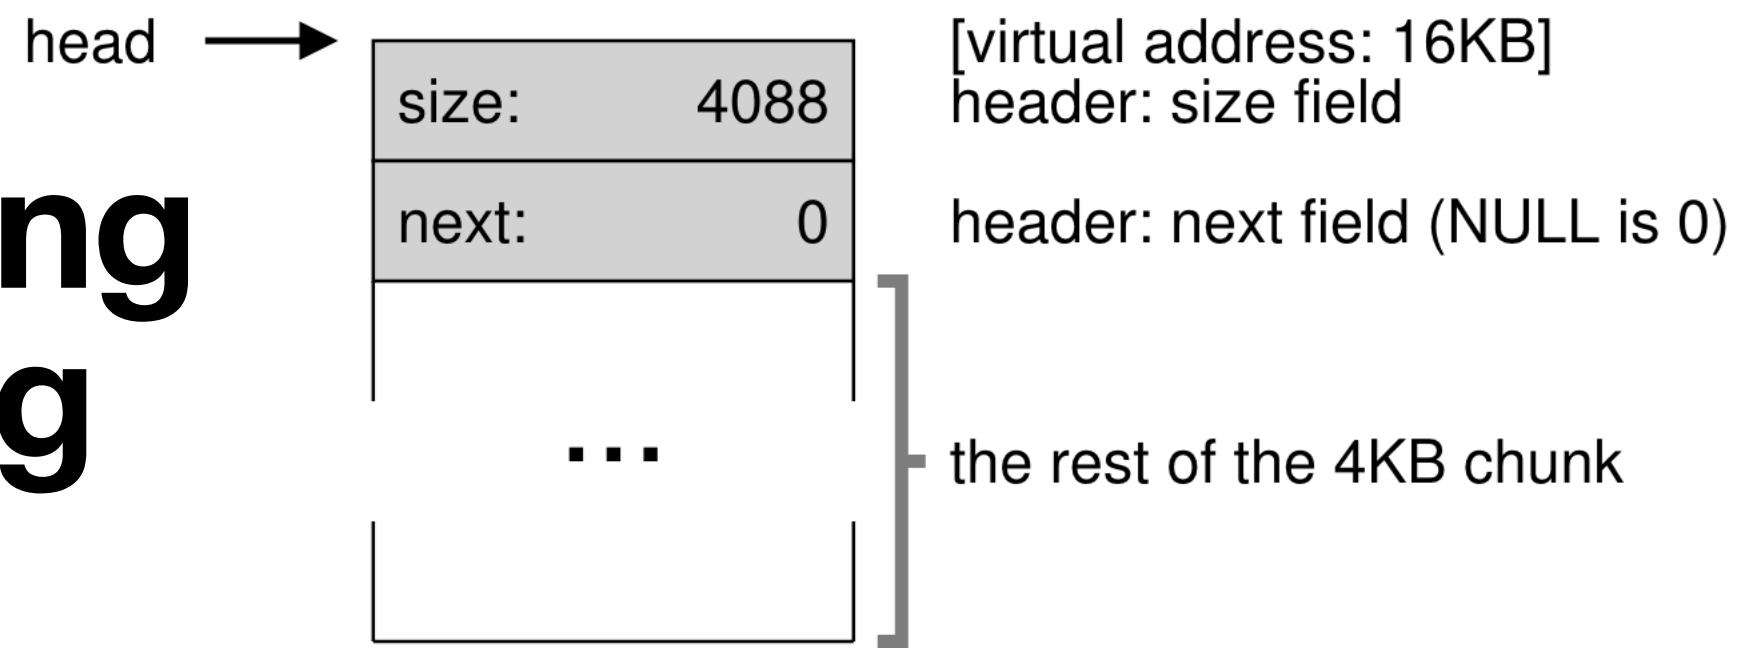


Figure 17.3: A Heap With One Free Chunk

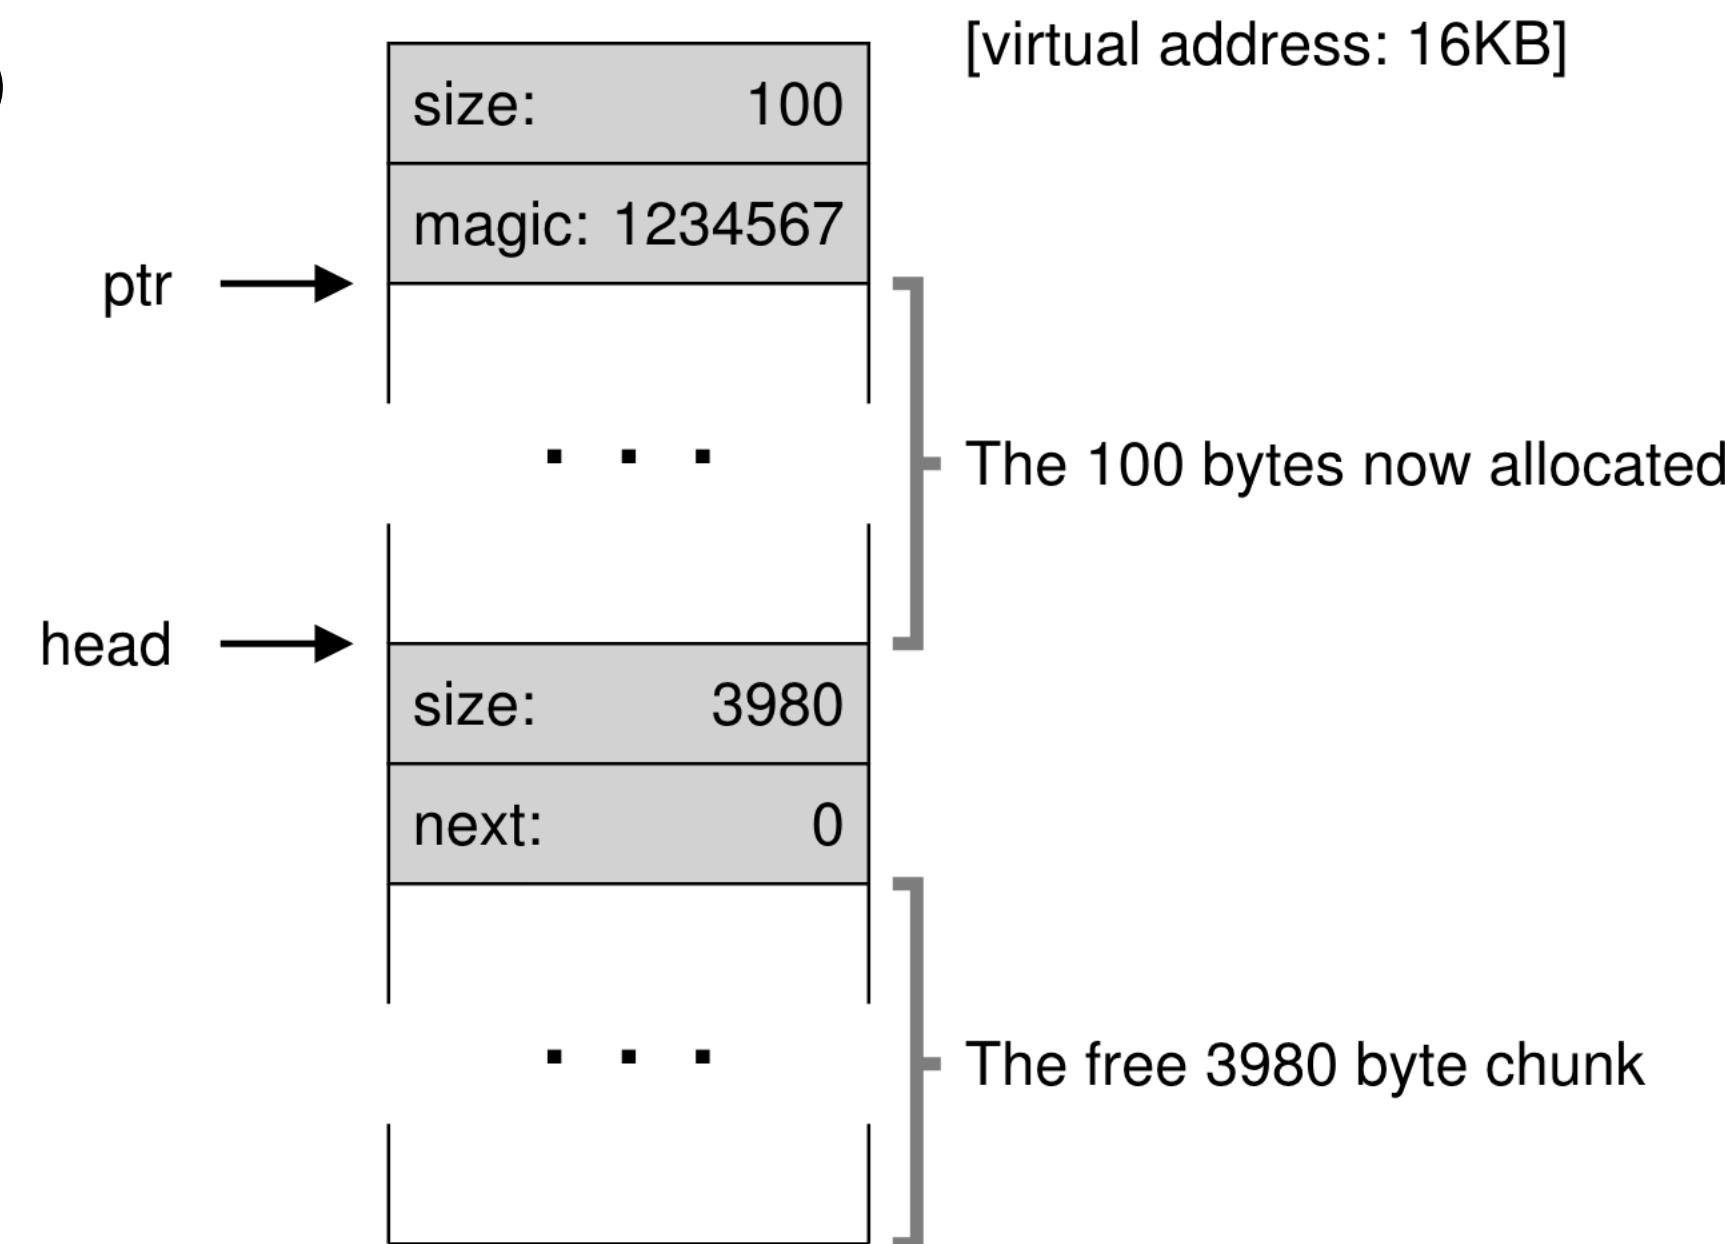


Figure 17.4: A Heap: After One Allocation

# Free list splitting and coalescing

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

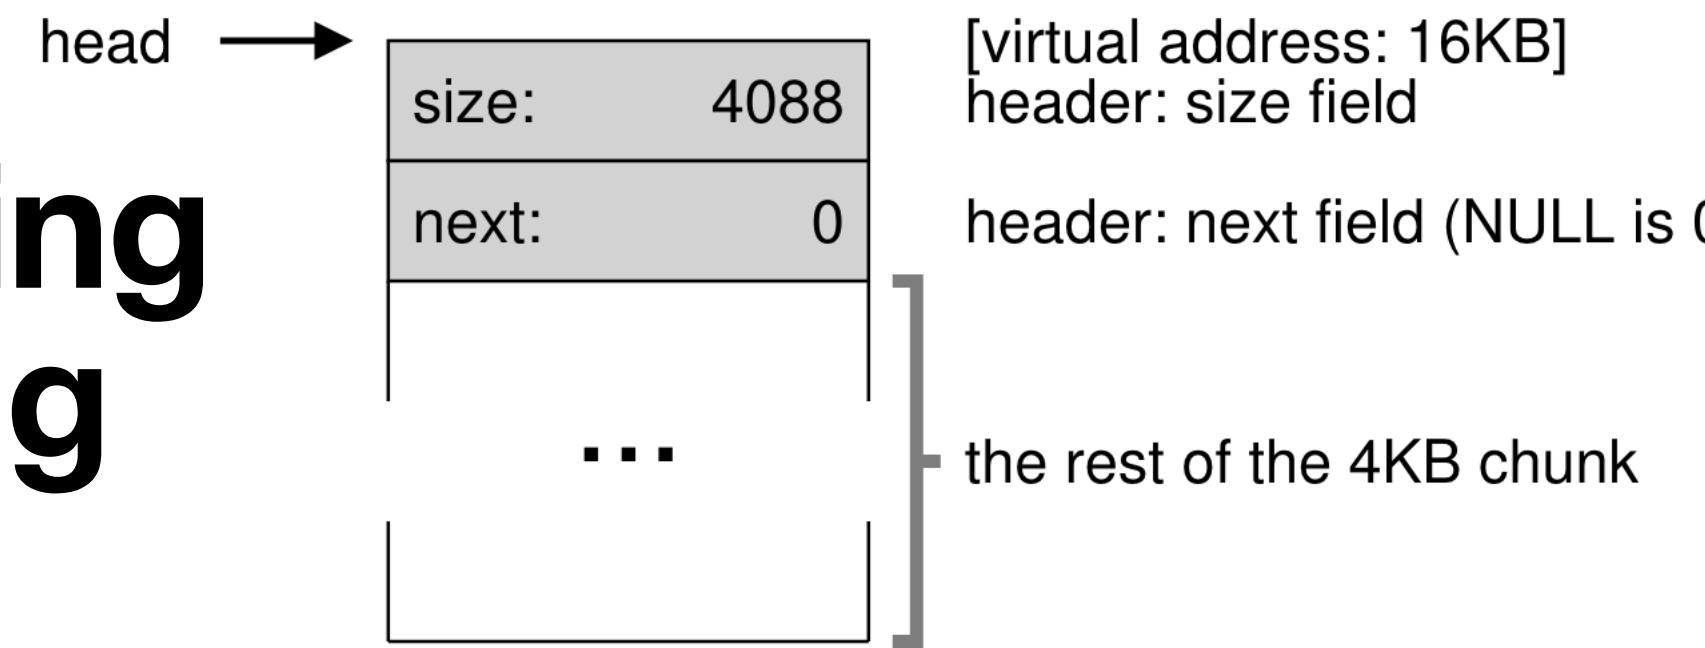


Figure 17.3: A Heap With One Free Chunk

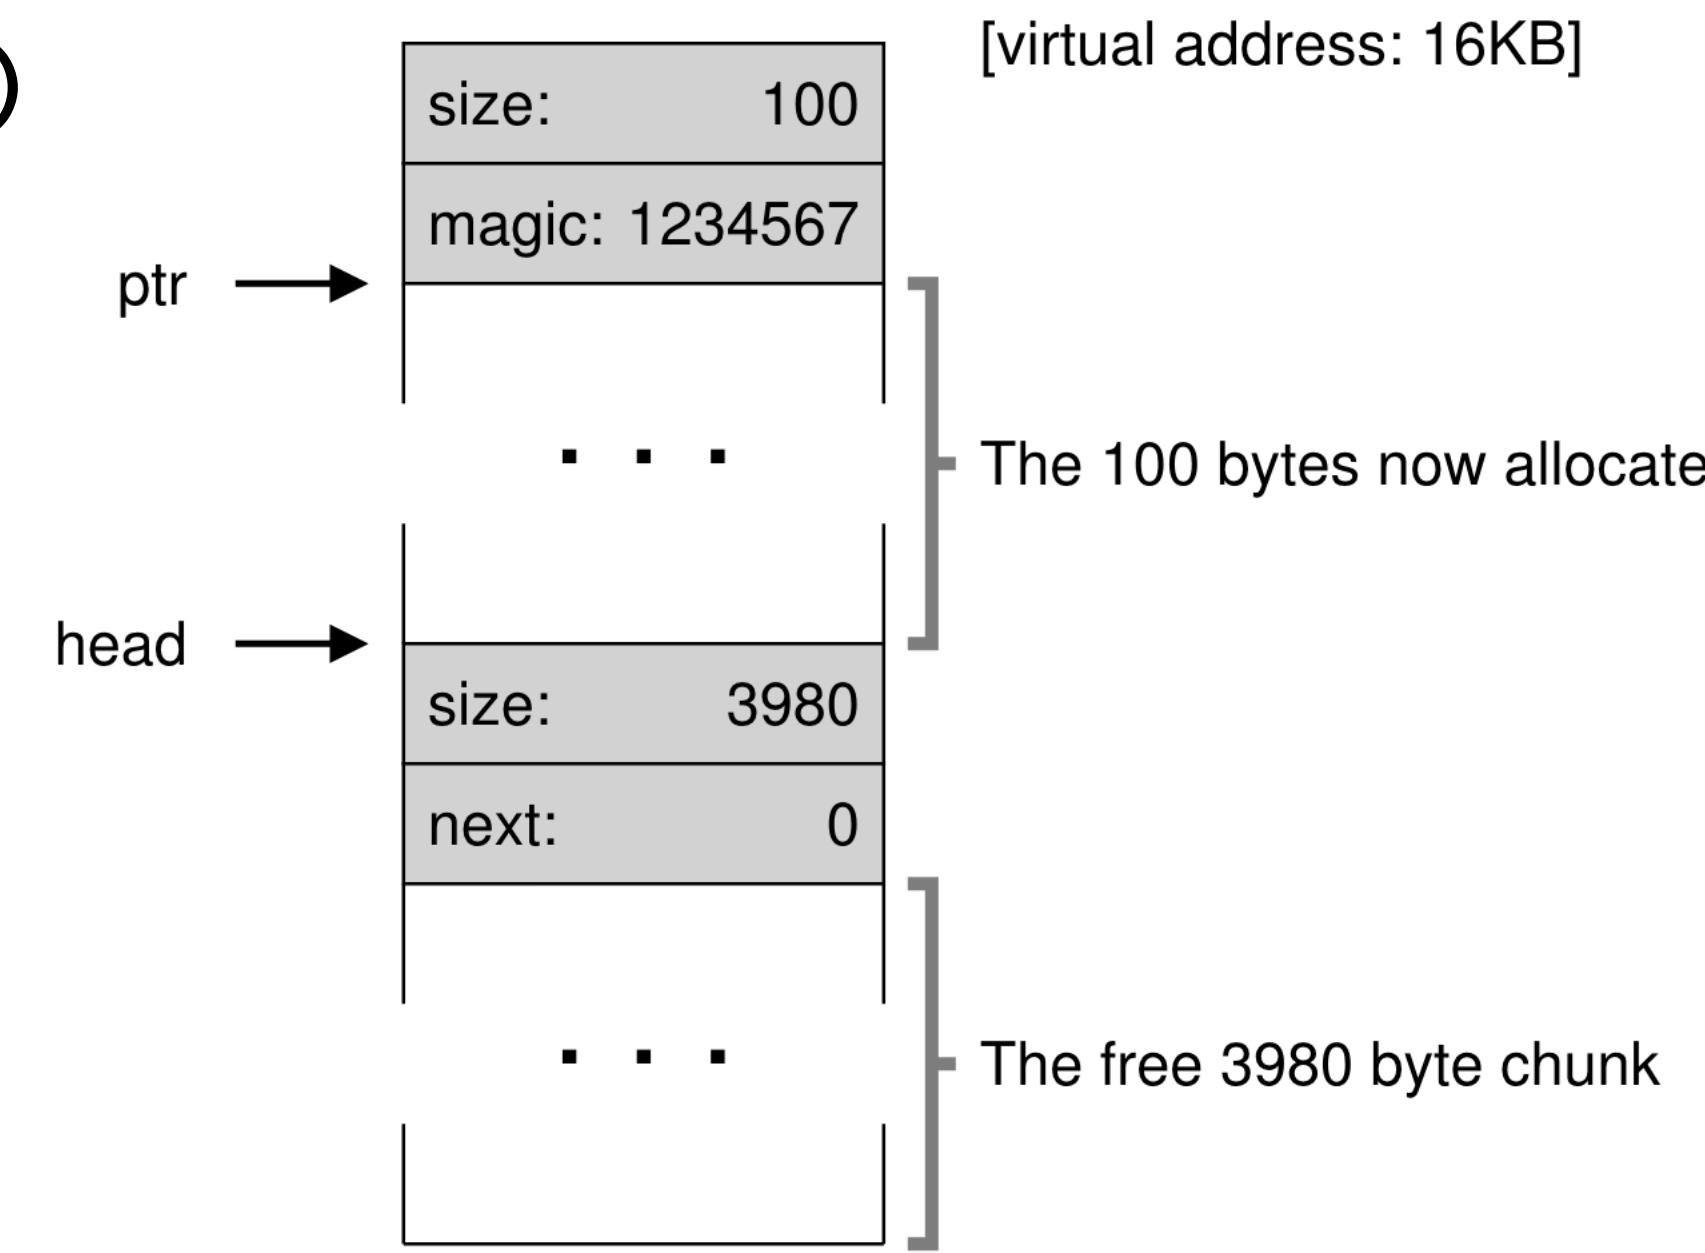


Figure 17.4: A Heap: After One Allocation

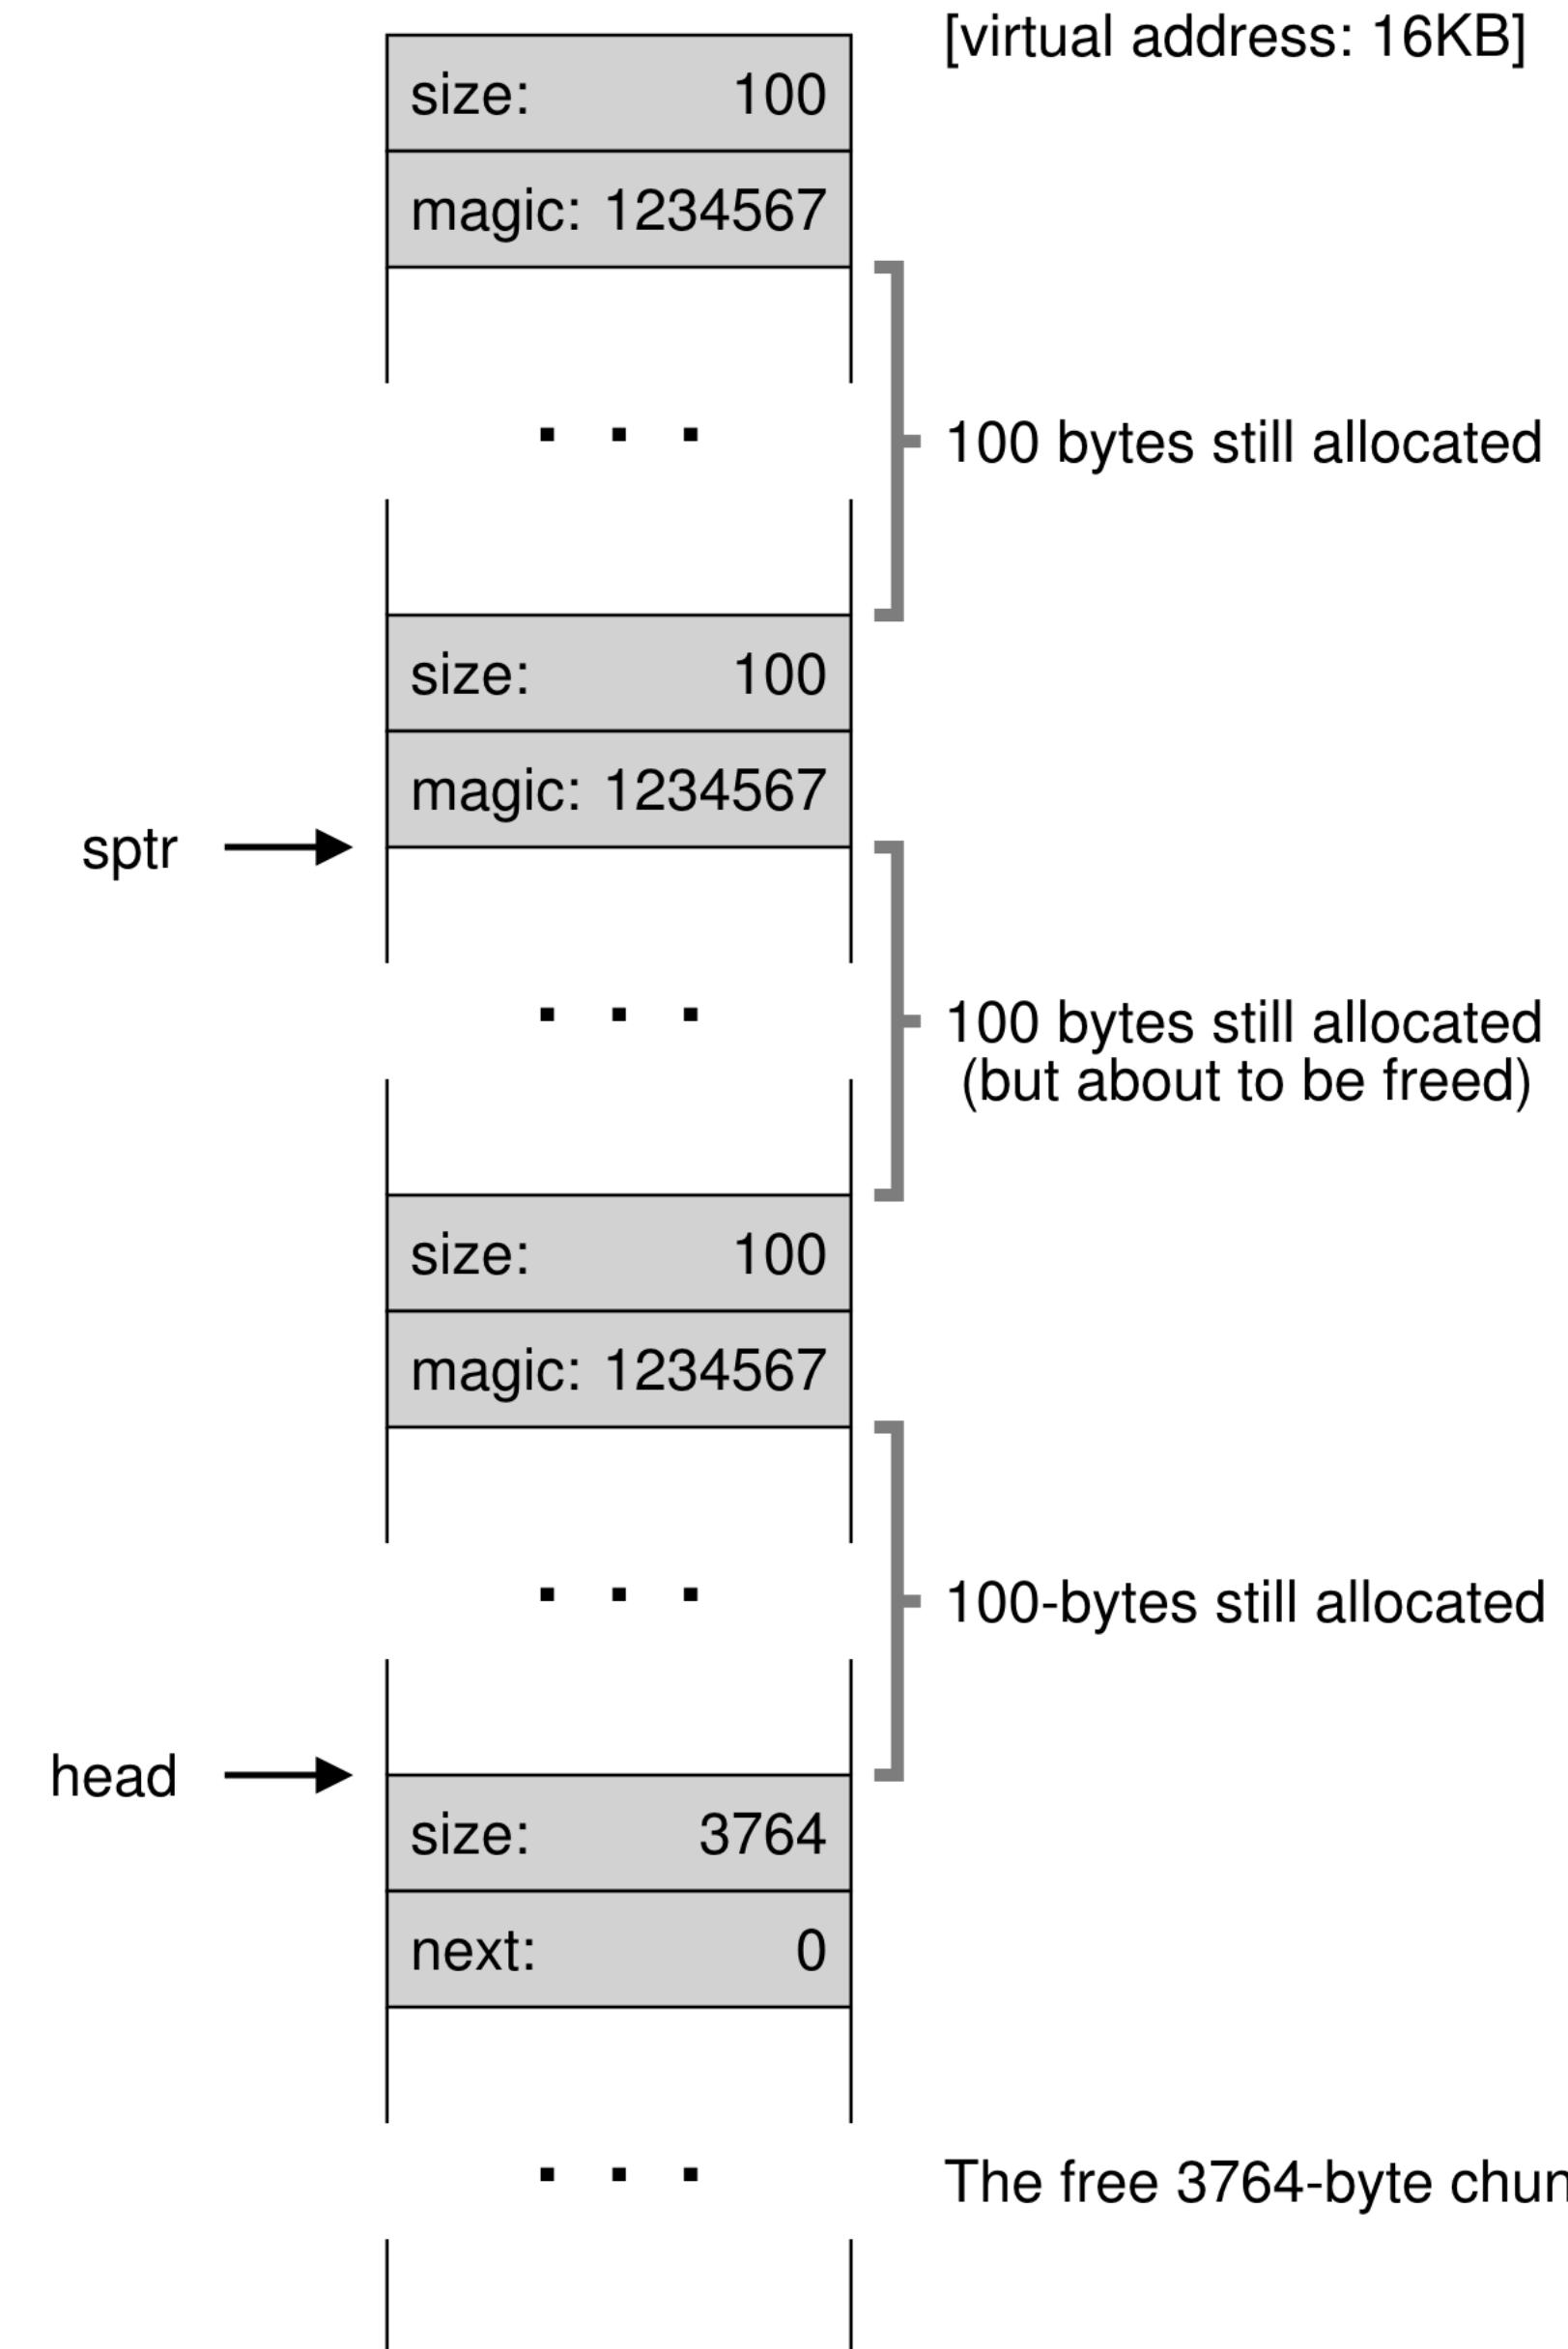


Figure 17.5: Free Space With Three Chunks Allocated

# Free list splitting and coalescing

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

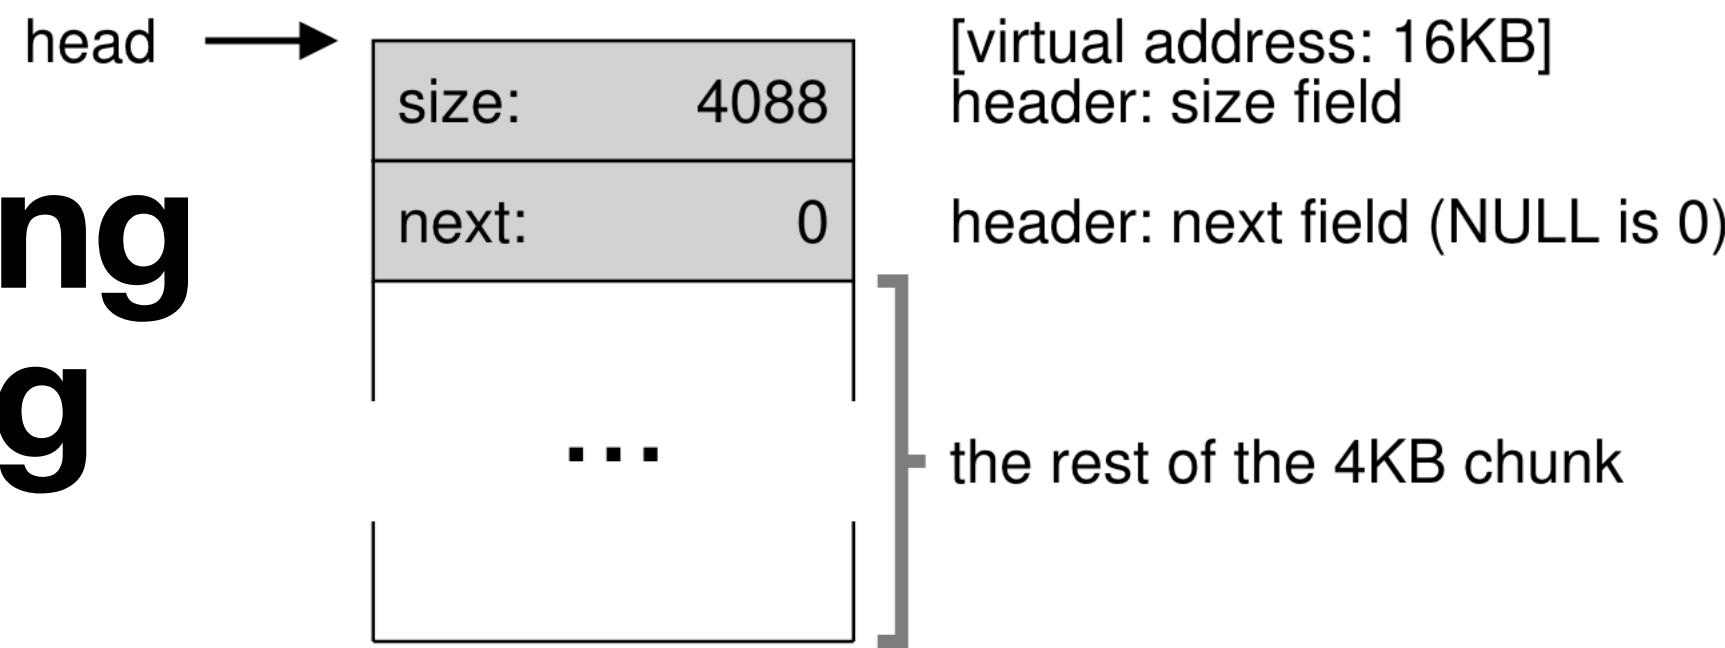


Figure 17.3: A Heap With One Free Chunk

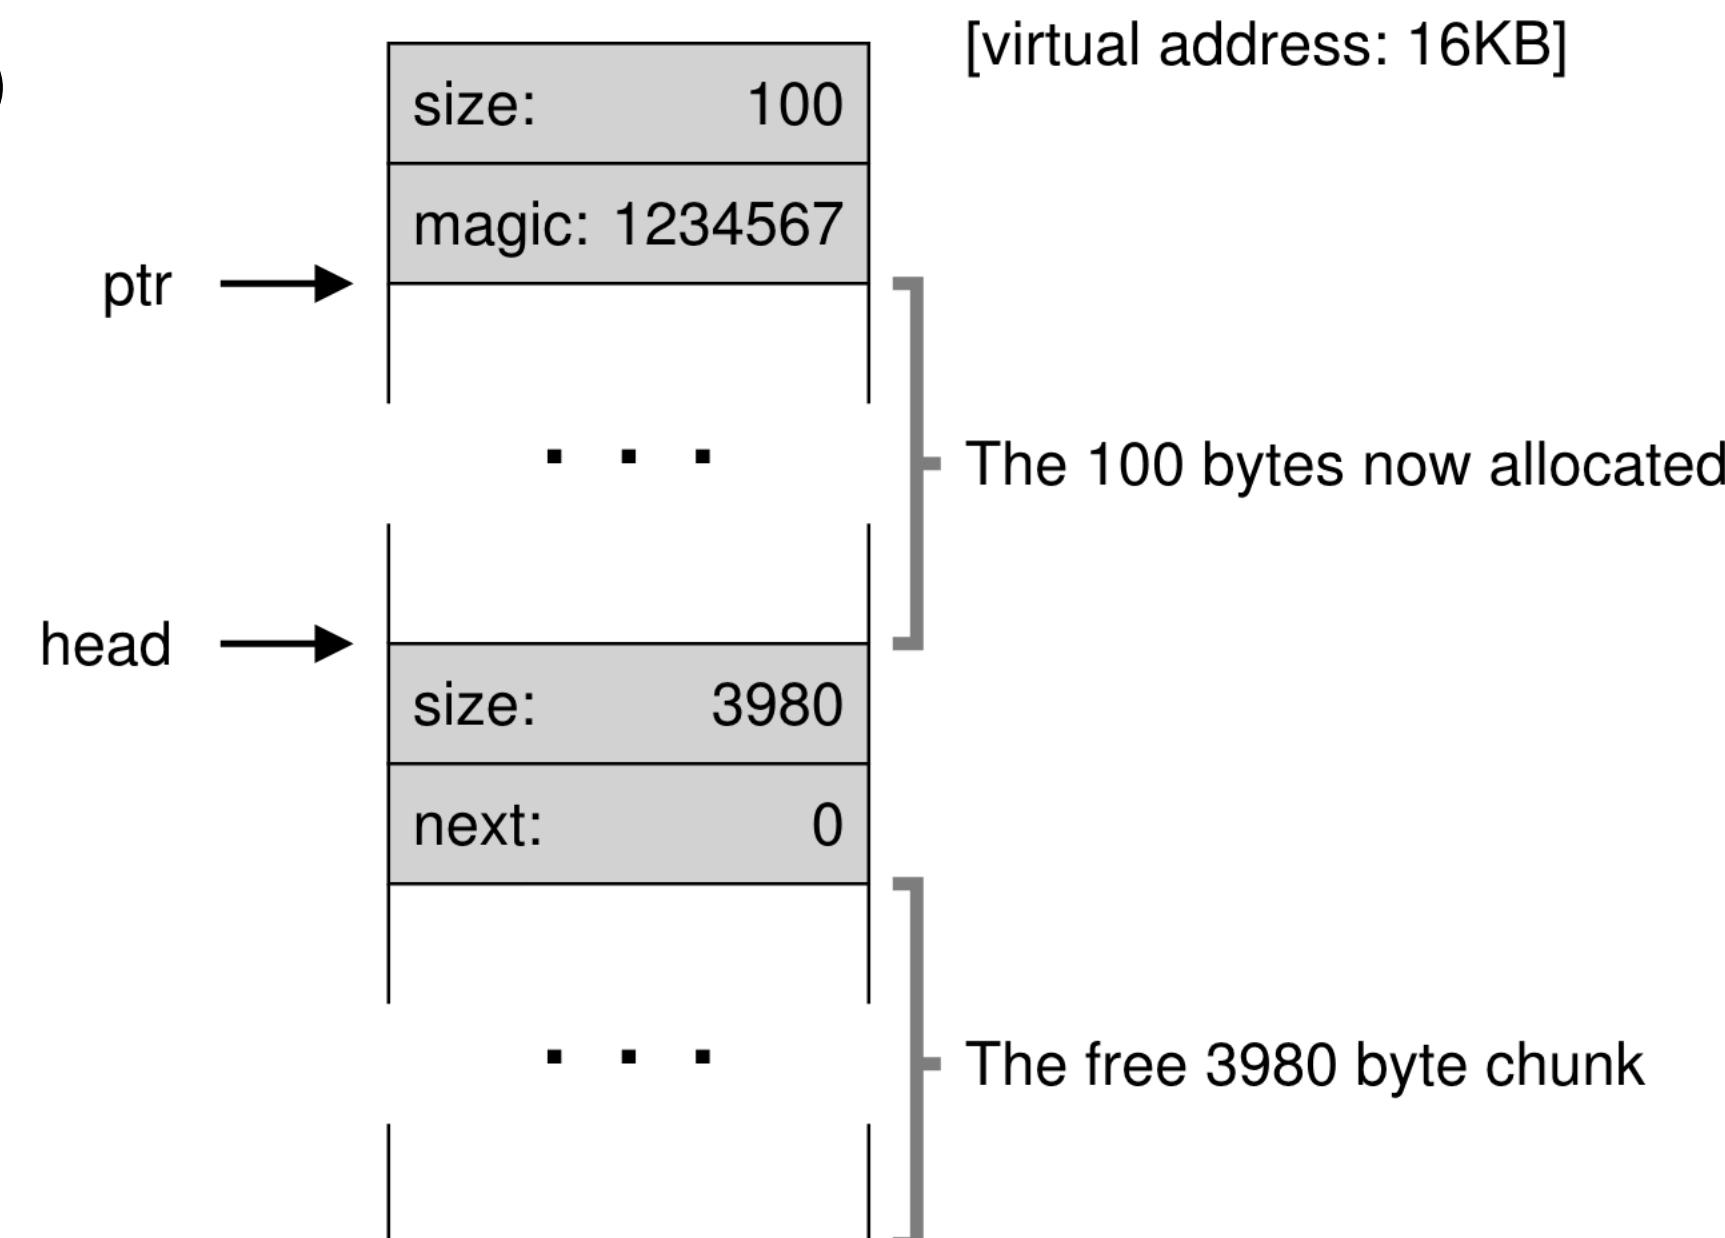


Figure 17.4: A Heap: After One Allocation

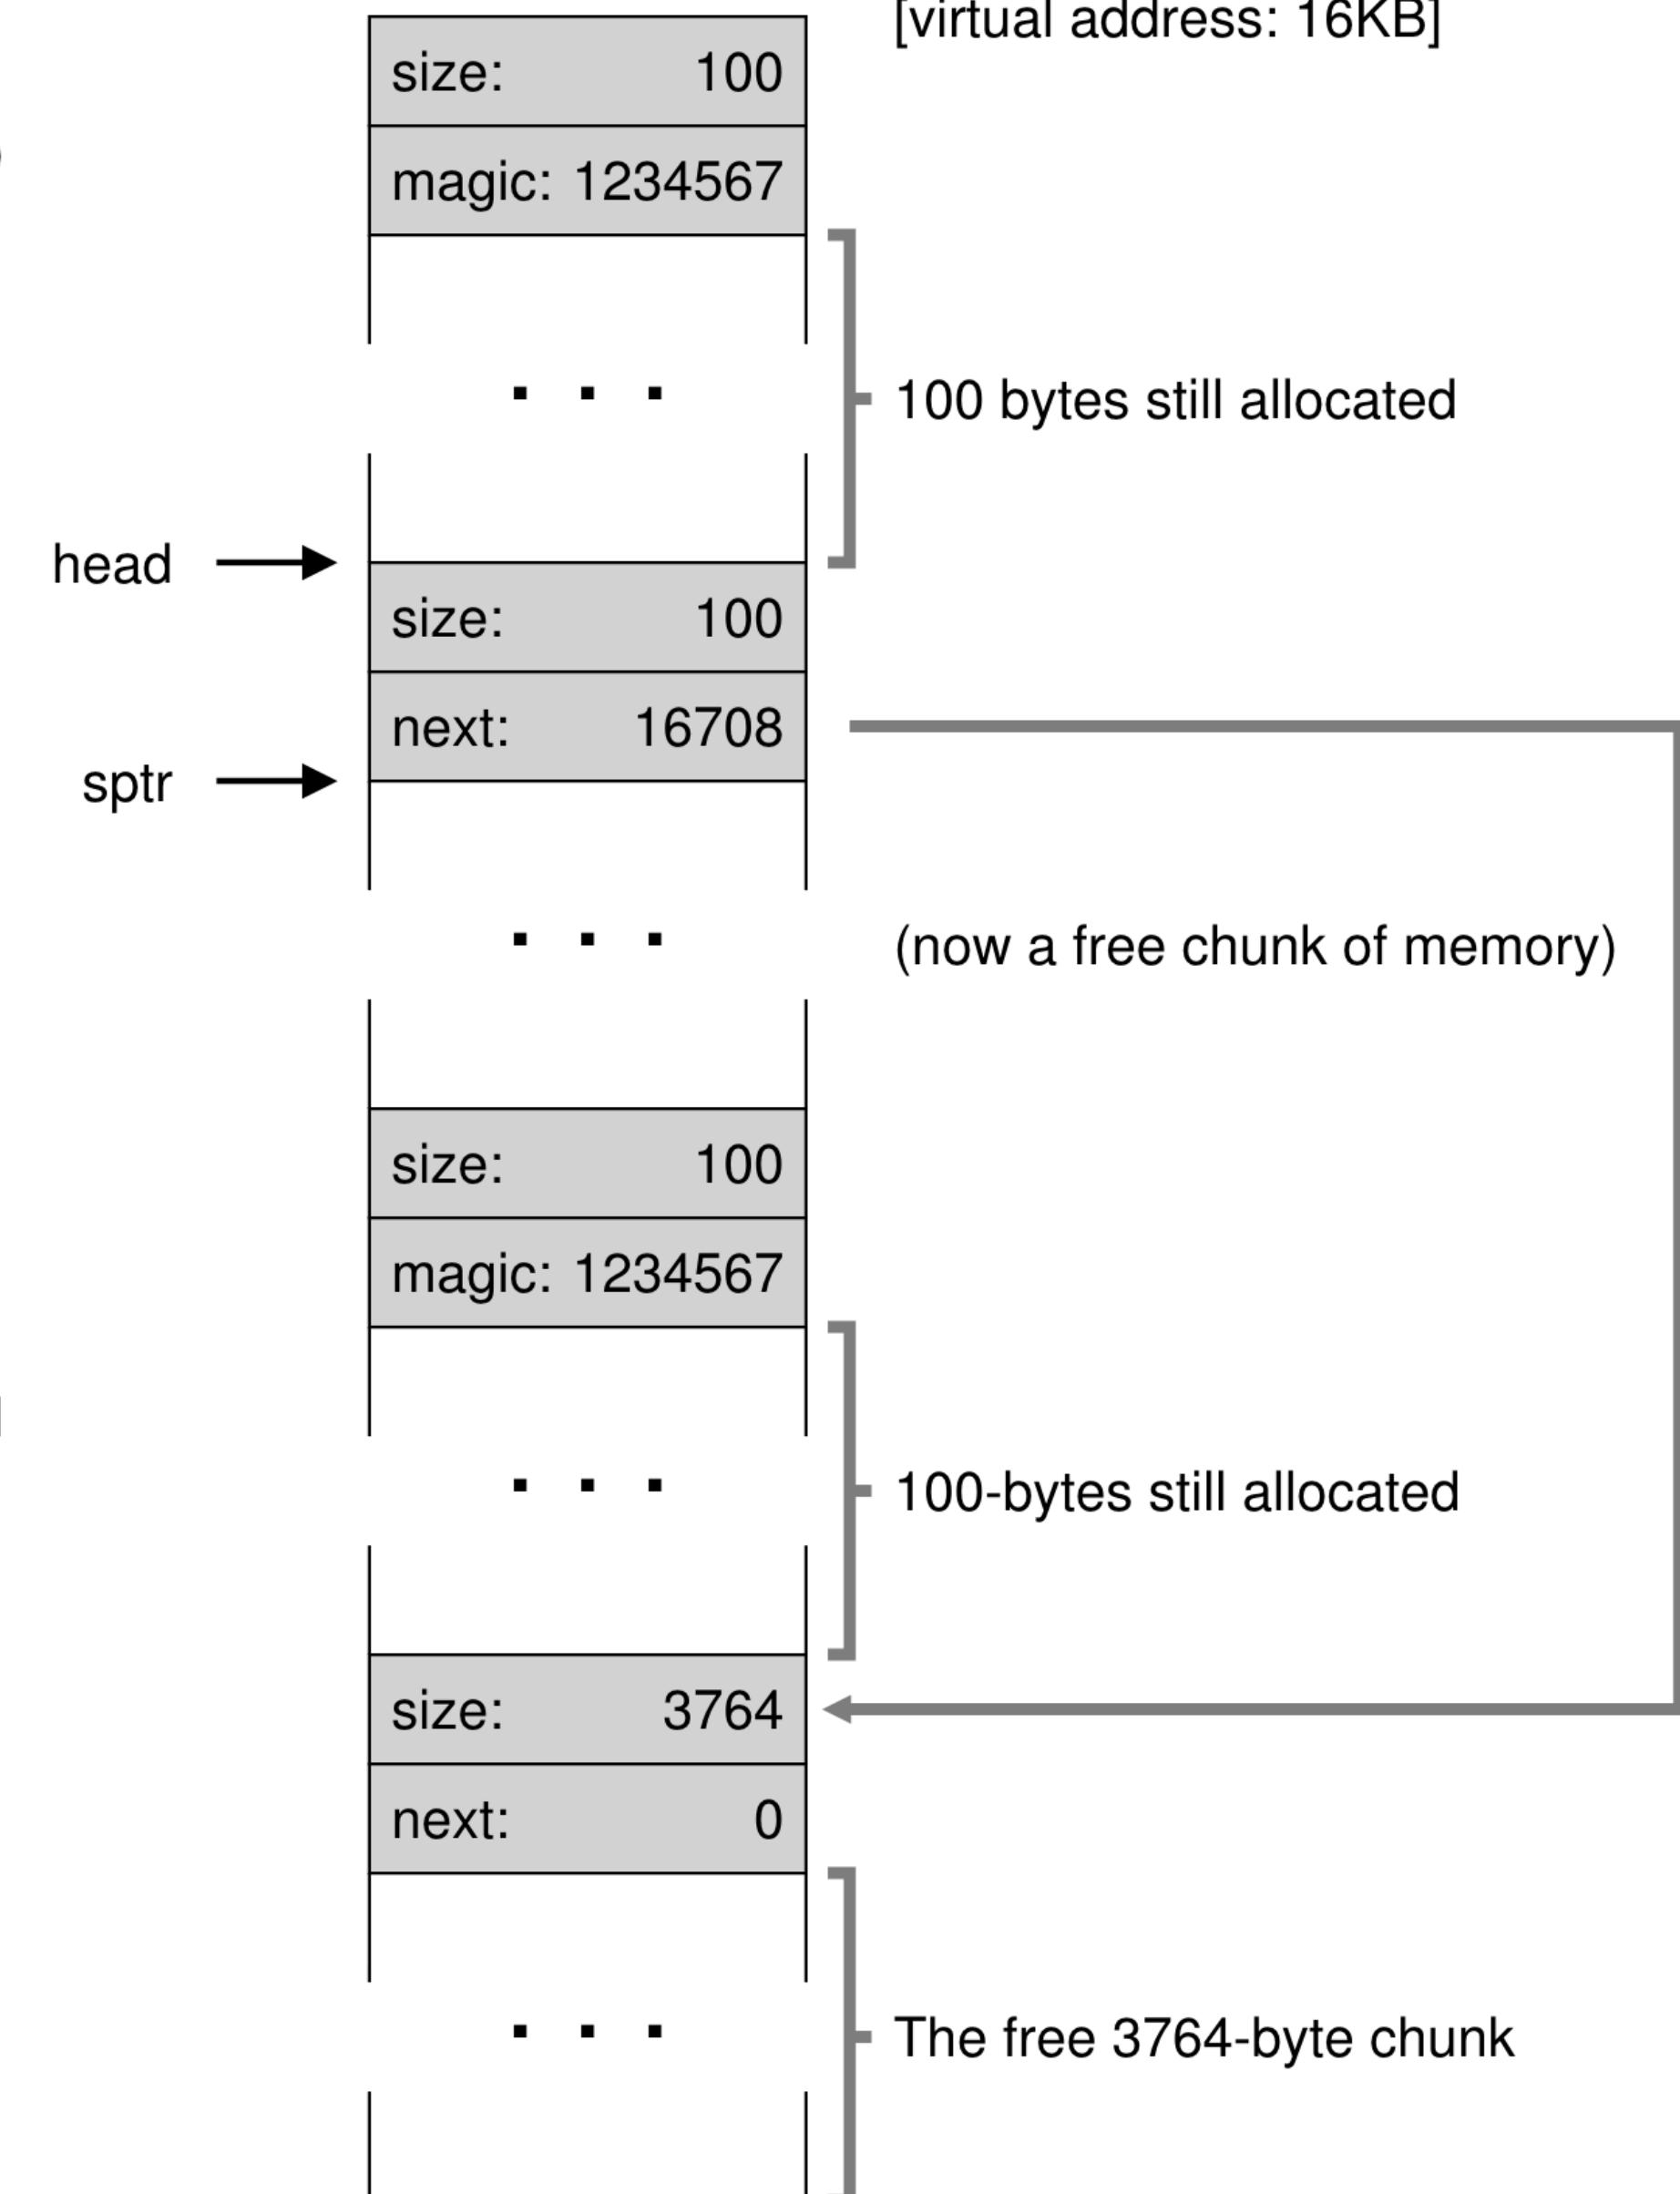


Figure 17.6: Free Space With Two Chunks Allocated

# Free list splitting and coalescing

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

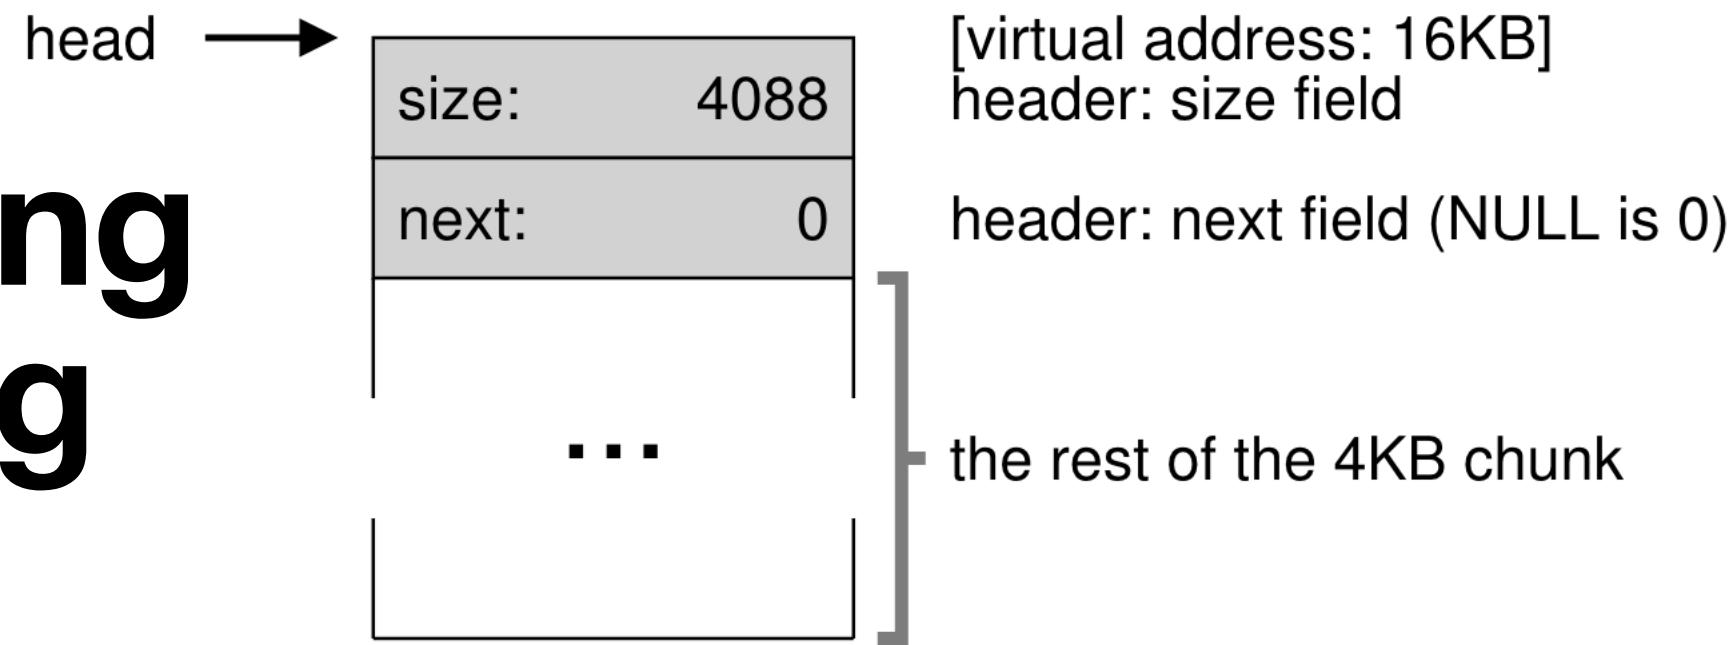


Figure 17.3: A Heap With One Free Chunk

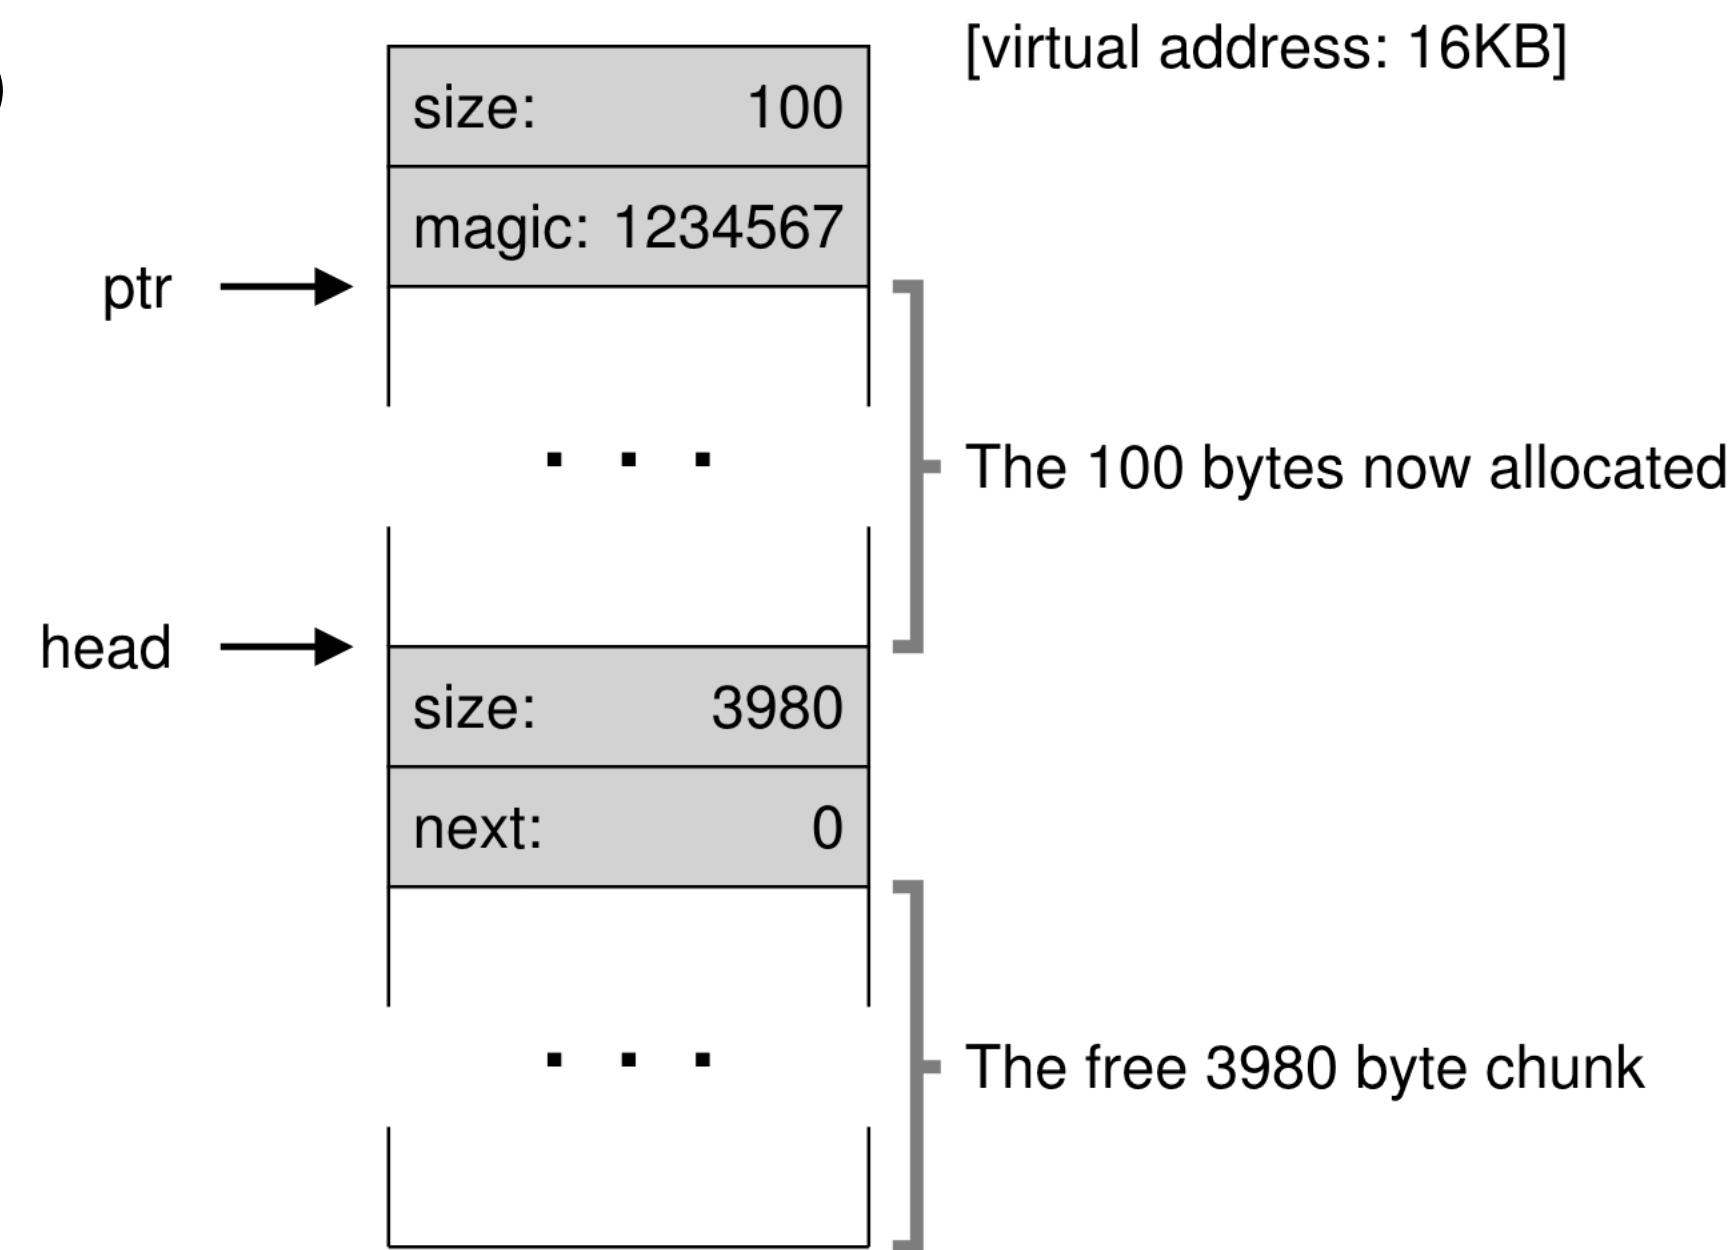


Figure 17.4: A Heap: After One Allocation

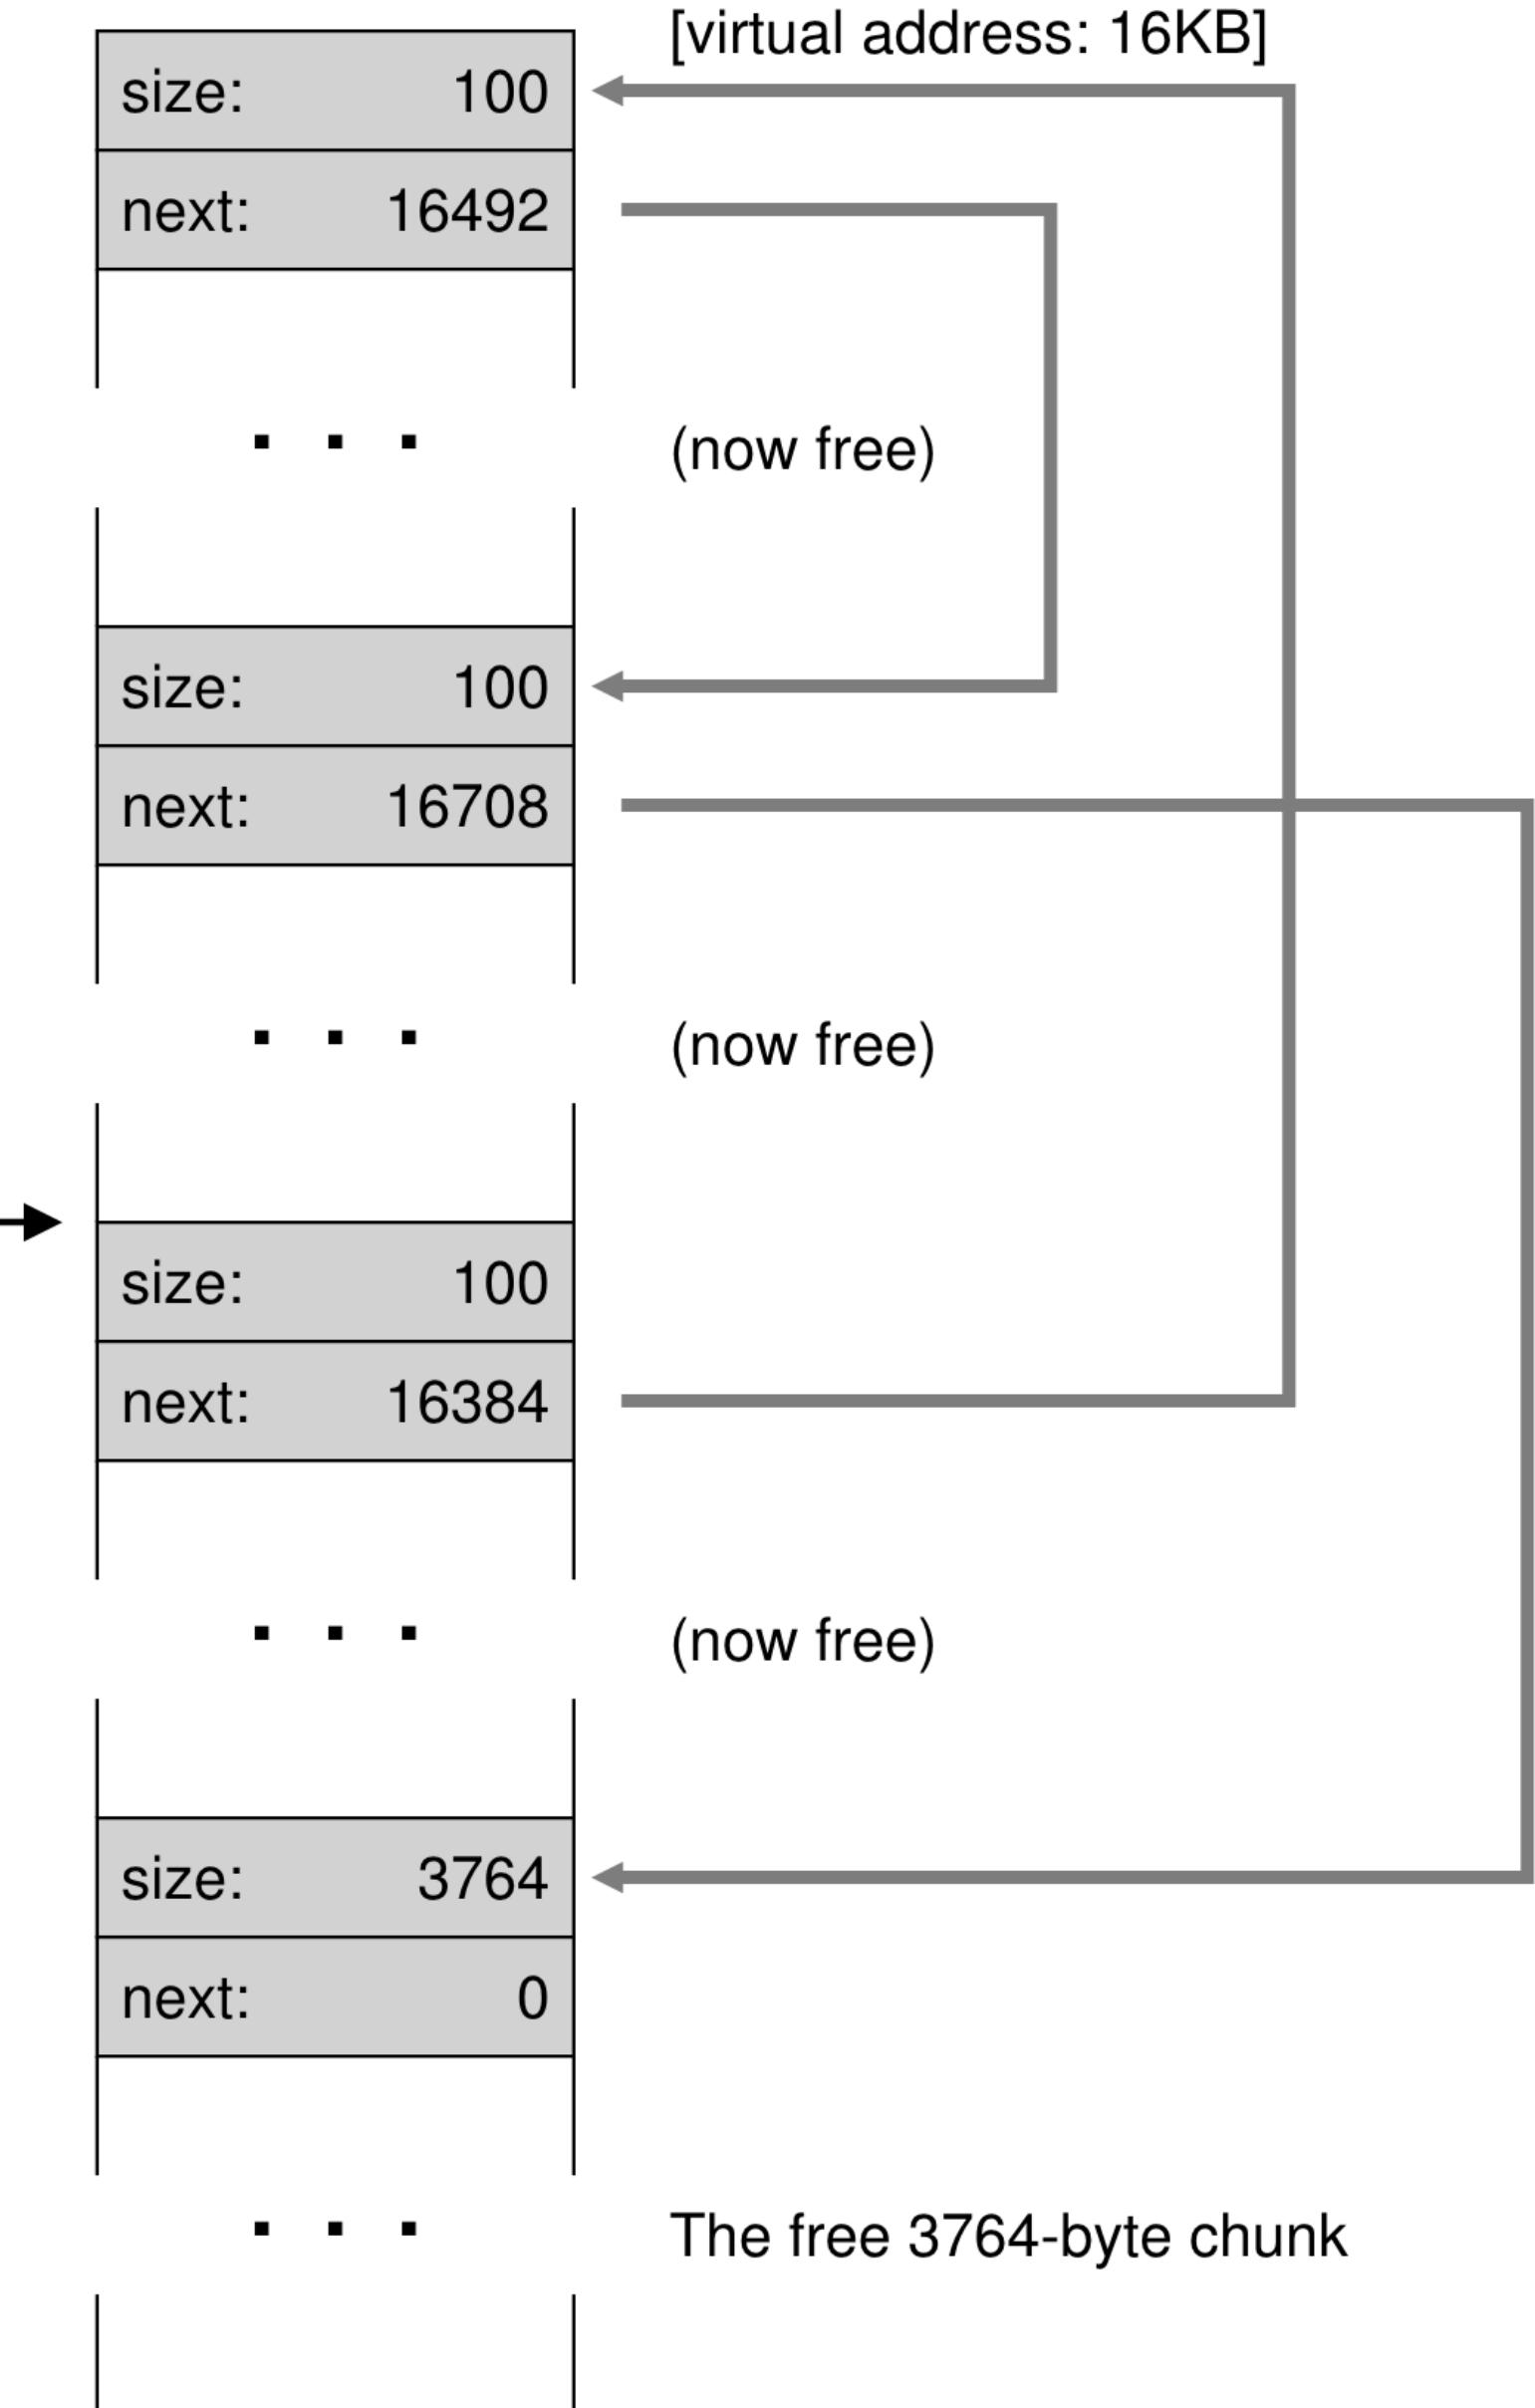


Figure 17.7: A Non-Coalesced Free List

# Free list splitting and coalescing

```
ptr = malloc(100)
```

```
sptr = malloc(100)
```

```
optr = malloc(100)
```

```
free(sptr)
```

```
free(ptr)
```

```
free(optr)
```

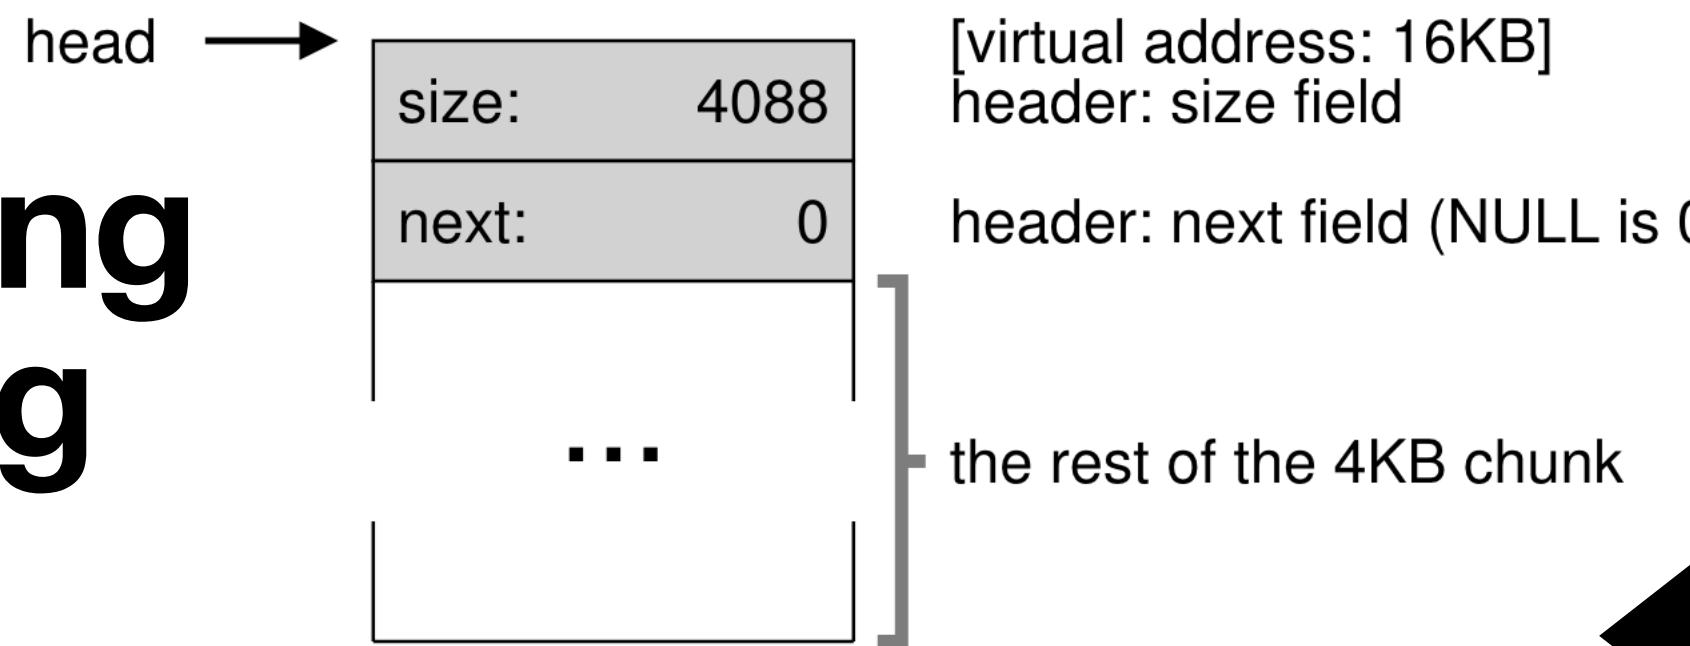


Figure 17.3: A Heap With One Free Chunk

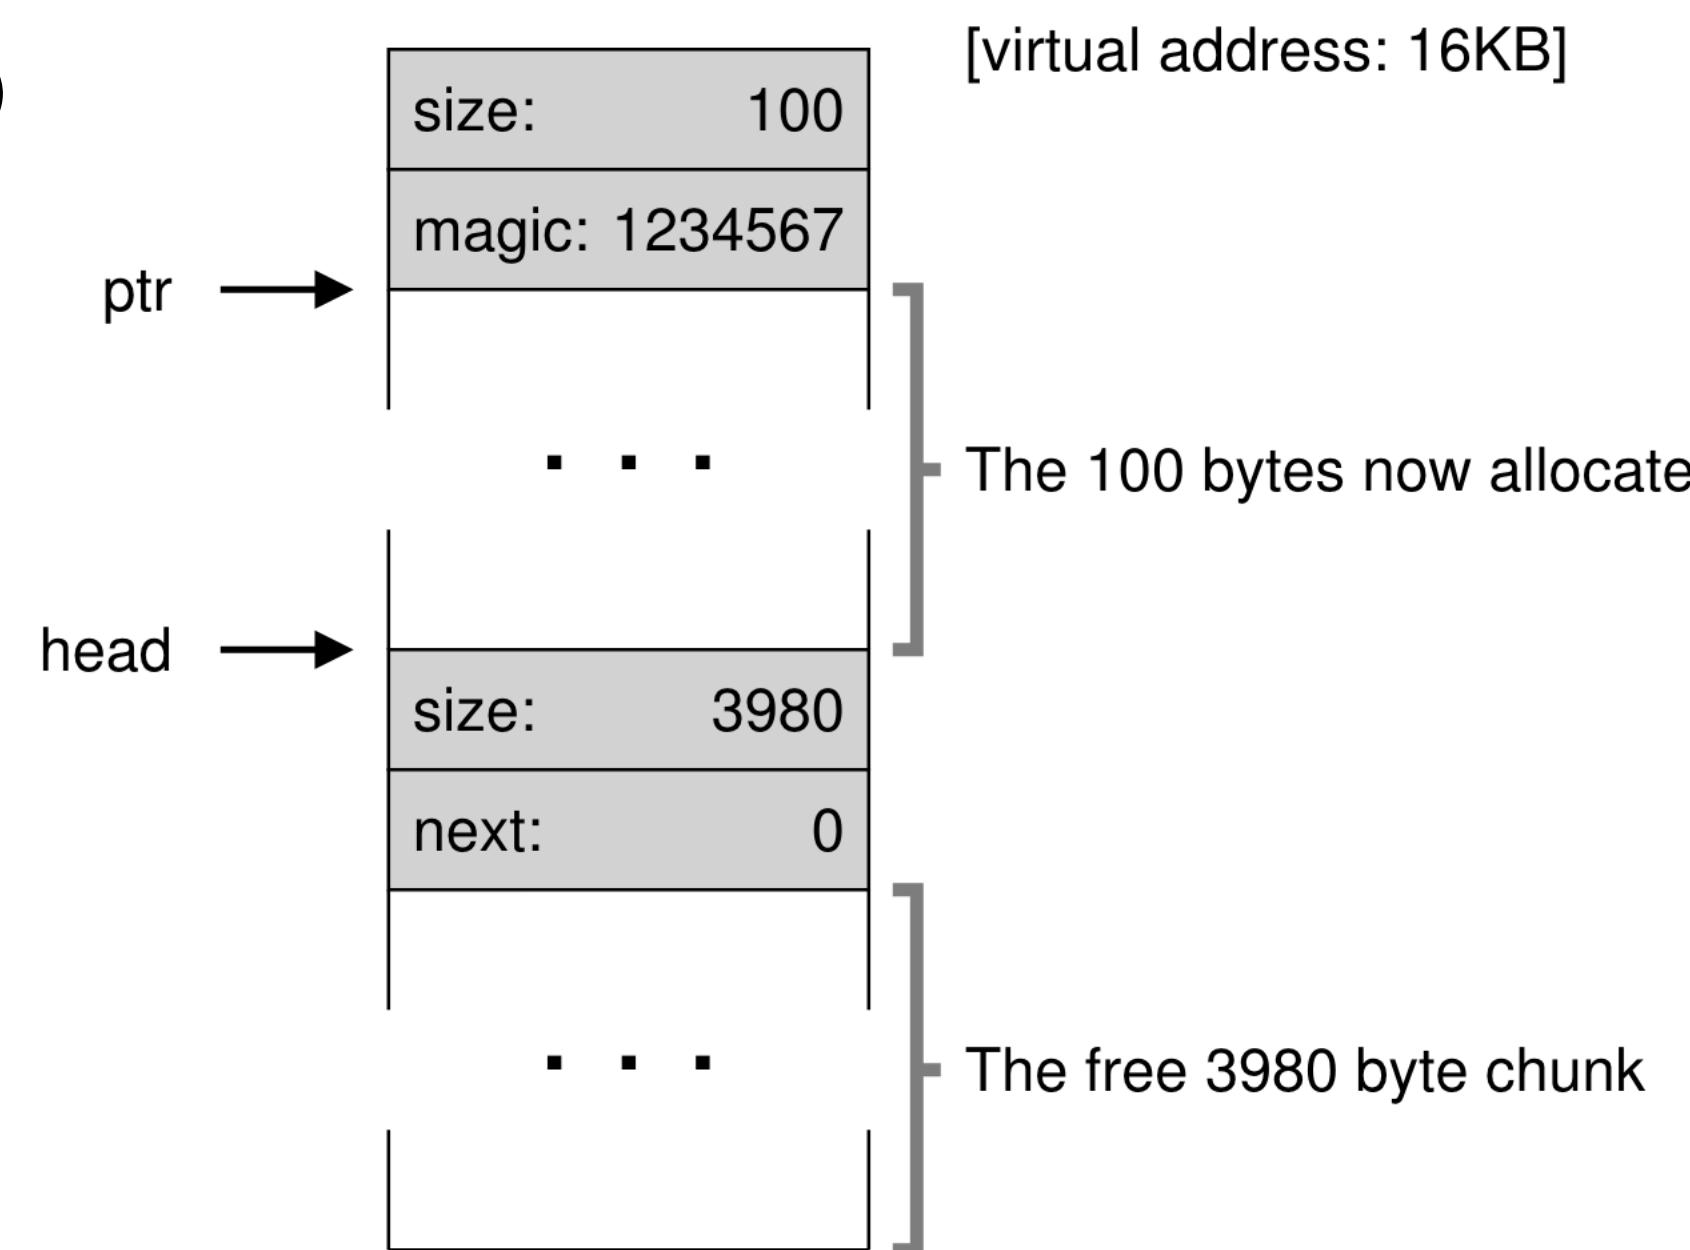


Figure 17.4: A Heap: After One Allocation

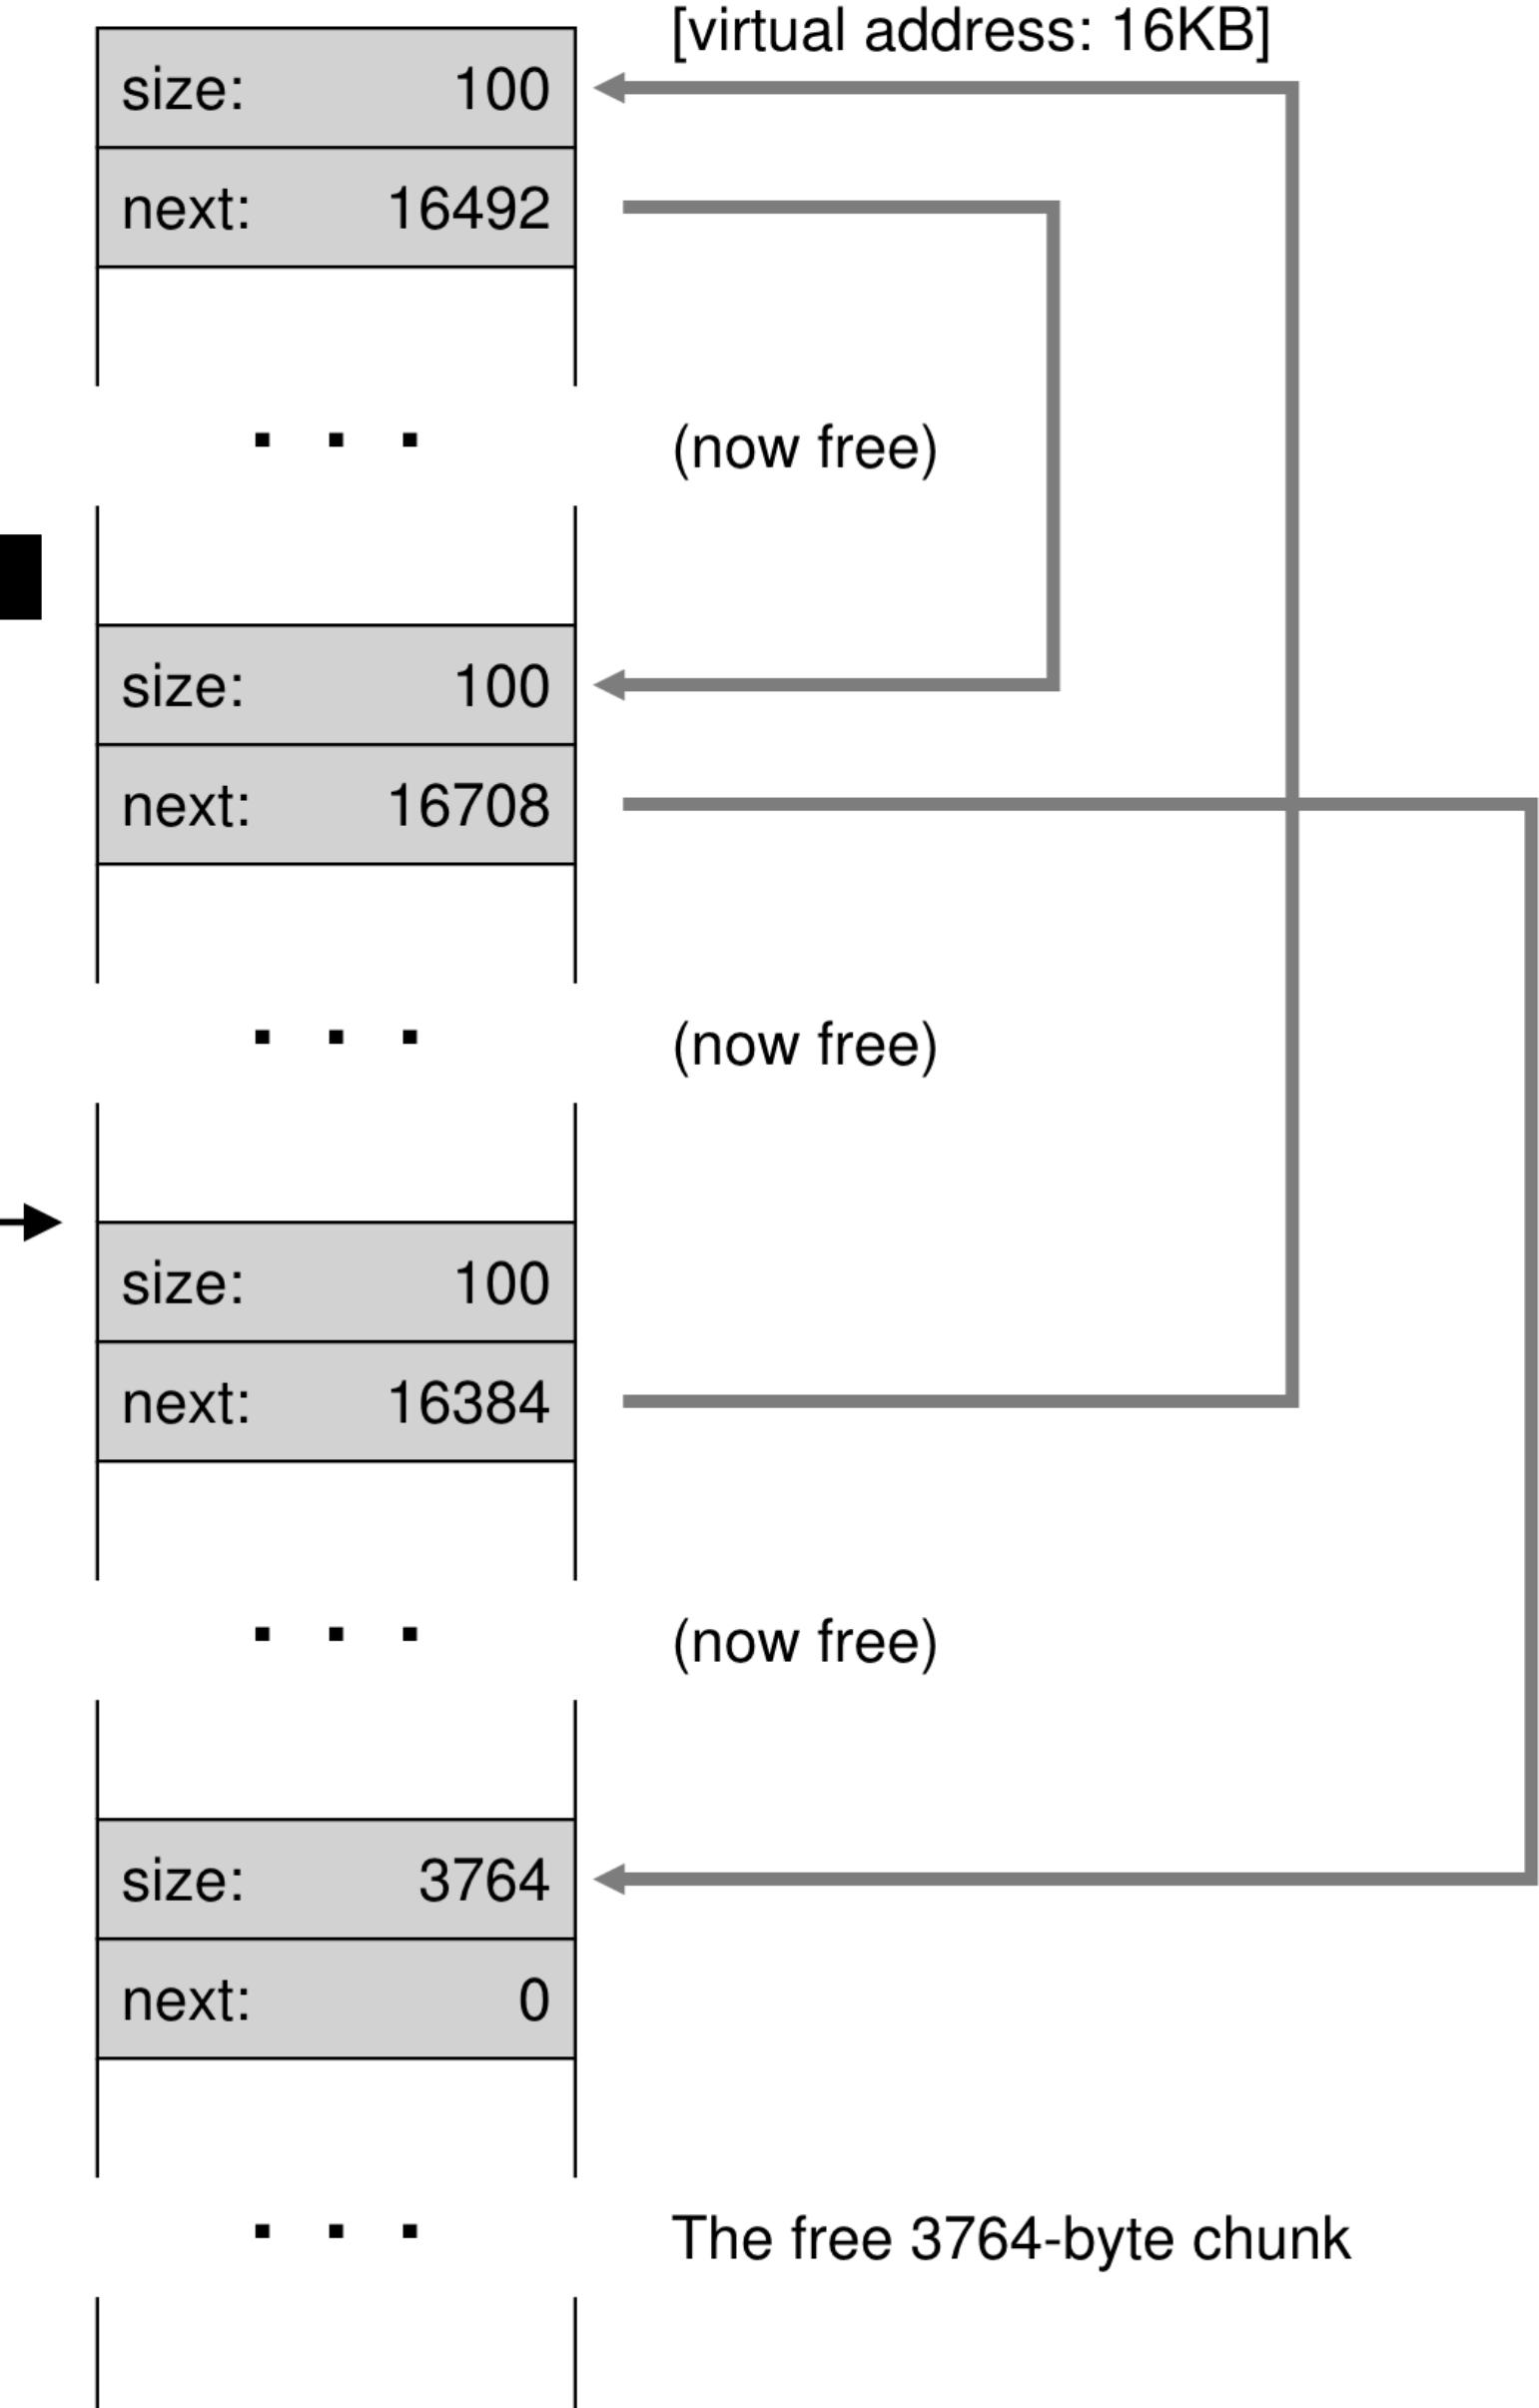
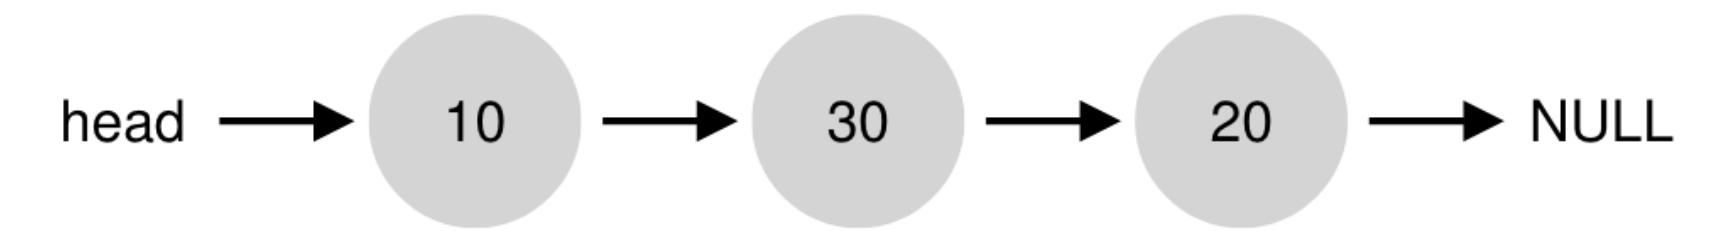


Figure 17.7: A Non-Coalesced Free List

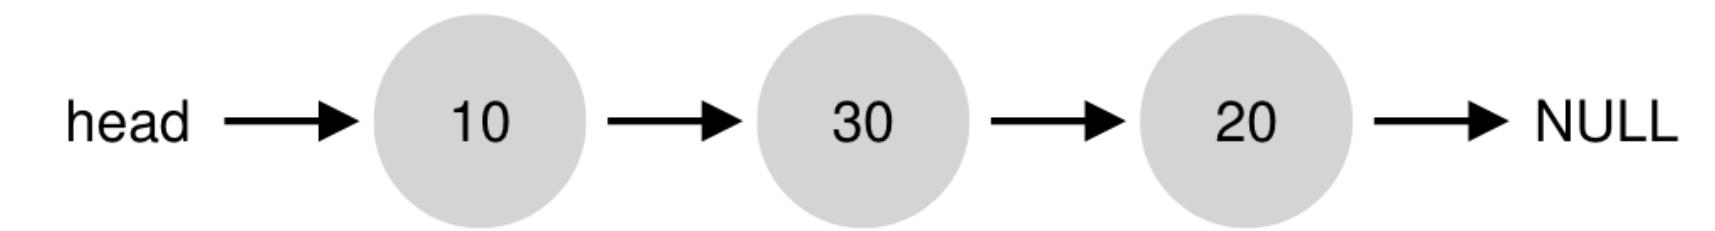
# Which block to allocate?

Example: malloc(15)



# Which block to allocate?

Example: malloc(15)



- Best fit
  - Slow. need to search the whole list

# Which block to allocate?

Example: malloc(15)

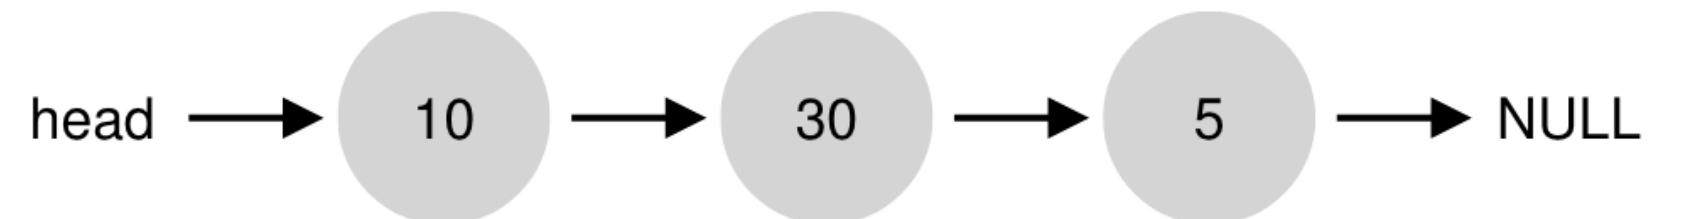
- Best fit
  - Slow. need to search the whole list



# Which block to allocate?

Example: malloc(15)

- Best fit
  - Slow. need to search the whole list
- First fit
  - Faster. (xv6: umalloc.c)



# Which block to allocate?

Example: malloc(15)

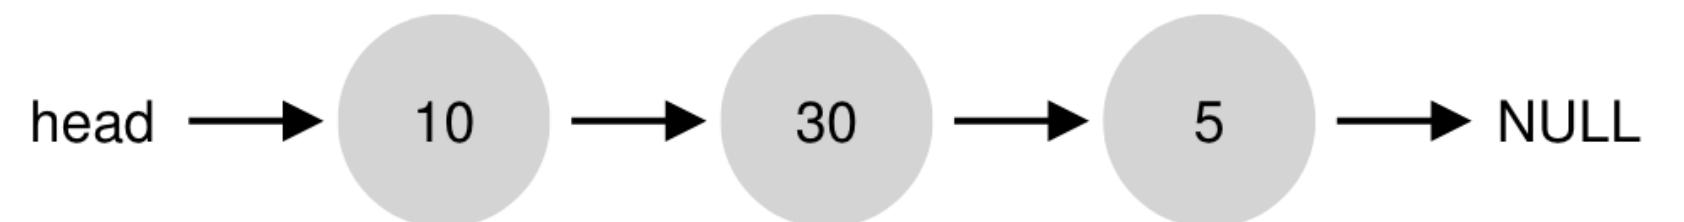
- Best fit
  - Slow. need to search the whole list
- First fit
  - Faster. (xv6: umalloc.c)



# Which block to allocate?

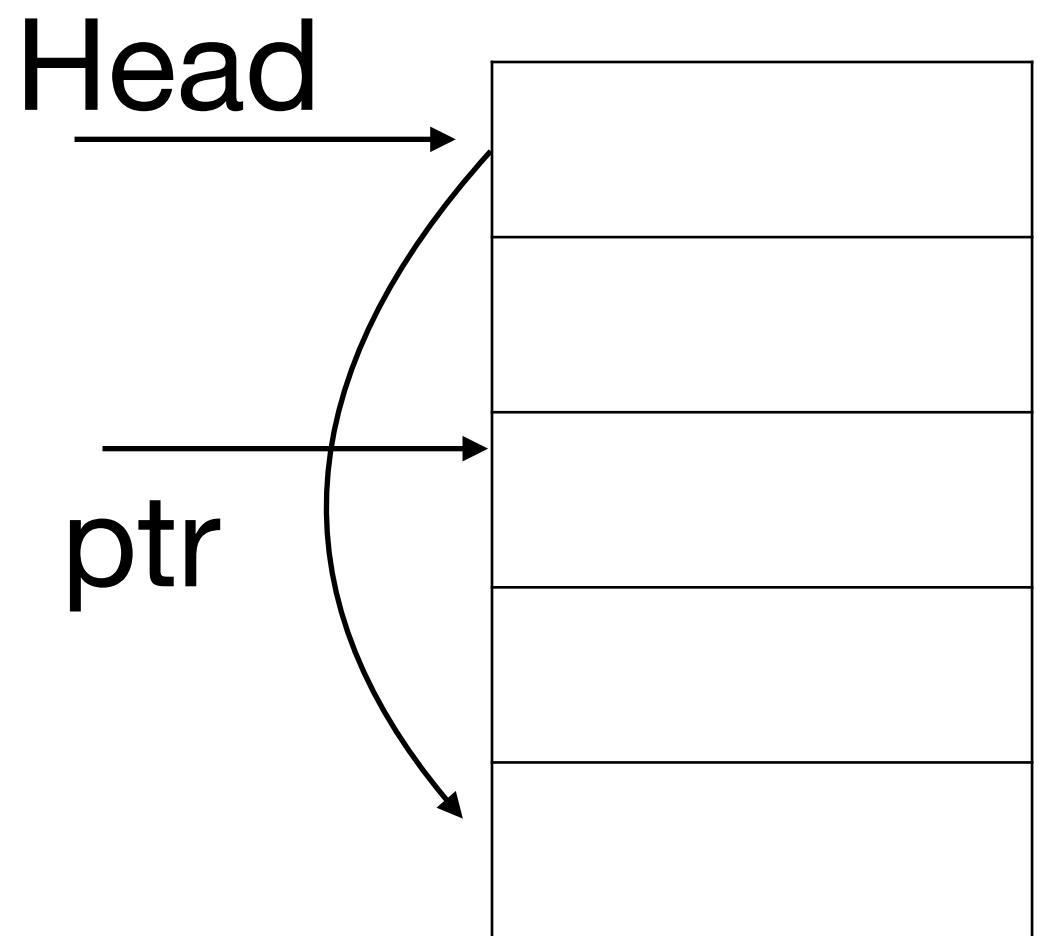
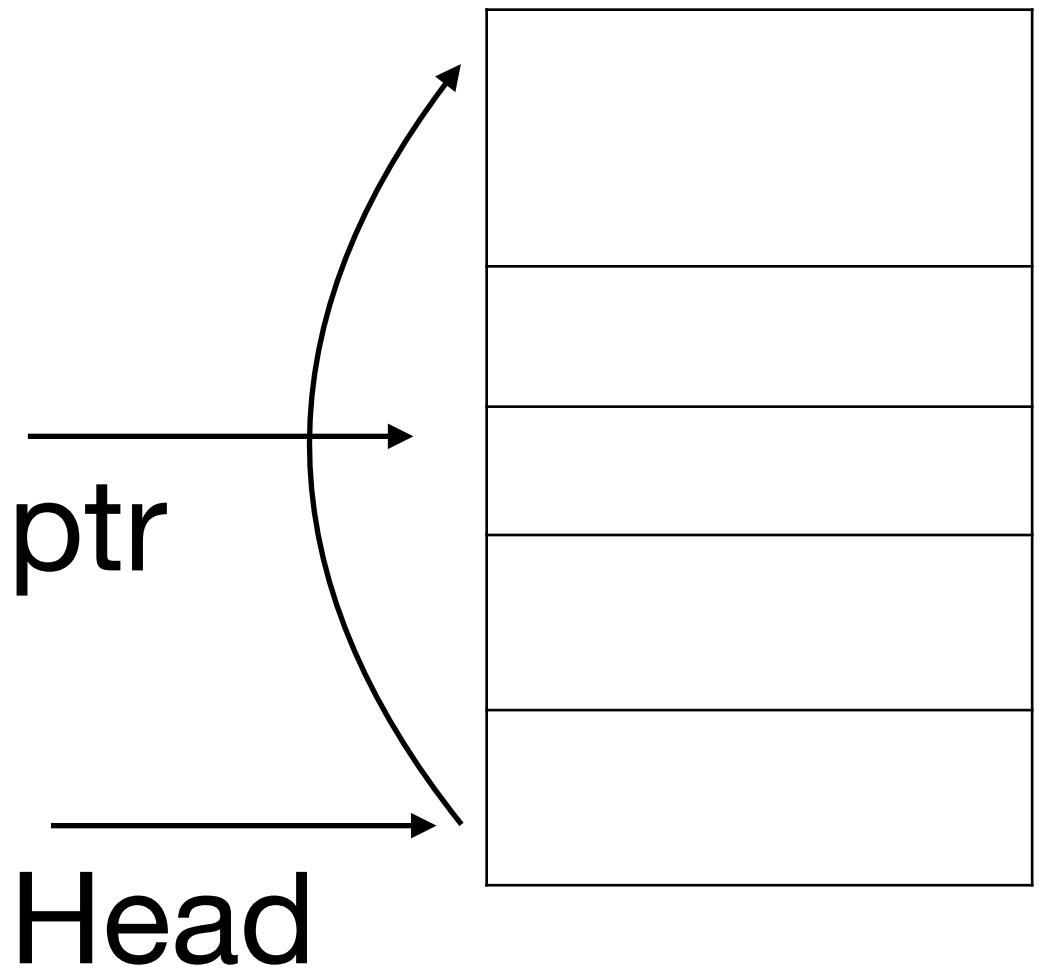
Example: malloc(15)

- Best fit
  - Slow. need to search the whole list
- First fit
  - Faster. (xv6: umalloc.c)
- Fragmentation
  - Example: malloc(25)



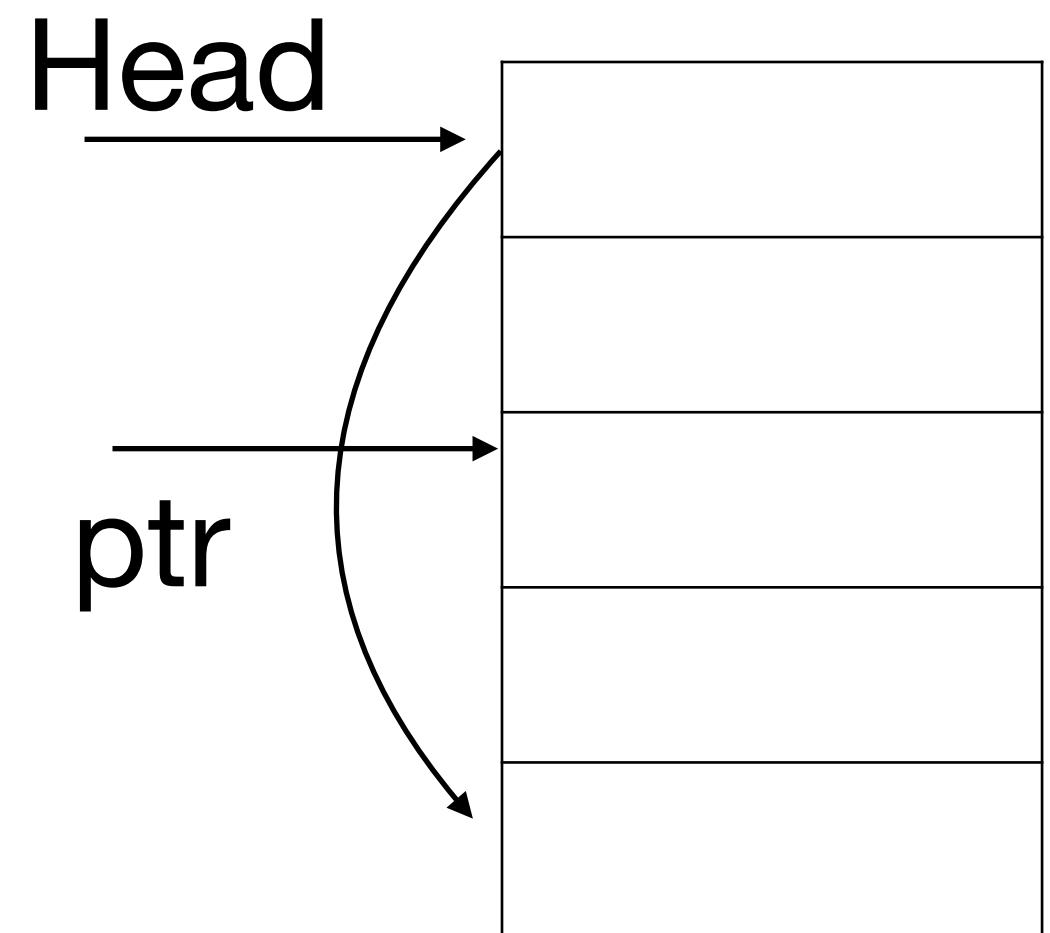
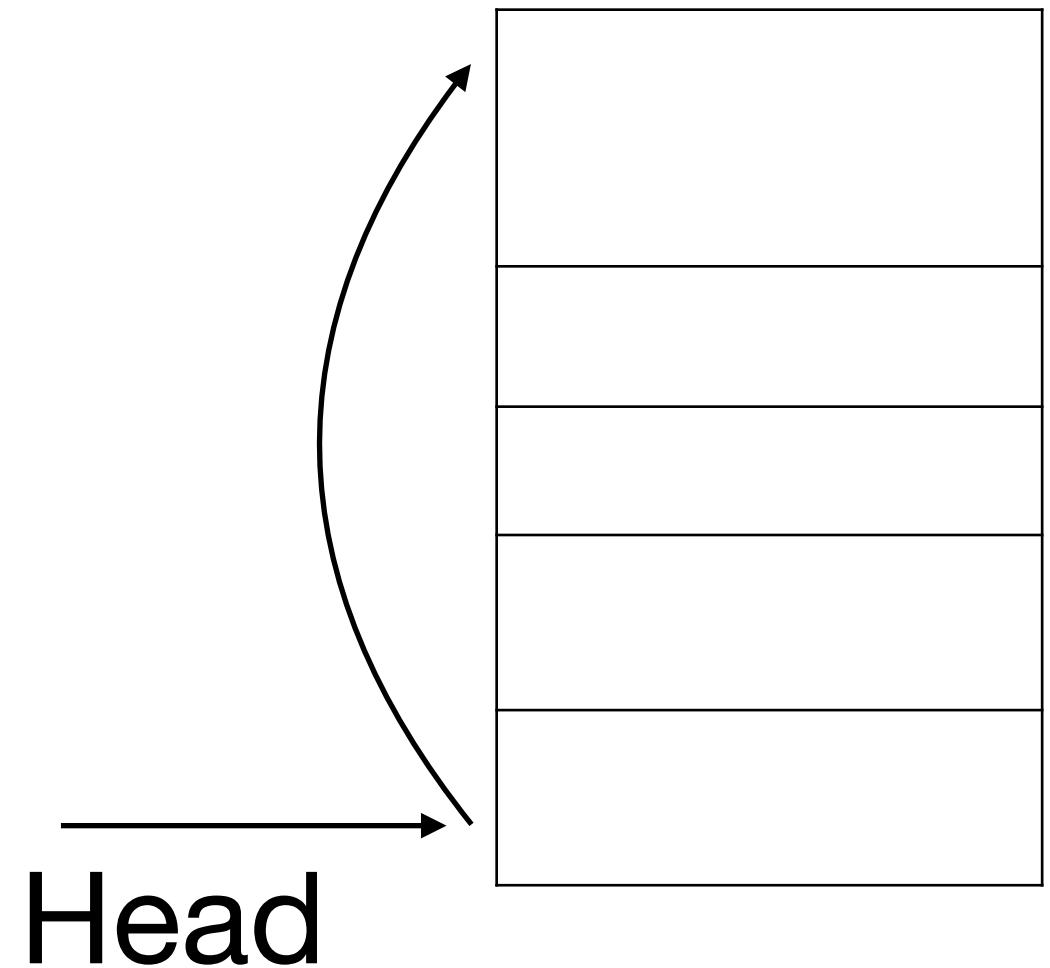
# In which order to maintain lists?

- (De)allocation order



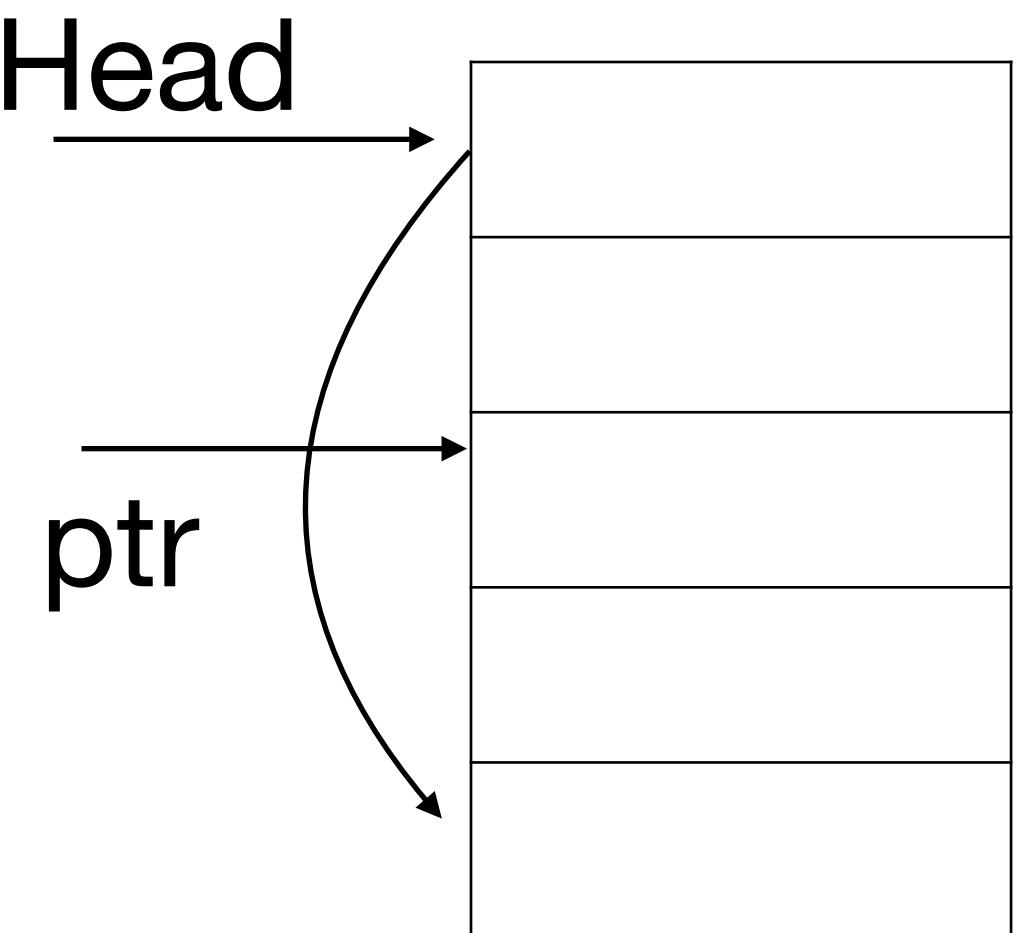
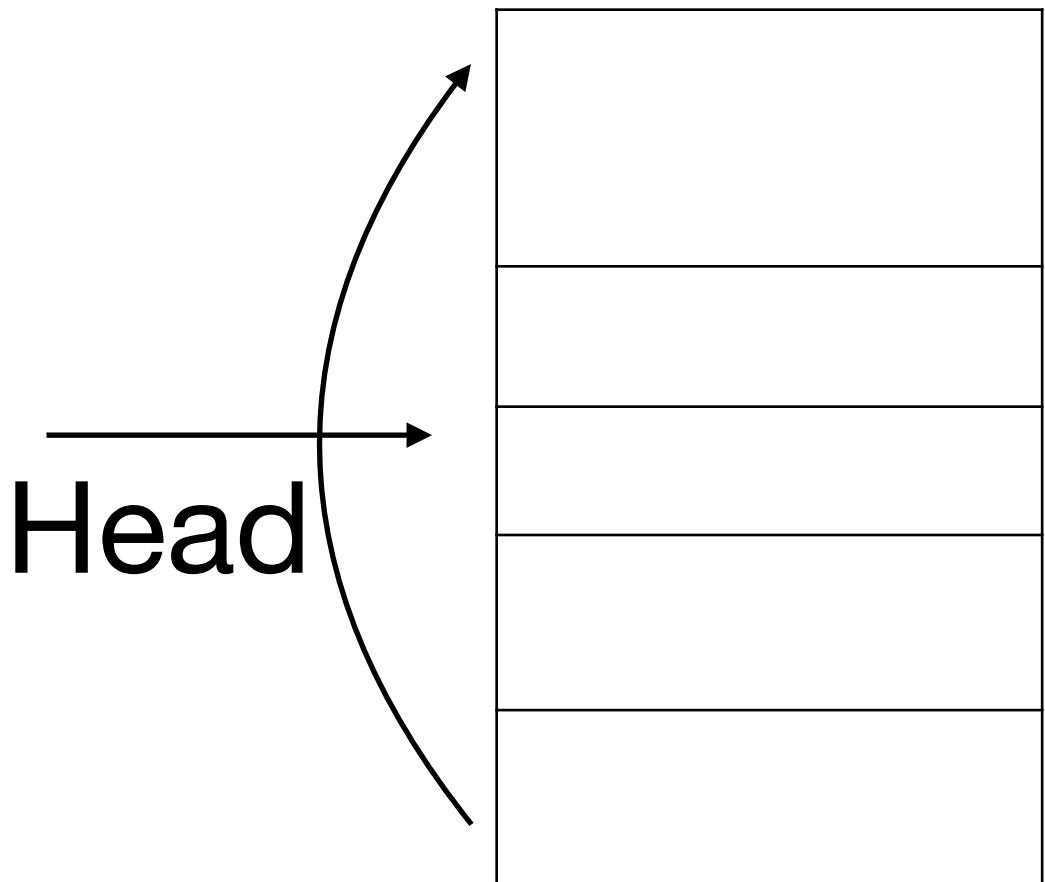
# In which order to maintain lists?

- (De)allocation order



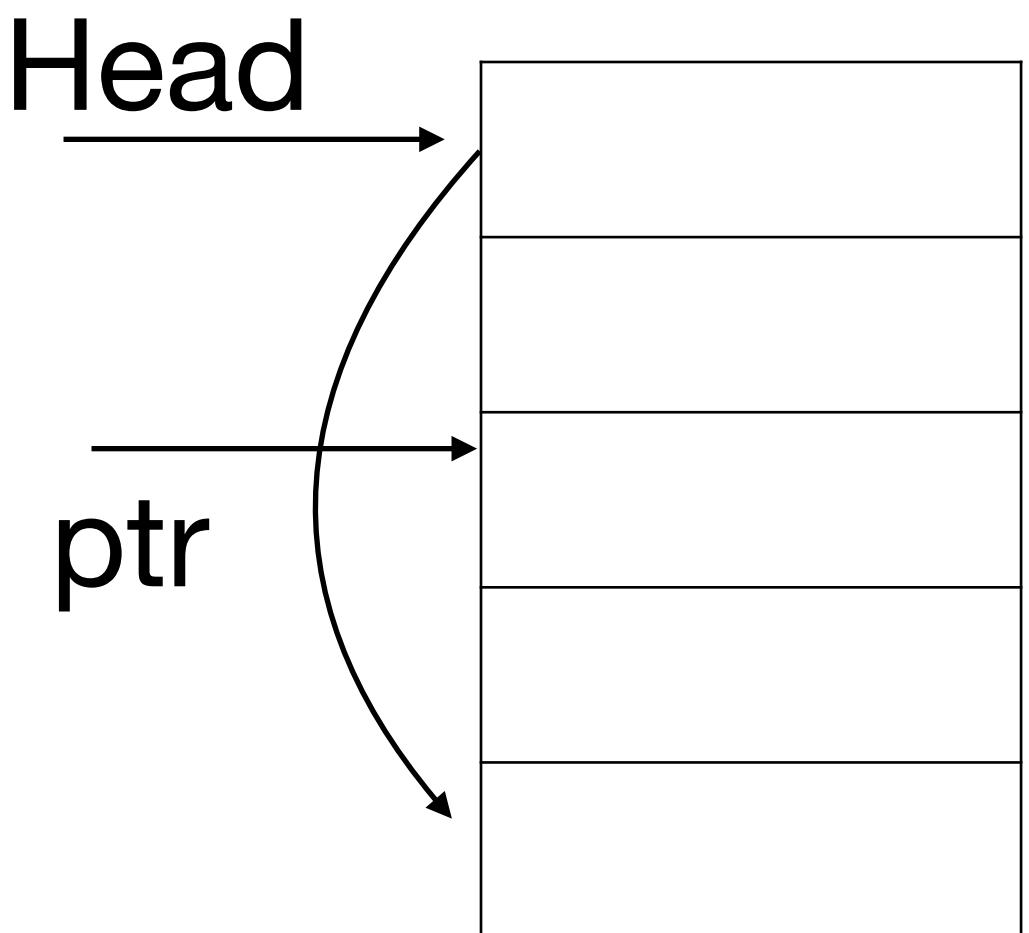
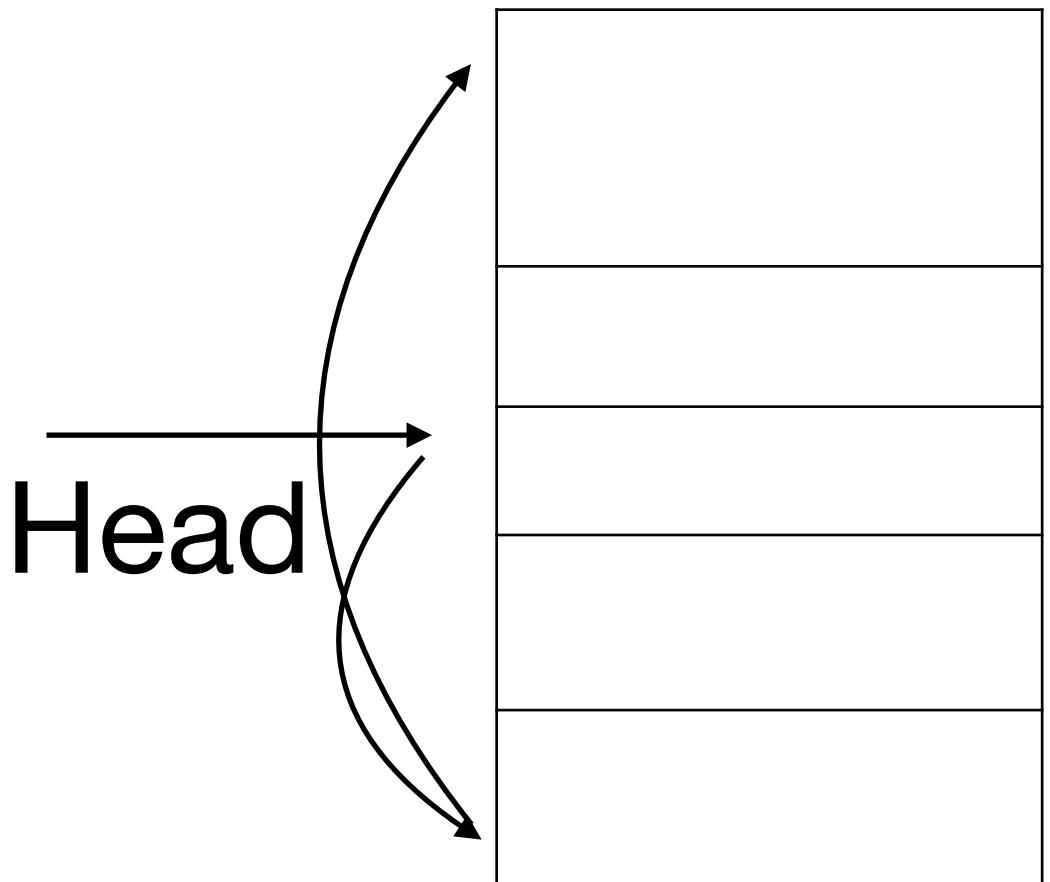
# In which order to maintain lists?

- (De)allocation order



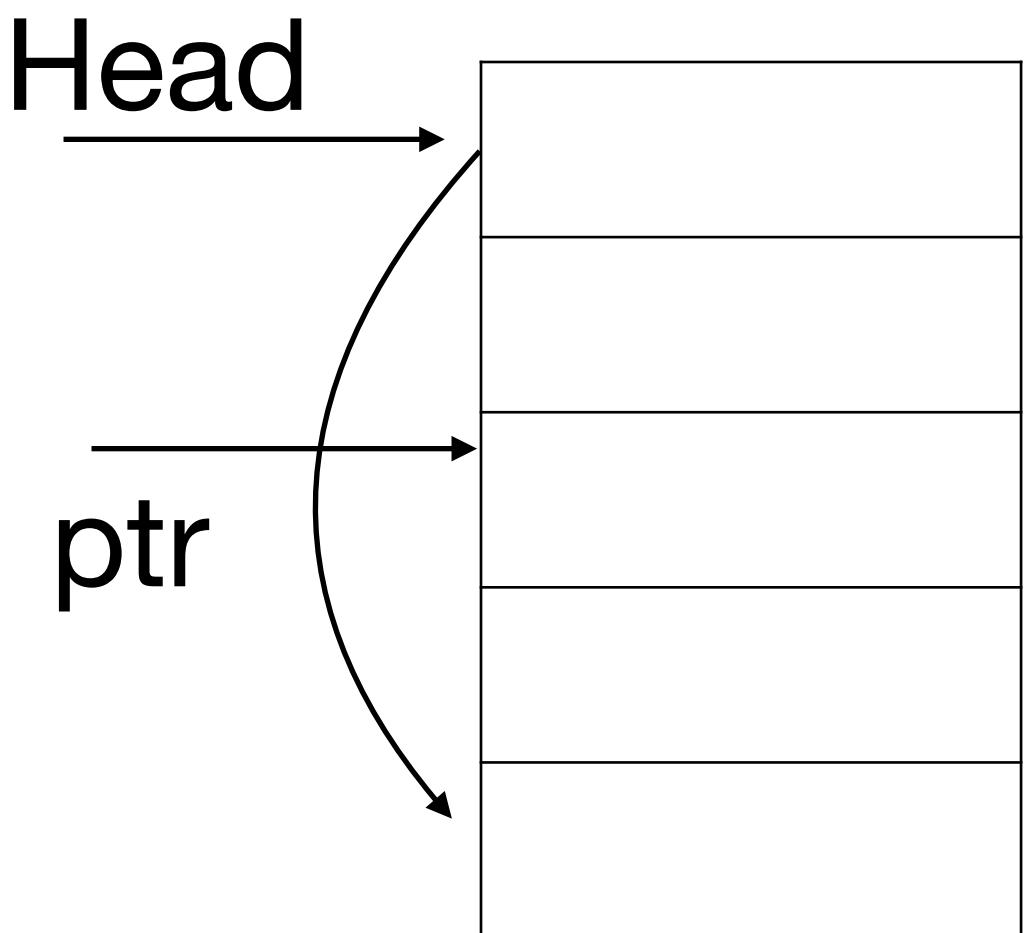
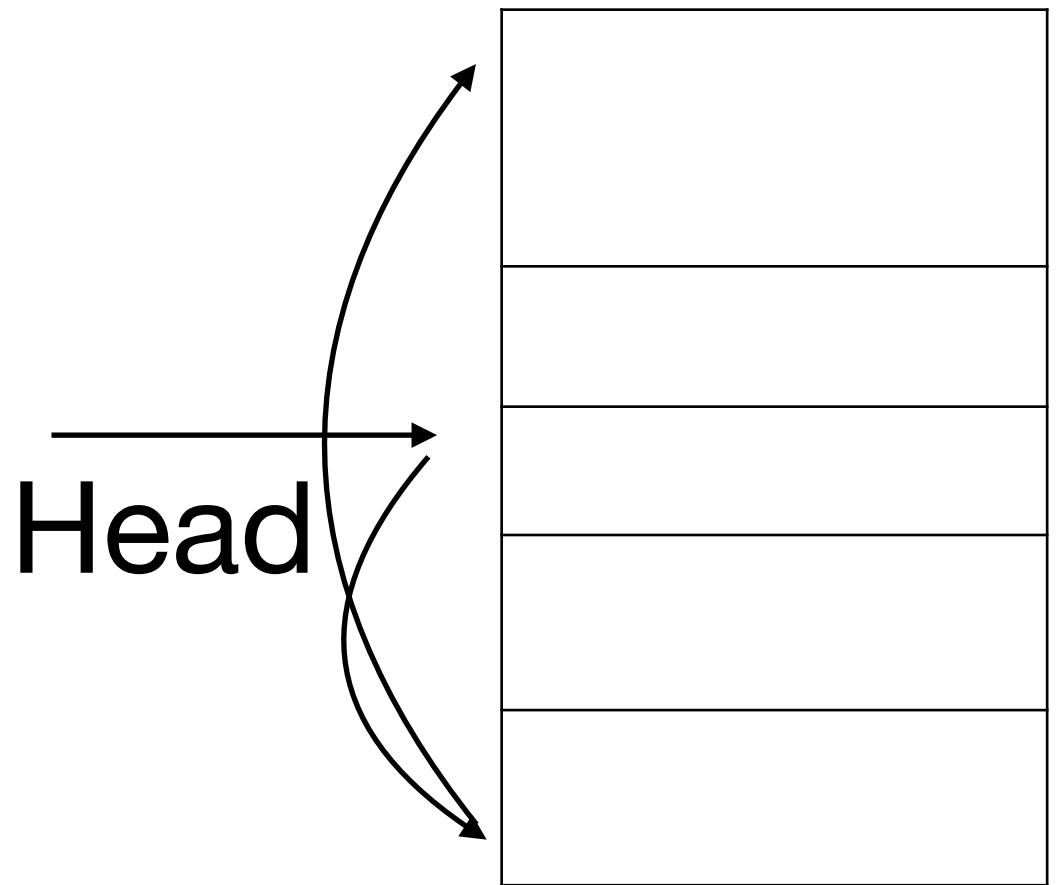
# In which order to maintain lists?

- (De)allocation order



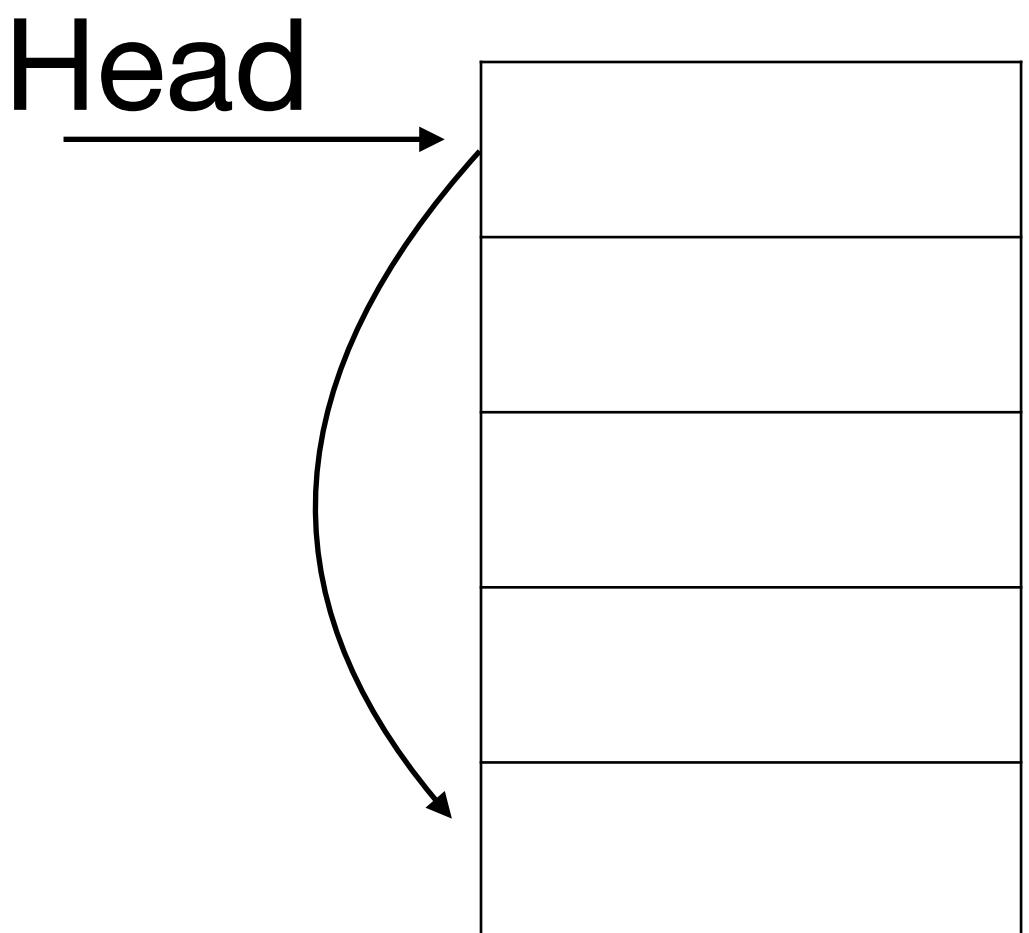
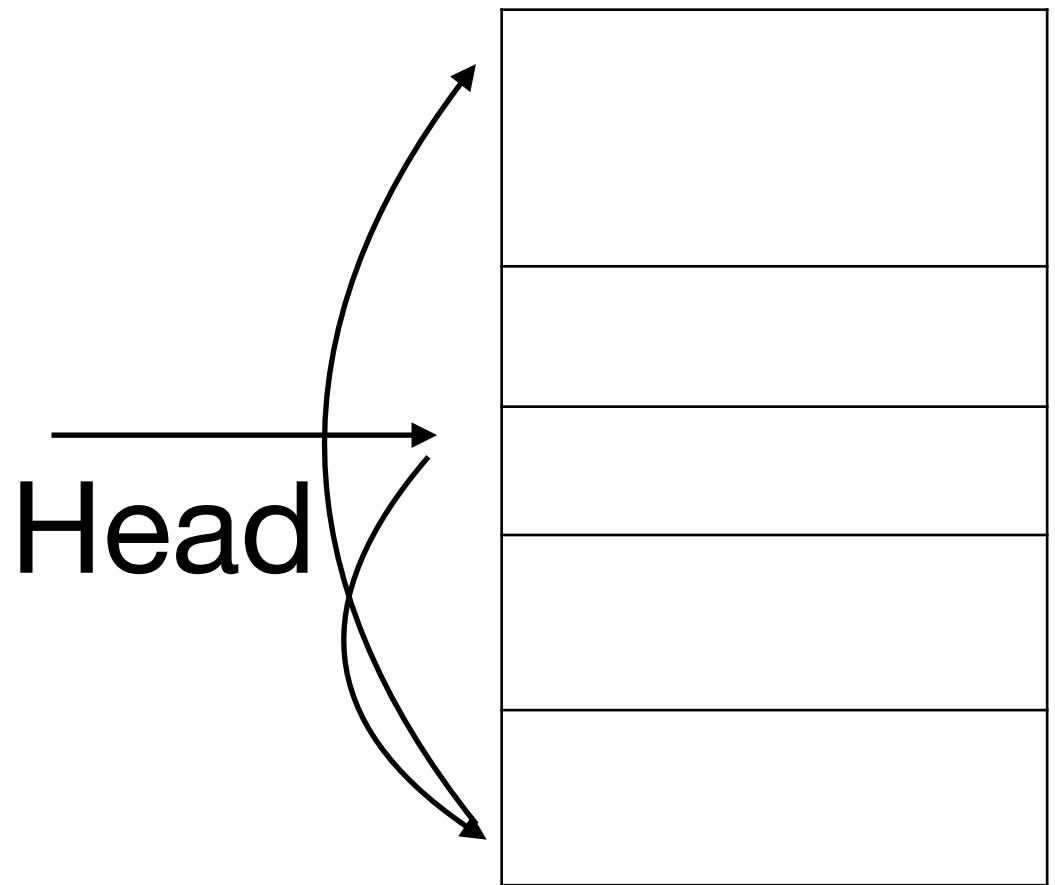
# In which order to maintain lists?

- (De)allocation order
- Address order



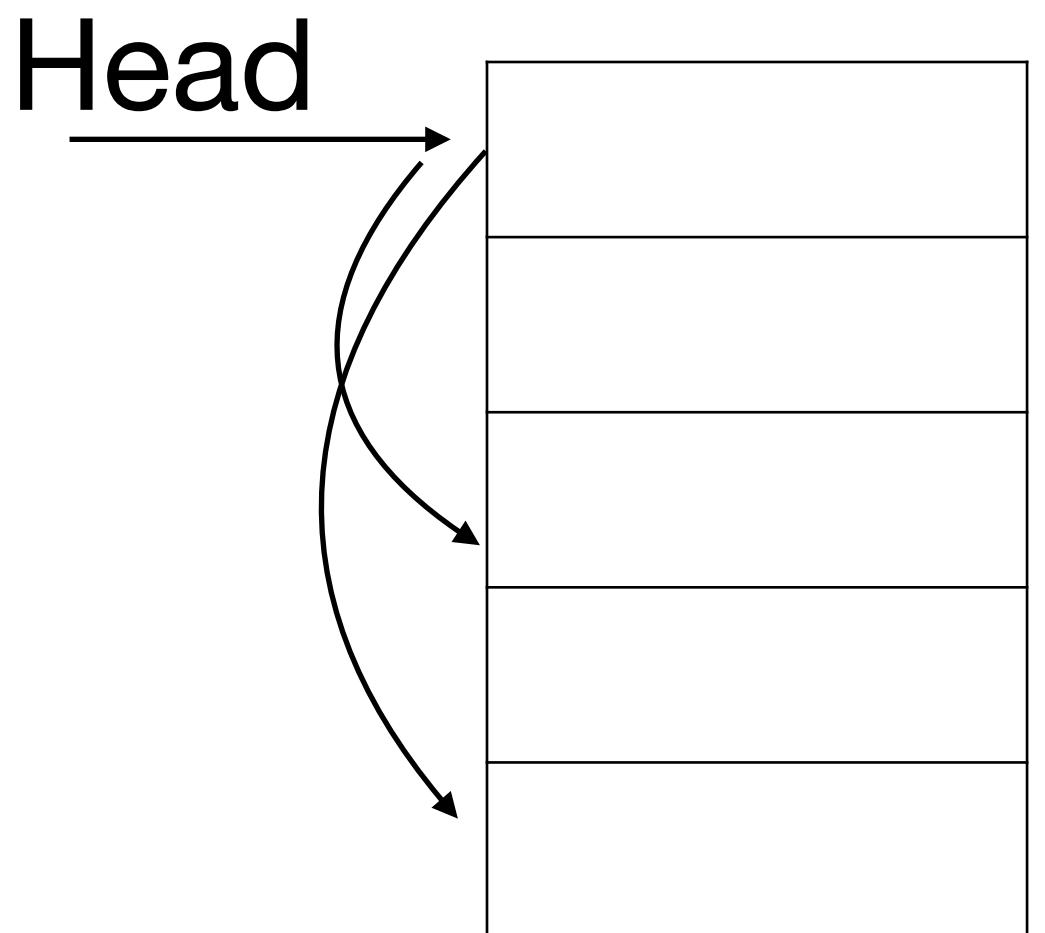
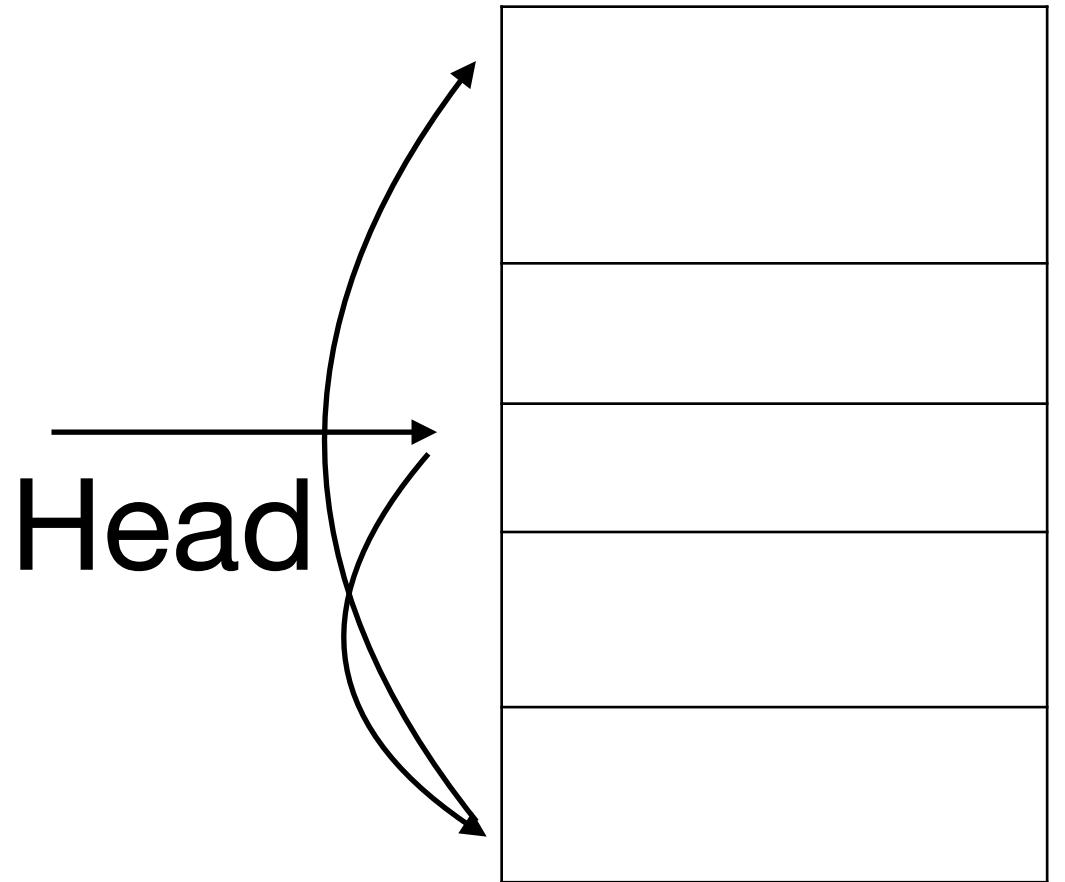
# In which order to maintain lists?

- (De)allocation order
- Address order



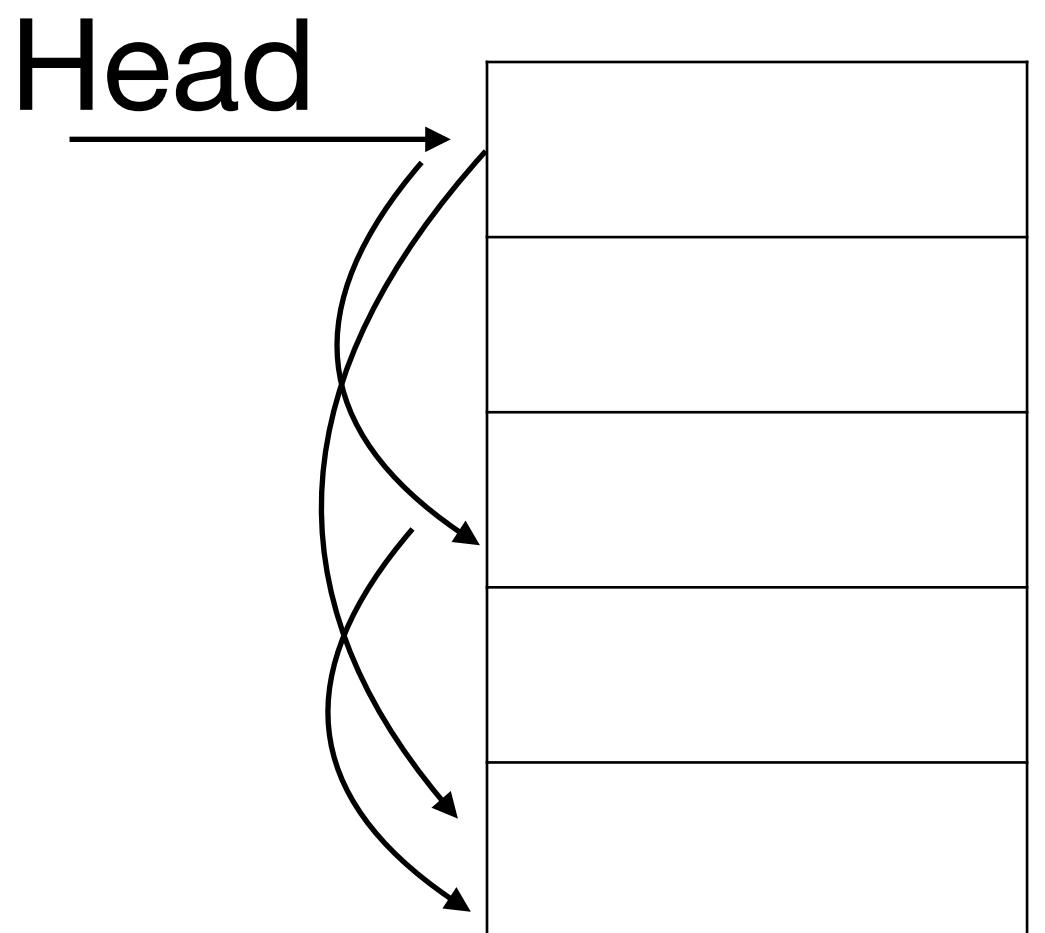
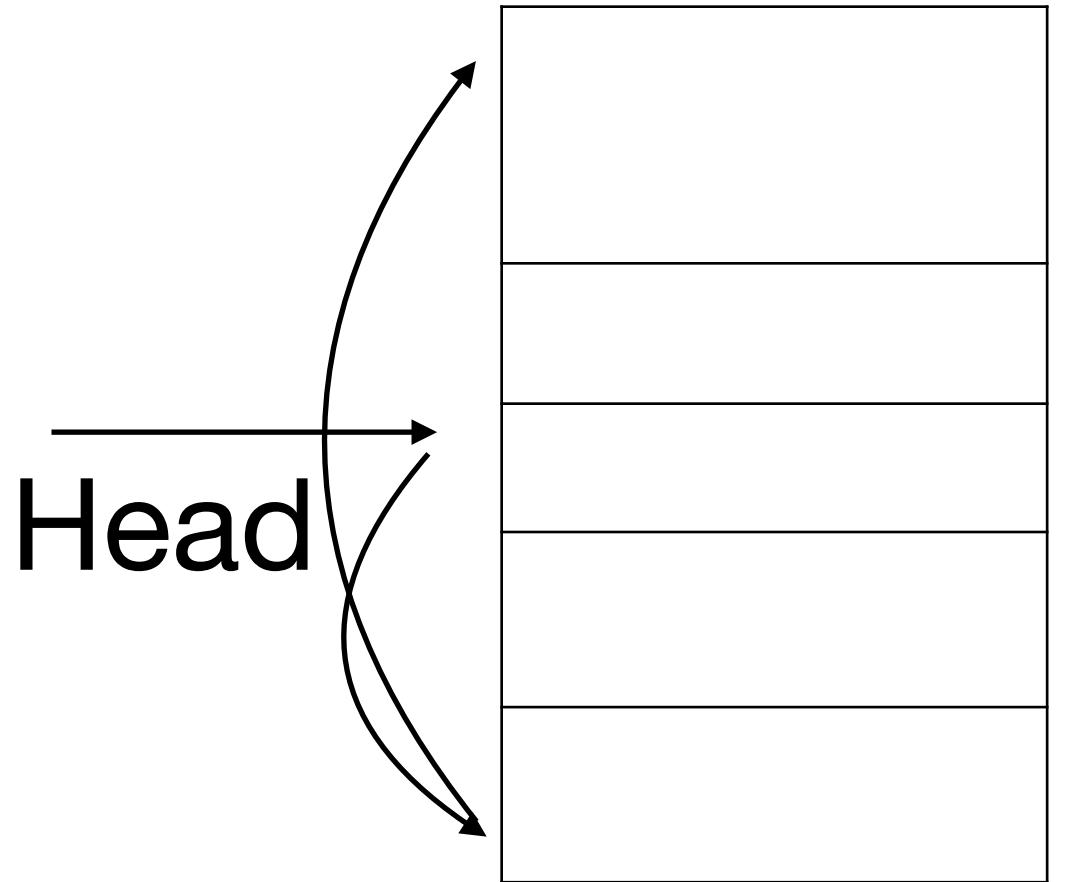
# In which order to maintain lists?

- (De)allocation order
- Address order



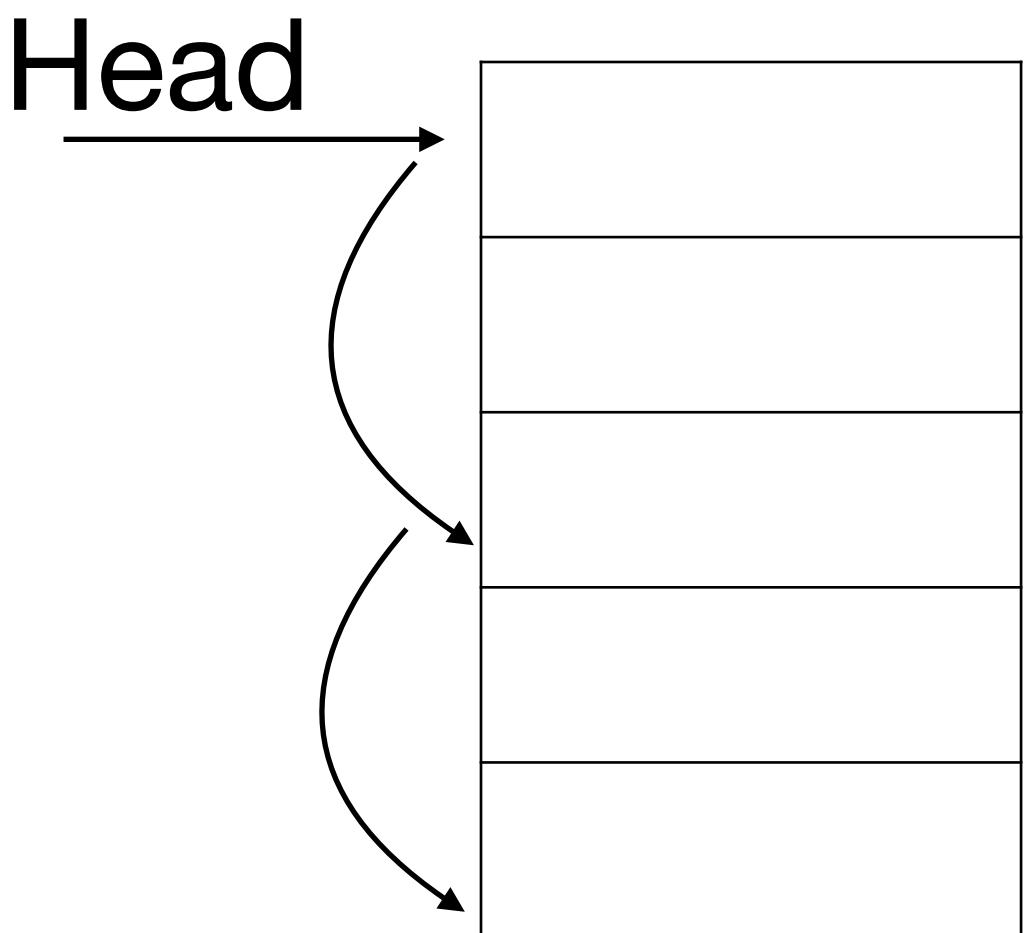
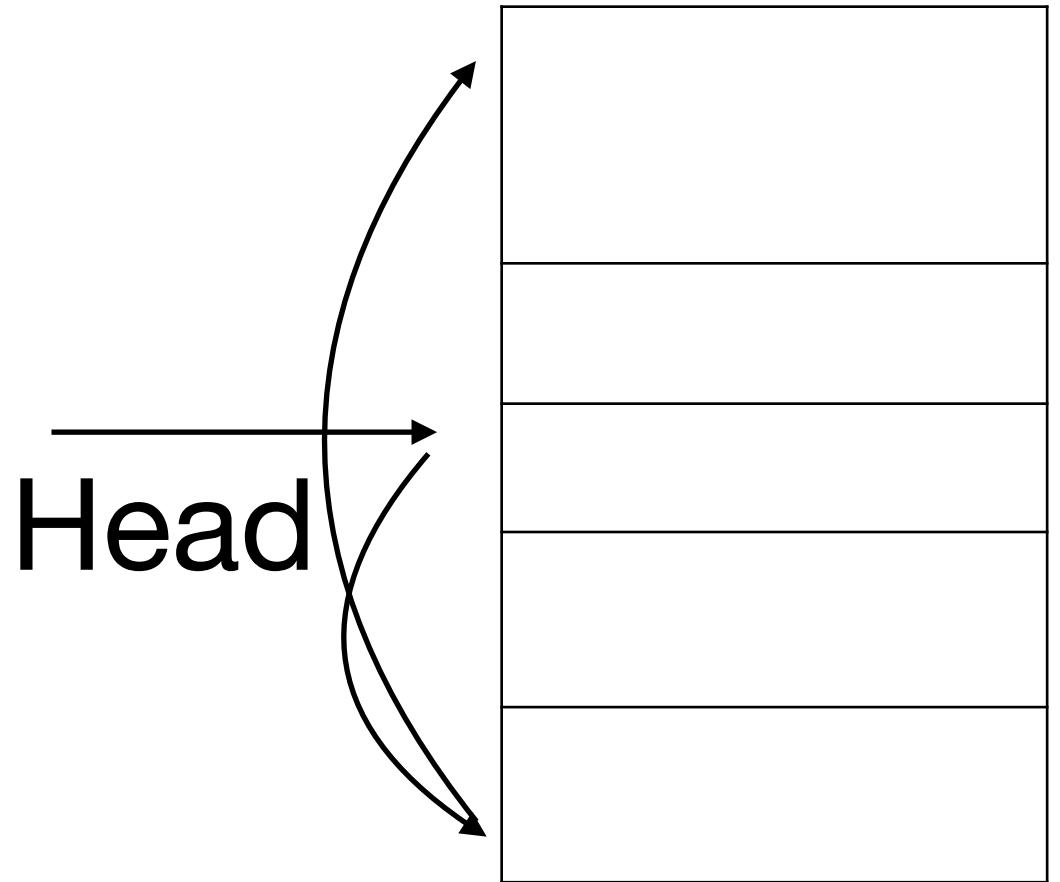
# In which order to maintain lists?

- (De)allocation order
- Address order



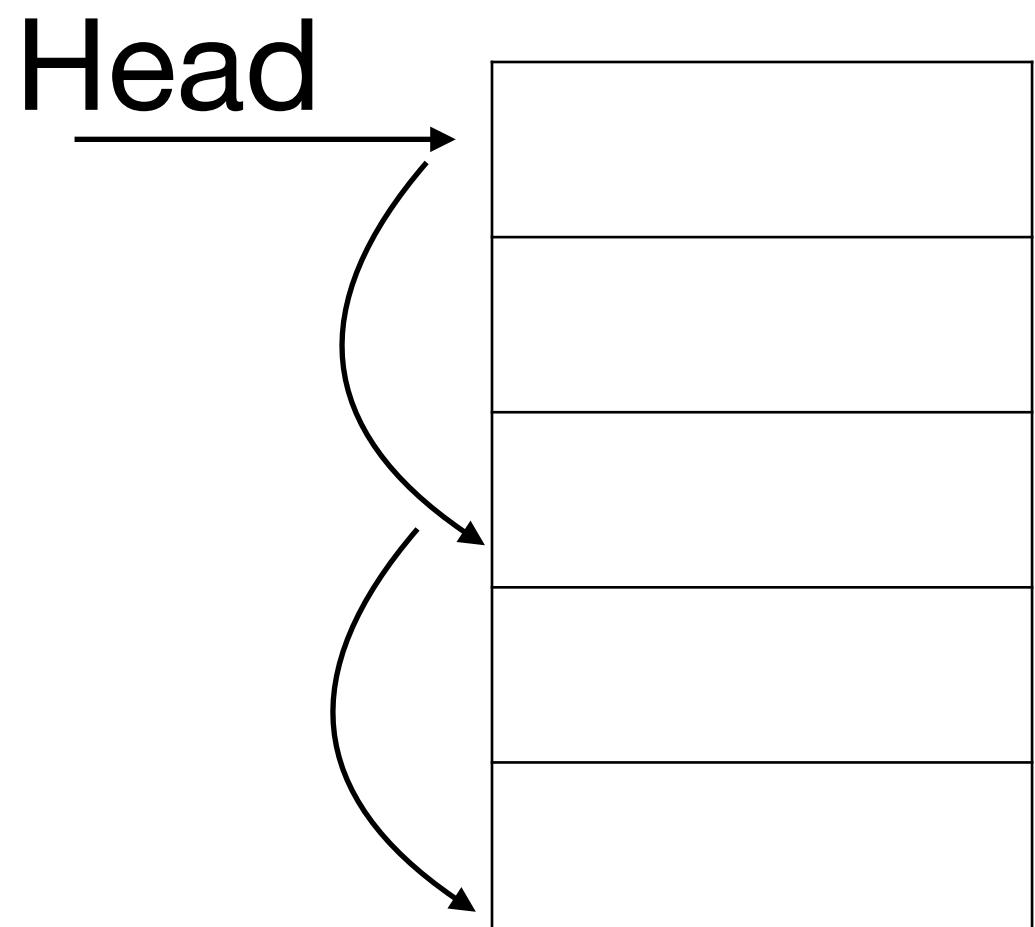
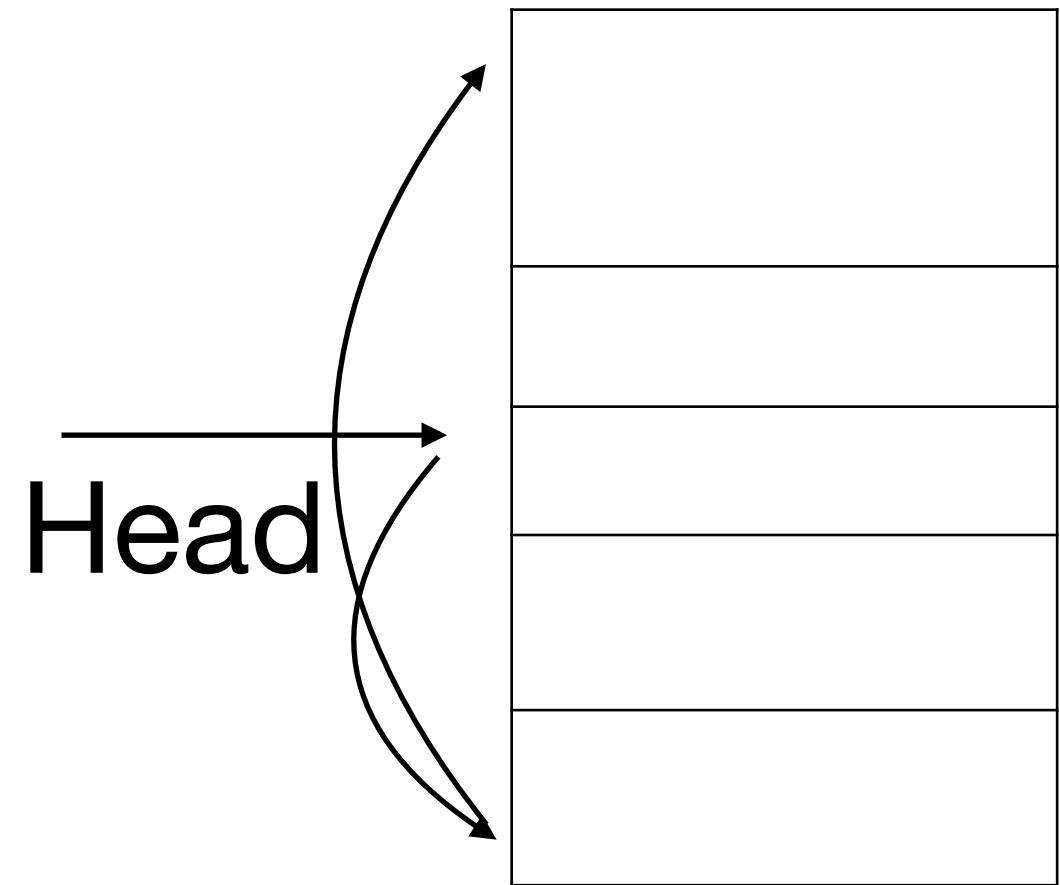
# In which order to maintain lists?

- (De)allocation order
- Address order



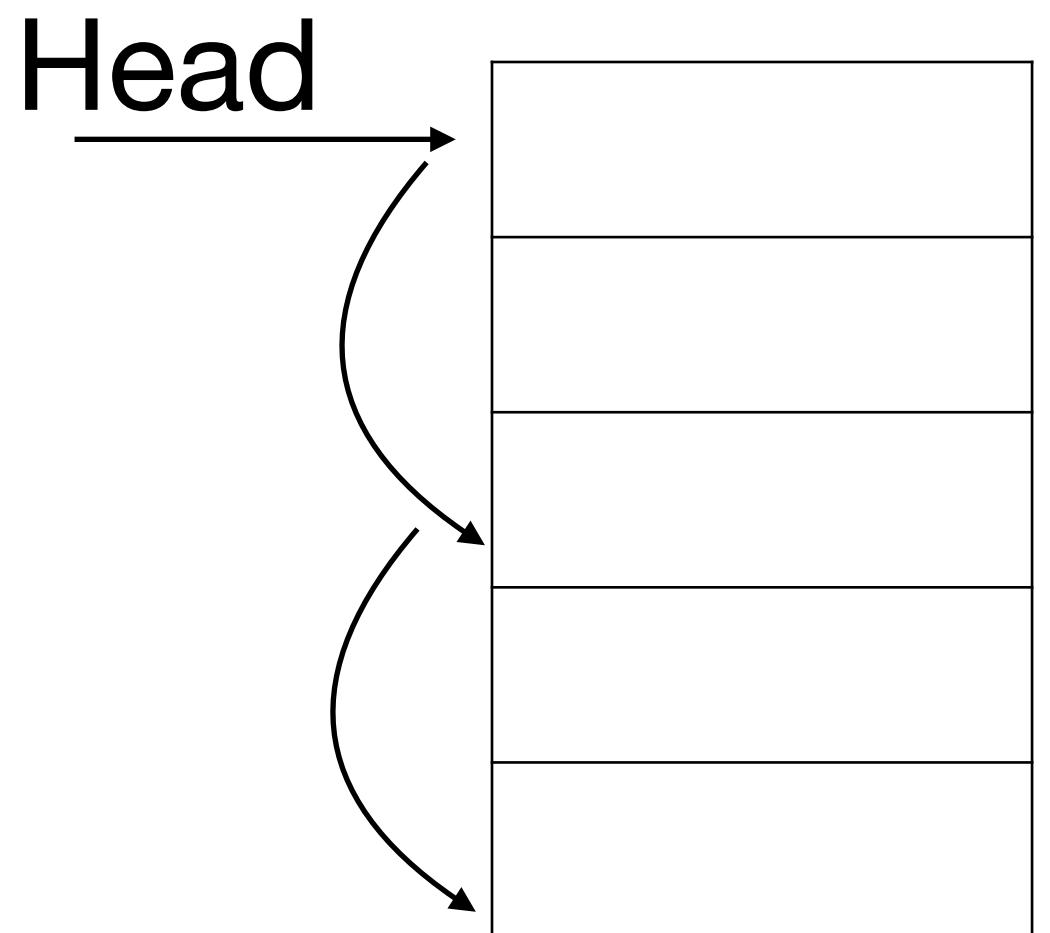
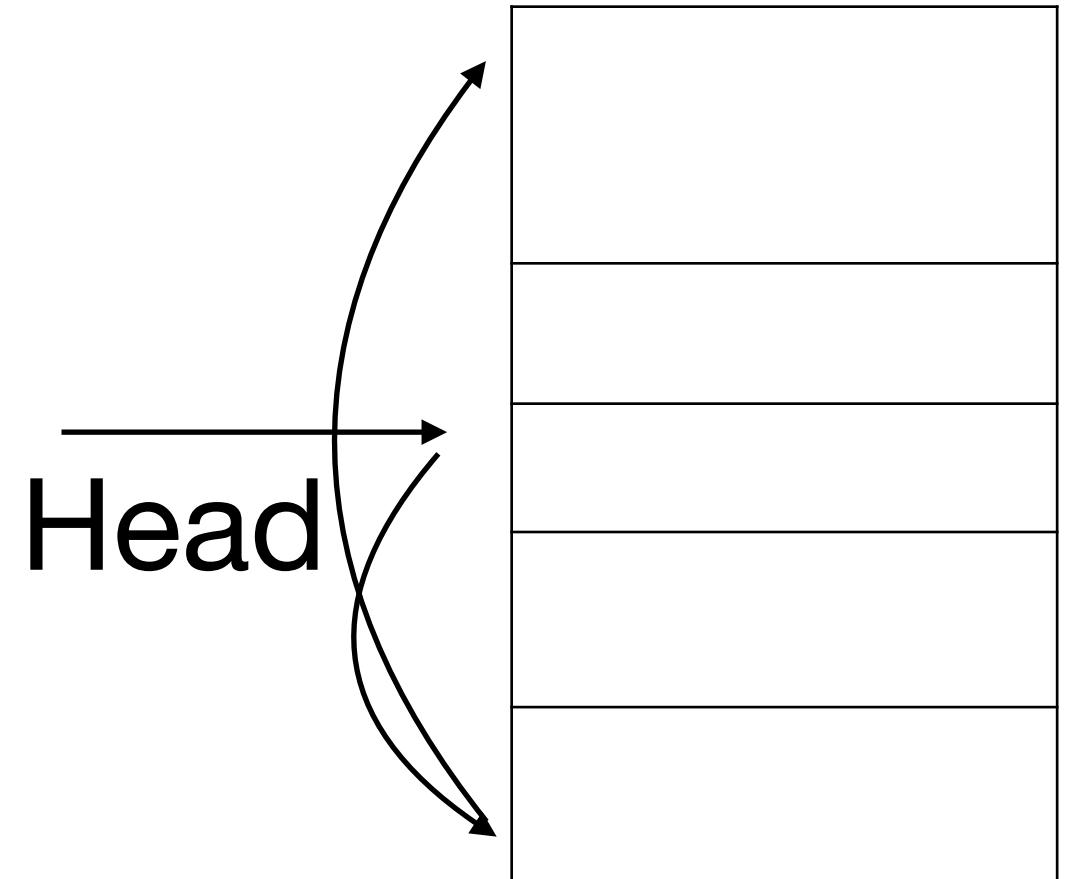
# In which order to maintain lists?

- (De)allocation order
- Address order
  - Slow frees: need to traverse the free list



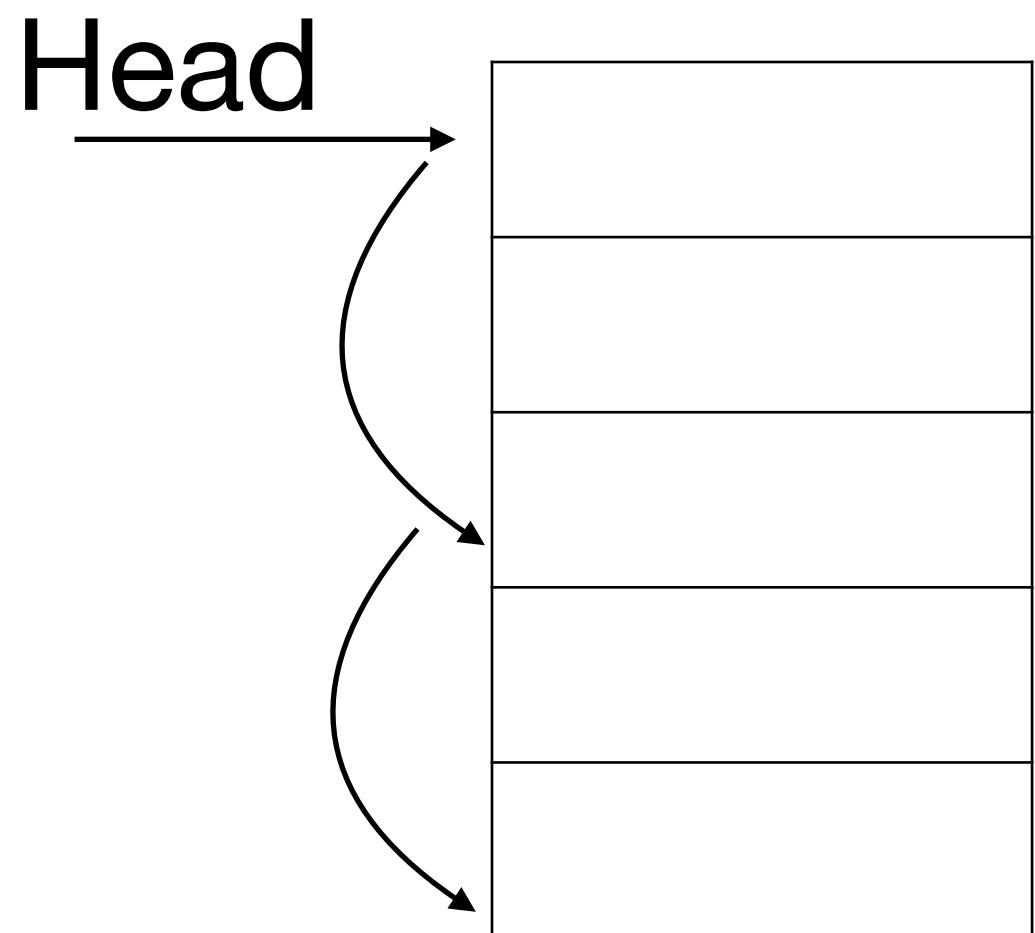
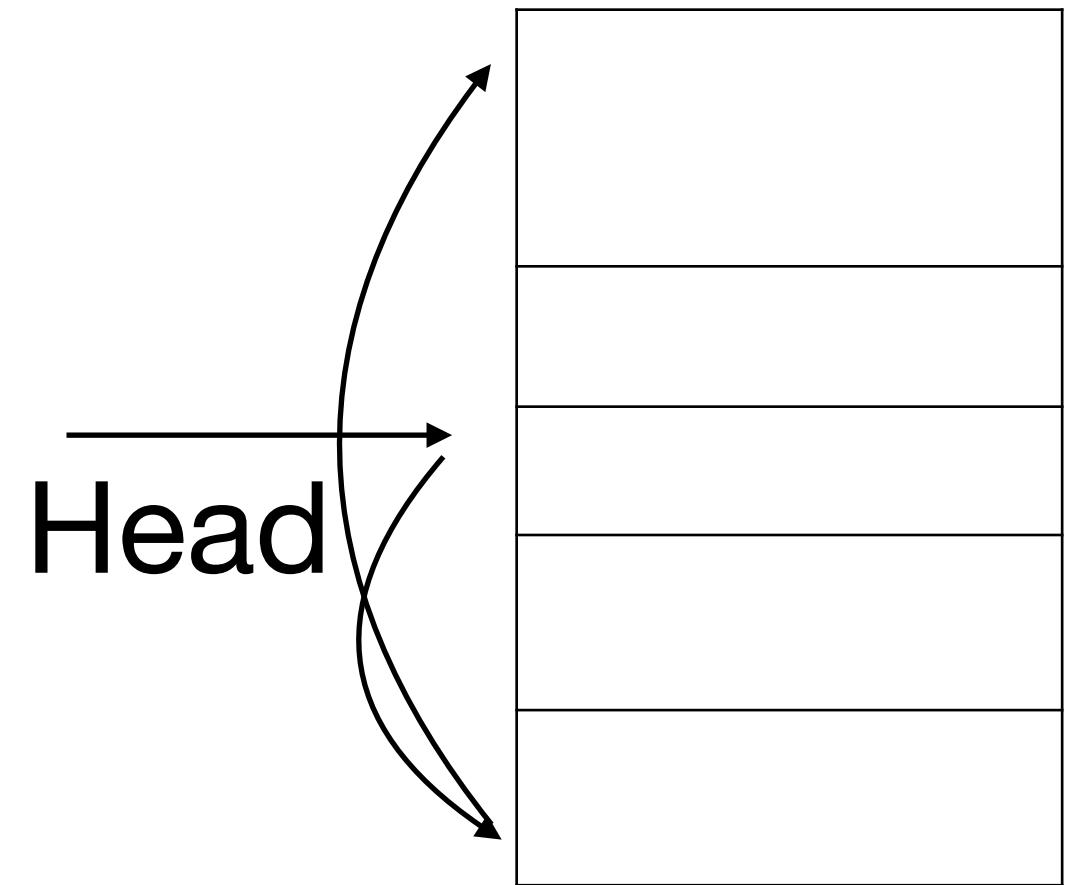
# In which order to maintain lists?

- (De)allocation order
- Address order
  - Slow frees: need to traverse the free list
  - (xv6: umalloc.c)



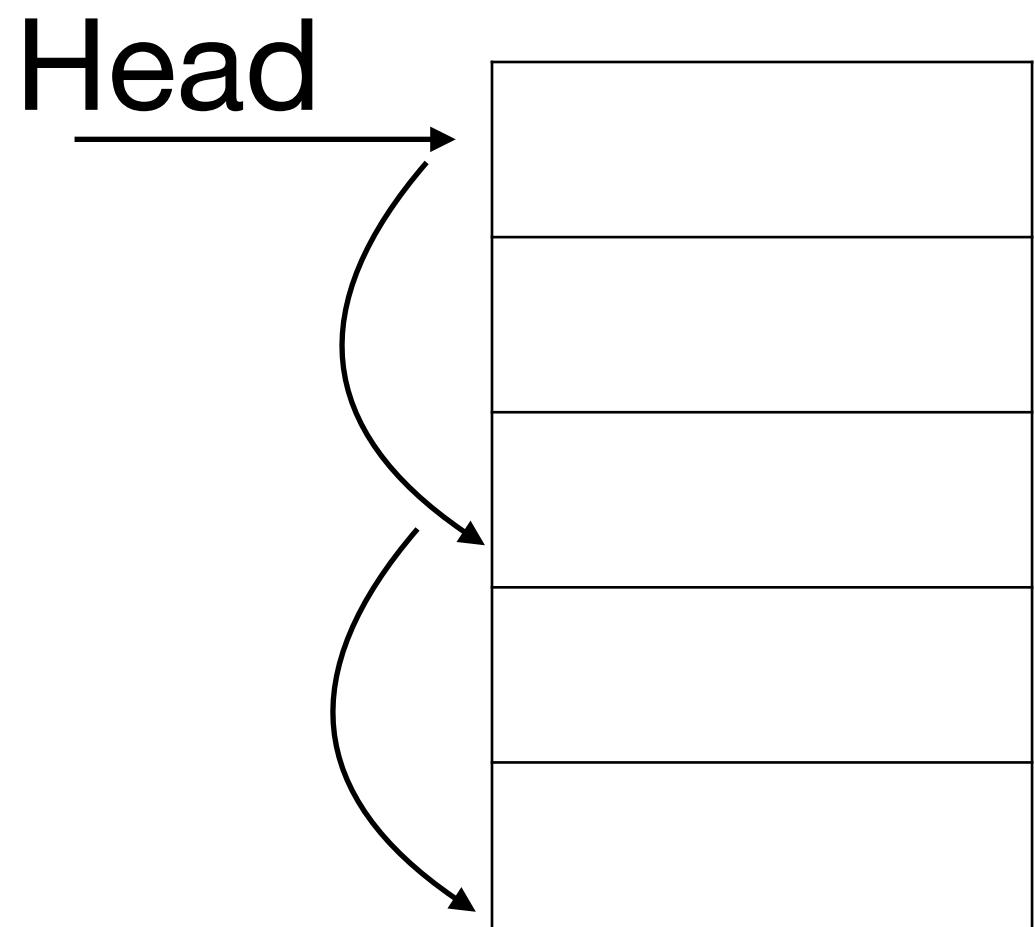
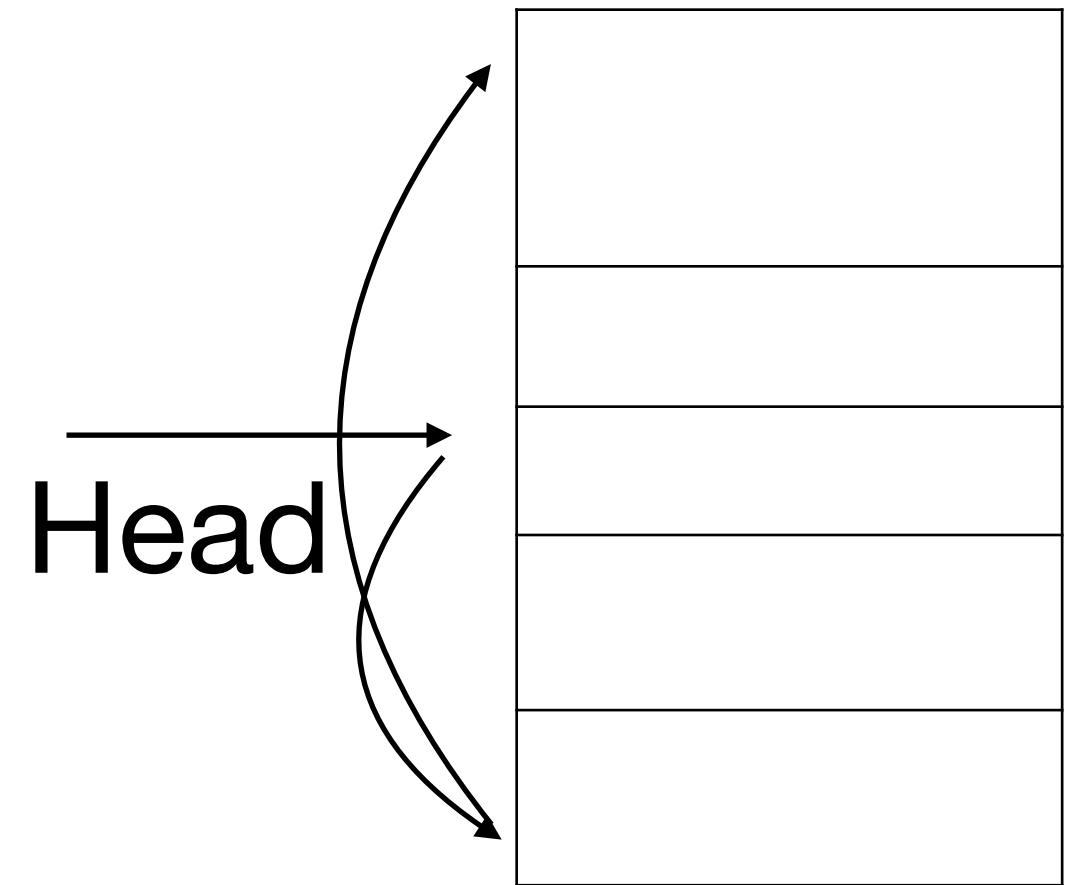
# In which order to maintain lists?

- (De)allocation order
- Address order
  - Slow frees: need to traverse the free list
  - (xv6: umalloc.c)
  - Address order, first fit will allocate back-to-back allocations contiguously.



# In which order to maintain lists?

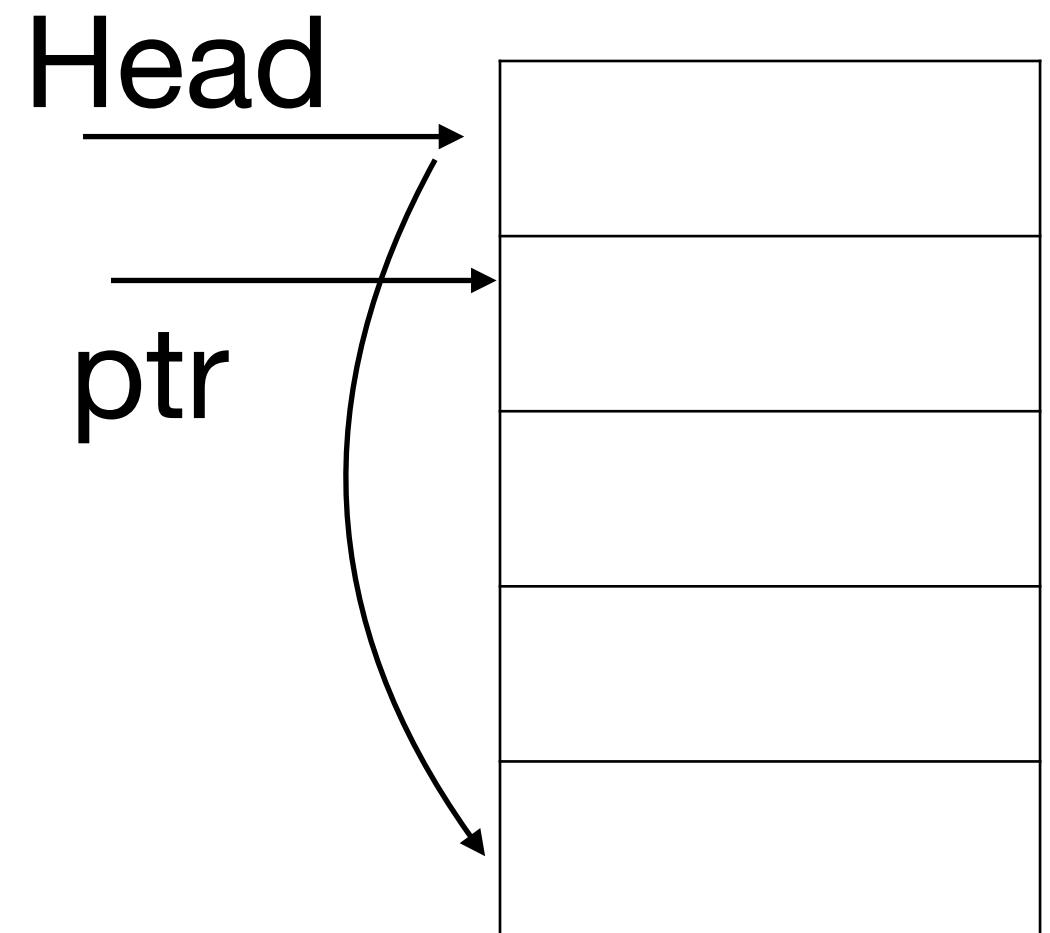
- (De)allocation order
- Address order
  - Slow frees: need to traverse the free list
  - (xv6: umalloc.c)
  - Address order, first fit will allocate back-to-back allocations contiguously.
    - Due to “clustered deaths”, we may get better chances of coalescing



# How to do coalescing?

**Example: free(ptr)**

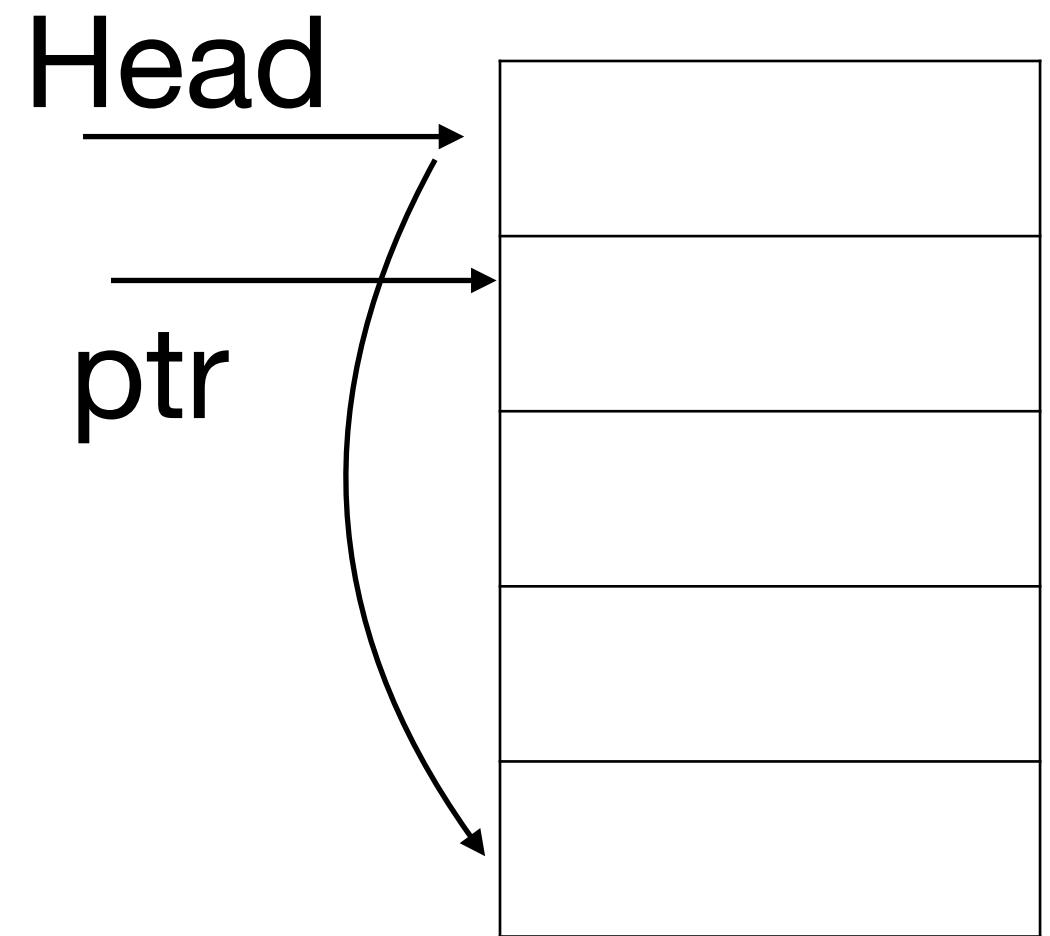
- Straightforward in address order since we are traversing the free list in address order



# How to do coalescing?

**Example: free(ptr)**

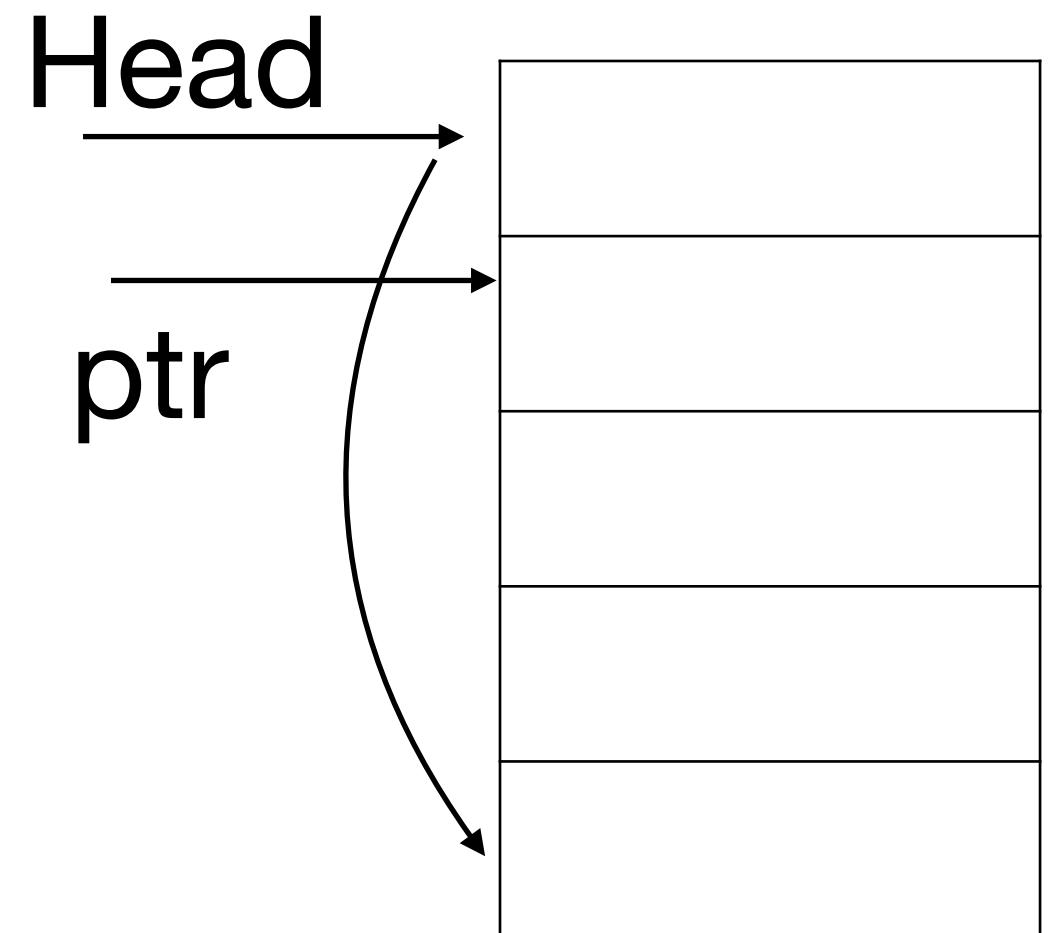
- Straightforward in address order since we are traversing the free list in address order
- In deallocation order: when an area is freed, check if the “boundary tag” is present in the footer above



# How to do coalescing?

**Example: free(ptr)**

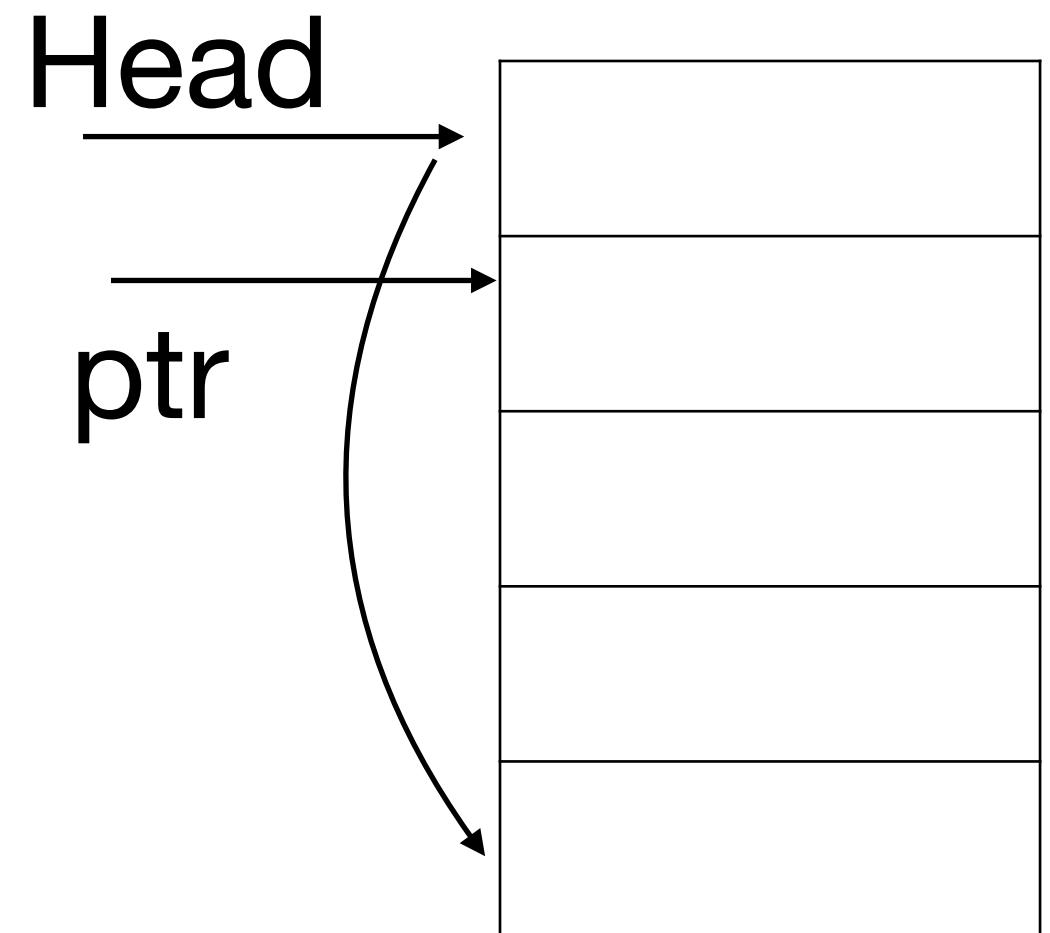
- Straightforward in address order since we are traversing the free list in address order
- In deallocation order: when an area is freed, check if the “boundary tag” is present in the footer above



# How to do coalescing?

**Example: free(ptr)**

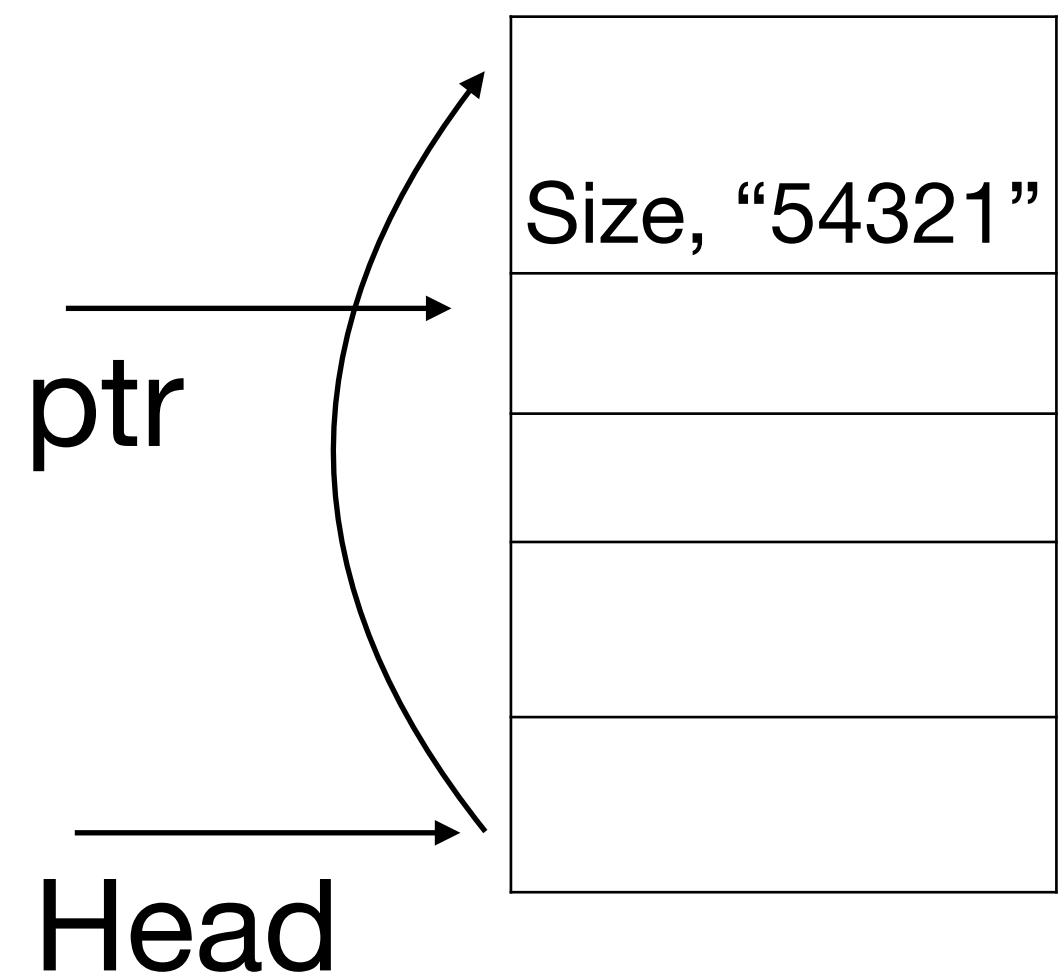
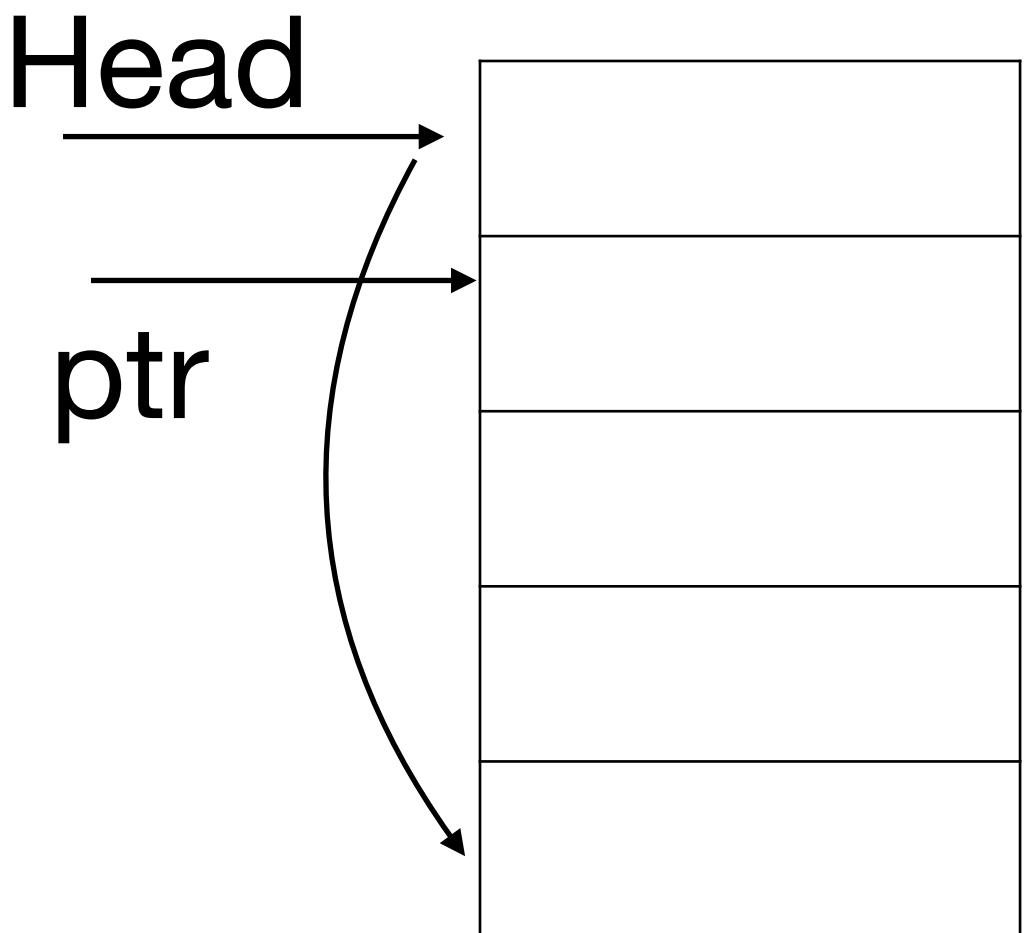
- Straightforward in address order since we are traversing the free list in address order
- In deallocation order: when an area is freed, check if the “boundary tag” is present in the footer above
- First fit and address order work well in clustered deaths, simplify coalescing (no boundary tag), but cause fragmentation and traverse free list at free calls



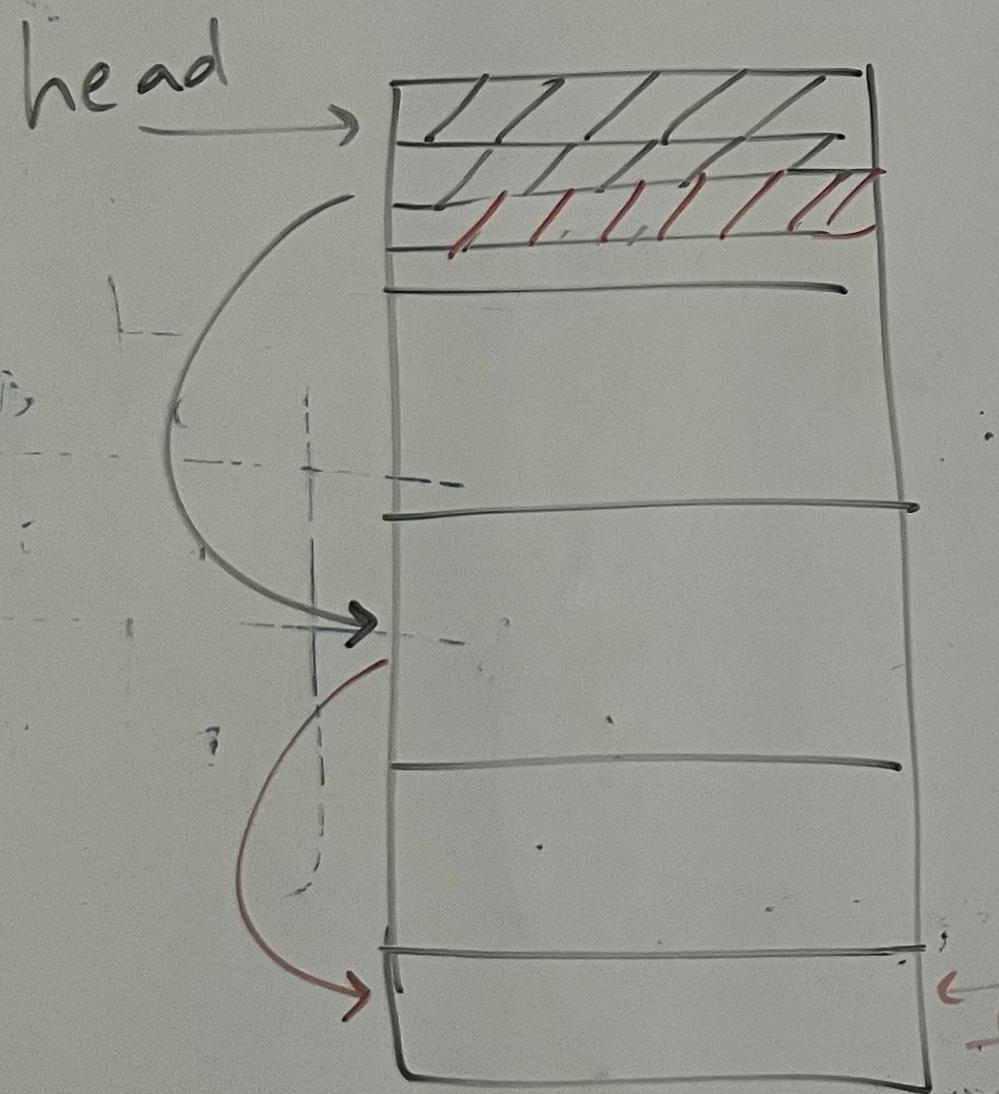
# How to do coalescing?

Example: `free(ptr)`

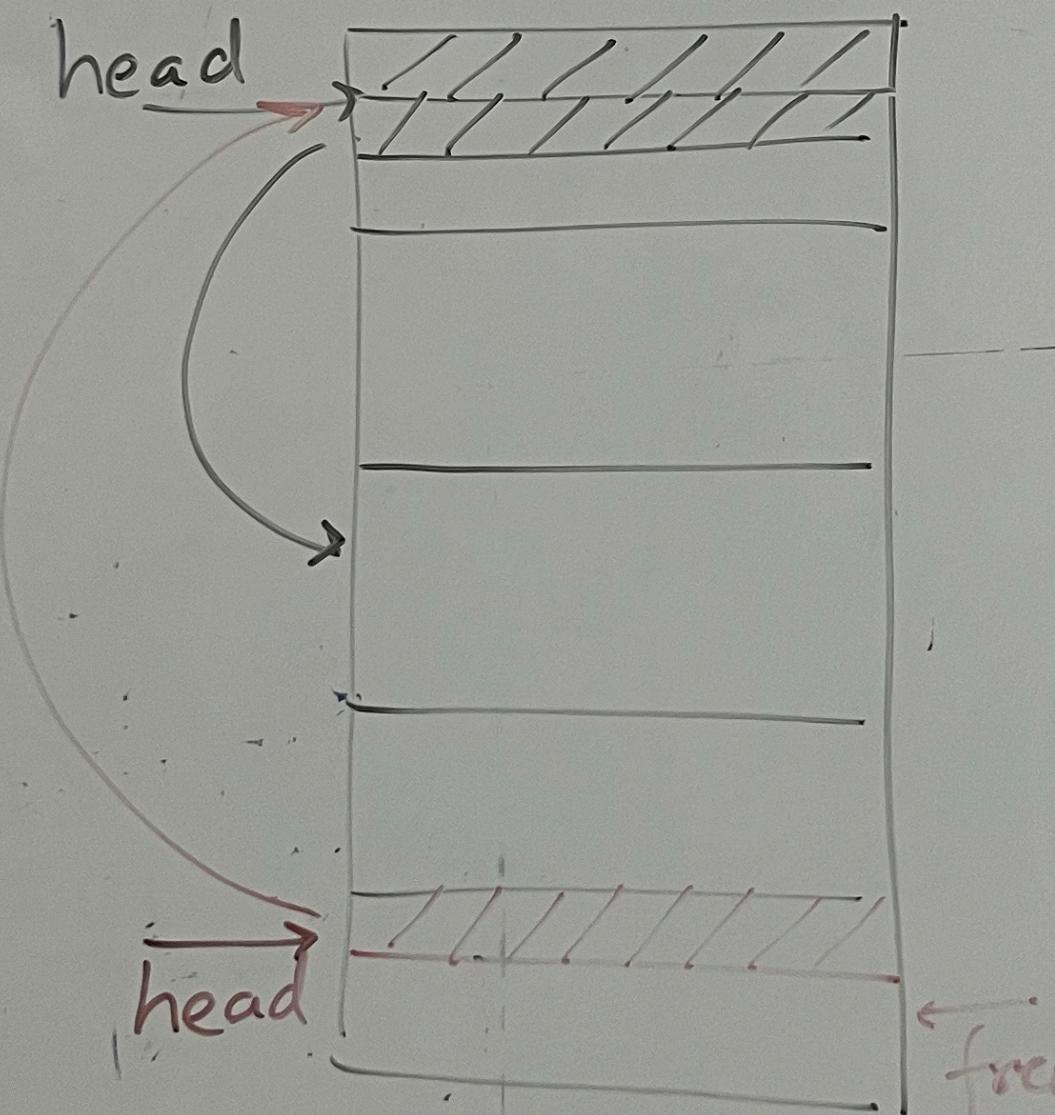
- Straightforward in address order since we are traversing the free list in address order
- In deallocation order: when an area is freed, check if the “boundary tag” is present in the footer above
- First fit and address order work well in clustered deaths, simplify coalescing (no boundary tag), but cause fragmentation and traverse free list at free calls



First fit  
address order



First fit  
(de) allocation order



Boundary tag

